

# CSCI 466: Networks

Lecture 5: Peer 2 Peer Networks (P2P), Content Distribution Networks (CDN)


Reese Pearsall  
Fall 2022

# Announcements

there are no announcements

# P2P Networks


- No always-on server
- Arbitrary end systems directly communicate
- Peers are intermittently connected and change IP addresses



Time to distribute a file of size  $F$  to  $N$  clients

$$D_{P2P} = \max \left\{ \frac{F}{u_s}, \frac{F}{d_{\min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

## Client-Server Architecture




- Server can upload data at rate  $u_s$
- Clients download data at rates  $d_1, d_2, \dots, d_N$

$$D_{CS} = \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{\min}} \right\}$$


Time to distribute a file of size  $F$  to  $N$  clients

# P2P Networks

## Client-Server Architecture




Existing clients have to share resources with new users



New peers are both a client and a server. There will not be a negative impact on current peers

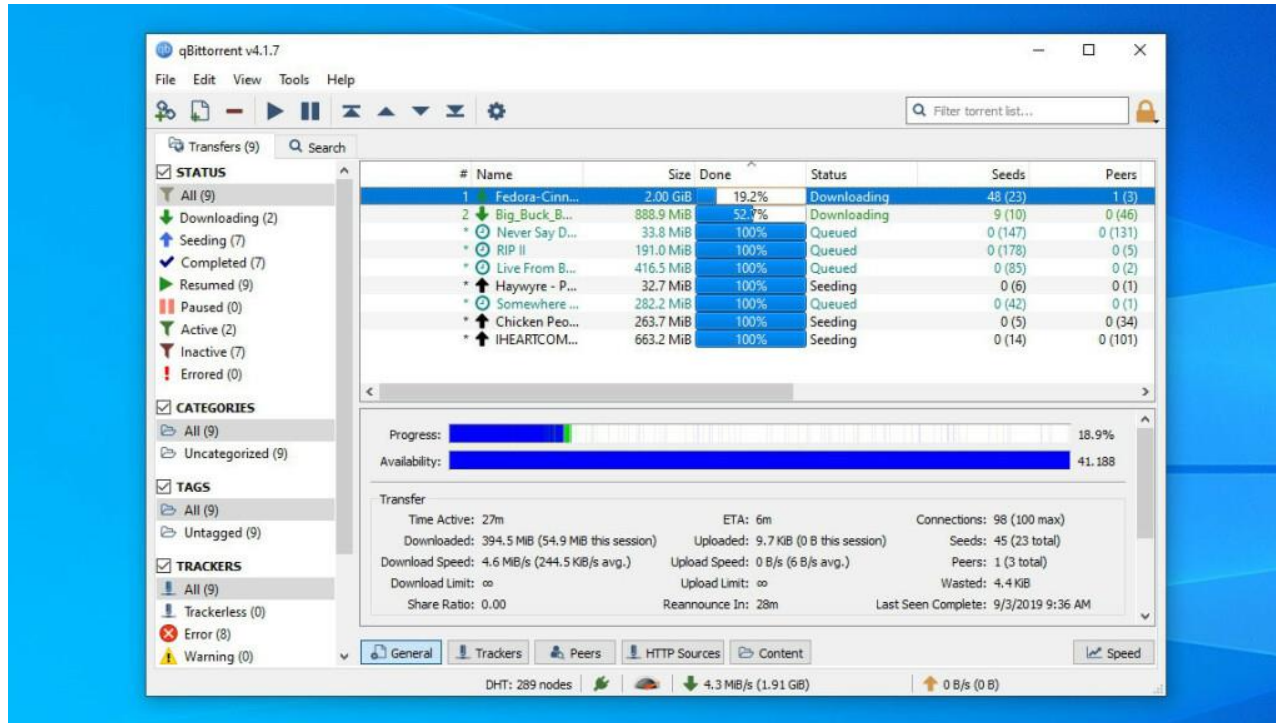
P2P architectures are *self-scaling*

# P2P Networks



- **Client server** distribution time grows with the number of clients
- **P2P** distribution time approaches 1 hour as number of clients grows

# P2P file distribution: BitTorrent




Service for sharing files  
over the internet in a  
decentralized fashion

# P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

*tracker*: tracks peers participating in torrent

*torrent*: group of peers exchanging chunks of a file




# P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

**tracker:** tracks peers participating in torrent

**torrent:** group of peers exchanging chunks of a file



## peer joining torrent:


- has no chunks, but will accumulate them over time from other peers
- registers with tracker to get list of peers, connects to subset of peers (“neighbors”)



# P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

**tracker:** tracks peers participating in torrent





Alice arrives ...  
... obtains list  
of peers from tracker  
... and begins exchanging  
file chunks with peers in torrent

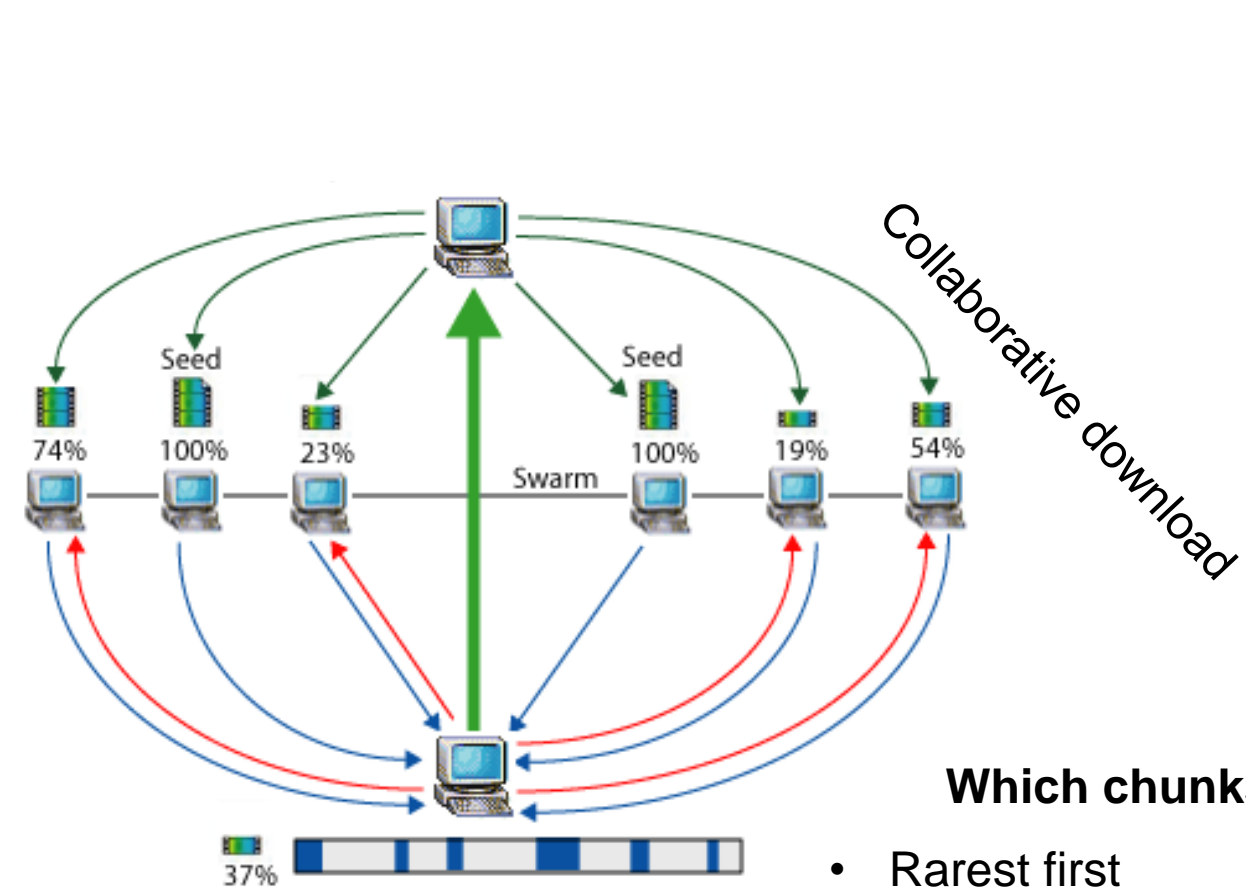
**torrent:** group of peers  
exchanging chunks of a file

- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn:** peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

# P2P file distribution: BitTorrent




# P2P file distribution: BitTorrent

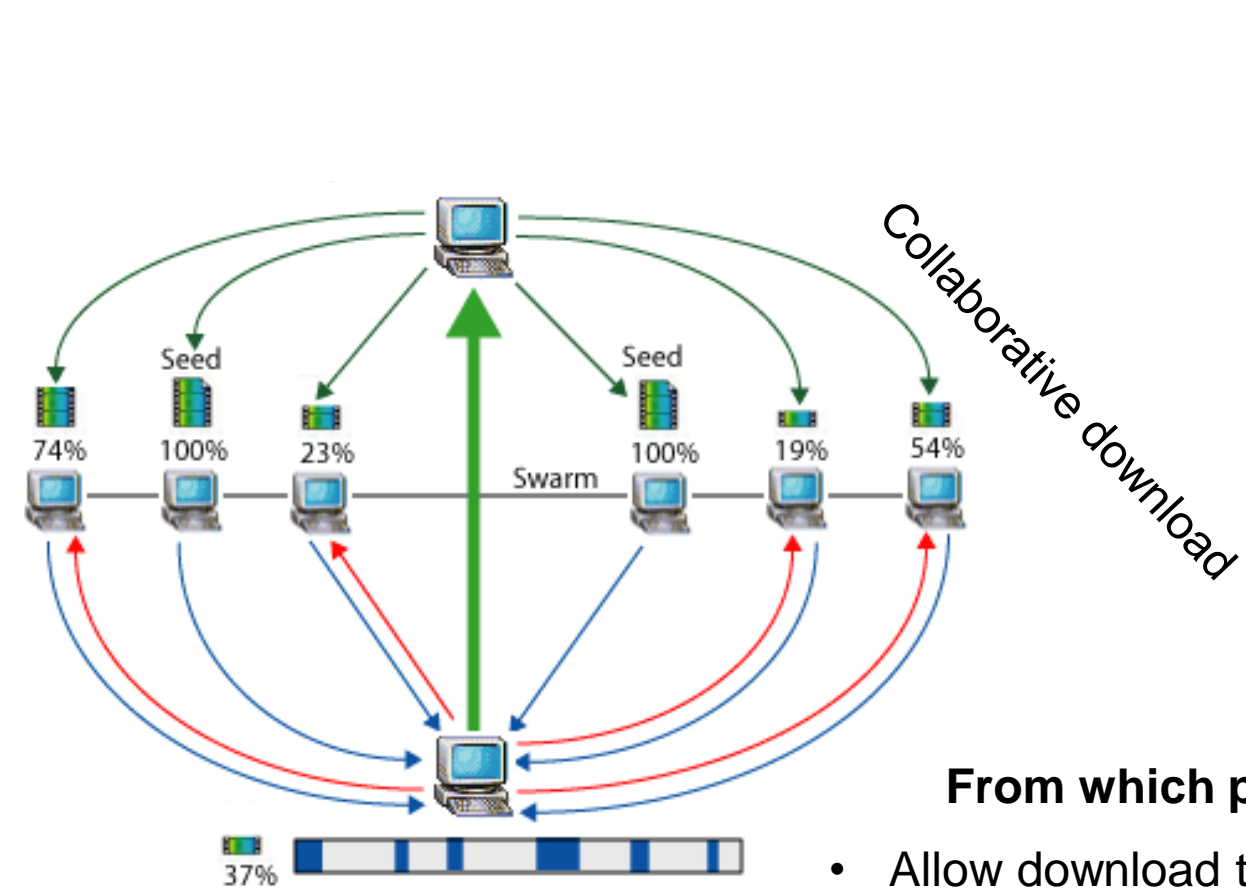


## Which chunks to download?

- Rarest first
- Try to equalize the number of copies of each chunk in the system




# P2P file distribution: BitTorrent



## From which peers to download?


- Allow download to N peers with highest upload rates
- Allows peers with similar upload rates to find each other




# P2P file distribution: BitTorrent

## *requesting chunks:*

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first



# P2P file distribution: BitTorrent



higher upload rate: find better trading partners, get file faster !

## *sending chunks: tit-for-tat*

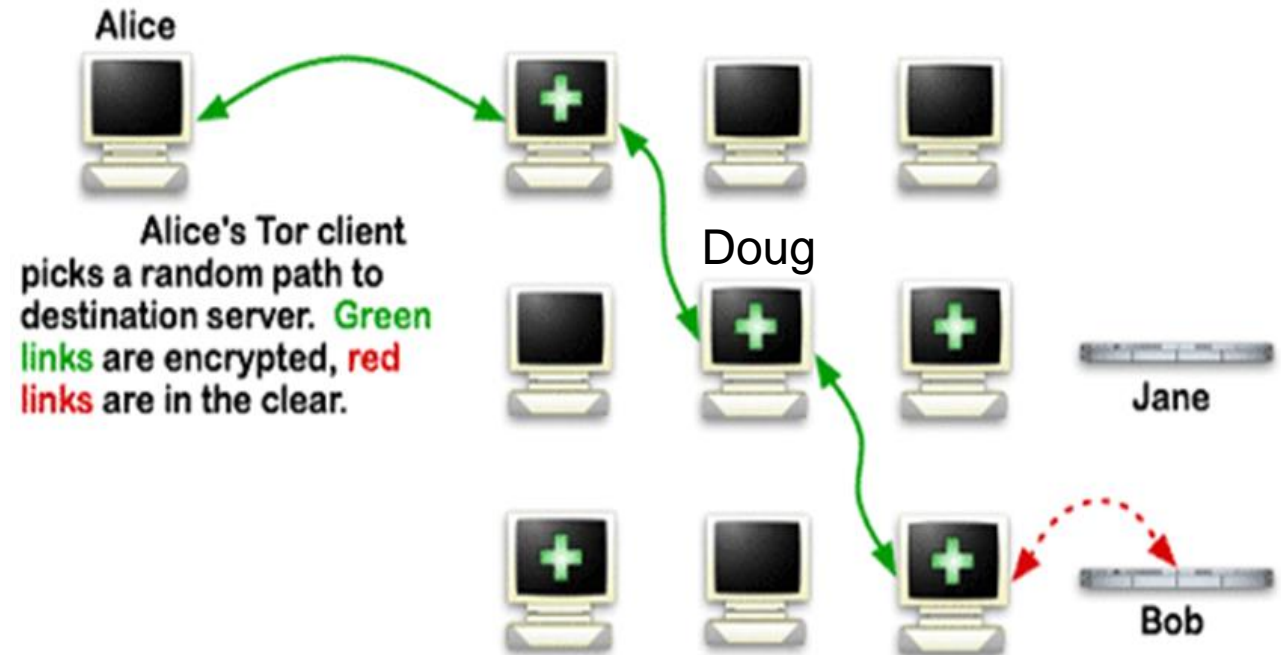

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
  - other peers are *choked* by Alice (do not receive chunks from her)
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

BitTorrent is referred to as a **unstructured P2P**

# P2P file distribution mechanisms

To help with anonymity and privacy, some P2P may implement an onion router/TOR router

- Functionality
  - Sender obtains a set of router keys
  - Each router only knows next hop
  - Intermediate routers cannot read message



# Structured P2P Example

Content is placed into the P2P network based on a Globally Unique Identifier (GUID)

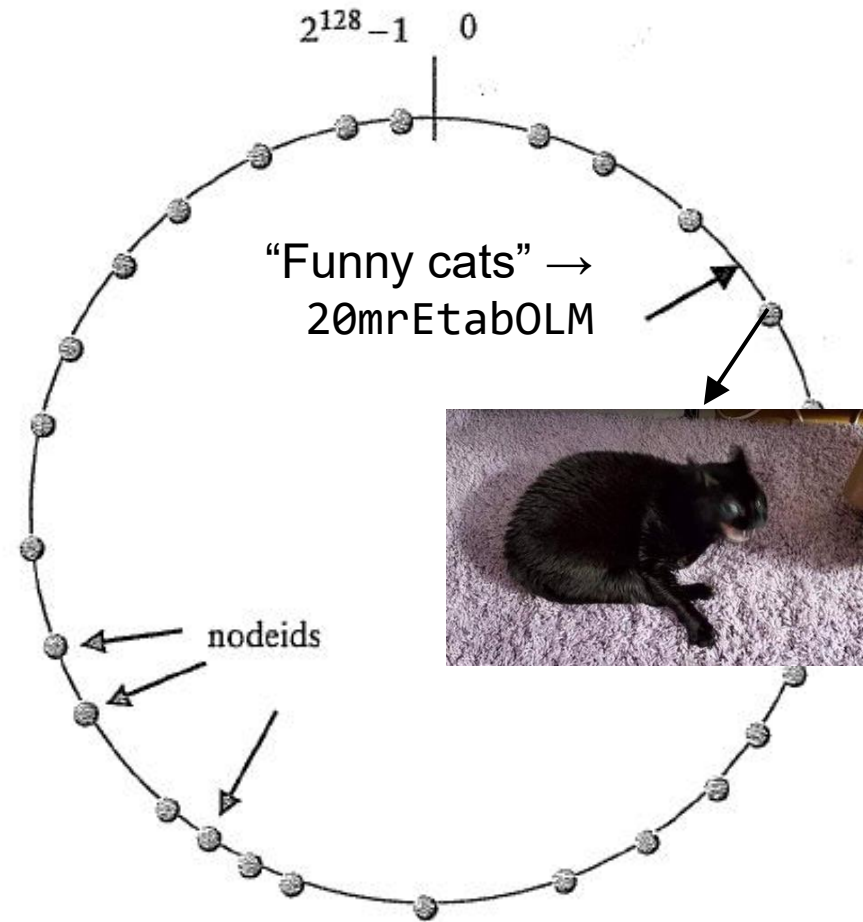
- $\text{GUID} = H(\text{key})$  where  $H$  is some hash function

$P( H(a) = H(b) )$  is very low

- Distributed <key, value> search
- Fully Distributed and self-organizing

## Distributed Hash Table (DHT)

- Join: join ring with GUID from public key
- Publish:  $\text{put}(\text{GUID}, \text{value})$
- Search/Fetch:  $\text{value} = \text{get}(\text{GUID})$



Each node maintains a link to the 4 **closest** neighbors and to nodes **half** way across the ring, **quarter** way across the ring, **eight** across the way...  $O(\log N)$



# Structured P2P Example

	Unstructured P2P	Structured P2P
Advantages	<ul style="list-style-type: none"><li>• Self-organizing</li><li>• Naturally resilient to node failure</li></ul>	<ul style="list-style-type: none"><li>• Guaranteed to locate objects (if exist)</li><li>• Time and complexity bounds</li><li>• Low message overhead</li></ul>
Disadvantages	<ul style="list-style-type: none"><li>• Cannot offer guarantees on locating objects</li><li>• Prone to excessive messaging that limits scalability</li></ul>	<ul style="list-style-type: none"><li>• Need to maintain complex overlays</li><li>• High control traffic overhead in dynamic environments</li></ul>


# CDNs

- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
- challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution*: distributed, application-level infrastructure



# CDNs


- Caching
  - Save previously delivered data
  - Subsequent requests served from cache on the browser, or in the access network
- Applications
  - Reduce response time for client request
  - Reduce ISP traffic costs
- Content distribution networks
  - Distributed caches
  - Web objects addressed to CDN server
  - CDN server fetches from content provider on first access



- Conditional GET
  - Cache: specify date of cached copy in HTTP request  
**If-modified-since: <date>**
  - Server: response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**

# CDNs

- Challenge:
  - How to stream content to millions of users?
- Option 1:
  - Single mega-datacenter
  - Pros: Simple
  - Cons:
    - Single point of failure
    - Point of network congestion
    - Long path to distant clients
    - Multiple copies of video sent over outgoing link



# CDNs

- Option 2:
  - Store/serve multiple copies of videos at multiple geographically distributed sites


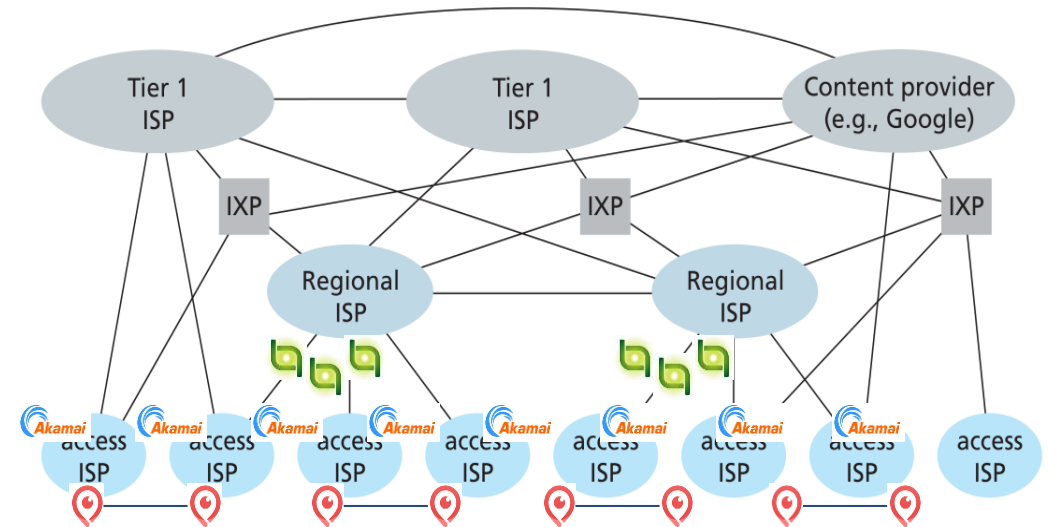
A **content delivery network (CDN)** refers to a geographically distributed group of cache servers which work together to provide fast delivery of Internet content. **(not a web host)**



34 DNS lookups

204 HTTP requests

520 KB of data downloaded



56% of domains resolve to a CDN

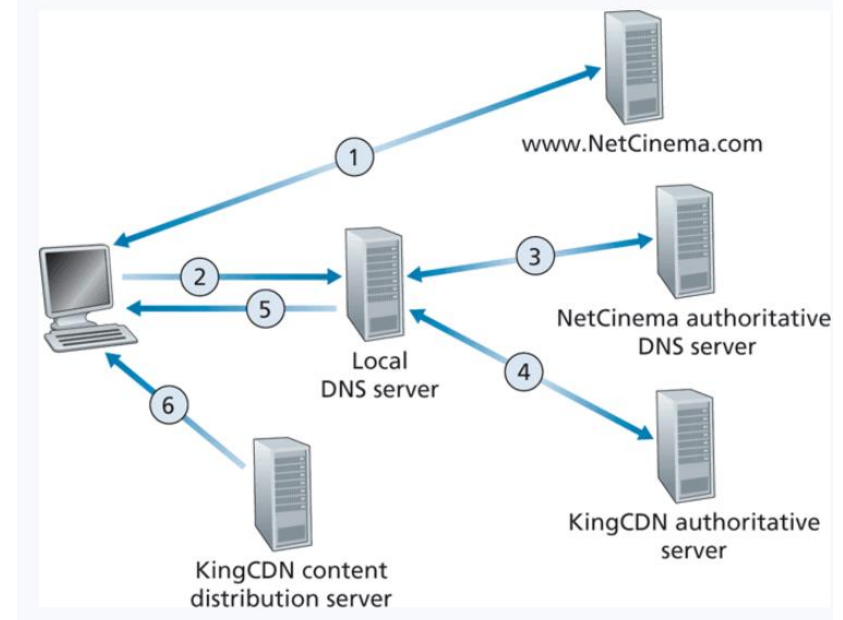
# CDNs



[https://\*\*cdn\*\*.discordapp.com/attachments/1017630358621143110/1017849417216294972/unknown.png](https://cdn.discordapp.com/attachments/1017630358621143110/1017849417216294972/unknown.png)

# CDNs

- Challenge: how does CDN DNS select “good” CDN node to stream to client
  - Pick CDN node **geographically** closest to client’s local DNS
  - Pick CDN node with **shortest delay** (or min # hops) to client (CDN nodes periodically ping access ISPs, reporting results to CDN DNS)



# CDNs

In HTTP steaming, video is stored at an HTTP server and retrieved with a GET request

Frames are sent to a client buffered and played back after a certain threshold

All clients receiving the same encoding of the video despite widely different bandwidth



# CDNs

In HTTP steaming, video is stored at an HTTP server and retrieved with a GET request

Frames are sent to a client buffered and played back after a certain threshold

All clients receiving the same encoding of the video despite widely different bandwidth

**Dynamic Adaptive Streaming over HTTP (DASH)**- video is encoded into several different versions

- Different encoding rate, different quality, etc

Clients dynamically request chunks of video segments every few seconds via GET requests

- Is the current bandwidth good? → Retrieve good quality
- Is the current bandwidth bad? → Retrieve ok quality

## CDNs


When streaming, you are consistently issuing GET requests

[illegible]

Name	Status	Type	Initiator	Size	Time	Waterfall
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	1.5 MB	48 ms	
watchtime?ns=yt&ei=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	30 ms	
qoe?fmt=399&afmt=251&cpn=D1QCo4_7KrXNdhs4...	204	xhr	VM685:201	39 B	28 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	256 kB	20 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	2.1 MB	50 ms	
log_event?alt=json&key=AlzaSyAO_FJ2SlqU8Q4STEH...	200	xhr	VM685:201	77 B	46 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	1.4 MB	56 ms	
watchtime?ns=yt&ei=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	31 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	2.1 MB	55 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	391 kB	23 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	1.6 MB	32 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	1.3 MB	53 ms	
watchtime?ns=yt&ei=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	32 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	2.1 MB	60 ms	
videoplayback?expire=1663023840&ei=gGYf7ICCoO...	200	fetch	VM685:223	1.6 MB	65 ms	
data:image/png;base...	200	png	Other	(memory cache)	0 ms	


18 requests | 14.8 MB transferred | 14.8 MB resources

# CDNs



Studio master versions are uploaded to a private Amazon cloud


# CDNs



Studio master versions are uploaded to a private Amazon cloud

Videos are processed into many different formats, allowing for DASH

# CDNs



Studio master versions are uploaded to a private Amazon cloud

Videos are processed into many different formats, allowing for DASH

Versions are uploaded to Netflix's CDNs