

CSCI 466: Networks

Lecture 6: UDP and RDT

Reese Pearsall
Fall 2022

Announcements

PA1 Due Monday September 26th

- Files must be pushed to a PA1 folder on your GitHub Repo
- Video demo is required
- Submit your repo link to D2L when finished

Next Friday will be a work day + help session. No Lecture

Announcements

PA1

PA1

UDP Example

OSI Model

Application Layer

Presentation Layer *

Session Layer *

Transport Layer

Network Layer

Data Link Layer

Physical Layer

Application Layer

Messages from Network Applications




Physical Layer

Bits being transmitted over some medium


**In the textbook, they condense it to a 5-layer model, but 7 layers is what is most used*

Transport Layer

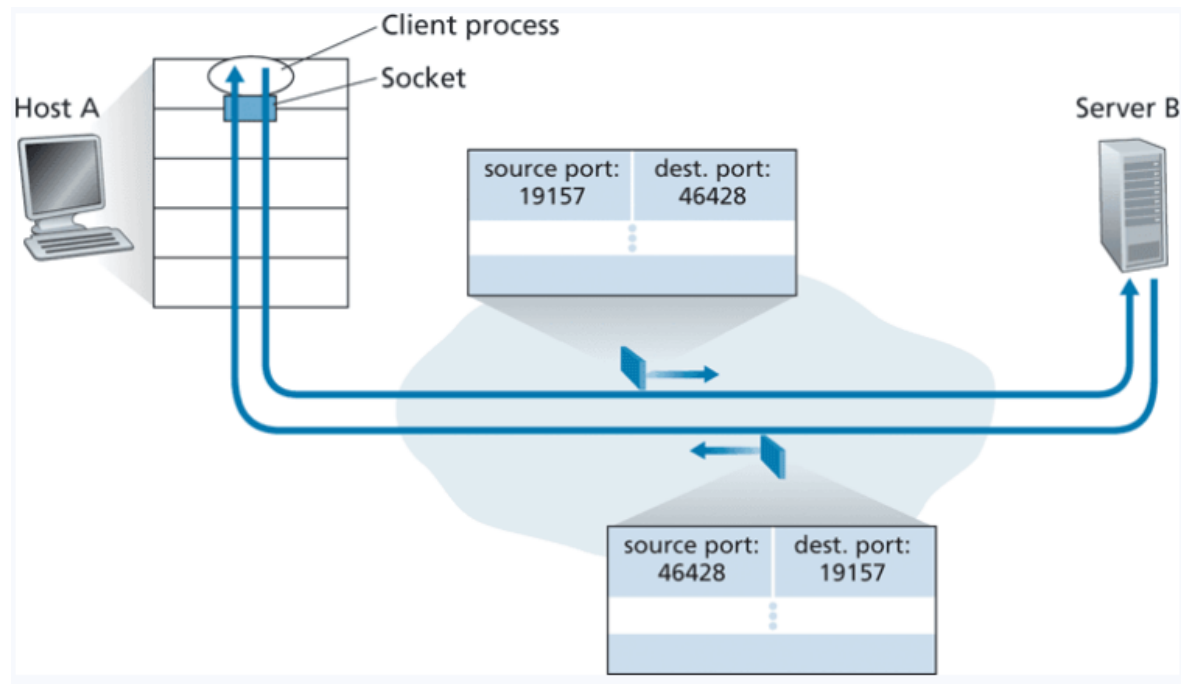
Multiplexing is the process of gathering chunks from sockets, encapsulating chunks with header information, and passing the segment into the network layer



Demultiplexing is the receiving segments from the transport layer and delivering the segment to the correct socket.



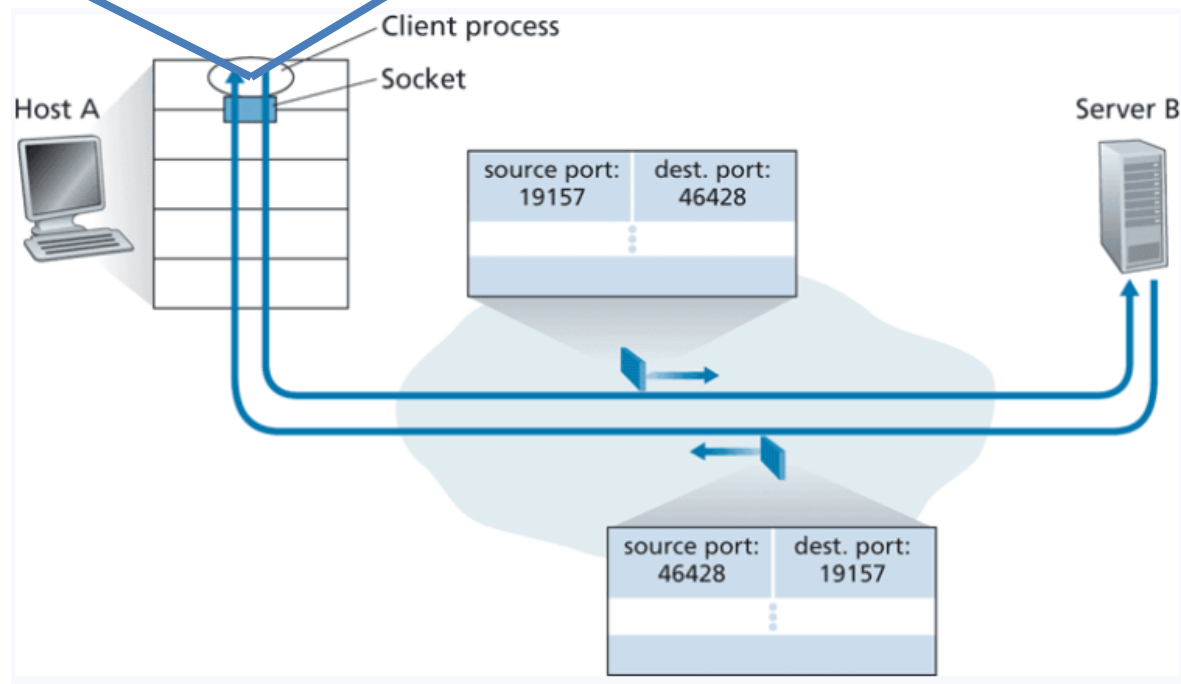
Transport Layer



Transport Layer

UDP sockets are identified by a two-tuple

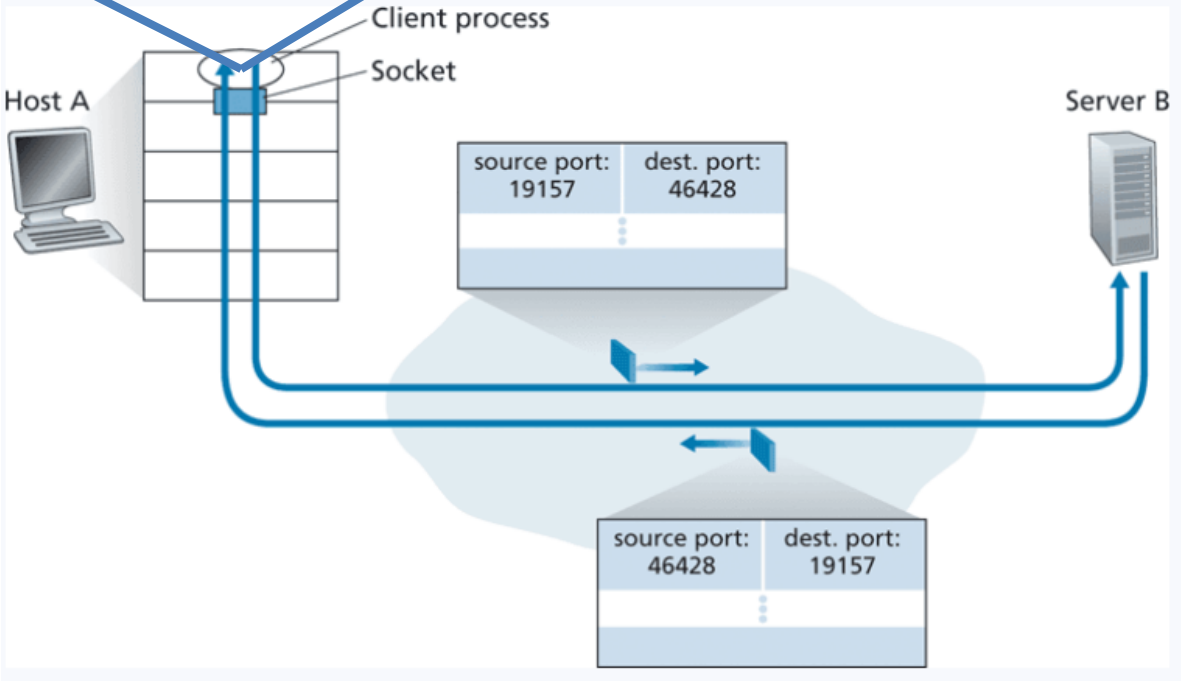
(destination IP address, destination port number)



Transport Layer

UDP sockets are identified by a two-tuple

```
(destination IP address, destination port number)
```



(92.7.32.223, 8000)


Transport Layer

UDP sockets are identified by a two-tuple


(destination IP address, destination port number)



92.7.32.223 8000



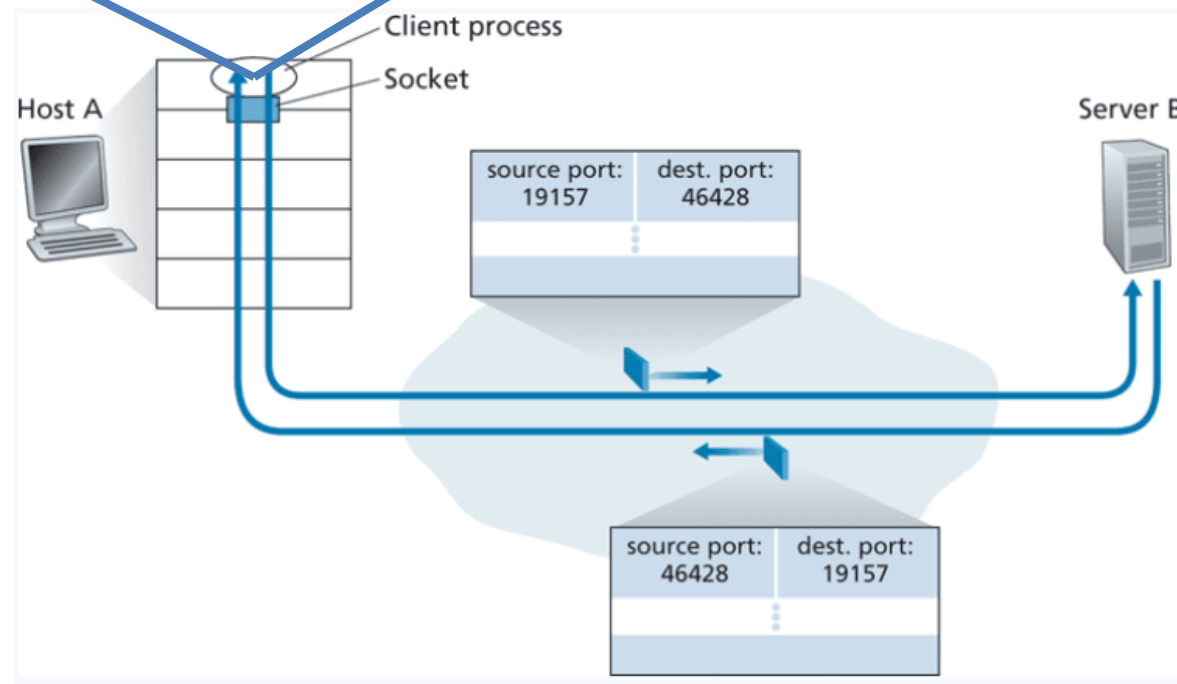
92.7.32.223 8000



(92.7.32.223, 8000)

The socket won't be able to distinguish between the two endpoints

The two segments will be directed to the same process

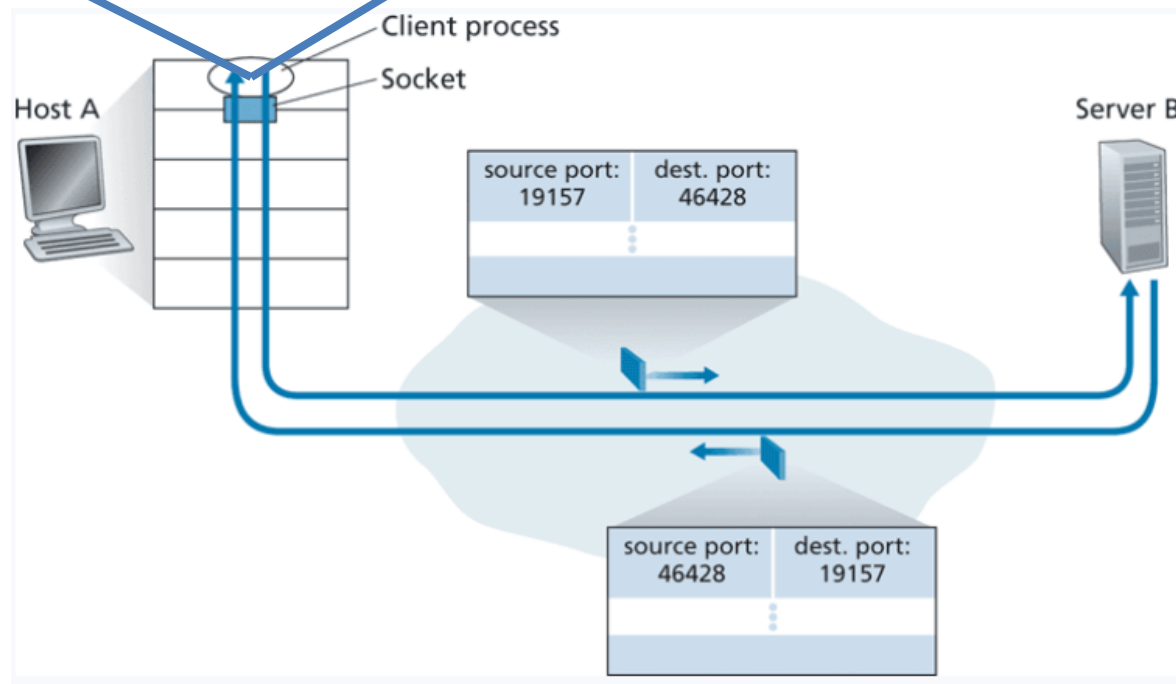


Transport Layer

TCP sockets are identified by a four-tuple

A host will demultiplex segments using all of these values

(source IP address, source port number, destination IP address, destination port number)

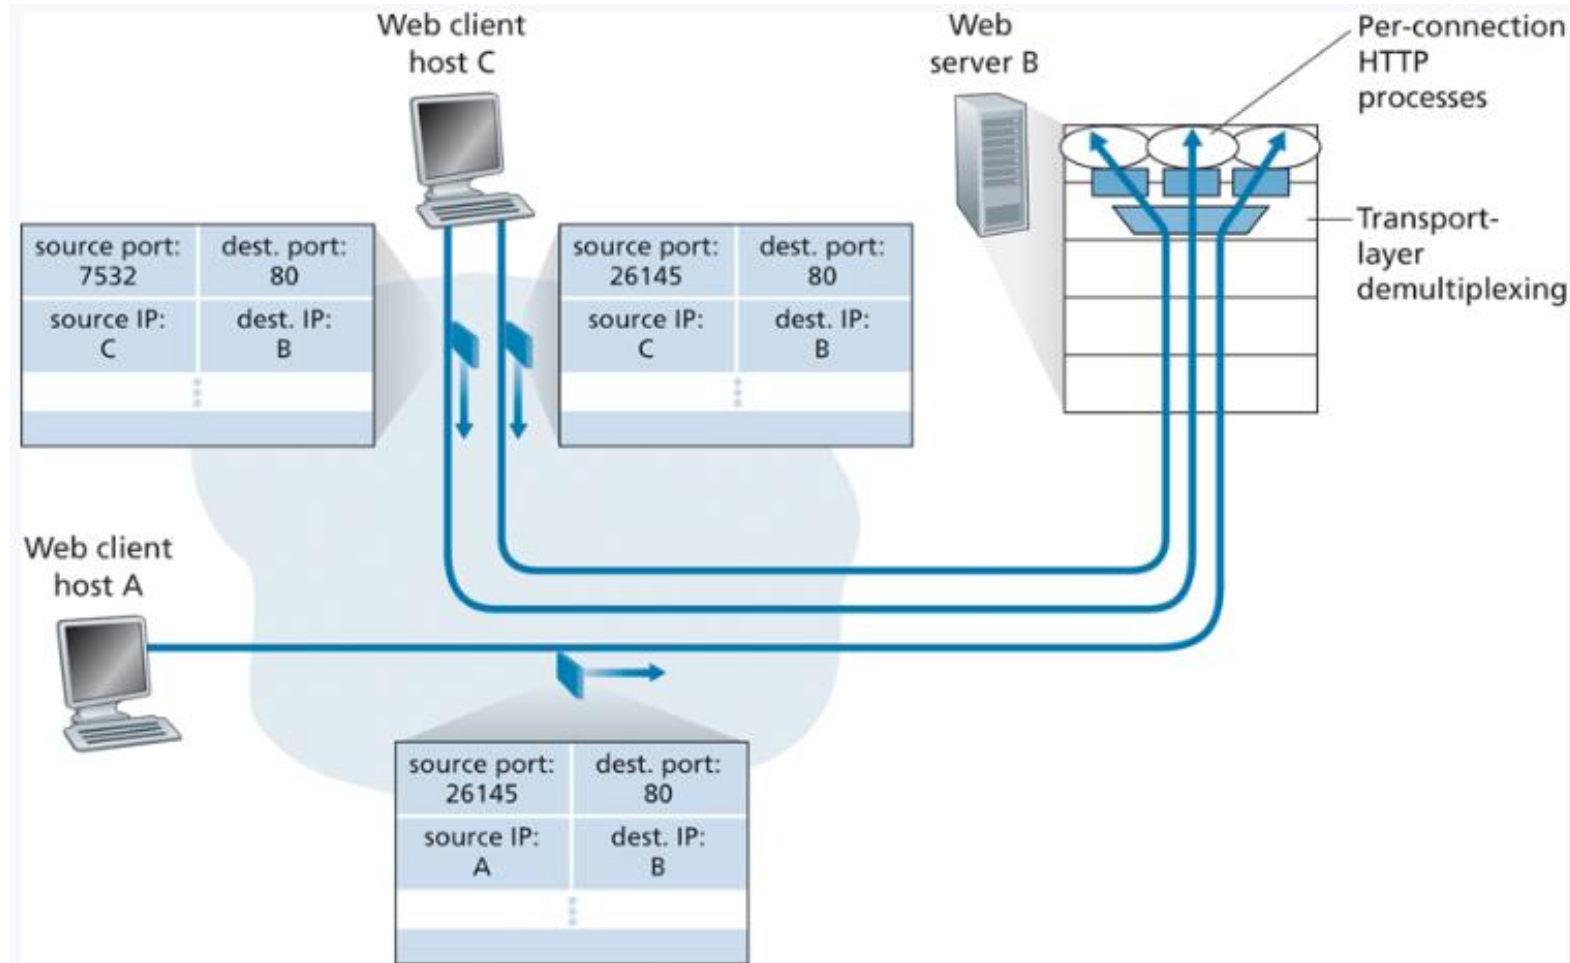


TCP servers will have a “welcoming socket” before creating the processes' socket

Transport Layer

(source IP address, source port number, destination IP address, destination port number)

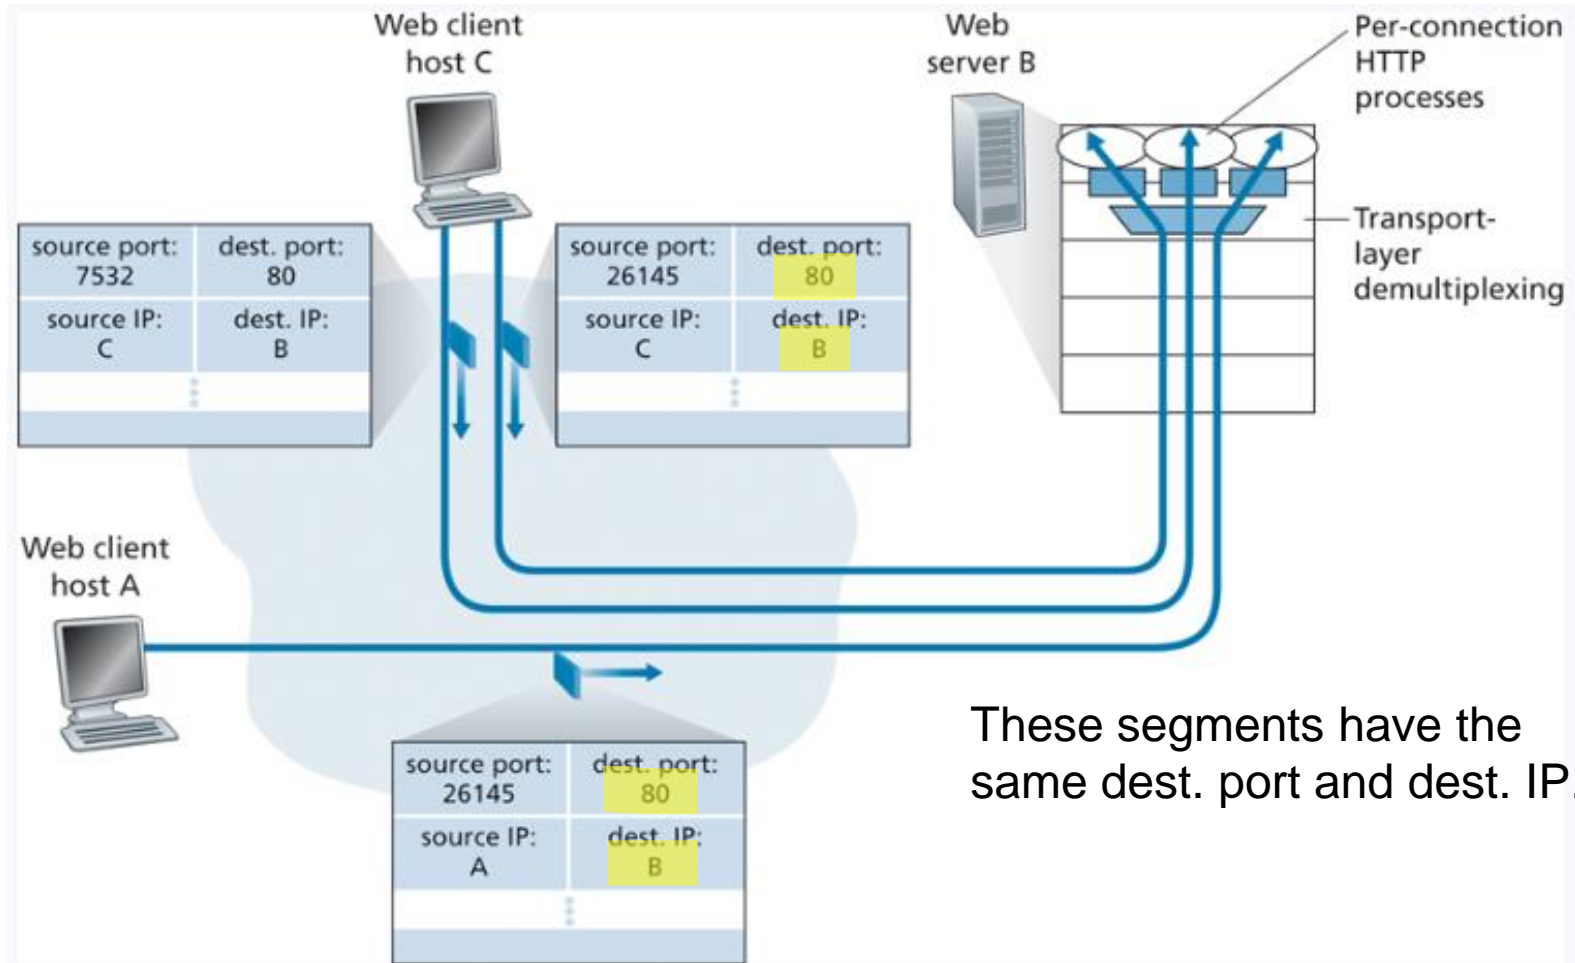
TCP sockets are identified by a four-tuple



Transport Layer

(source IP address, source port number, destination IP address, destination port number)

TCP sockets are identified by a four-tuple

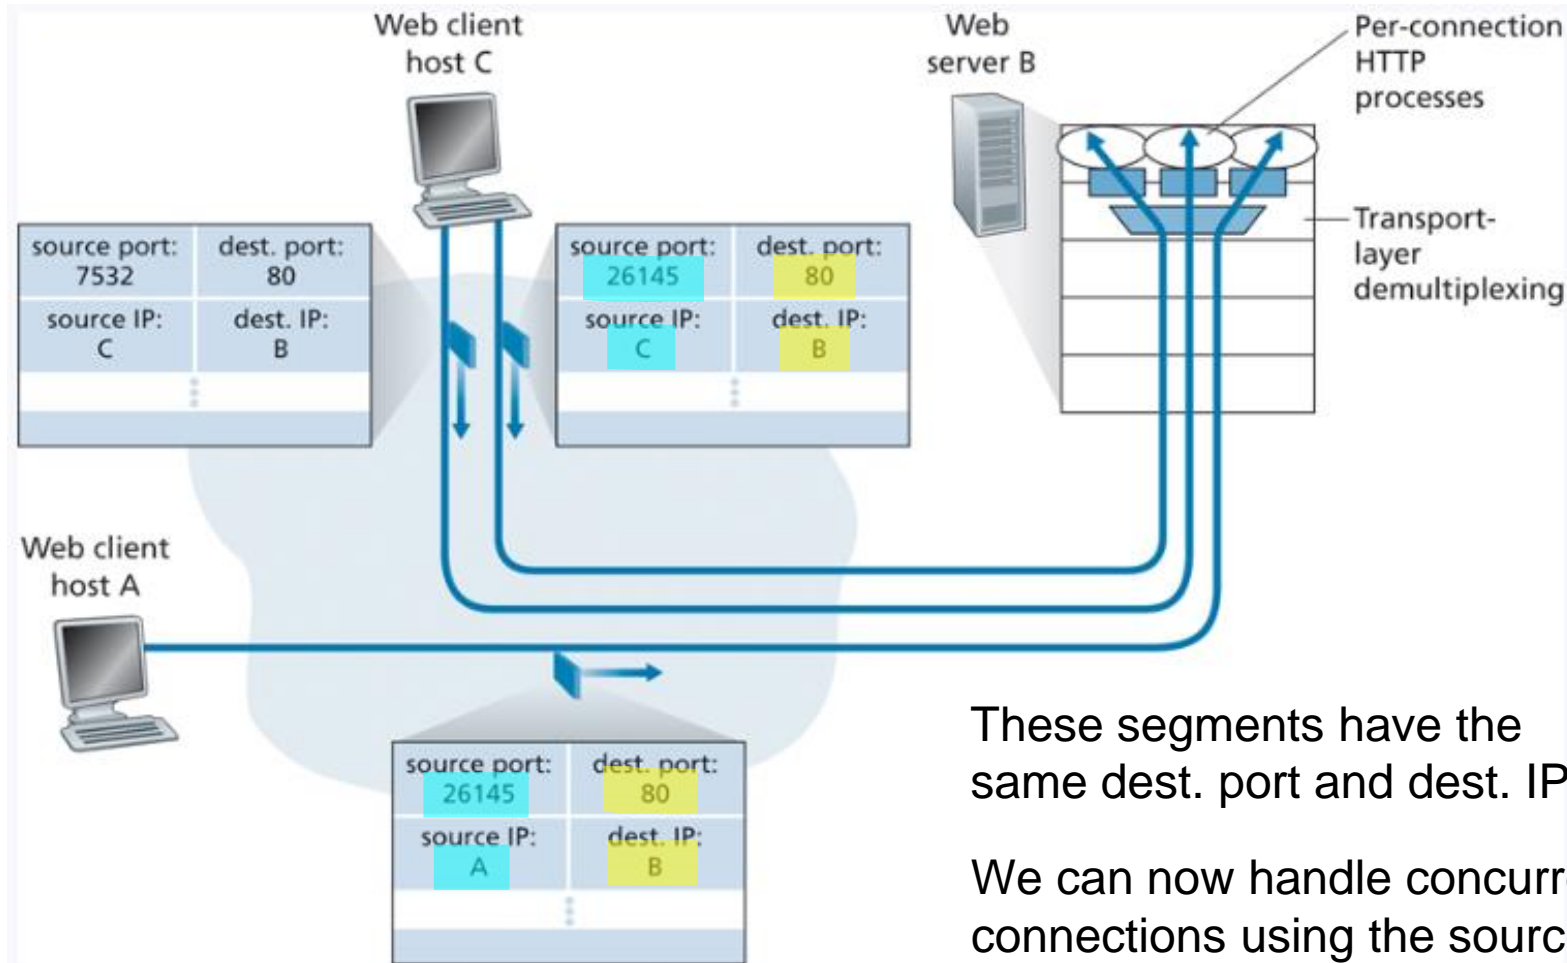


These segments have the same dest. port and dest. IP.

Transport Layer

(source IP address, source port number, destination IP address, destination port number)

TCP sockets are identified by a four-tuple



These segments have the same dest. port and dest. IP.

We can now handle concurrent connections using the source information

Transport Layer

Application	Application-Layer Protocol	Underlying Transport Protocol
Electronic mail	SMTP	TCP
Remote terminal access	Telnet	TCP
Secure remote terminal access	SSH	TCP
Web	HTTP, HTTP/3	TCP (for HTTP), UDP (for HTTP/3)
File transfer	FTP	TCP
Remote file server	NFS	Typically UDP
Streaming multimedia	DASH	TCP
Internet telephony	typically proprietary	UDP or TCP
Network management	SNMP	Typically UDP
Name translation	DNS	Typically UDP

“Do as little as possible and give it our best effort”

- Bare bones and connectionless

Advantages of UDP:

- Immediate transmission
 - No connection establishment
- Lower memory requirements
 - No connection state, 8B for UDP vs 20B for TCP

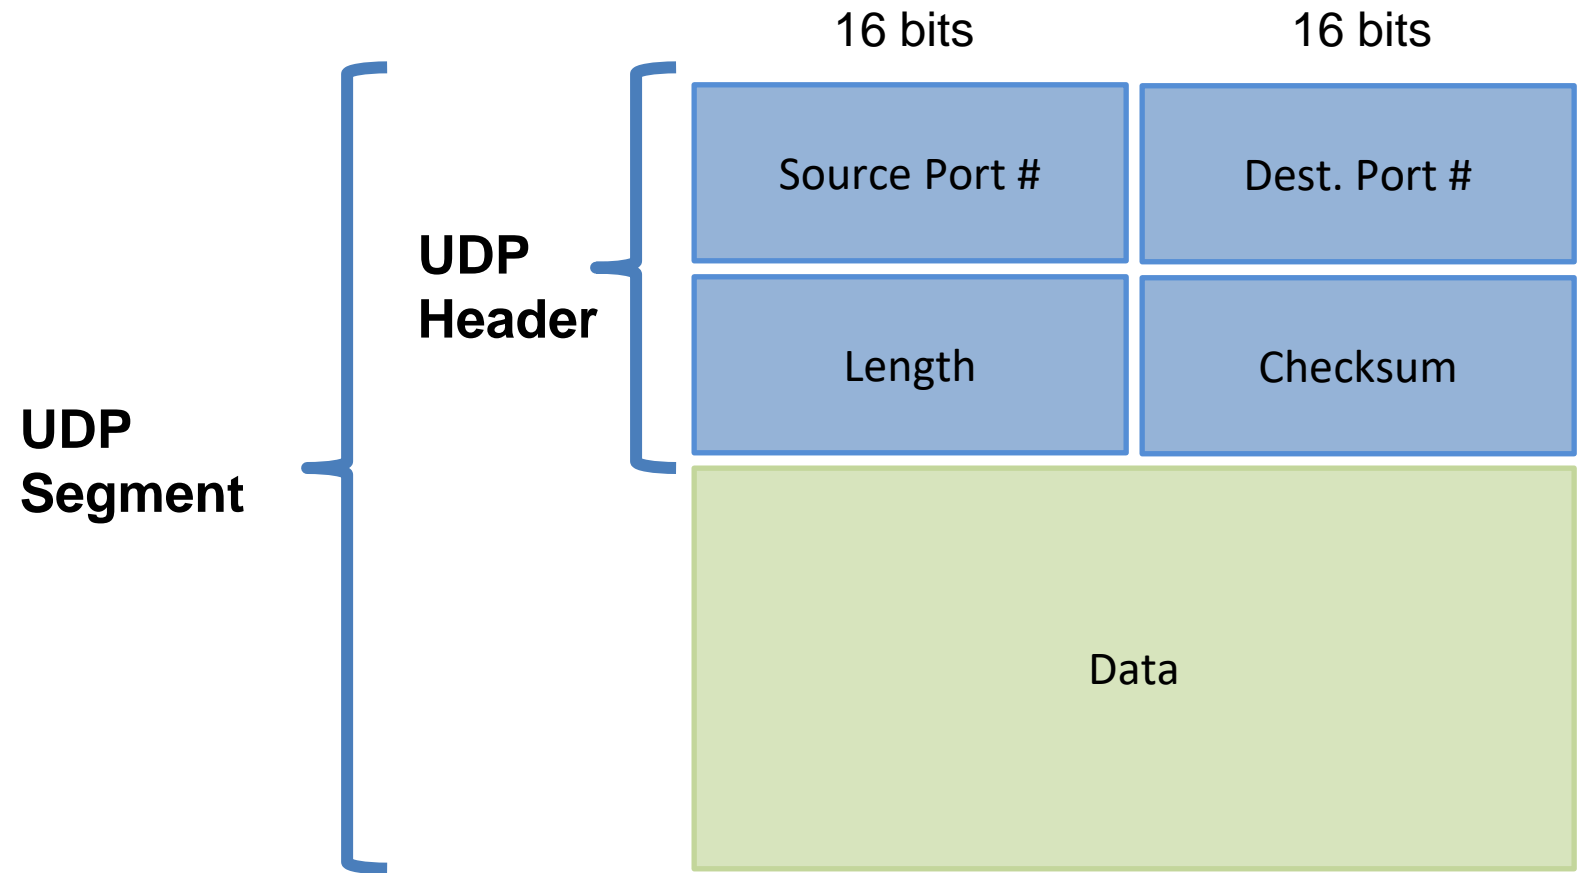
Disadvantages

- No congestion control
- No guarantee for in-order delivery
- Reliability and Flow control

QUIC (Quick UDP Internet Connection) is a transport layer protocol that adds reliability to UDP

Transport Layer

UDP “Do as little as possible and give it our best effort”

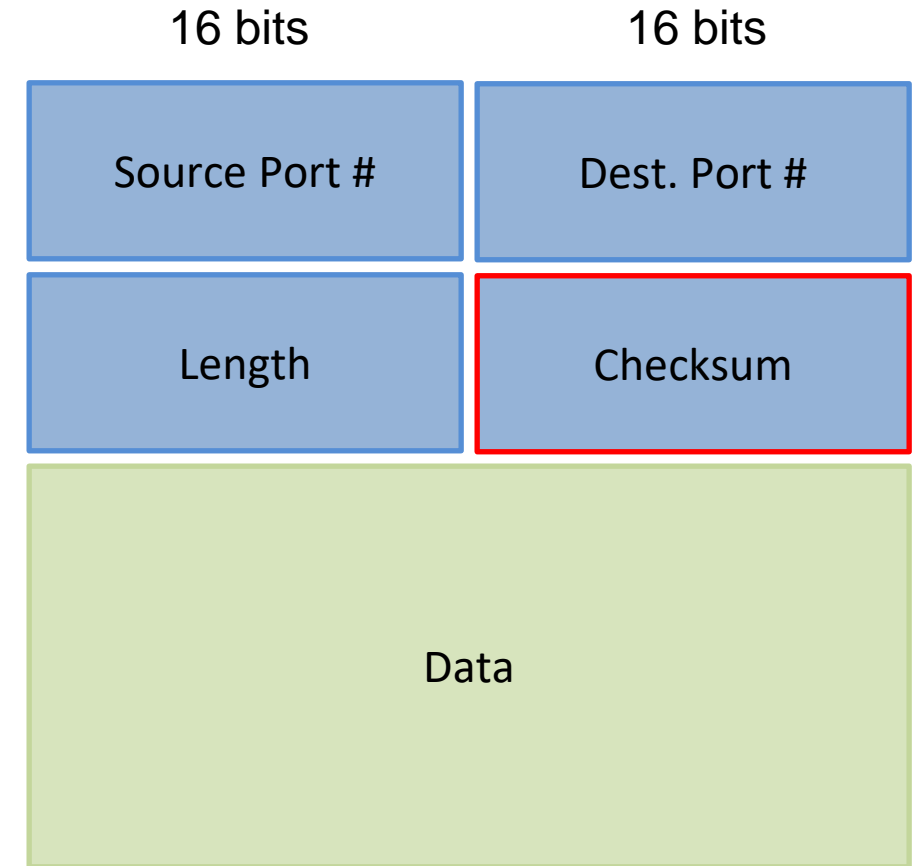


Transport Layer

UDP “Do as little as possible and give it our best effort”

UDP provides a **checksum** that is used to determine whether bits within the UDP segment have been altered

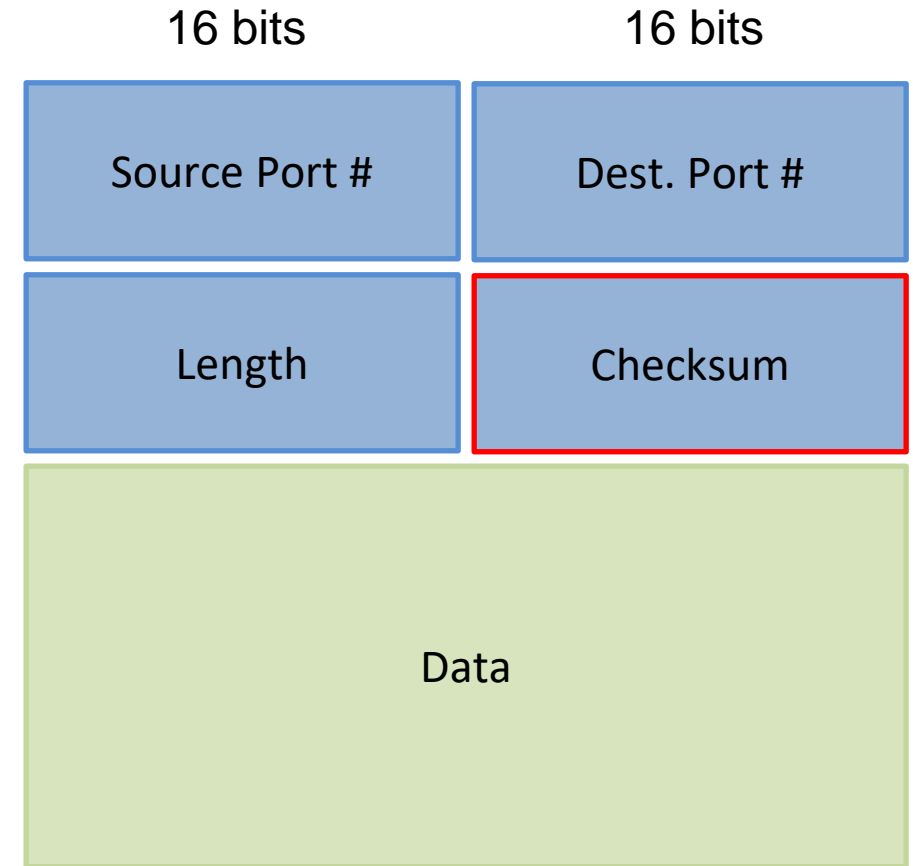
It is the sum of all n bit words in the segment followed by the 1s complement



Transport Layer

0110011001100000
0101010101010101
1000111100001100

UDP “Do as little as possible and give it our best effort”




Transport Layer

UDP “Do as little as possible and give it our best effort”

0110011001100000
0101010101010101
1000111100001100

0110011001100000
0101010101010101

1011101110110101



Transport Layer

UDP “Do as little as possible and give it our best effort”

0110011001100000


0101010101010101

1000111100001100

1011101110110101

1000111100001100

0100101011000010



Transport Layer

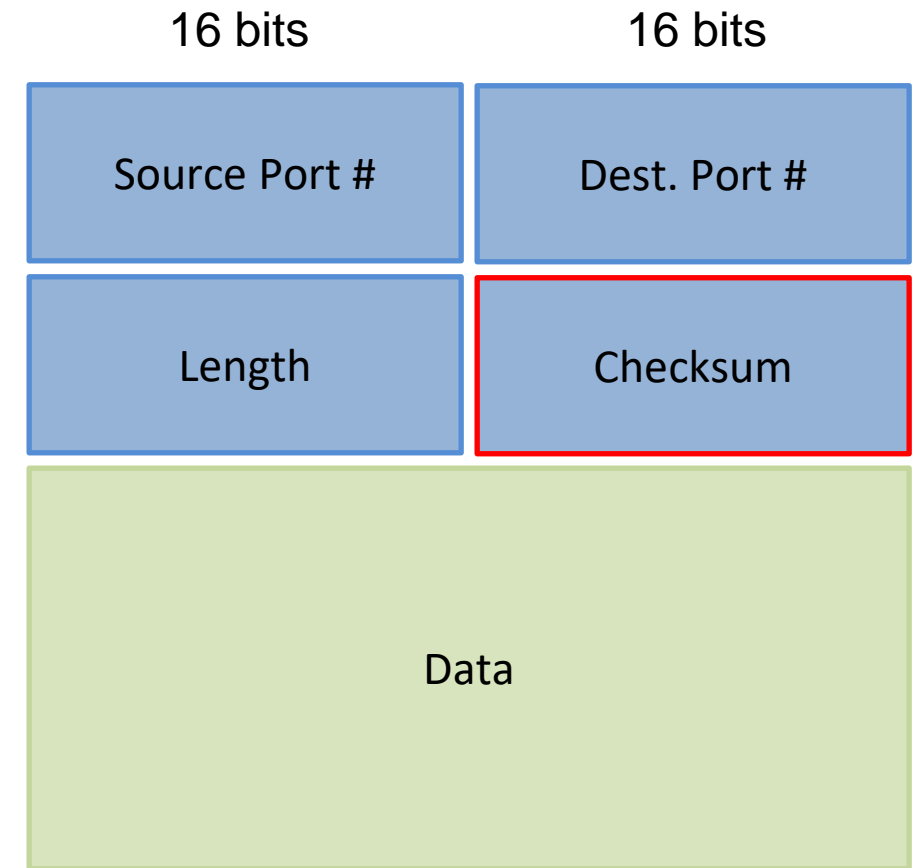
UDP “Do as little as possible and give it our best effort”

0110011001100000
0101010101010101
1000111100001100

1011101110110101
1000111100001100

0100101011000010 (one's complement)

1011010100111101 = checksum




Transport Layer

UDP “Do as little as possible and give it our best effort”

0110011001100000
0101010101010101
1000111100001100
1011010100111101

Receiving packet: Sum up all 16-bit words → 1111111111111111

If there are zeros, errors were introduced into the packet



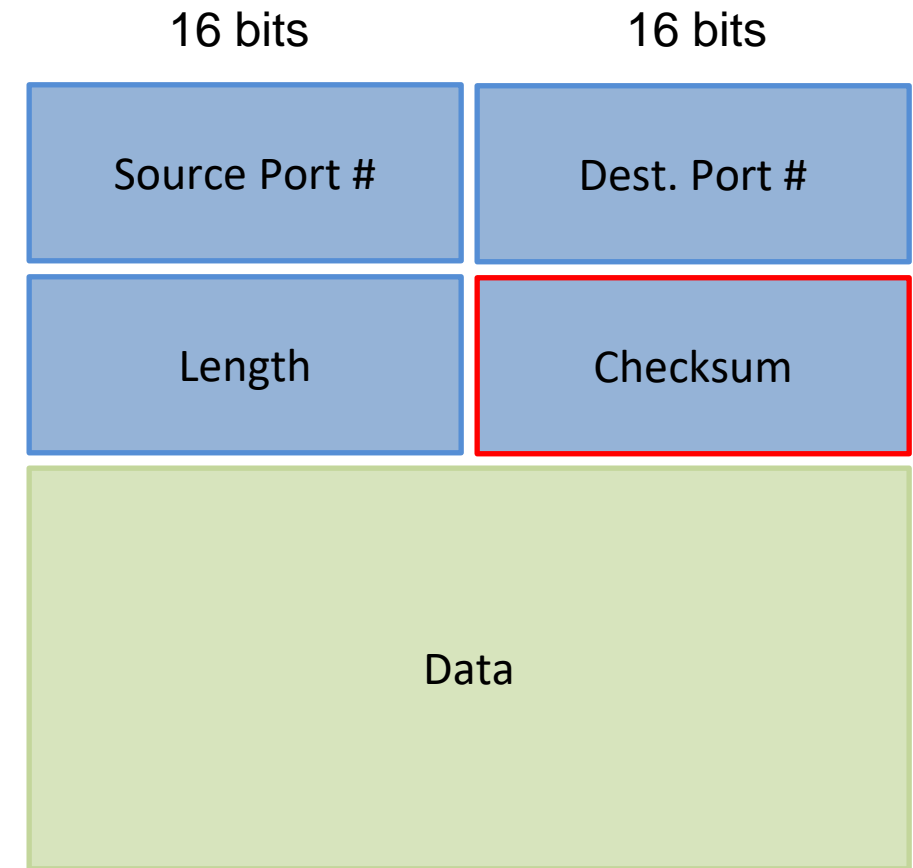
Transport Layer

UDP “Do as little as possible and give it our best effort”

Why do error checking here?

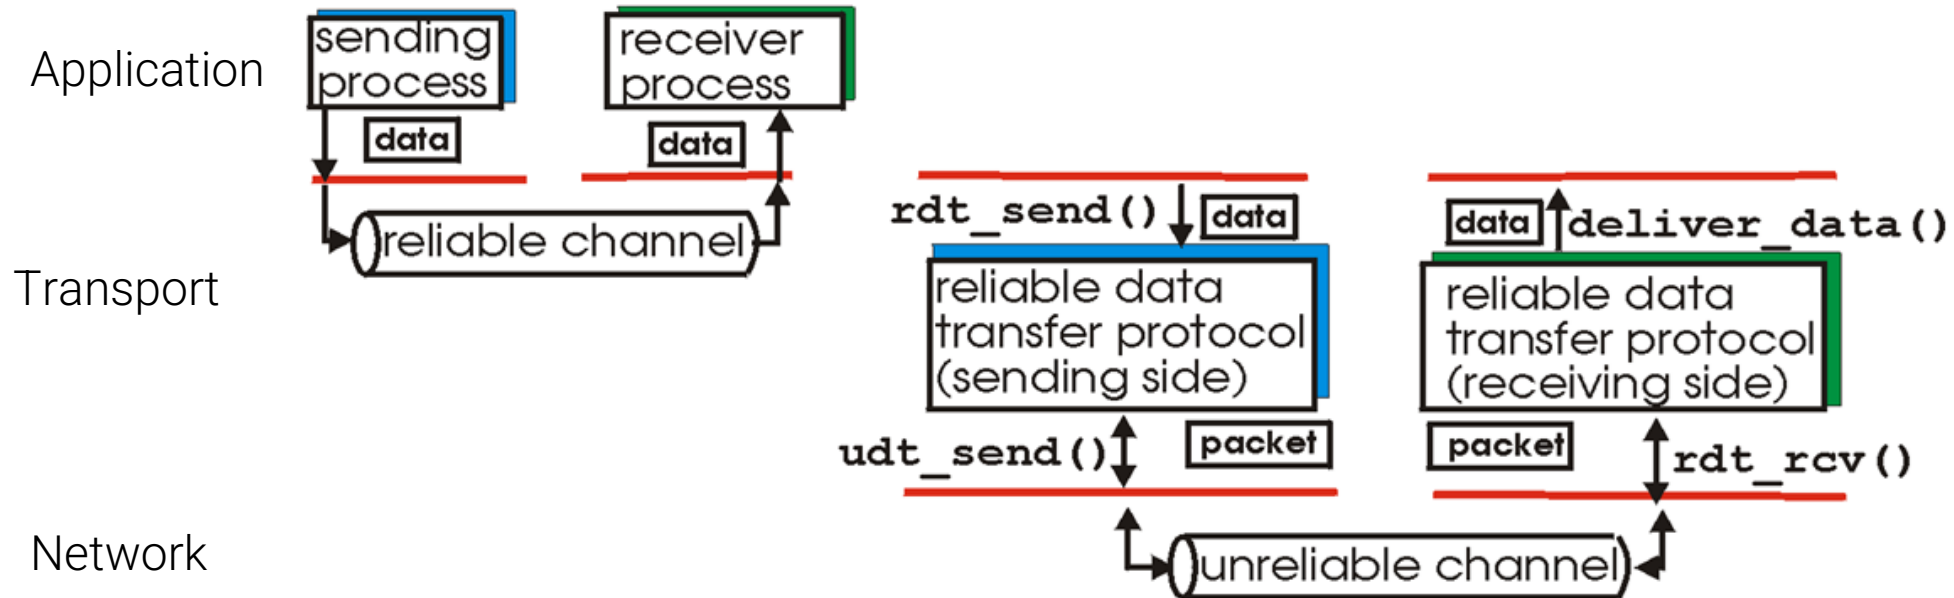
End-to-end principle states that since certain functionality such as error detection, must be implemented on an end-end bases

Functionality places at the lower levels may be redundant or of little value when compared to the cost of providing them at a higher level



Transport Layer

Reliable Data Transfer



Characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

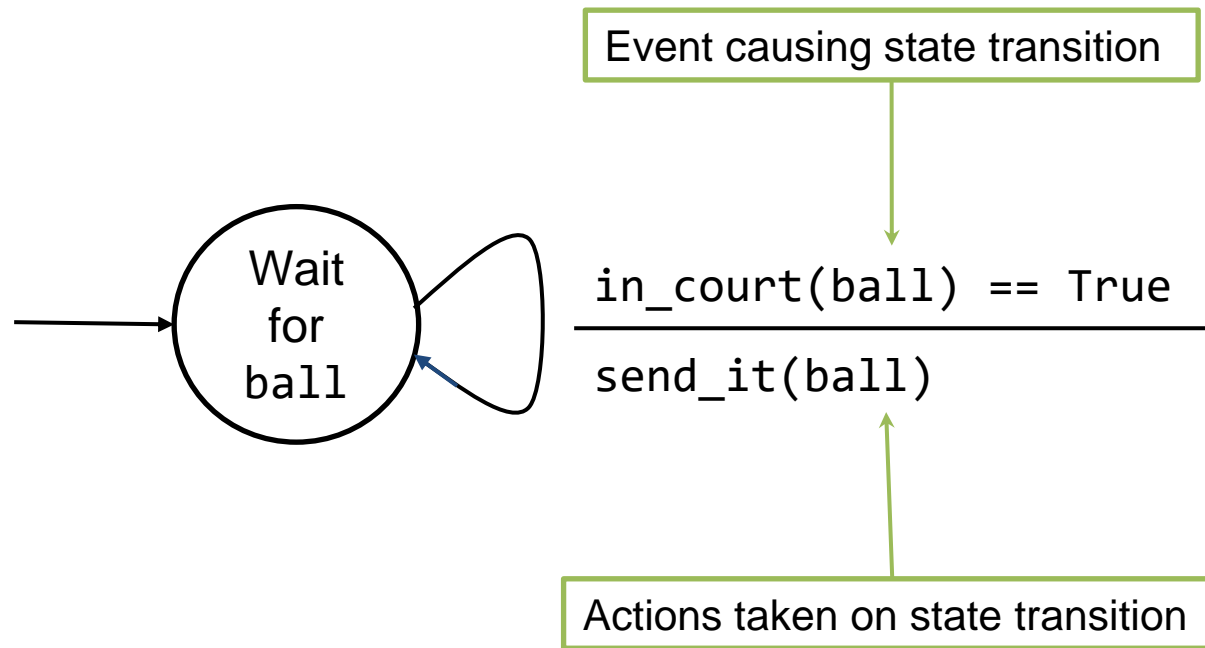
What are some ways in which the network channel can be unreliable?

Things to consider:

- Corruption
- Loss of packets
- Duplicate delivery



Bruce Lee FSM



Transport Layer

Reliable Data Transfer 1.0

RDT 1.0

Assumptions:

- Unidirectional long data flows
- Perfectly reliable channel
- No bit errors
- No packet loss
- No packet reordering




Transport Layer

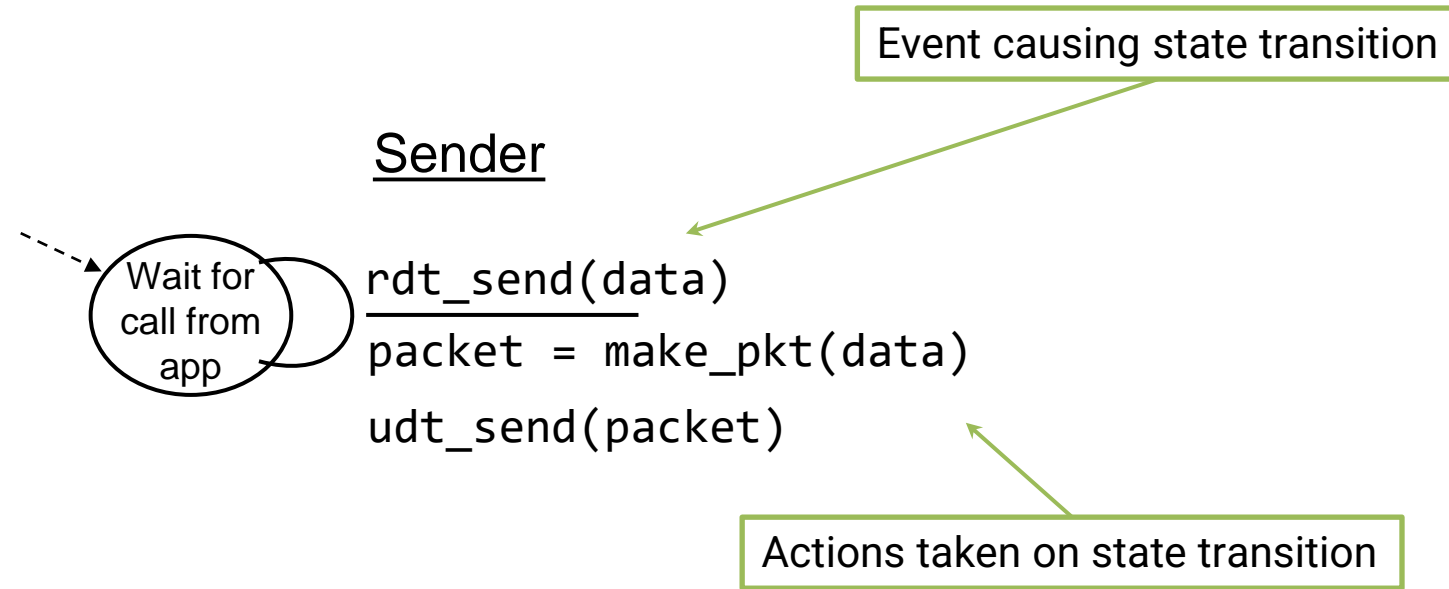
RDT 1.0

Assumptions:

- Unidirectional long data flows
- Perfectly reliable channel
- No bit errors
- No packet loss
- No packet reordering



Reliable Data Transfer 1.0




Transport Layer

RDT 1.0

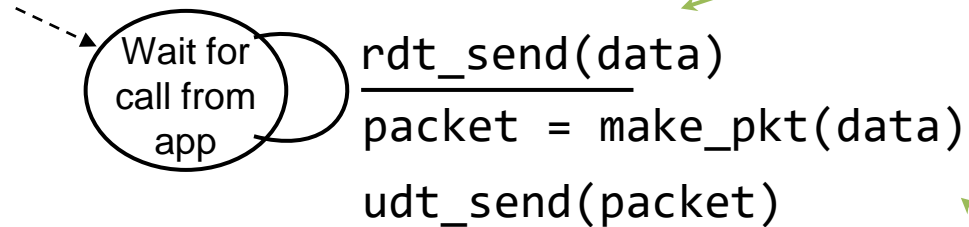
Assumptions:

- Unidirectional long data flows
- Perfectly reliable channel
- No bit errors
- No packet loss
- No packet reordering




Reliable Data Transfer 1.0

Sender



Receiver




Transport Layer

Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

How can we detect errors?



Transport Layer


Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

How can we detect errors?

-Checksum




RDT 2.0

Potential for bit errors

How can we detect errors?

-Checksum

What is a good way to handle and prevent errors?



RDT 2.0


Potential for bit errors

How can we detect errors?

- Checksum

What is a good way to handle and prevent errors?

- Acknowledged packet, Ask for retransmit if needed



Transport Layer

Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

How can we detect errors?

-Checksum

What is a good way to handle and prevent errors?
- Acknowledged packet, Ask for retransmit if needed

Stop-and-wait: sender sends one packet, then waits for receiver response

Sender

`rdt_send(data)`

`sndpkt = make_pkt(data, checksum)`

`udt_send(sndpkt)`

Wait for call
from appl



Transport Layer

Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

How can we detect errors?

-Checksum

What is a good way to handle and prevent errors?
- Acknowledged packet, Ask for retransmit if needed

Stop-and-wait: sender sends one packet, then waits for receiver response

Sender

rdt_send(data)

sndpkt = make_pkt(data, checksum)

udt_send(sndpkt)

Wait for call
from appl

rdt_rcv(rcvpkt) &&

isNAK(rcvpkt)

udt_send(sndpkt)

Wait for ACK
or NAK

Transport Layer

Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

How can we detect errors?

-Checksum

What is a good way to handle and prevent errors?
- Acknowledged packet, Ask for retransmit if needed

Stop-and-wait: sender sends one packet, then waits for receiver response

Sender

rdt_send(data)

sndpkt = make_pkt(data, checksum)

udt_send(sndpkt)

Wait for call
from appl

Wait for ACK
or NAK

rdt_rcv(rcvpkt) &&

isNAK(rcvpkt)

udt_send(sndpkt)

rdt_rcv(rcvpkt) && isACK(rcvpkt)

Λ

Transport Layer

Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

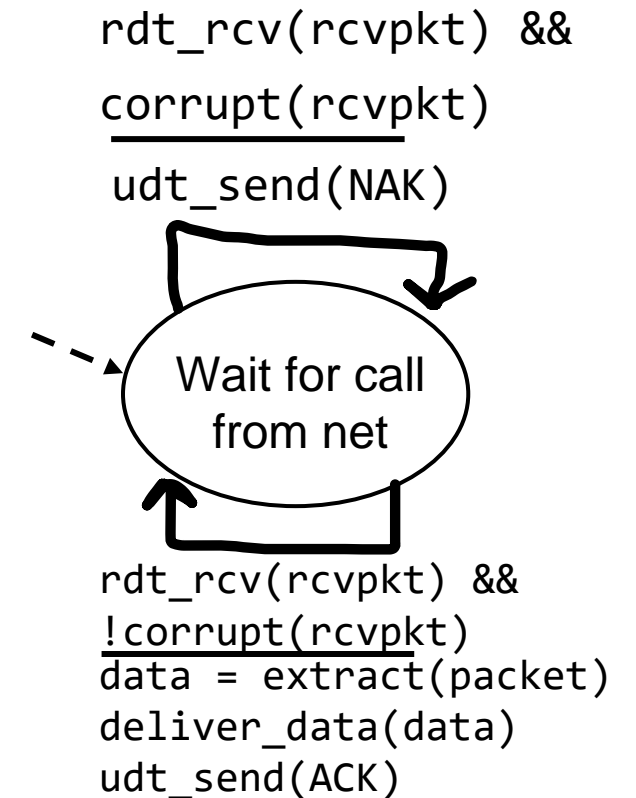
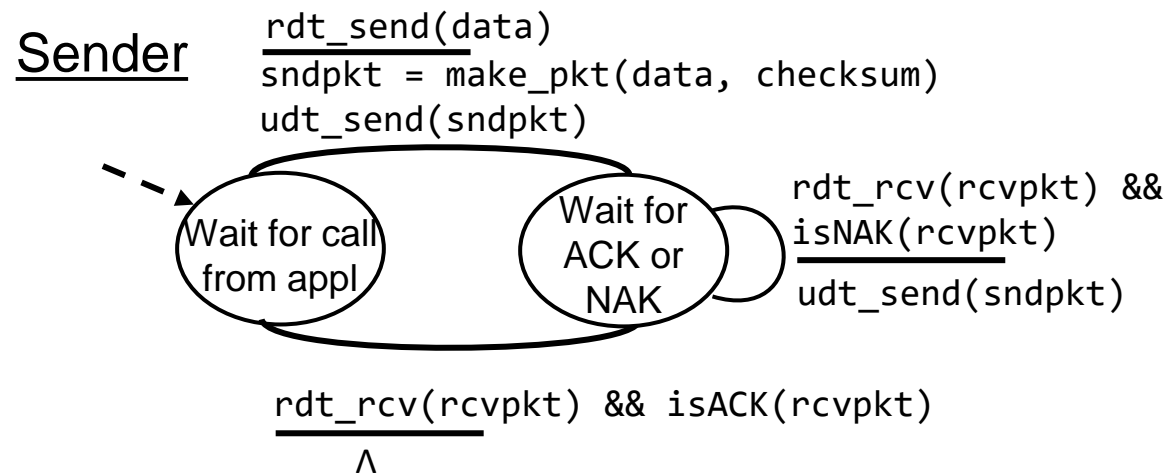
How can we detect errors?

-Checksum

What is a good way to handle and prevent errors?
- Acknowledged packet, Ask for retransmit if needed

Stop-and-wait: sender sends one packet, then waits for receiver response

Receiver



Transport Layer

Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

How can we detect errors?

-Checksum


What is a good way to handle and prevent errors?
- Acknowledged packet, Ask for retransmit if needed

Stop-and-wait: sender sends one packet, then waits for receiver response

Receiver

Sender

rdt_send(data)
sndpkt = make_pkt(data, checksum)
udt_send(sndpkt)



rdt_rcv(rcvpkt) && isACK(rcvpkt)
Λ

What happens if ACK/NAK Corrupted?

→ Duplicate delivery, or no retransmission


Solution?

→ Retransmit is CORRUPT packet received

How to deal we duplicate Packets?

→ ???

rdt_rcv(rcvpkt) &&
corrupt(rcvpkt)
udt_send(NAK)



rdt_rcv(rcvpkt) &&
!corrupt(rcvpkt)
data = extract(packet)
deliver_data(data)
udt_send(ACK)

Transport Layer

Reliable Data Transfer 2.0

RDT 2.0

Potential for bit errors

How can we detect errors?

-Checksum


What is a good way to handle and prevent errors?
- Acknowledged packet, Ask for retransmit if needed

Stop-and-wait: sender sends one packet, then waits for receiver response

Receiver

Sender

rdt_send(data)
sndpkt = make_pkt(data, checksum)
udt_send(sndpkt)



rdt_rcv(rcvpkt) && isACK(rcvpkt)
Λ

What happens if ACK/NAK Corrupted?

→ Duplicate delivery, or no retransmission


Solution?

→ Retransmit is CORRUPT packet received

How to deal we duplicate Packets?

→ Sequence Number

rdt_rcv(rcvpkt) &&
corrupt(rcvpkt)
udt_send(NAK)



rdt_rcv(rcvpkt) &&
!corrupt(rcvpkt)
data = extract(packet)
deliver_data(data)
udt_send(ACK)

Transport Layer

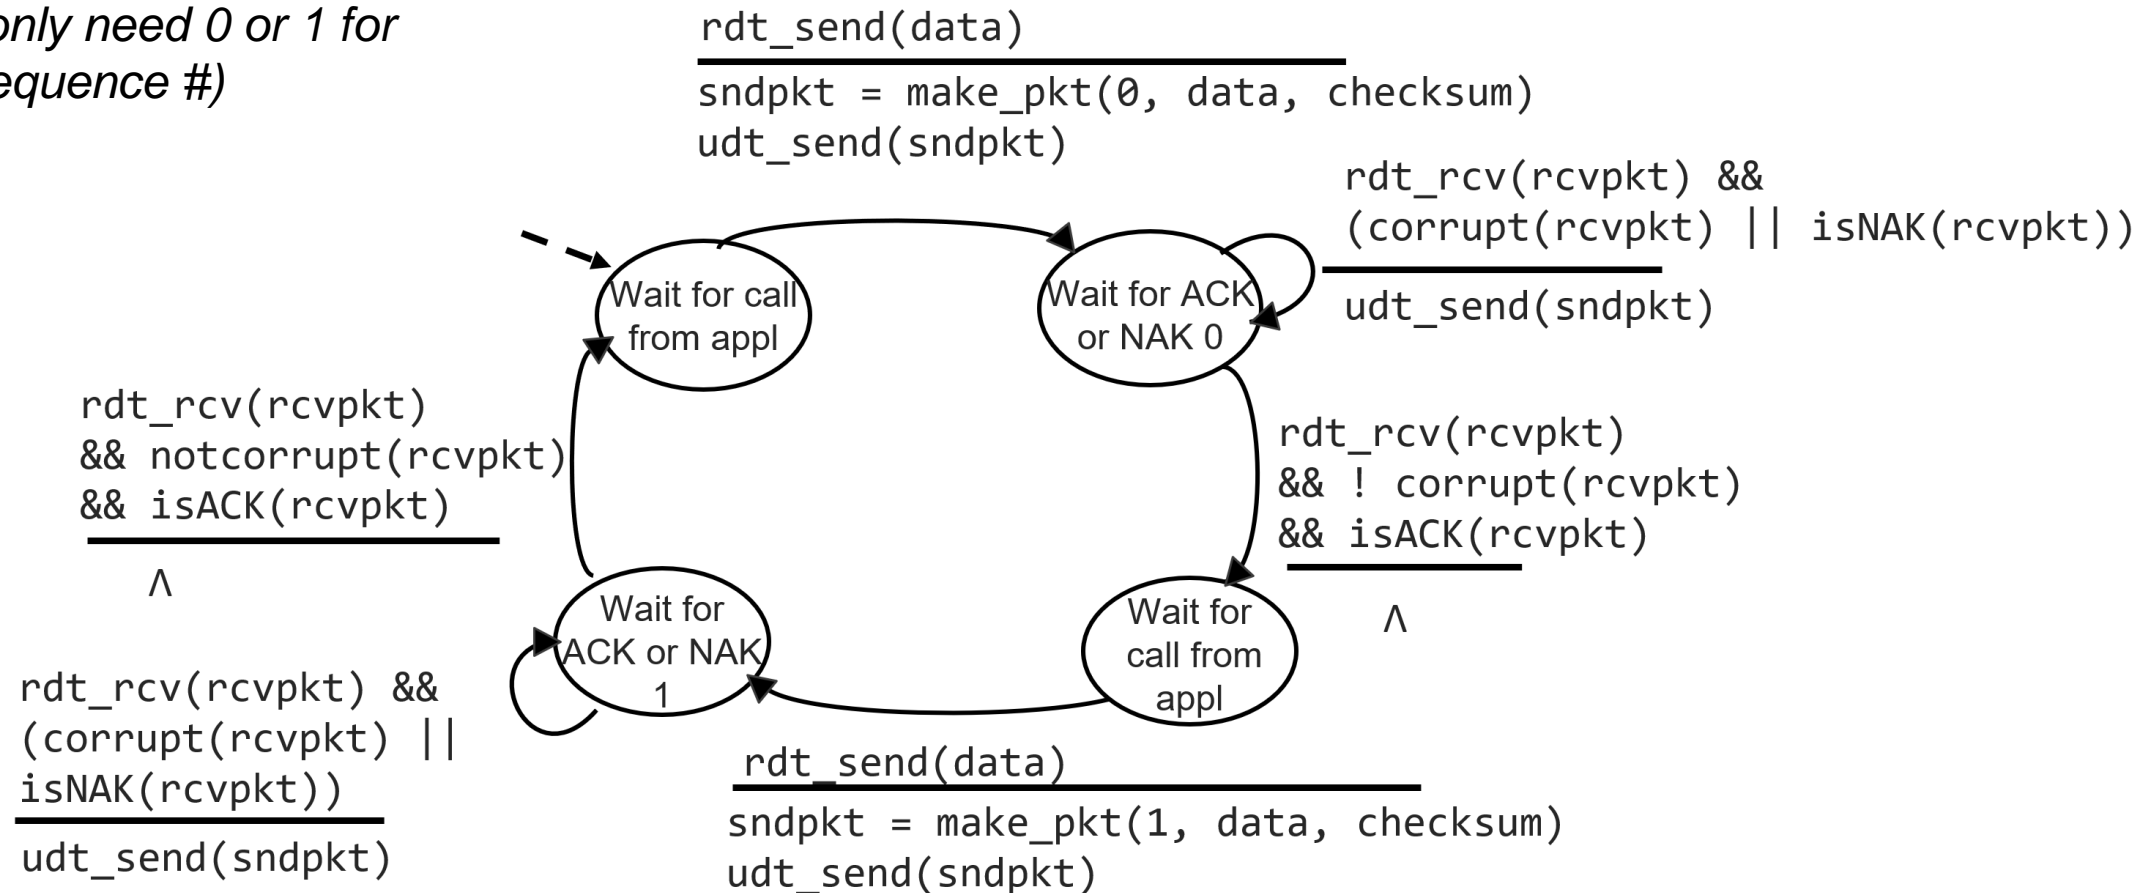
Reliable Data Transfer 2.1

RDT 2.1

Potential for bit errors and garbled ACKs

Stop-and-wait: sender sends one packet, then waits for receiver response

(We only need 0 or 1 for the sequence #)



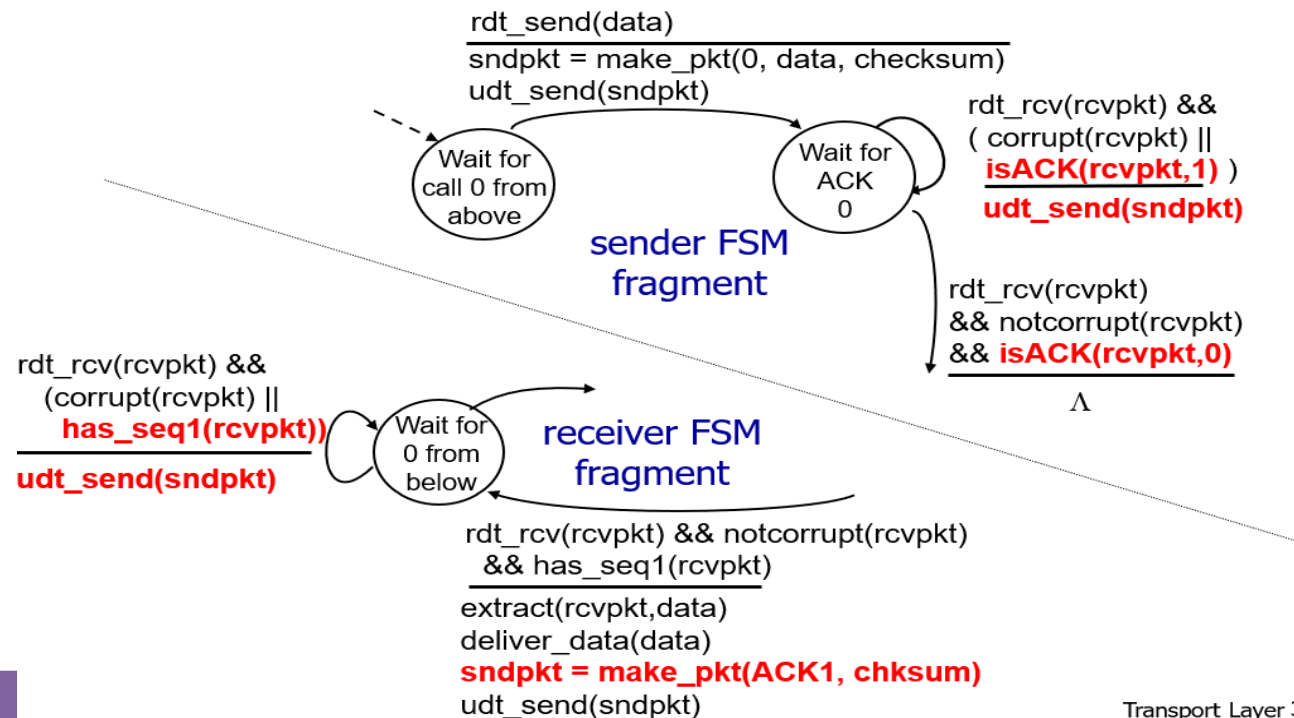
RDT 2.2

Same as rdt2.1, using only ACKs

Instead of NAK, receiver send the ACK for last pkt received successfully.

A.K.A Receiver must explicitly include seq # of pkt being ACKed

Duplicate ACK at sender results in same action as NAK: retransmit current pkt



RDT 3.0

Packets can get lost or dropped

We will still need checksums, seq #, ACKS, but we need more

What if an ACK gets dropped? Sender is stuck waiting forever

Solution?

RDT 3.0

Packets can get lost or dropped

We will still need checksums, seq #, ACKS, but we need more

What if an ACK gets dropped? Sender is stuck waiting forever

Sender should wait a “reasonable” amount of time for an ack

A timer!!!

Retransmit if no ACK received in X amount of time

What if the ACK is just taking a really long time to arrive?

RDT 3.0

Packets can get lost or dropped

We will still need checksums, seq #, ACKS, but we need more

What if an ACK gets dropped? Sender is stuck waiting forever

Sender should wait a “reasonable” amount of time for an ack

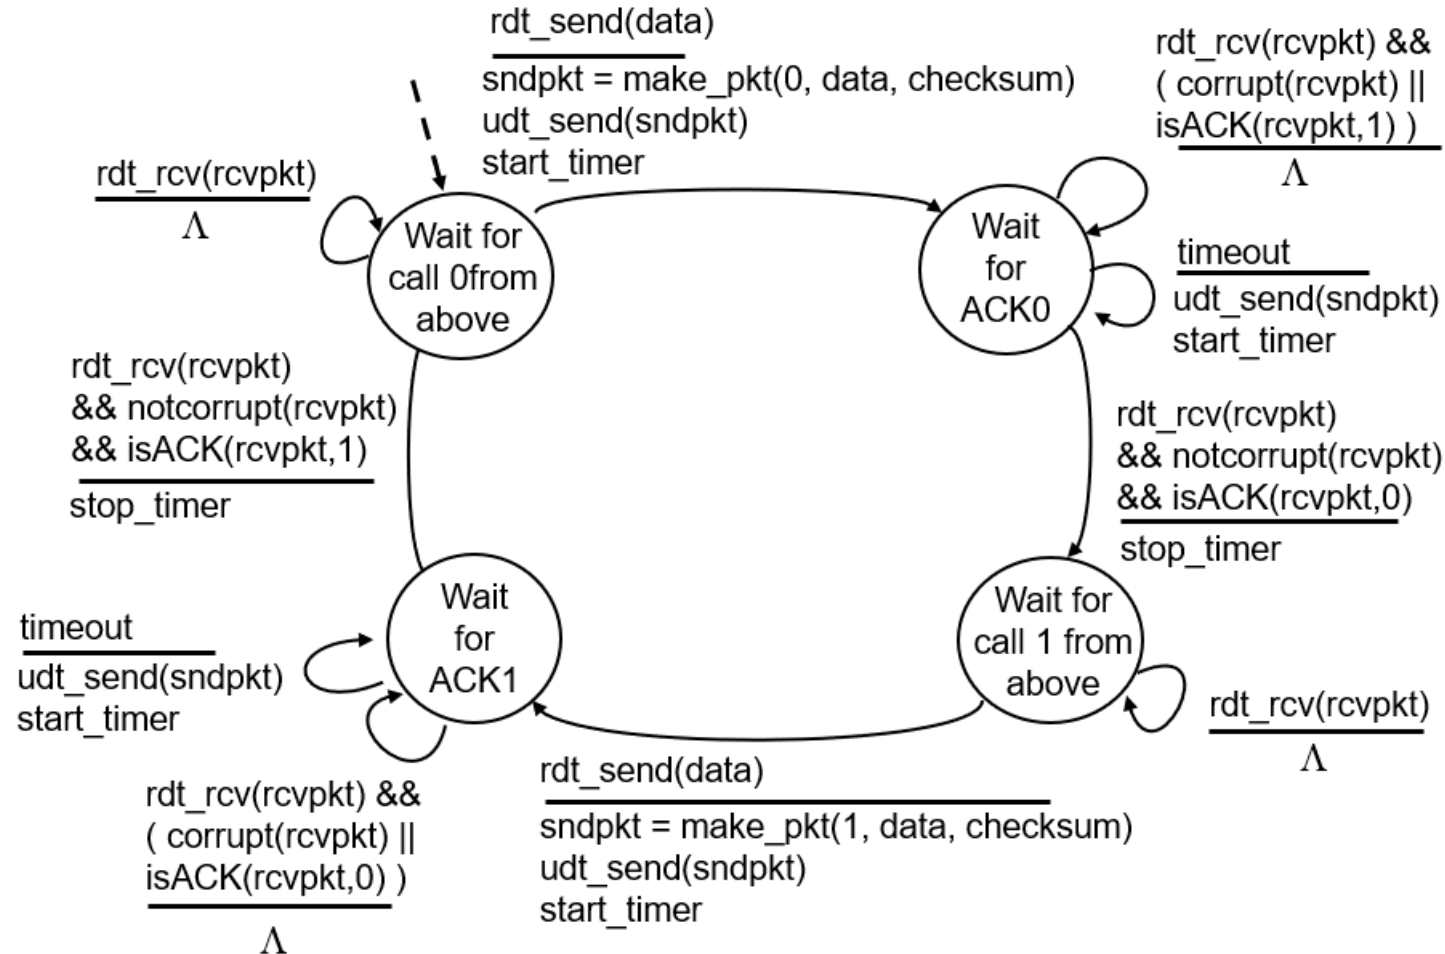
A timer!!!

Retransmit if no ACK received in X amount of time

What if the ACK is just taking a really long time to arrive? This will be duplicate data, but we have Seq # to handle that.

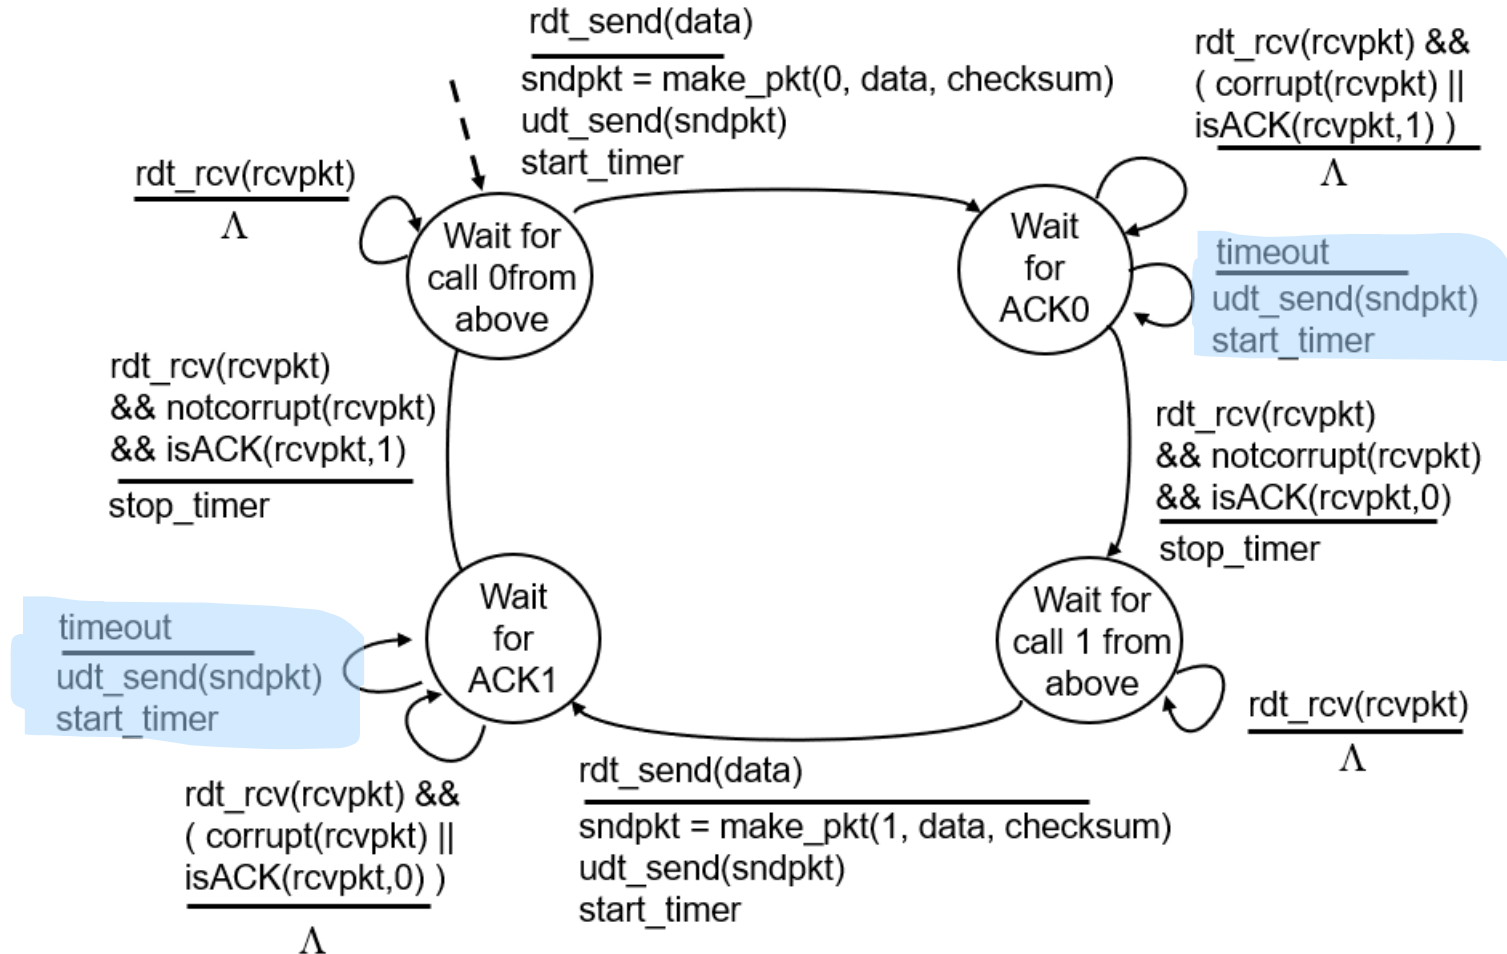
RDT 3.0

Packets can get lost or dropped



RDT 3.0

Packets can get lost or dropped

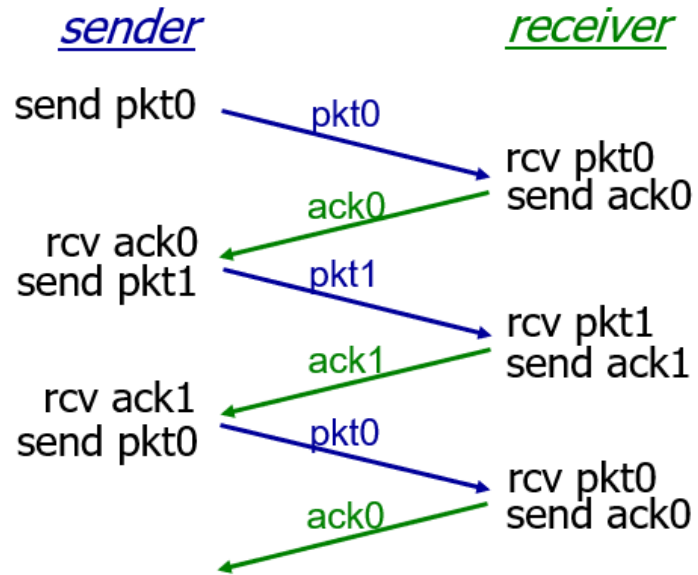


Transport Layer

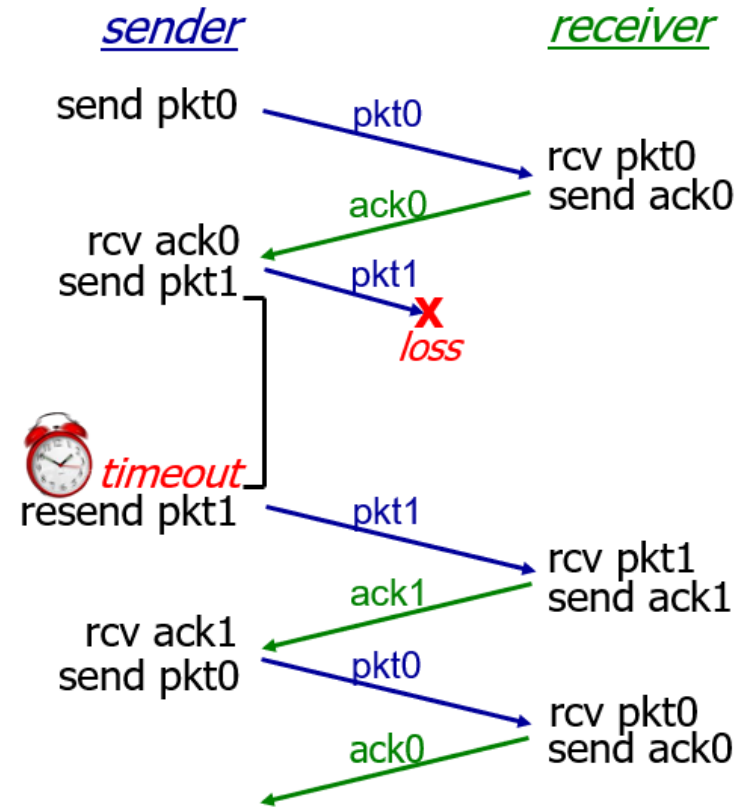
Reliable Data Transfer 3.0

RDT 3.0

Packets can get lost or dropped



(a) no loss




(b) packet loss

Transport Layer

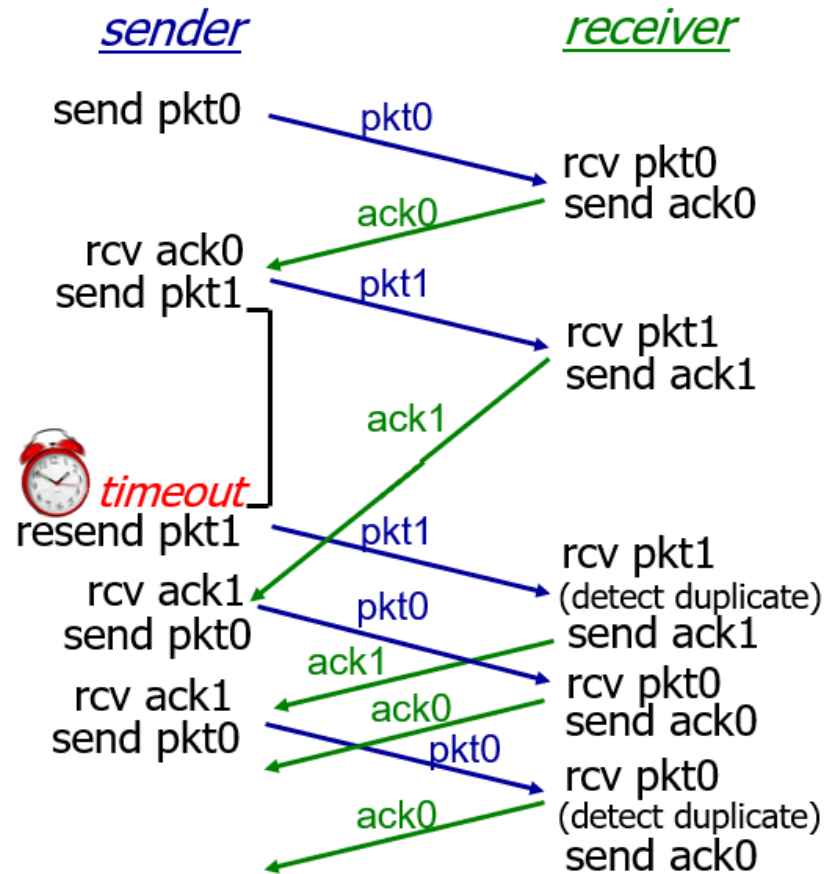
Reliable Data Transfer 3.0

RDT 3.0

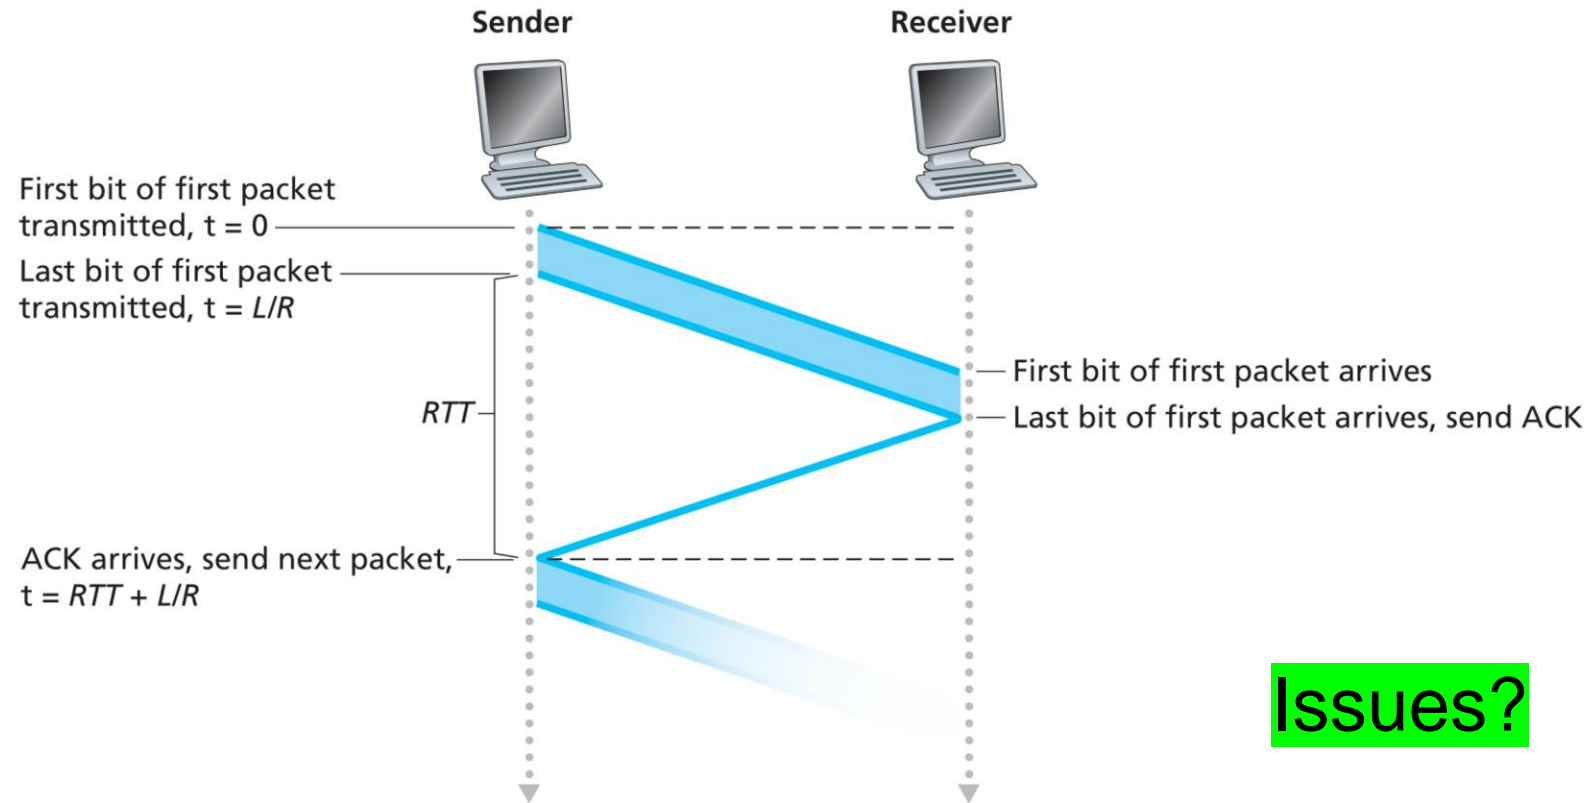
Packets can get lost or dropped



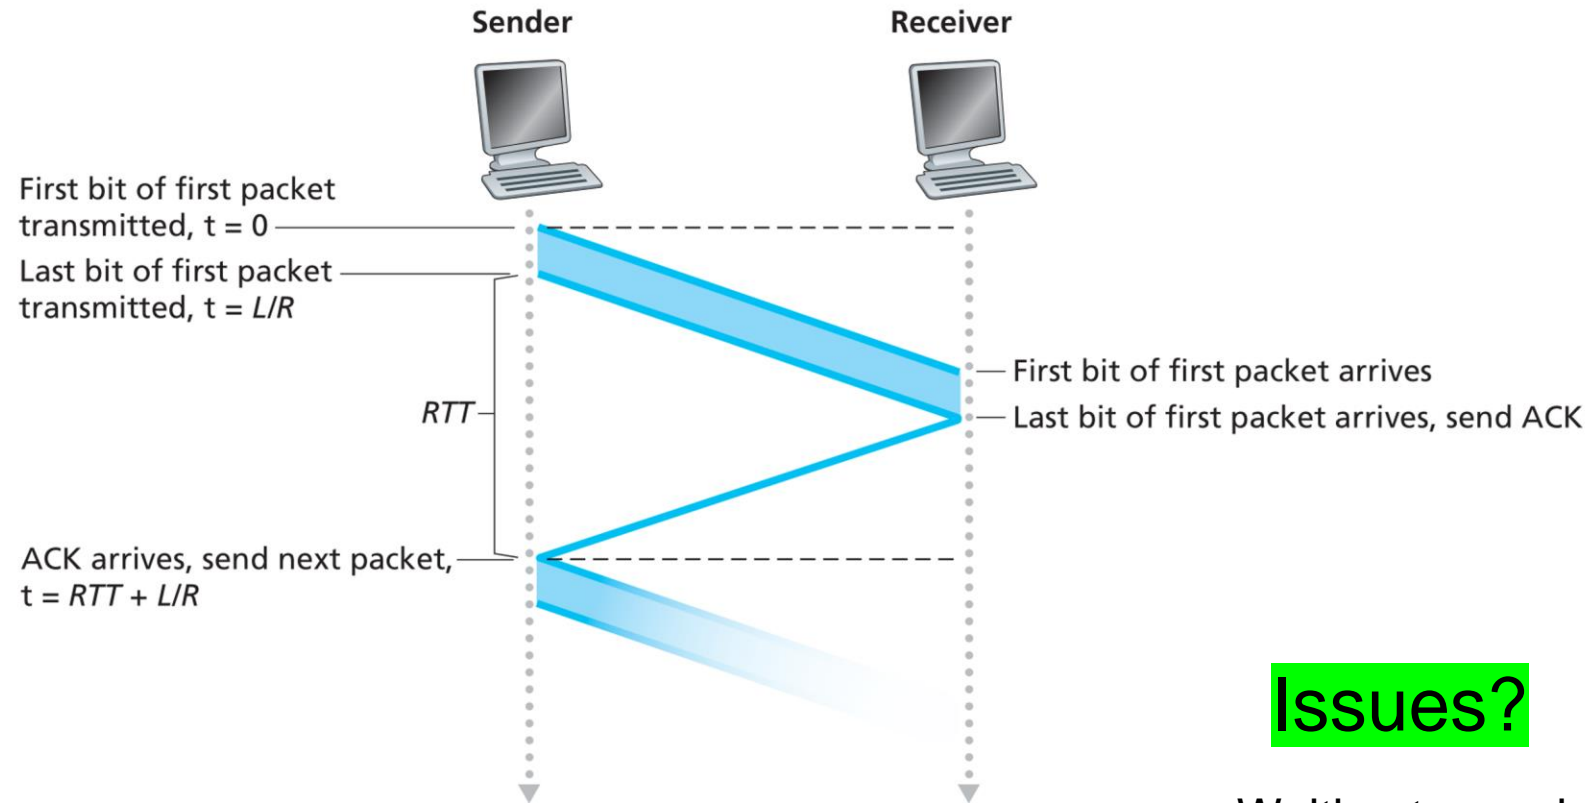
(c) ACK loss



(d) premature timeout/ delayed ACK




Issues?



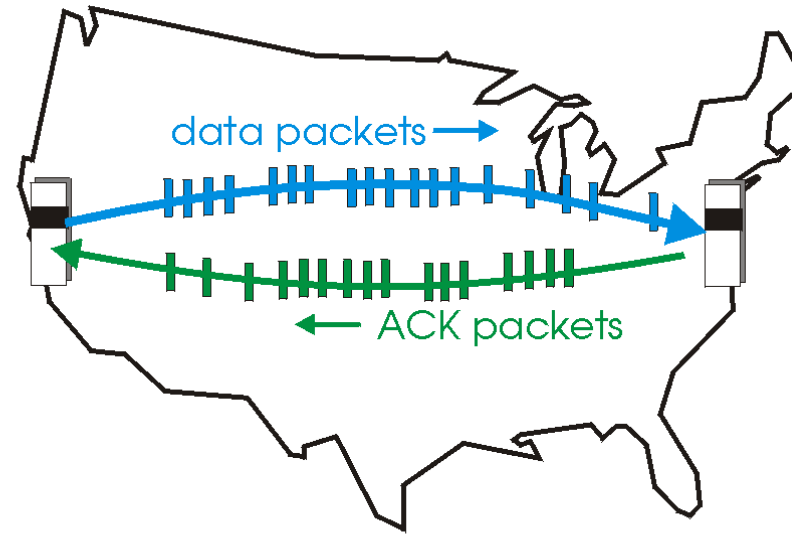
Issues?

Waiting to receive an ACK is inefficient

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledges pkts

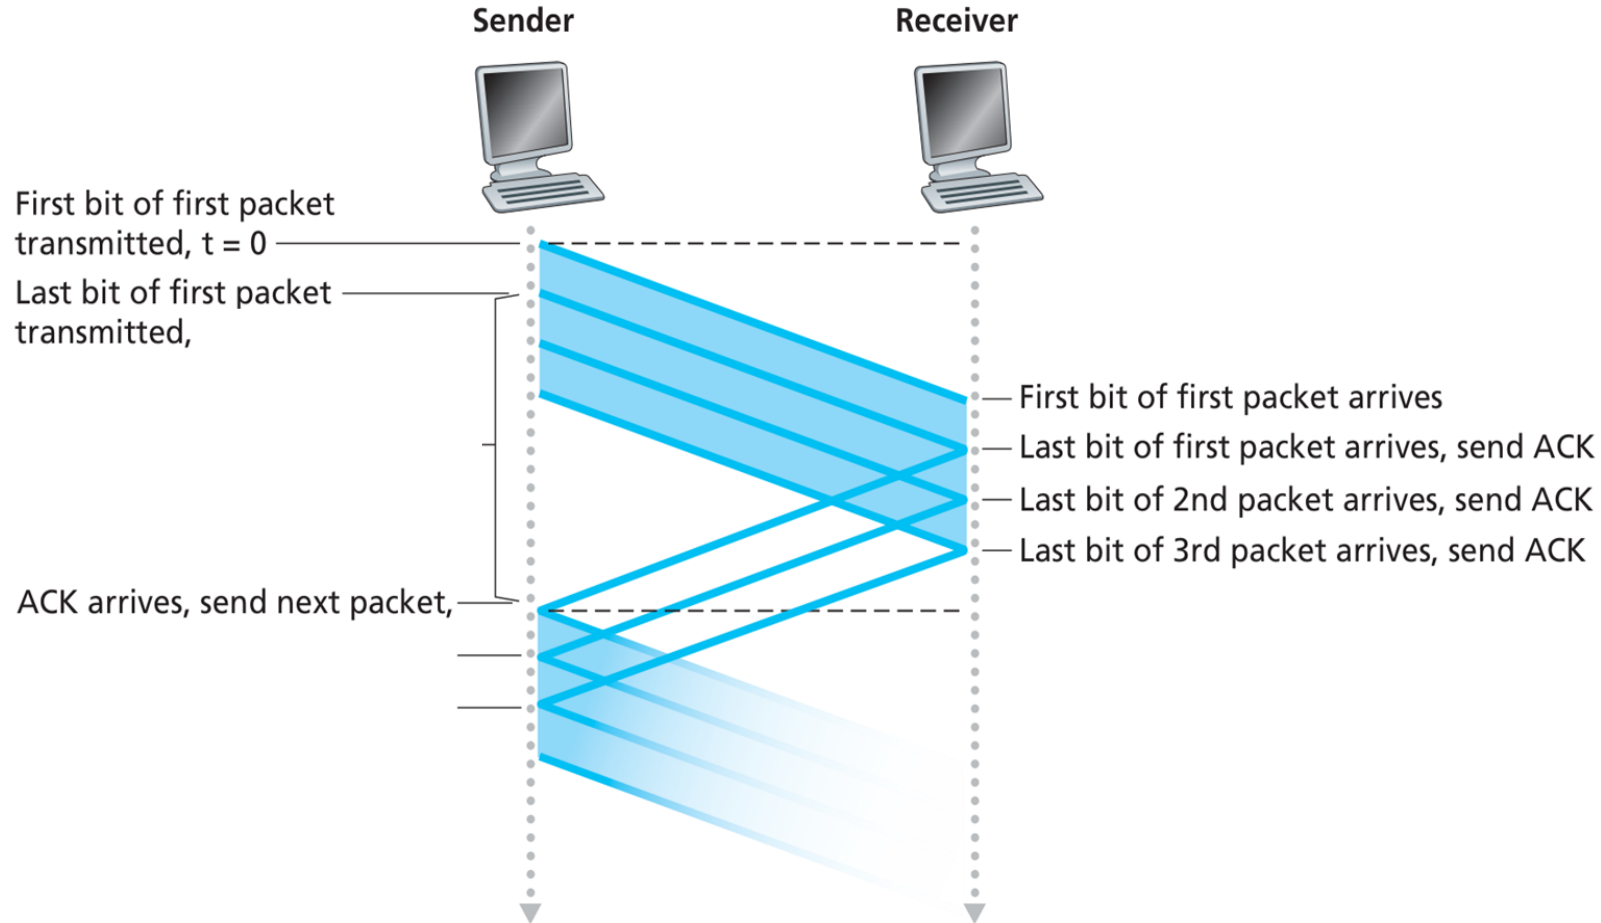


(a) a stop-and-wait protocol in operation




(b) a pipelined protocol in operation

Pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledges pkts



Go-back-N (GBN)

- Sender can have up to N unacked packets in pipeline
- Receiver only sends **cumulative** ack
 - ACK for last contiguous packet
 - No ACK for new packets past a sequence number gap
- Sender has timer for oldest unacked packet
 - When timer expires, retransmit all unacked packets



Transport Layer

GBN

sender window (N=4)

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8
0 1 2 3 4 5 6 7 8

sender

send pkt0
send pkt1
send pkt2
send pkt3
(wait)

rcv ack0, send pkt4
rcv ack1, send pkt5

ignore duplicate ACK



pkt 2 timeout

send pkt2
send pkt3
send pkt4
send pkt5

receiver

receive pkt0, send ack0
receive pkt1, send ack1

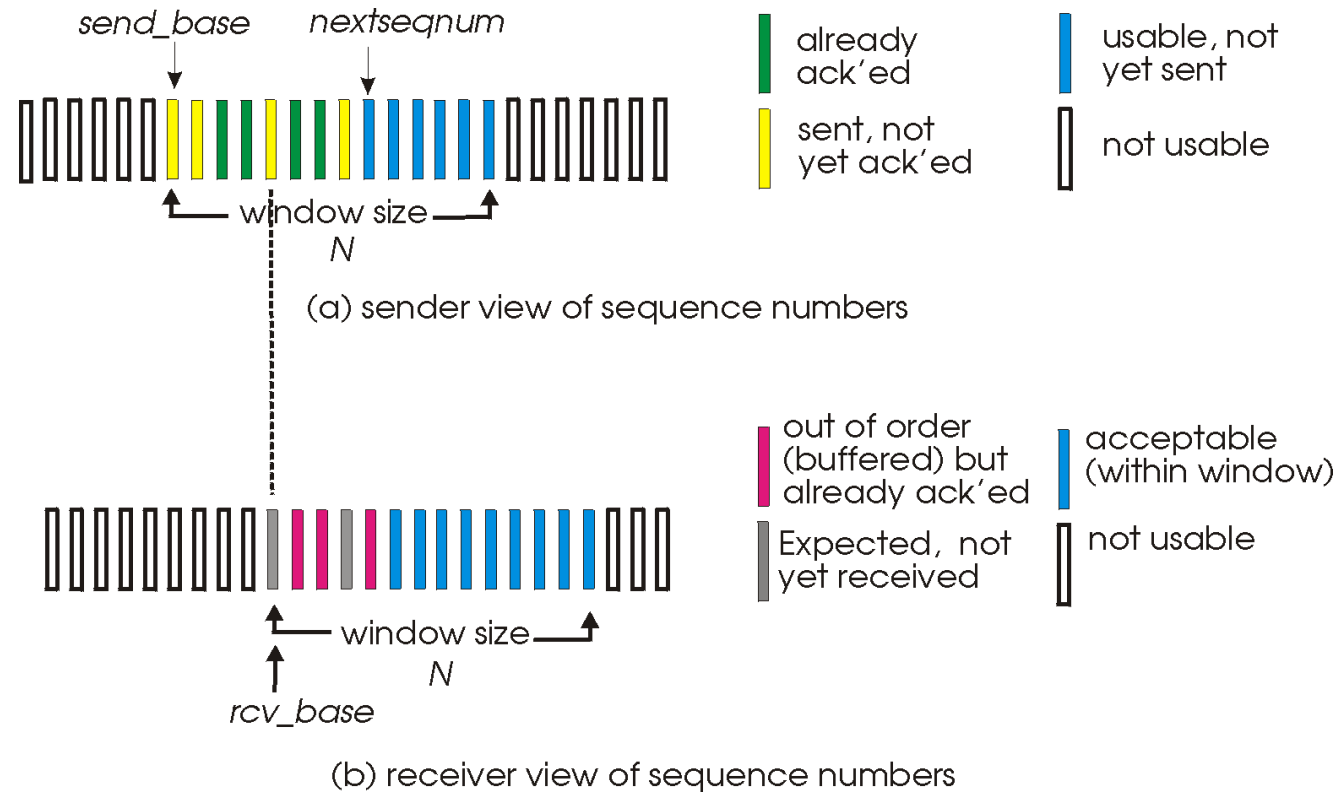
receive pkt3, discard,
(re)send ack1

receive pkt4, discard,
(re)send ack1

receive pkt5, discard,
(re)send ack1

rcv pkt2, deliver, send ack2
rcv pkt3, deliver, send ack3
rcv pkt4, deliver, send ack4
rcv pkt5, deliver, send ack5

- receiver *individually* acknowledges all correctly received pkts
 - buffers pkts, as needed, for eventual in-order delivery to upper layer
- sender only resends pkts for which ACK not received
 - sender timer for each unACKed pkt
- sender window
 - N consecutive seq #'s
 - limits seq #'s of sent, unACKed pkts



Checksum- Used to detect bit errors in transmitted pacets

Checksum- Used to detect bit errors in transmitted packets

Timer- Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

Checksum- Used to detect bit errors in transmitted packets

Timer- Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

Sequence Number- Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

Checksum- Used to detect bit errors in transmitted packets

Timer- Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

Sequence Number- Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

Acknowledgement- Used by the receiver to tell the sender that a packet or set of packets has been received correctly. ACKs will typically carry the sequence # of the packet being acknowledged

Checksum- Used to detect bit errors in transmitted packets

Timer- Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

Sequence Number- Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

Acknowledgement- Used by the receiver to tell the sender that a packet or set of packets has been received correctly. ACKs will typically carry the sequence # of the packet being acknowledged

Negative Acknowledgement- Used by the receiver to tell the sender that a packet has not been received correctly. Negative acknowledgements will typically carry the sequence number of the packet that was not received correctly

Checksum- Used to detect bit errors in transmitted packets

Timer- Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

Sequence Number- Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

Acknowledgement- Used by the receiver to tell the sender that a packet or set of packets has been received correctly. ACKs will typically carry the sequence # of the packet being acknowledged


Negative Acknowledgement- Used by the receiver to tell the sender that a packet has not been received correctly. Negative acknowledgements will typically carry the sequence number of the packet that was not received correctly

Window, pipelining- The sender may be restricted to sending only packets with sequence numbers that fall within a given range. By allowing multiple packets to be transmitted but not yet acknowledged, sender utilization can be increased over a stop-and-wait mode of operation.

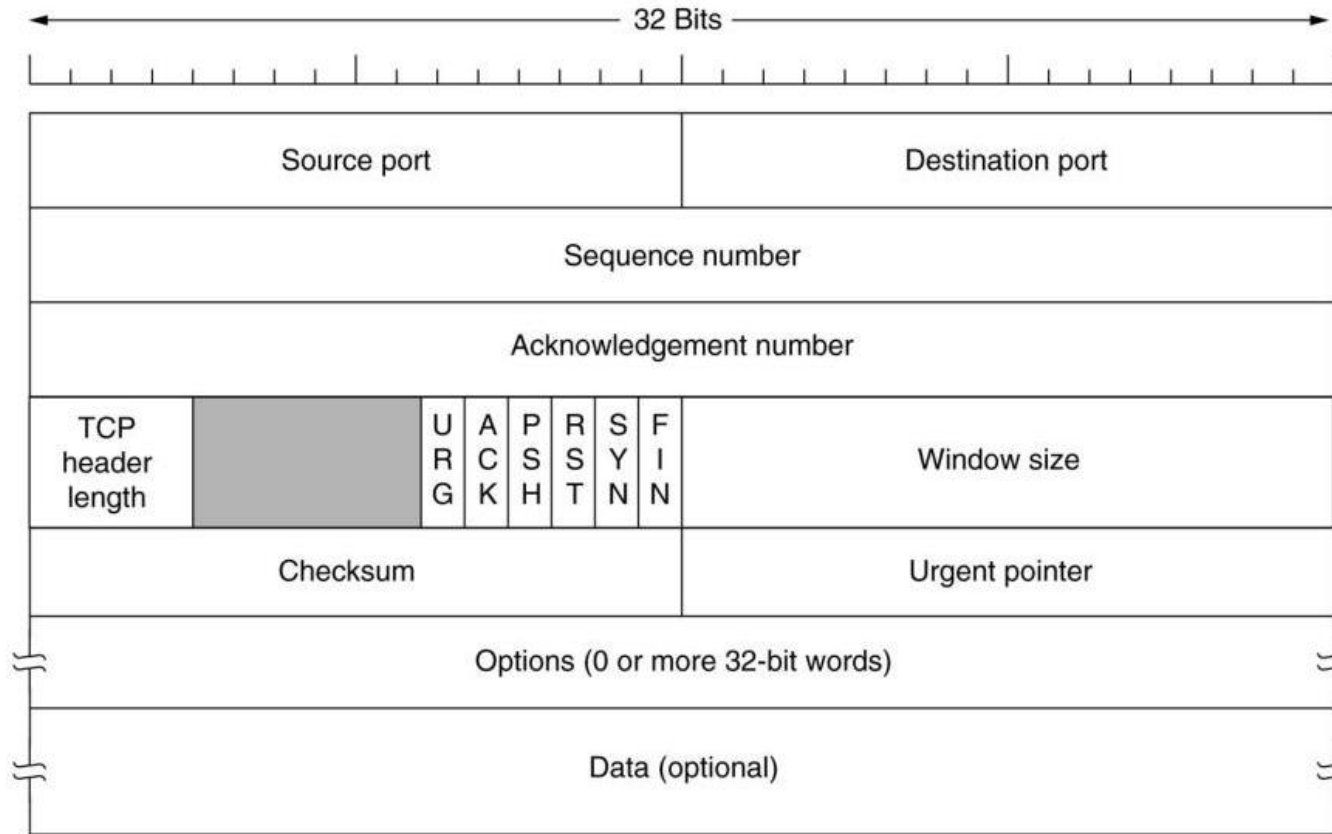
Transport Layer

TCP (Transmission Control Protocol)

- **point-to-point:**
 - one sender, one receiver
- **reliable, in-order *byte stream*:**
- **pipelined:**
 - TCP congestion and flow control set window size



- **full duplex data:**
 - bi-directional data flow in same connection
- **connection-oriented:**
 - handshaking (exchange of control msgs) in its sender, receiver state before data exchange
- **flow controlled:**
 - sender will not overwhelm receiver

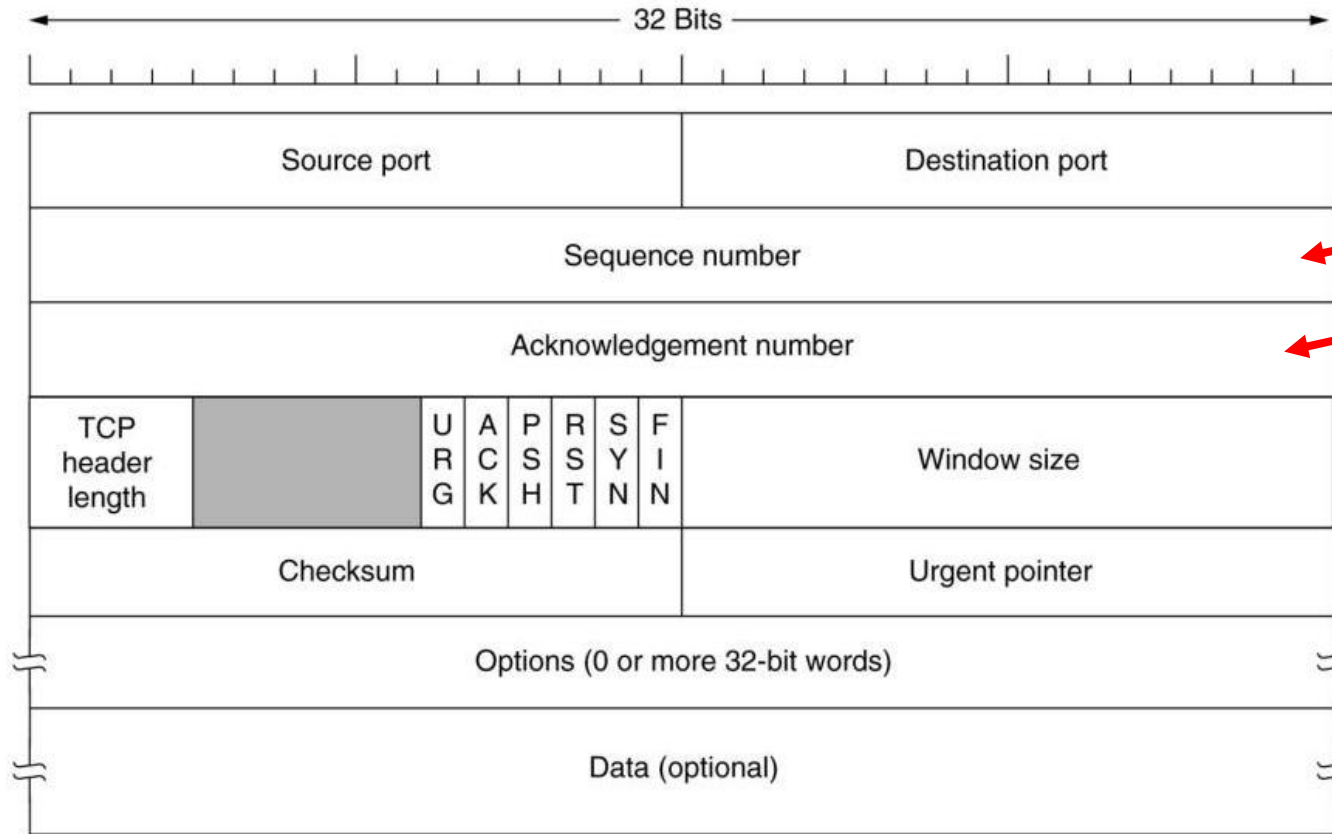


TCP header length in 32 bit words,

URG-urgent, ACK- ack number is valid, PSH-push, RST-reset connection,

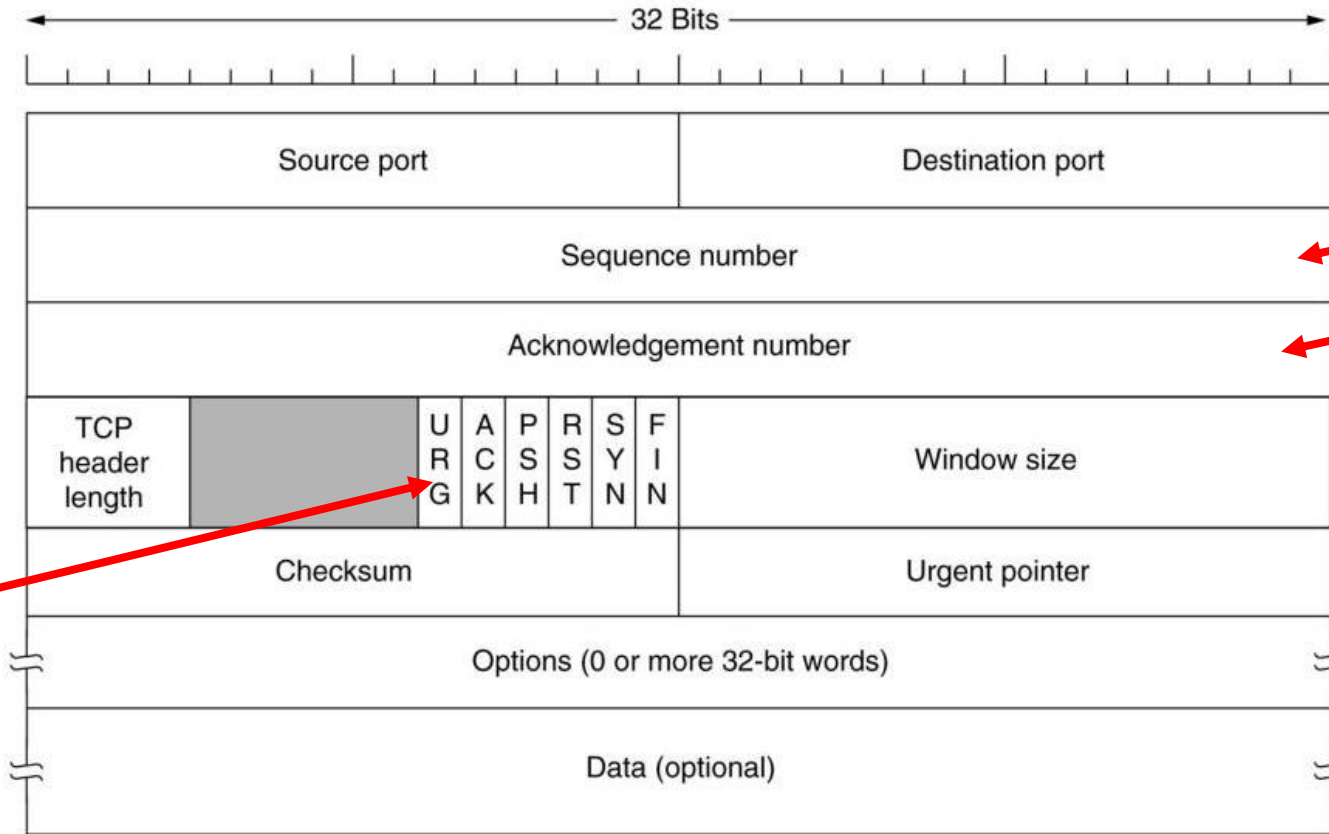
SYN-used to establish connection, FIN-used to release connection

TCP Segment Structure



Counts by bytes, not segment

TCP header length in 32 bit words,
URG-urgent, ACK- ack number is valid, PSH-push, RST-reset connection,
SYN-used to establish connection, FIN-used to release connection



Counts by bytes, not segment

TCP header length in 32 bit words,
URG-urgent, ACK- ack number is valid, PSH-push, RST-reset connection,
SYN-used to establish connection, FIN-used to release connection

sequence numbers:

- byte stream “number” of first byte in segment’s data

acknowledgements:


- seq # of next byte expected from other side
- cumulative ACK

Q: how receiver handles out-of-order segments

- **A:** TCP spec doesn’t say, - up to implementor

outgoing segment from sender

source port #	dest port #
sequence number	
acknowledgement number	
	rwnd
checksum	urg pointer



incoming segment to sender

source port #	dest port #
sequence number	
acknowledgement number	
	A
checksum	urg pointer

No class on Wednesday 10/5 (go to the career fair)

Fill out the short survey on discord for extra credit

You vs the guy she tells you not to worry about.

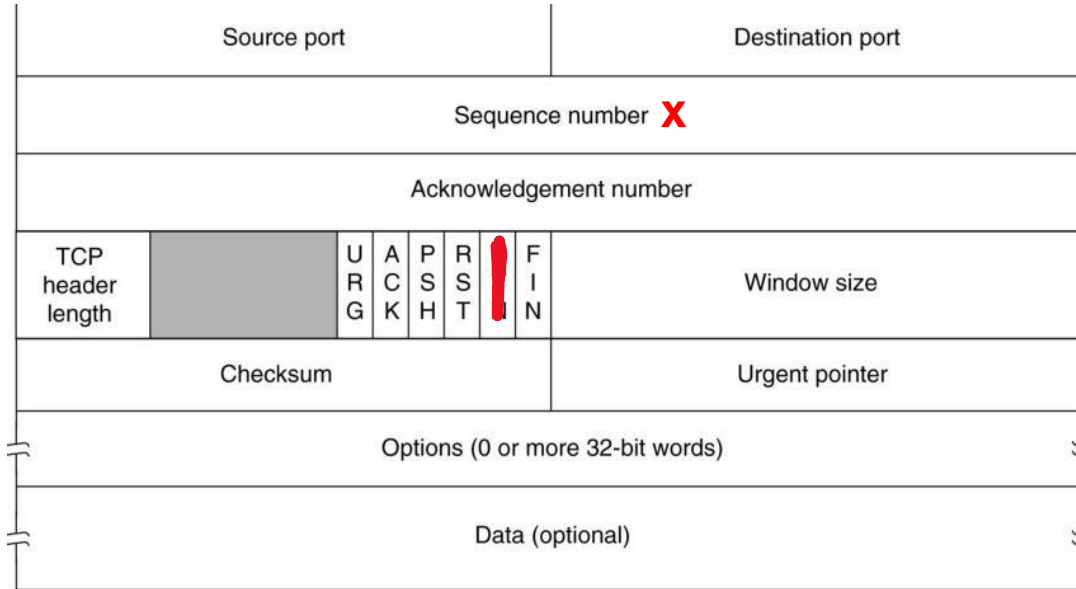
UDP	TCP
Unreliable	Reliable
Connectionless	Connection-oriented
No windowing or retransmission	Segment retransmission and flow control through windowing
No sequencing	Segment sequencing
No acknowledgement	Acknowledge segments

TCP Segment Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags		Window Size			
128	Header and Data Checksum				Urgent Pointer			
160...	Options							

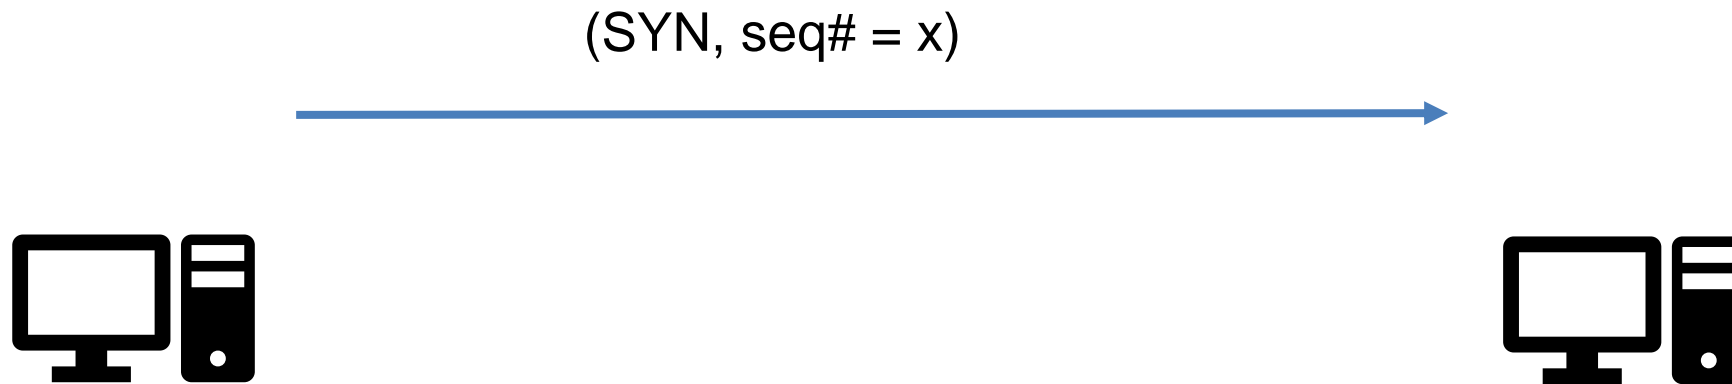
UDP Datagram Header Format

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Length				Header and Data Checksum			



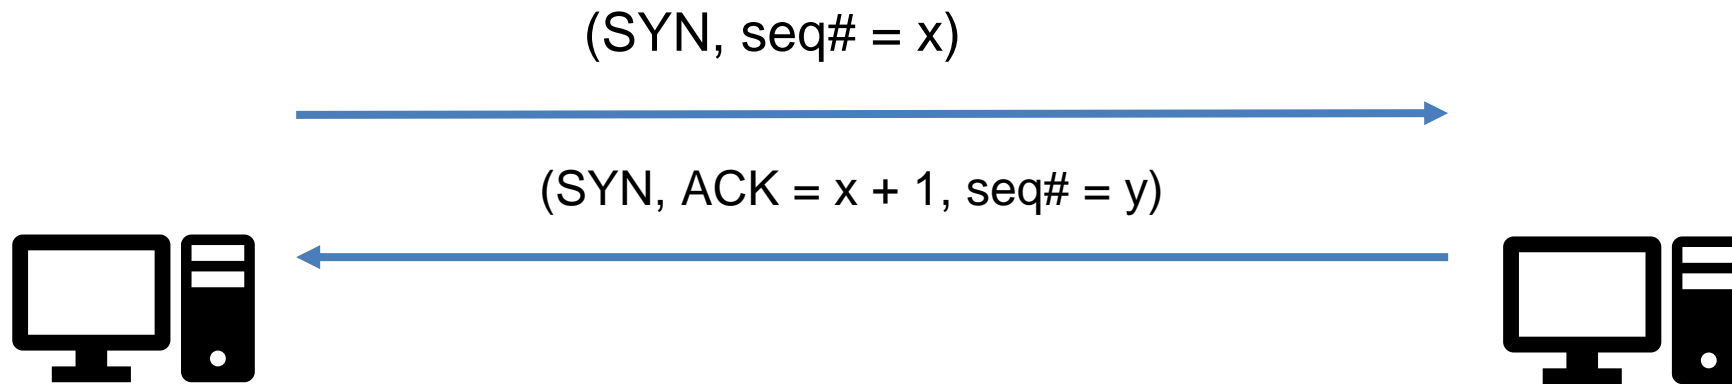
- When establishing the connection, enable the **SYN** flag (set to 1)
- Set an initial sequence number





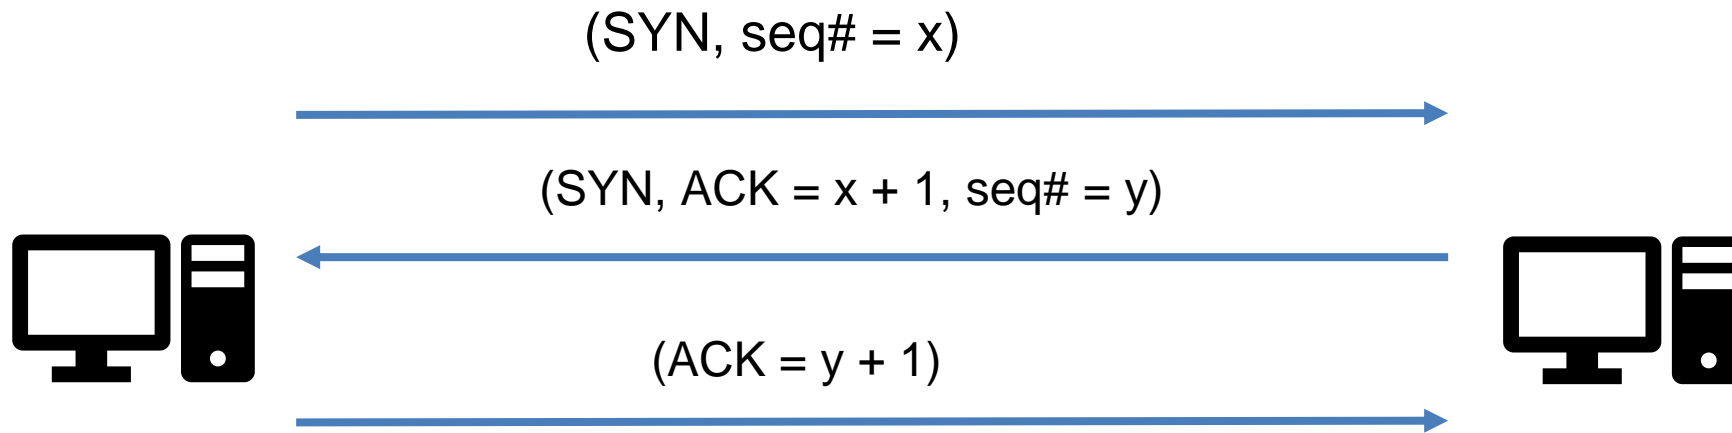
- When establishing the connection, enable the **SYN** flag (set to 1)
- Set an initial sequence number





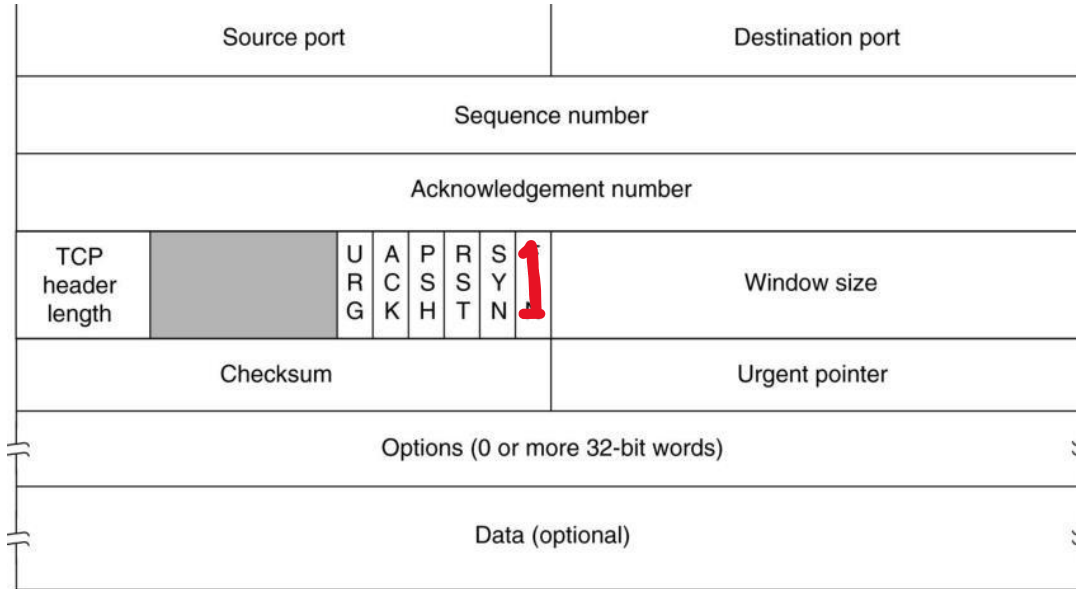
- Acknowledge message and increment sequence number






- Acknowledge the acknowledgement

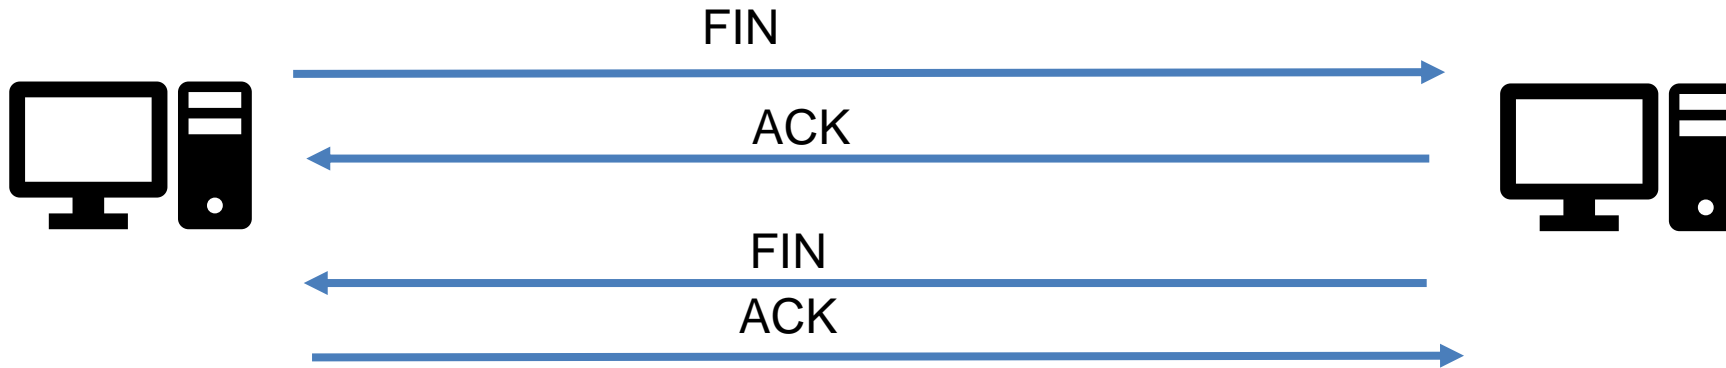




- The end communication, set the FIN flag




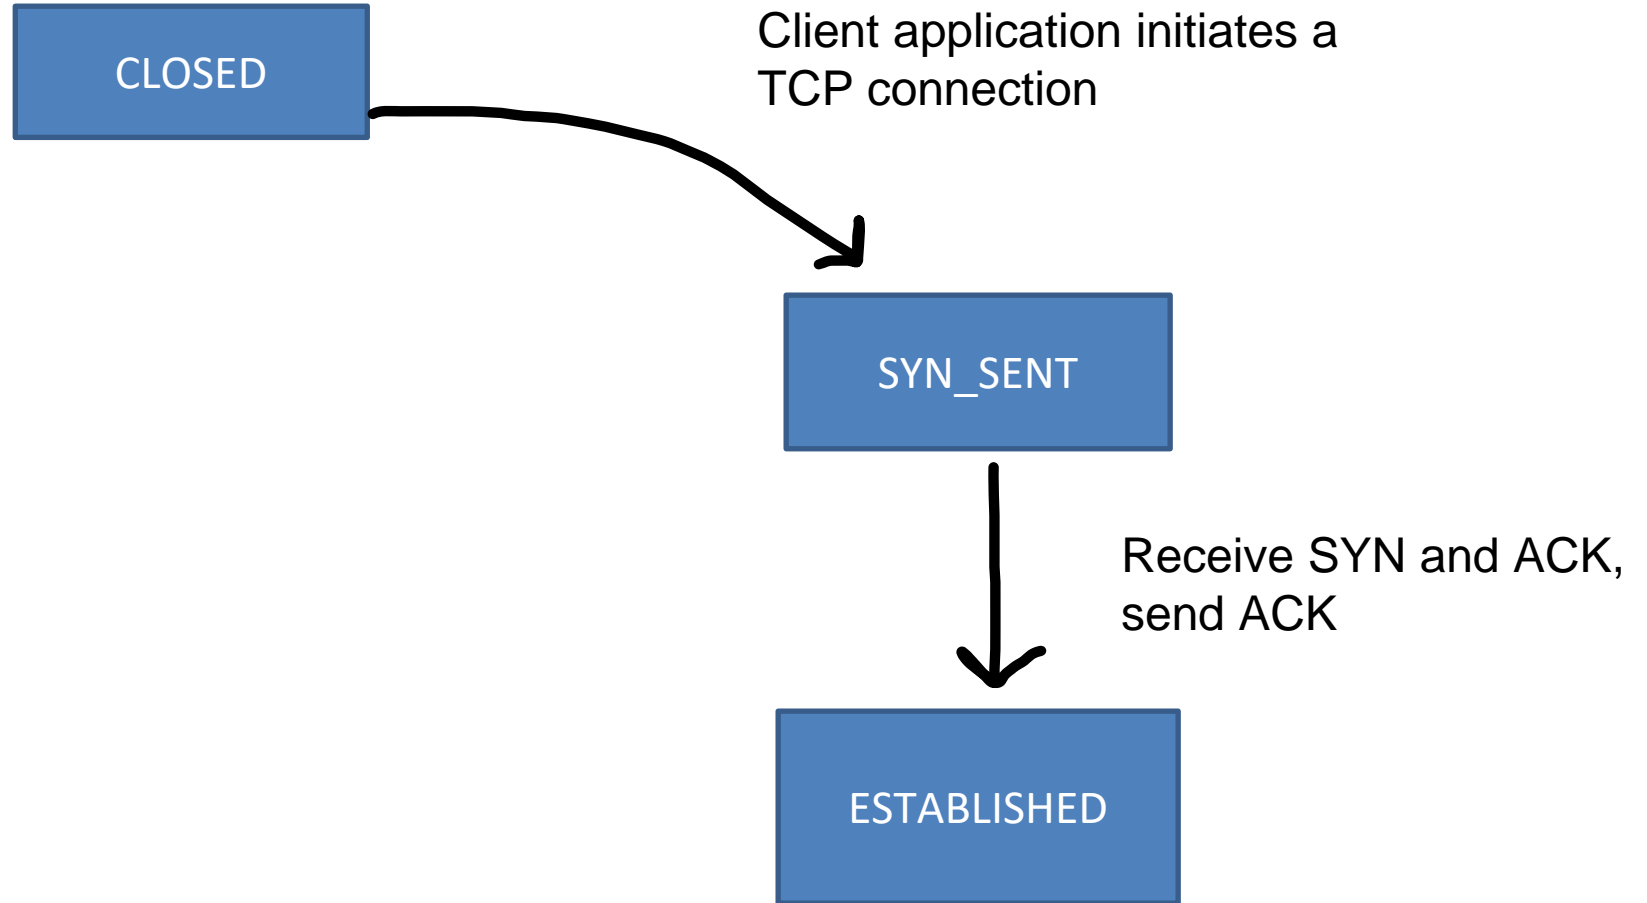


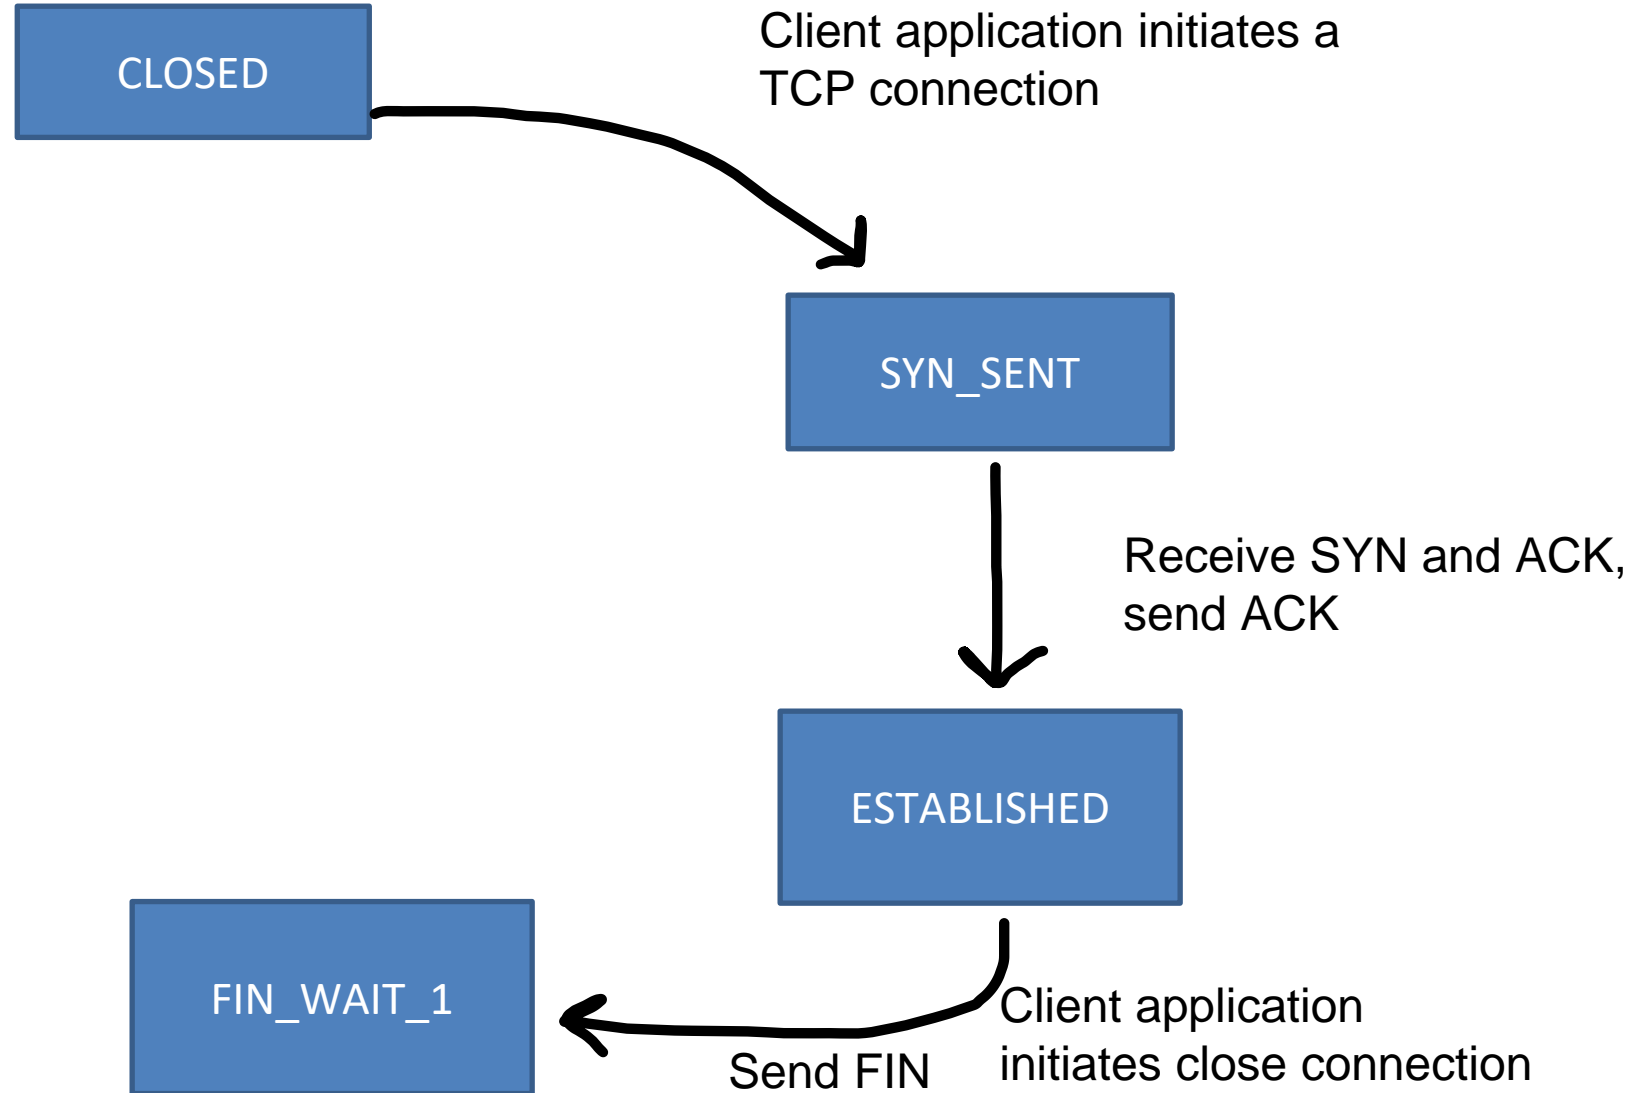


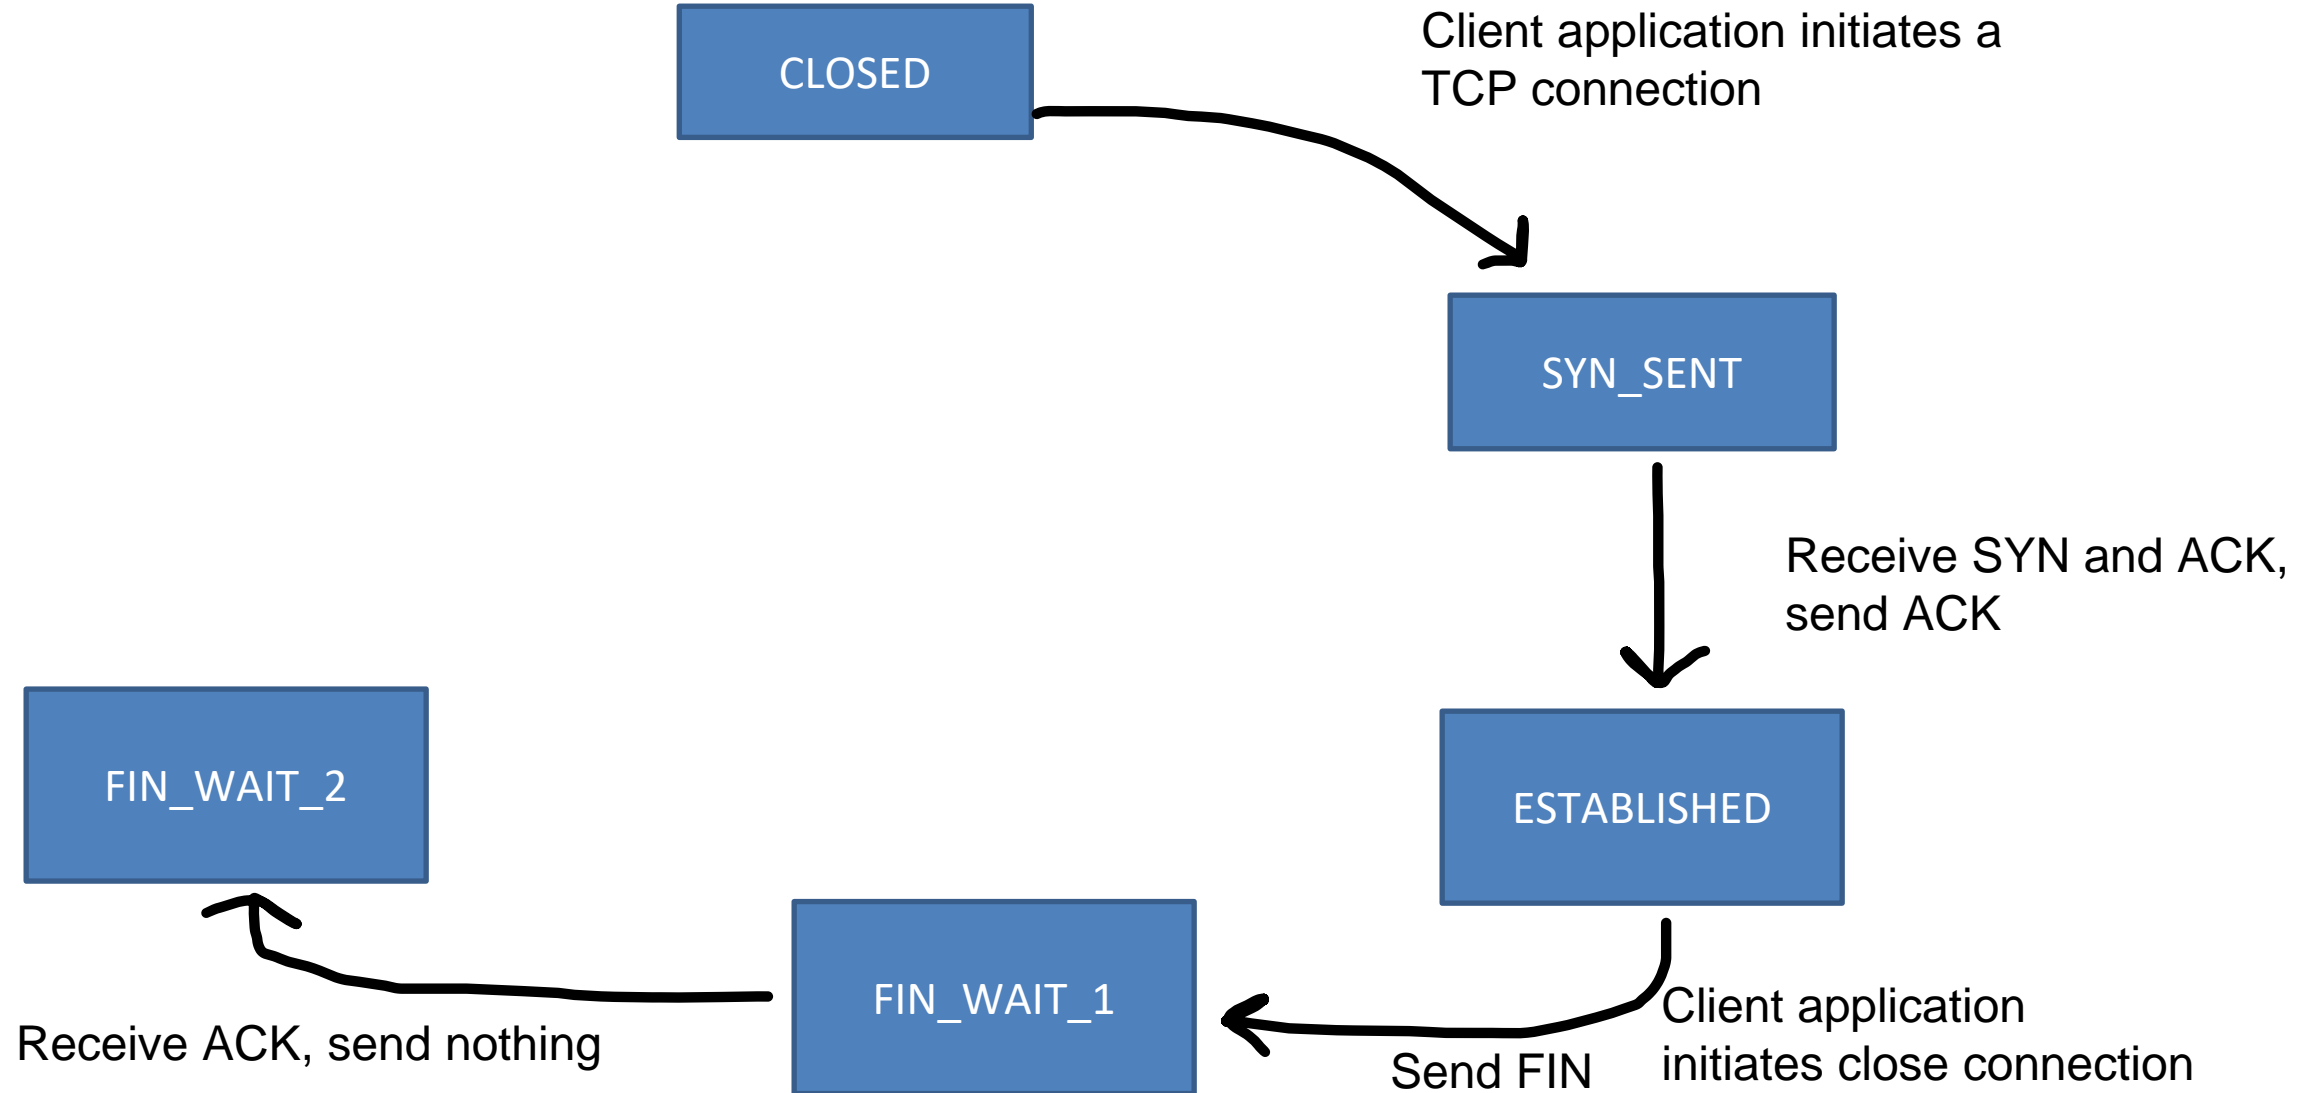
CLOSED

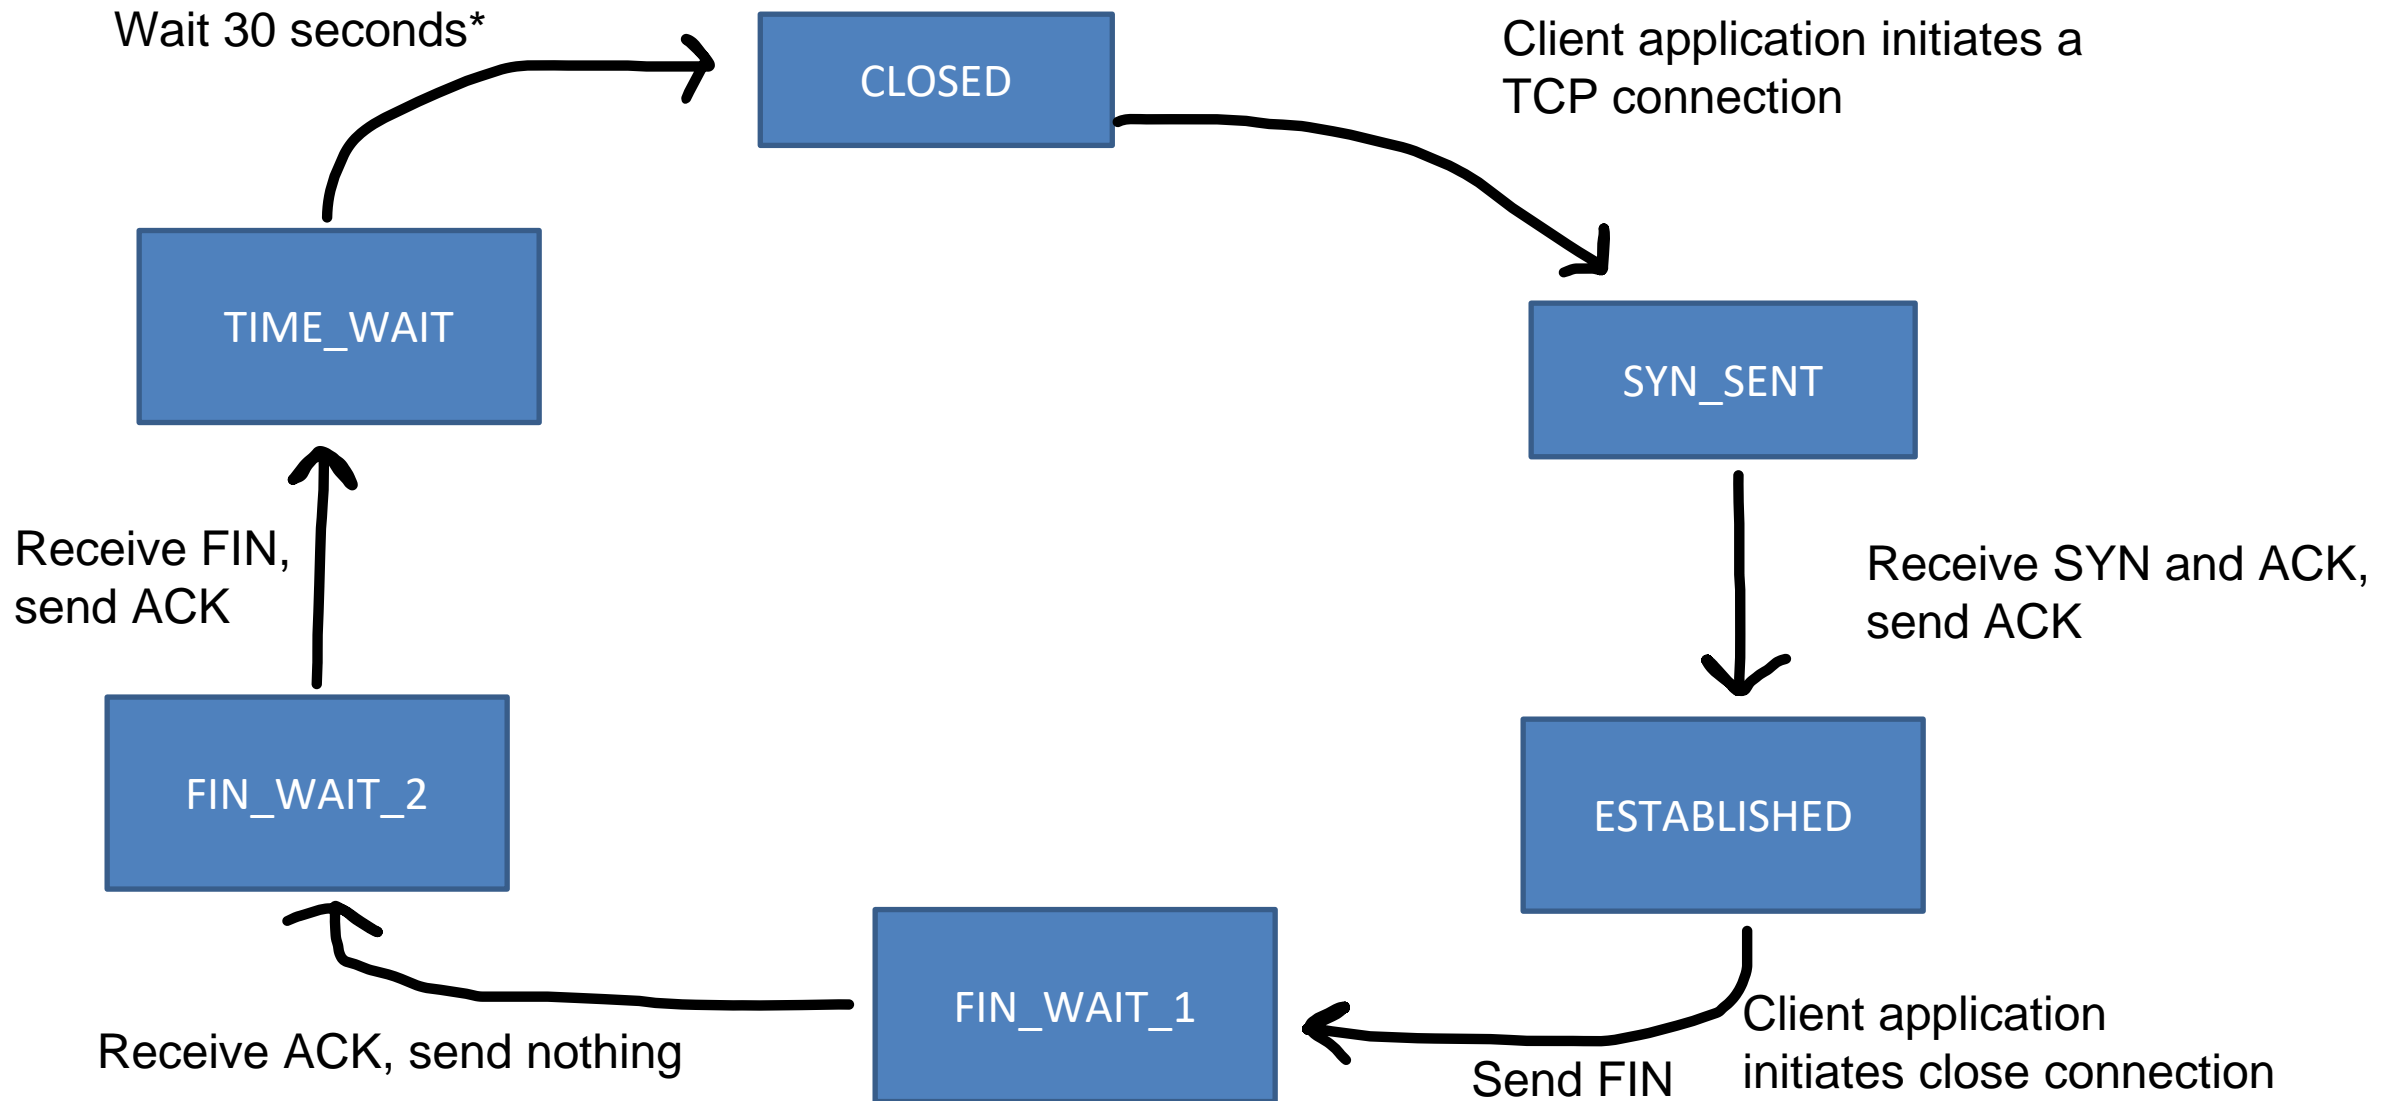
During the life of a TCP connection, the TCP protocol makes transitions through various **TCP states**













What if we receive a packet that has an invalid port number?

TCP Packet → Send a TCP segment with RST flag on

UDP Packet → Send an **ICMP** datagram (network layer thing)




Application layer messages are split into smaller chunks called **segments**



The size of these segments is determined by the **maximum segment size (MSS)**

Transport Layer

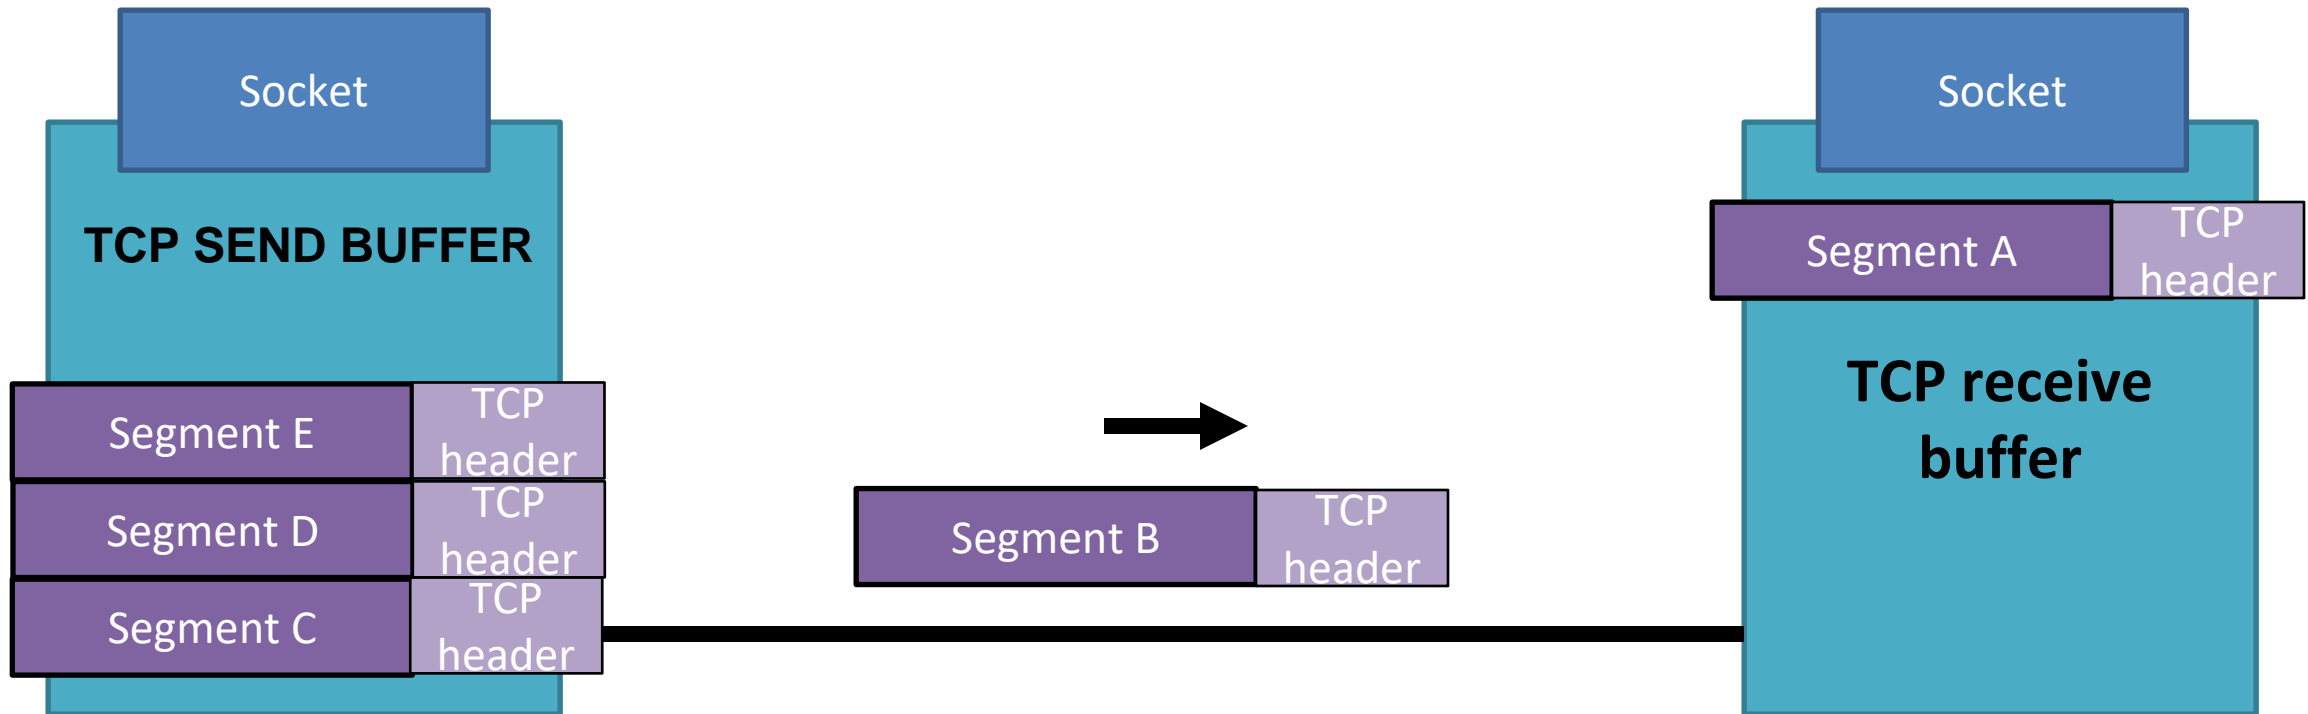
TCP Flow Control



Applications read streams of data from a **TCP buffer**

Transport Layer

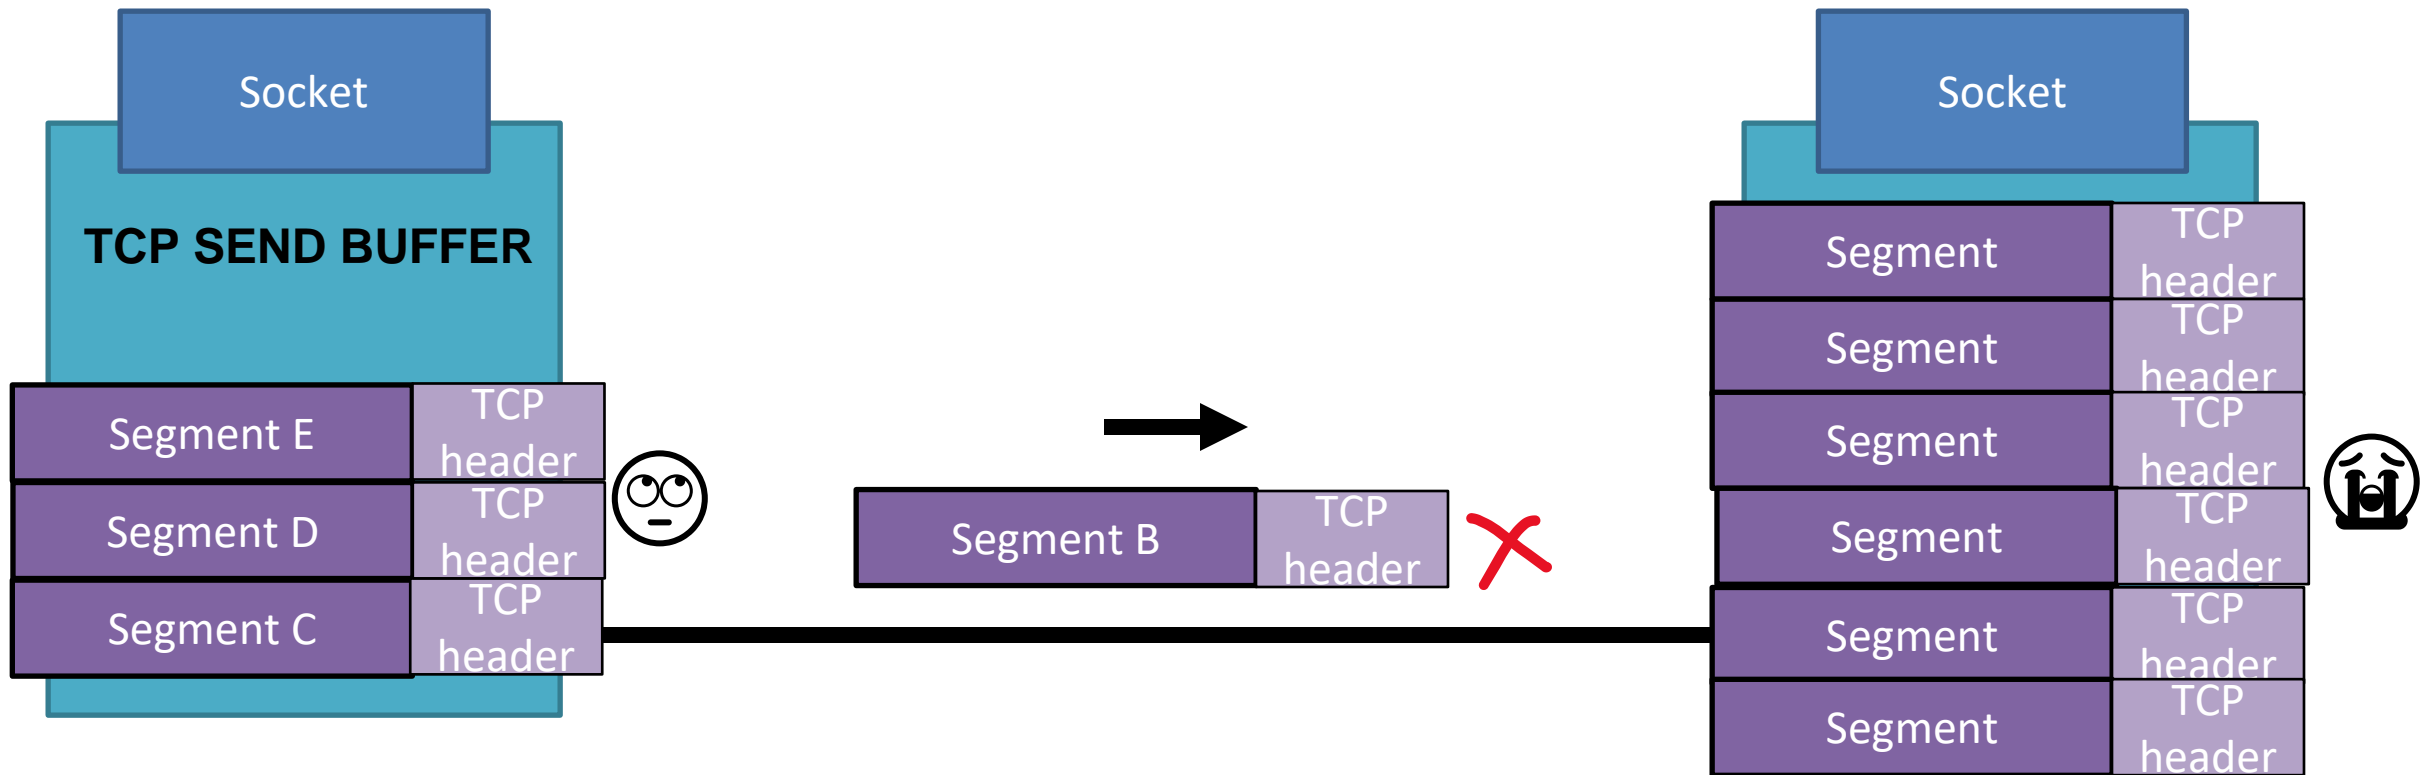
TCP Flow Control



Applications read streams of data from a **TCP buffer**

Transport Layer

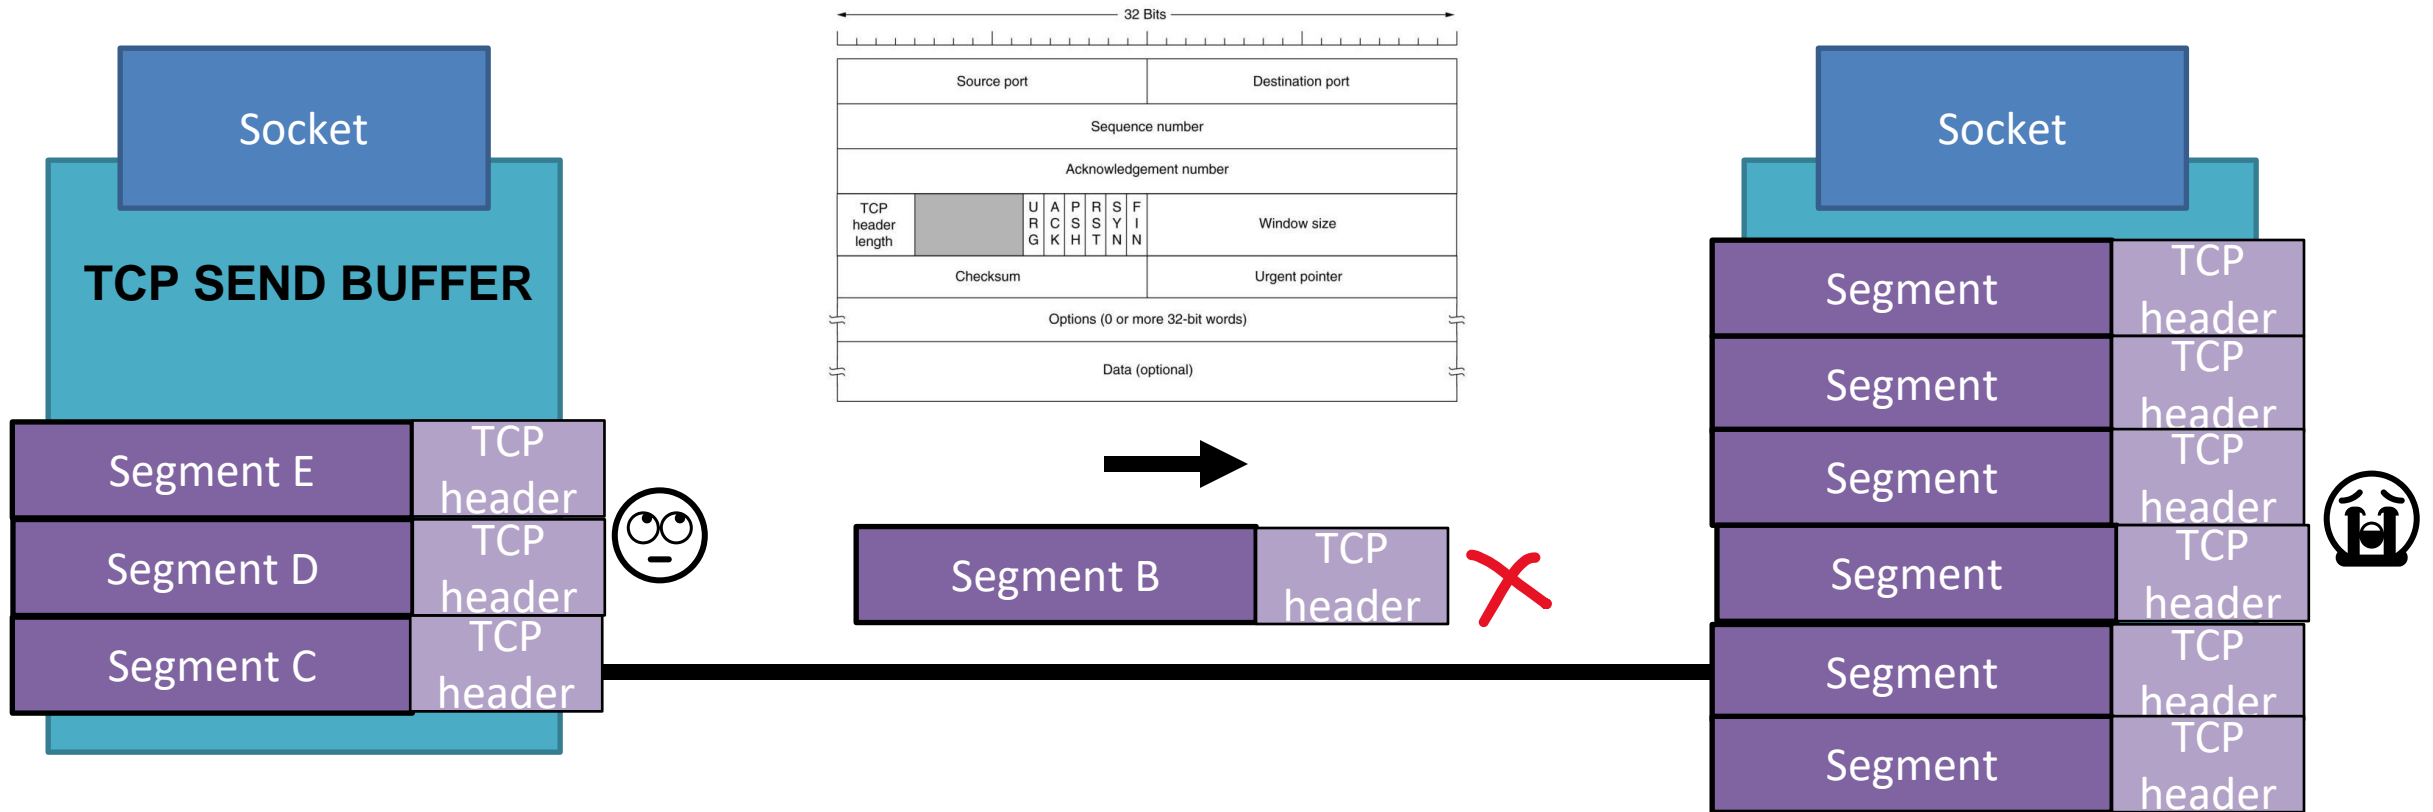
TCP Flow Control



How could we prevent something like this from happening?

Transport Layer


TCP Flow Control



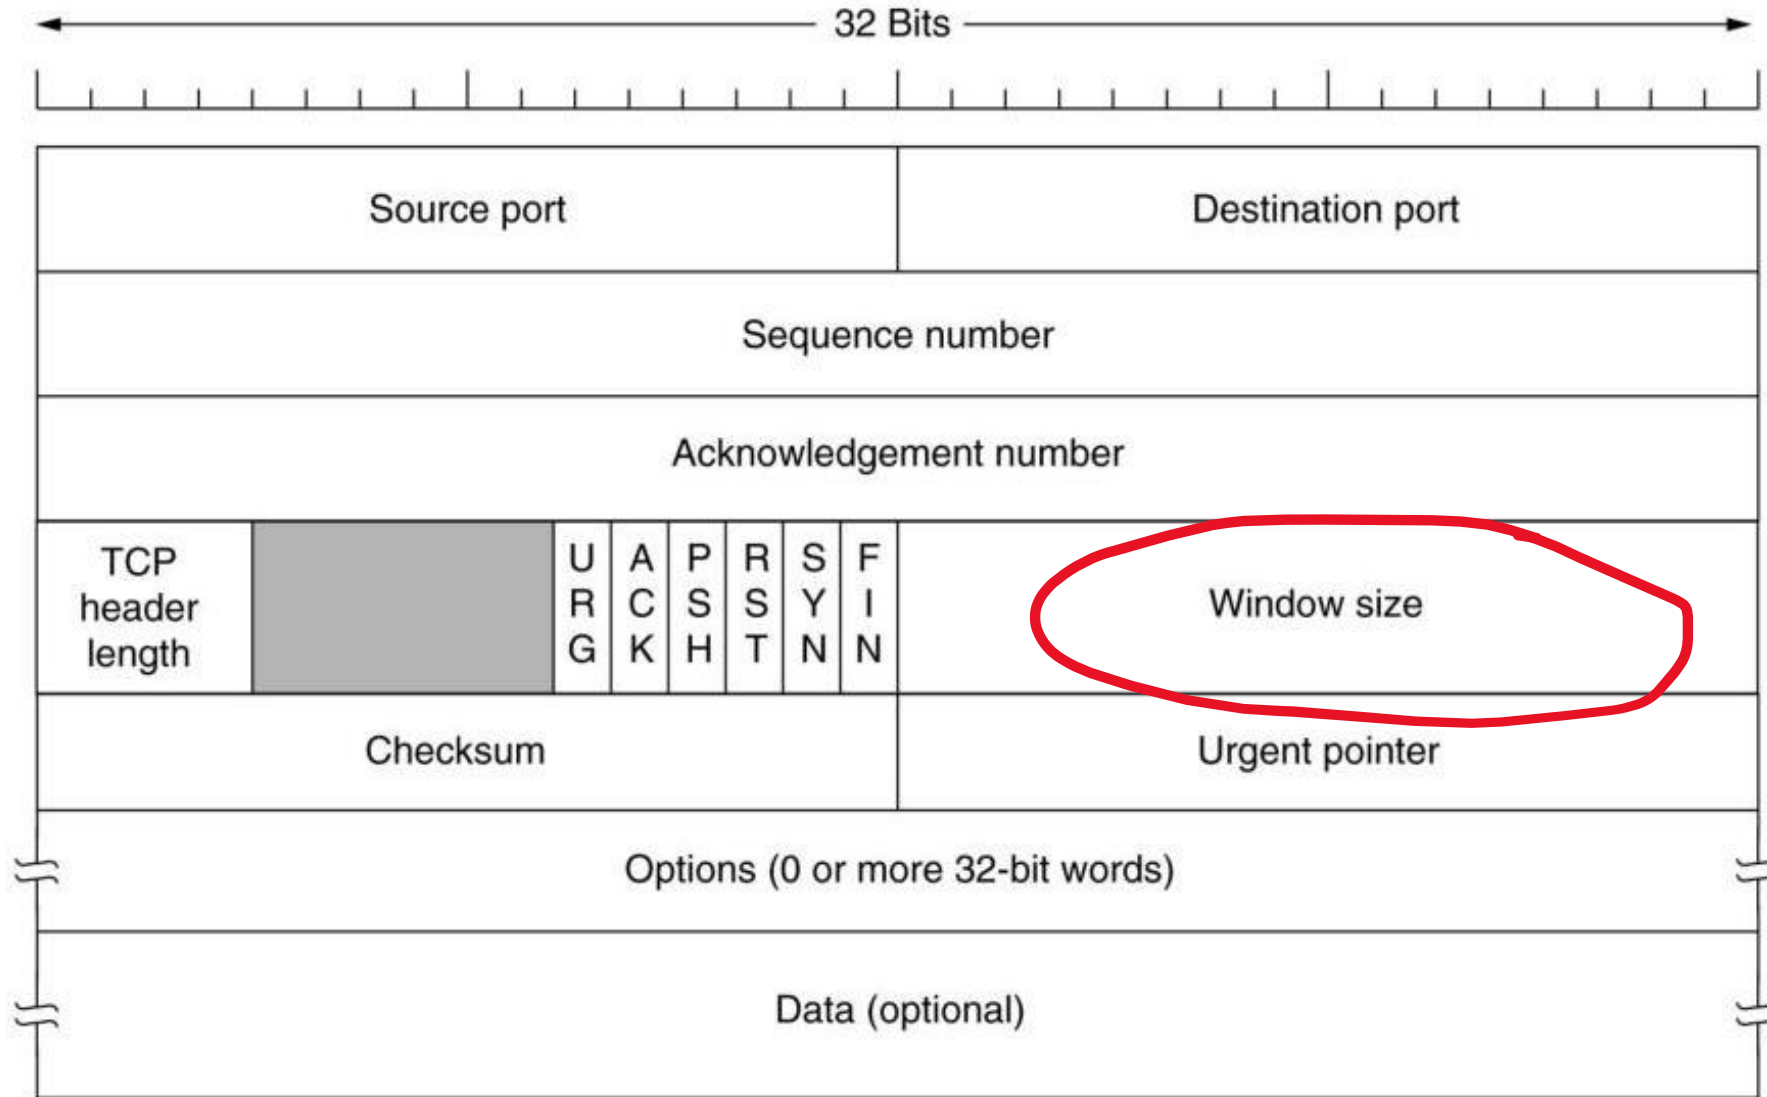
How could we prevent something like this from happening?

Transport Layer

TCP Flow Control

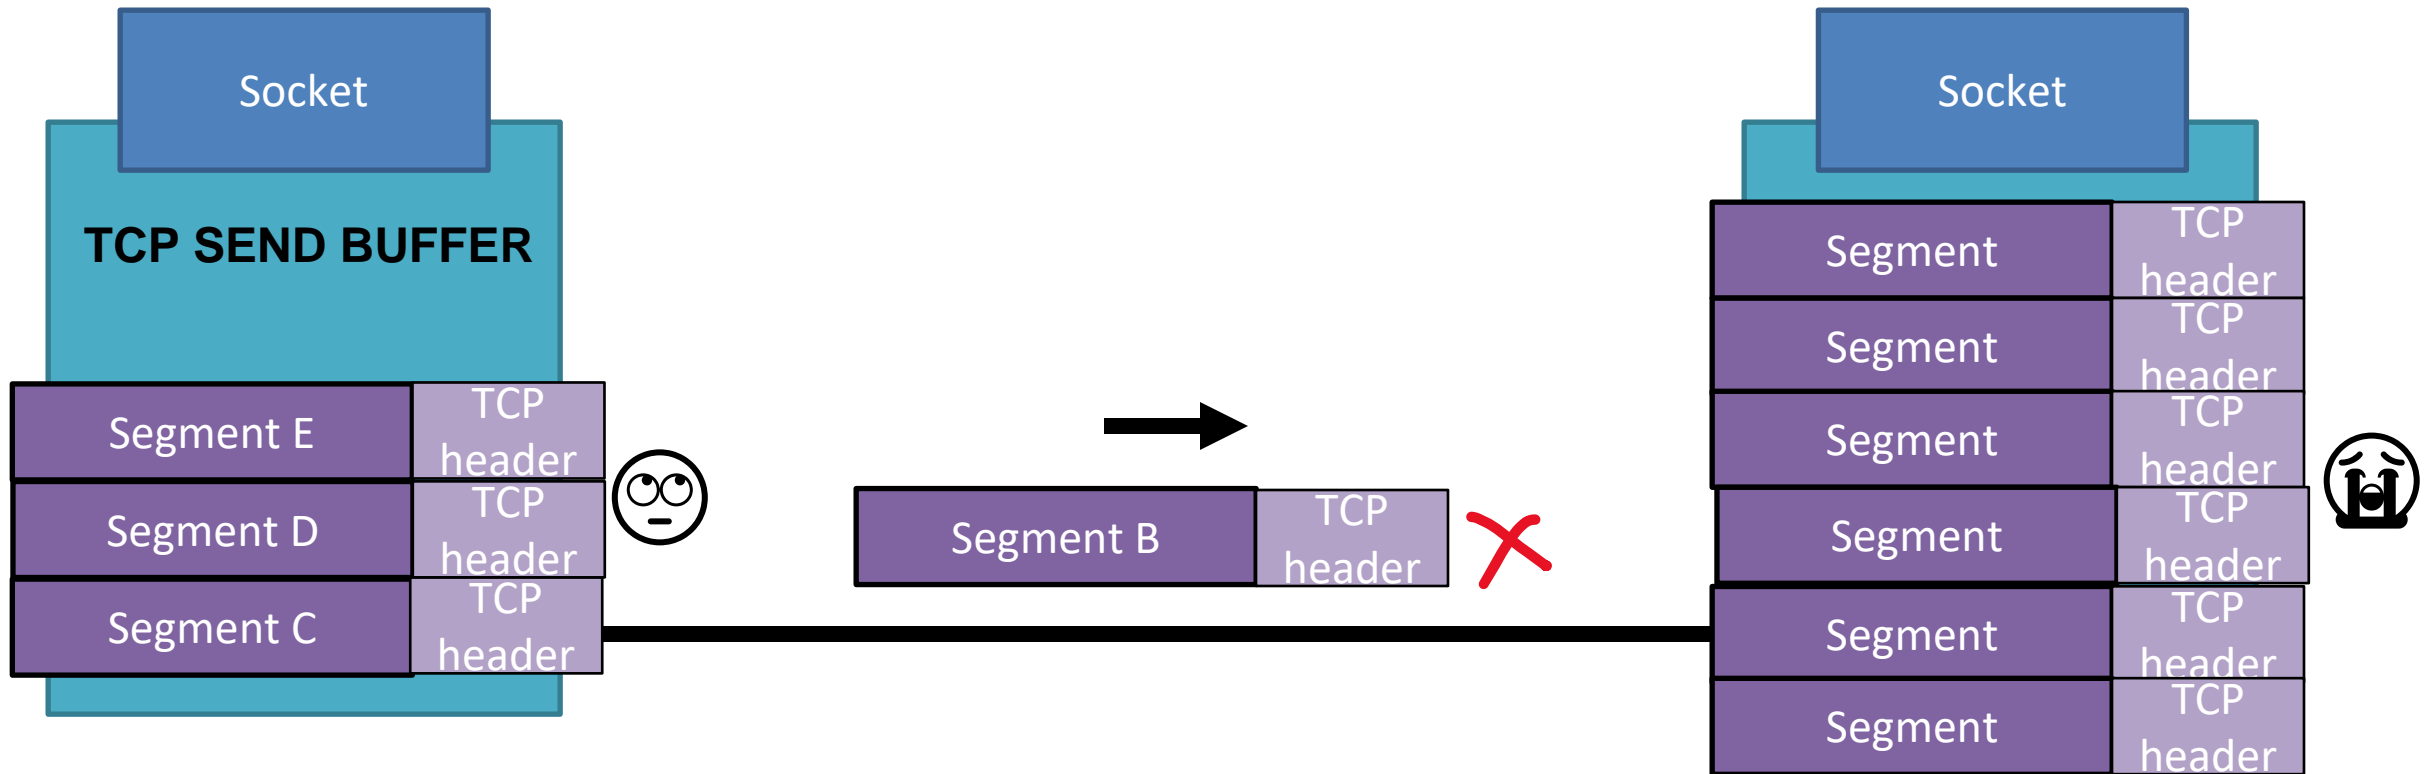


We could send back to the sender how much available space we have in our buffer!



Transport Layer

TCP Flow Control



We could send back to the sender how much available space we have in our buffer!

https://wps.pearsoned.com/ecs_kurose_compnetw_6/216/55463/14198702.cw/index.html

What is a good way to determine when to timeout? (aka the length of timer)

1. Too short: premature timeout, unnecessary retransmissions
2. Too long: slow reaction to segment loss


The TCP timeout value should be around the same time it takes to


What is a good way to determine when to timeout? (aka the length of timer)

1. Too short: premature timeout, unnecessary retransmissions
2. Too long: slow reaction to segment loss


The TCP timeout value should be around the same time it takes to receive an acknowledgement on a sent packet (on average)


Let's consider setting it to be a dynamic value!






11 seconds into the capture, the RTT captured 500ms





$$\text{EstimatedRTT} = (1 - \alpha) \cdot \text{EstimatedRTT} + \alpha \cdot \text{SampleRTT}$$

$$\alpha = 0.125$$



In addition, we also want some kind of safety margin

$$\text{DevRTT} = (1-\beta) * \text{DevRTT} + \beta * |\text{SampleRTT} - \text{EstimatedRTT}|$$

(typically, $\beta = 0.25$)


TimeoutInterval =

$$\text{EstimatedRTT} + 4 * \text{DevRTT}$$


(safety margin)

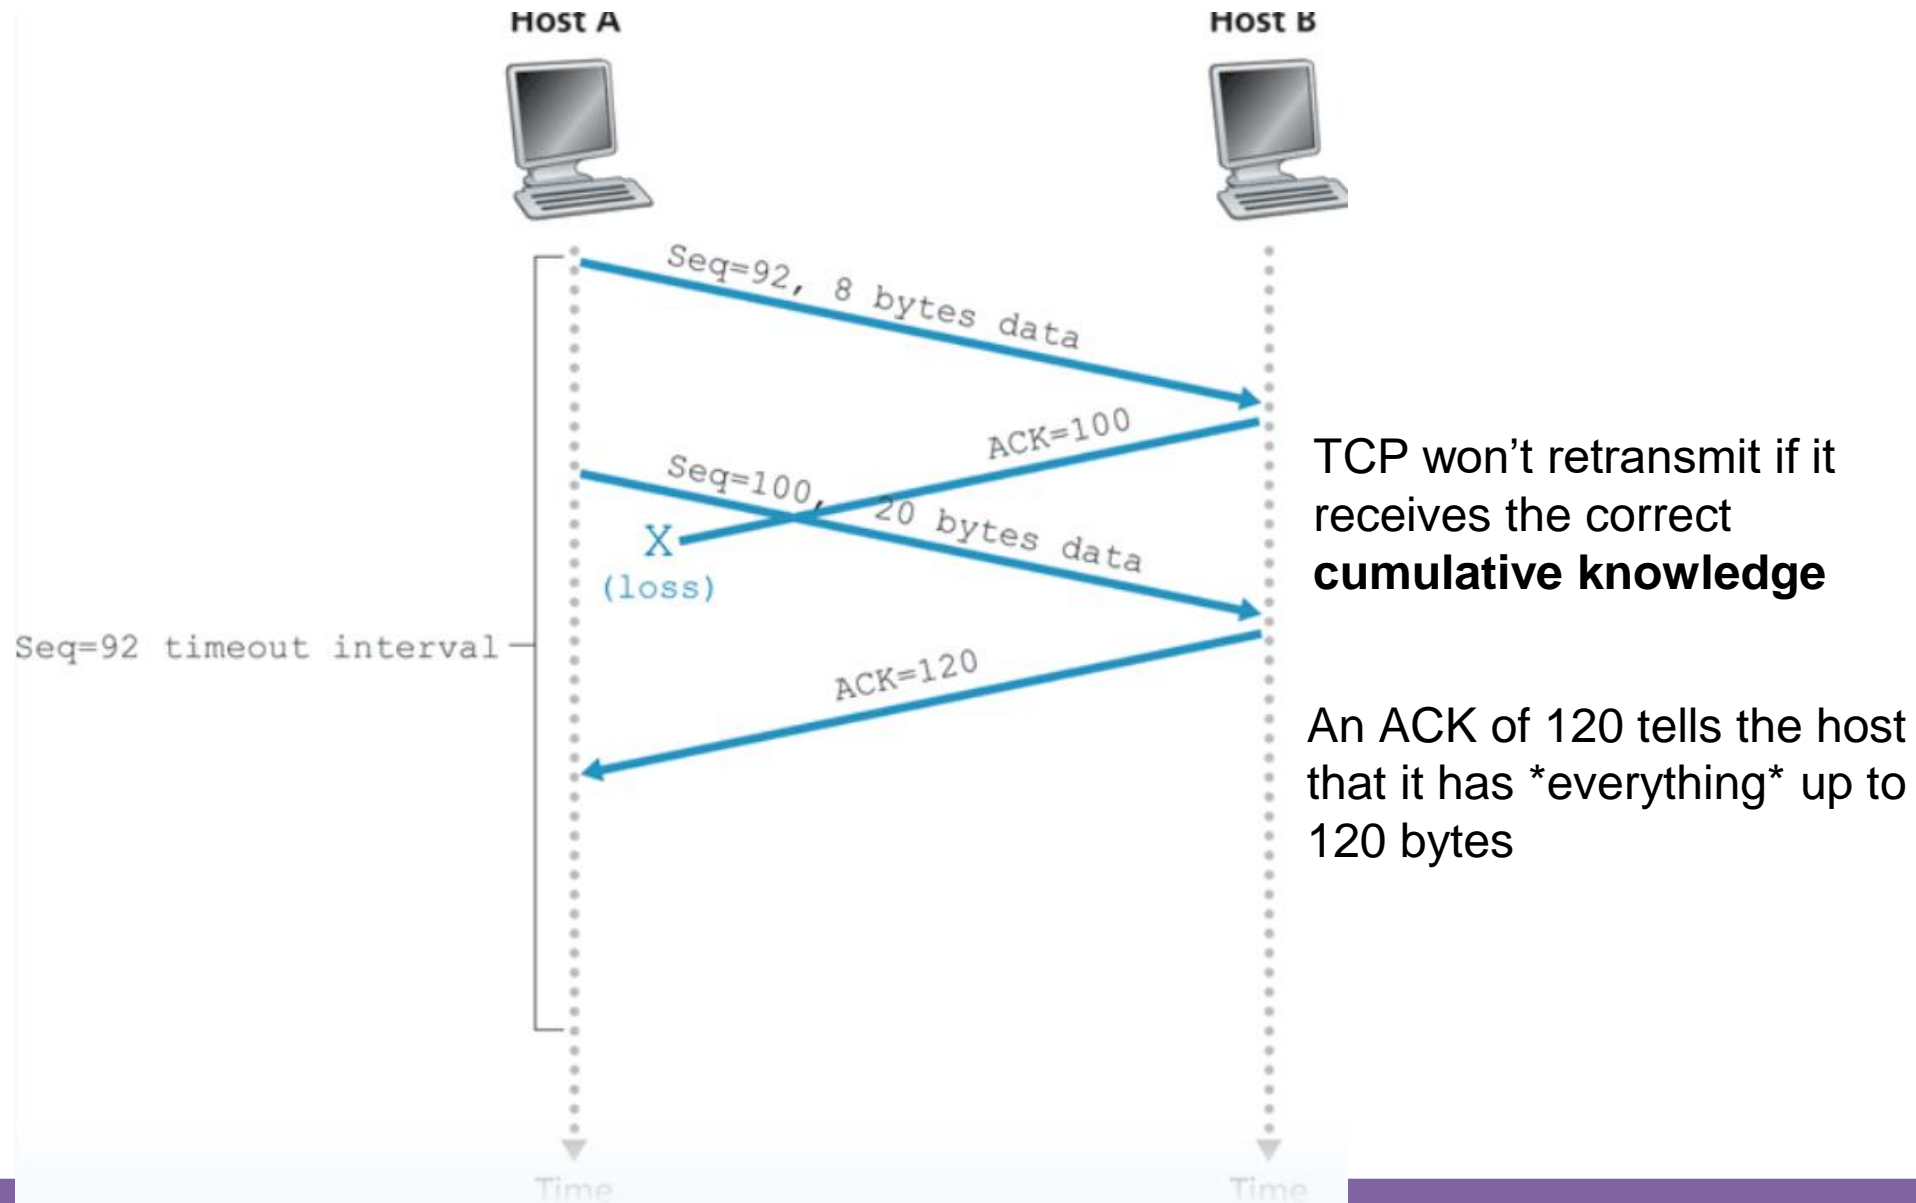
$$\text{EstimatedRTT} = (1 - \alpha) \cdot \text{EstimatedRTT} + \alpha \cdot \text{SampleRTT}$$

$$\alpha = 0.125$$



TCP retransmits on ACK loss





Event	TCP Receiver Action
Arrival of in-order segment with expected sequence number. All data up to expected sequence number already acknowledged.	Delayed ACK. Wait up to 500 msec for arrival of another in-order segment. If next in-order segment does not arrive in this interval, send an ACK.
Arrival of in-order segment with expected sequence number. One other in-order segment waiting for ACK transmission.	Immediately send single cumulative ACK, ACKing both in-order segments.
Arrival of out-of-order segment with higher-than-expected sequence number. Gap detected.	Immediately send duplicate ACK, indicating sequence number of next expected byte (which is the lower end of the gap).
Arrival of segment that partially or completely fills in gap in received data.	Immediately send ACK, provided that segment starts at the lower end of gap.

Specifics about when/how/why to send ACKs are described in TCP Congestion Control's **RFC (request for comments)**

9293

5681

Transport Layer

RFCs

Network Working Group
Request for Comments: 1149

D. Waitzman
BBN STC
1 April 1990


A Standard for the Transmission of IP Datagrams on Avian Carriers

Status of this Memo

This memo describes an experimental method for the encapsulation of IP datagrams in avian carriers. This specification is primarily useful in Metropolitan Area Networks. This is an experimental, not recommended standard. Distribution of this memo is unlimited.


Overview and Rational

<https://www.rfc-editor.org/>



Transmission Control Protocol

- Segment Structure




TCP header length in 32 bit words,

URG-urgent, ACK- ack number is valid, PSH-push, RST-reset connection,

SYN-used to establish connection, FIN-used to release connection

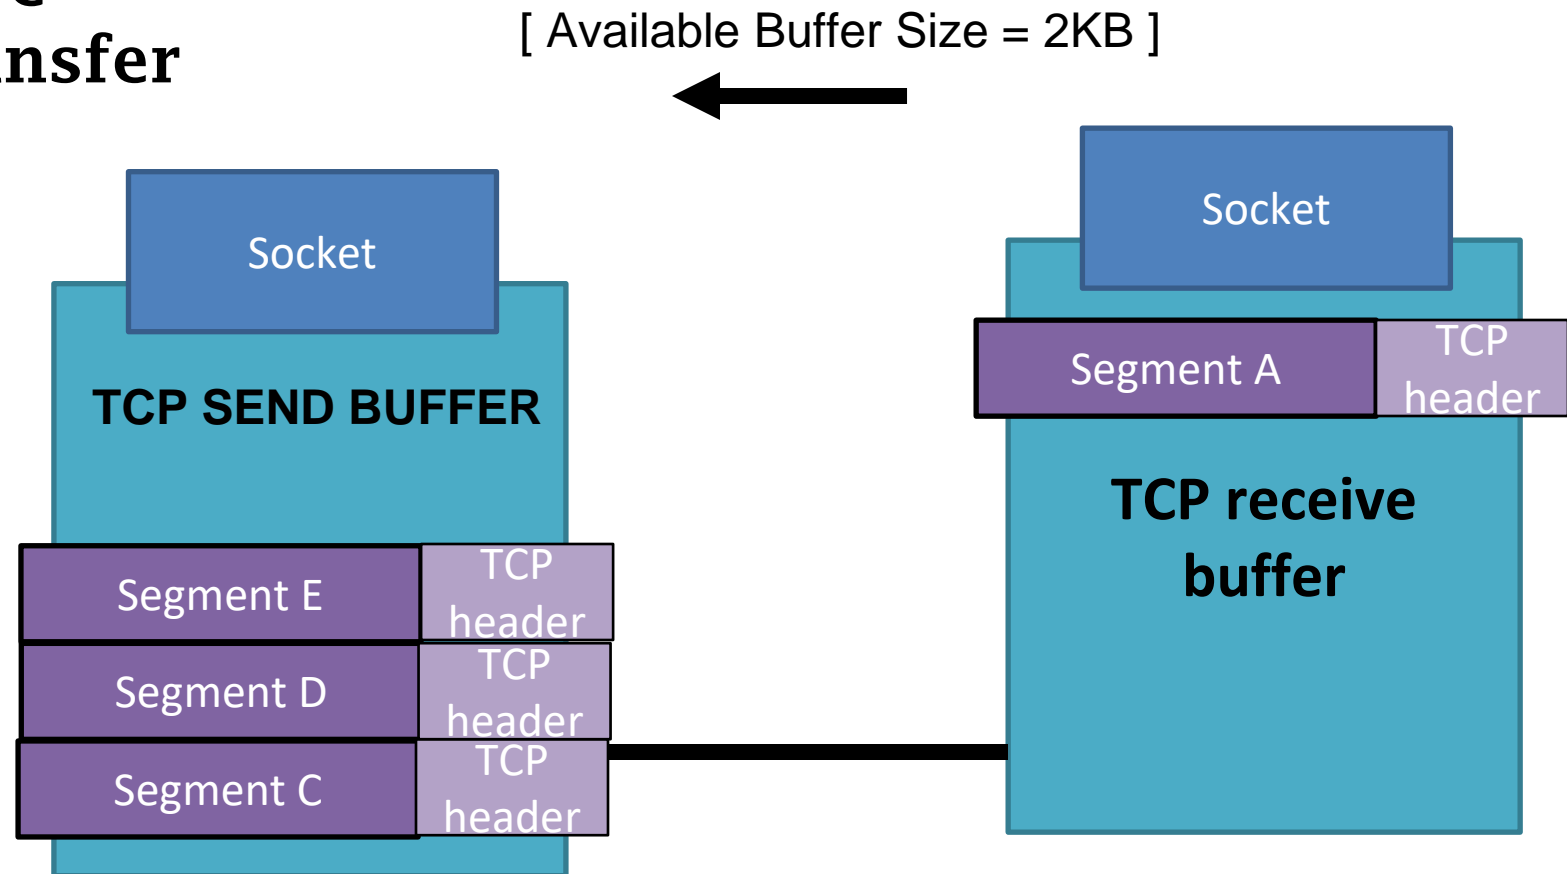
Transmission Control Protocol

- Segment Structure
- Reliable Data Transfer



Transmission Control Protocol

- Segment Structure
- Reliable Data Transfer
- Flow Control



Transmission Control Protocol

- Segment Structure
- Reliable Data Transfer
- Flow Control
- Connection Management

