

CSCI 466: Networks

Lecture 7: Transport Layer Wrap-up

Reese Pearsall
Fall 2022

Announcements

PA2 Due Wednesday October 19th

- Files must be pushed to a PA2 folder on your GitHub Repo
- **(NEW) BOTH MEMBERS MUST SUBMIT A REPO LINK TO D2L**
- Video demo with commentary is required
- GLHF

October is very busy for Reese


Reese is out of the country 10/13 – 10/18

- No class next Friday (10/14) → Work day for PA2
- No class next next Monday (10/17)
- I will still be reachable by email and discord


Extra Credit survey still available !!

- Virtual class on Monday 10/23

TCP Congestion Control




TCP sends back amount of available buffer space in the receiver
This helps make sure we don't overwhelm the receiver




Issues:

- If the network is congested, we want to slow down our sending rate
- If the network is not congested, we should try to send more stuff

TCP Congestion Control



TCP sends back amount of available buffer space in the receiver
This helps make sure we don't overwhelm the receiver




Issues:

- If the network is congested, we want to slow down our sending rate
- If the network is not congested, we should try to send more stuff

From the sender perspective, how could we measure how congested the network is?


TCP Congestion Control



TCP sends back amount of available buffer space in the receiver
This helps make sure we don't overwhelm the receiver

Some ways we could measure how congested the network is

- See how many dropped packets we are getting
- Amount of duplicate ACKs received
- Amount of UnAcked packets



Issues:

- If the network is congested, we want to slow down our sending rate
- If the network is not congested, we should try to send more stuff


TCP Congestion Control

TCP is **self-clocking**

TCP sender also has a **congestion window**, which controls the amount of unAck'd that can be sent out

The amount of unacknowledged data at a sender may not exceed the *minimum* of the congestion window and receiving window

$$\text{LastByteSent} - \text{LastByteAcked} \leq \min\{\text{cwnd}, \text{rwnd}\}$$




TCP Congestion Control

TCP is **self-clocking**

TCP sender also has a **congestion window**, which controls the amount of unAck'd that can be sent out

If TCP notices a **loss event**, it will slow down the sending rate




TCP Congestion Control Algorithm

TCP Algorithm to prevent network congestion

- **Slow Start**
- Congestion Avoidance
- Fast recovery

Start sending slow, but exponentially grows up to a *threshold*

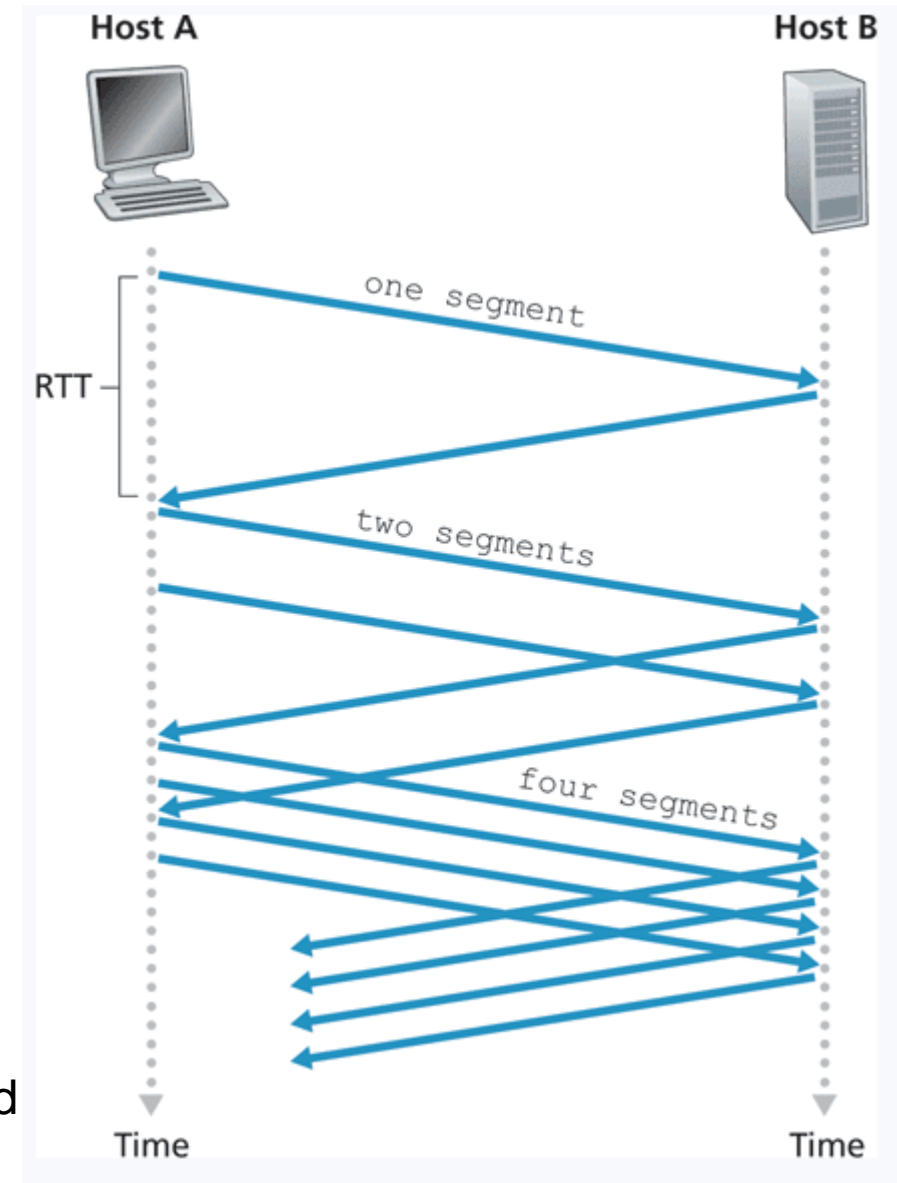


TCP Congestion Control Algorithm

TCP Algorithm to prevent network congestion

- Slow Start
- **Congestion Avoidance**
- Fast recovery

Linearly increase congestion window for each ACK received




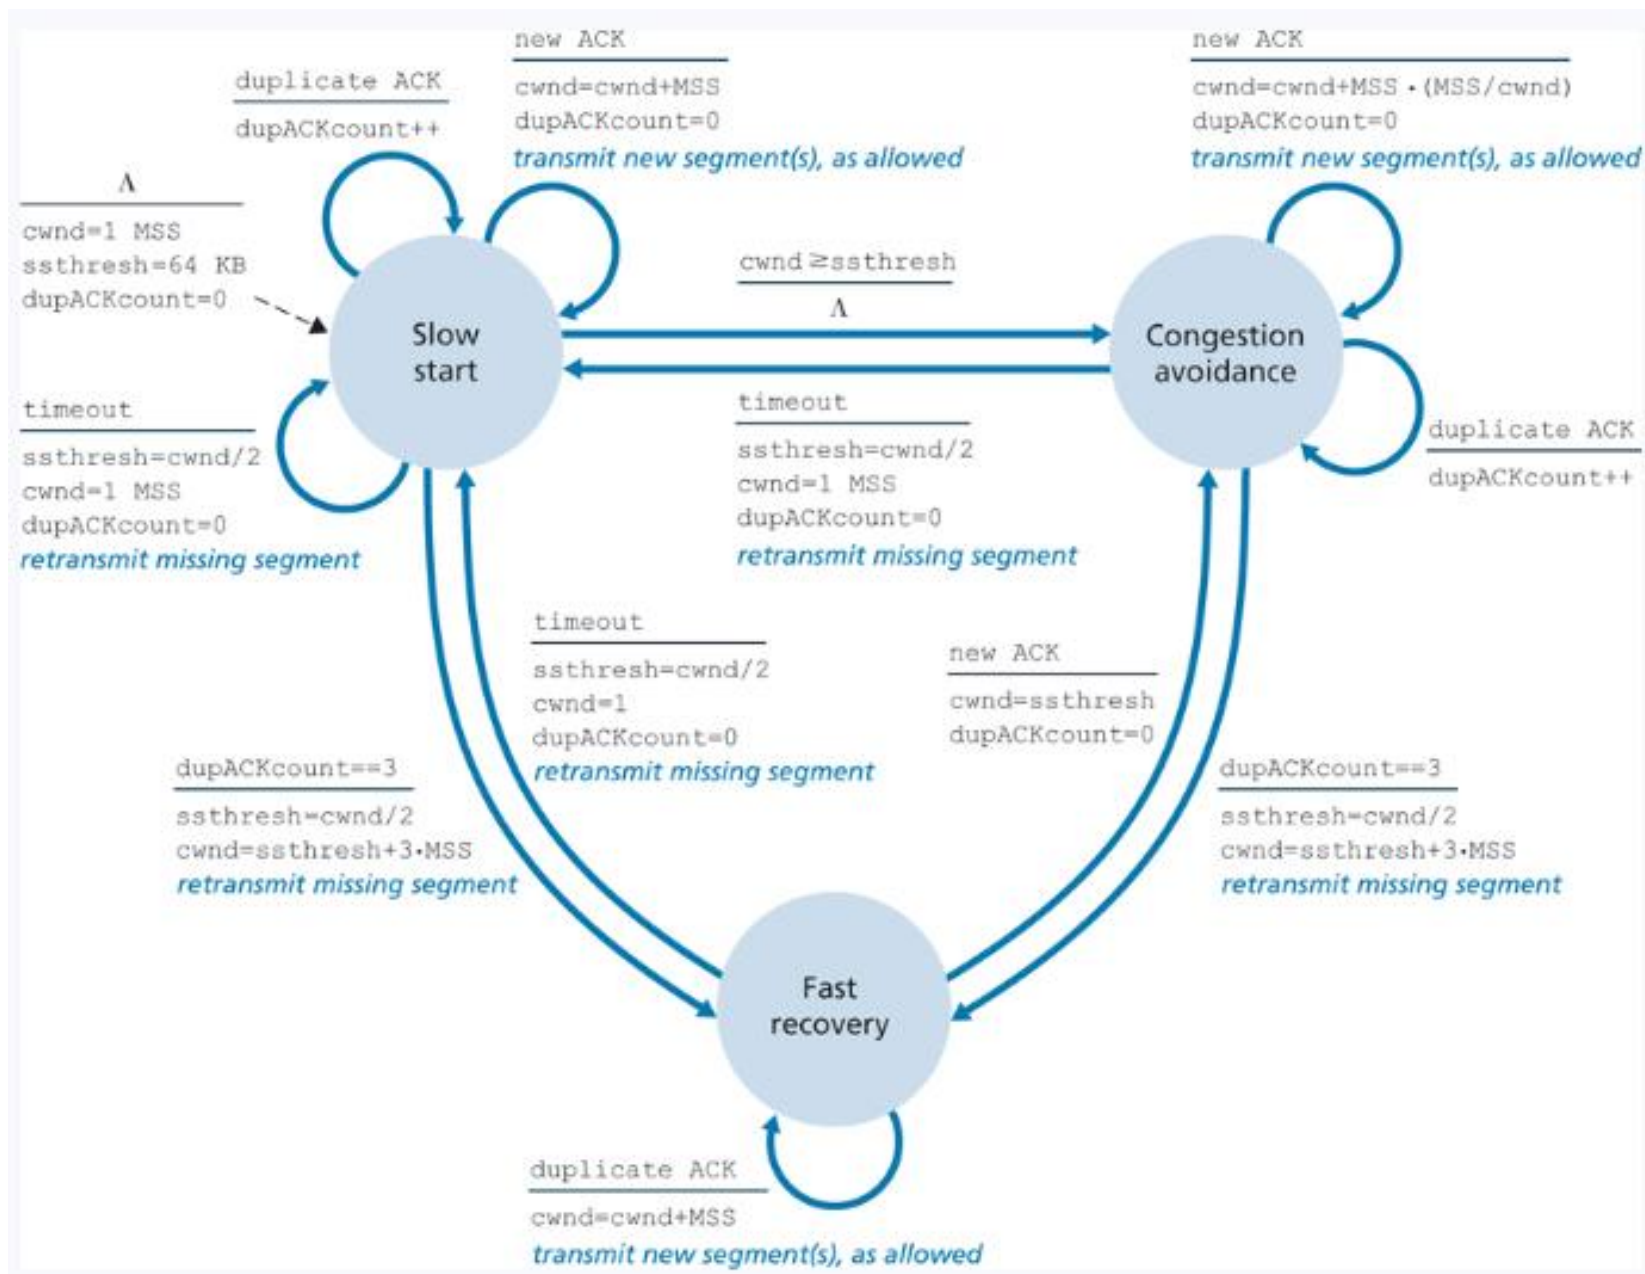
TCP Congestion Control Algorithm

TCP Algorithm to prevent network congestion

- Slow Start
- Congestion Avoidance
- **Fast recovery**

Upon knowledge of packet loss, throttle the TCP connection and start off slow again





TCP Congestion Control Algorithm

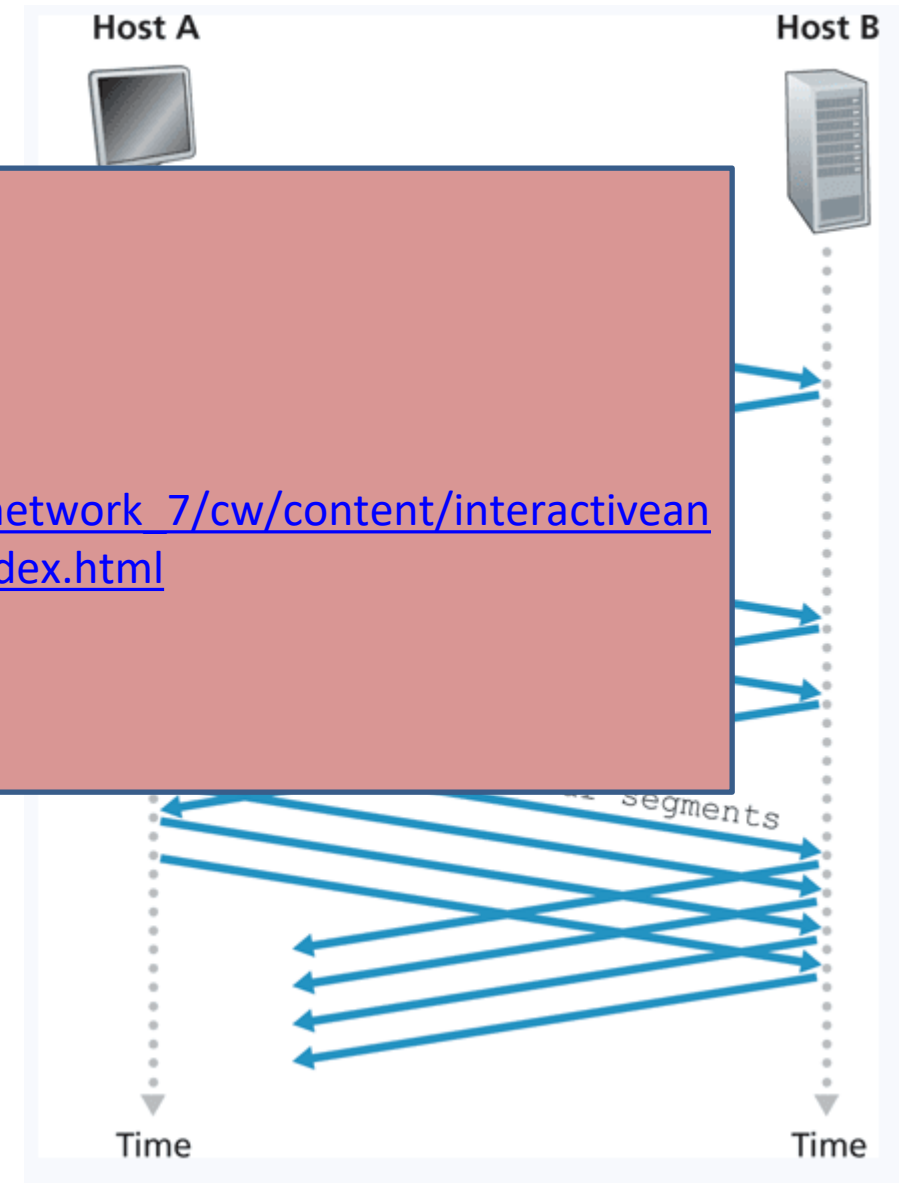
TCP Alg

- Slow
- Cong
- Fast


Upon knowledge of packet loss, throttle the TCP connection and start off slow again

Animation time!

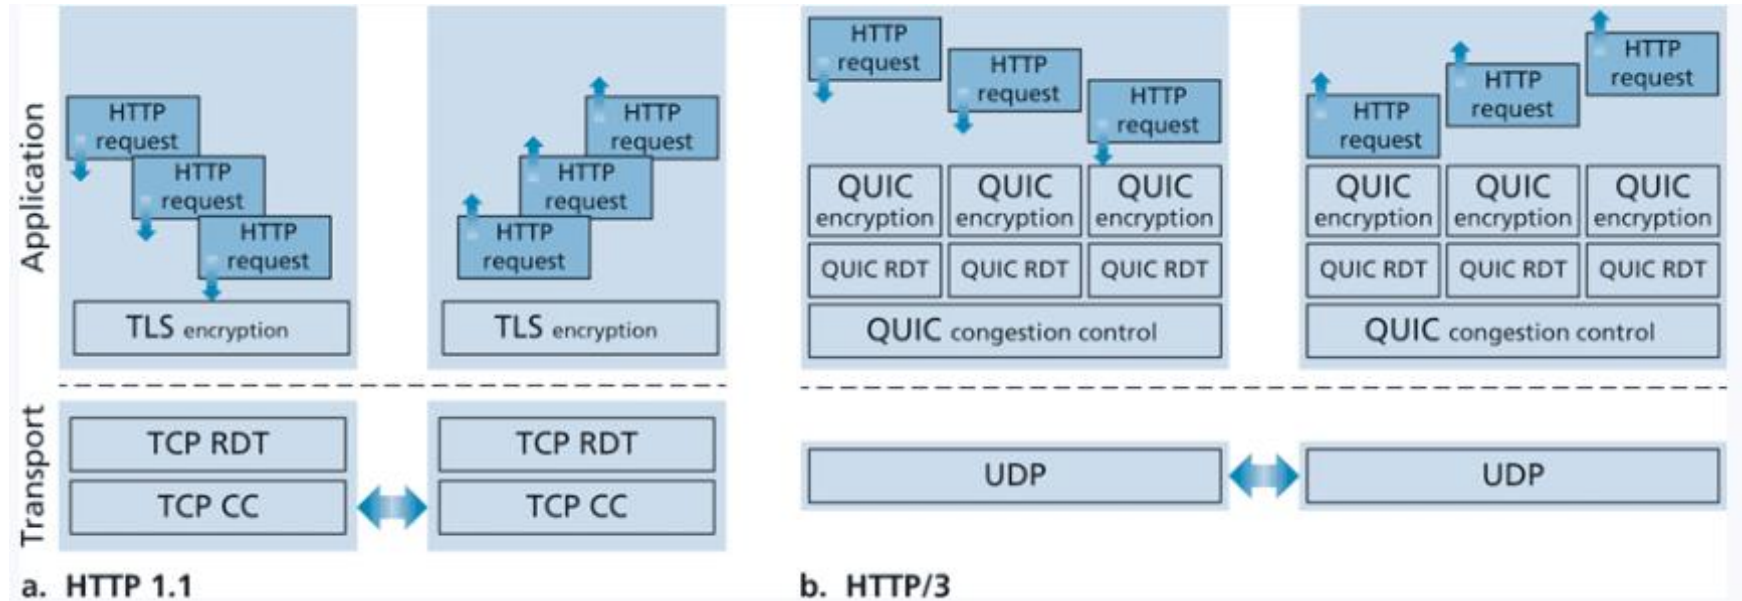
https://media.pearsoncmg.com/aw/ecs_kurose_compnetwork_7/cw/content/interactiveanimations/tcp-congestion/index.html



Current transport layer implementation



Transport layer protocols and congestion control is still a heavily researched area!



FIN

