

CSCI 466: Networks

Network Layer – Data Plane

Reese Pearsall
Fall 2022

Announcements

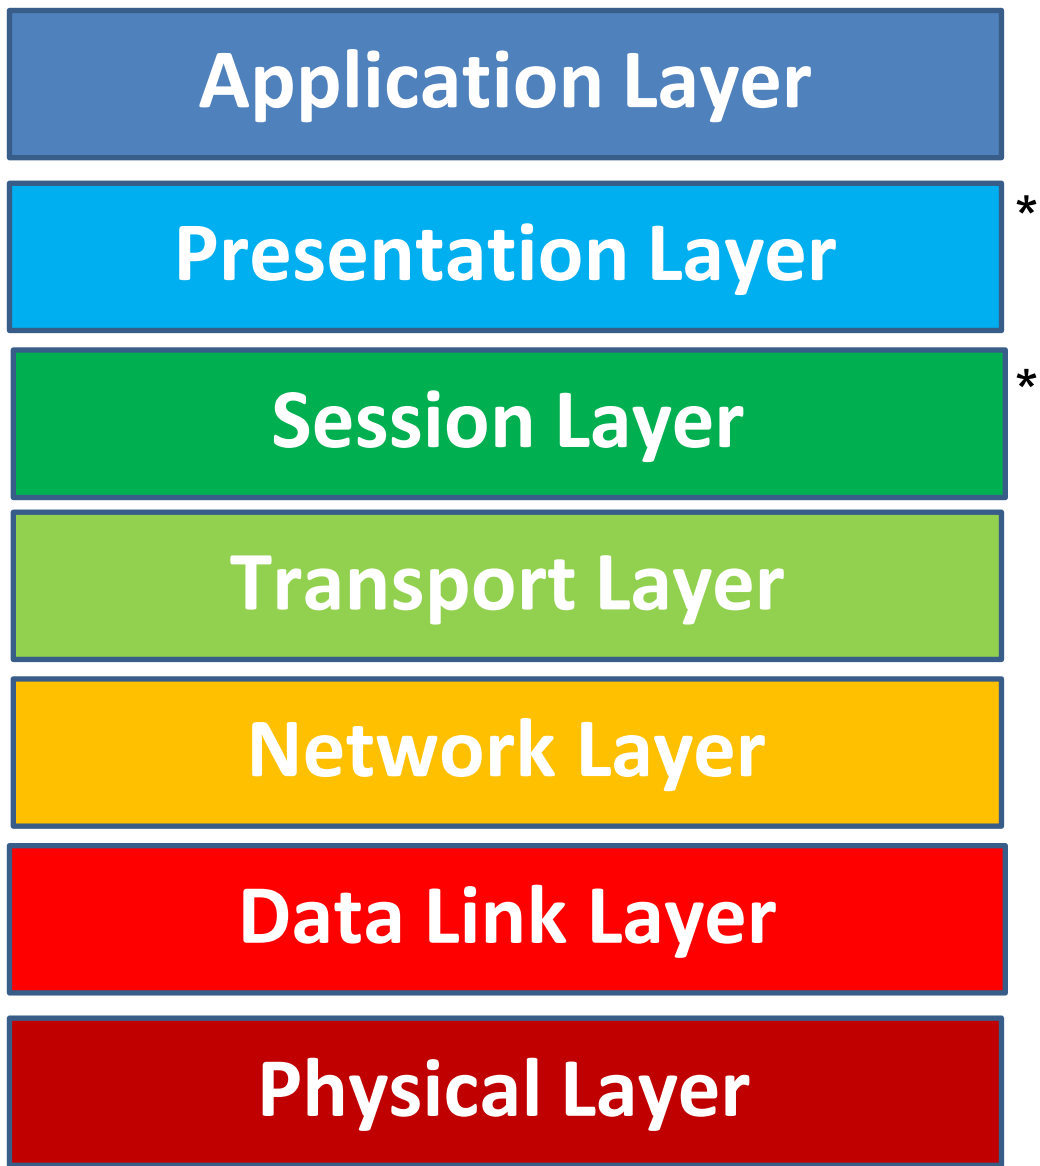
NO CLASS this Friday (10/14) and next Monday (10/17)

- You can always email/DM me if you have questions


Wireshark Lab 1 due TONIGHT @ 11:59 PM

- If you work with a partner, both need to submit to D2L!

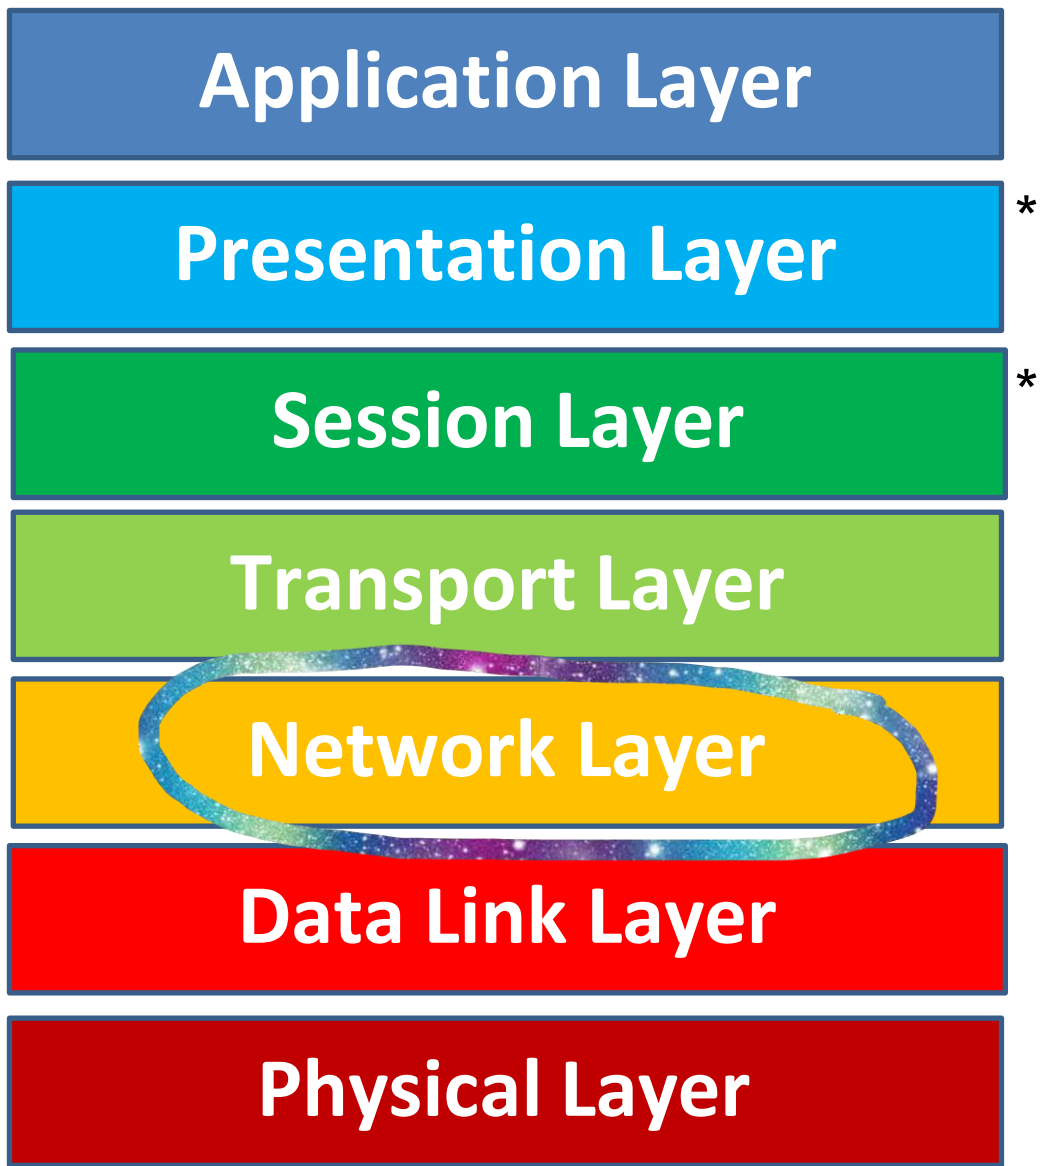
Extra Credit
survey still
available !!




OSI Model



**In the textbook, they condense it to a 5-layer model, but 7 layers is what is most used*



OSI Model



**In the textbook, they condense it to a 5-layer model, but 7 layers is what is most used*

Application Layer

Transport Layer

Our packet of information so far...

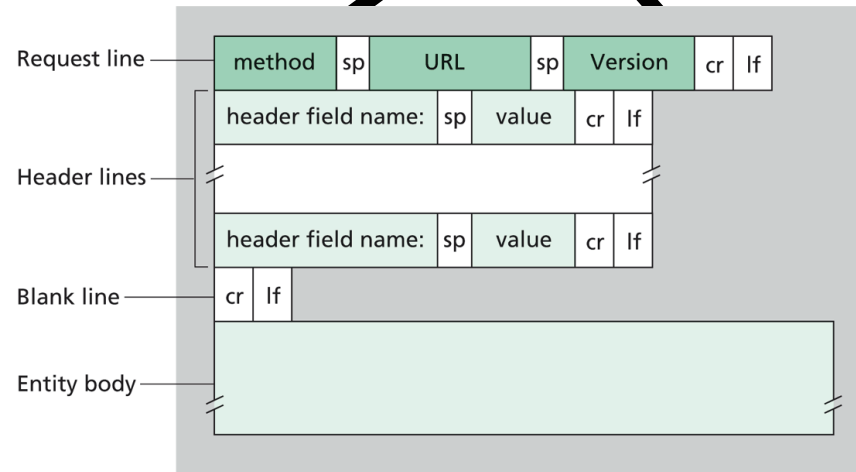


“Segment”



0				8				16				24				32			
Source Port								Destination Port											
Sequence Number																			
Acknowledgment Number																			
Data Offset		Reserved		C	E	U	A	P	R	S	F	Window Size							
		W	R	E	G	K	H	T	N										
Checksum								Urgent Pointer											
Options (if any)																			

Transport Layer Protocol
(ex. **TCP**)



Application Layer Protocol
(ex. **HTTP**)

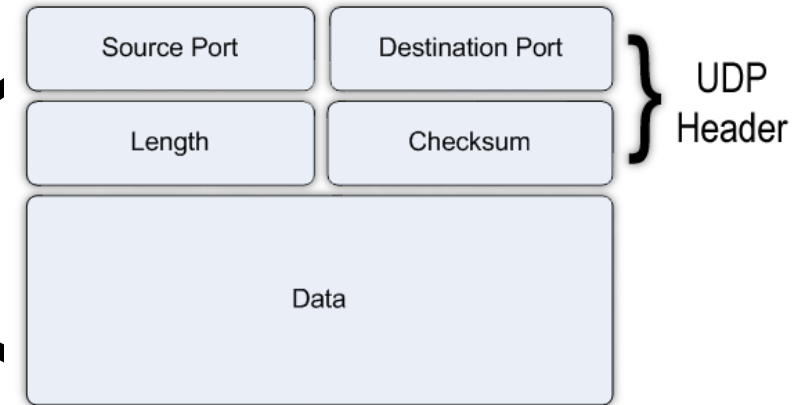
Application Layer

Transport Layer

Our packet of information so far...



“Segment”




Transport Layer Protocol
(ex. **UDP**)

Header Format

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
ID															
QR	Opcode				AA	TC	RD	RA	Z			RCODE			
QDCOUNT															
ANCOUNT															
NSCOUNT															
ARCOUNT															



Application Layer Protocol
(ex. **DNS**)




Application Layer


Transport Layer

We've only looked at protocols
that are running on some **host**

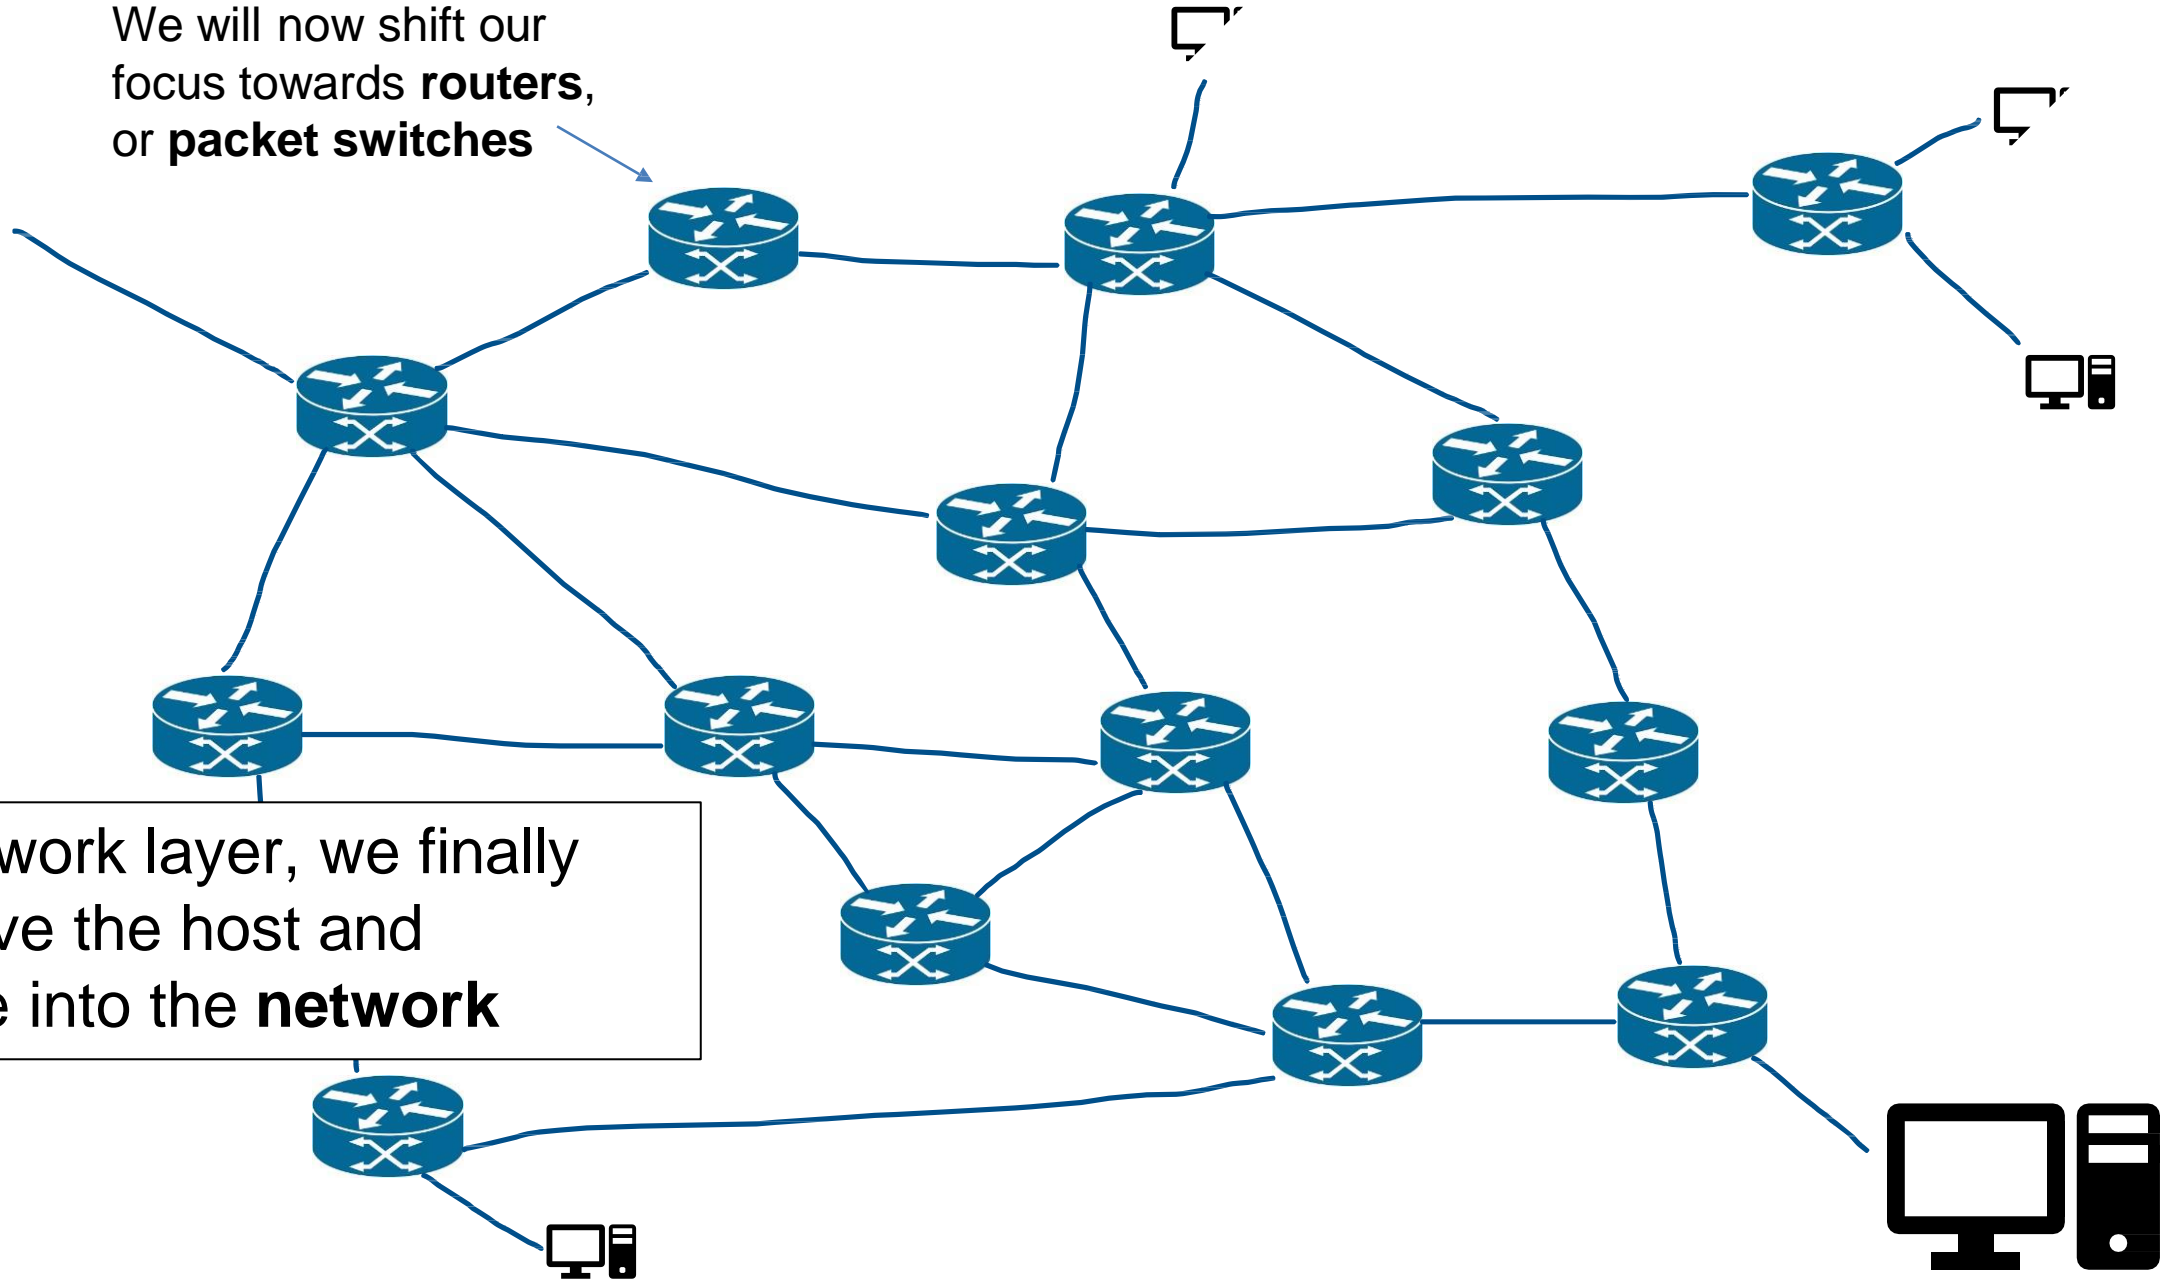


Network Layer

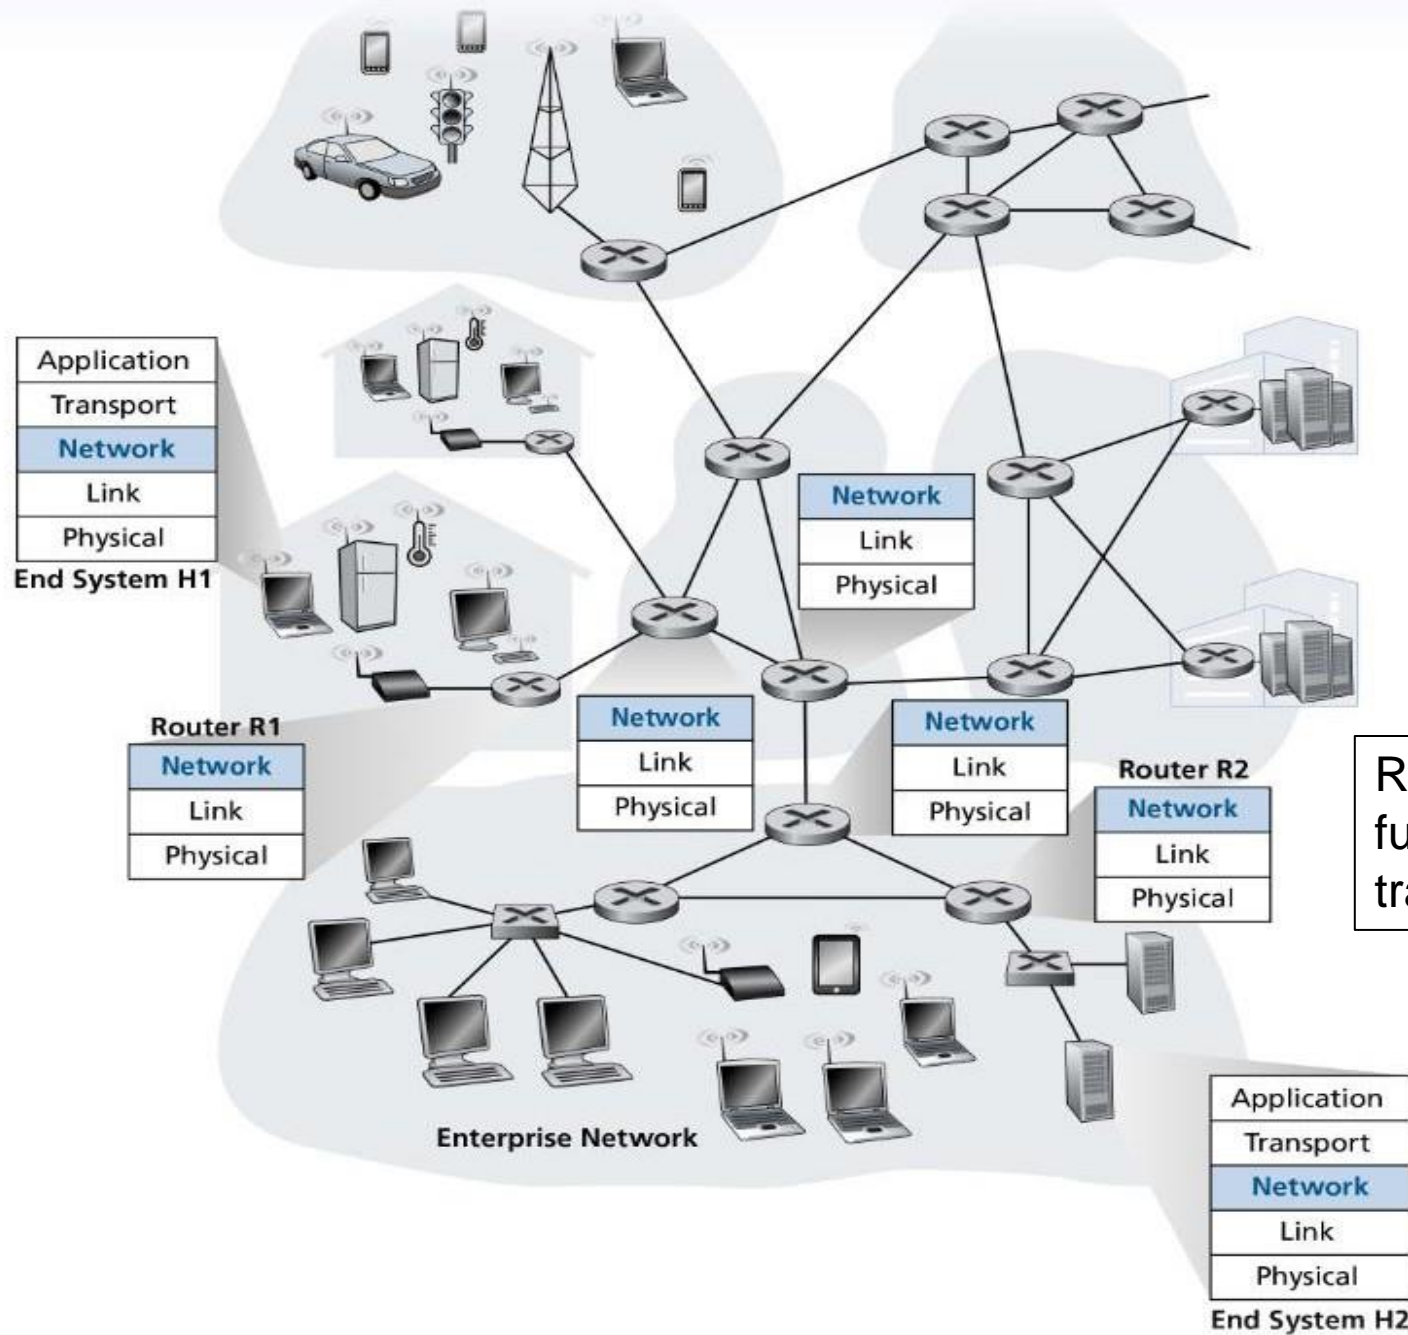





We will now shift our focus towards **routers**, or **packet switches**




In the network layer, we finally get to leave the host and adventure into the **network**







We will now shift our focus towards **routers**, or **packet switches**



Routing: determine route taken by packets from source to destination





We will now shift our focus towards **routers**, or **packet switches**



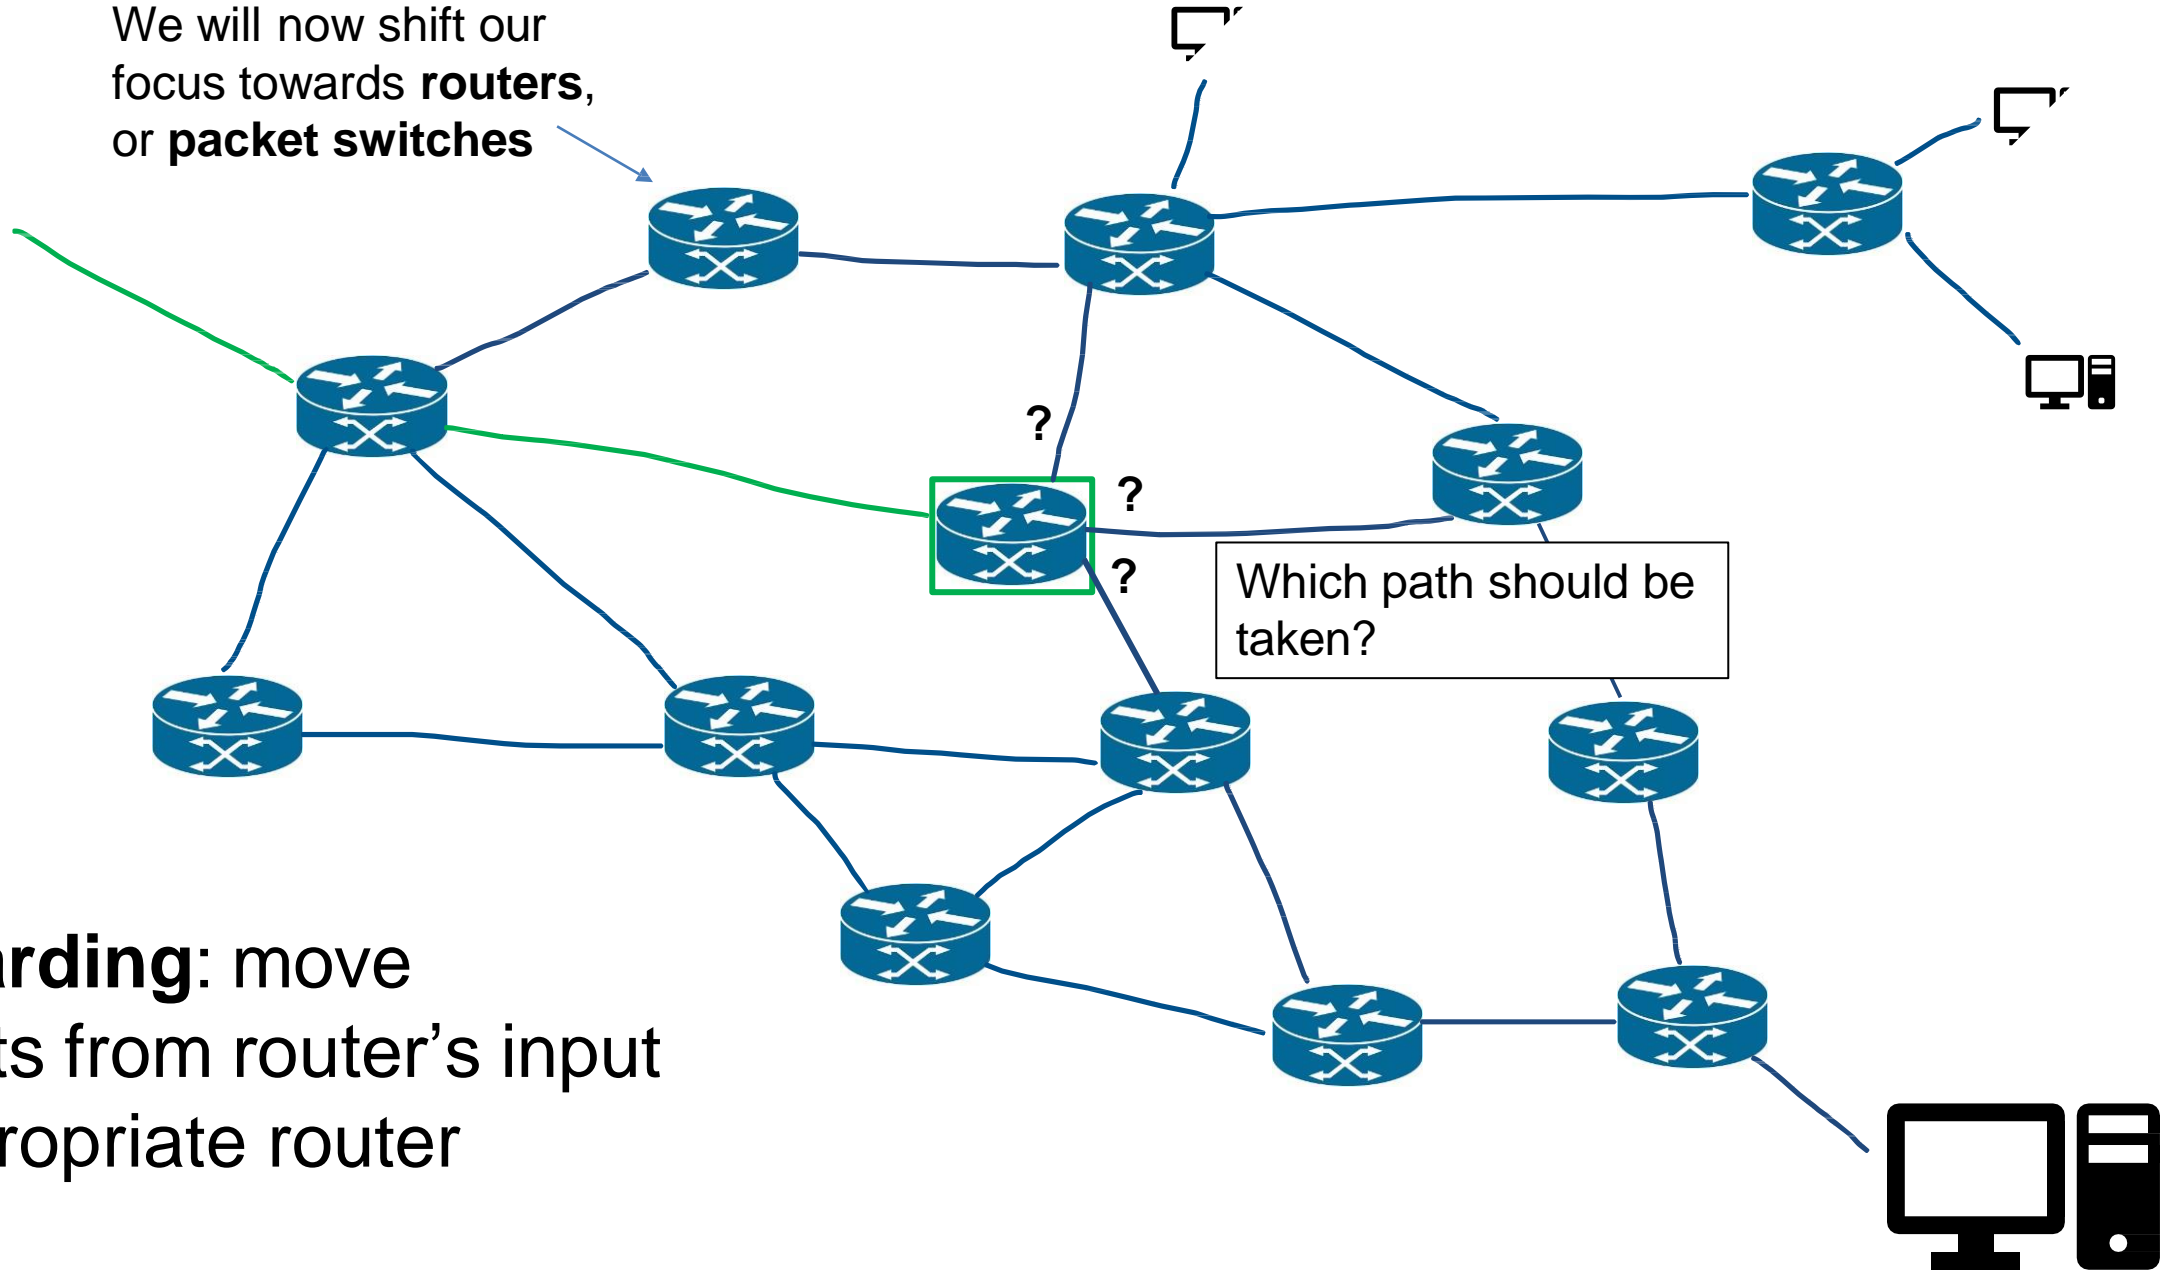
Routing: determine route taken by packets from source to destination

There is likely very many paths from one host to another! Which is the most efficient path?





We will now shift our focus towards **routers**, or **packet switches**



Forwarding: move packets from router's input to appropriate router output

Network Layer

Responsible for the delivery of data through a network


Forwarding

Data Plane

Routing

Control Plane

Router Architecture Overview



1. Destination-based forwarding

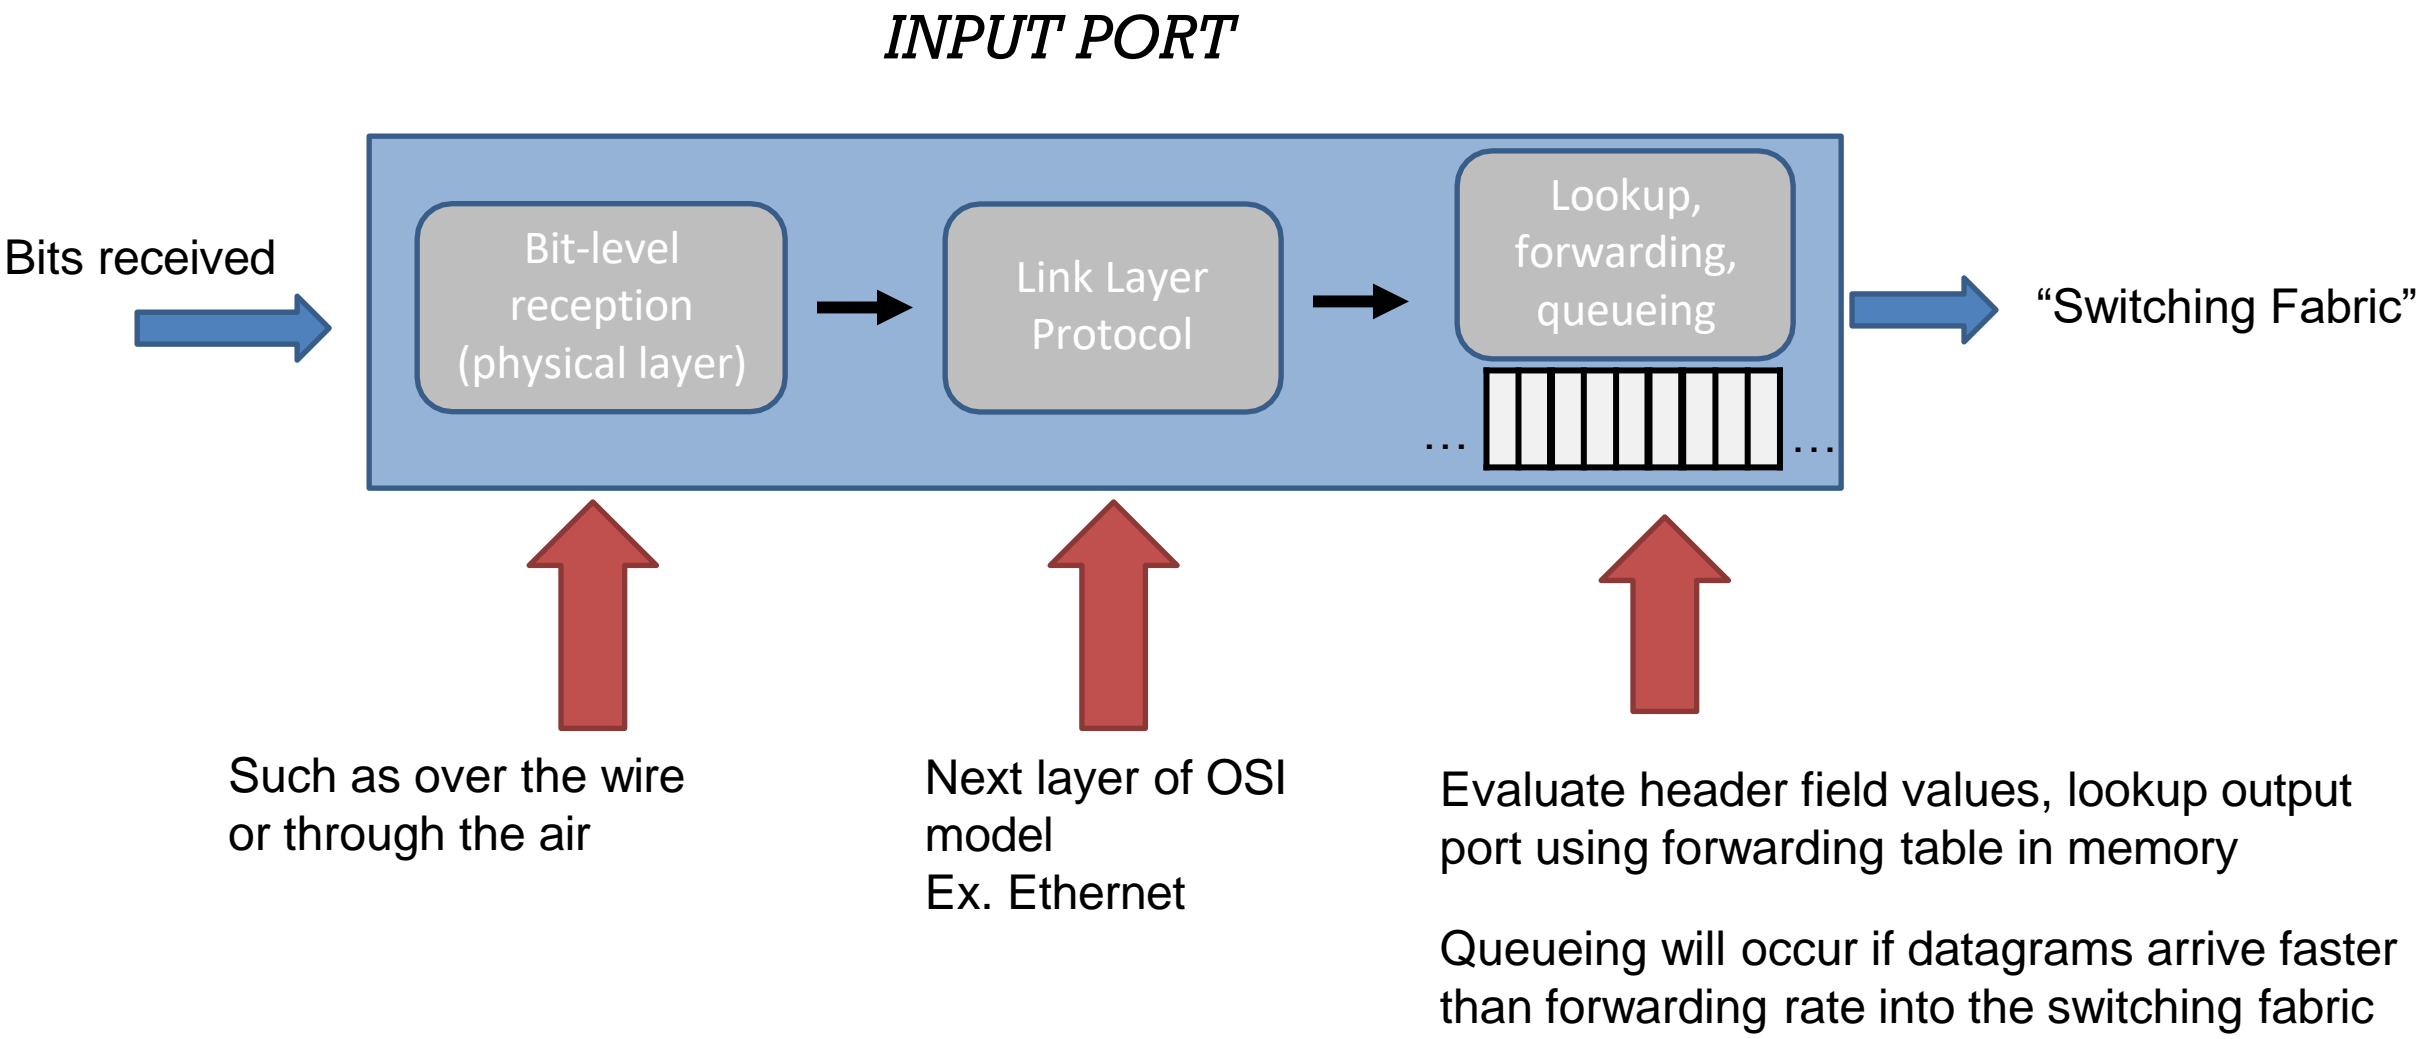
- Forwarding decisions are based on the **destination** of the packet

2. Generalized forwarding

- Forwarding decisions based on any set of header field values




Router Architecture Overview



Router Architecture Overview

OUTPUT PORT

“Switching Fabric”



To destination

Forwarding Techniques

- Connection-less
- does not require startup,
- Has no idea where the final destination is

Routing Table

Address range	Interface (output link)
128.11.52.0 – 128.11.52.255	1
153.90.2.0 – 153.90.2.255	2
153.90.2.87 – 153.90.2.89	3

This routing table could get very big...
IP addresses need 32/64 bits of memory each

Lookup,
forwarding,
queueing



Forwarding Techniques




Lookup,
forwarding,
queueing

...

Routing Table


Address range	Interface (output link)
128.11.52.0 – 128.11.52.255	1
153.90.2.0 – 153.90.2.255	2
153.90.2.87 – 153.90.2.89	3

Virtual Circuit Networks

VC identifier	Interface (output link)
22 	1
13 	2
5 	3

- Connection-less
- does not require startup,
- Has no idea where the final destination is

- Connection-oriented
- Requires startup



Costly ☹️

Forwarding Techniques

Longest prefix matching

Address range	Interface (output link)
11001000 00010111 00010*** *****	1
11001000 00010111 00011000 *****	2
11001000 00010111 00011*** *****	3
otherwise	4

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address

examples:

DA: 11001000 00010111 00010110 10100001
DA: 11001000 00010111 00011000 10101010

which interface?
which interface?

Longest prefix matching

Address range	Interface (output link)
11001000 00010111 00010*** *****	1
11001000 00010111 00011000 *****	2
11001000 00010111 00011*** *****	3
otherwise	4

Lookup,
forwarding,
queueing



when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address




These lookups need to happen in nanoseconds for our network to function

Ternary content addressable memories (TCAMs) are used in modern routers to do **LPM**
Cisco routers can carry millions of TCAM entries in their routers

Fabric Switch and Switching

Switching fabric: Mechanism that forwards data from an input port to output port

Switching via memory:




Handled by a CPU (routing processor)

Cannot forward data in parallel

Fabric Switch and Switching

Switching fabric: Mechanism that forwards data from an input port to output port


Switching via memory:



Handled by a CPU (routing processor)

Cannot forward data in parallel

Switching via bus:



Datagrams are prepended with a header


A “bus” transports input port datagrams to output ports

“keep going around the roundabout until you find your port”

Fabric Switch and Switching


Switching fabric: Mechanism that forwards data from an input port to output port

Switching via memory:



Handled by a CPU (routing processor)
Cannot forward data in parallel

Switching via bus:




Datagrams are prepended with a header
A “bus” transports input port datagrams to output ports

“keep going around the roundabout until you find your port”

One shared bus, but can only fit one datagram

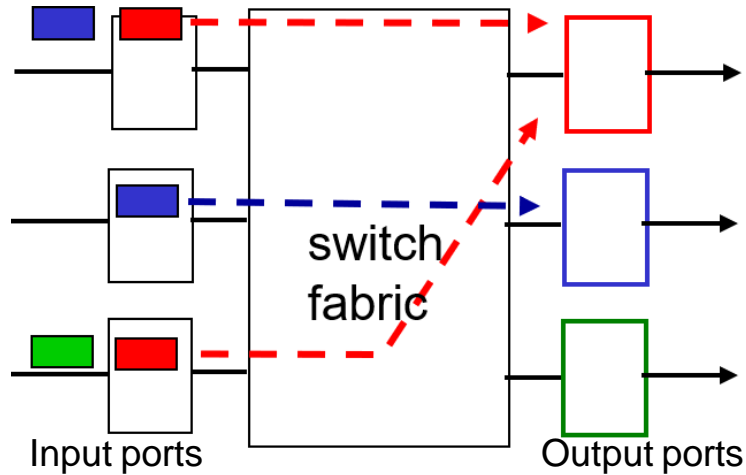
Switching via network:



A network of buses


Input Queueing

fabric slower than input ports combined → queueing may occur at input queues



output port contention:
only one red datagram can be
transferred.

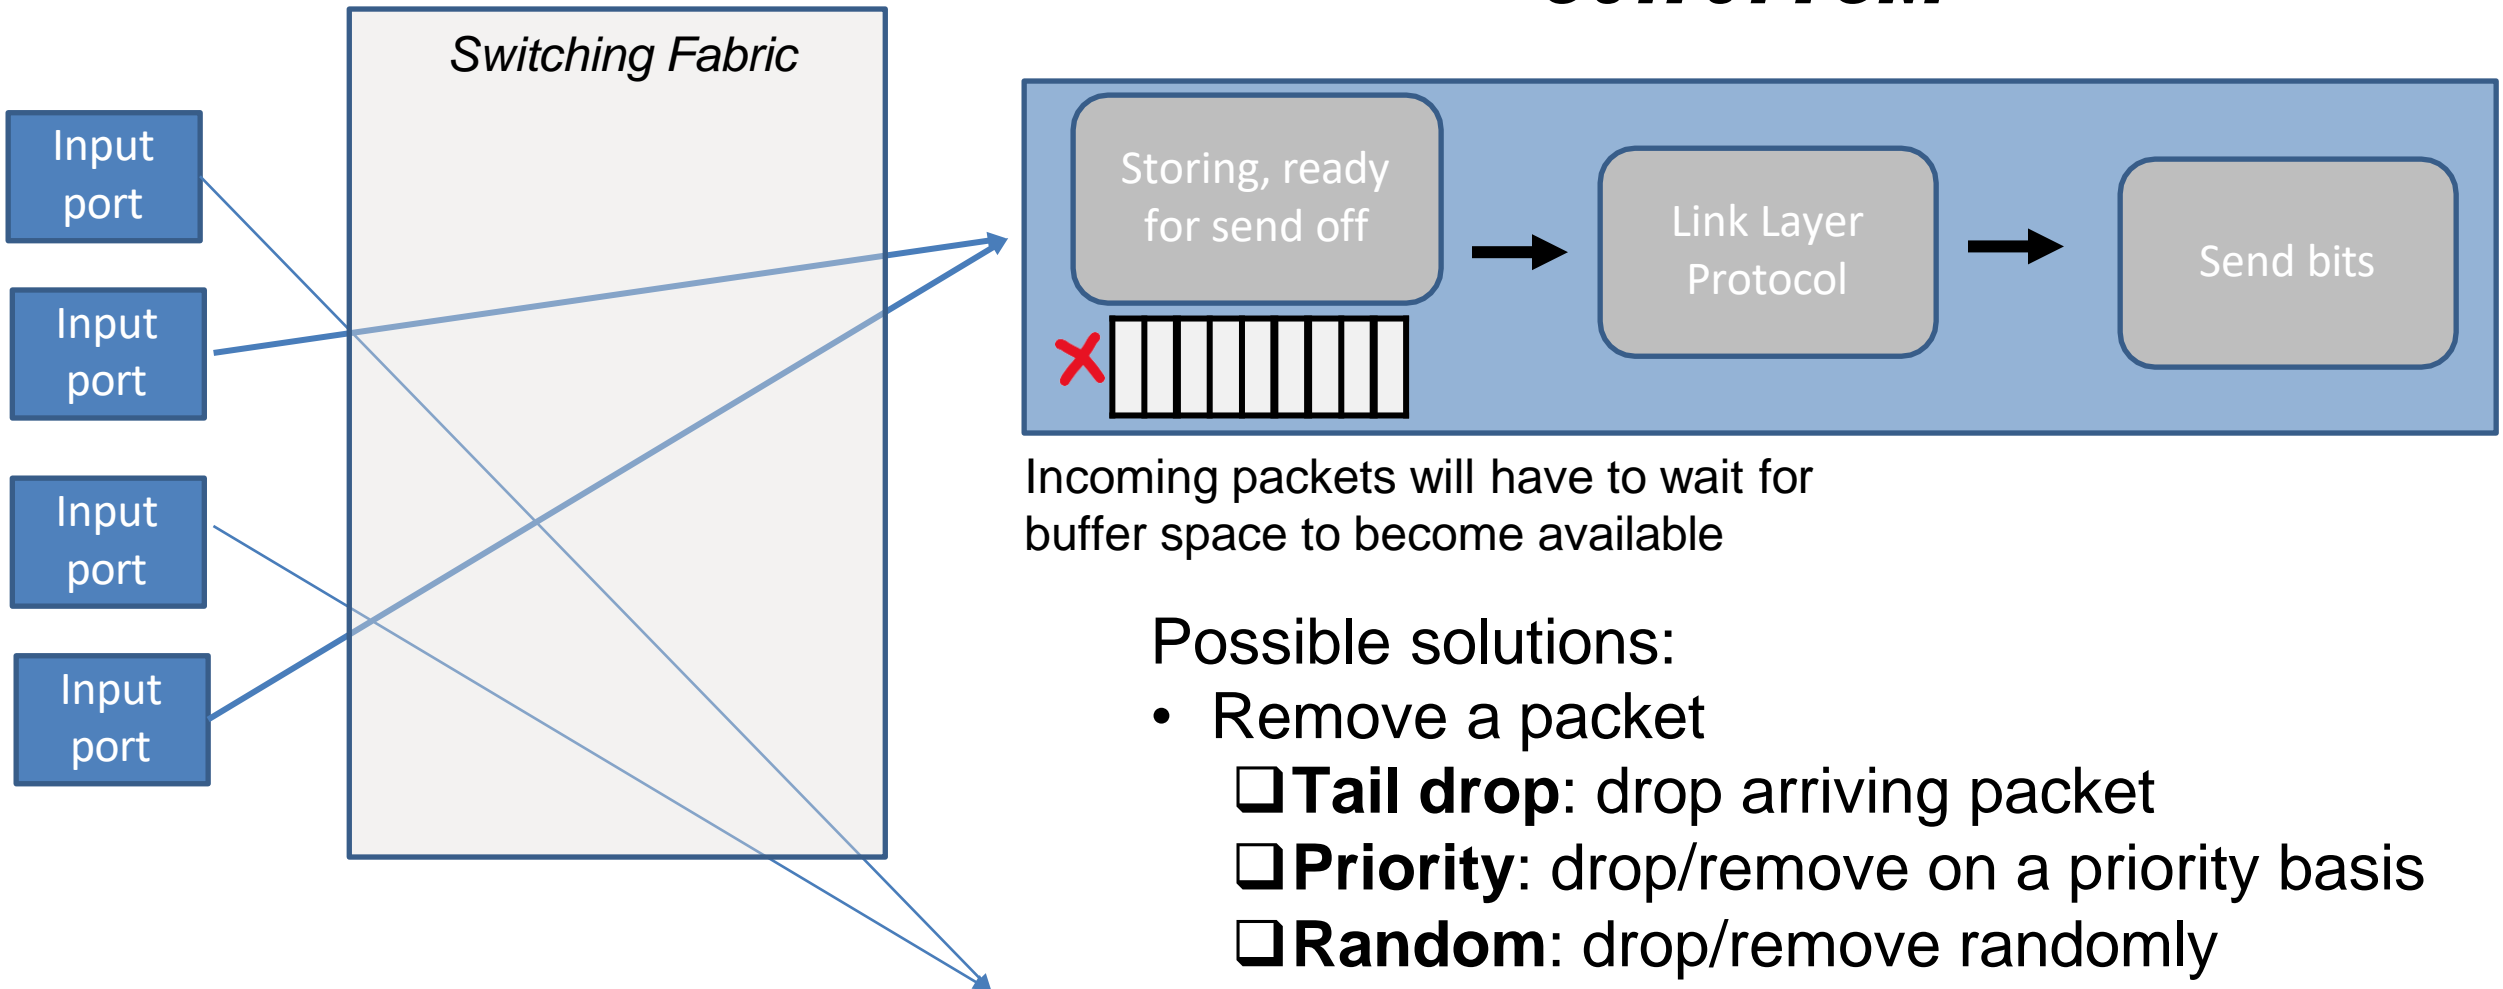
lower red packet is blocked



one packet time later:
green packet
experiences HOL
blocking

Head-of-the-Line (HOL) blocking: queued datagram at front of queue prevents others in queue from moving forward


Output Queueing



Packet Scheduling

Scheduling is used to determine the next packet to send on the link


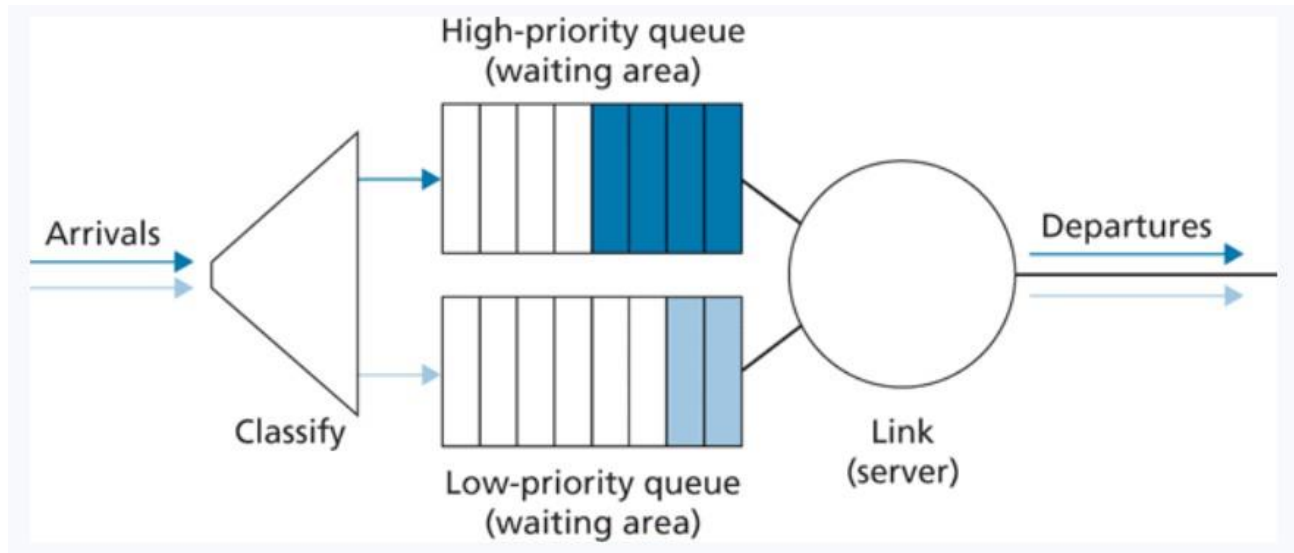
FIFO (first in first out): Send in order of arrival to queue



Packet Scheduling

Scheduling is used to determine the next packet to send on the link


Priority: packets are classified into priority classes. High priority = sent over link first



Packet Scheduling

Scheduling is used to determine the next packet to send on the link

Round robin weighted fair queueing: packets are classified into priority classes. Each class gets to send one packet during a “cycle”



Packet from the priority group gets sent

Packet from the middle class gets sent

Packet from the plebian class gets sent

IP PROTOCOL

IP addressing, IPv4, and IPv6

<https://www.rfc-editor.org/rfc/rfc791>

Packets traversing through the network layer are referred to as a **datagram**. Each packet gets an IPv4/IPv6 header

IPv4: 32-bit addresses (decimal)

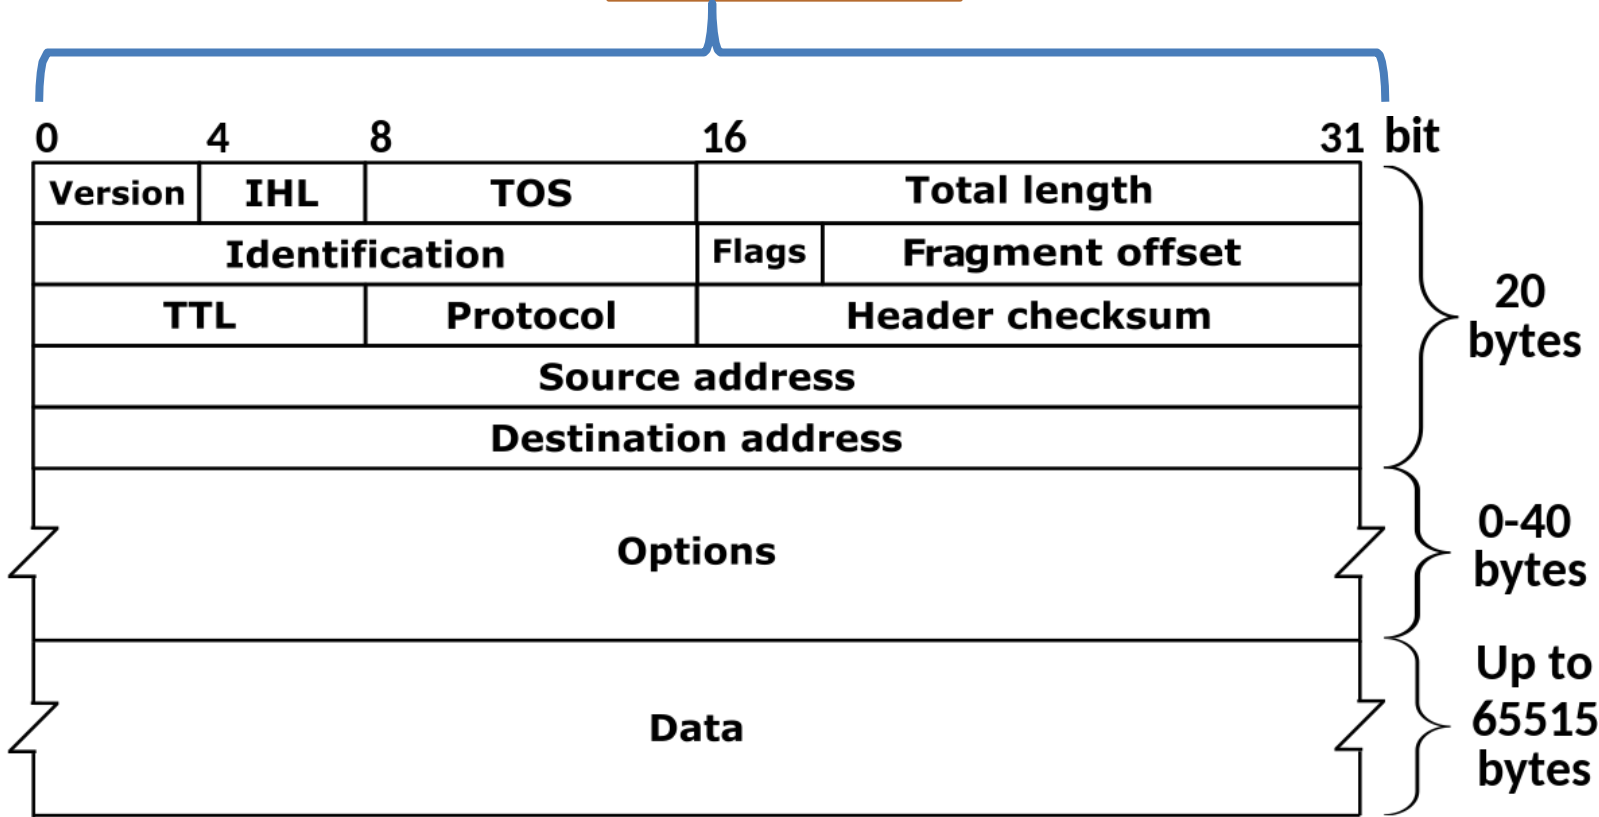

192.149.252.76

IPv6: 64-bit addresses (hexademical)

3ffe:1900:fe21:4545::

Why do we need IPv6?

IPv4 Header



Announcements


NO CLASS this Friday (10/14) and next Monday (10/17)

- You can always email/DM me if you have questions

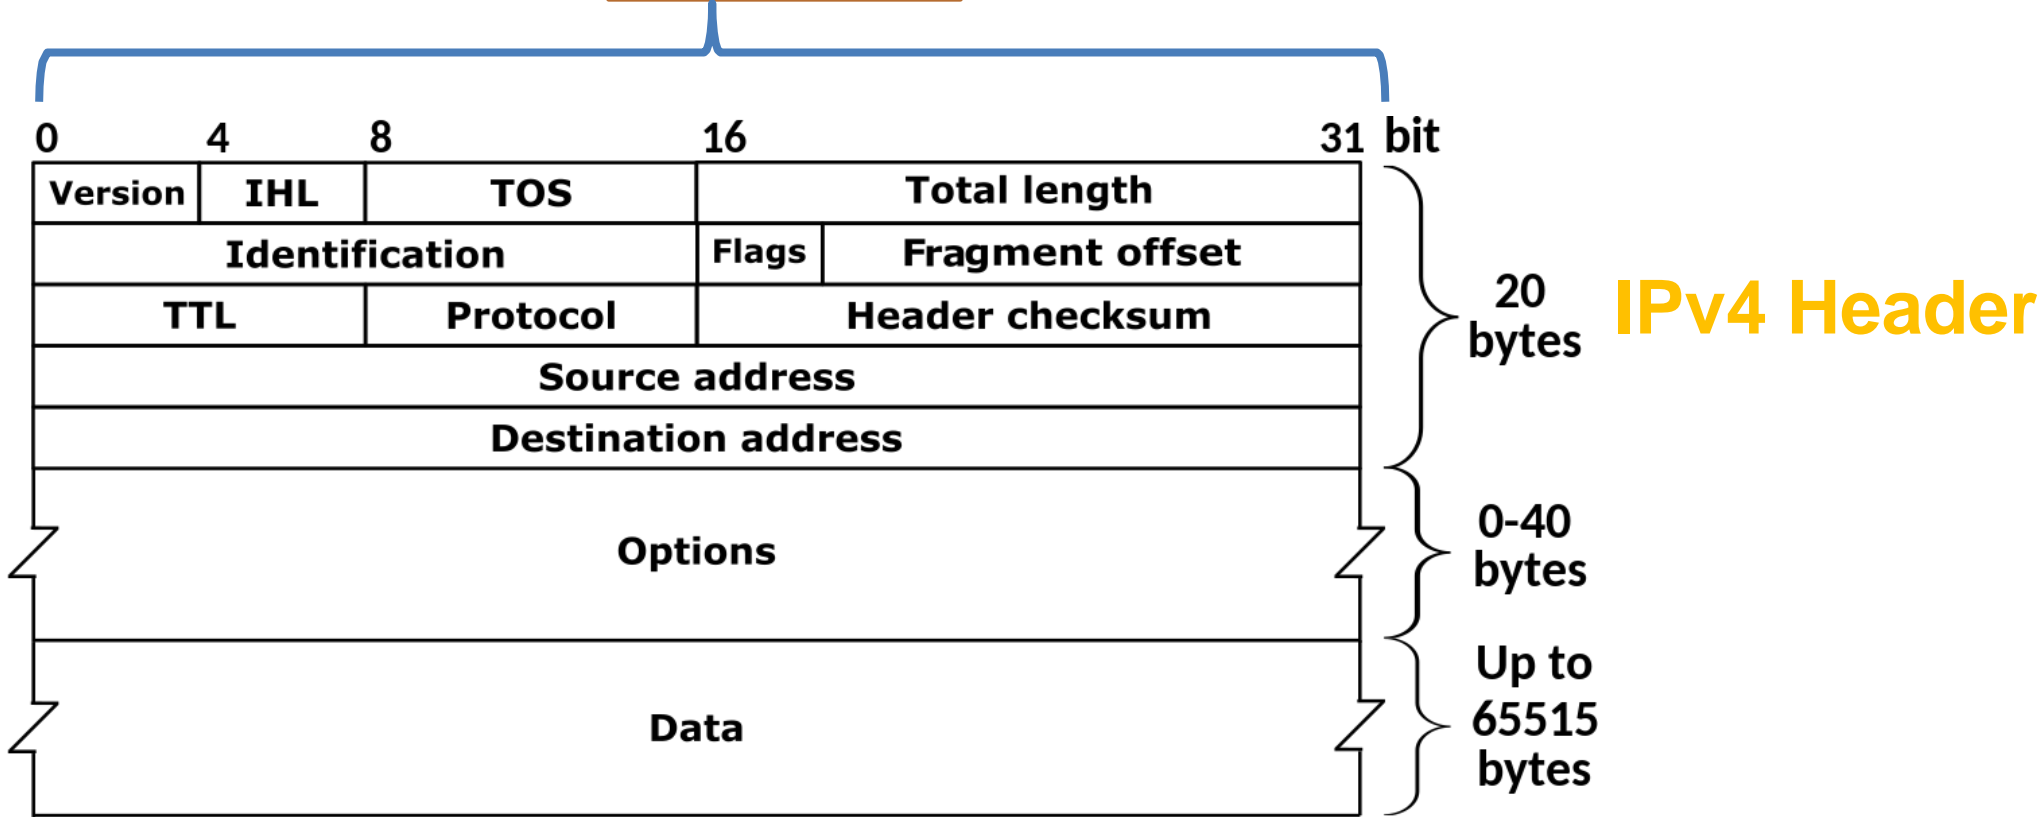
PA2 due in one week (10/19)

Put `time.sleep()` commands after you `send()` to prevent weird timing errors

Packets traversing through the network layer are referred to as a **datagram**. Each packet gets an IPv4/IPv6 header



TCP header = 20 bytes
IP header = 20 bytes



IP Address: Globally unique* 32 bit (4 byte) **dotted decimal** number assigned to interfaces on hosts and routers

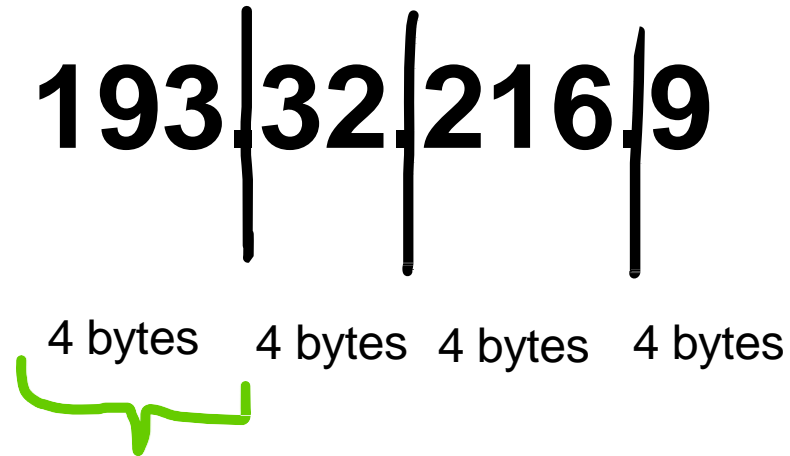
193.32.216.9

IP Address: Globally unique* 32 bit (4 byte) **dotted decimal** number assigned to interfaces on hosts and routers

(1 byte = 8 bits)

193|32|216|9

4 bytes 4 bytes 4 bytes 4 bytes



128	64	32	16	8	4	2	1
1	1	0	0	0	0	0	1

$$128 + 64 + 1 = \mathbf{193}$$

193 = 11000001 in binary

IP Address: Globally unique* 32 bit (4 byte) **dotted decimal** number assigned to interfaces on hosts and routers

(1 byte = 8 bits)

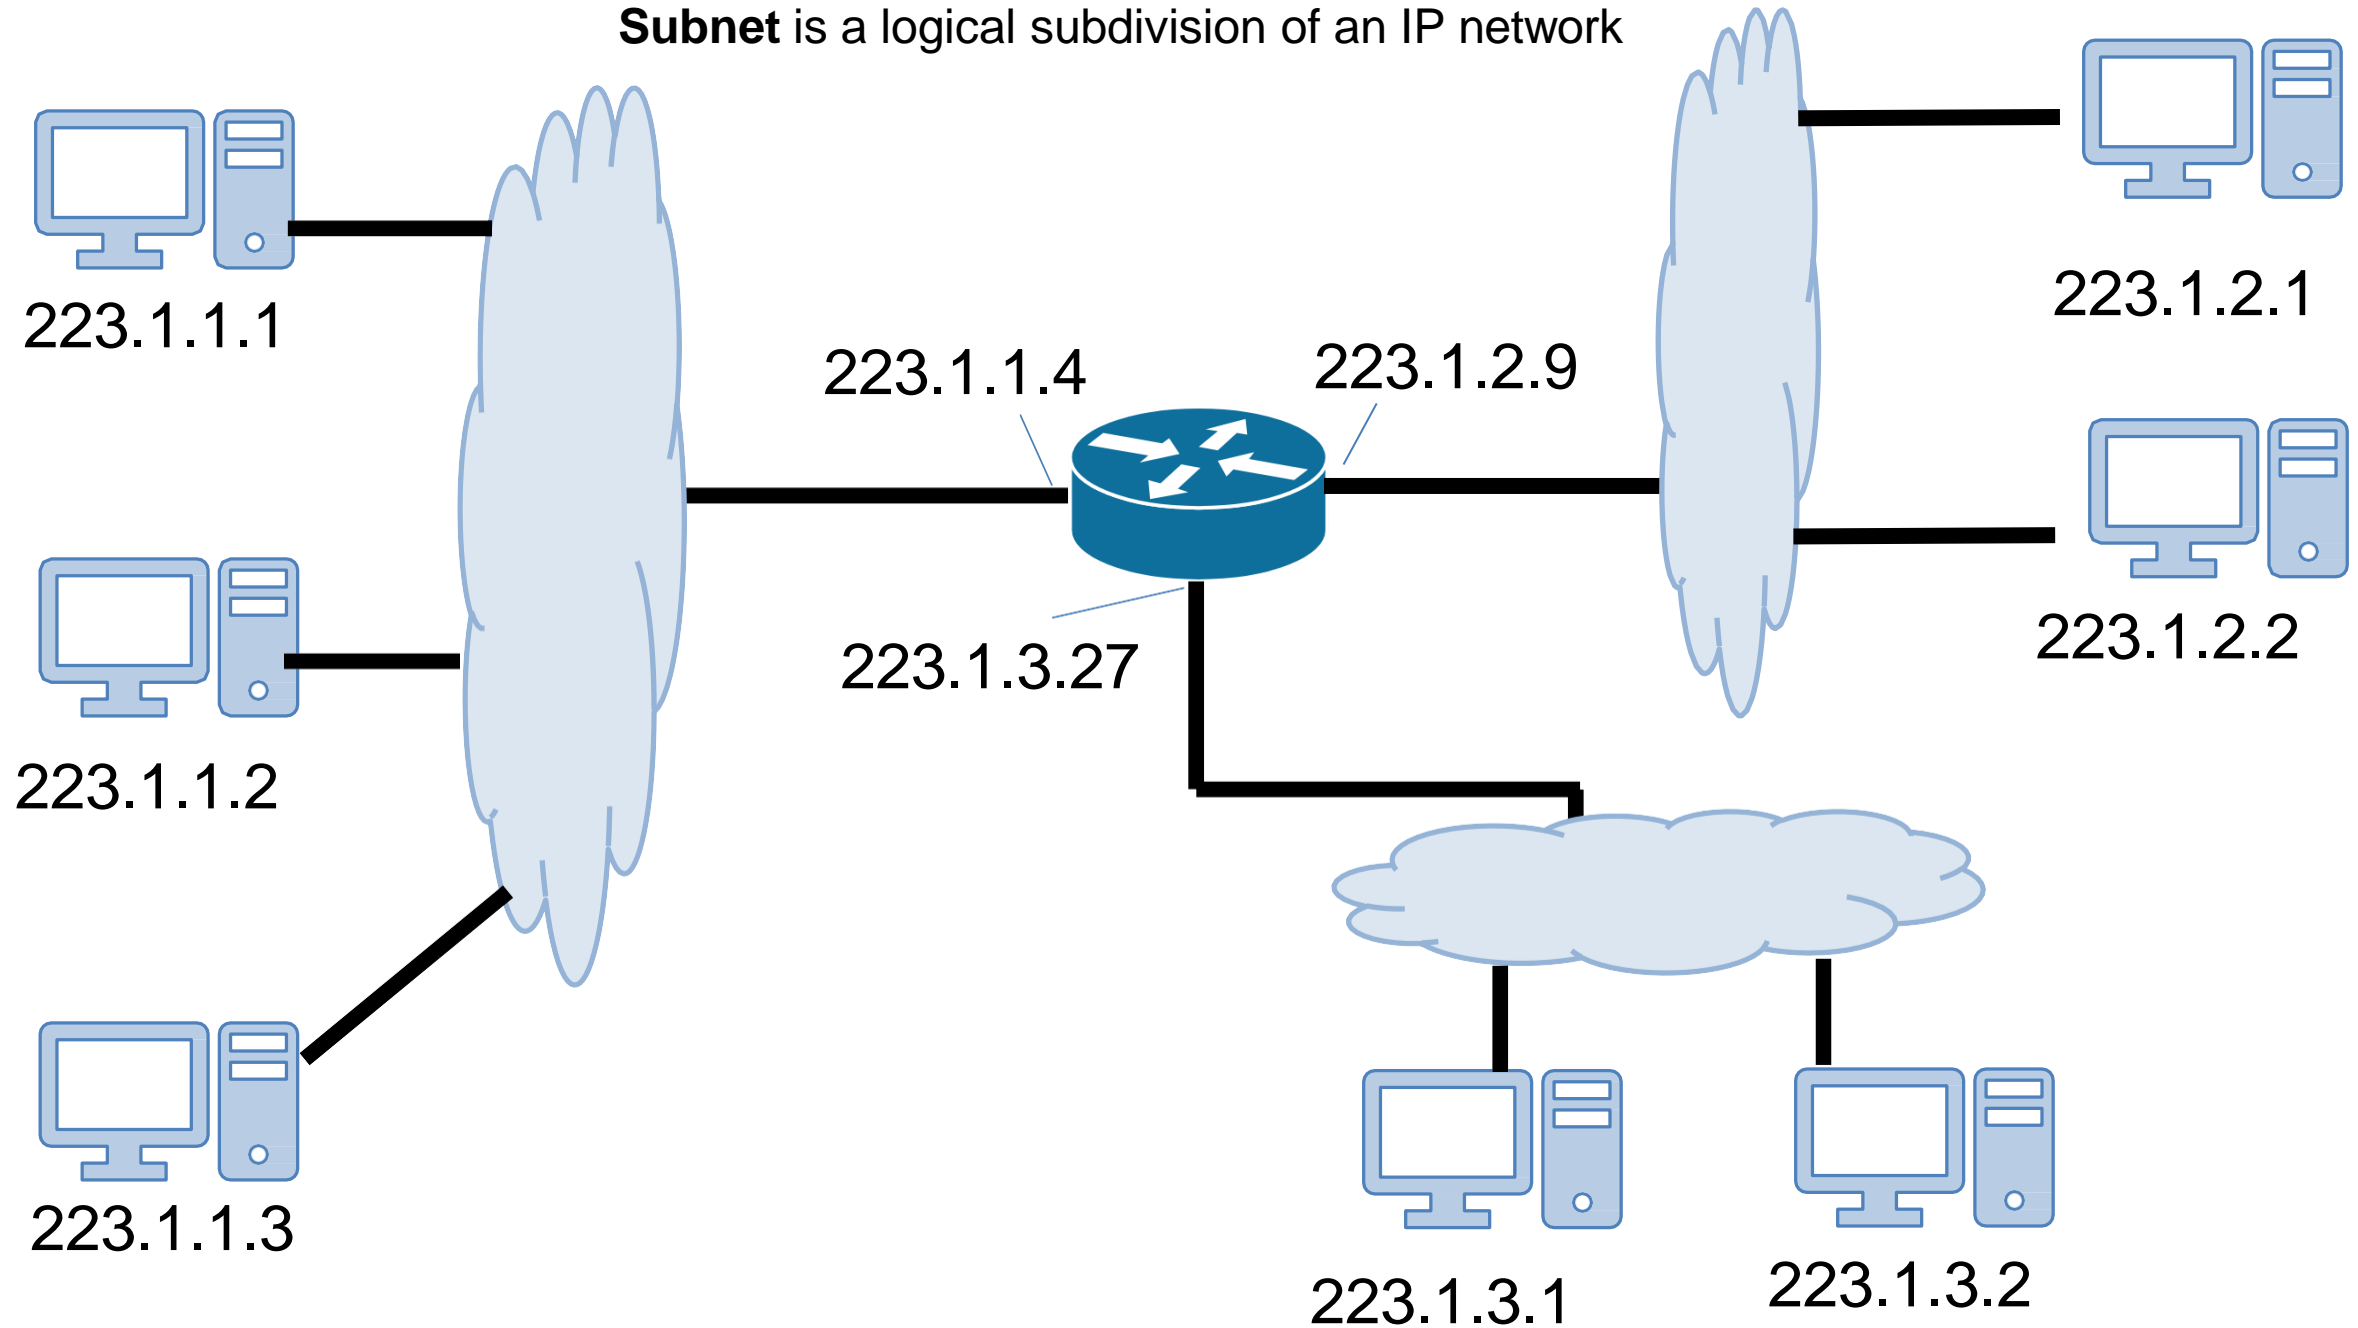
193.32.216.9

=

11000001 00100000 11011000 00001001

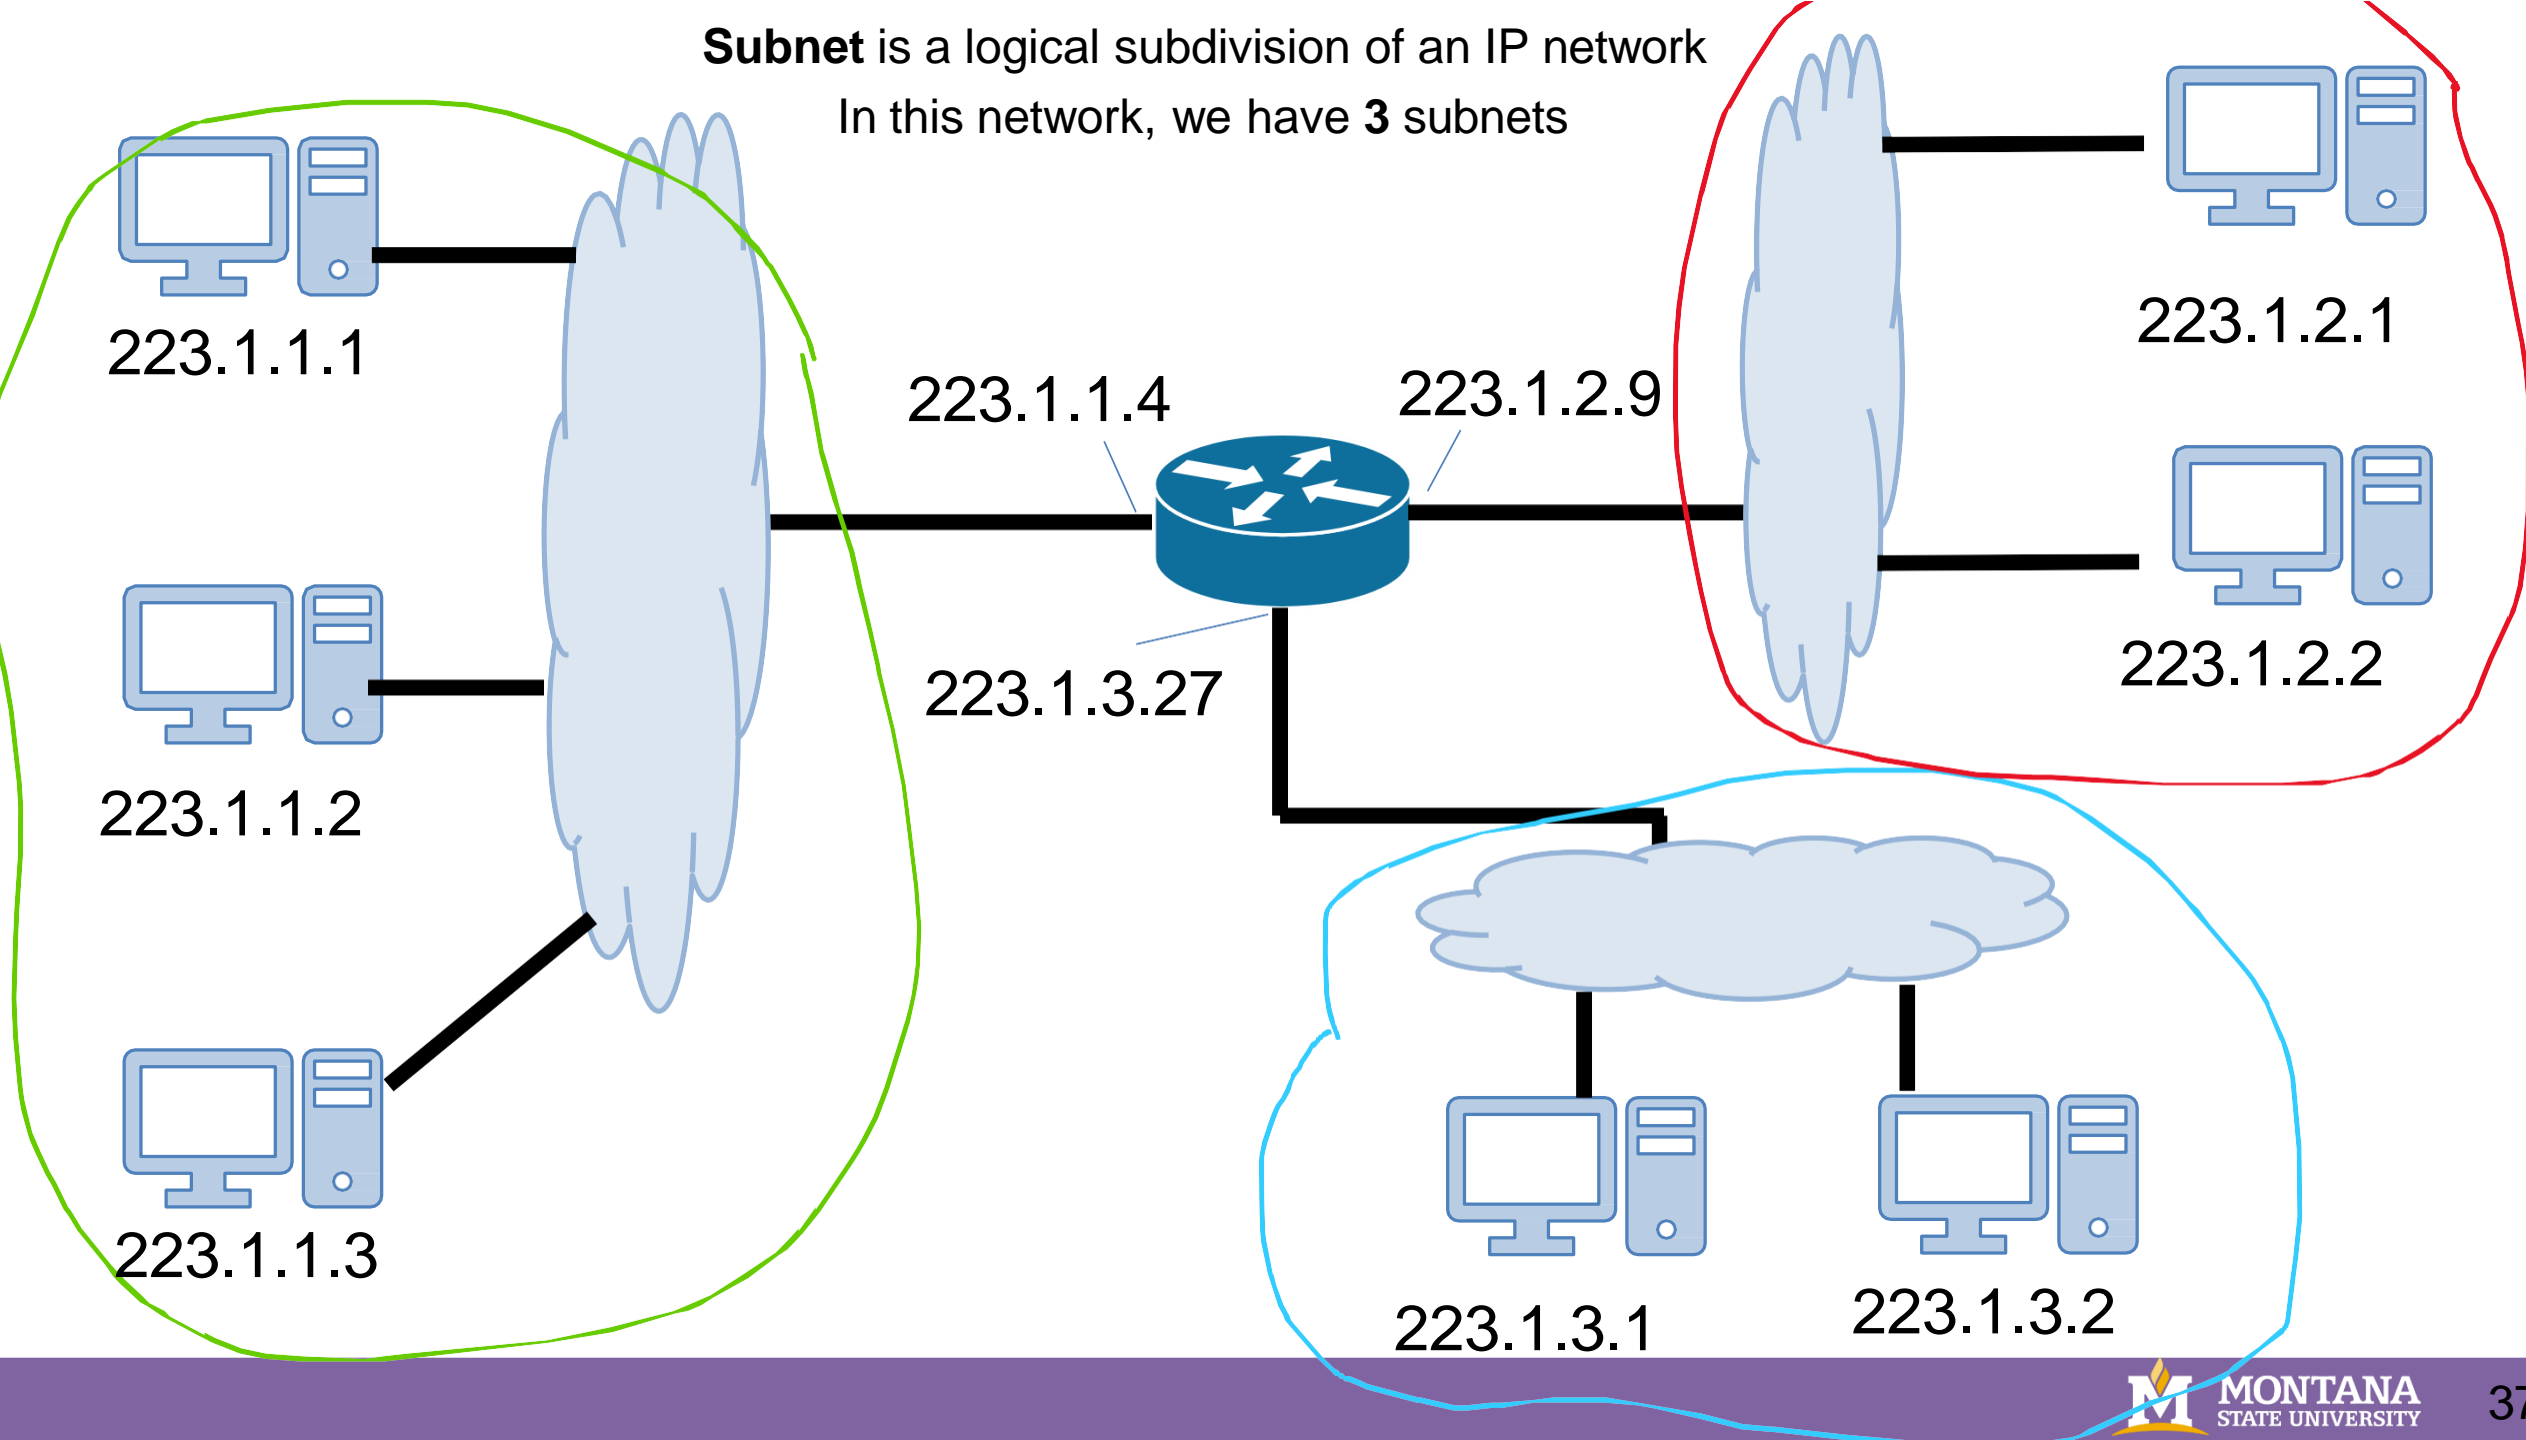
How many addresses are possible with a 32 bit number? ~4 billion possible IPv4 addresses

Subnet is a logical subdivision of an IP network



Subnet is a logical subdivision of an IP network

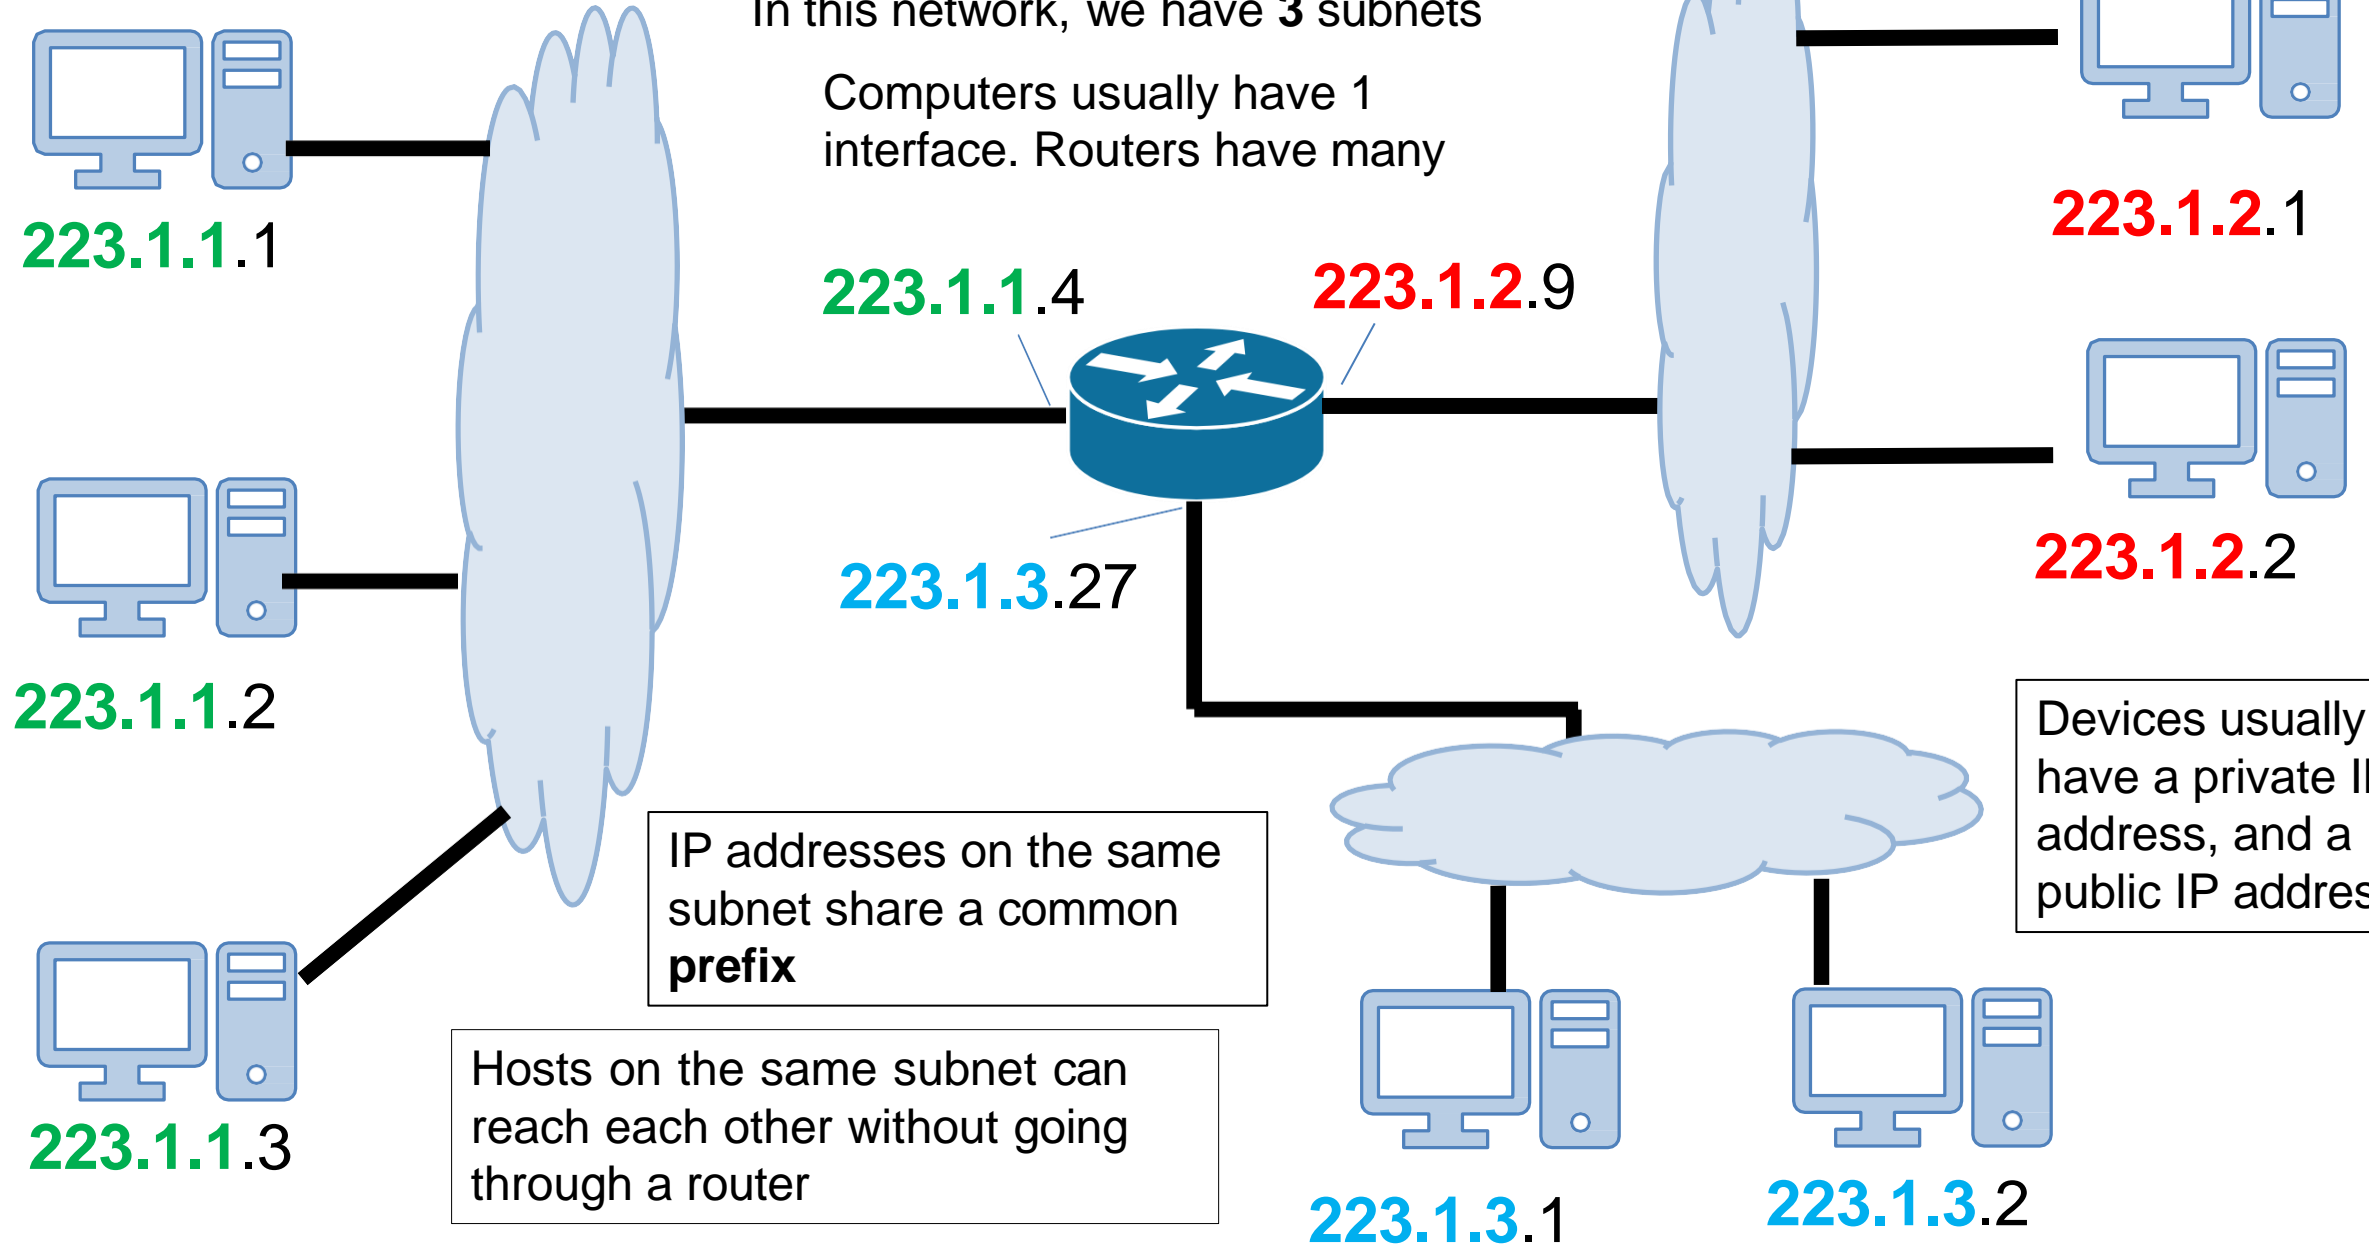
In this network, we have **3** subnets




Subnet is a logical subdivision of an IP network

In this network, we have **3** subnets

Computers usually have 1 interface. Routers have many



It is very common to have a **range** of IP addresses assigned to you (random assignment would be chaos)



Subnet mask

223.1.1.0 / 24

The leftmost 24 bits represent the prefix of the subnet

11111111 11111111 11111111 XXXXXXXX = 255.255.255.0

223.1.1.67 ✓

223.2.1.67 ✗

223.2.1.255 ✗

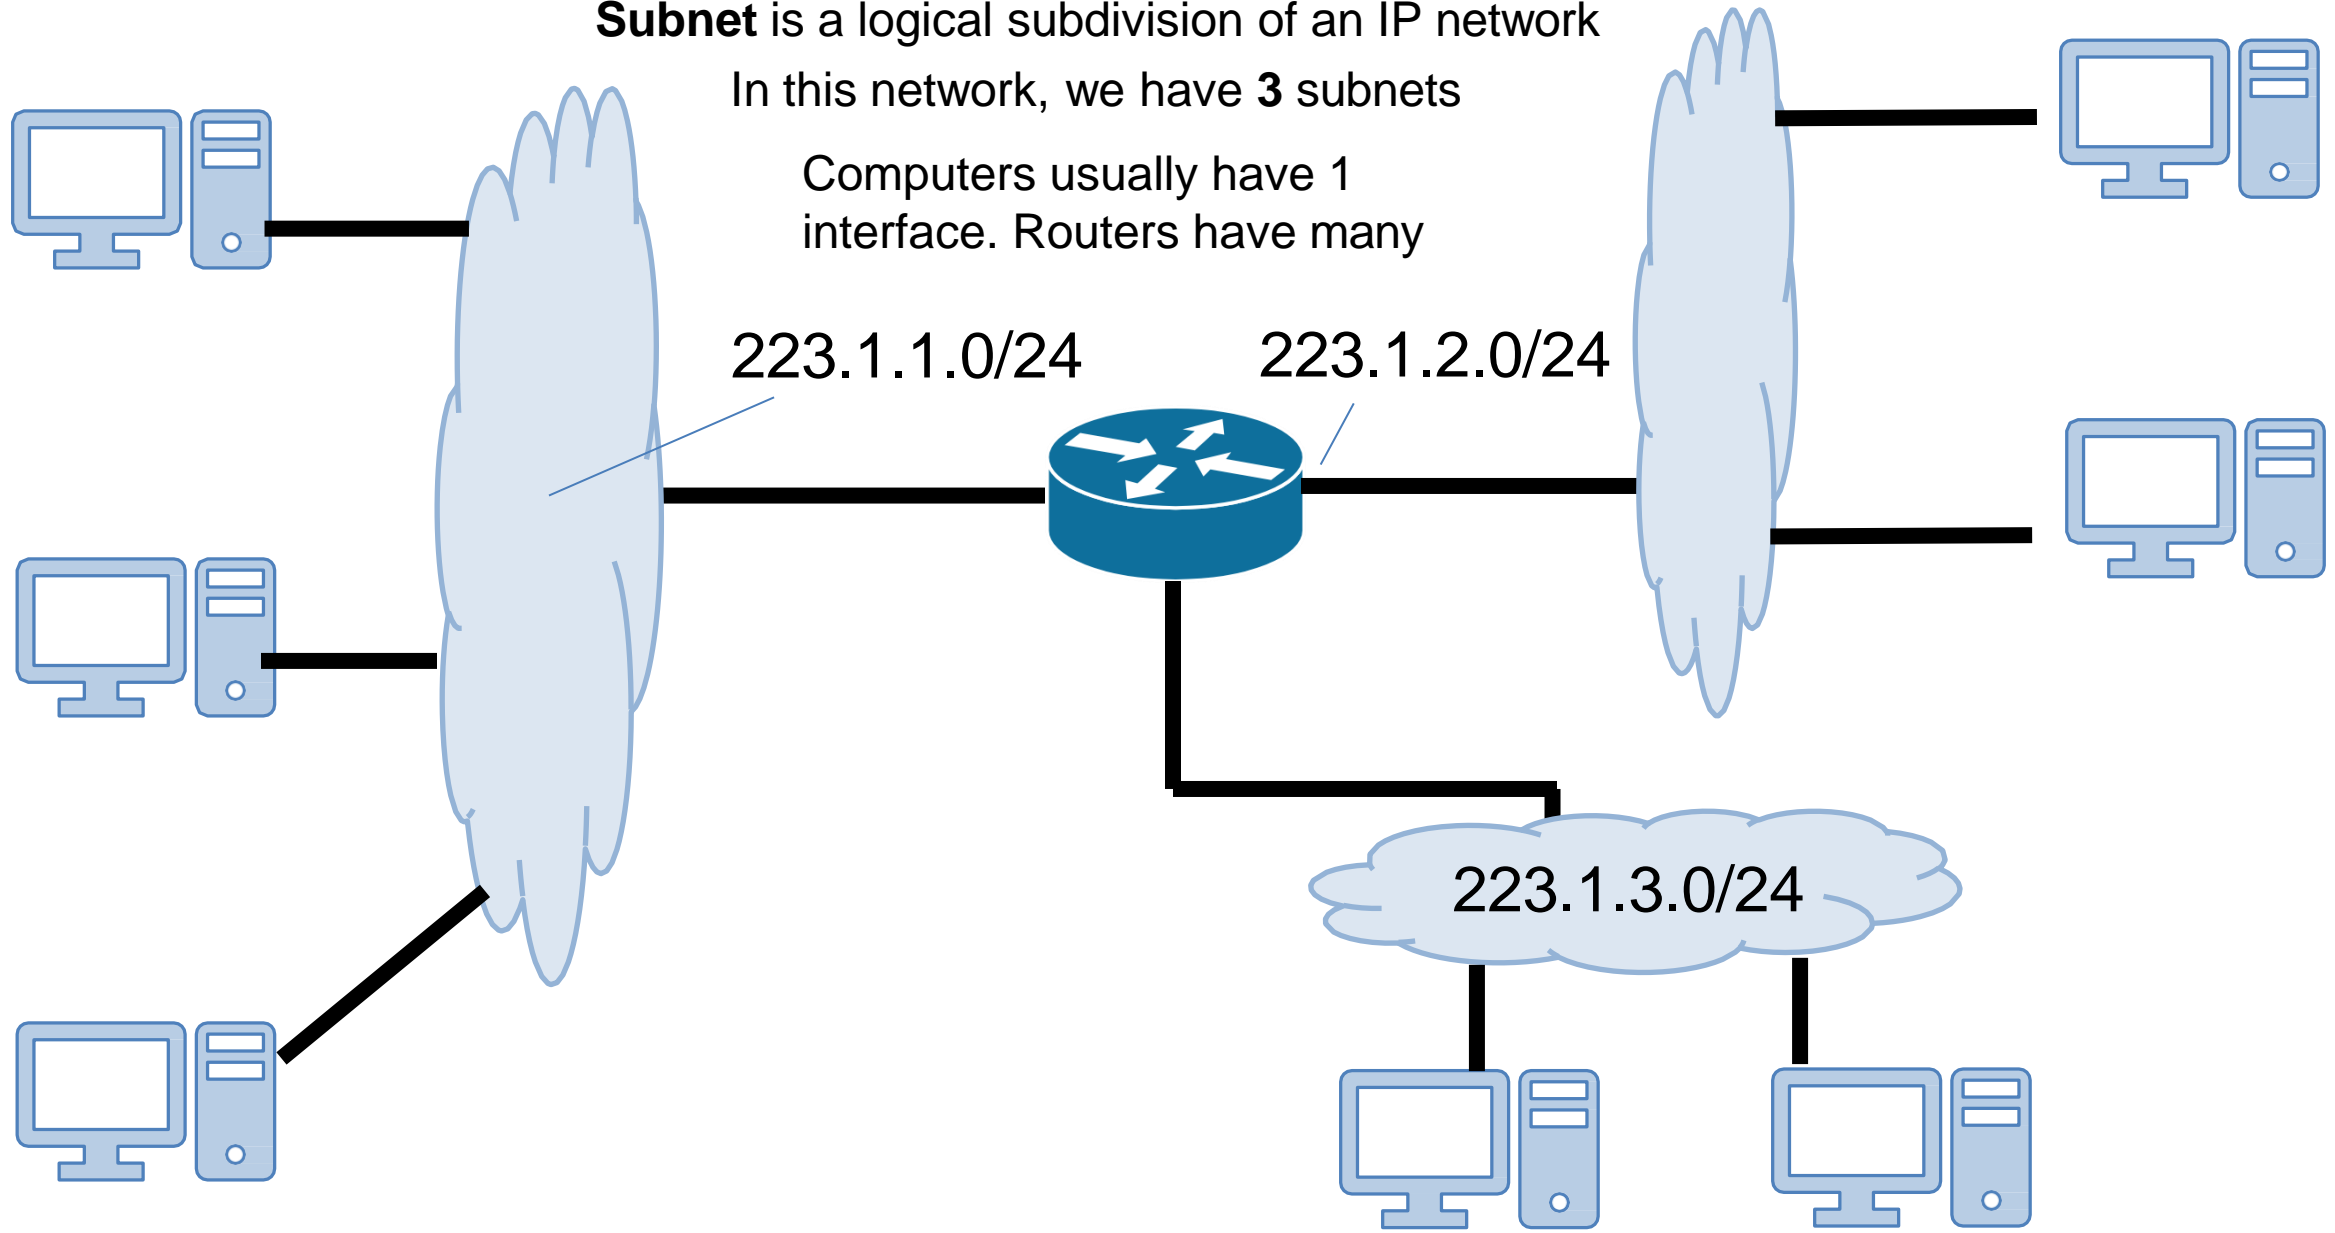
Network bits **193.32.216.9** **Host bits**

11000001 00100000 11011000 00001001

Subnet is a logical subdivision of an IP network


In this network, we have **3** subnets

Computers usually have 1 interface. Routers have many



6 subnets

1. 223.1.1.0/24
2. 223.1.7.0/24
3. 223.1.3.0/24
4. 223.1.8.0/24
5. 223.1.2.0/24
6. 223.1.9.0/23



Managing and Assigning Subnets

The subnet `223.1.1.0/24` contains the following addresses

`223.1.1.1`

`223.1.1.2`

`223.1.1.3`

`223.1.1.4`

`223.1.1.5`

...

`223.1.1.252`

`223.1.1.253`

`223.1.1.254`

How many addresses does a /24 subnet provide?

0 and 255 are reserved for special services (?)

Managing and Assigning Subnets

The subnet `223.1.1.0/24` contains the following addresses

223.1.1.1

223.1.1.2

223.1.1.3

223.1.1.4

223.1.1.5

...

223.1.1.252

223.1.1.253

223.1.1.254

How many addresses does a /24 subnet provide?

$$2^8 - 2 = \mathbf{253}$$



Class C network

0 and 255 are reserved for special services (?)

Managing and Assigning Subnets

Subnet Type	Slash	Subnet Mask	Available Addresses
Class A	/8	255.0.0.0	$2^{24} - 2 = 16777216$
Class B	/16	255.255.0.0	$2^{16} - 2 = 65,634$
Class C	/24	255.255.255.0	$2^8 - 2 = 254$

Issues with this type of assignment?

Managing and Assigning Subnets

Subnet Type	Slash	Subnet Mask	Available Addresses
Class A	/8	255.0.0.0	$2^{24} - 2 = 16777216$
Class B	/16	255.255.0.0	$2^{16} - 2 = 65,634$
Class C	/24	255.255.255.0	$2^8 - 2 = 254$

If I have 2000 devices, I will either need 8 class C subnets or 1 class B subnet (but waste 63,000ish IP addresses)

We need a better way to subnet!

Managing and Assigning Subnets

Our only available subnet masks have been:

255.255.255.0	11111111 11111111 11111111 00000000	/24
255.255.0.0	11111111 11111111 00000000 00000000	/16
255.0.0.0	11111111 00000000 00000000 00000000	/8

CIDR introduces a more flexible way for subnetting

We can any number of bits for our mask

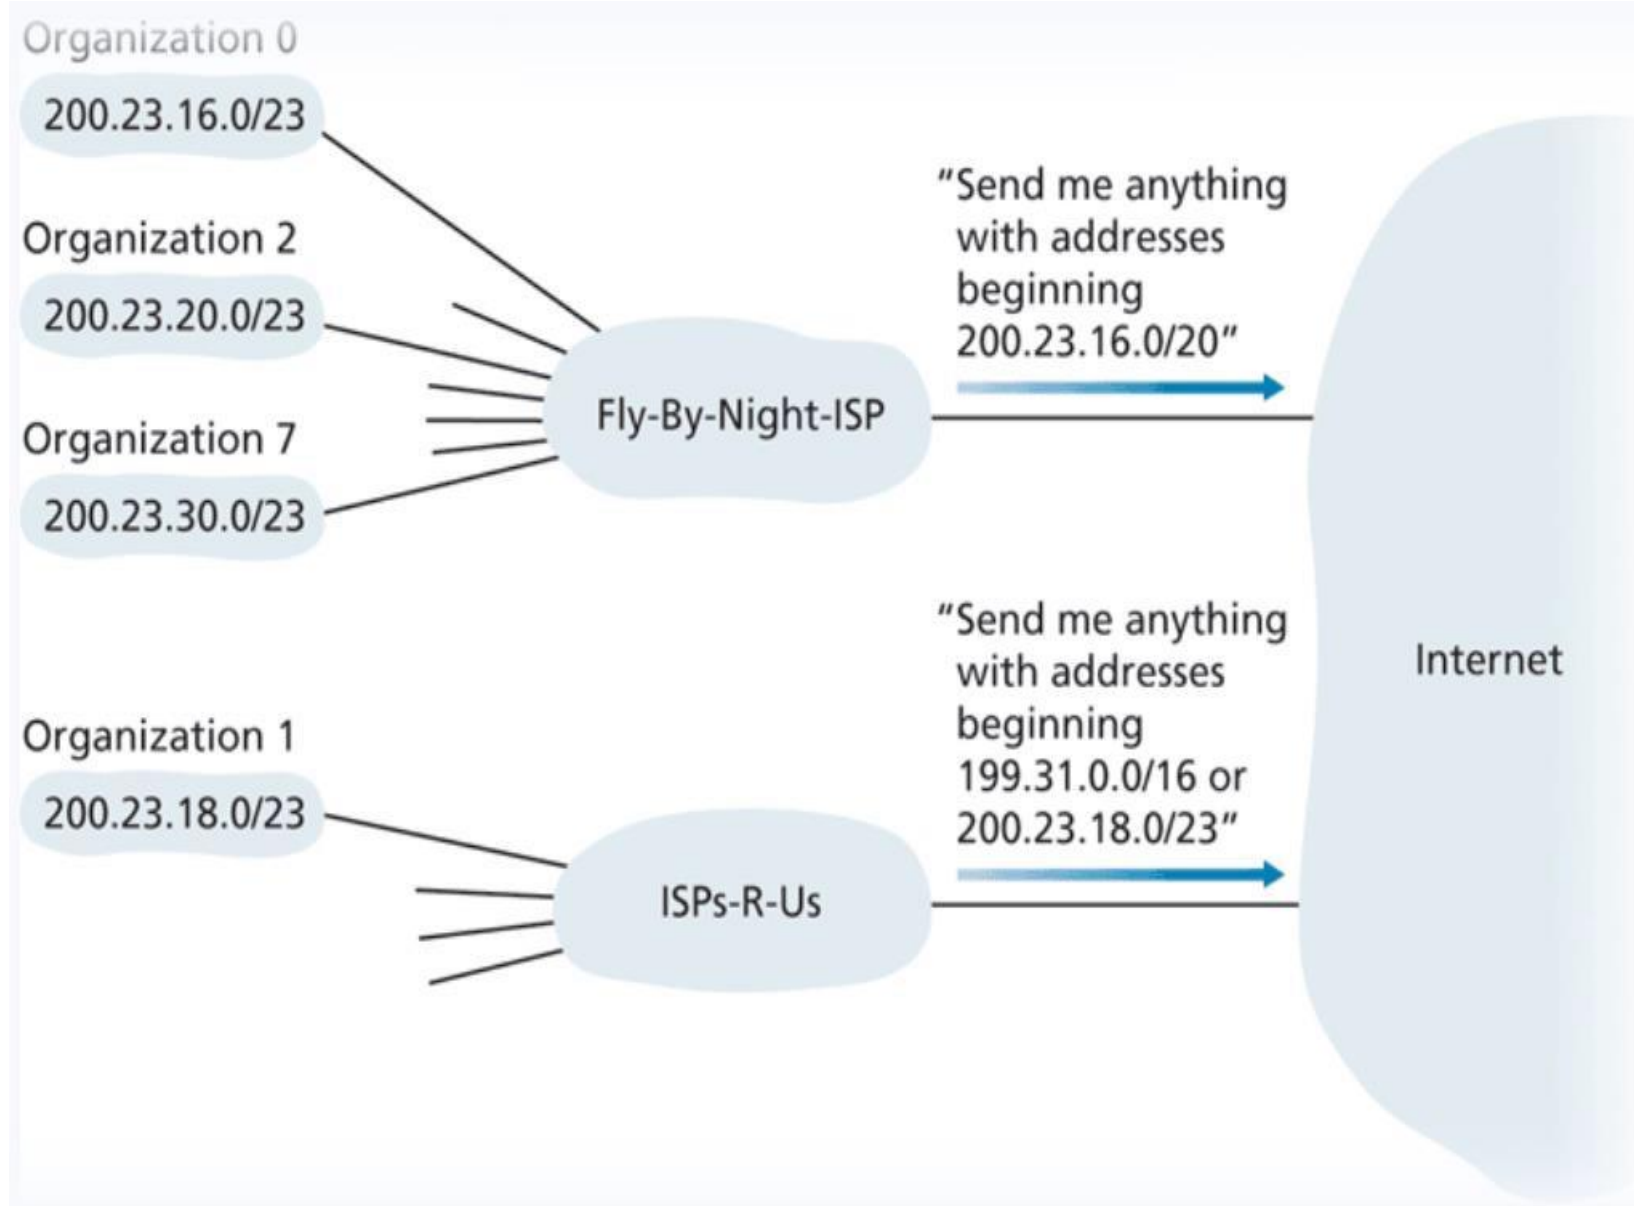
Host bits!

200.23.16.0/20 = 11001000.00010111.00010000.00000000

Mask = 11111111.11111111.11110000.00000000

ISPs can now group and advertise organizations by IP blocks

We once again have another hierarchy ☺




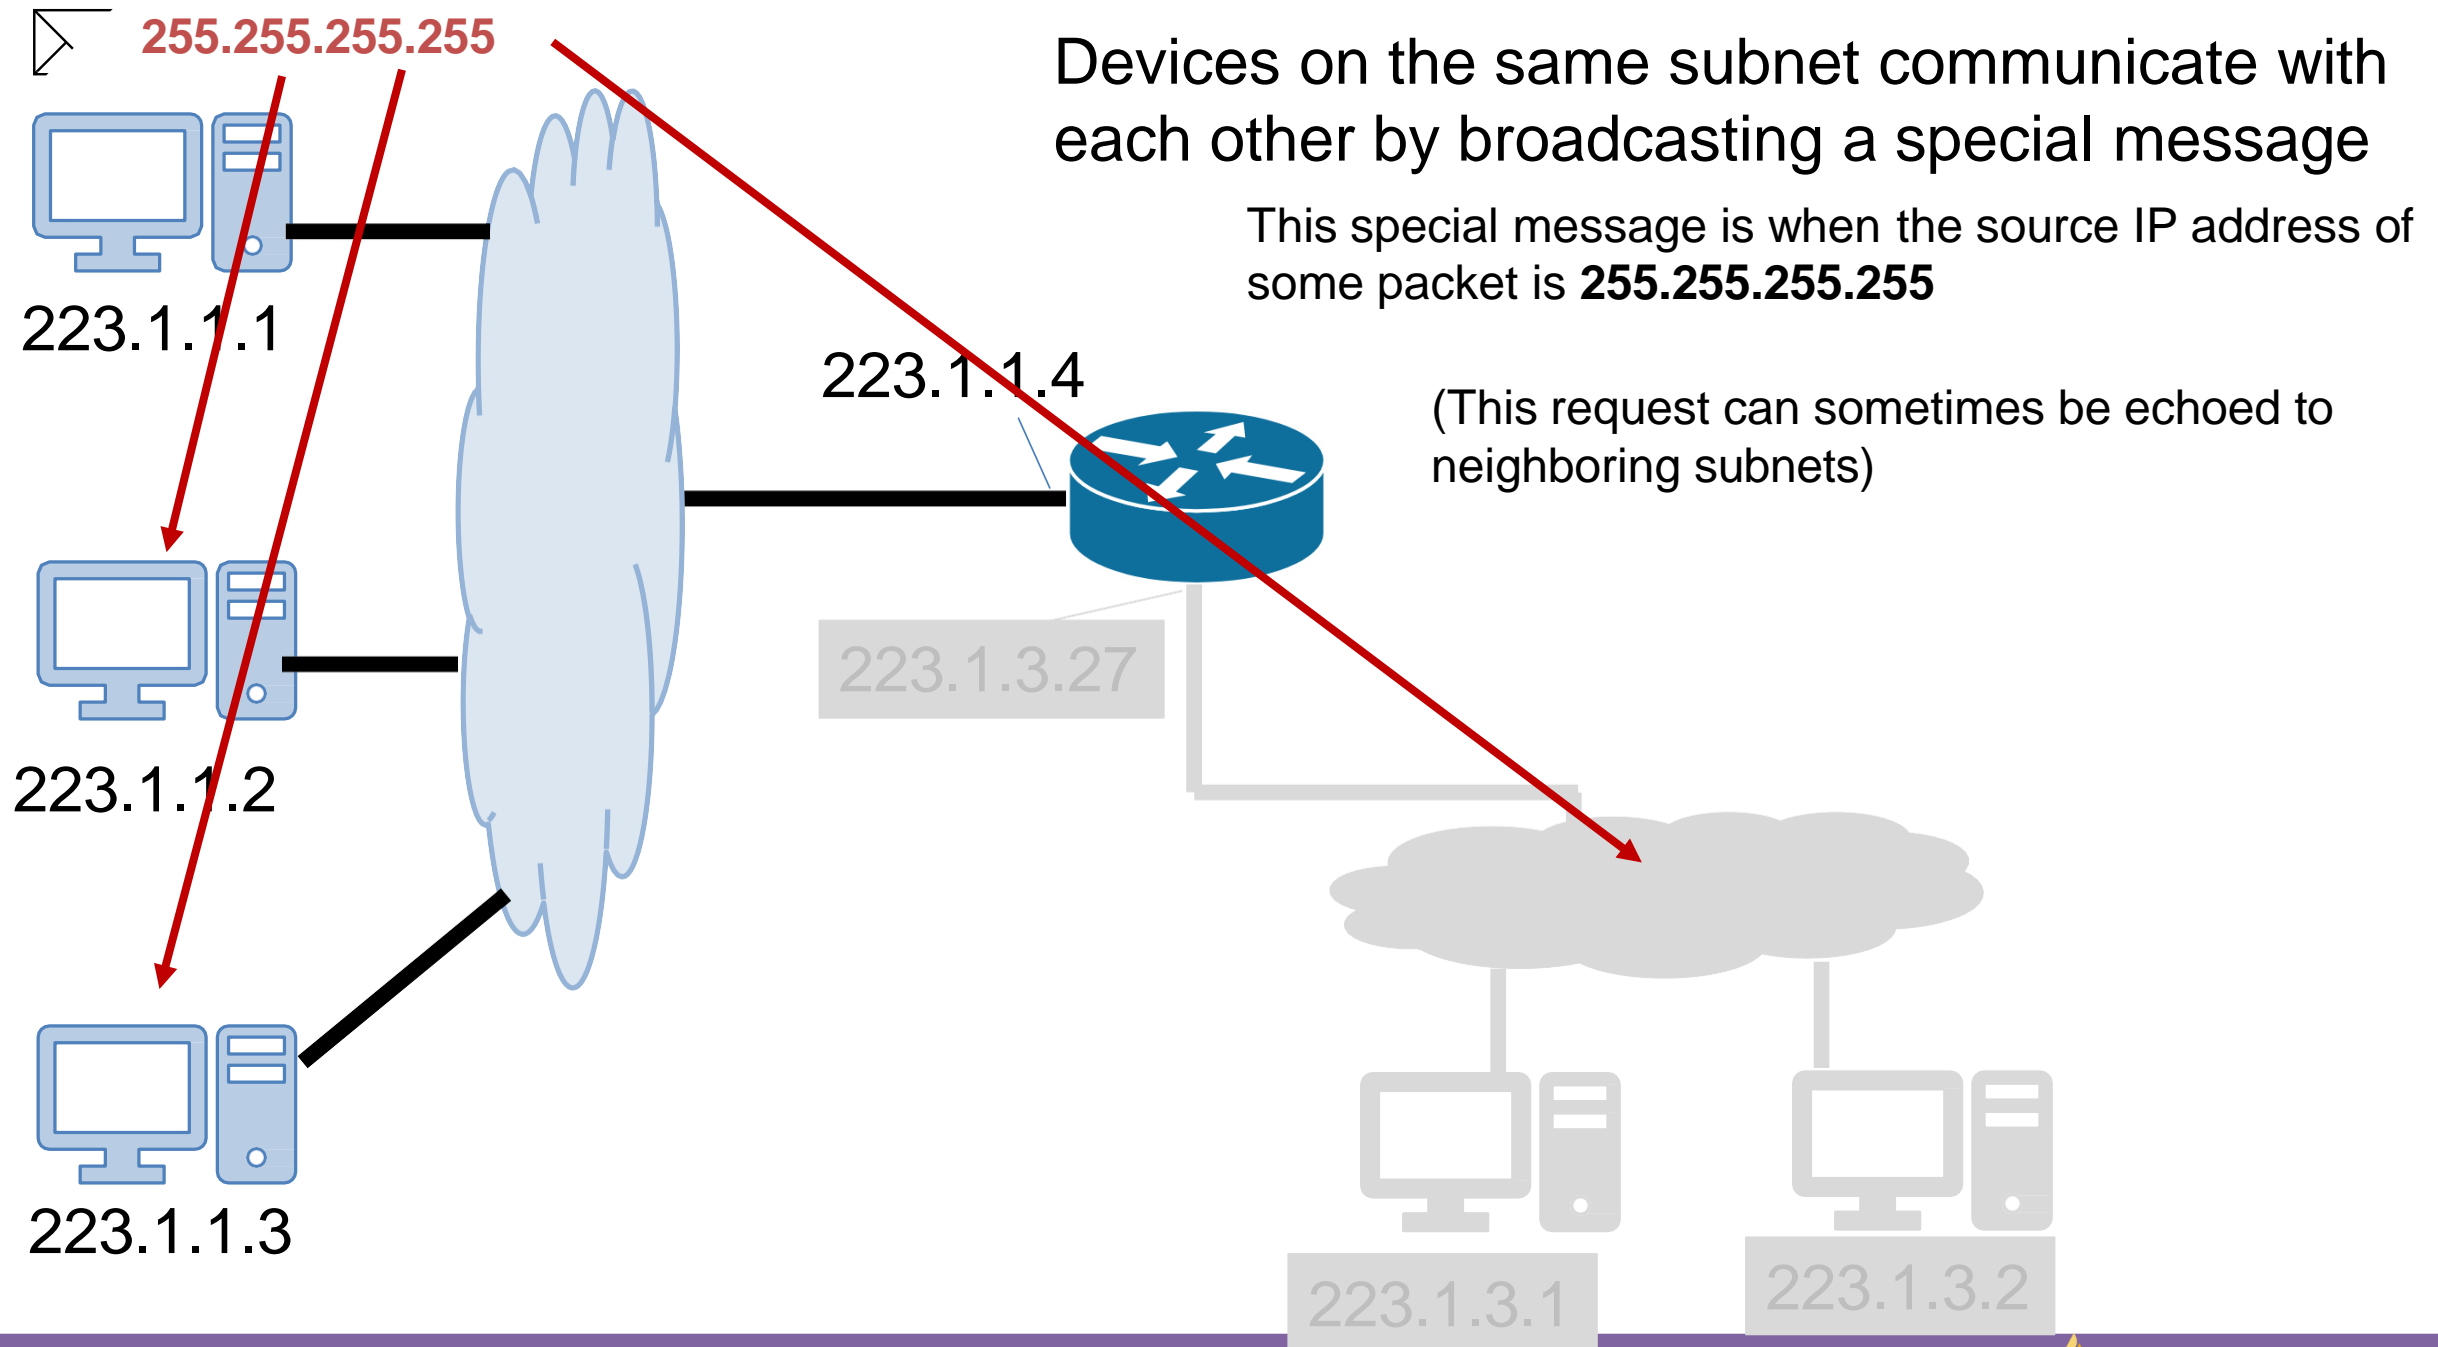
One possible way an ISP could organize IP blocks

ISP's block	200.23.16.0/20	<u>11001000 00010111 0001</u> 0000 00000000
Organization 0	200.23.16.0/23	<u>11001000 00010111 0001</u> 0000 00000000
Organization 1	200.23.18.0/23	<u>11001000 00010111 0001</u> 0010 00000000
Organization 2	200.23.20.0/23	<u>11001000 00010111 0001</u> 0100 00000000
...		
Organization 7	200.23.30.0.23	<u>11001000 00010111 0001</u> 1110 00000000

Devices on the same subnet communicate with each other by broadcasting a special message

This special message is when the source IP address of some packet is **255.255.255.255**





How do IPs get obtained/assigned?

The Internet Corporation for Assigned Names and Numbers (ICANN) is responsible for managing and allocating IP address space for ISPs and organizations

(they also manage the DNS root servers!)

When an organization gets a range of IP addresses to use, how to we give assign them to devices?

Do we do it manually?


we could....

Dynamic Host Configuration Protocol (DHCP) is a **plug-and-play**, client-server protocol that allows a host to obtain an IP address automatically

When a host is automatically assigned an IP address, it might keep that one forever, or the IP addresses can be temporary

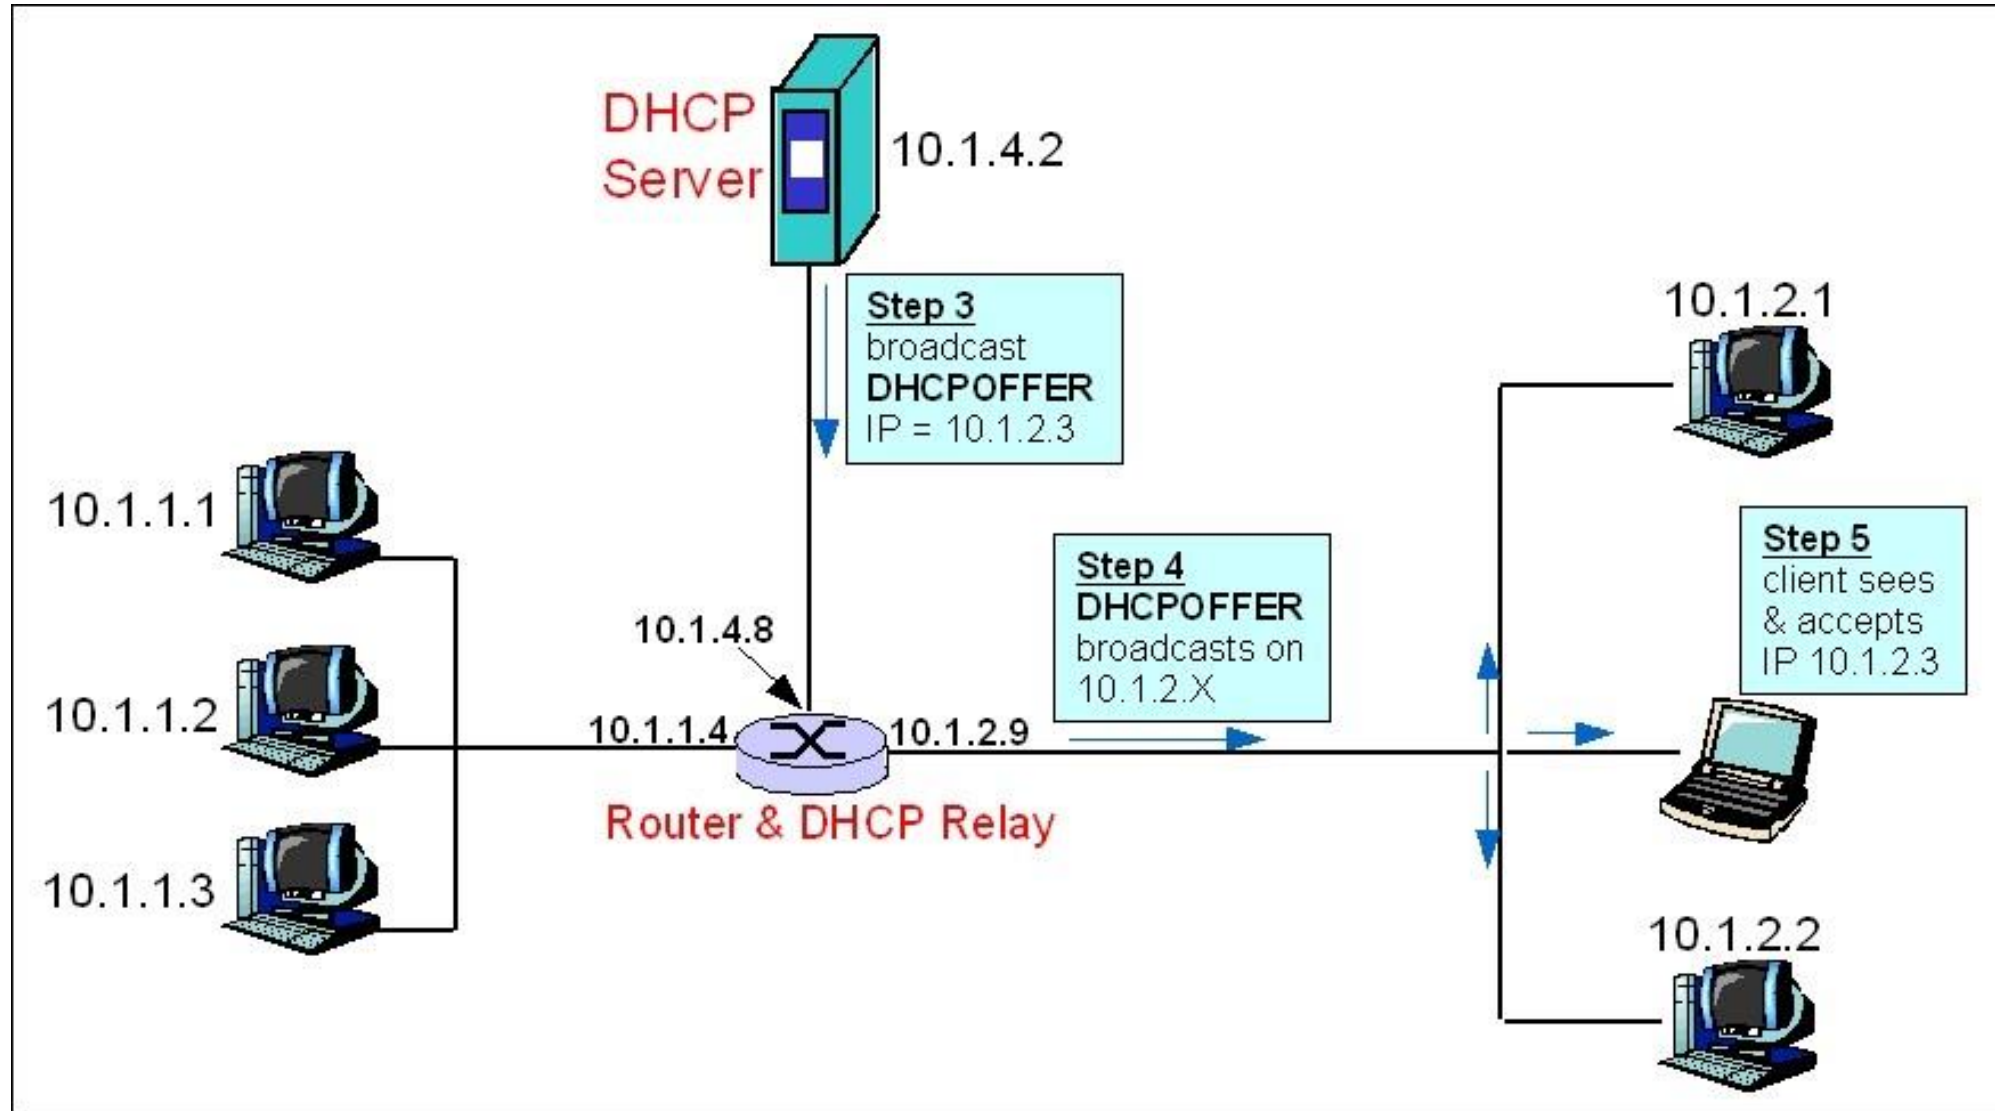
(more common)

This process is similar to a TCP handshake!




But devices that are **not** the dhcp server will ignore it

This process is similar to a TCP handshake!



This process is similar to a TCP handshake!




Announcements

If I haven't responded to a DM/email, please poke me about it

PA2 due on Friday 10/21

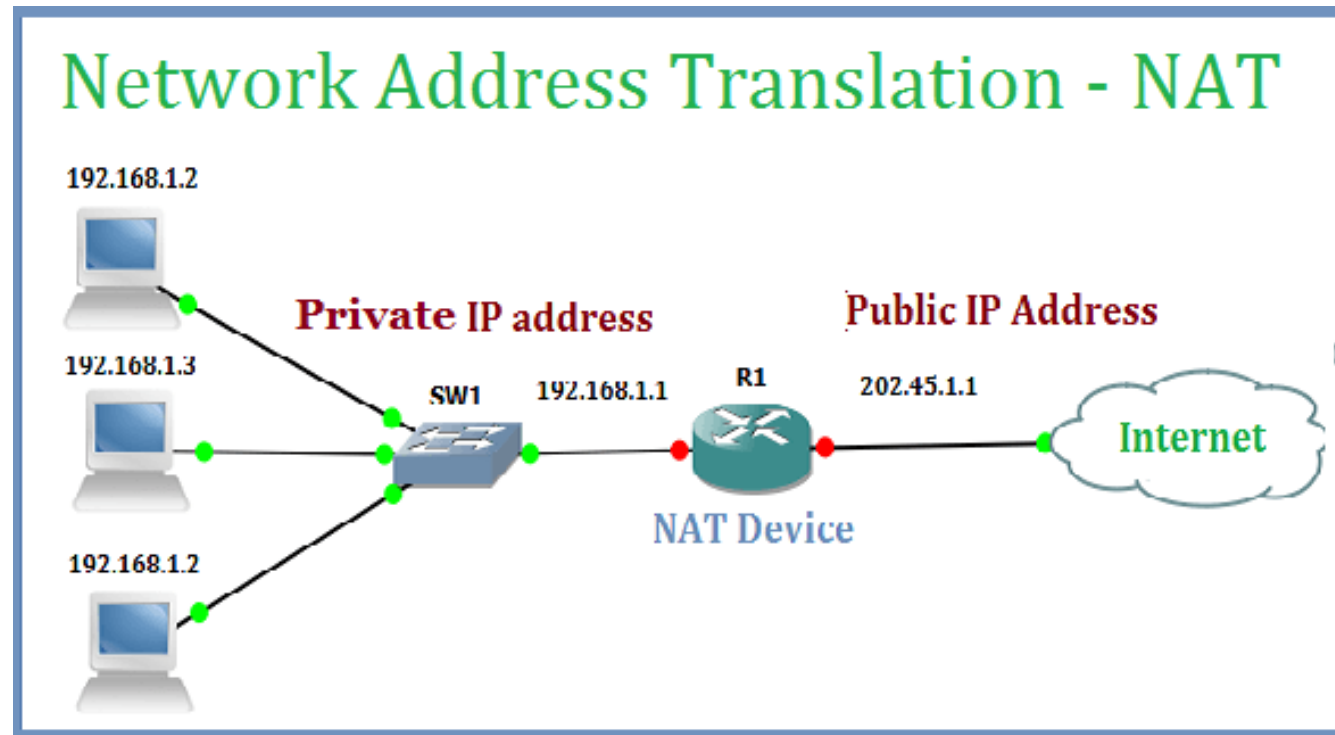
- I will allow you to submit it any time this weekend with no late penalty (no late pass required!)
- Lecture will be asynchronous (recorded) on Monday

PA2: Flow Demo, FIN ACKs



NAT (Network Address Translation)

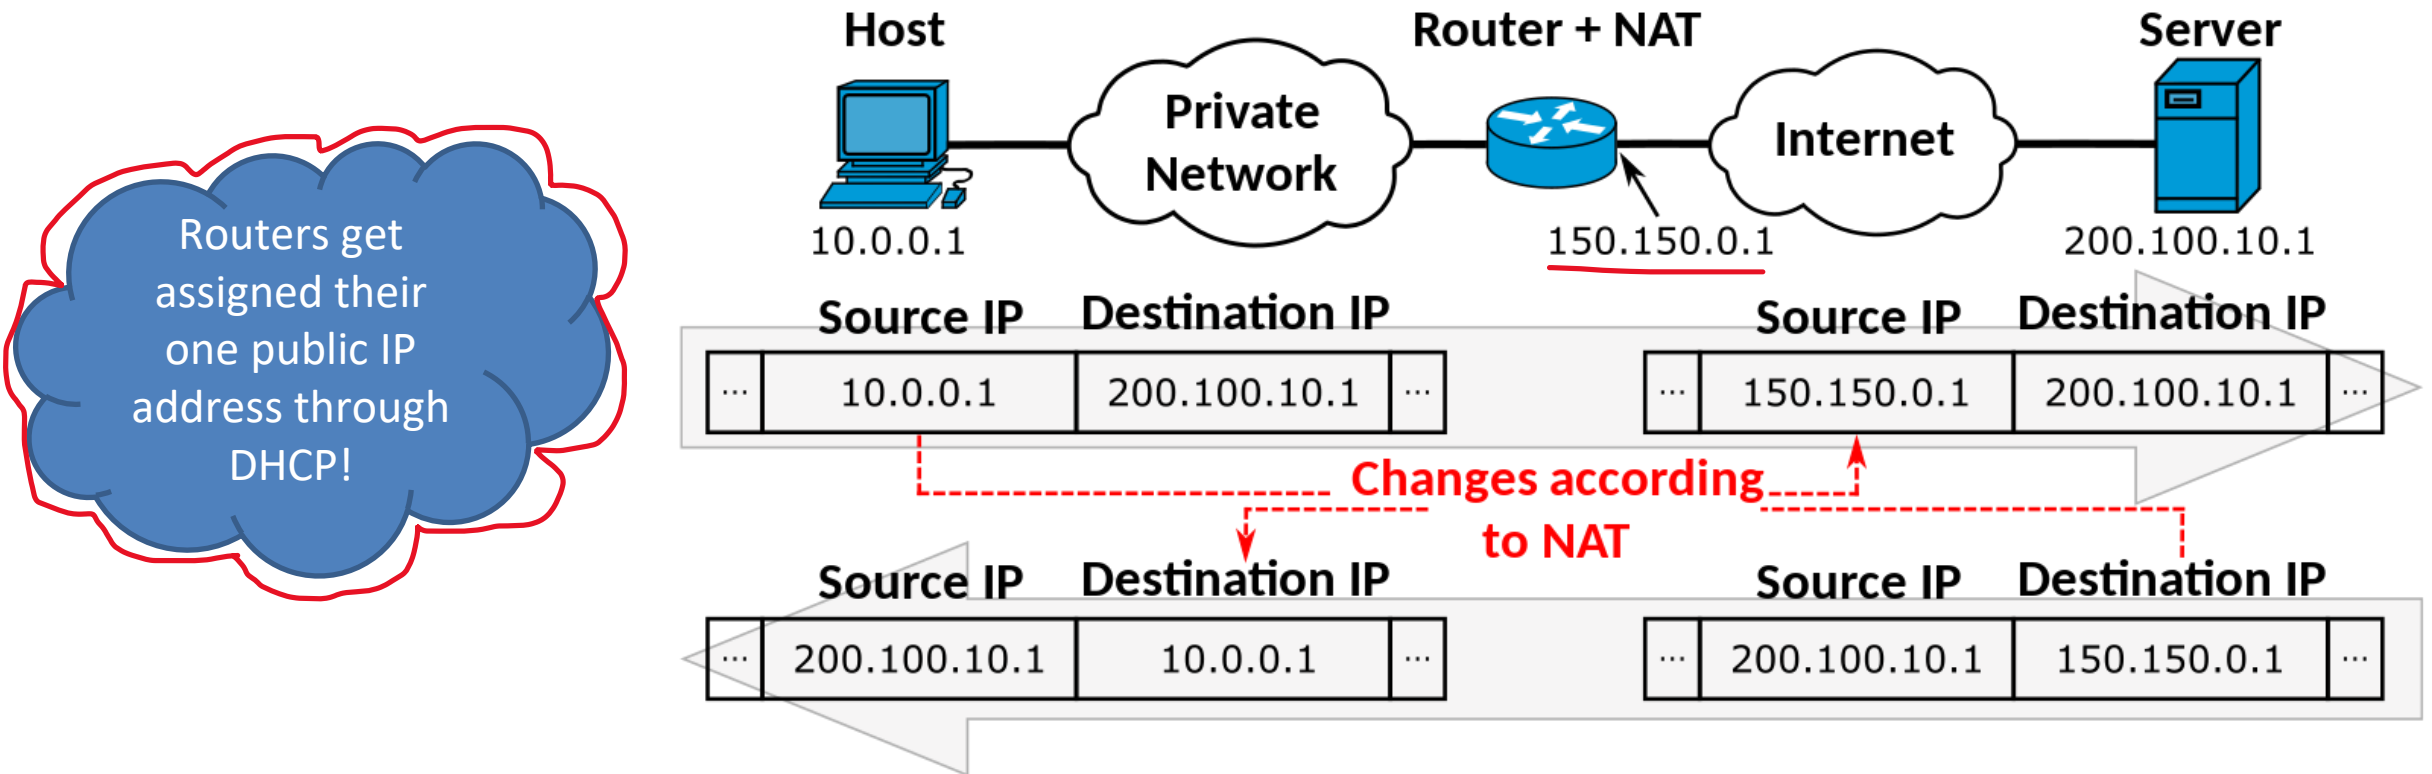
NAT is a translation of multiple private IP addresses to one single public IP address



NAT (Network Address Translation)

NAT is a translation of multiple private IP addresses to one single public IP address

- Hides details of inner home network from outside world
- All incoming traffic will have same public IP, all outgoing will have same public IP



Packets traversing through the network layer are referred to as a **datagram**. Each packet gets an IPv4/IPv6 header



Version	Traffic class ✖	Flow label ✖
Payload length	Next header	Hop limit
Source address (128 bits)		
Destination address (128 bits)		


IPv6 = **128 bits** (not 64)

IPv6 Header


The goal is for our internet to be only using IPv6 in the near future? **Issues?**

Tunneling

logical view:



physical view:



So far, a router takes input from input port, and then outputs on some output port

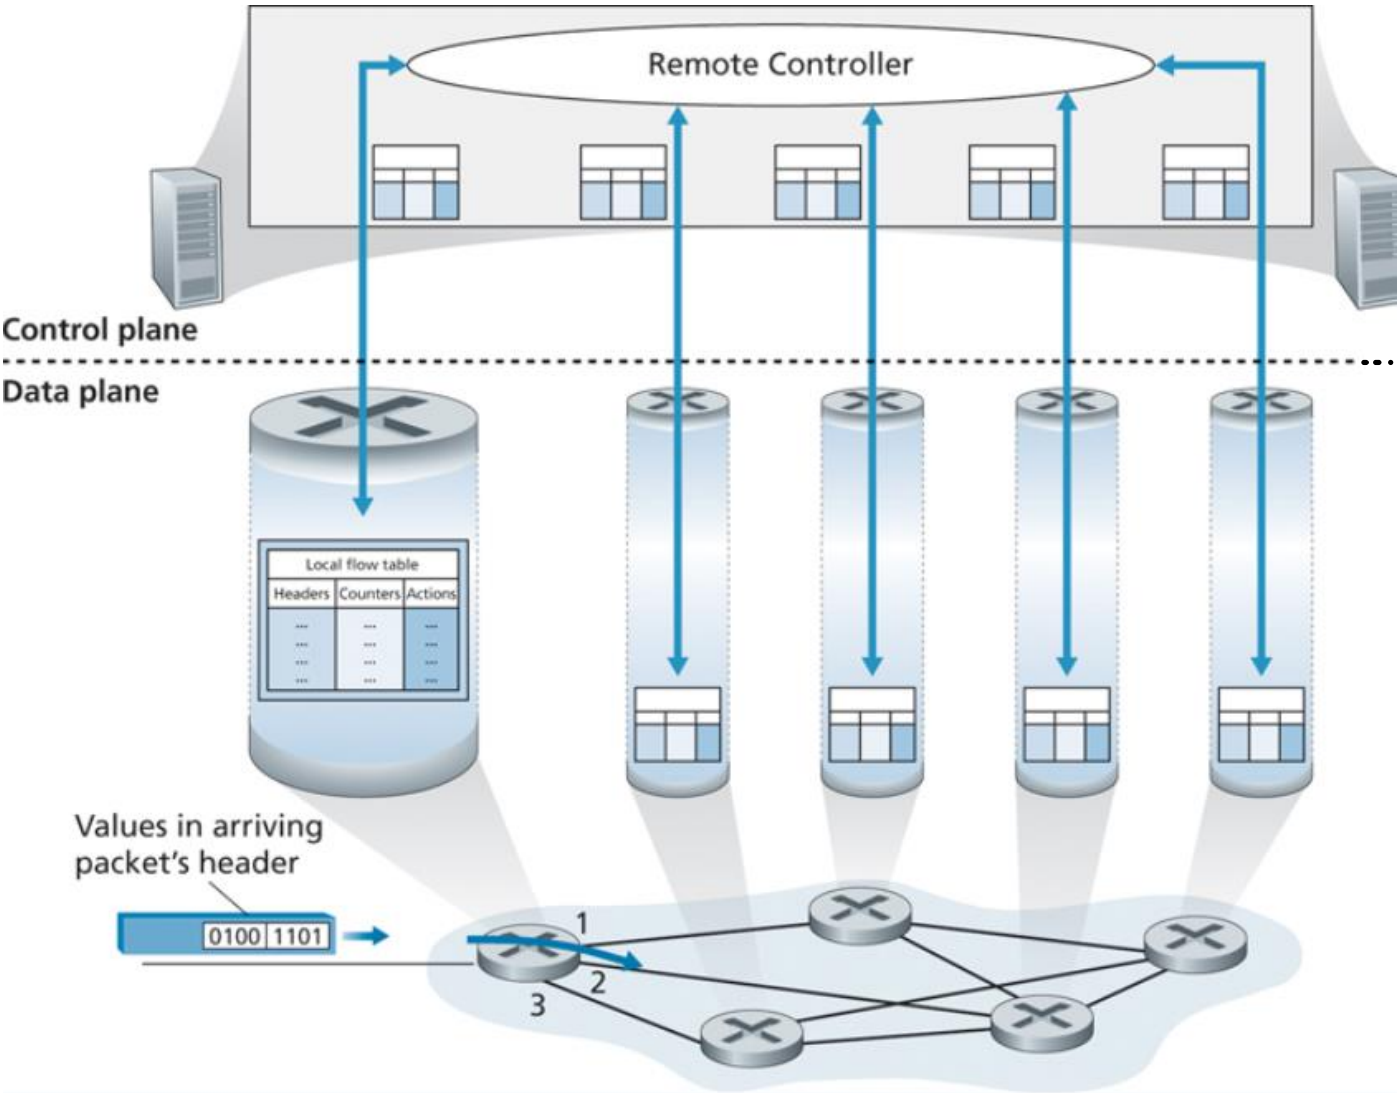
What else might a router need to do?

So far, a router takes input from input port, and then outputs on some output port

What else might a router need to do?

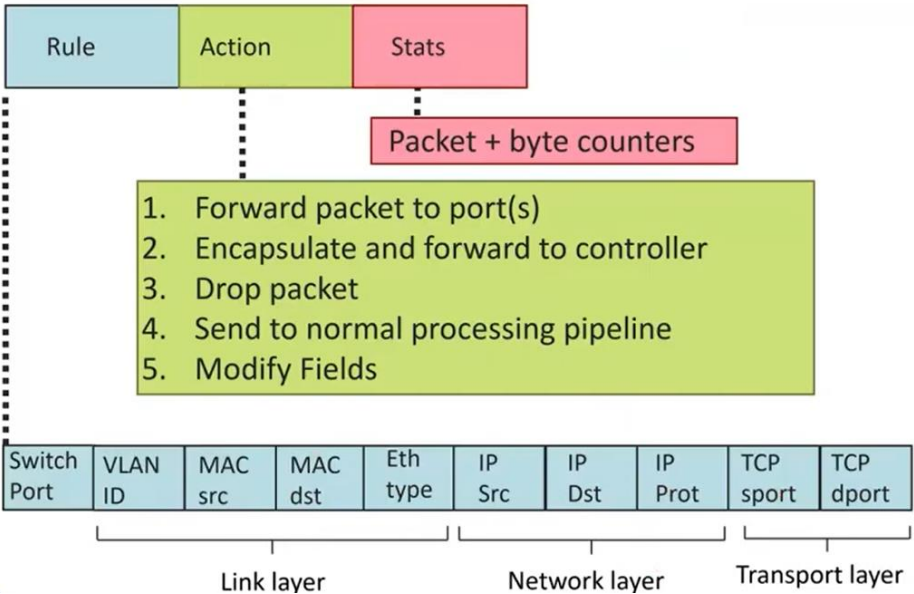
Forward, Drop stuff, Modify, Load balance

We need more flexibility and functionality with our forwarding!!




We need **headers/rules**, which are going to the values the remote controller is going to evaluate

We need **actions** to do based on some pattern match



Generalized Forwarding and Software Defines Network (SDN)




Match	Action
Ingress Port = 1 ; IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	<u>Forward(4)</u>

Match	Action
IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	Forward(3)

Match	Action
Ingress port = 2 ; IP Dst = 10.2.0.3	Forward(3)
Ingress port = 2 ; IP Dst = 10.2.0.4	Forward(4)

Generalized Forwarding and Software Defines Network (SDN)




Match	Action
Ingress Port = 1 ; IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	Forward(4)

Match	Action
IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	<u>Forward(3)</u>

Match	Action
Ingress port = 2 ; IP Dst = 10.2.0.3	Forward(3)
Ingress port = 2 ; IP Dst = 10.2.0.4	Forward(4)

Generalized Forwarding and Software Defines Network (SDN)




Match	Action
Ingress Port = 1 ; IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	Forward(4)

Match	Action
IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	<u>Forward(3)</u>

Match	Action
Ingress port = 2 ; IP Dst = 10.2.0.3	Forward(3)
Ingress port = 2 ; IP Dst = 10.2.0.4	Forward(4)

Generalized Forwarding and Software Defines Network (SDN)



Match	Action
Ingress Port = 1 ; IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	Forward(4)

Match	Action
IP Src = 10.3.*.* ; IP Dst = 10.2.*.*	Forward(3)

Match	Action
Ingress port = 2 ; IP Dst = 10.2.0.3	Forward(3)
Ingress port = 2 ; IP Dst = 10.2.0.4	Forward(4)

load balancing

Generalized Forwarding and Software Defines Network (SDN)

Destination-based forwarding:

Pattern →

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	51.6.0.8	*	*	*	port6

← Action

IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6

Firewall:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Forward
*	*	*	*	*	*	*	*	*	22	drop

do not forward (block) all datagrams destined to TCP port 22

Destination-based layer 2 (switch) forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	22:A7:23:11:E1:02	*	*	*	*	*	*	*	*	port3

layer 2 frames from MAC address 22:A7:23:11:E1:02 should be forwarded to output port 6