

# CSCI 466: Networks

Peer 2 Peer Networks (P2P), Content Distribution Networks (CDN)

Reese Pearsall  
Fall 2023

# Announcements

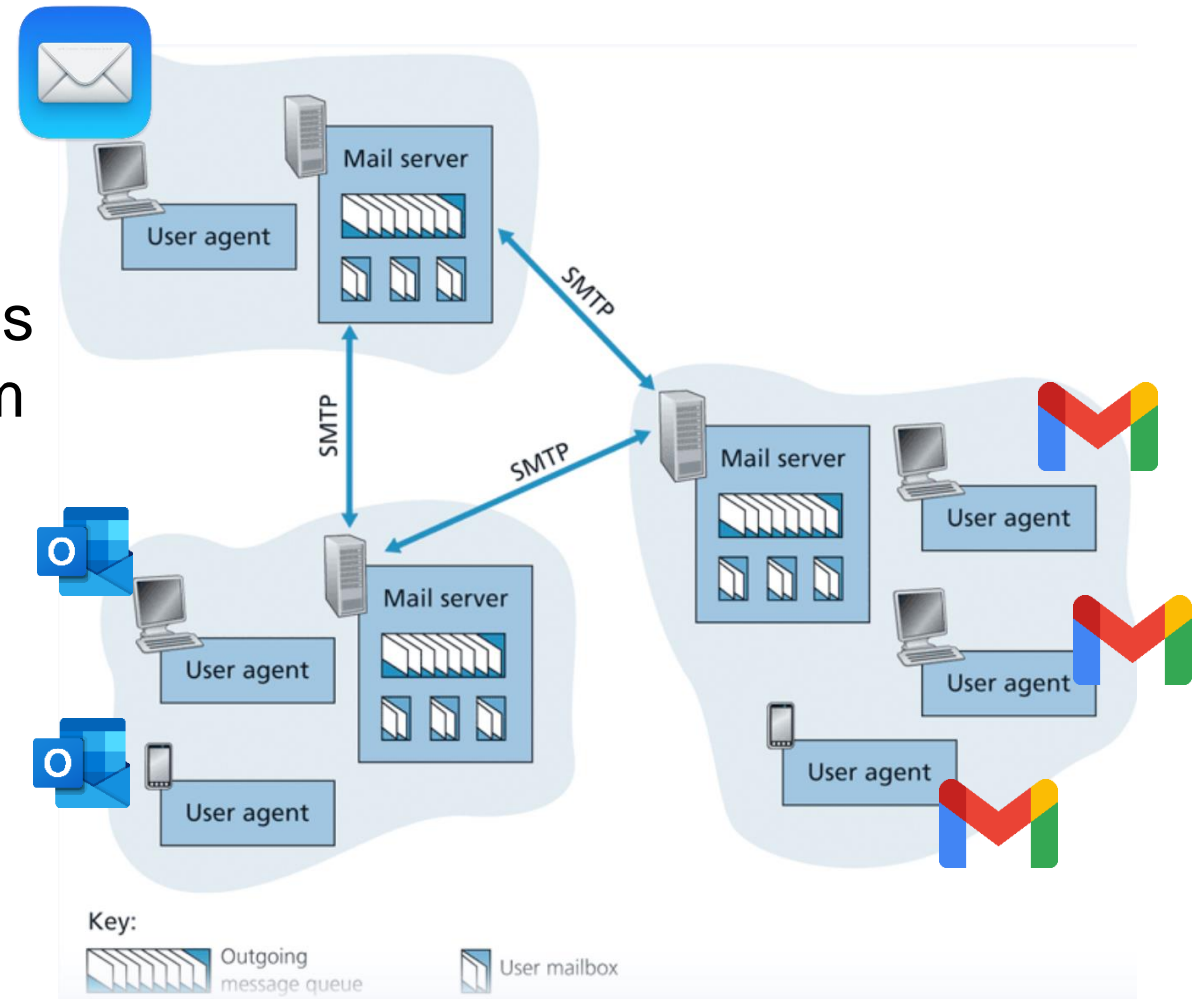
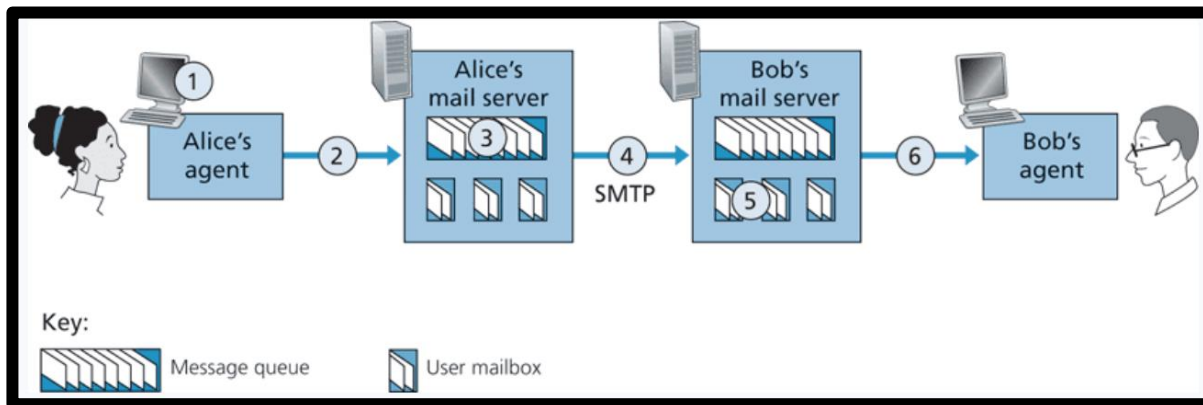
Wireshark Lab 1 due **tonight** at 11:59 PM

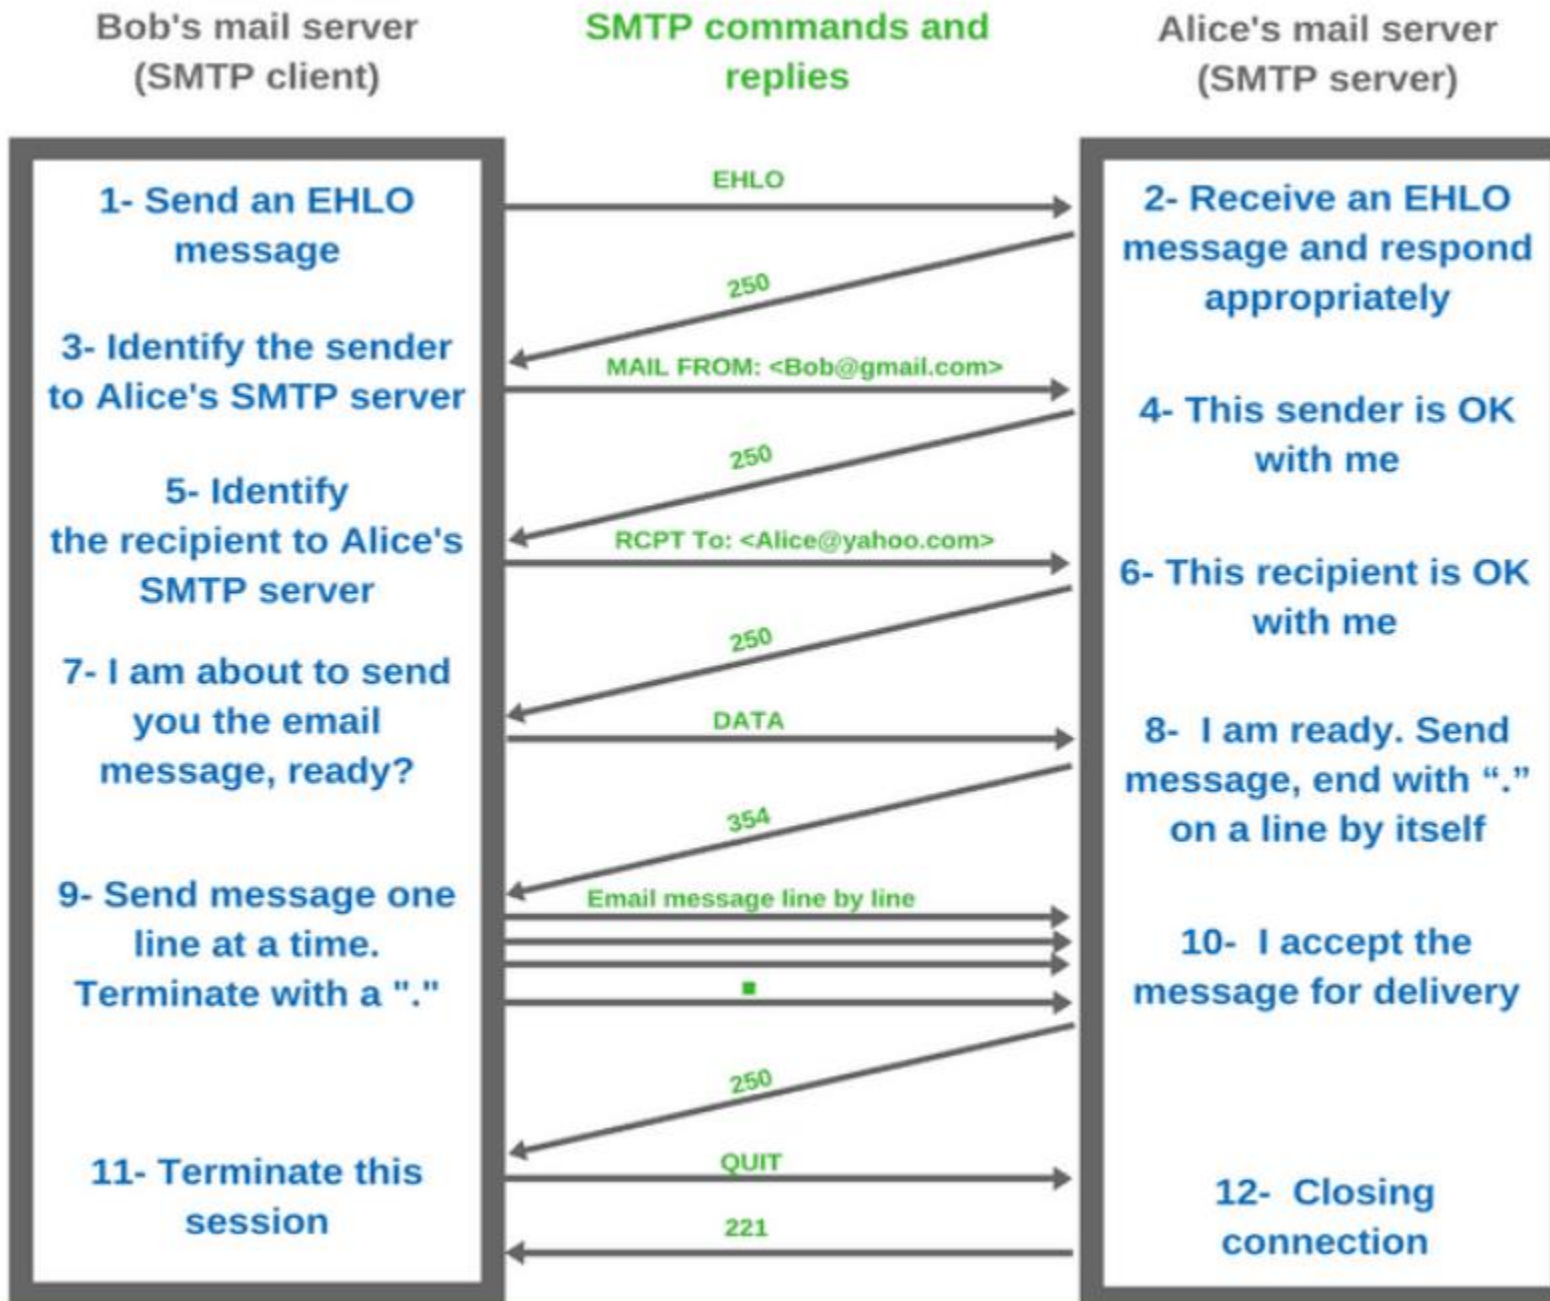
Friday will be a program 1 work day (**no lecture**)

Office hours are cancelled tomorrow

**Simple Mail Transfer Protocol (SMTP)** is the protocol used for sending e-mails from one server to another

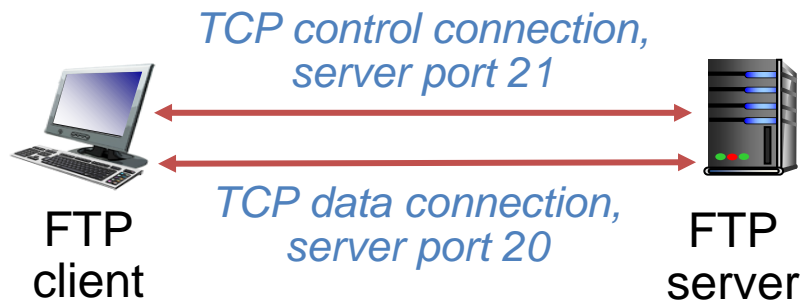
This is not a protocol for *retrieving* emails





# FTP

**File Transfer Protocol (FTP)**- protocol used for transferring files from server to client



- FTP communicates over two connections
  - Port 21 for control information
  - Port 20 for data
- Differences from HTTP
  - Control communication “out-of-band”
  - Server maintains per client state: authentication, current directory

- **FTP procedure:**


1. FTP client contacts FTP server at port 21, using TCP
2. Client authorized over control connection
3. Client browses remote directory, sends commands over control connection
4. When server receives file transfer command, server opens 2nd TCP data connection (for file) to client
5. After transferring one file, server closes data connection

Why use a separate control connection?

WinSCP

# P2P Networks


- No always-on server
- Arbitrary end systems directly communicate
- Peers are intermittently connected and change IP addresses



Time to distribute a file of size  $F$  to  $N$  clients

$$D_{P2P} = \max \left\{ \frac{F}{u_s}, \frac{F}{d_{\min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

## Client-Server Architecture




- Server can upload data at rate  $u_s$
- Clients download data at rates  $d_1, d_2, \dots, d_N$

$$D_{CS} = \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{\min}} \right\}$$


Time to distribute a file of size  $F$  to  $N$  clients

# P2P Networks

## Client-Server Architecture



Existing clients have to share resources with new users

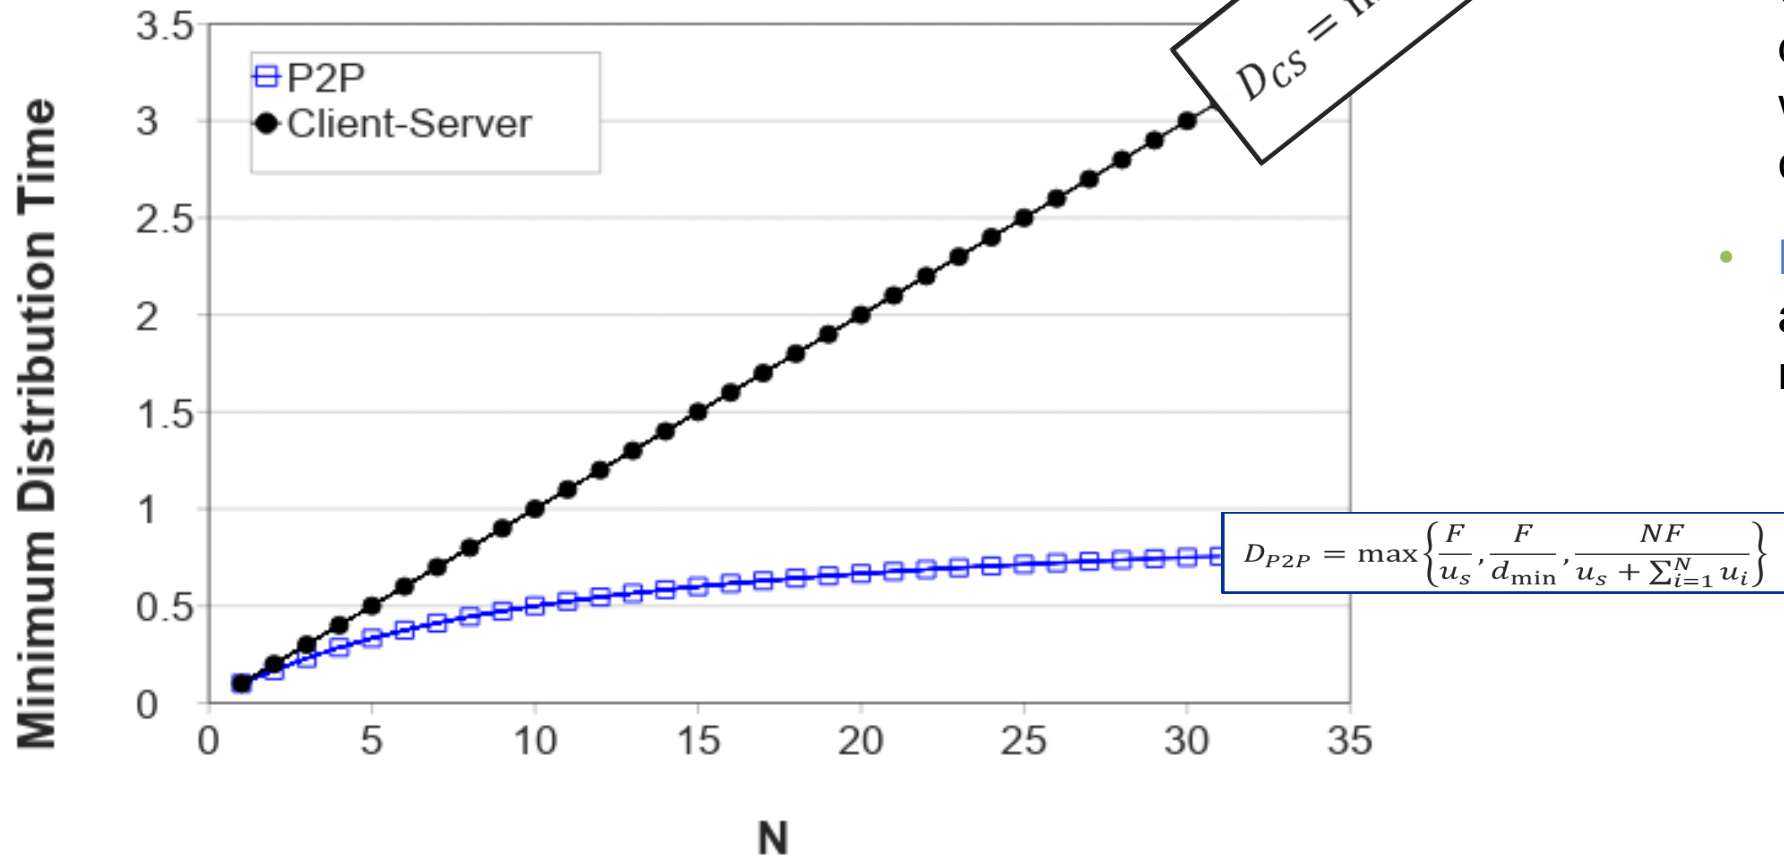


New peers are both a client and a server. There will not be a negative impact on current peers

P2P architectures are *self-scaling*

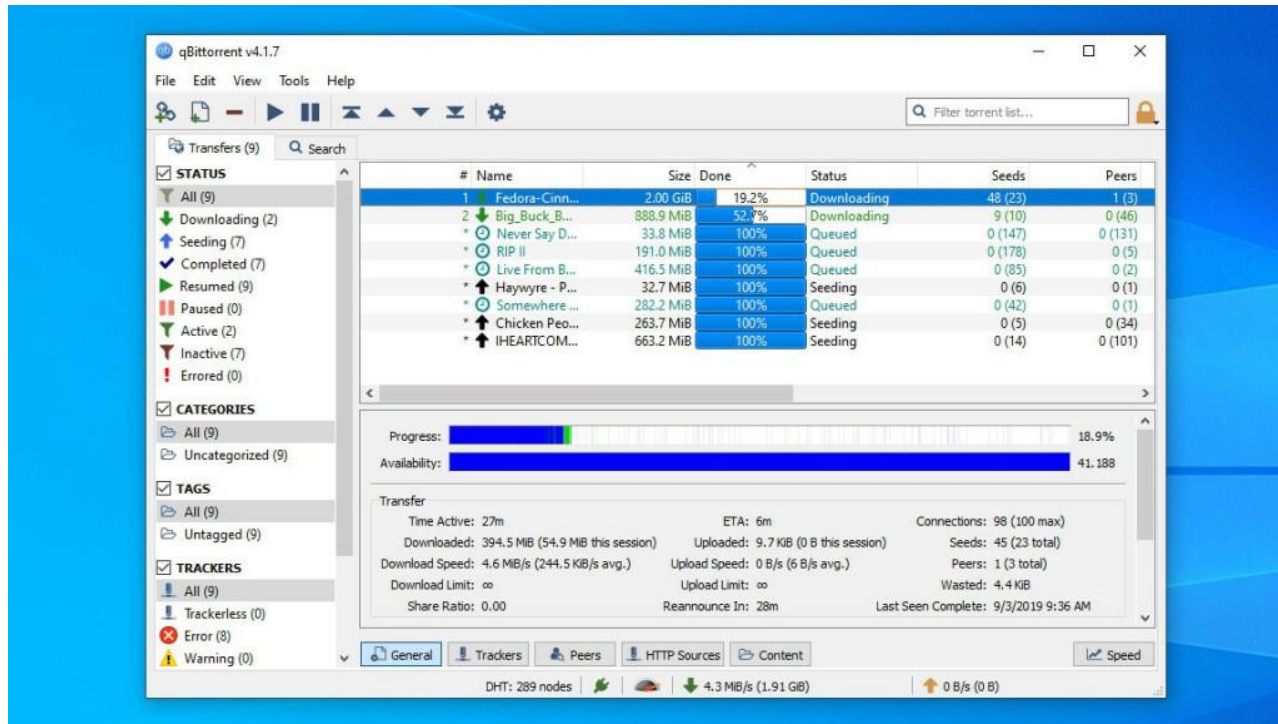


# P2P Networks



- **Client server** distribution time grows with the number of clients
- **P2P** distribution time approaches 1 hour as number of clients grows

# P2P file distribution: BitTorrent




Service for sharing files  
over the internet in a  
decentralized fashion

# P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

*tracker*: tracks peers participating in torrent

*torrent*: group of peers exchanging chunks of a file




# P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

**tracker:** tracks peers participating in torrent

**torrent:** group of peers exchanging chunks of a file




## peer joining torrent:

- has no chunks, but will accumulate them over time from other peers
- registers with tracker to get list of peers, connects to subset of peers (“neighbors”)

# P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

**tracker:** tracks peers participating in torrent


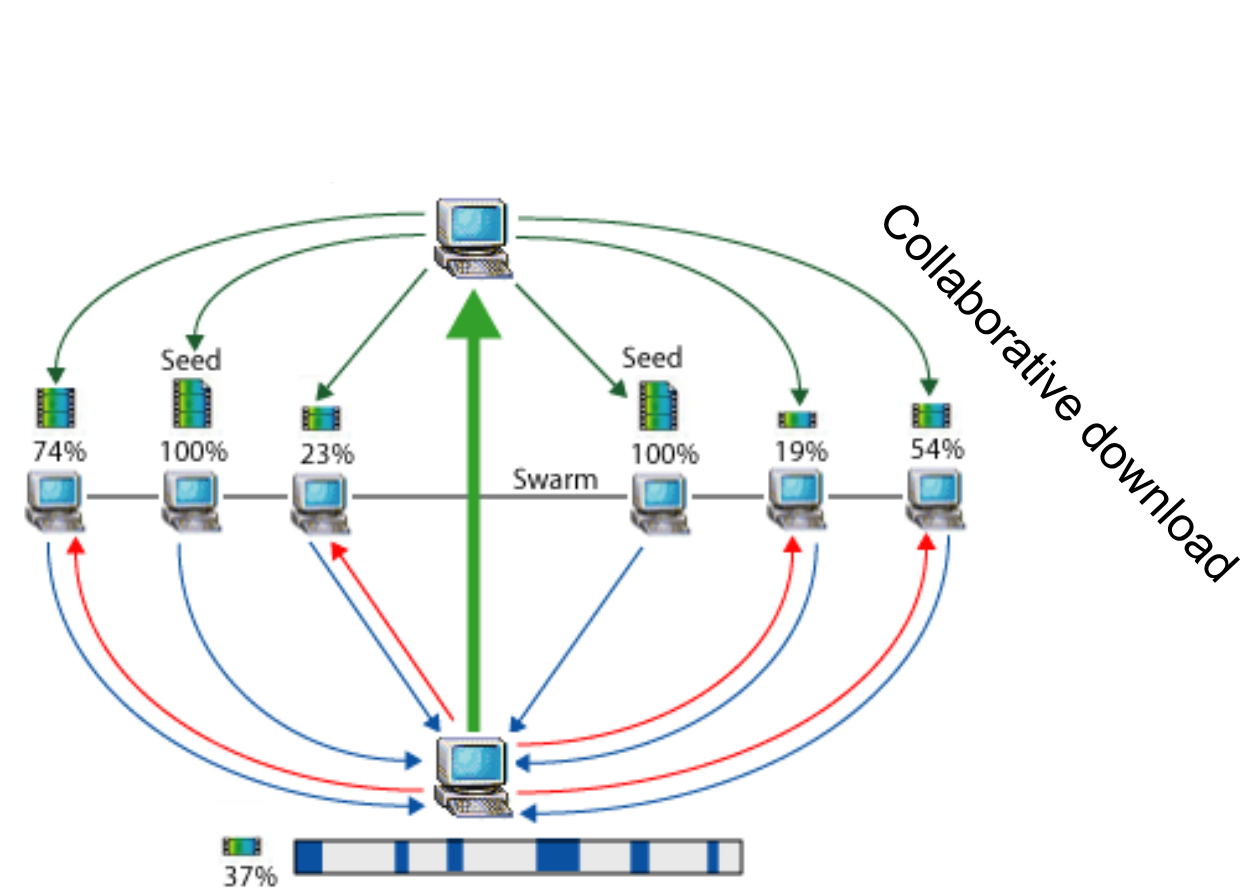


Alice arrives ...  
... obtains list  
of peers from tracker  
... and begins exchanging  
file chunks with peers in torrent

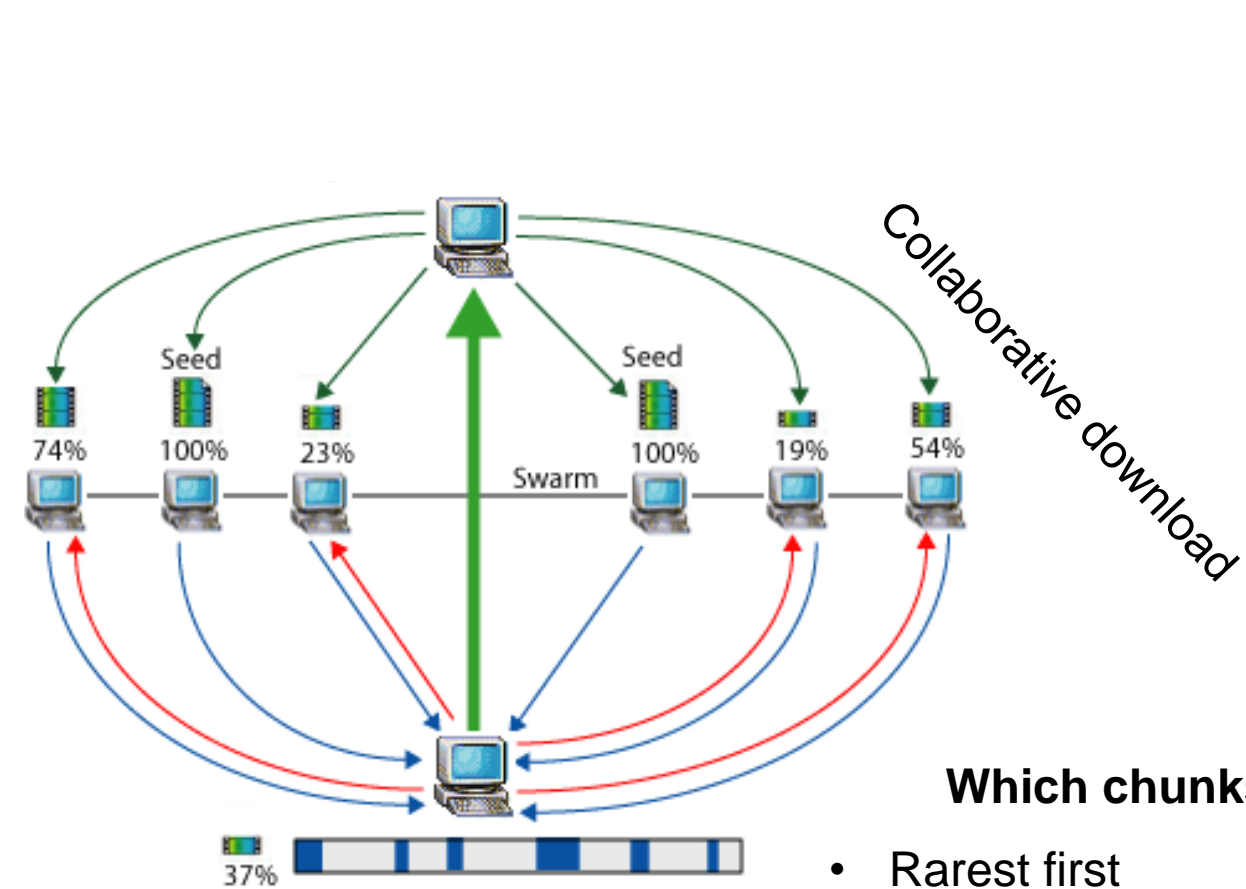
**torrent:** group of peers  
exchanging chunks of a file

- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn:** peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

# P2P file distribution: BitTorrent




# P2P file distribution: BitTorrent

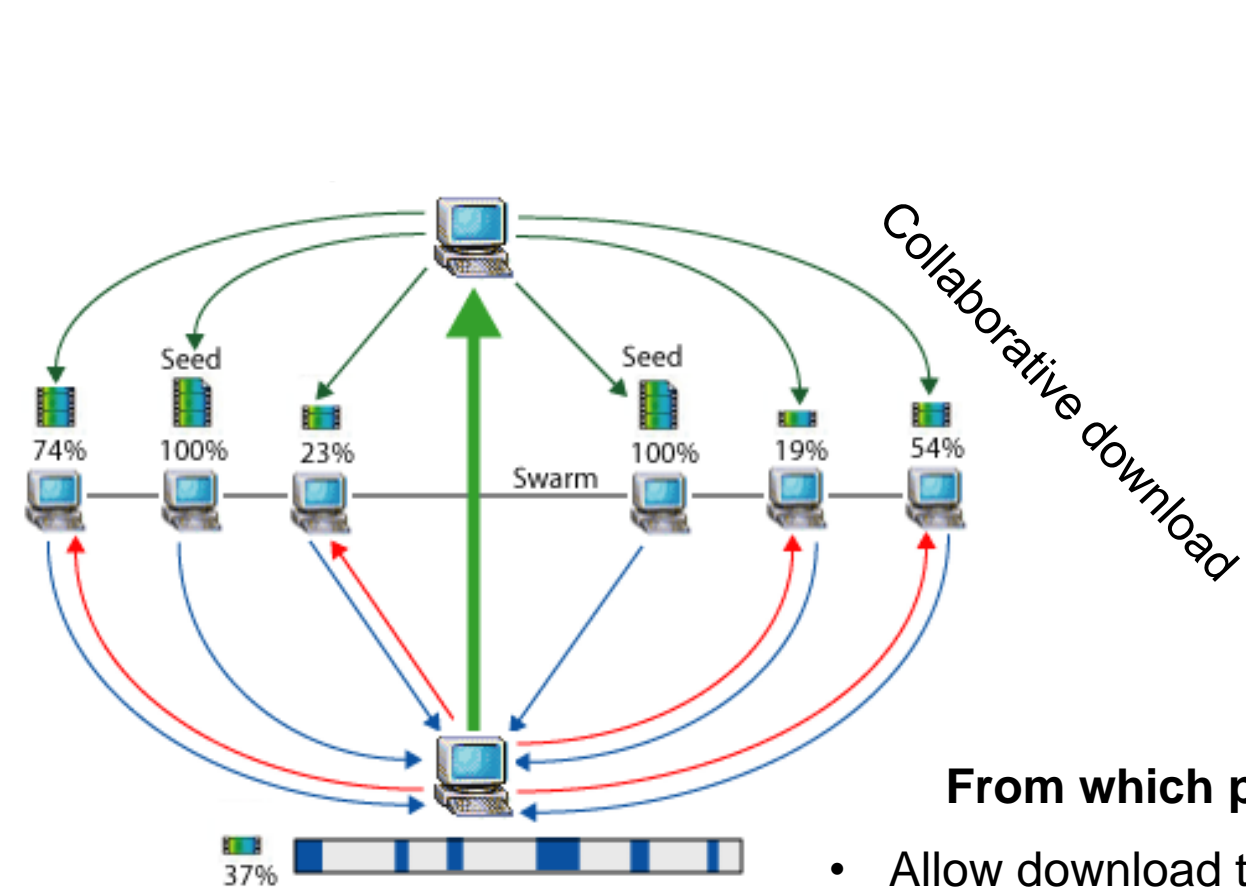


## Which chunks to download?

- Rarest first
- Try to equalize the number of copies of each chunk in the system




# P2P file distribution: BitTorrent



## From which peers to download?

- Allow download to N peers with highest upload rates
- Allows peers with similar upload rates to find each other







# P2P file distribution: BitTorrent

## *requesting chunks:*

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first



# P2P file distribution: BitTorrent



*higher upload rate: find better trading partners, get file faster !*

## *sending chunks: tit-for-tat*


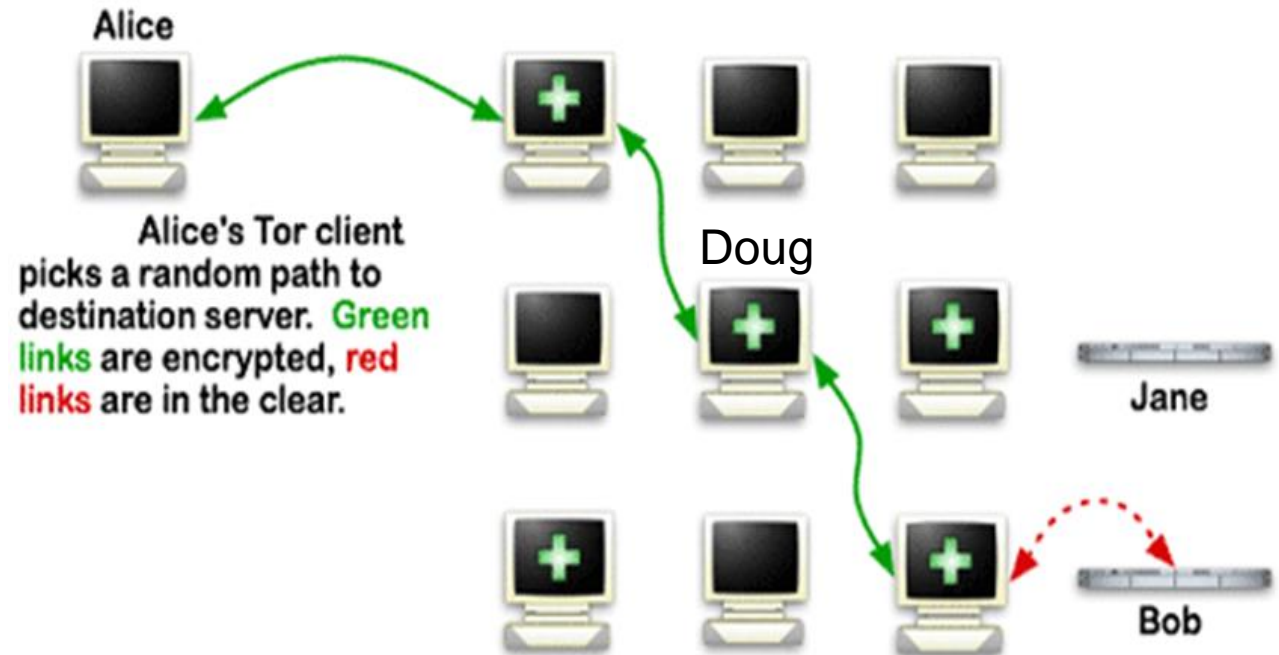
- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
  - other peers are *choked* by Alice (do not receive chunks from her)
  - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
  - “optimistically unchoke” this peer
  - newly chosen peer may join top 4

BitTorrent is referred to as a **unstructured P2P**

# P2P file distribution mechanisms

To help with anonymity and privacy, some P2P may implement an onion router/TOR router

- Functionality
  - Sender obtains a set of router keys
  - Each router only knows next hop
  - Intermediate routers cannot read message




# CDNs

- video traffic: major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
  - single mega-video server won't work (why?)
- challenge: heterogeneity
  - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution*: distributed, application-level infrastructure



# CDNs


- Caching
  - Save previously delivered data
  - Subsequent requests served from cache on the browser, or in the access network
- Applications
  - Reduce response time for client request
  - Reduce ISP traffic costs
- Content distribution networks
  - Distributed caches
  - Web objects addressed to CDN server
  - CDN server fetches from content provider on first access



- Conditional GET
  - Cache: specify date of cached copy in HTTP request  
**If-modified-since: <date>**
  - Server: response contains no object if cached copy is up-to-date:  
**HTTP/1.0 304 Not Modified**

# CDNs

- Challenge:
  - How to stream content to millions of users?
- Option 1:
  - Single mega-datacenter
  - Pros: Simple
  - Cons:
    - Single point of failure
    - Point of network congestion
    - Long path to distant clients
    - Multiple copies of video sent over outgoing link



# CDNs

- Option 2:
  - Store/serve multiple copies of videos at multiple geographically distributed sites


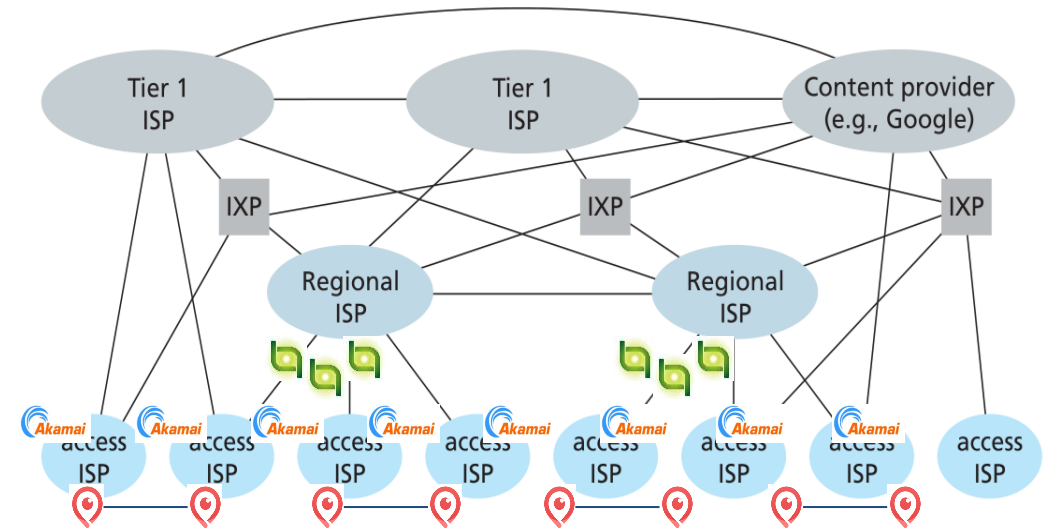
A **content delivery network (CDN)** refers to a geographically distributed group of cache servers which work together to provide fast delivery of Internet content. **(not a web host)**



34 DNS lookups

204 HTTP requests

520 KB of data downloaded



56% of domains resolve to a CDN

# CDNs

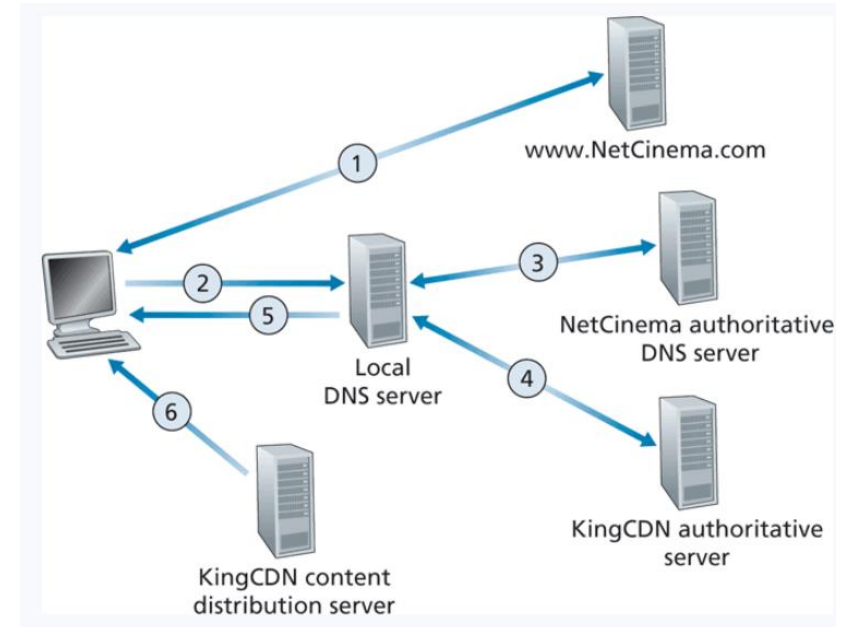


[https://\*\*cdn\*\*.discordapp.com/attachments/1017630358621143110/1017849417216294972/unknown.png](https://cdn.discordapp.com/attachments/1017630358621143110/1017849417216294972/unknown.png)



# CDNs

- Challenge: how does CDN DNS select “good” CDN node to stream to client
  - Pick CDN node **geographically** closest to client’s local DNS
  - Pick CDN node with **shortest delay** (or min # hops) to client (CDN nodes periodically ping access ISPs, reporting results to CDN DNS)



# CDNs

In HTTP streaming, video is stored at an HTTP server and retrieved with a GET request

Frames are sent to a client buffered and played back after a certain threshold

All clients receiving the same encoding of the video despite widely different bandwidth

# CDNs

In HTTP steaming, video is stored at an HTTP server and retrieved with a GET request

Frames are sent to a client buffered and played back after a certain threshold

All clients receiving the same encoding of the video despite widely different bandwidth

**Dynamic Adaptive Streaming over HTTP (DASH)**- video is encoded into several different versions

- Different encoding rate, different quality, etc

Clients dynamically request chunks of video segments every few seconds via GET requests

- Is the current bandwidth good? → Retrieve good quality
- Is the current bandwidth bad? → Retrieve ok quality

## CDNs

When streaming, you are consistently issuing GET requests

[illegible]

Chromium DevTools Network tab showing a waterfall chart and a table of network requests.


**Waterfall Chart:** A horizontal bar chart at the top showing the timing of 18 requests. The x-axis represents time in milliseconds, ranging from 0 to 28,000 ms. The bars are color-coded by type: blue for fetch, green for xhr, and grey for other. The chart shows that most requests are completed within the first 10,000 ms, with a few requests taking up to 28,000 ms.

**Network Requests Table:**

Name	Status	Type	Initiator	Size	Time	Waterfall
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.5 MB	48 ms	
watchtime?ns=yt&el=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	30 ms	
qoe?fmt=399&afmt=251&cpn=D1QCo4_7KrXNdhs4...	204	xhr	VM685:201	39 B	28 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	256 kB	20 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	2.1 MB	50 ms	
log_event?alt=json&key=AlzaSyAO_FJ2SlqU8Q4STEH...	200	xhr	VM685:201	77 B	46 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.4 MB	56 ms	
watchtime?ns=yt&el=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	31 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	2.1 MB	55 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	391 kB	23 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.6 MB	32 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.3 MB	53 ms	
watchtime?ns=yt&el=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	32 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	2.1 MB	60 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.6 MB	65 ms	
data:image/png;base...	200	png	Other	(memory cache)	0 ms	


**Summary:** 18 requests | 14.8 MB transferred | 14.8 MB resources

# CDNs



Studio master versions are uploaded to a private Amazon cloud


# CDNs



Studio master versions are uploaded to a private Amazon cloud

Videos are processed into many different formats, allowing for DASH

# CDNs



Studio master versions are uploaded to a private Amazon cloud

Videos are processed into many different formats, allowing for DASH

Versions are uploaded to Netflix's CDNs