

# CSCI 466: Networks

UDP and TCP

Reese Pearsall  
Fall 2023

## Announcements

PA 2 Posted. Due Wednesday October 18th

# OSI Model

Application Layer

Presentation Layer \*

Session Layer \*

Transport Layer

Network Layer

Data Link Layer

Physical Layer

Application Layer

Messages from Network Applications

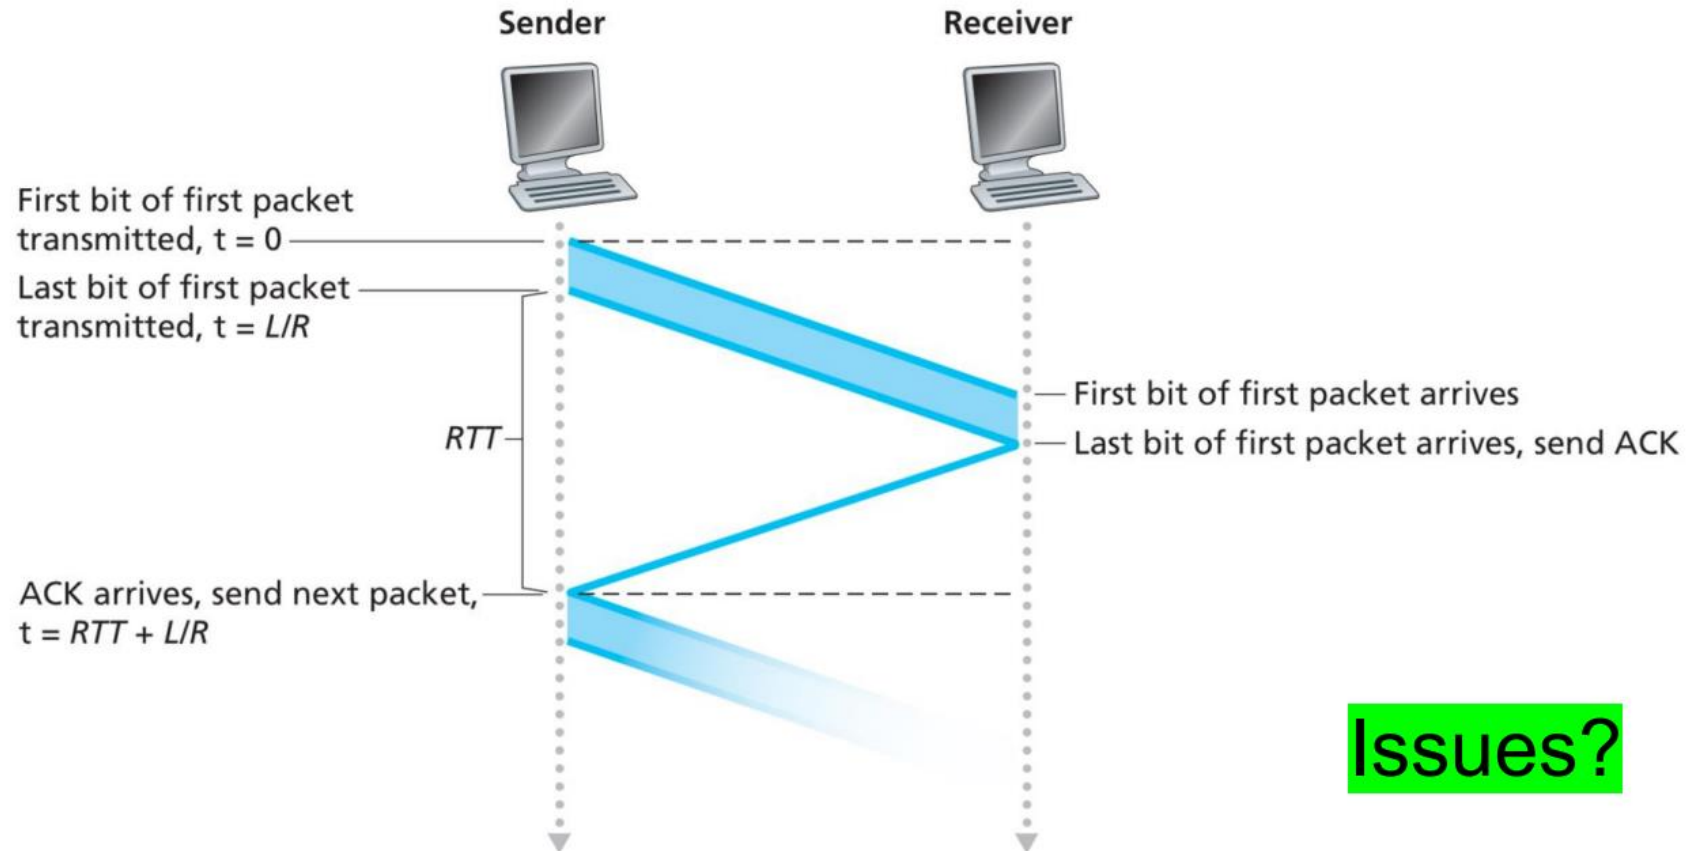


Physical Layer

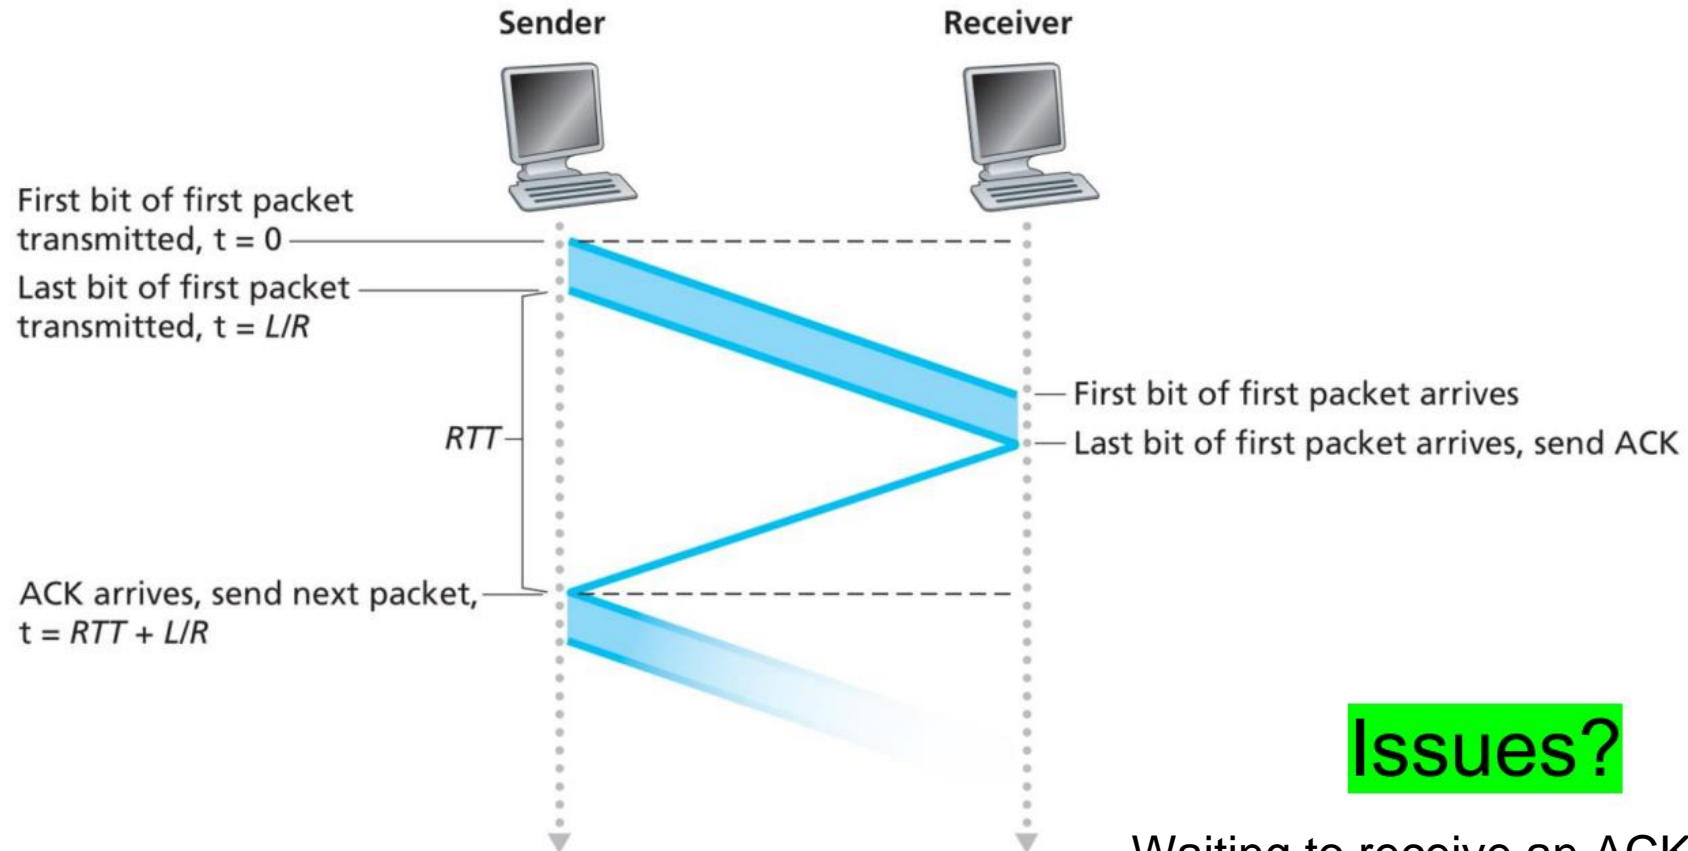
Bits being transmitted over some medium

*\*In the textbook, they condense it to a 5-layer model, but 7 layers is what is most used*

So far, our reliable data transfer protocols have been **stop-and-wait**



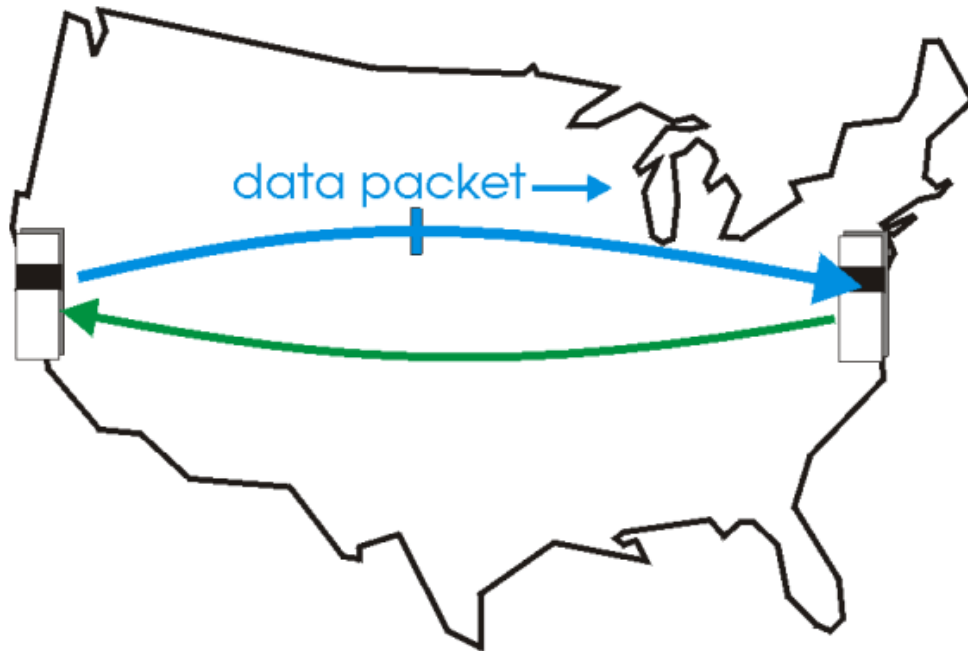
So far, our reliable data transfer protocols have been **stop-and-wait**



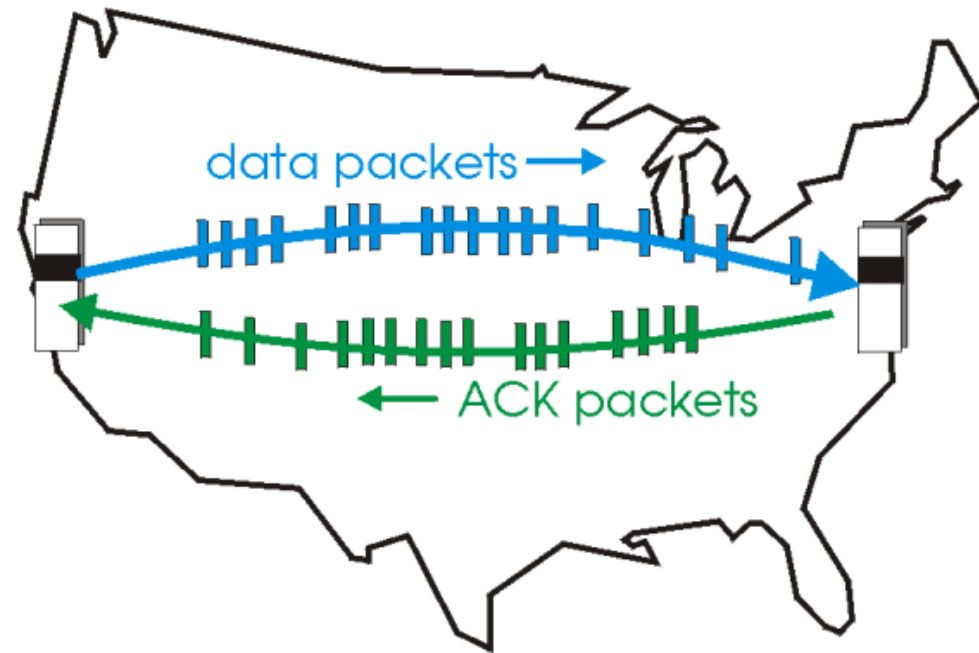
Issues?

Waiting to receive an ACK is inefficient  
We spend a *lot* of time waiting

**Pipelining:** sender allows multiple, “in-flight”, yet-to-be acknowledged packets




(a) a stop-and-wait protocol in operation



(b) a pipelined protocol in operation


**Pipelining:** sender allows multiple, “in-flight”, yet-to-be acknowledged packets



Consequences:

- Need to use a wider range of sequence numbers

**Pipelining:** sender allows multiple, “in-flight”, yet-to-be acknowledged packets



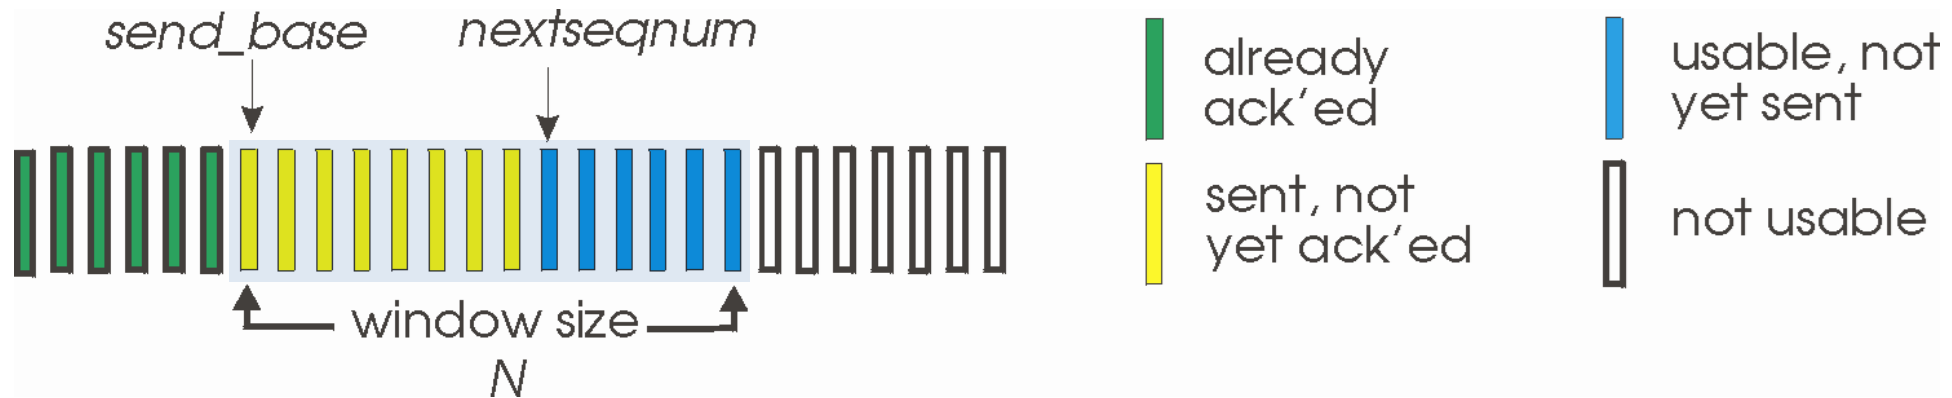
Consequences:

- Need to use a wider range of sequence numbers
- Sender and receiver may need to **buffer** more than one packet



# Go-Back-N: sender

- sender: “window” of up to  $N$ , consecutive transmitted but unACKed pkts
  - $k$ -bit seq # in pkt header

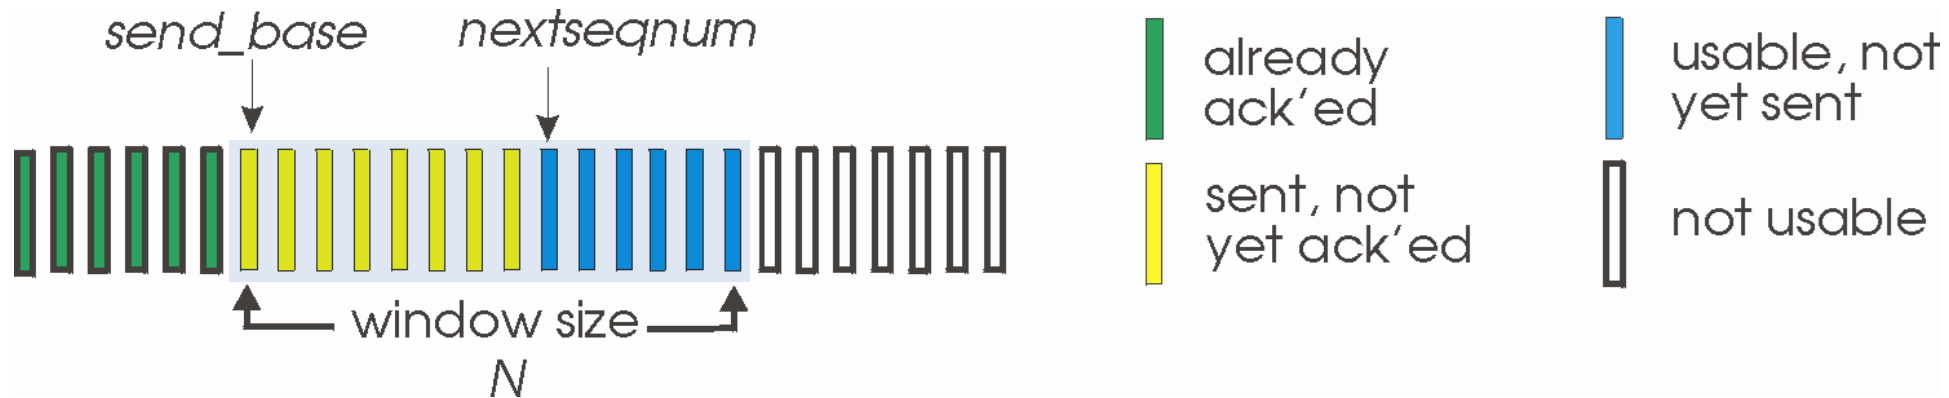


- ***cumulative ACK***:  $ACK(n)$ : ACKs all packets up to, including seq #  $n$ 
  - on receiving  $ACK(n)$ : move window forward to begin at  $n+1$
- timer for oldest in-flight packet
- ***timeout(n)***: retransmit packet  $n$  and all higher seq # packets in window

# Go-Back-N: sender

(Why not have an infinite window size?)

- sender: “window” of up to  $N$ , consecutive transmitted but unACKed pkts
  - $k$ -bit seq # in pkt header

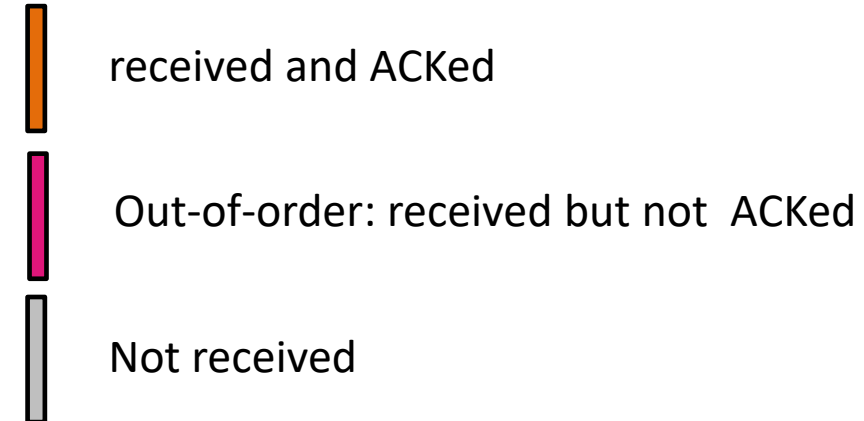
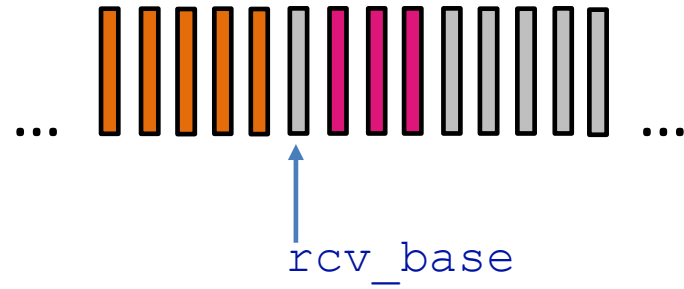


- **cumulative ACK:**  $ACK(n)$ : ACKs all packets up to, including seq #  $n$ 
  - on receiving  $ACK(n)$ : move window forward to begin at  $n+1$
- timer for oldest in-flight packet
- **timeout( $n$ ):** retransmit packet  $n$  and all higher seq # packets in window

# Go-Back-N: receiver

- ACK-only: always send ACK for correctly-received packet so far, with highest *in-order* seq #
  - may generate duplicate ACKs
  - need only remember `rcv_base`
- on receipt of out-of-order packet:
  - can discard (don't buffer) or buffer: an implementation decision
  - re-ACK pkt with highest in-order seq #

Receiver view of sequence number space:



sender window (N=4)

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8  
0 1 2 3 4 5 6 7 8

sender

send pkt0  
send pkt1  
send pkt2  
send pkt3  
(wait)

rcv ack0, send pkt4  
rcv ack1, send pkt5

ignore duplicate ACK



*pkt 2 timeout*

send pkt2  
send pkt3  
send pkt4  
send pkt5

receiver

receive pkt0, send ack0  
receive pkt1, send ack1

receive pkt3, discard,  
(re)send ack1

receive pkt4, discard,  
(re)send ack1

receive pkt5, discard,  
(re)send ack1

rcv pkt2, deliver, send ack2  
rcv pkt3, deliver, send ack3  
rcv pkt4, deliver, send ack4  
rcv pkt5, deliver, send ack5

[https://www2.tkn.tu-berlin.de/teaching/rn/animations/gbn\\_sr/](https://www2.tkn.tu-berlin.de/teaching/rn/animations/gbn_sr/)

# Selective repeat: the approach

- *pipelining*: multiple packets in flight
- *receiver individually ACKs* all correctly received packets
  - buffers packets, as needed, for in-order delivery to upper layer
- sender:
  - maintains (conceptually) a timer for each unACKed pkt
    - timeout: retransmits single unACKed packet associated with timeout
  - maintains (conceptually) “window” over *N* consecutive seq #s
    - limits pipelined, “in flight” packets to be within this window

[https://www2.tkn.tu-berlin.de/teaching/rn/animations/gbn\\_sr/](https://www2.tkn.tu-berlin.de/teaching/rn/animations/gbn_sr/)

# Selective repeat: sender and receiver

## sender

### data from above:

- if next available seq # in window, send packet

### timeout( $n$ ):

- resend packet  $n$ , restart timer

### ACK( $n$ ) in [sendbase, sendbase+N-1]:

- mark packet  $n$  as received
- if  $n$  smallest unACKed packet, advance window base to next unACKed seq #

## receiver

### packet $n$ in [rcvbase, rcvbase+N-1]

- send ACK( $n$ )
- out-of-order: buffer
- in-order: deliver (also deliver buffered, in-order packets), advance window to next not-yet-received packet

### packet $n$ in [rcvbase-N, rcvbase-1]

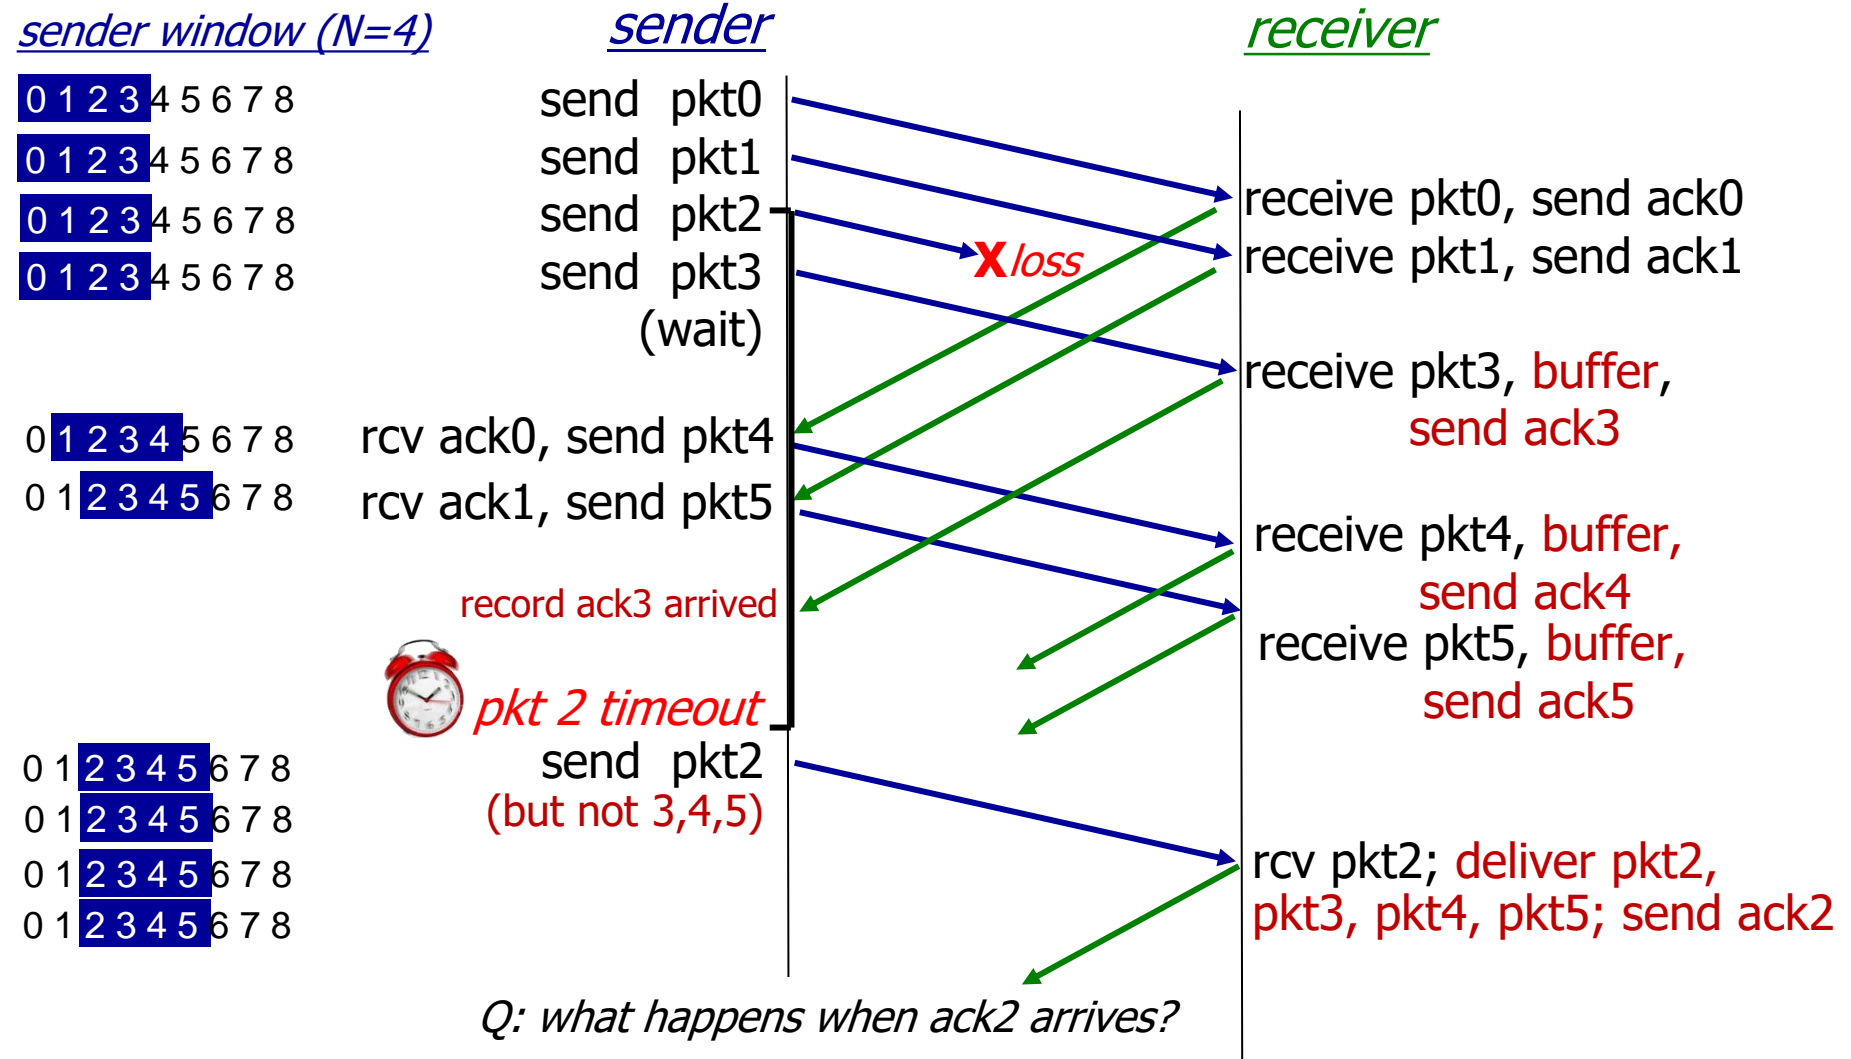
- ACK( $n$ )

### otherwise:

- ignore



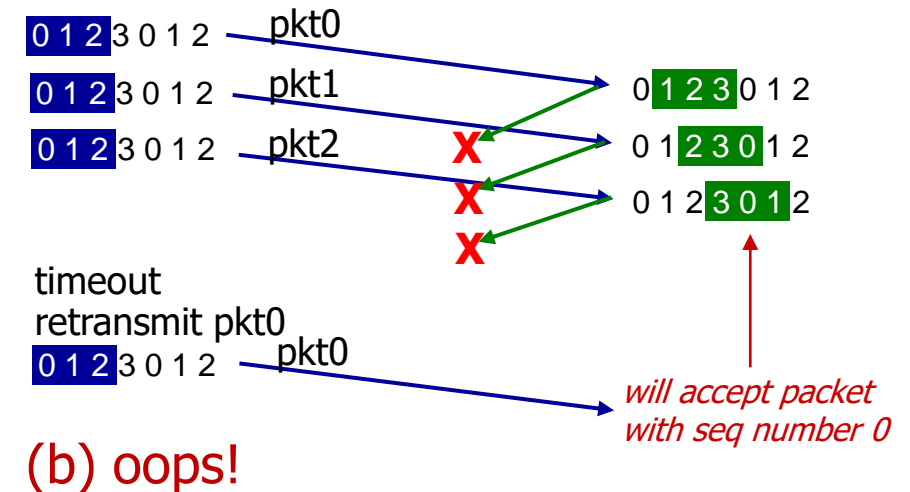
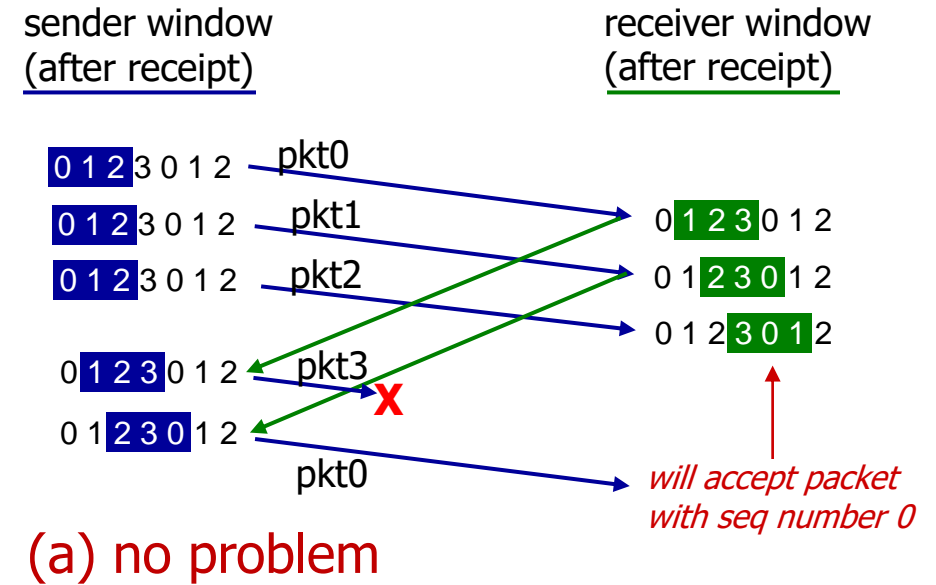
# Selective Repeat in action



# Selective repeat: a dilemma!

example:

- seq #s: 0, 1, 2, 3 (base 4 counting)
- window size=3

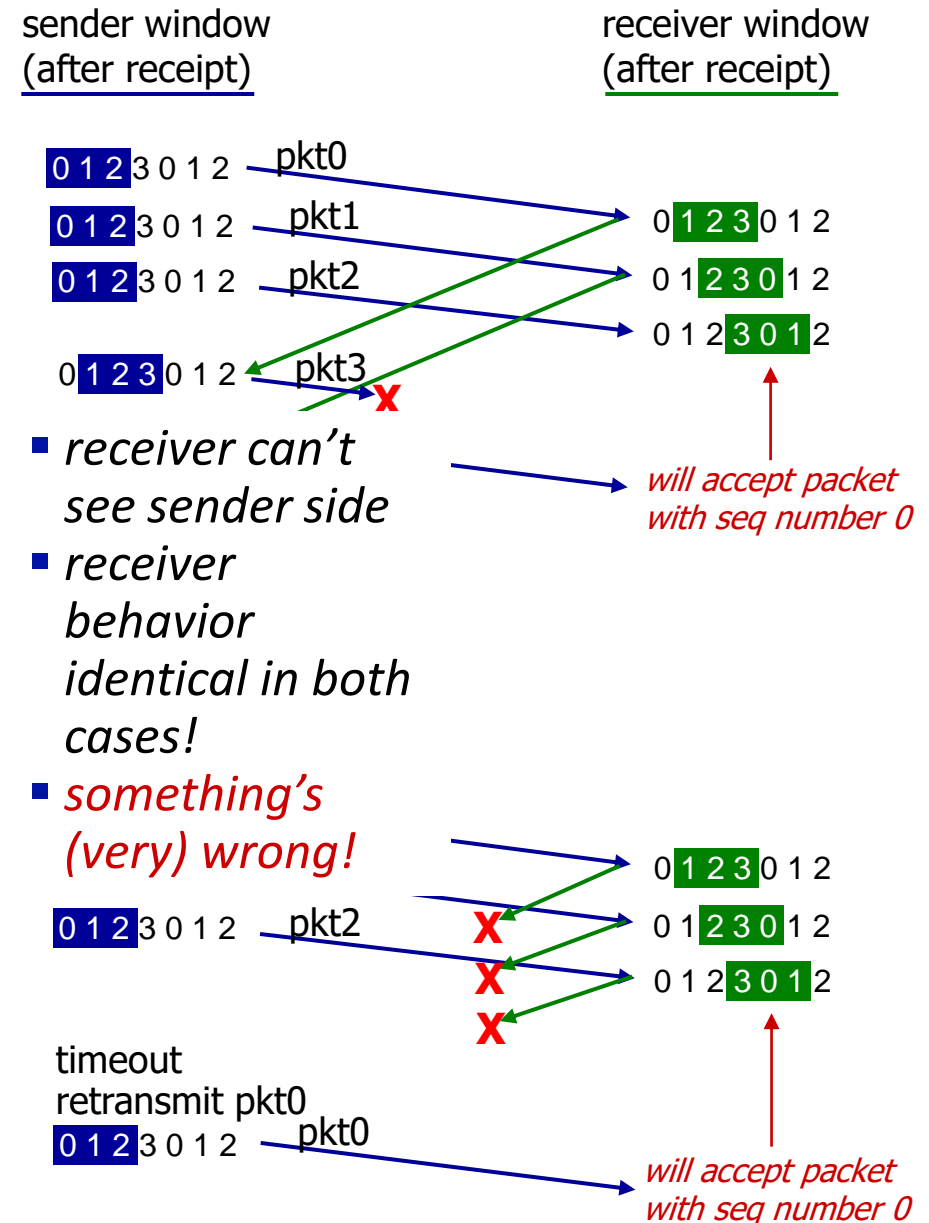


# Selective repeat: a dilemma!

example:

- seq #s: 0, 1, 2, 3 (base 4 counting)
- window size=3

**Q:** what relationship is needed between sequence # size and window size to avoid problem in scenario (b)?



## Transport Layer

# RDT Principles

**Checksum-** Used to detect bit errors in transmitted packets

**Checksum-** Used to detect bit errors in transmitted packets

**Timer-** Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

**Checksum-** Used to detect bit errors in transmitted packets

**Timer-** Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

**Sequence Number-** Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

**Checksum-** Used to detect bit errors in transmitted packets

**Timer-** Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

**Sequence Number-** Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

**Acknowledgement-** Used by the receiver to tell the sender that a packet or set of packets has been received correctly. ACKs will typically carry the sequence # of the packet being acknowledged

**Checksum-** Used to detect bit errors in transmitted packets

**Timer-** Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

**Sequence Number-** Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

**Acknowledgement-** Used by the receiver to tell the sender that a packet or set of packets has been received correctly. ACKs will typically carry the sequence # of the packet being acknowledged

**Negative Acknowledgement-** Used by the receiver to tell the sender that a packet has not been received correctly. Negative acknowledgements will typically carry the sequence number of the packet that was not received correctly



**Checksum-** Used to detect bit errors in transmitted packets

**Timer-** Used to timeout/retransmit a packet, possibly because the packet (or its ACK) was lost within a channel

**Sequence Number-** Used for sequential numbering of packets of data flowing from sender to receiver. Gaps in sequence number of packets allow the receiver to detect a lost or duplicate packet

**Acknowledgement-** Used by the receiver to tell the sender that a packet or set of packets has been received correctly. ACKs will typically carry the sequence # of the packet being acknowledged

**Negative Acknowledgement-** Used by the receiver to tell the sender that a packet has not been received correctly. Negative acknowledgements will typically carry the sequence number of the packet that was not received correctly

**Window, pipelining-** The sender may be restricted to sending only packets with sequence numbers that fall within a given range. By allowing multiple packets to be transmitted but not yet acknowledged, sender utilization can be increased over a stop-and-wait mode of operation

## Transport Layer Protocols:

1. **Transmission Control Protocol (TCP)**
2. **User Datagram Protocol (UDP)**

- “no frills,” “bare bones” Internet transport protocol
- “best effort” service, UDP segments may be:
  - lost
  - delivered out-of-order to app
- *connectionless*:
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others


### Why is there a UDP?

- no connection establishment (which can add RTT delay)
- simple: no connection state at sender, receiver
- small header size
- no congestion control
  - UDP can blast away as fast as desired!
  - can function in the face of congestion

# Transport Layer


## UDP

The UDP header is very small!!  
(4 bytes, 32 bits)




- Connection oriented, point-to-point (1 to 1)  
→ **TCP Handshake** must occur before data is being transmitted

*A logical connection*



- Reliable, in order, data transfer

- Cumulative ACKs
- Pipelining  
→ TCP Congestion and flow control set window size
- Flow controlled  
→ Sender will not overwhelm receiver
- Full-duplex service



The size of the segment is determined by the maximum segment size (MSS)

*(Roughly 1500 bytes– size of a link layer frame)*


A TCP connection is transmitting a **byte stream**



Sequence numbers are based on *how much data has been sent*  
Acknowledgement numbers are based on *how much data has been successfully received*

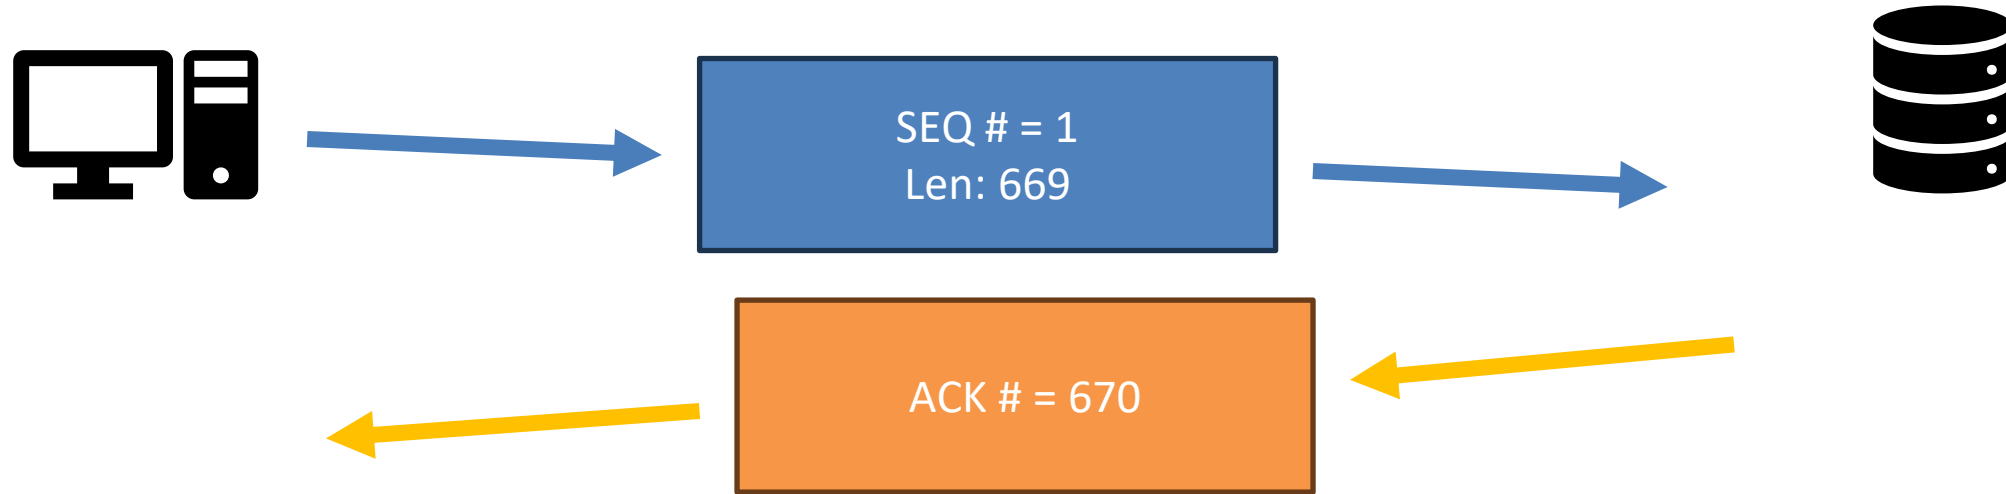


A TCP connection is transmitting a **byte stream**

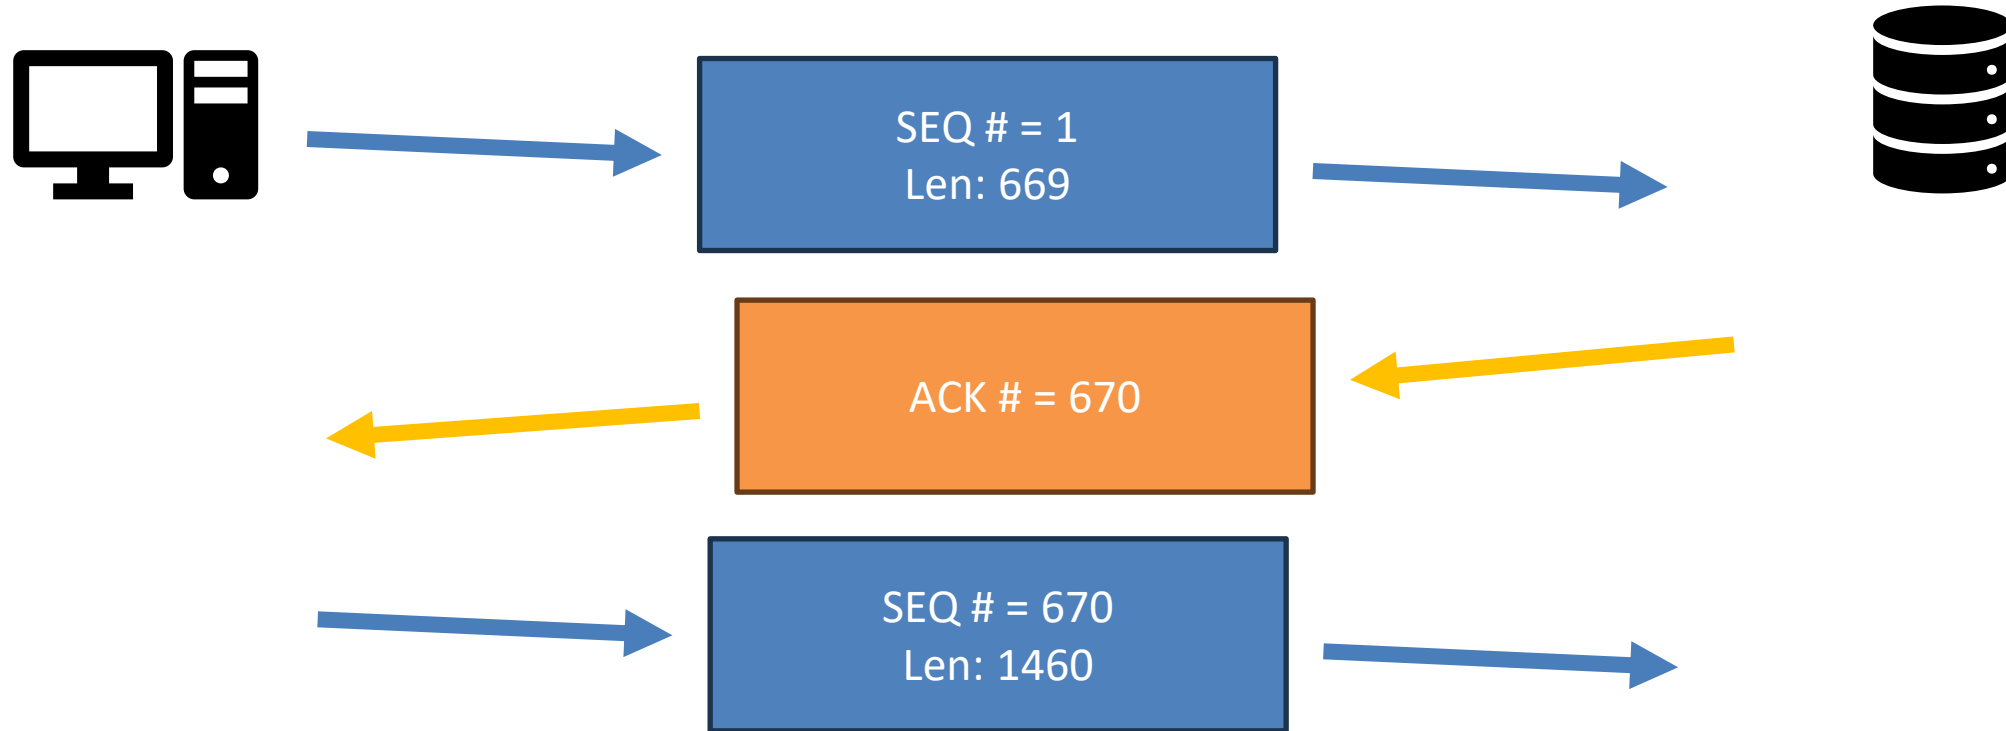




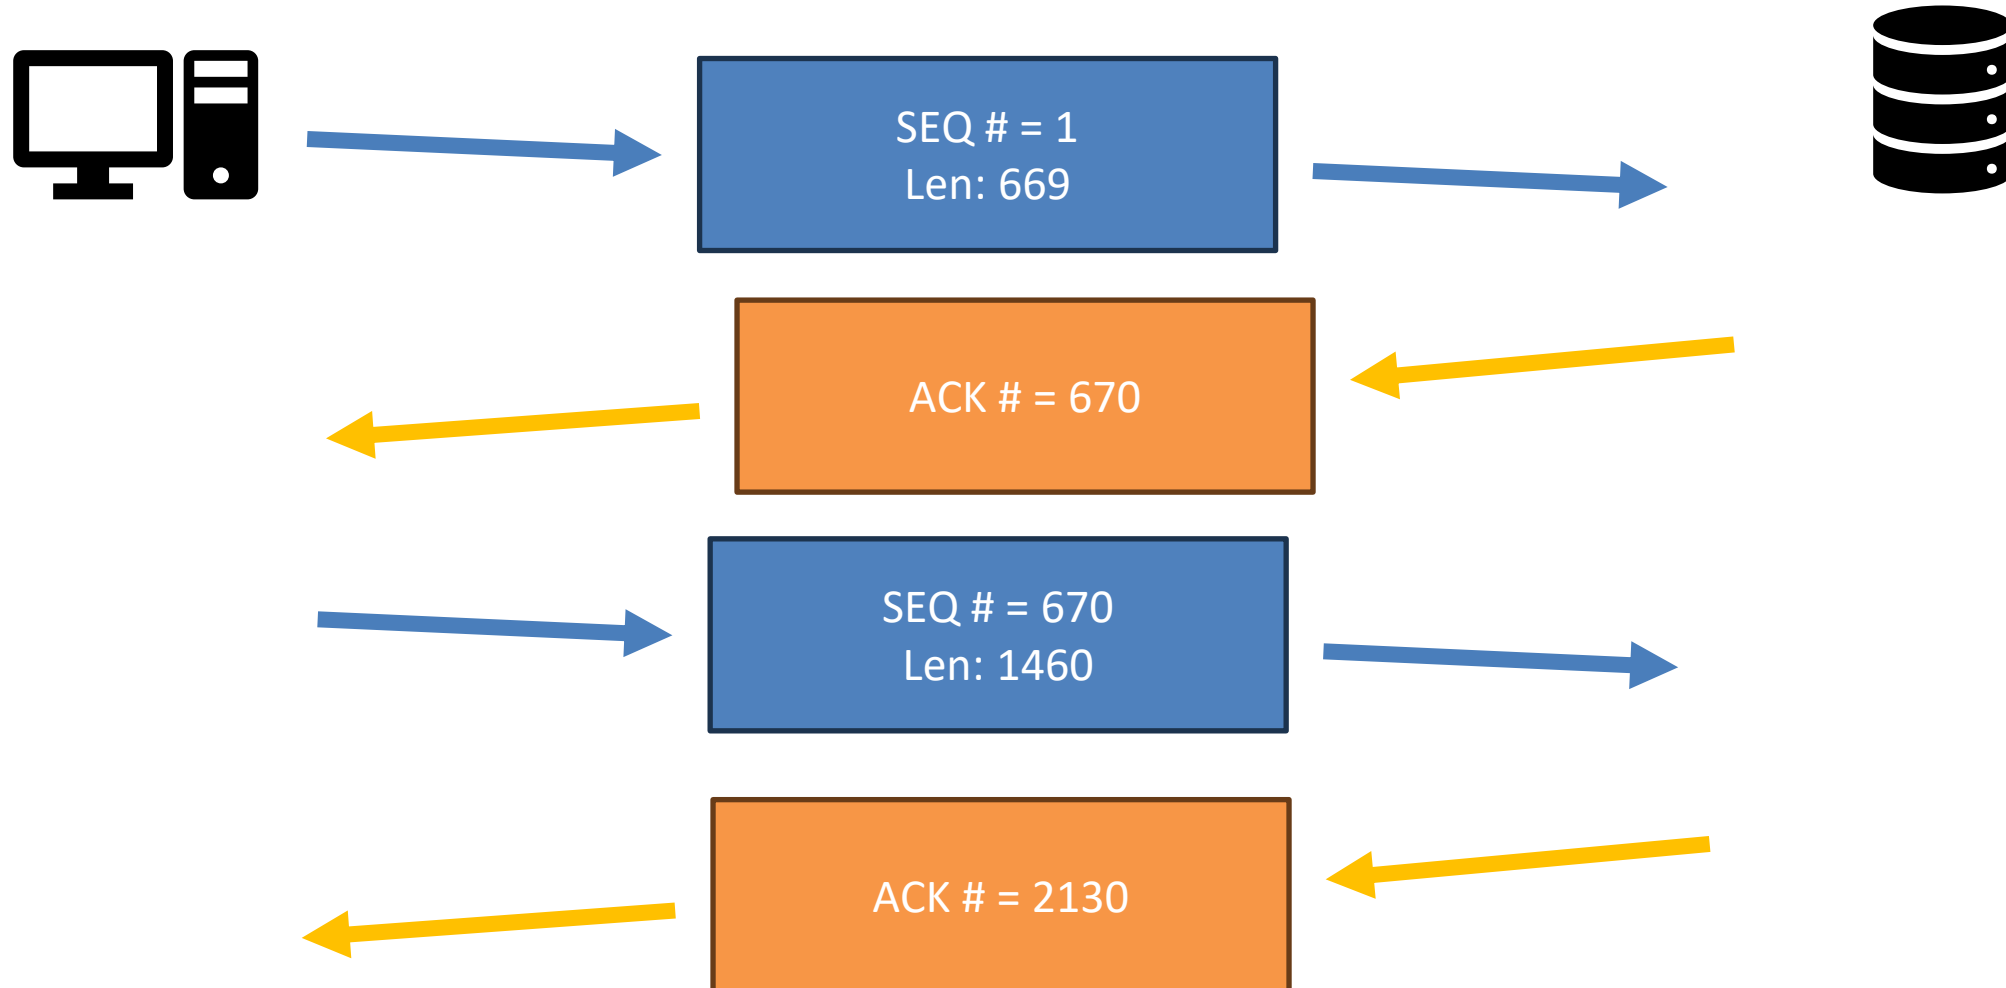
A TCP connection is transmitting a **byte stream**




A TCP connection is transmitting a **byte stream**



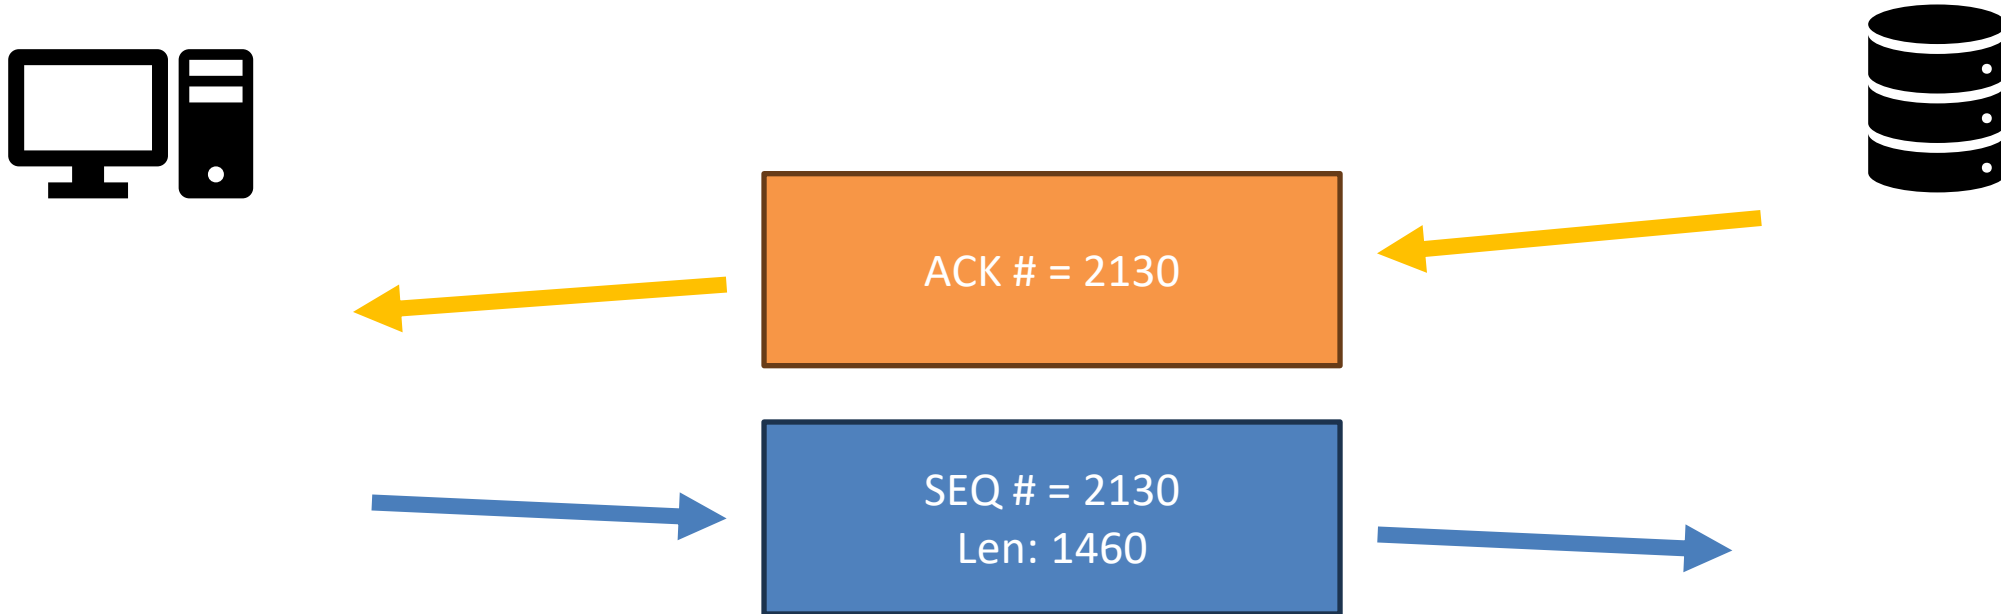
A TCP connection is transmitting a **byte stream**



A TCP connection is transmitting a **byte stream**



A TCP connection is transmitting a **byte stream**



A TCP connection is transmitting a **byte stream**

