

CSCI 476: Computer Security

Lecture 3: Operating Systems, Processes, and `forking()`

Reese Pearsall
Fall 2023

Announcements

If you have an M1/M2 chip, or if you are still struggling with your VM,
check in with me this week

Lab 0 due on Sunday 8/10 @ 11:59 PM
(All assignments will be due on Sundays)

We will start using our VM on Thursday

Might have to cancel a class next week



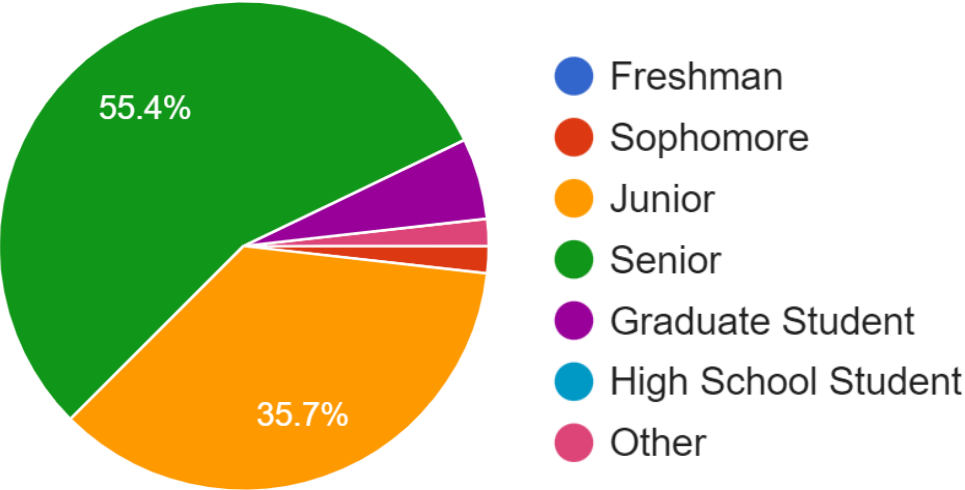
WELCOME HACKERCATS

Meeting Information
Wednesday 09/06 at 6 pm in NAH 137

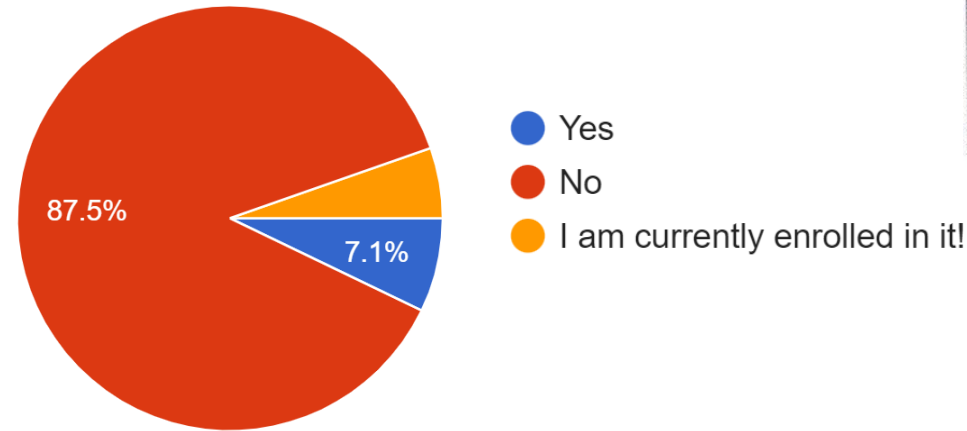
Topics

- Officer introductions
- Cybersecurity in a nutshell
- OSI interactive challenge

Course Questionnaire Results



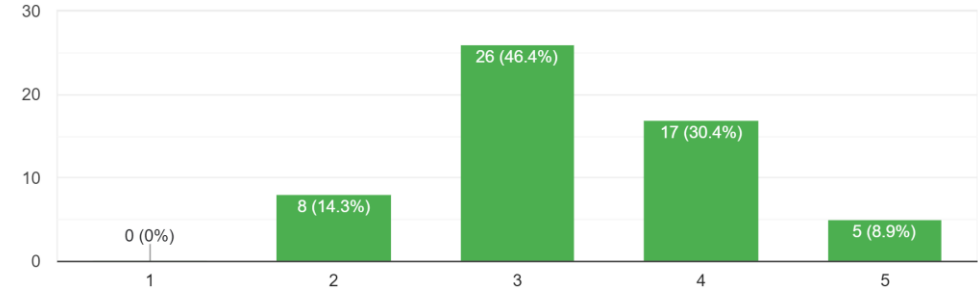
Have you taken Operating Systems (CSCI 460)



How comfortable are you C?

56 responses

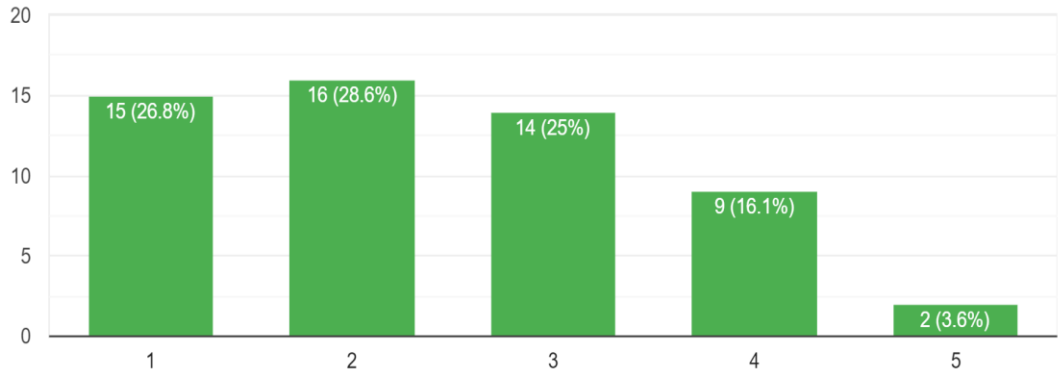
[Copy](#)



How comfortable are you with reading assembly code?

56 responses

[Copy](#)



Course Questionnaire Results

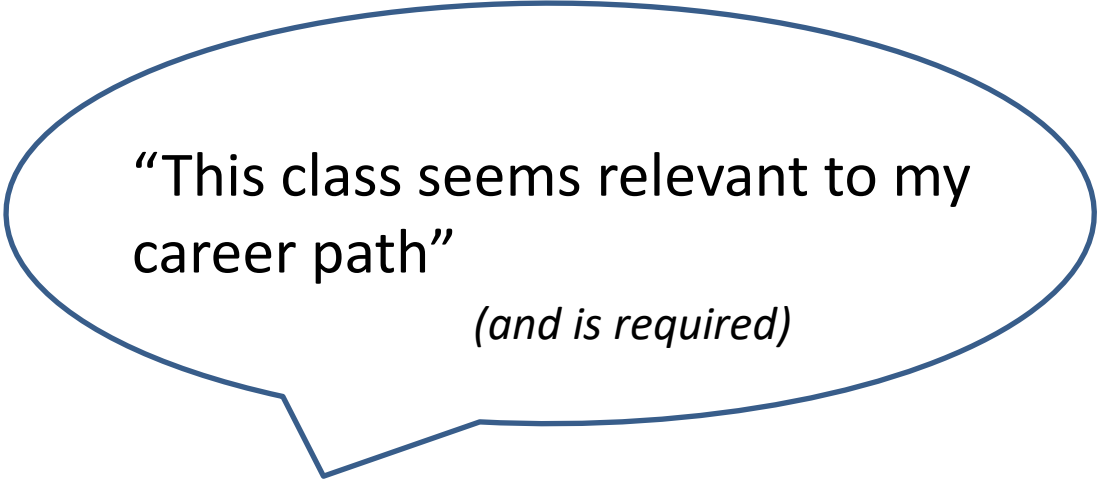


"I am a big procrastinator"

Course Questionnaire Results



"I am a big procrastinator"



"This class seems relevant to my
career path"
(and is required)

Course Questionnaire Results

“I am a big procrastinator”

“Im interested in learning
about cryptography”

“This class seems relevant to my
career path”

(and is required)

Course Questionnaire Results

“I am a big procrastinator”

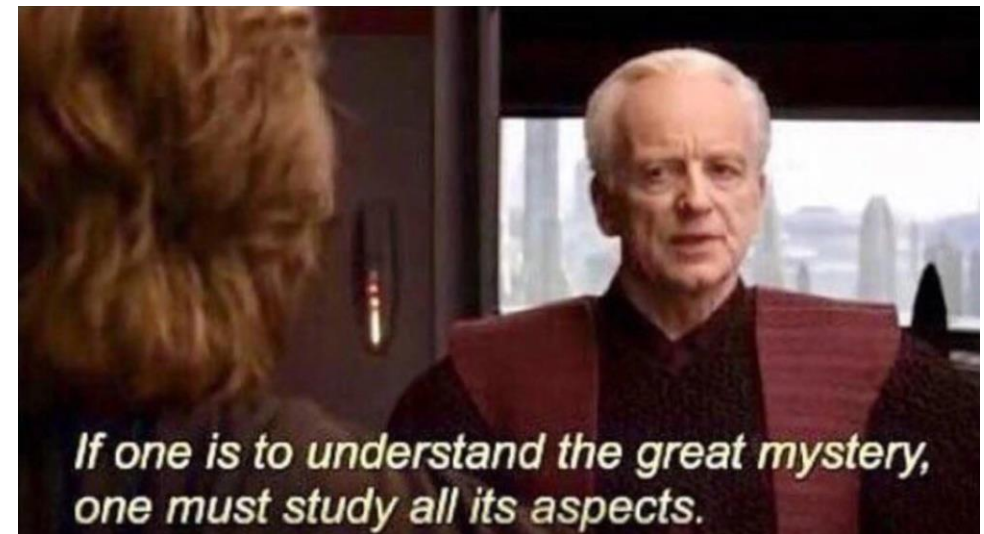
“Im interested in learning
about cryptography”

“This class seems relevant to my
career path”
(and is required)

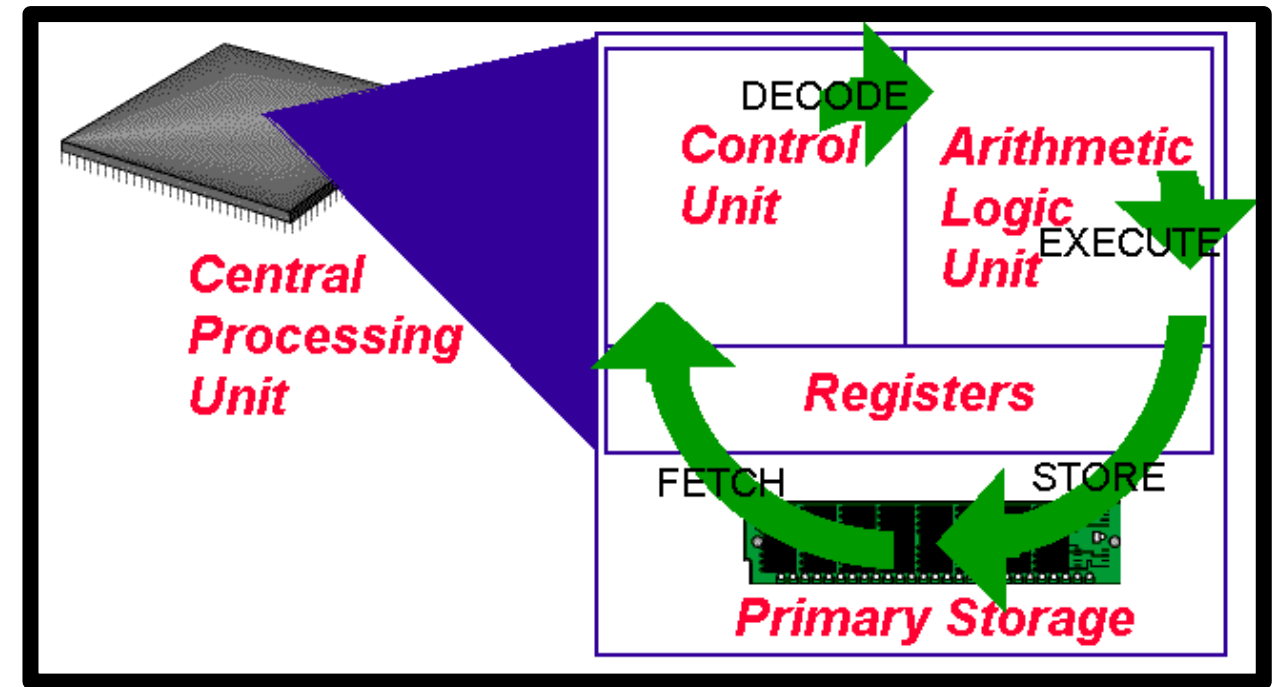
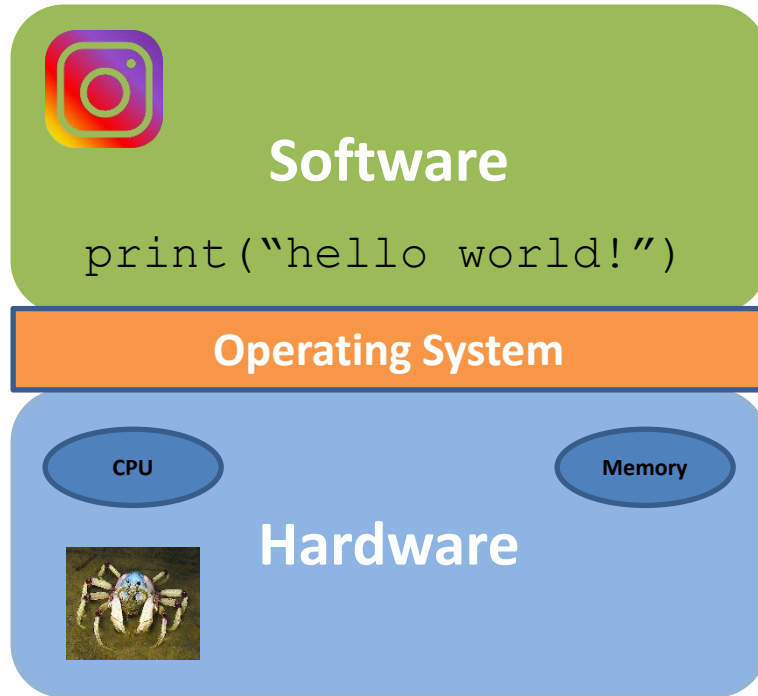
“The best cereal is *just milk*”

To understand the technical aspects of security, we must have a good understanding of how ~~computers~~ work

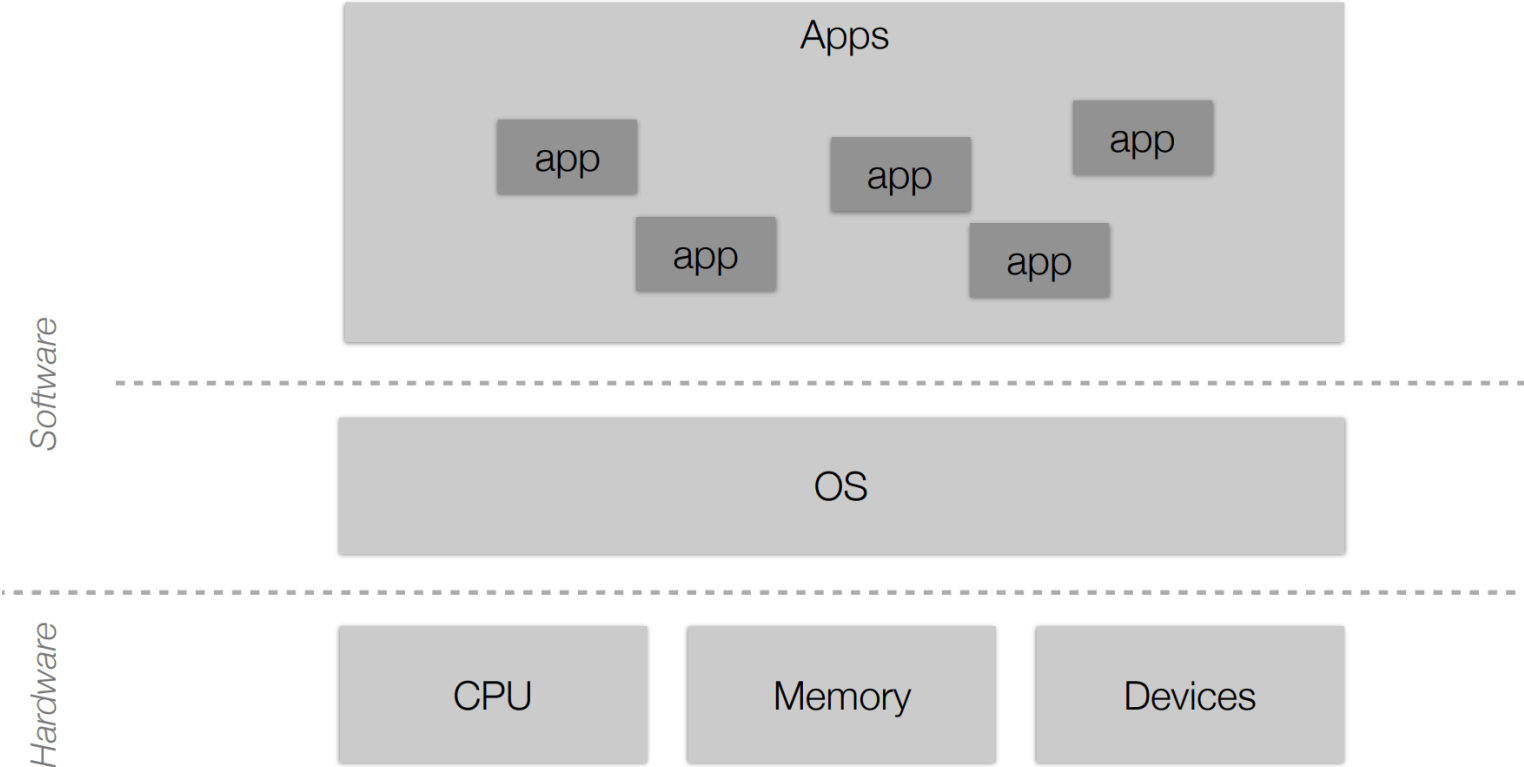
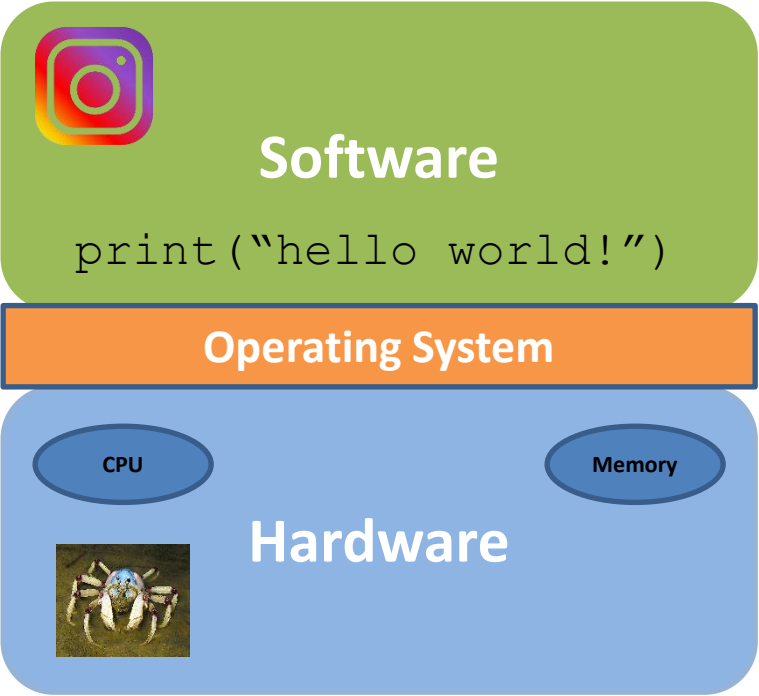
operating systems



The Operating System



The Operating System



The jobs of an Operating System

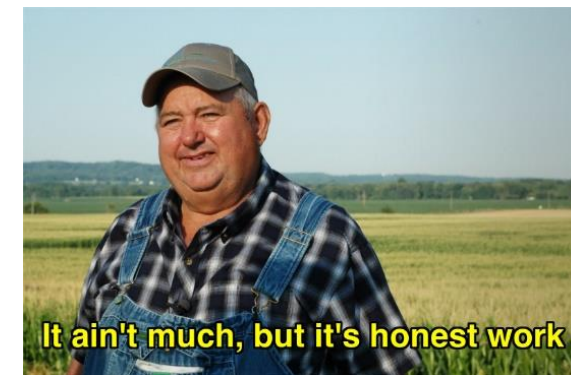
1. Process Manager
“The Coach”

2. Interface Manager
“The Bouncer”

3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”



The jobs of an Operating System

1. Process Manager
“The Coach”

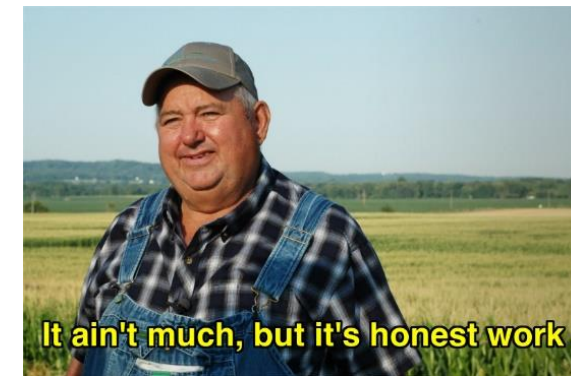
2. Interface Manager
“The Bouncer”

3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”

*This will be the focus
of the first half of
lecture*



Source code to binary

```
#include <stdio.h>

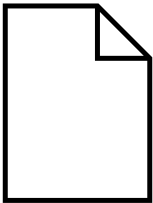
int main() {
    printf("Hello W0rld! \n");

    int x = 0;
    int y = 3;

    int z = x + y;

    printf("%d %d %d \n",x,y,z);
    return 0;
}
```

Preprocessor



- Removal of comments
- Expand Macros

Compiler

```
0000000000000000 <main>:
0: f3 0f 1e fa          endbr64
4: 55                   push    %rbp
5: 48 89 e5             mov     %rsp,%rbp
8: 48 83 ec 10          sub     $0x10,%rsp
c: 48 8d 3d 00 00 00 00 lea     0x0(%rip),%rdi    # 13 <main+0x13>
13: e8 00 00 00 00      callq  18 <main+0x18>
18: c7 45 f4 00 00 00 00 movl    $0x0,-0xc(%rbp)
1f: c7 45 f8 03 00 00 00 movl    $0x3,-0x8(%rbp)
26: 8b 55 f4             mov     -0xc(%rbp),%edx
29: 8b 45 f8             mov     -0x8(%rbp),%eax
2c: 01 d0               add     %edx,%eax
2e: 89 45 fc             mov     %eax,-0x4(%rbp)
31: 8b 4d fc             mov     -0x4(%rbp),%ecx
34: 8b 55 f8             mov     -0x8(%rbp),%edx
37: 8b 45 f4             mov     -0xc(%rbp),%eax
3a: 89 c6               mov     %eax,%esi
3c: 48 8d 3d 00 00 00 00 lea     0x0(%rip),%rdi    # 43 <main+0x43>
43: b8 00 00 00 00      mov     $0x0,%eax
48: e8 00 00 00 00      callq  4d <main+0x4d>
4d: b8 00 00 00 00      mov     $0x0,%eax
52: c9                   leaveq  %eax
53: c3                   retq
```

Assembler

- Converted to assembly code
- .s file

./hello_world



.exe

Linker

1	00000000	00000100	0000000000000000
2	01011110	00001100	11000010 0000000000000010
3	11101111	00010110	00000000000000101
4	11101111	10011110	0000000000001011
5	11111000	10101101	11011111 000000000010010
6	01100010	11011111	0000000000010101
7	11101111	00000010	11111011 000000000010111
8	11110100	10101101	11011111 000000000011110
9	00000011	10100010	11011111 000000000100001
10	11101111	00000010	11111011 000000000100100
11	01111110	11110100	10101101
12	11111000	10101110	11000101 000000000101011
13	00000110	10100010	11111011 000000000110001
14	11101111	00000010	11111011 000000000110100
15	01010000	11010100	000000000111011
16	00000100	00000000	00000111101

What happens when we run `./hello_world`

It gets turned into a **process**

A **process** is an instance of a running program on a computer

Processes									
Performance App history Startup Users Details Services									
Name	Status	37% CPU	54% Memory	1% Disk	1% Network	17% GPU	GPU engine	Power usage	Power usage t...
> Firefox (42)		6.5%	1,304.5 MB	0.5 MB/s	3.1 Mbps	9.0%	GPU 0 - Video Decode	High	Very low
> Google Chrome (14)		0.8%	484.9 MB	0 MB/s	0 Mbps	0%		Low	Very low
> Discord (32 bit) (6)		4.3%	328.8 MB	0 MB/s	8.7 Mbps	6.6%	GPU 0 - Video Encode	Moderate	Very low
> Search		5.0%	185.9 MB	0.2 MB/s	0.8 Mbps	0%	GPU 0 - 3D	Moderate	Very low
> Antimalware Service Executable		3.8%	178.2 MB	0.1 MB/s	0 Mbps	0%		Moderate	Very low
Google Chrome		0%	175.4 MB	0 MB/s	0 Mbps	0%		Very low	Very low
Slack		0%	95.5 MB	0 MB/s	0 Mbps	0%		Very low	Very low
Steam Client WebHelper		0%	89.1 MB	0 MB/s	0 Mbps	0%		Very low	Very low
Google Chrome		0%	82.6 MB	0 MB/s	0 Mbps	0%		Very low	Very low
> Microsoft PowerPoint (32 bit) (2)		0.1%	69.3 MB	0 MB/s	0 Mbps	0%		Very low	Very low
SteelSeries GG Core		0.2%	67.7 MB	0 MB/s	0 Mbps	0%		Very low	Very low
Steam Client WebHelper		0%	66.1 MB	0 MB/s	0 Mbps	0%		Very low	Very low

A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

1. Executable Code

2. Associated Data

3. Execution Context/Bookkeeping information

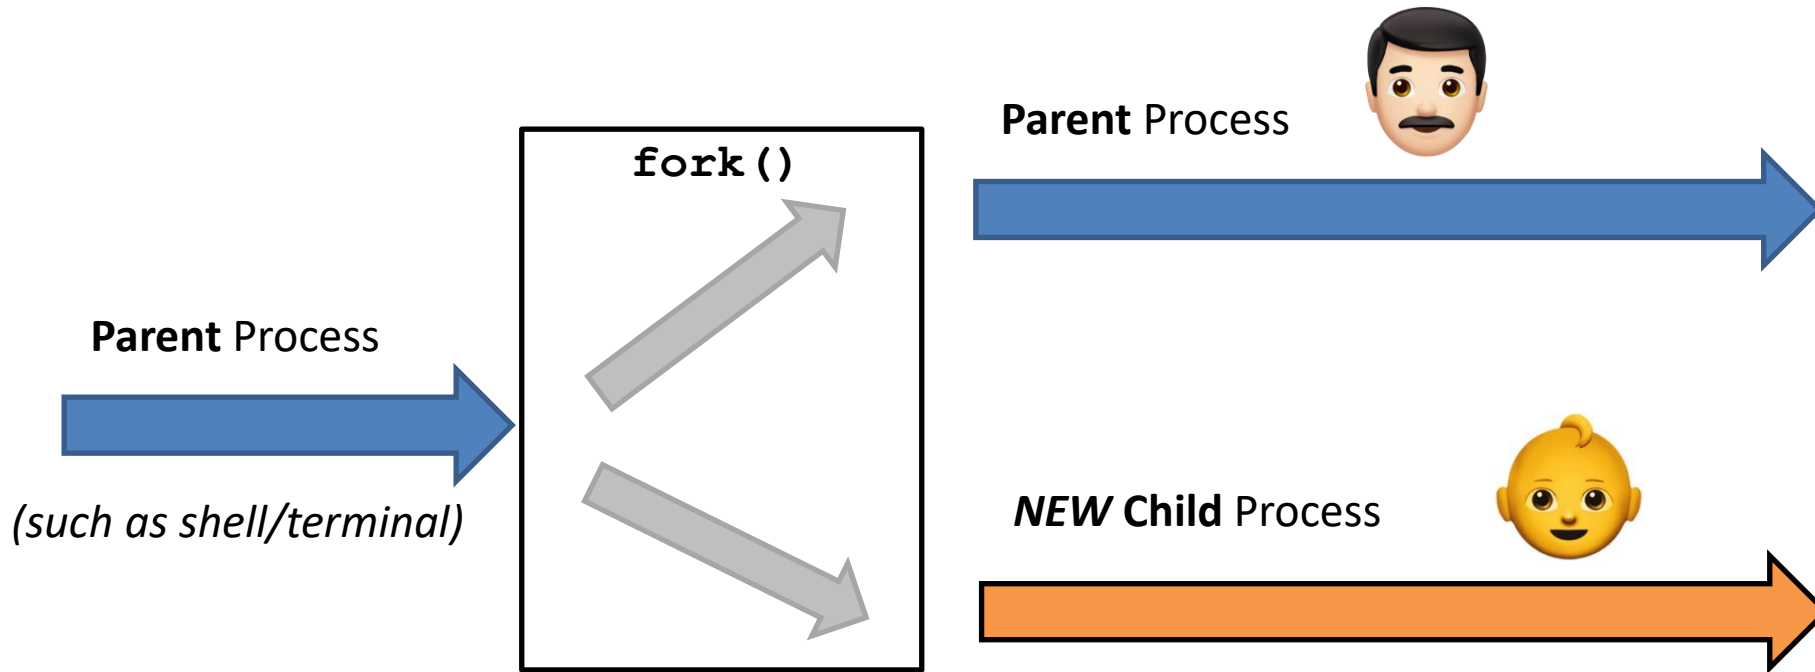
(info that the OS needs to handle the process)

Main Memory

Process A Information
Process A Data
Process A Executable Code
Process B Information
Process B Data
Process B Executable Code

Ok, but how do we *actually* create a process?

- In the Unix family (and others), we use **fork ()** to create a new process



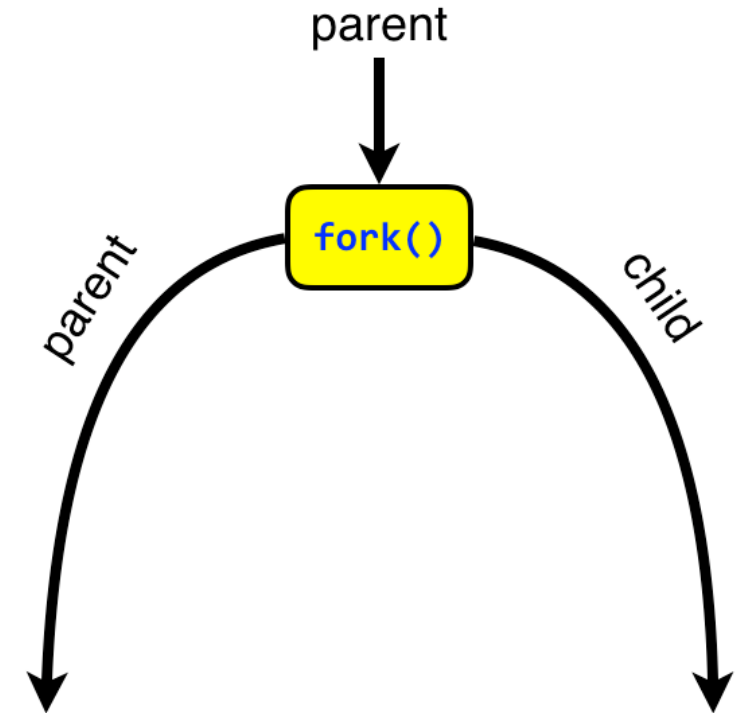
fork () duplicates a process so that instead of one process, you get two!

fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;  
  
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }  
  
    sleep(1);  
    printf("I'm the parent.");  
  
    return 0;  
}
```

We check the return value
of **fork()**!



fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;
```

We check the return value
of **fork()**!

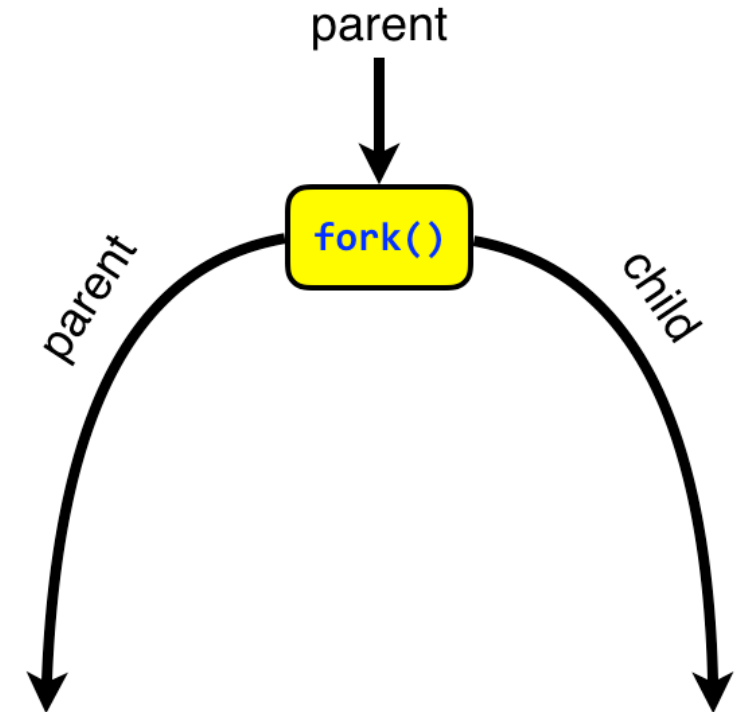
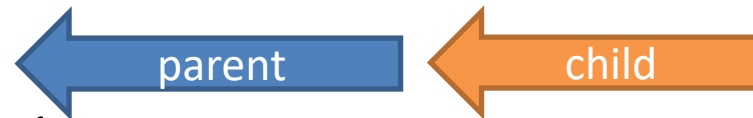
```
    pid = fork();
```

```
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }
```

```
    sleep(1);  
    printf("I'm the parent.);
```

```
    return 0;
```

```
}
```



1. Remember, **fork()** creates two process that are both actively running

fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;
```

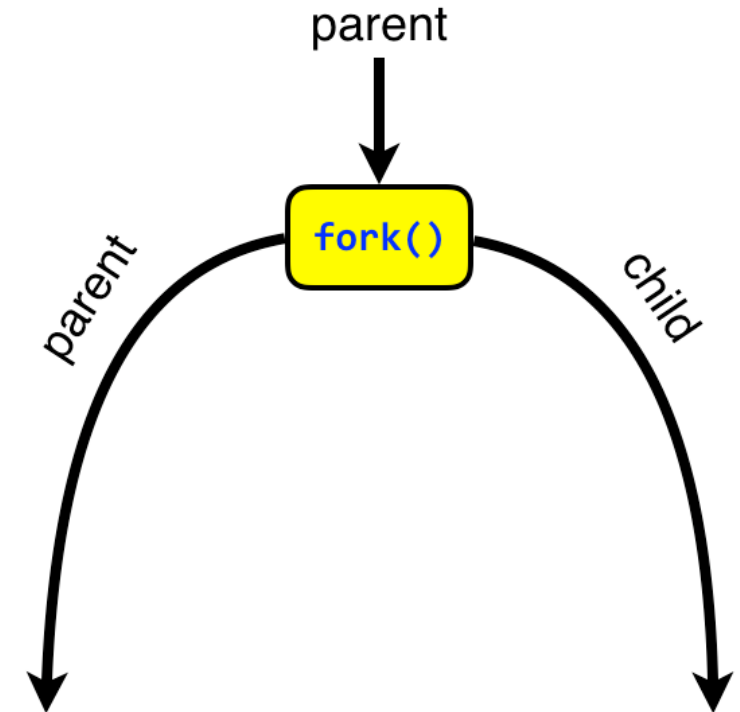
```
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }
```

```
    sleep(1);  
    printf("I'm the parent.);
```

```
    return 0;
```

```
}
```

We check the return value
of **fork()** !



2. **fork()** always returns 0 for the child process, the parent process jumps to the code after the if statement

fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;
```

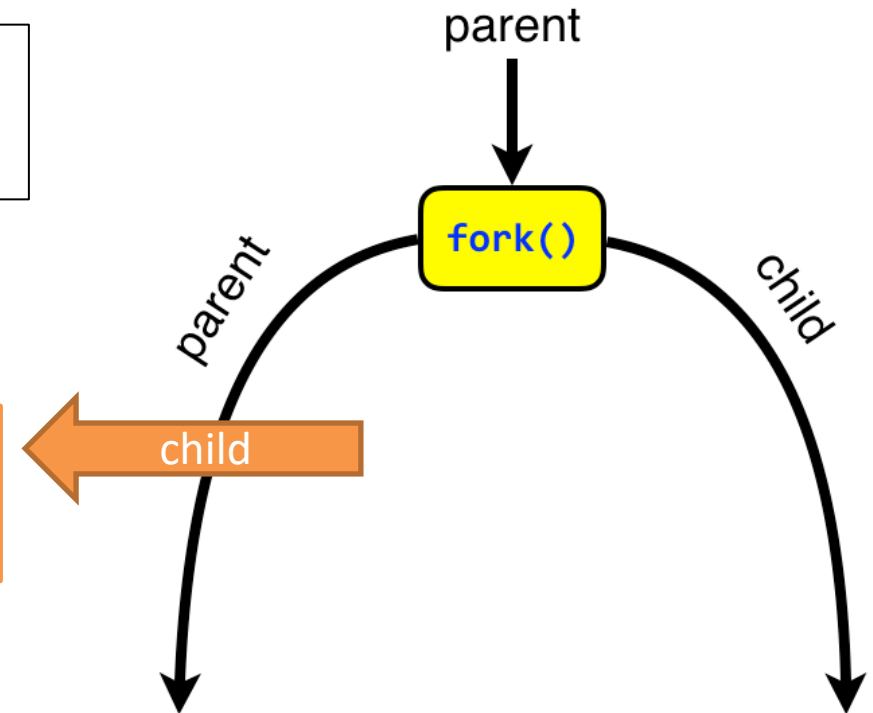
We check the return value
of **fork()**!

```
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }
```

```
    sleep(1);  
    printf("I'm the parent.);
```

```
    return 0;
```

```
}
```



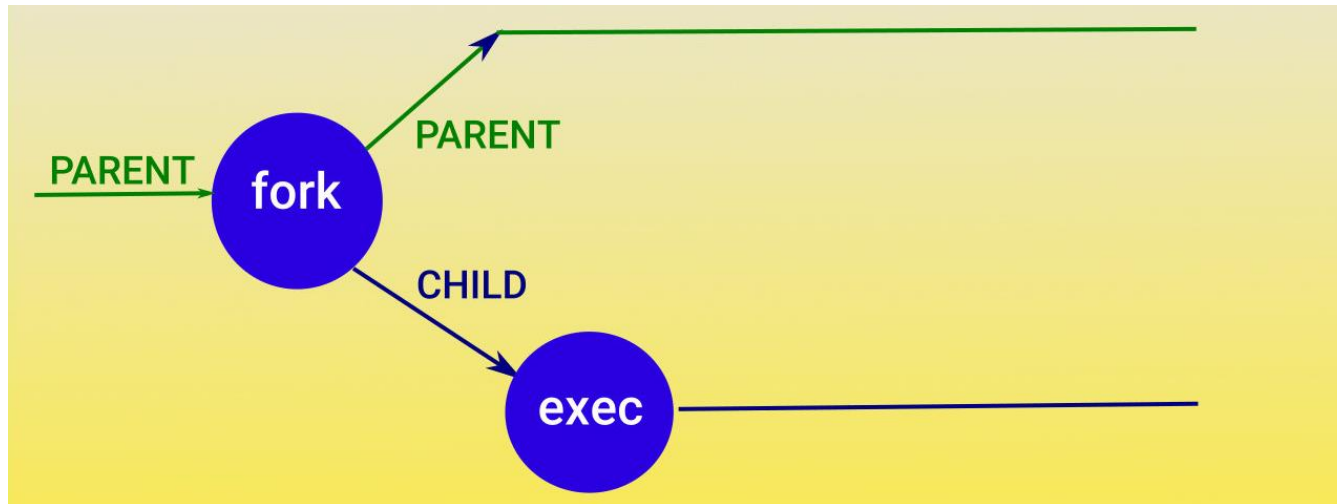
3. **fork()** always returns 0 for the child process, so the child process will execute the code in the if statement

Demo?

fork1.c

Issue: We want our child process to run an entirely new program
(hello_world c program)

We use the **exec()** family of functions to execute a different program

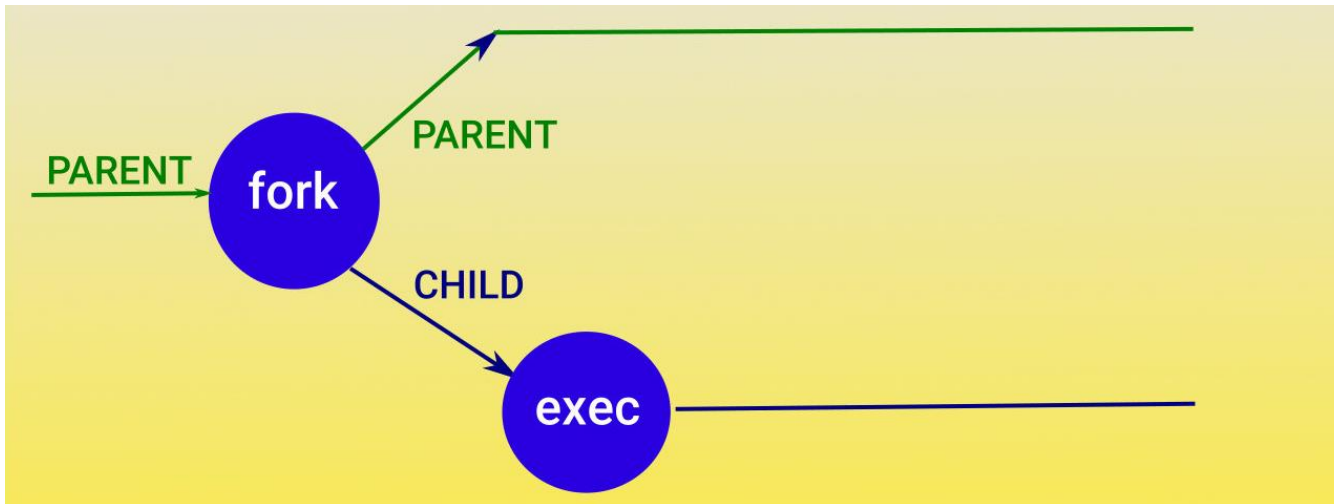


There are many different forms of the **exec()** function call

```
char *name[2];  
name[0] = "./hello";  
name[1] = NULL;  
execve(name[0], name, NULL);
```

Issue: We want our child process to run an entirely new program
(hello_world c program)

We use the **exec()** family of functions to execute a different program



There are many different forms of the **exec()** function call

```
char *name[2];  
name[0] = "./hello";  
name[1] = NULL;  
execve(name[0], name, NULL);
```

This will invoke a program called `hello`

Fork() and Exec()

```
int main(void) {  
    int pid;  
  
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
  
        char *name[2];  
        name[0] = "./hello";  
        name[1] = NULL;  
        execve(name[0], name, NULL);  
  
        _exit(0);  
    }  
    sleep(1);  
    printf("I'm the parent. My child has pid %d\n", pid);  
  
    return 0;  
}
```

Fork() and Exec()

```
int main(void) {  
    int pid;
```

```
    pid = fork();  
    if (0 == pid) {
```

```
        // I'm the child
```

```
        char *name[2];  
        name[0] = "./hello";  
        name[1] = NULL;  
        execve(name[0], name, NULL);
```

```
        _exit(0);
```

```
    }
```

```
    sleep(1);  
    printf("I'm the parent. My child has pid %d\n", pid);
```

```
    return 0;
```

```
}
```

Child code

Parent code

Fork() and Exec()

```
int main(void) {  
    int pid;  
  
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
  
        char *name[2];  
        name[0] = "./hello";  
        name[1] = NULL;  
        execve(name[0], name, NULL);  
  
        _exit(0);  
    }  
    sleep(1);  
    printf("I'm the parent. My child has pid %d\n", pid);  
  
    return 0;  
}
```

output

```
[01/25/23] seed@VM:~$ ./forkexec  
Hello from the C program!  
I'm the parent. My child has pid 33578
```

Demo?

forkandexec.c

Tl;dr

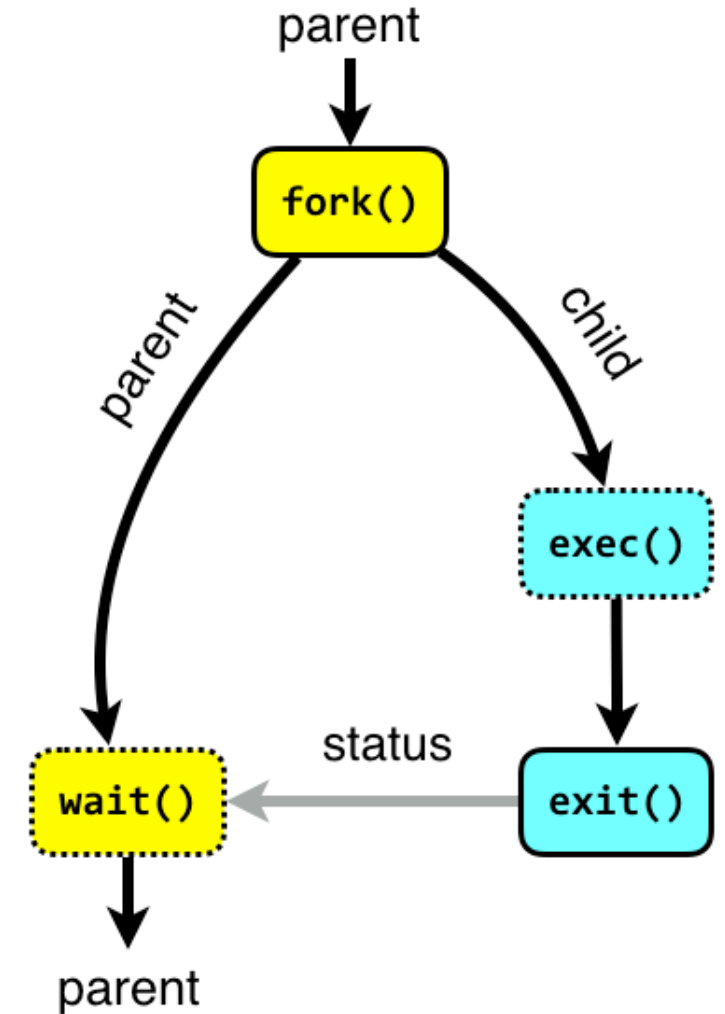
The programs we run get turned into a **process**

fork() is used to create a new process

- The parent process is typically the shell/terminal, and waits for the child process to finish
- The child process runs **exec()** to run our program

Contents	
9.4	Process Primitives
9.4.1	Having Children
9.4.2	Watching Your Children Die
9.4.3	Running New Programs
9.4.4	A Bit of History: vfork()
9.4.5	Killing Yourself
9.4.6	Killing Others
9.4.7	Dumping Core
9.5	Simple Children

you can kill children with the `kill()` function or `kill` command



```
#include <sys/types.h>
#include <unistd.h>
```

```
int main()
{
    while(1) {
        fork();
    }
    return 0;
}
```

Any ideas what might happen?

```
#include <sys/types.h>
#include <unistd.h>
```

```
int main()
{
    while(1) {
        fork();
    }
    return 0;
}
```



“Oh, these forks() aren’t homemade. They were made in factory. A **fork()** **bomb** factory. This is a **fork() bomb**”

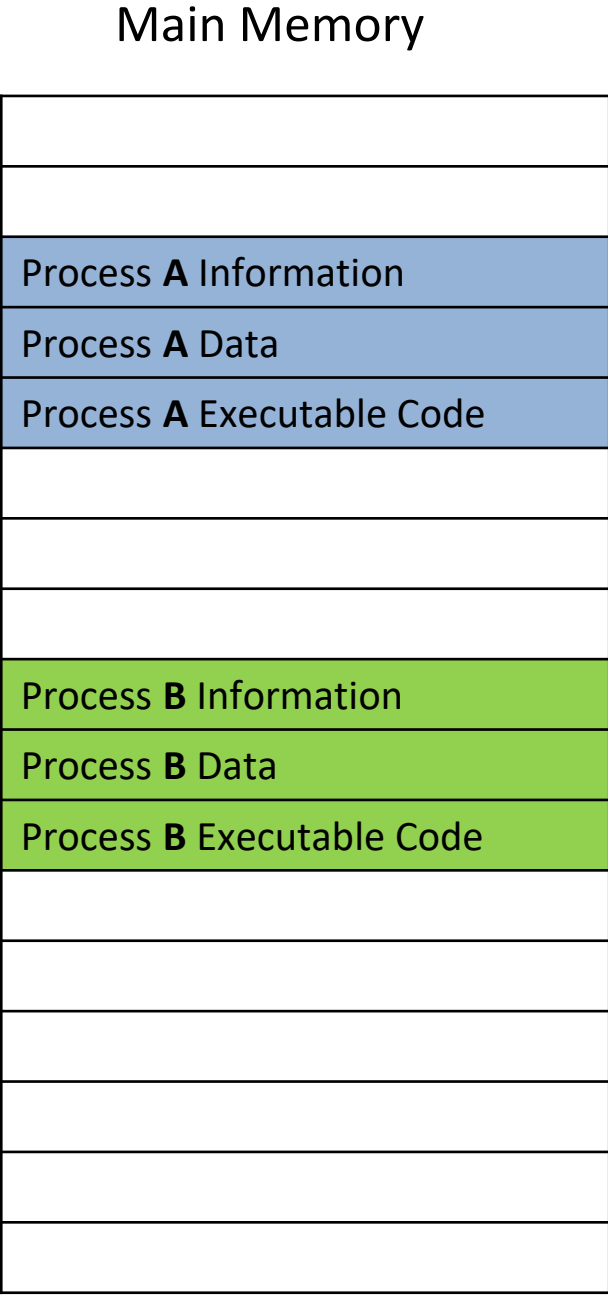


A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

- 1. Executable Code
- 2. Associated Data
- 3. Execution Context/Bookkeeping information

(info that the OS needs to handle the process)



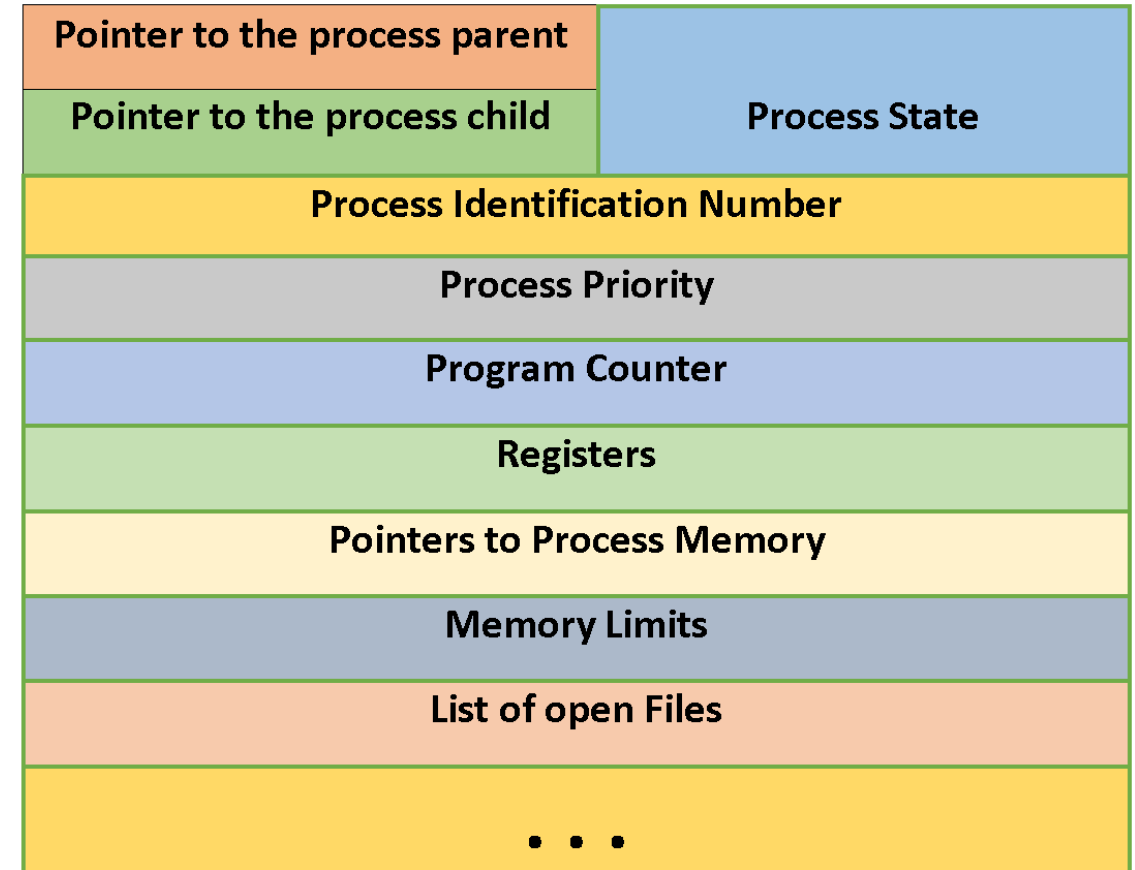
3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

Example PCB:



Created by NotesJam

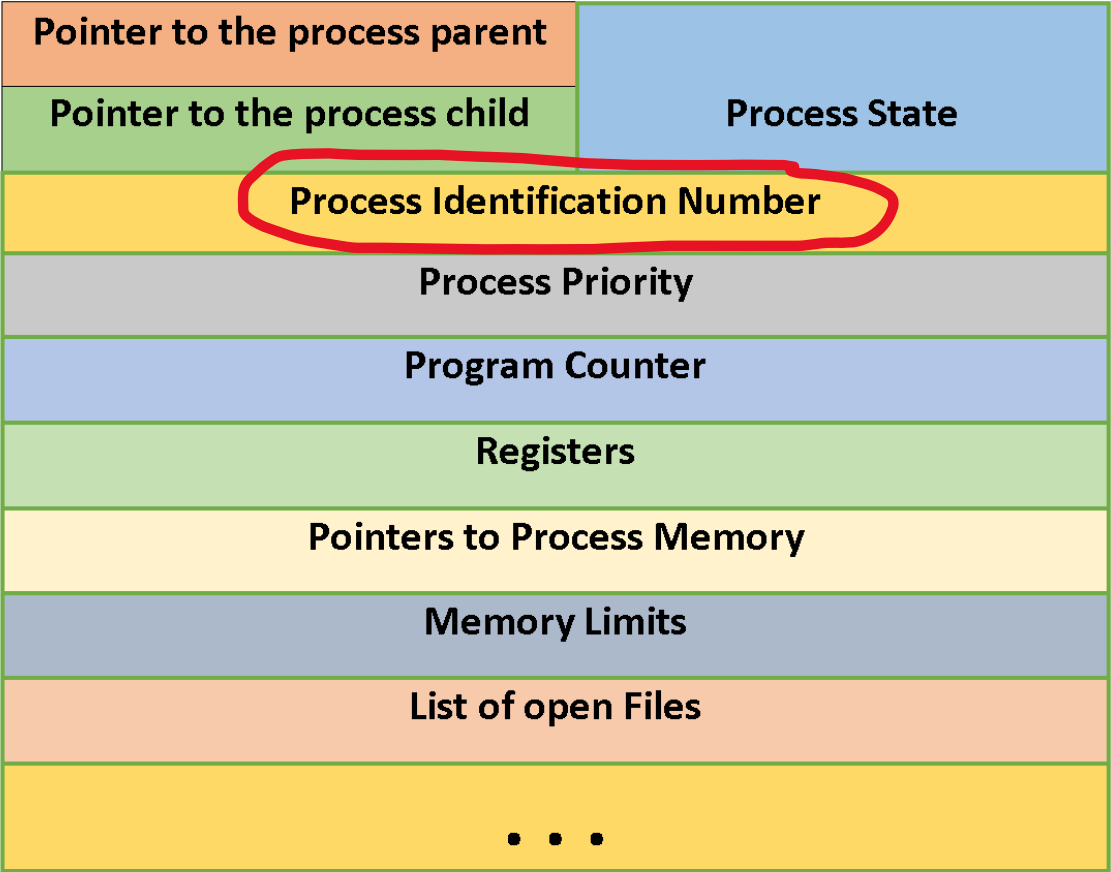
3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**
 - Simply just a data structure that holds information
 - The name of this varies by OS

Every process has a unique process ID (PID)

Process Name	User	% CPU	ID	Memory	Disk read toti D
at-spi2-registryd	seed	0	1870	196.0 KiB	120.0 KiB
at-spi-bus-launcher	seed	0	1779	292.0 KiB	28.0 KiB
bash	seed	0	16245	1.6 MiB	3.1 MiB
bash	seed	0	20664	1.8 MiB	72.7 MiB
dbus-daemon	seed	0	1560	1.5 MiB	420.0 KiB

Example PCB:



We can use the PID to search for process, kill process, fork new process, etc

Created by NotesJam

3. Execution Context/Bookkeeping information

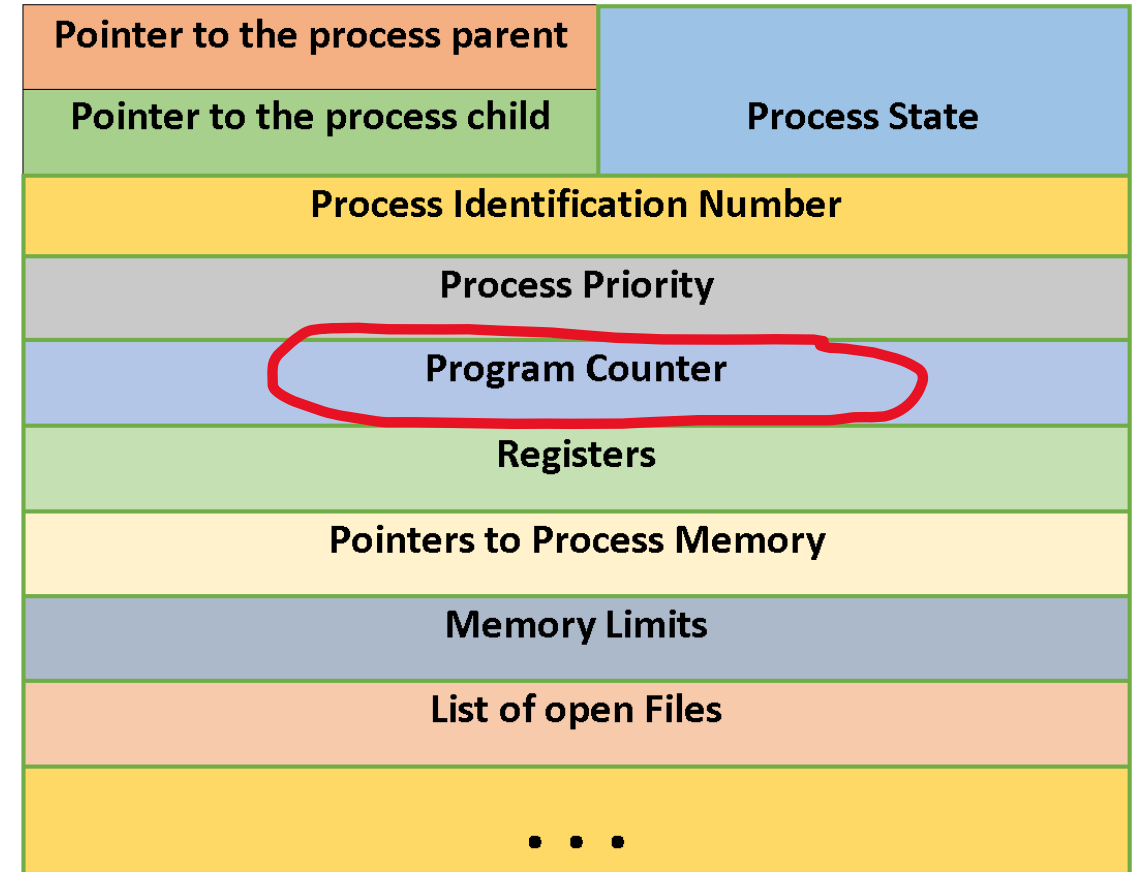
- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

Each process has a program counter (PC), which tells the CPU the next instruction to run in the process

Example PCB:



Created by NotesJam

3. Execution Context/Bookkeeping information

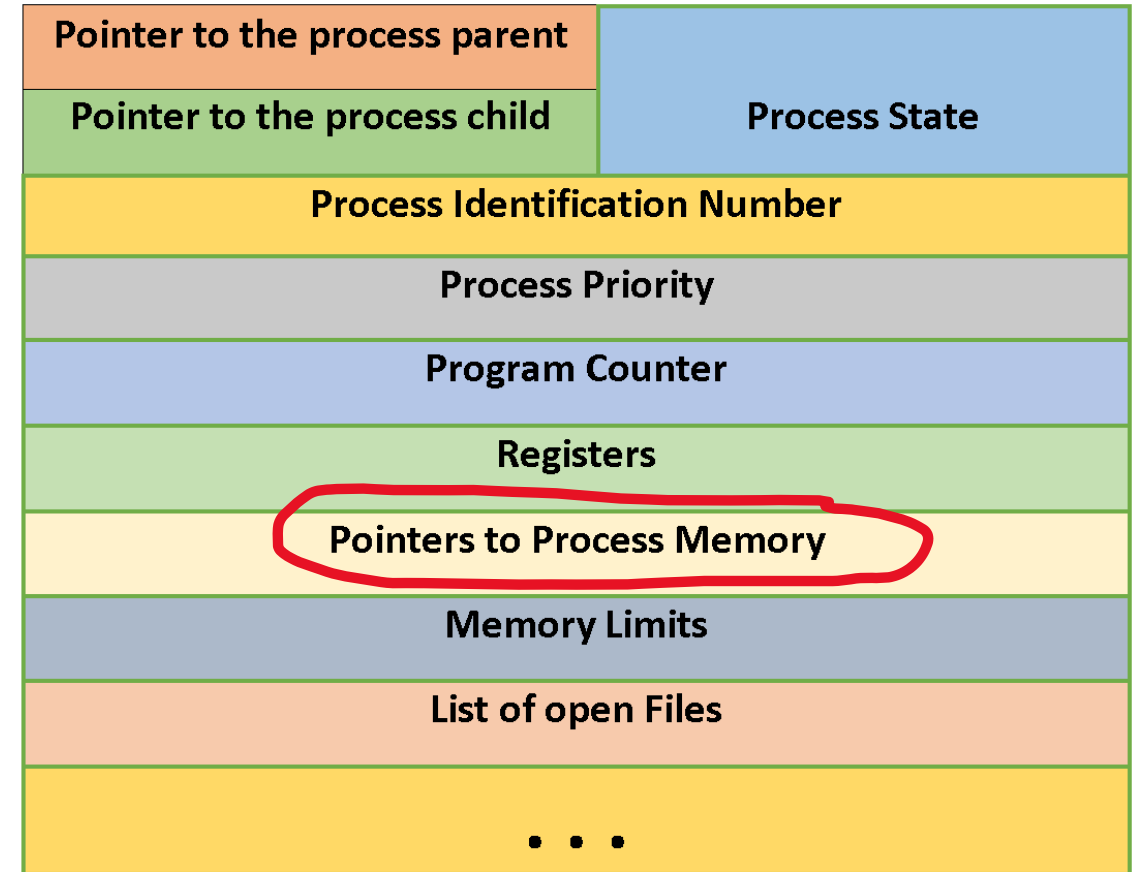
- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

PCB also maintains locations for the process Data and Code

Example PCB:



Created by NotesJam

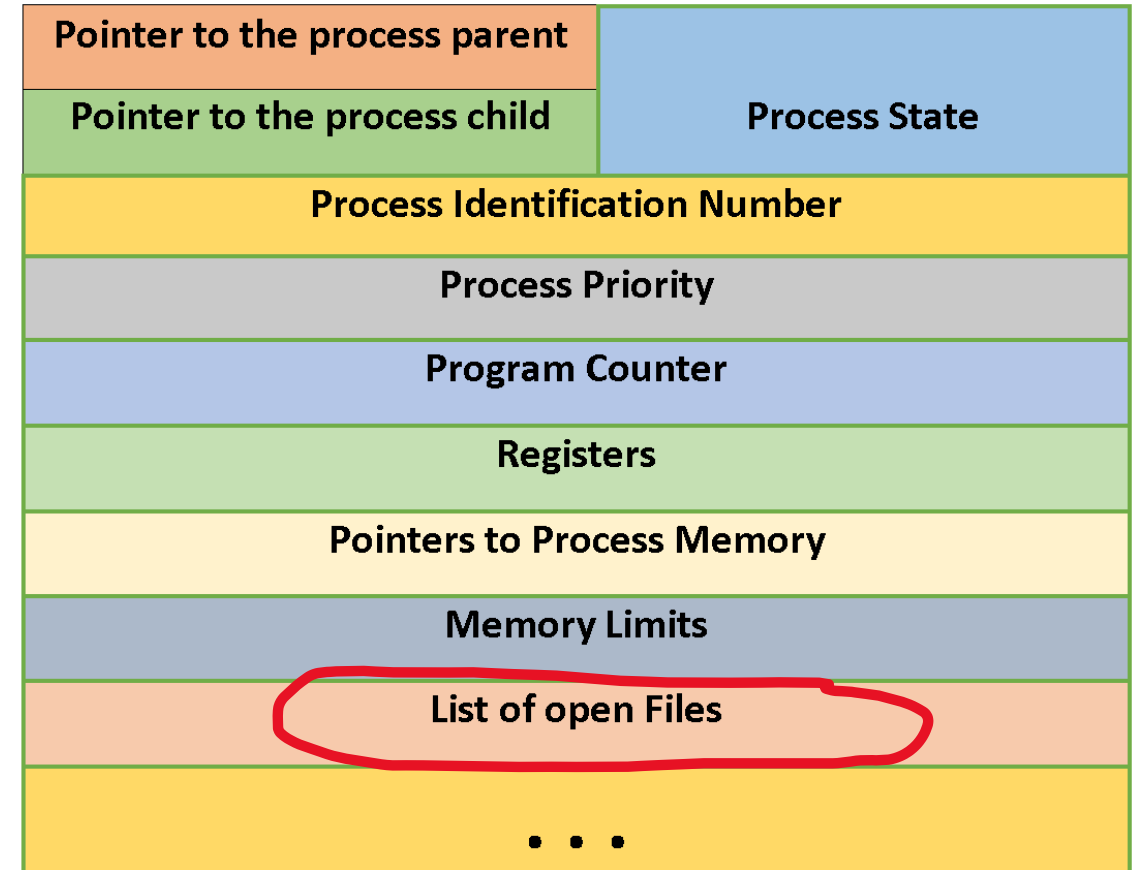
3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

Example PCB:



Created by NotesJam

3. Execution Context/Bookkeeping information

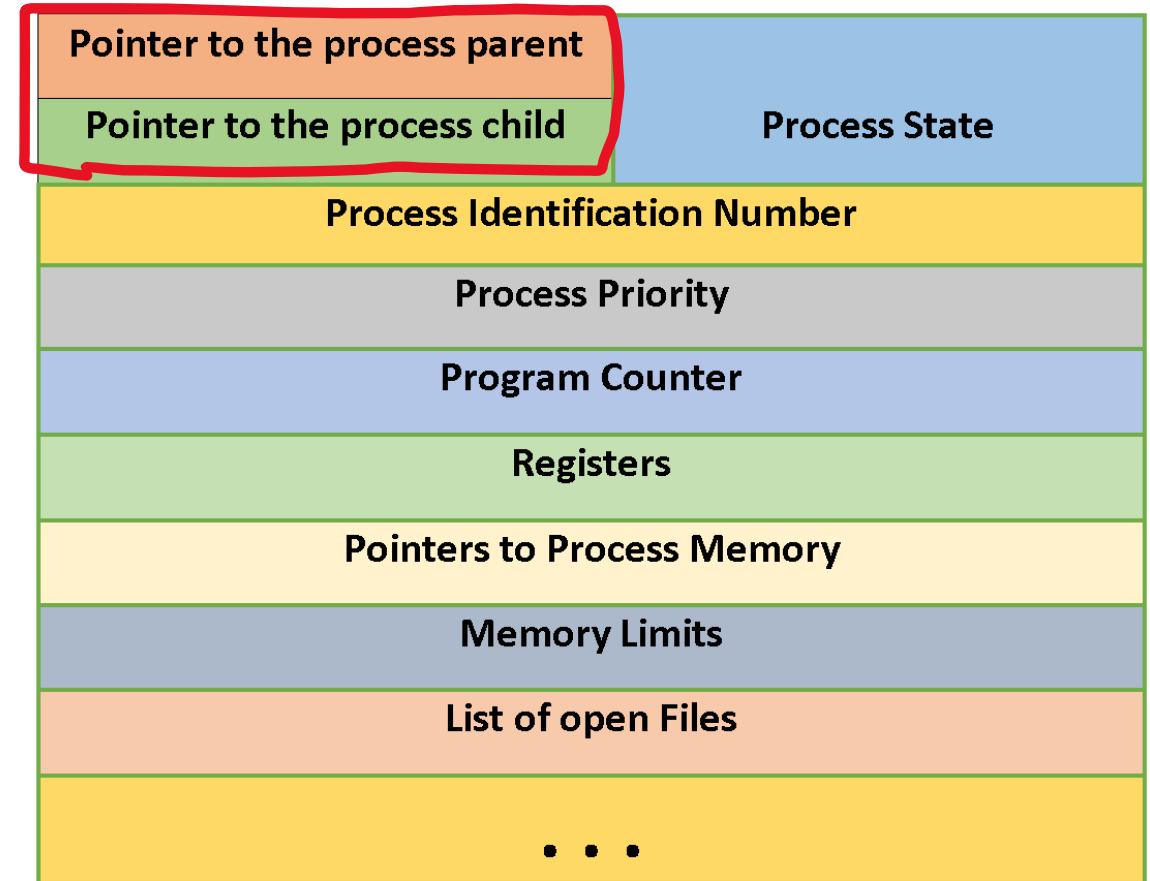
- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

PCB keeps track of who their parent is, and any child process (good parenting)

Example PCB:



Created by NotesJam

3. Execution Context/Bookkeeping information

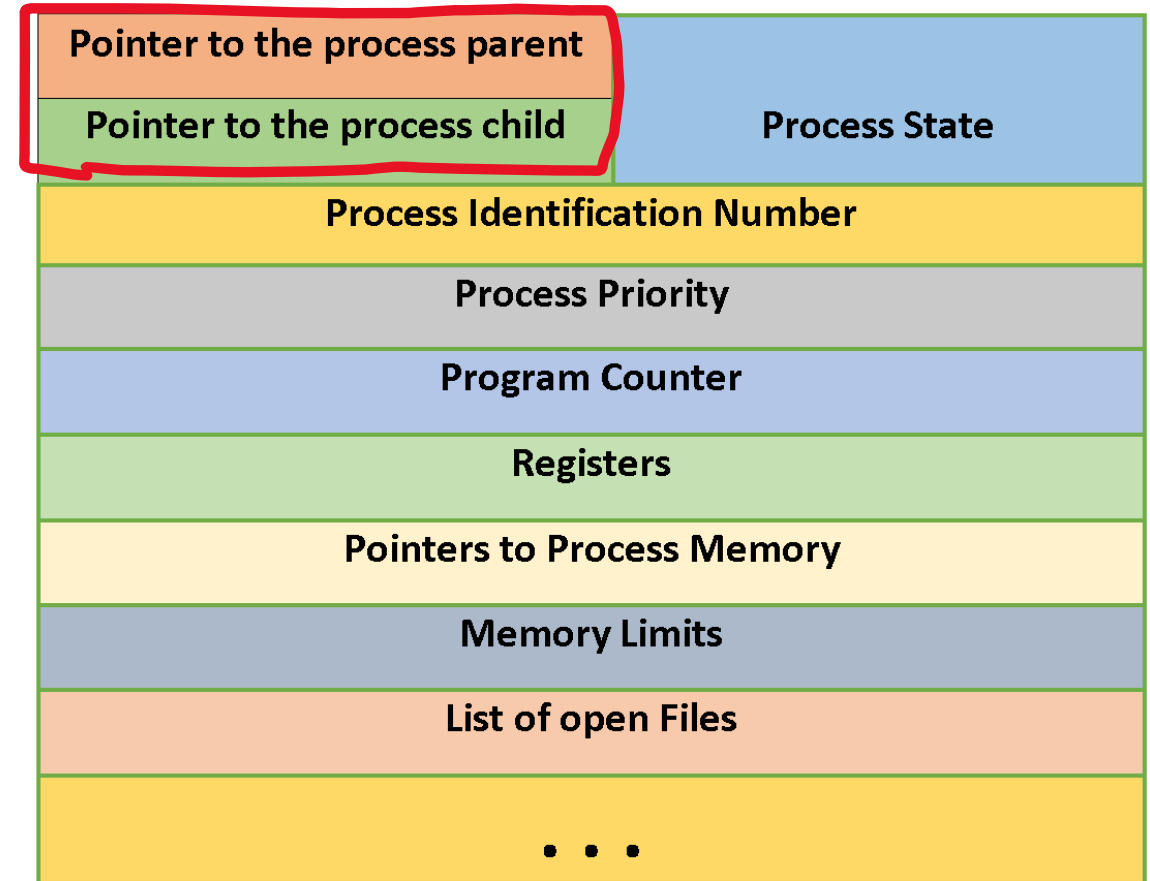
- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

PCB keeps track of who their parent is, and any child process (good parenting)

Example PCB:



Created by NotesJam

3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**

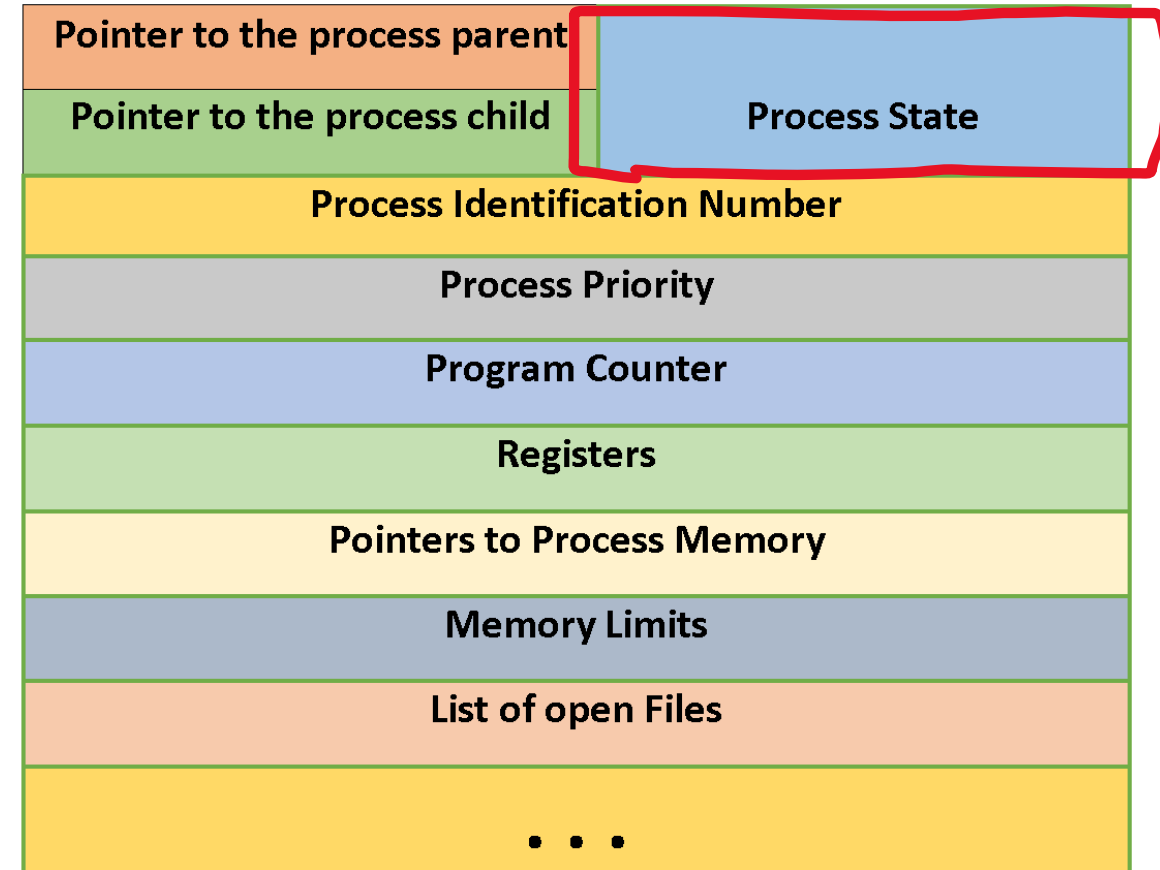
→ Simply just a data structure that holds information

→ The name of this varies by OS

A process goes through many **states**

- Active (running)**
- Blocked**
- Waiting**
- Suspended**

Example PCB:



Created by NotesJam

A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

1. Executable Code

2. Associated Data

3. Execution Context/Bookkeeping information

(info that the OS needs to handle the process)

We will talk about what goes here shortly

Pointer to the process parent	Process State
Pointer to the process child	
Process Identification Number	
Process Priority	
Program Counter	
Registers	
Pointers to Process Memory	
Memory Limits	
List of open Files	
...	

Created by NotesJam

The jobs of an Operating System

1. Process Manager

“The Coach”

The OS manages many active processes all at once, and they must create processes, manage current process, and control which processes do what



`./hello_world`



Fork() and exec()



Program is now
running as a
process

A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

1. Executable Code

2. Associated Data

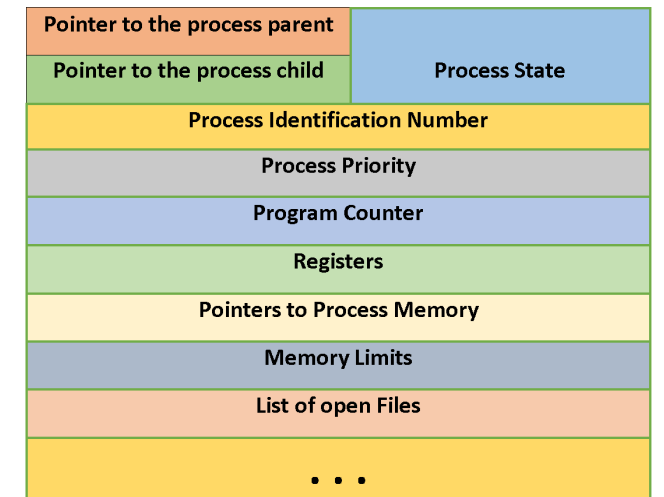
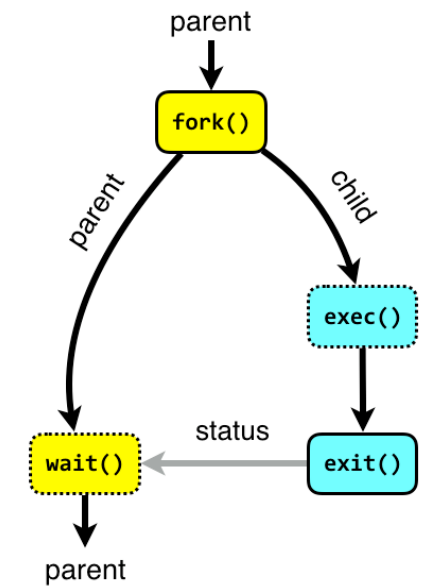
3. Execution Context/Bookkeeping information

(info that the OS needs to handle the process)

`./hello_world`

Fork() and exec()

Program is now
running as a
process



Created by NotesJam

Demo time!

```
int main(void) {
    int pid;

    pid = fork();
    if (0 == pid) {
        // I'm the child
        printf("Hi, I'm the child. \n");
    }

    sleep(1);
    // we could wait() here
    printf("I'm the parent.");

    return 0;
}
```

```
int main(void) {
    int pid;

    pid = fork();
    if (0 == pid) {
        // I'm the child

        char *name[2];
        name[0] = "./hello";
        name[1] = NULL;
        execve(name[0], name, NULL);

        _exit(0);
    }
    sleep(1);
    printf("I'm the parent. My child

    return 0;
}
```

The jobs of an Operating System

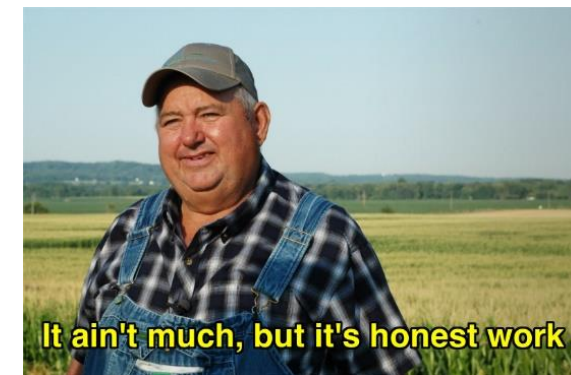
1. Process Manager
“The Coach”

2. Interface Manager
“The Bouncer”

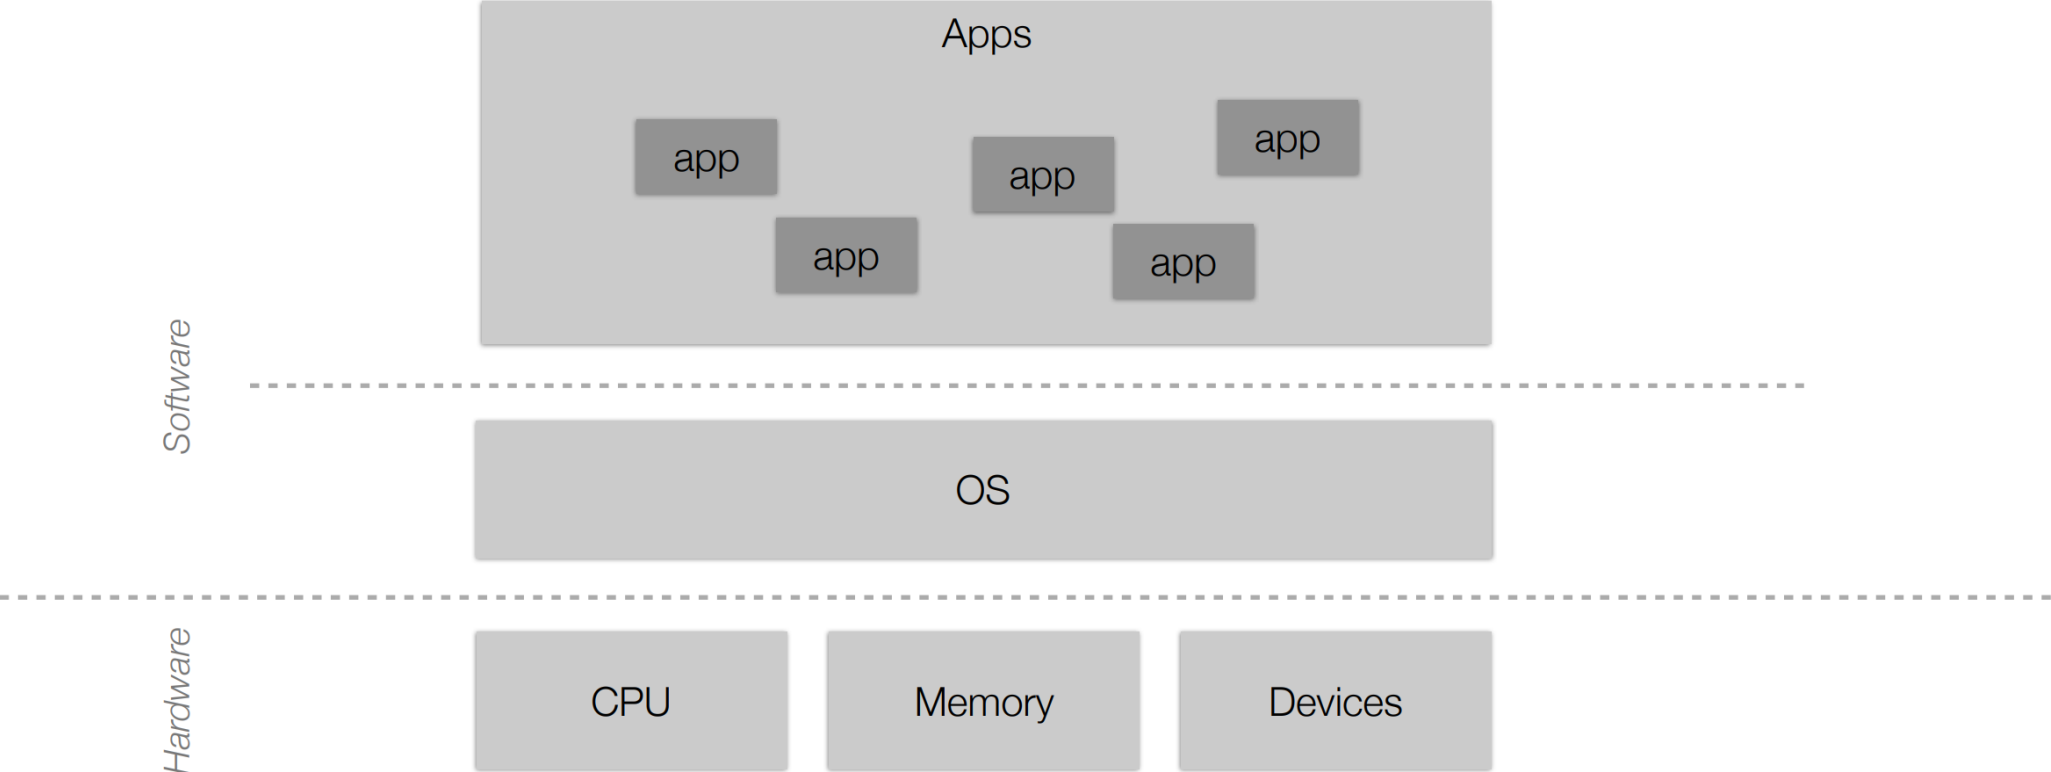
3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”



Operating Systems Review

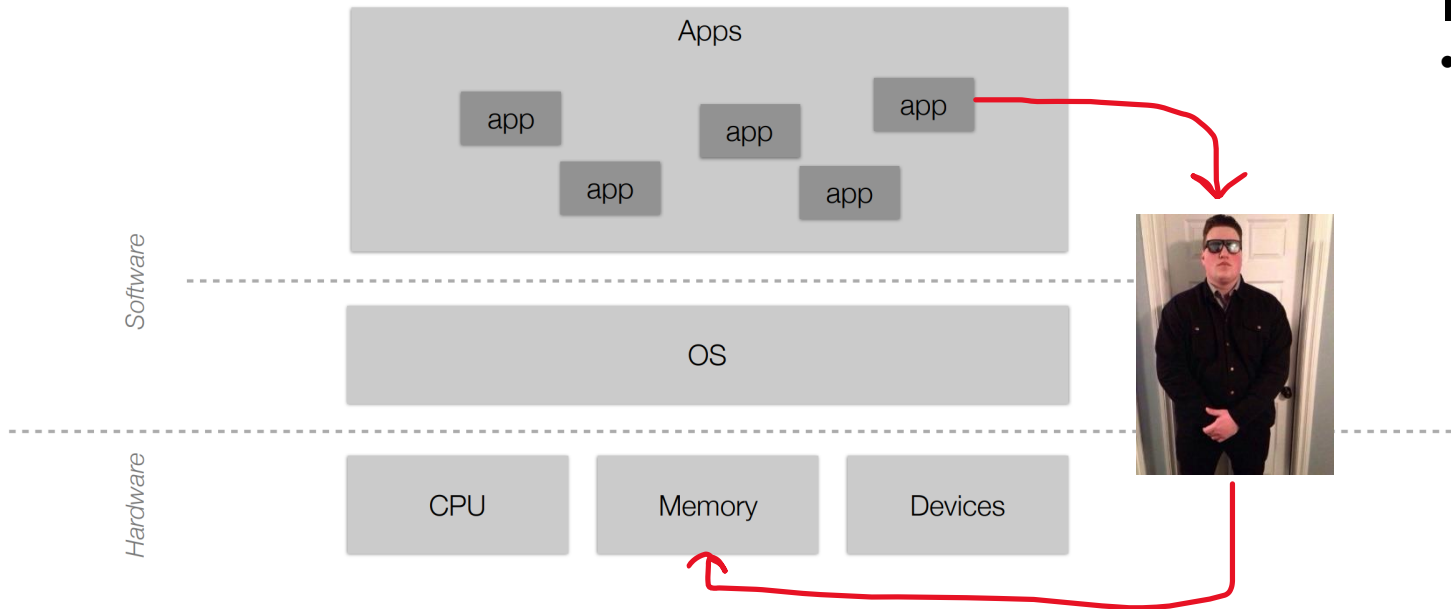


Operating Systems Review

Responsibilities of the OS?

Interface Manager

- Manages communication between apps and hardware



Operating Systems Review

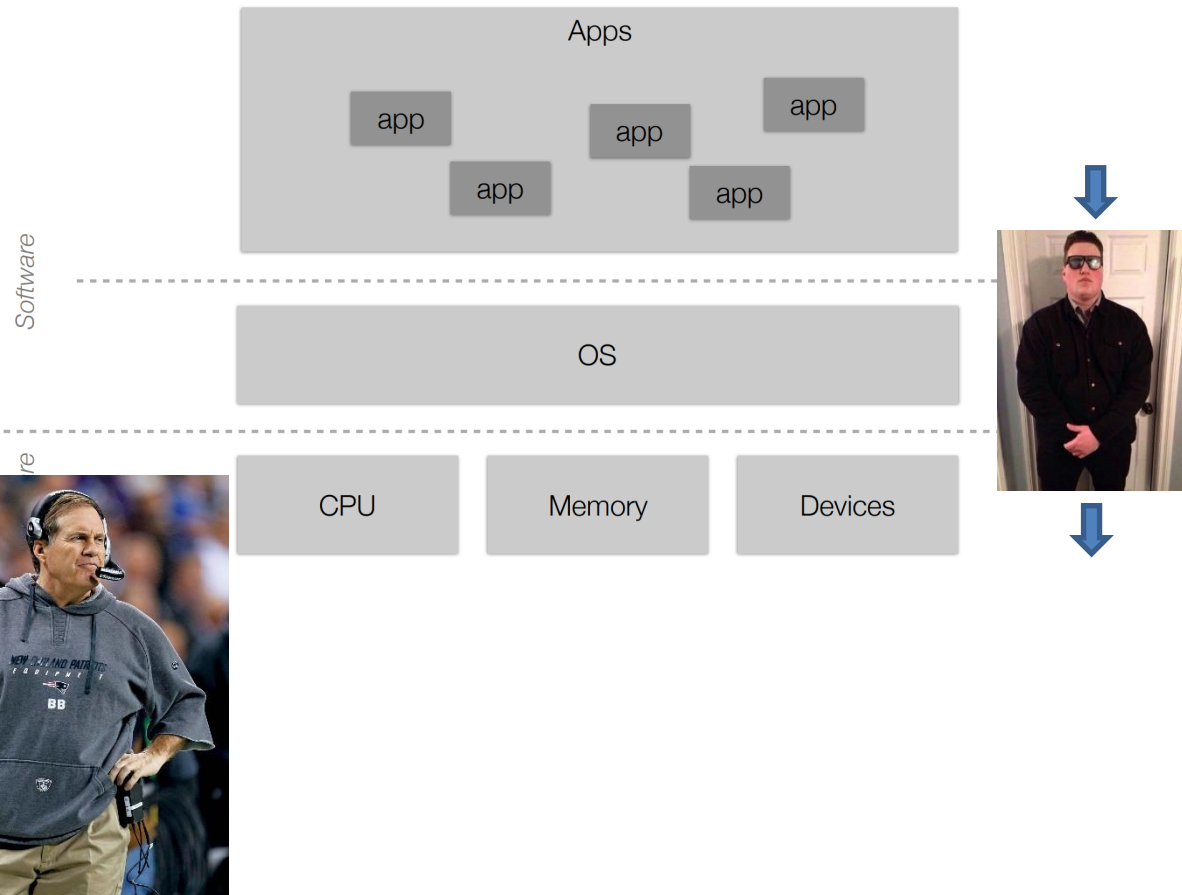
Responsibilities of the OS?

Interface Manager

- Manages communication between apps and hardware

Process Manager

- Manages how processes are structured and how to handle many processes running at once



Operating Systems Review

Responsibilities of the OS?

Interface Manager

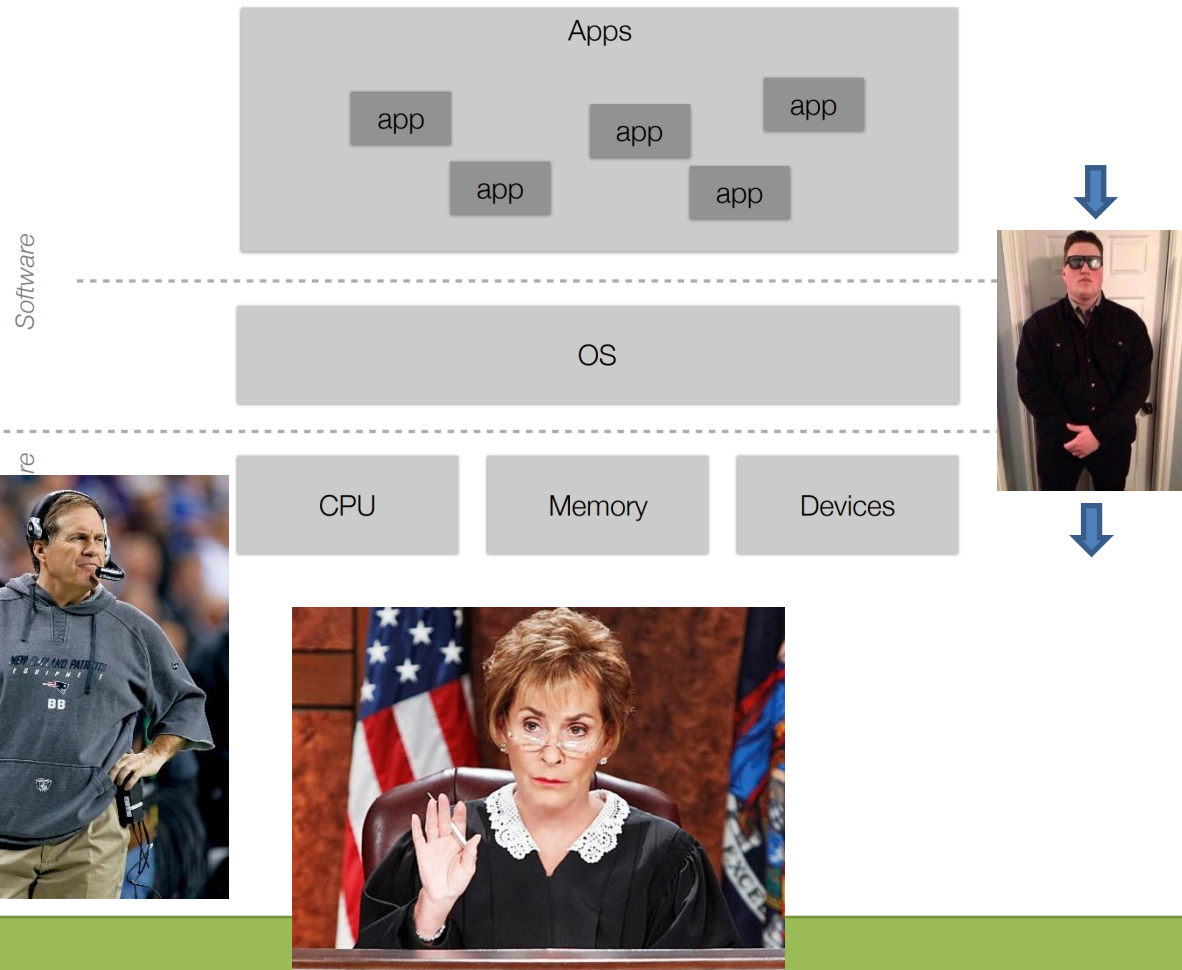
- Manages communication between apps and hardware

Process Manager

- Manages how processes are structured and how to handle many processes running at once

Traffic Manager

- Manages which programs should be executed by the CPU



Operating Systems Review

Responsibilities of the OS?

Interface Manager

- Manages communication between apps and hardware

Process Manager

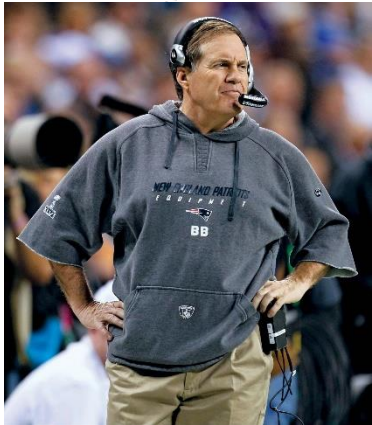
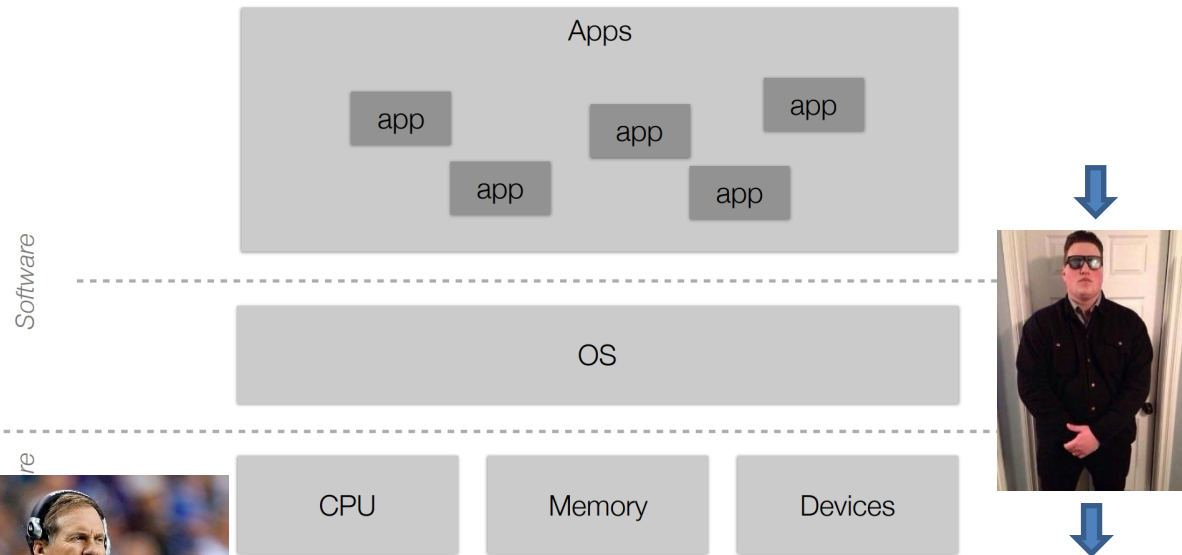
- Manages how processes are structured and how to handle many processes running at once

Traffic Manager

- Manages which programs should be executed by the CPU

Memory Manager

- Manages how physical memory is utilized



It ain't much, but it's honest work

Operating Systems Review

Responsibilities of the OS?



Interf

- M
- ha

Pro

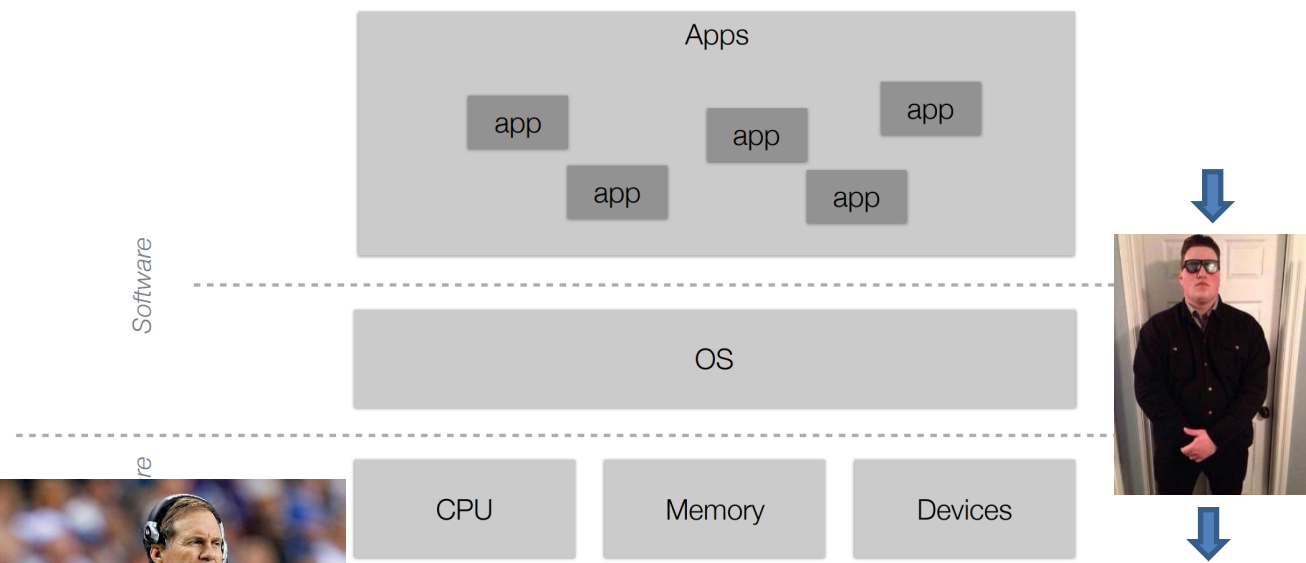
- M
- to

Traffi

- Mar
- the

Mem

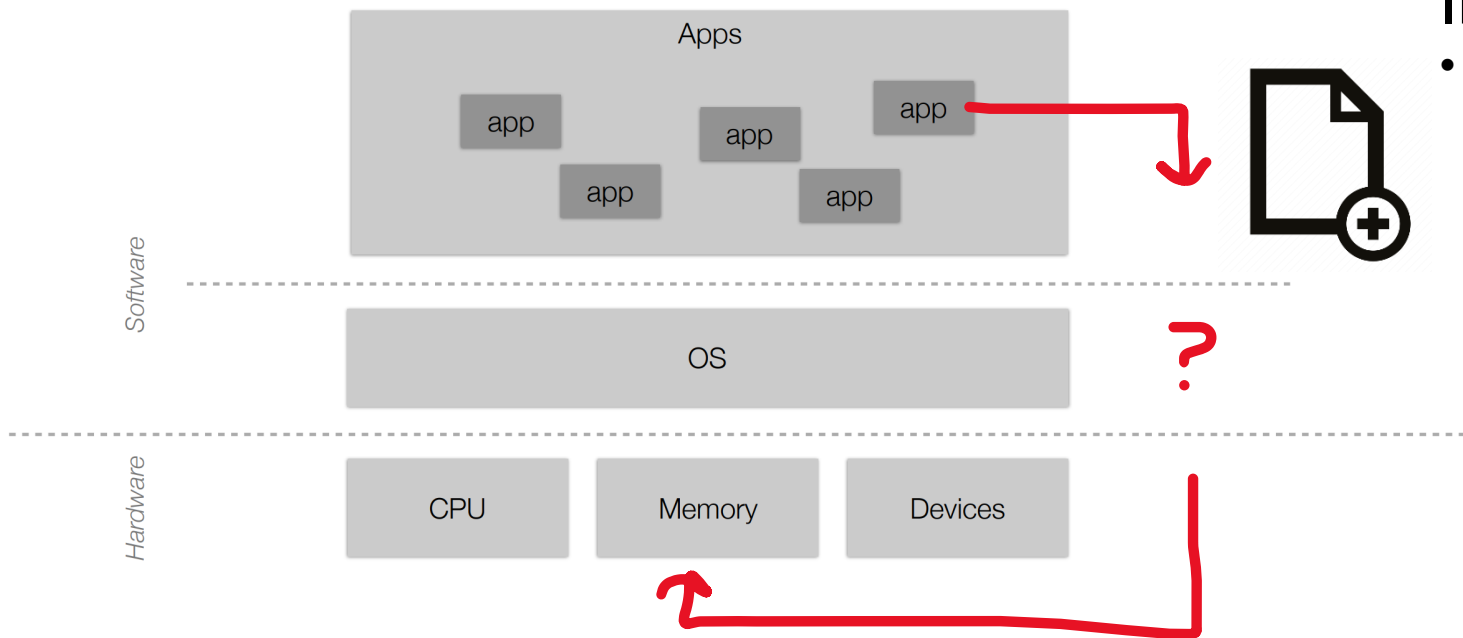
- Mar



It ain't much, but it's honest work

Operating Systems Review

Responsibilities of the OS?



Interface Manager

- Manages communication between apps and hardware

How does an application get access to a computer's resources?



Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
int main(void)
{
    printf("Hello, World!\n");

    return 0;
}
```


Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

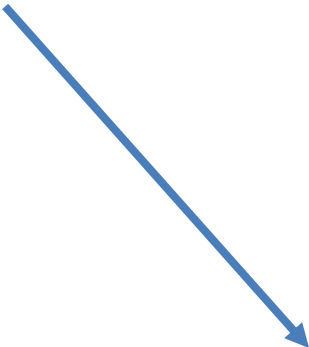
```
int main(void)
{
    printf("Hello, World!\n");

    return 0;
}
```



```
int main(void)
{
    write(1, "Hello, World!\n", 14);

    return 0;
}
```



```
int main(void)
{
    syscall(SYS_write, 1, "Hello, World!\n", 14);

    return 0;
}
```

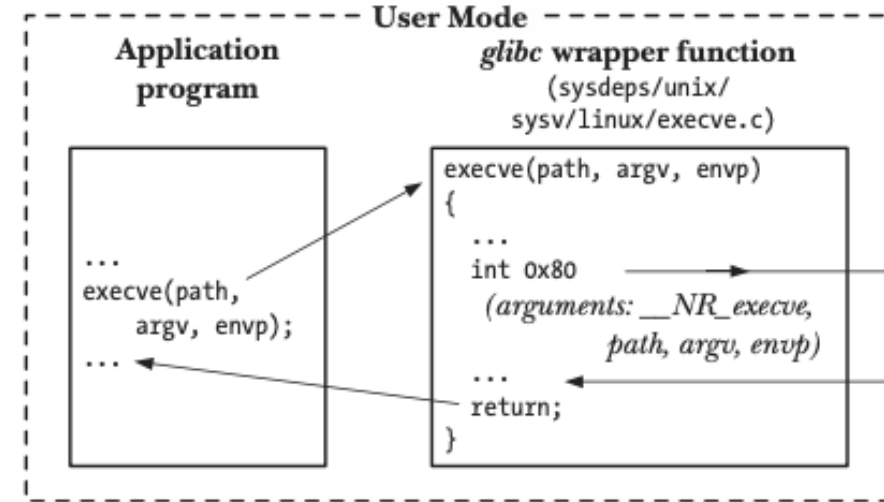
Number	Name	Description
1	exit	terminate process execution
2	fork	fork a child process
3	read	read data from a file or socket
4	write	write data to a file or socket
5	open	open a file or socket
6	close	close a file or socket
37	kill	send a kill signal
90	old_mmap	map memory
91	munmap	unmap memory
301	socket	create a socket
303	connect	connect a socket

Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

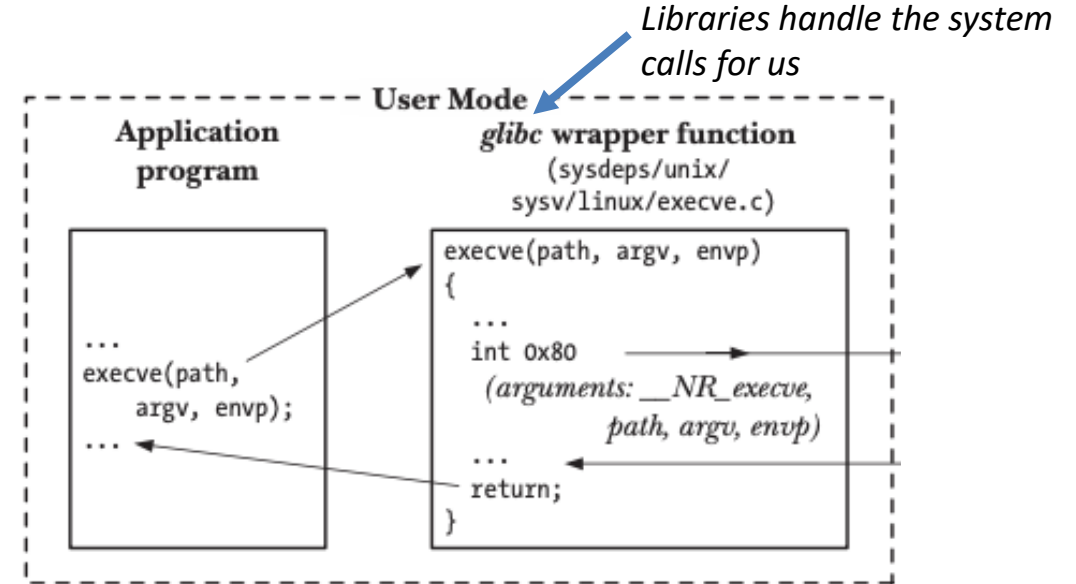


Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

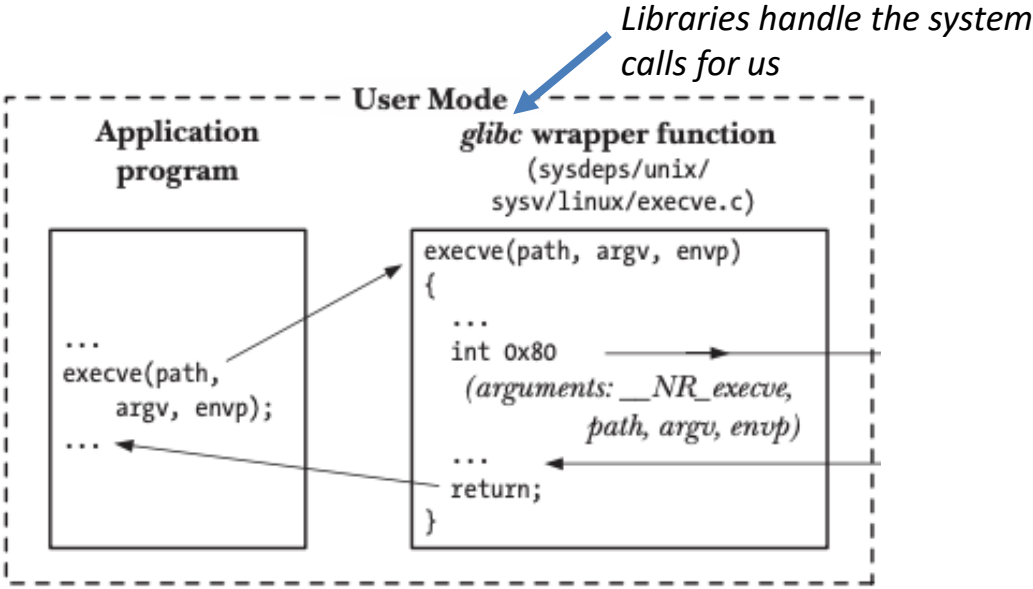


Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```



The operating system have hundreds of different syscalls, and different syscalls have different parameters, we need a way to distinguish them

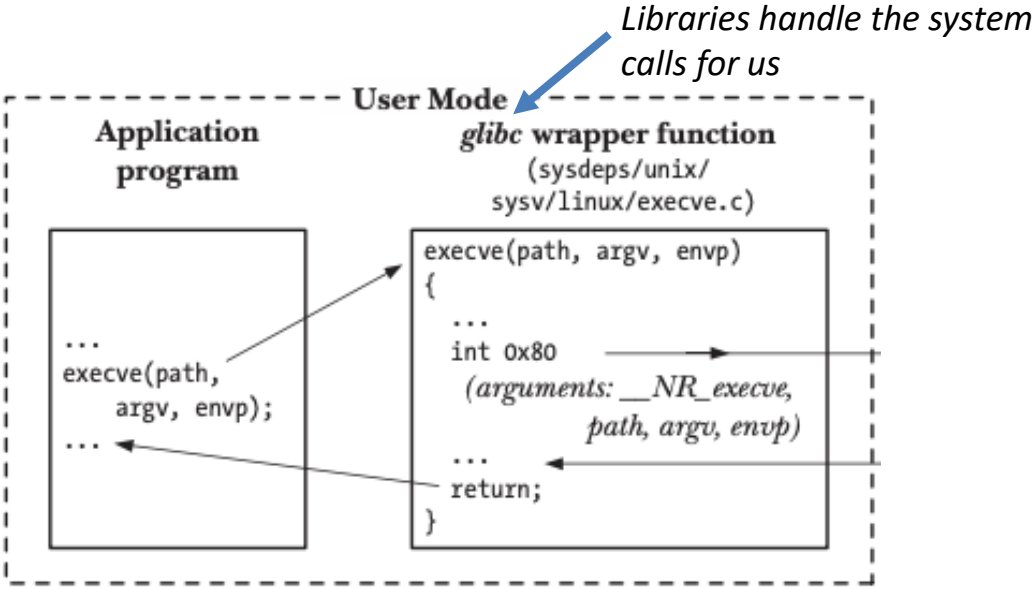
- EAX
- EBX
- ECX
- EDX

System calls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```



EAX	System Call Number
EBX	Address of “/bin/bc”
ECX	0 or 1 Environment variables
EDX	INT 0x80 send trap to kernel and invoke the syscall

The operating system have hundreds of different syscalls, and different syscalls have different parameters, we need a way to distinguish them

The OS will look at the values at certain registers!

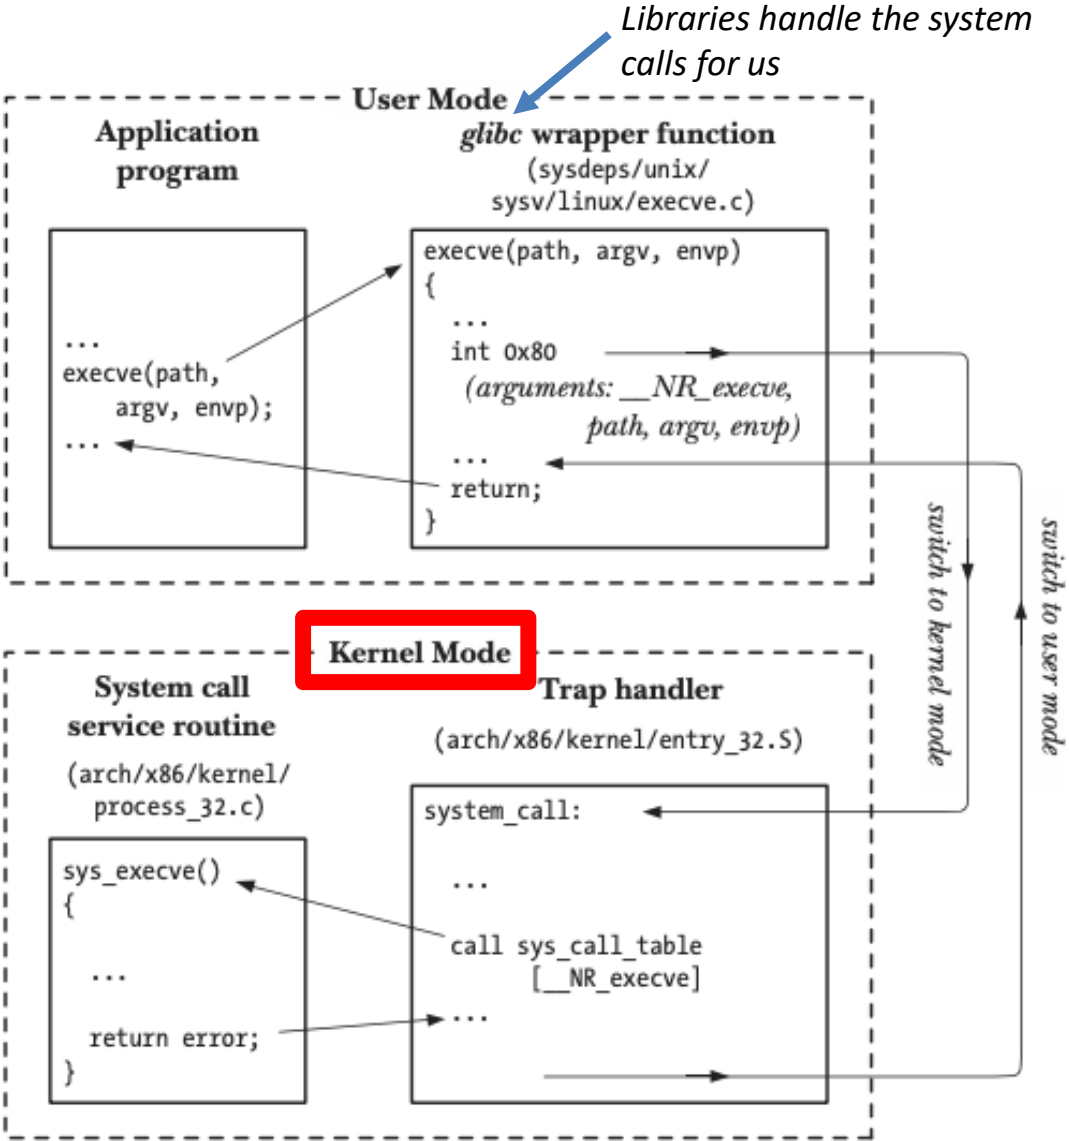
Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

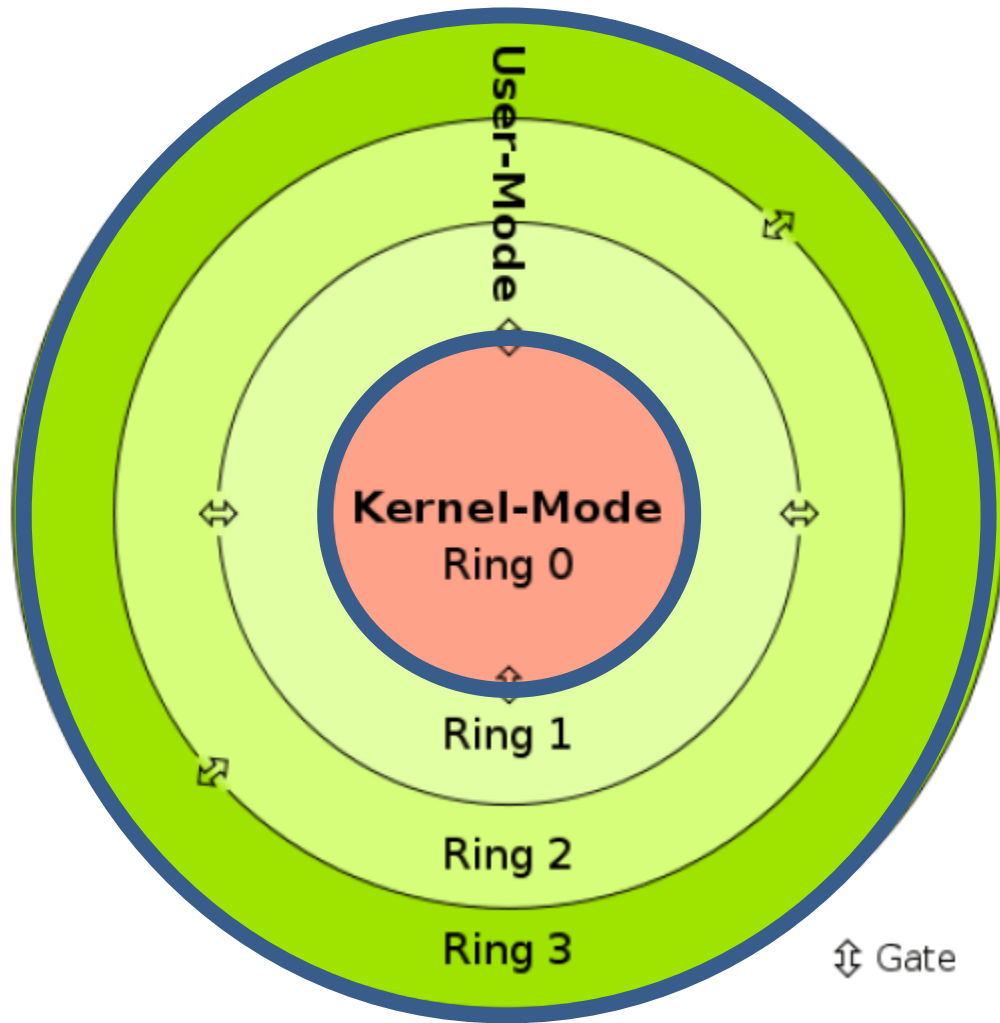
```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

- EAX** System Call Number
- EBX** Address of “/bin/bc”
- ECX** 0 or 1 Environment variables
- EDX** INT 0x80 send trap to kernel and invoke the syscall



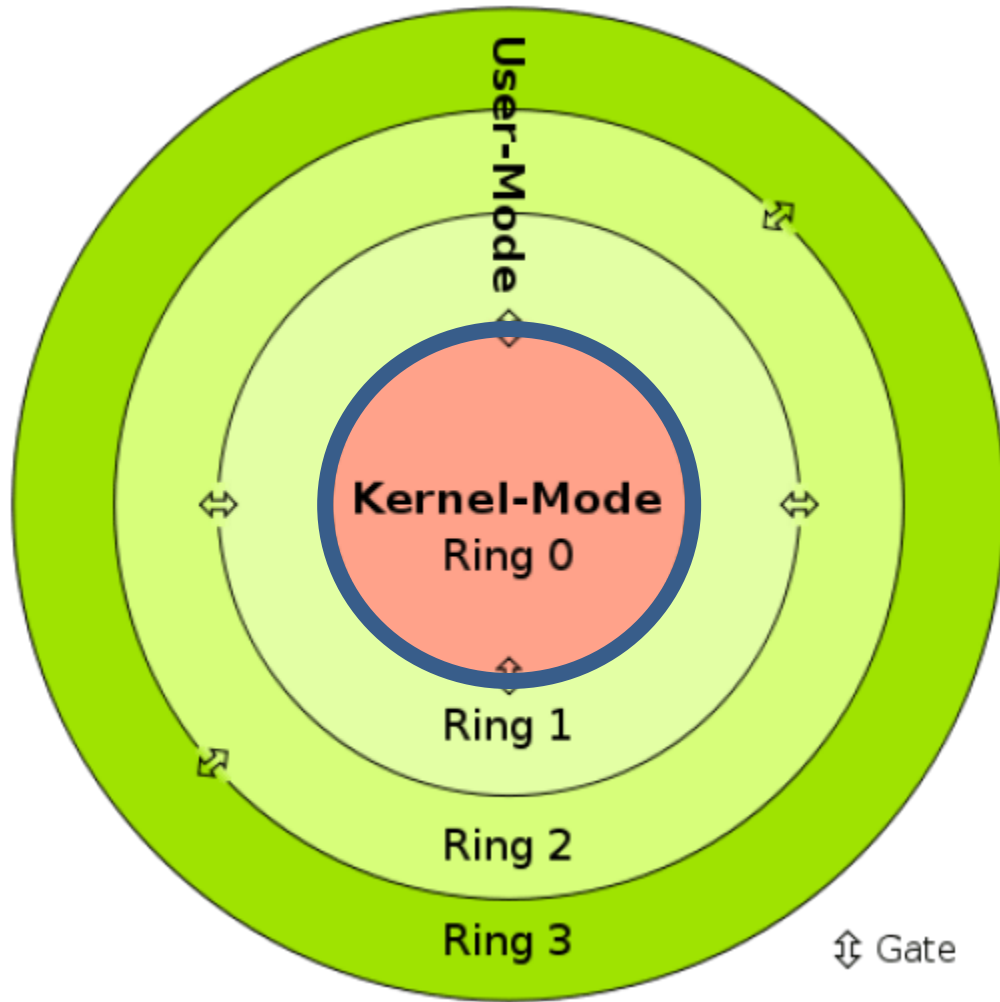
Syscalls



All applications run in user mode.

The code has no ability to directly access hardware
Code running in user mode must use API/syscalls to access hardware and memory

Syscalls



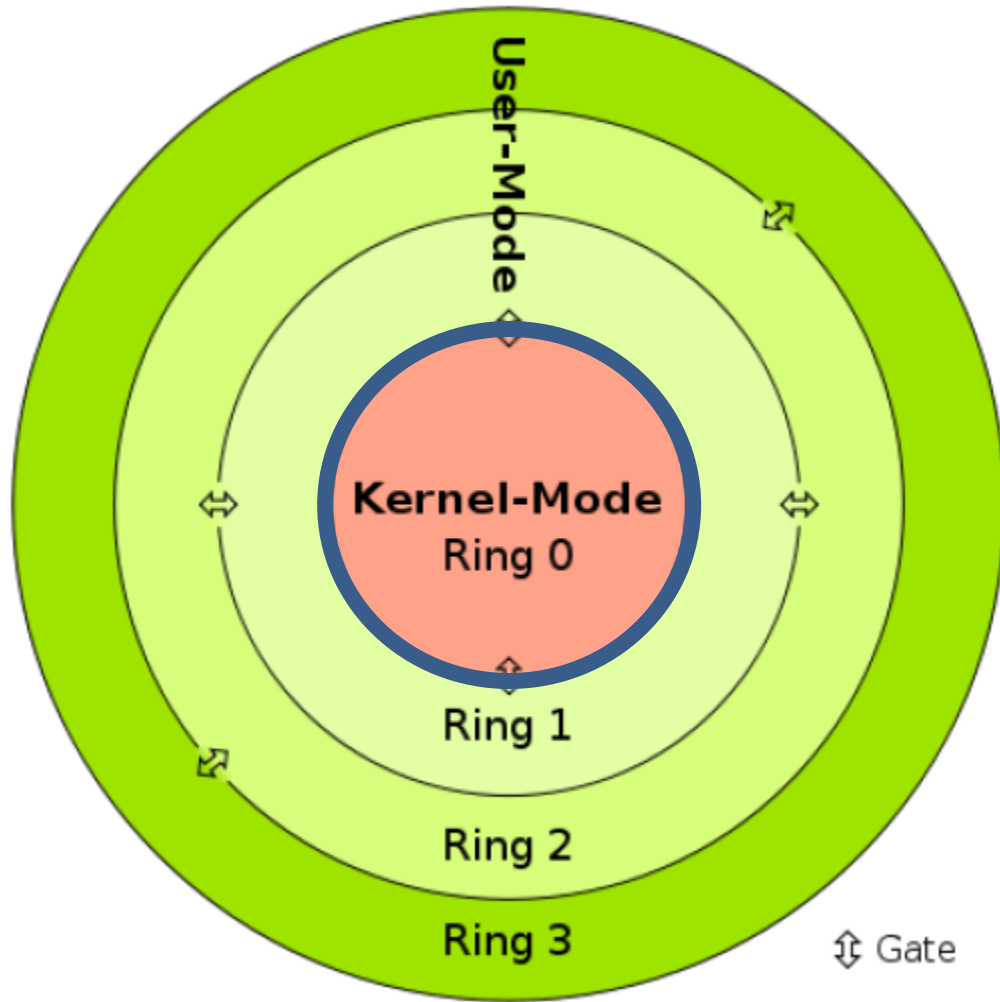
All applications run in user mode.

The code has no ability to directly access hardware
Code running in user mode must use API/syscalls to access hardware and memory

Code running in kernel-mode has complete, unrestricted access to computer resources

Reserved for the lowest-level trusted functions of the operating system

Syscalls



The collective functionality and services of the OS that manages the computer and its resources is called the **kernel**

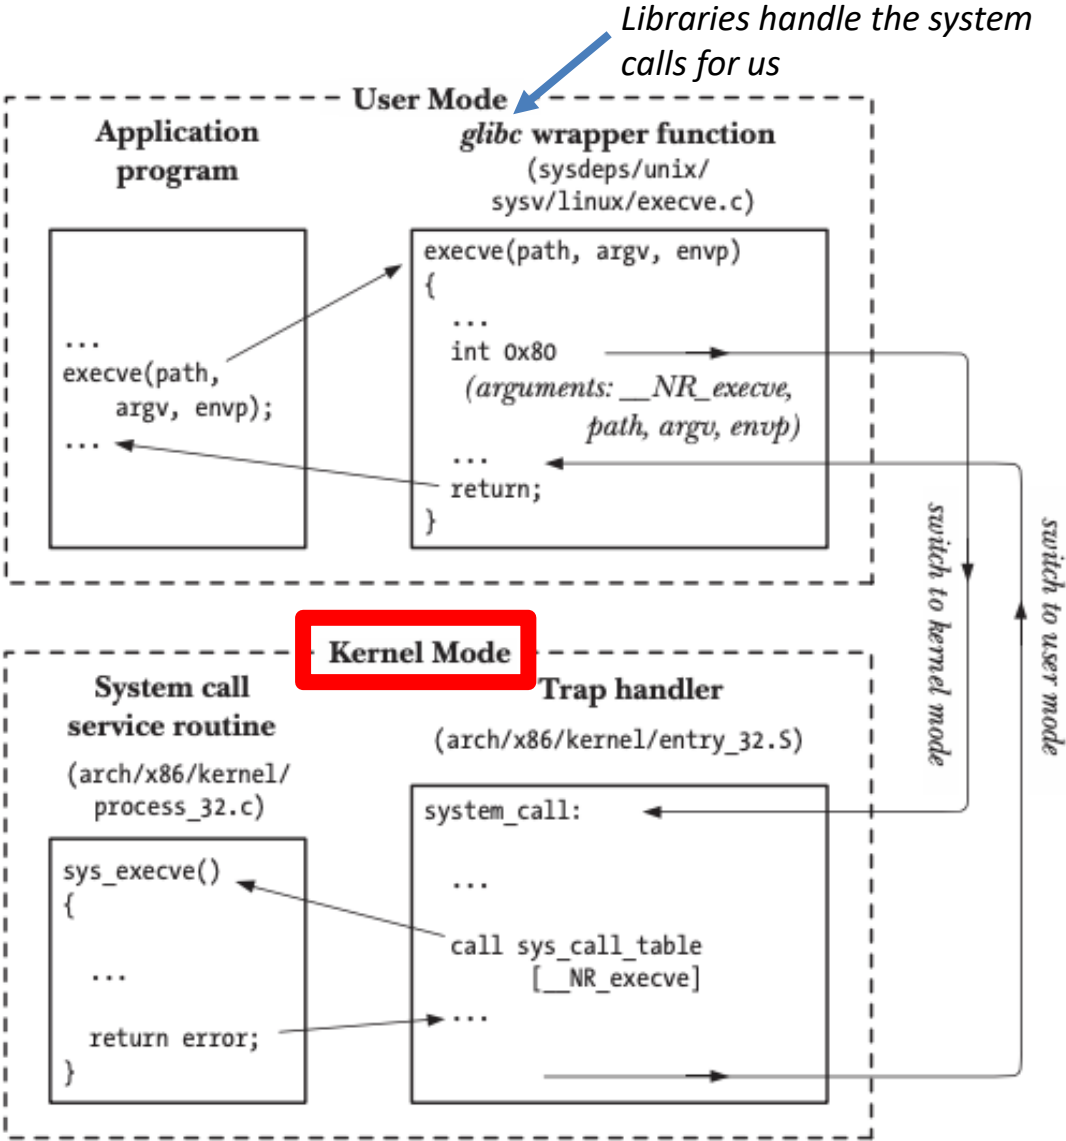
Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

- EAX** System Call Number
- EBX** Address of “/bin/bc”
- ECX** 0 or 1 Environment variables
- EDX** INT 0x80 send trap to kernel and invoke the syscall



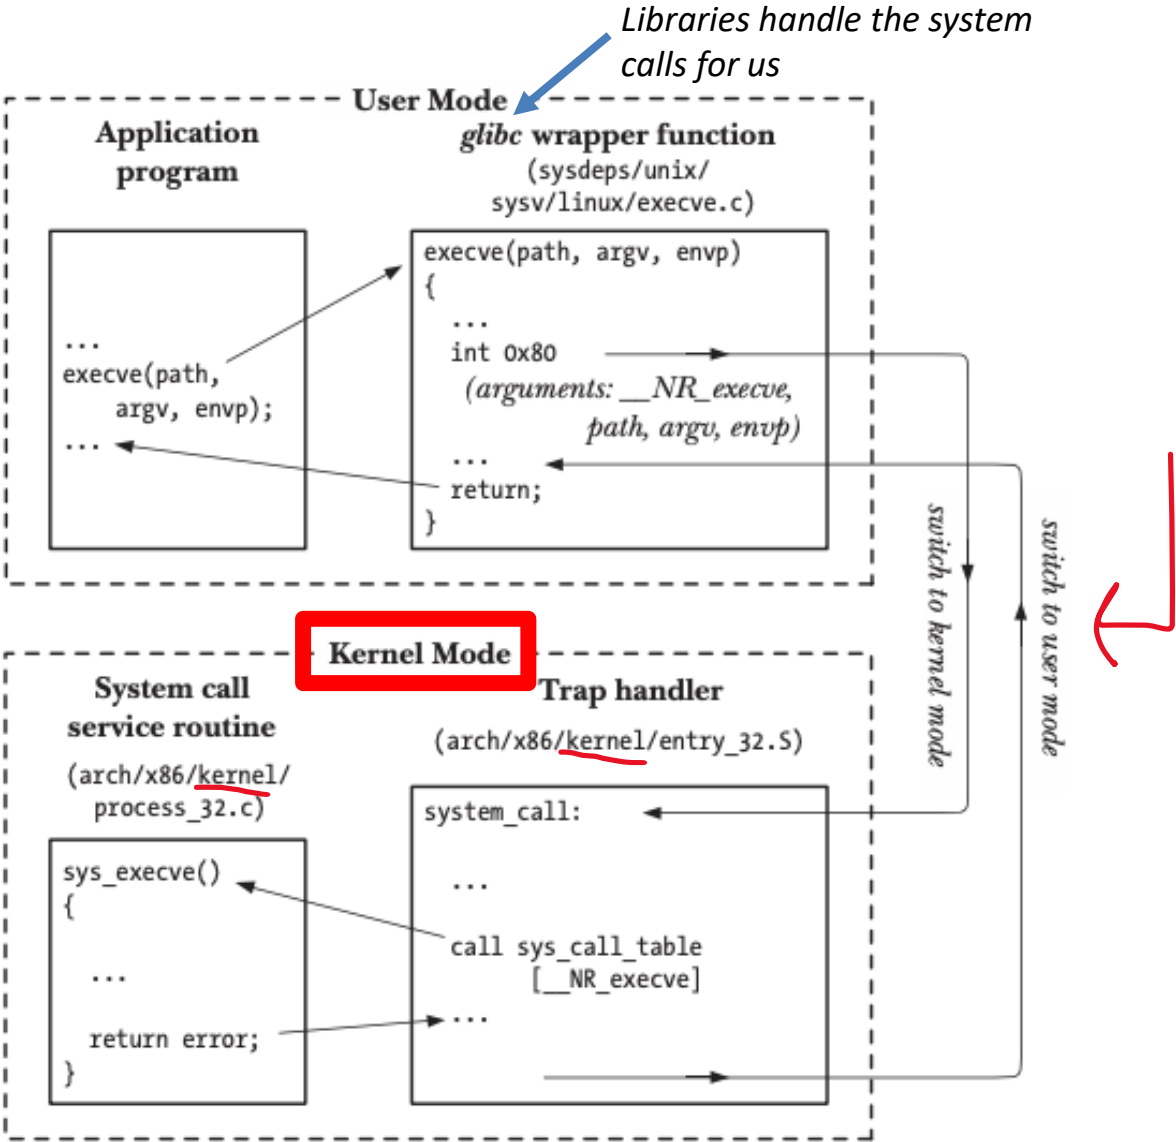
Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

- EAX** System Call Number
- EBX** Address of “/bin/bc”
- ECX** 0 or 1 Environment variables
- EDX** INT 0x80 send trap to kernel and invoke the syscall



Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

NR	syscall name	references	%eax	arg0 (%ebx)	arg1 (%ecx)	arg2 (%edx)	arg3 (%esi)	arg4 (%edi)	arg5 (%ebp)
0	restart_syscall	man/ cs/	0x00	-	-	-	-	-	-
1	exit	man/ cs/	0x01	int error_code	-	-	-	-	-
2	fork	man/ cs/	0x02	-	-	-	-	-	-
3	read	man/ cs/	0x03	unsigned int fd	char *buf	size_t count	-	-	-
4	write	man/ cs/	0x04	unsigned int fd	const char *buf	size_t count	-	-	-
5	open	man/ cs/	0x05	const char *filename	int flags	umode_t mode	-	-	-
6	close	man/ cs/	0x06	unsigned int fd	-	-	-	-	-
7	waitpid	man/ cs/	0x07	pid_t pid	int *stat_addr	int options	-	-	-
8	creat	man/ cs/	0x08	const char *pathname	umode_t mode	-	-	-	-
9	link	man/ cs/	0x09	const char *oldname	const char *newname	-	-	-	-
10	unlink	man/ cs/	0x0a	const char *pathname	-	-	-	-	-
11	execve	man/ cs/	0x0b	const char *filename	const char *const *argv	const char *const *envp	-	-	-
12	chdir	man/ cs/	0x0c	const char *filename	-	-	-	-	-

EDX

INT 0x80

send trap to kernel and
invoke the syscall

Systemcalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

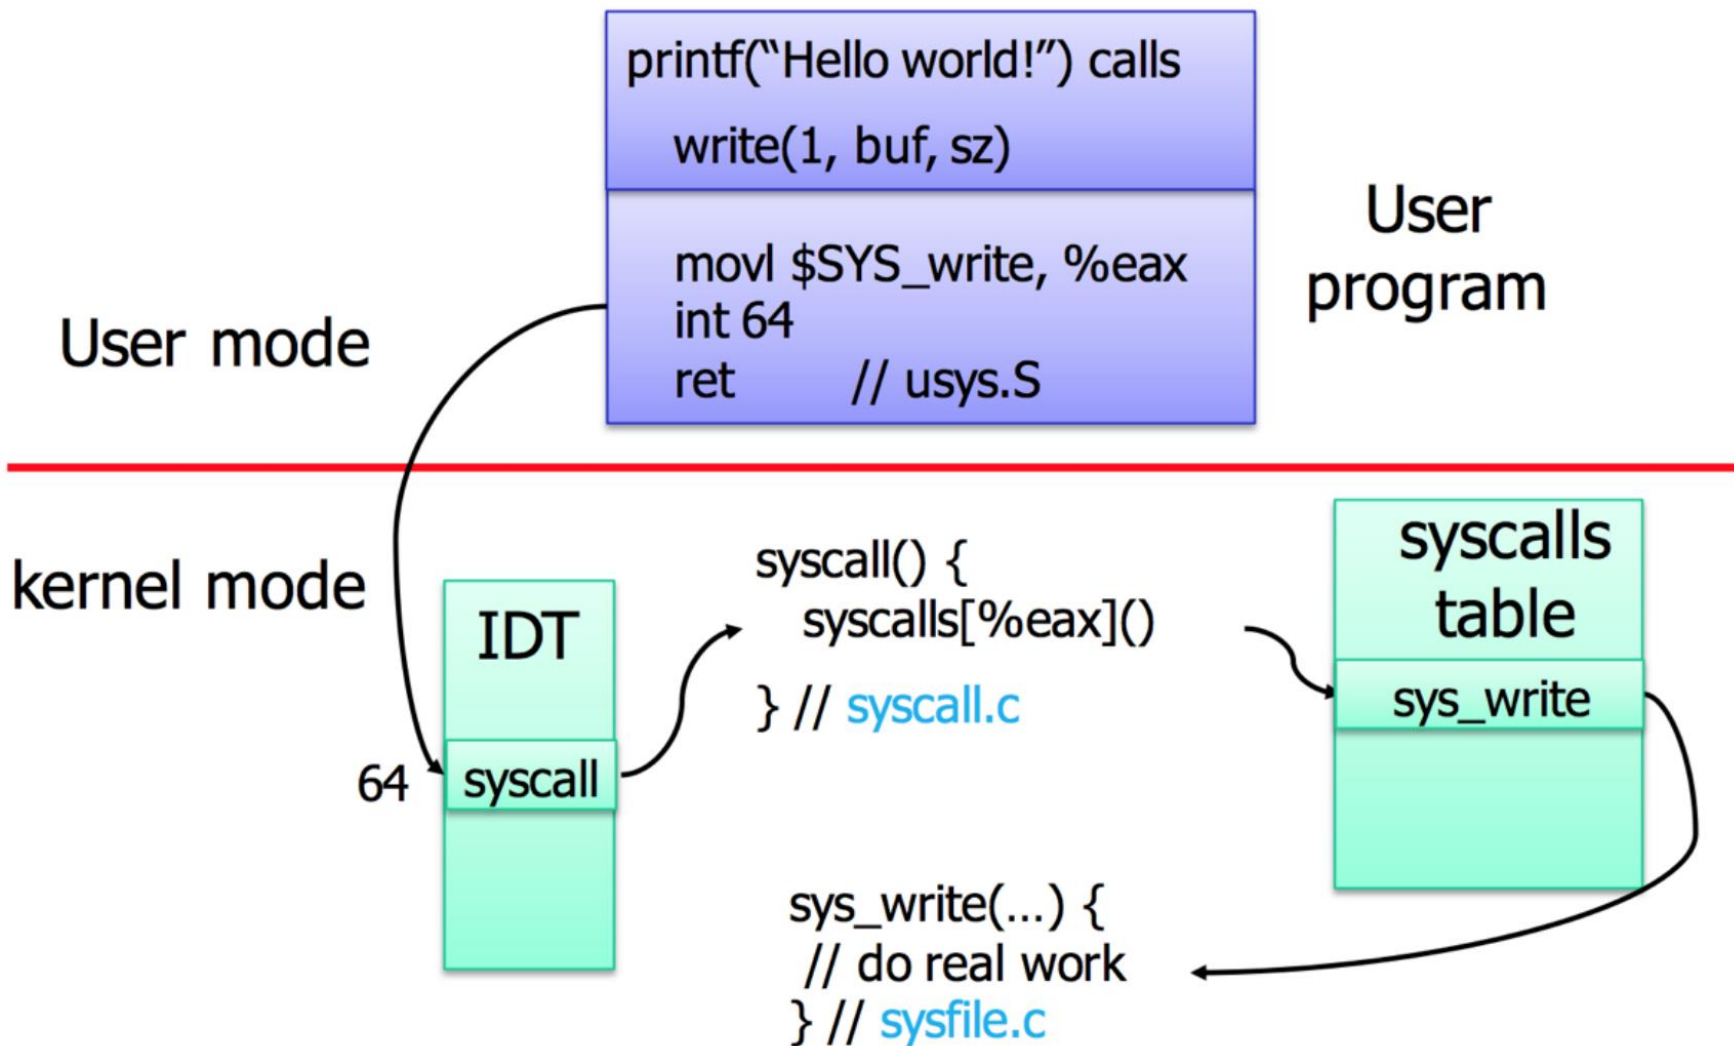
NR	syscall name	references	%eax	arg0 (%ebx)	arg1 (%ecx)	arg2 (%edx)	arg3 (%esi)	arg4 (%edi)	arg5 (%ebp)
0	restart_syscall	man/ cs/	0x00	-	-	-	-	-	-
1	exit	man/ cs/	0x01	int error_code	-	-	-	-	-
2	fork	man/ cs/	0x02	-	-	-	-	-	-
3	read	man/ cs/	0x03	unsigned int fd	char *buf	size_t count	-	-	-
4	https://chromium.googlesource.com/chromiumos/docs/+/master/constants/syscalls.md#x86-32_bit								
5									
6									
7									
8									
9									
10	unlink	man/ cs/	0x0a	const char *pathname	-	-	-	-	-
11	execve	man/ cs/	0x0b	const char *filename	const char *const *argv	const char *const *envp	-	-	-
12	chdir	man/ cs/	0x0c	const char *filename	-	-	-	-	-

EDX

INT 0x80

send trap to kernel and invoke the syscall

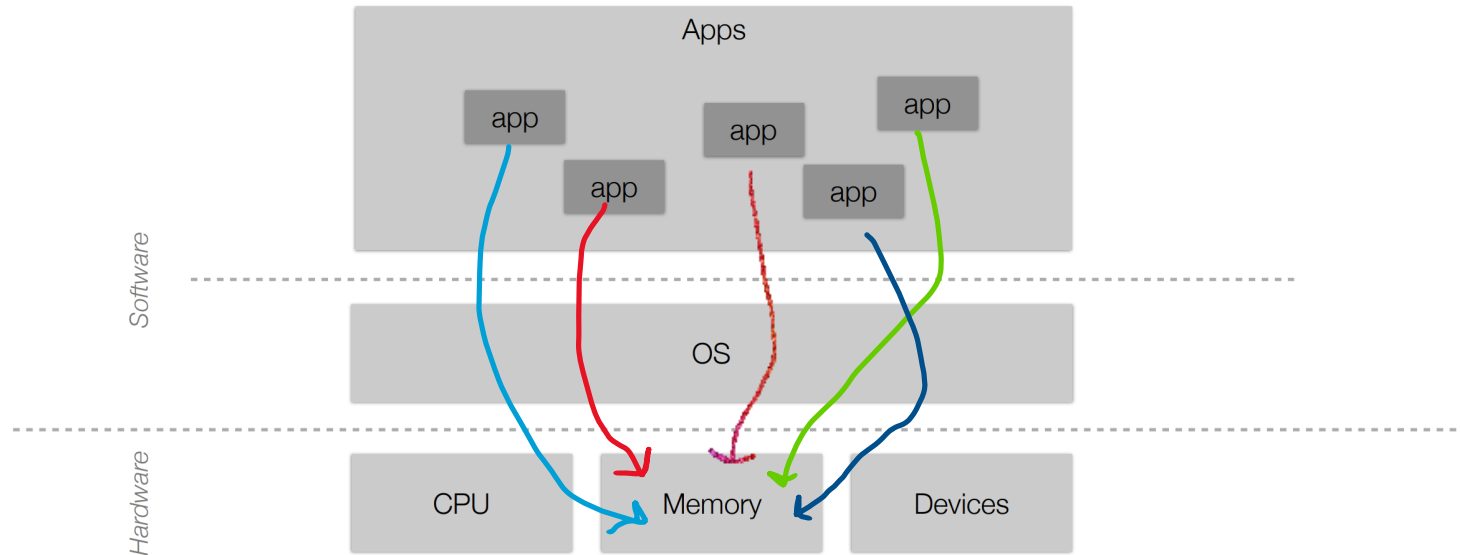
Syscalls



Applications Layout in Memory

Process Manager

- Manages how processes are structured and how to handle many processes running at once

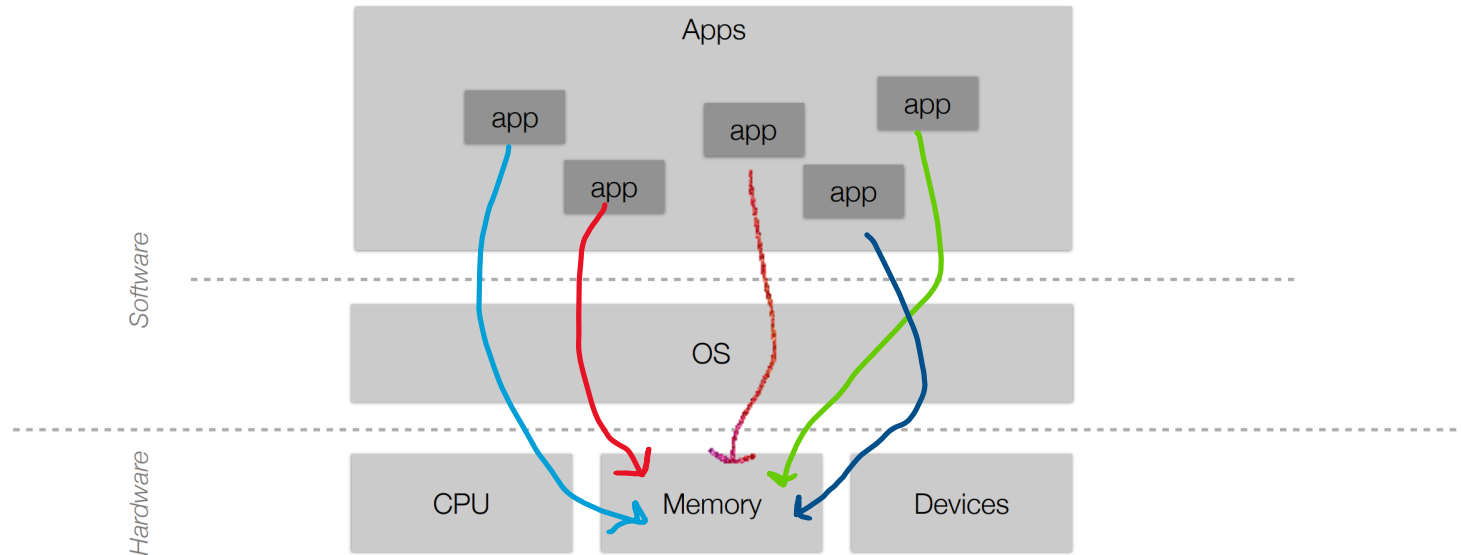


How does a **program** get loaded into memory?

Applications Layout in Memory

Process Manager

- Manages how processes are structured and how to handle many processes running at once



How does a **program** get loaded into memory?

An active program running on a computer is called a **process**

Applications Layout in Memory

What does this look like?

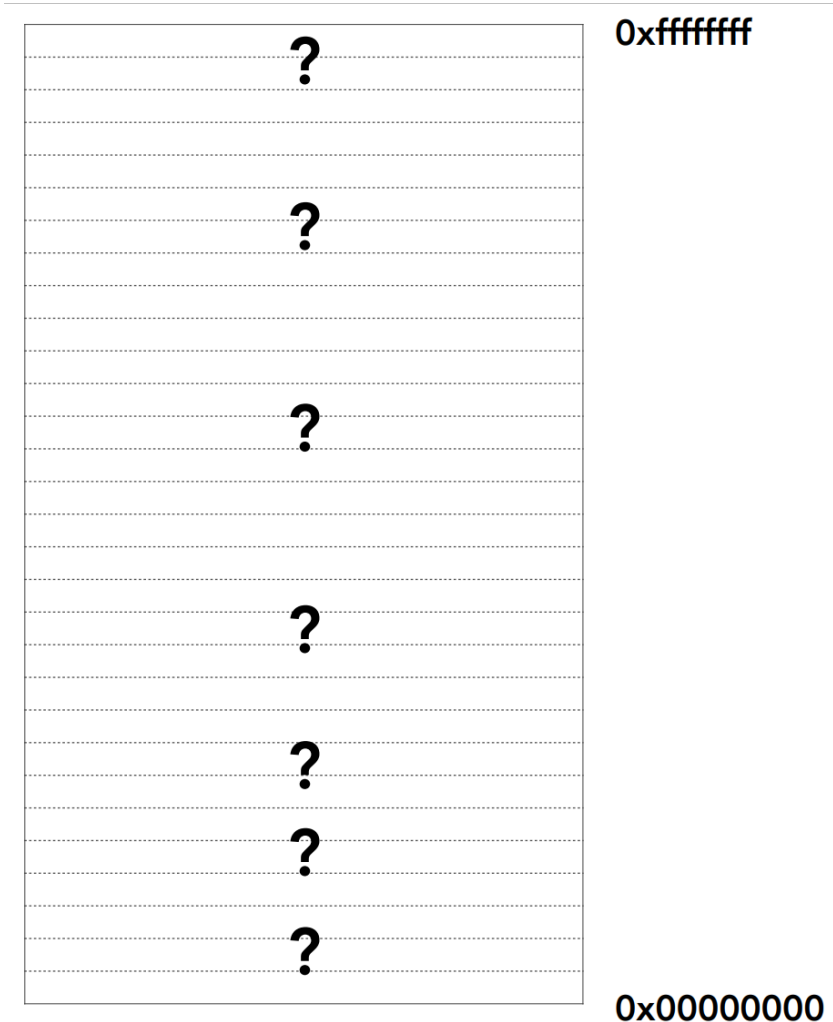
1. Executable Code

2. Associated Data

3. Execution Context/Bookkeeping information

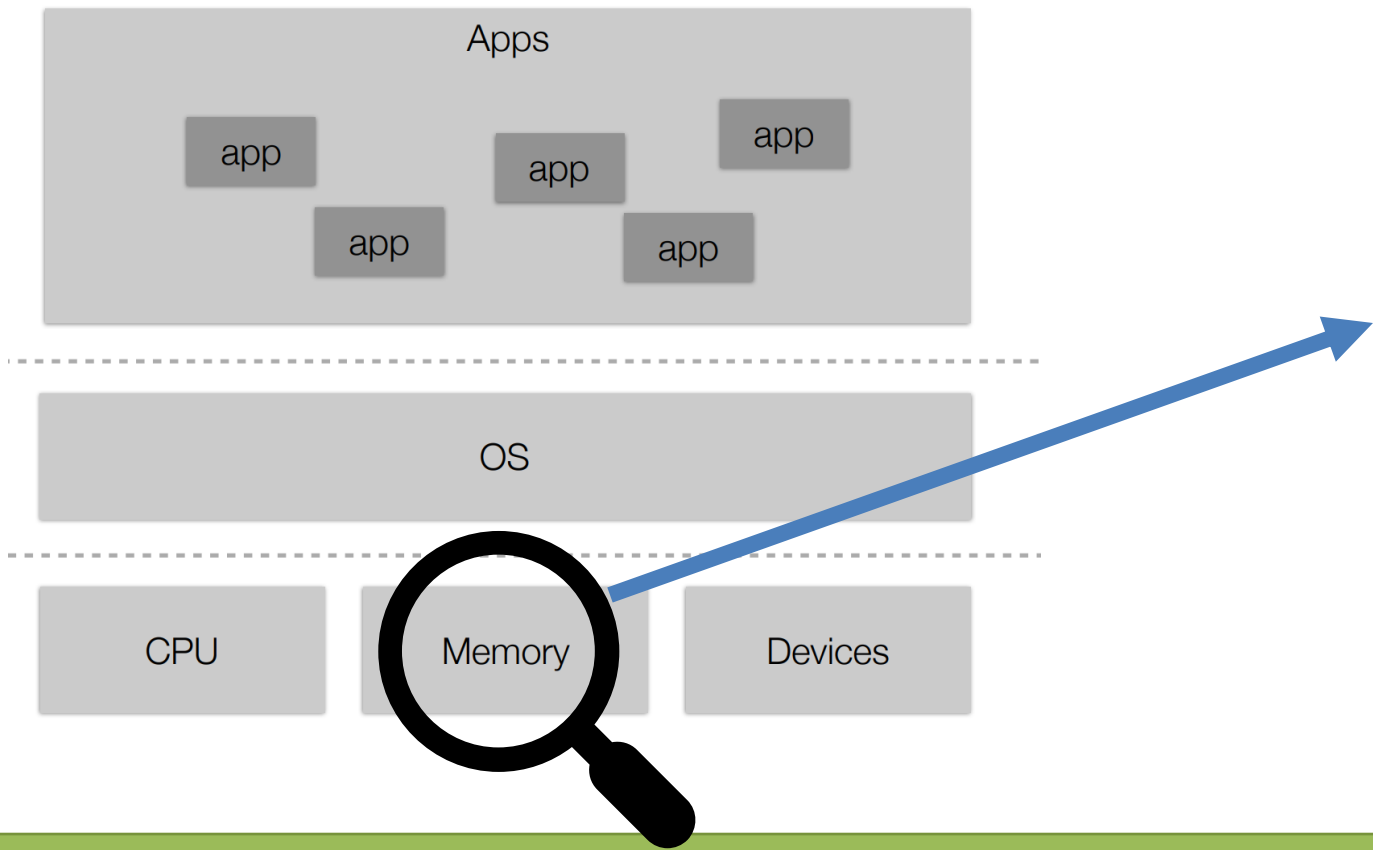
Process Manager

- Manages how processes are structured and how to handle many processes running at once



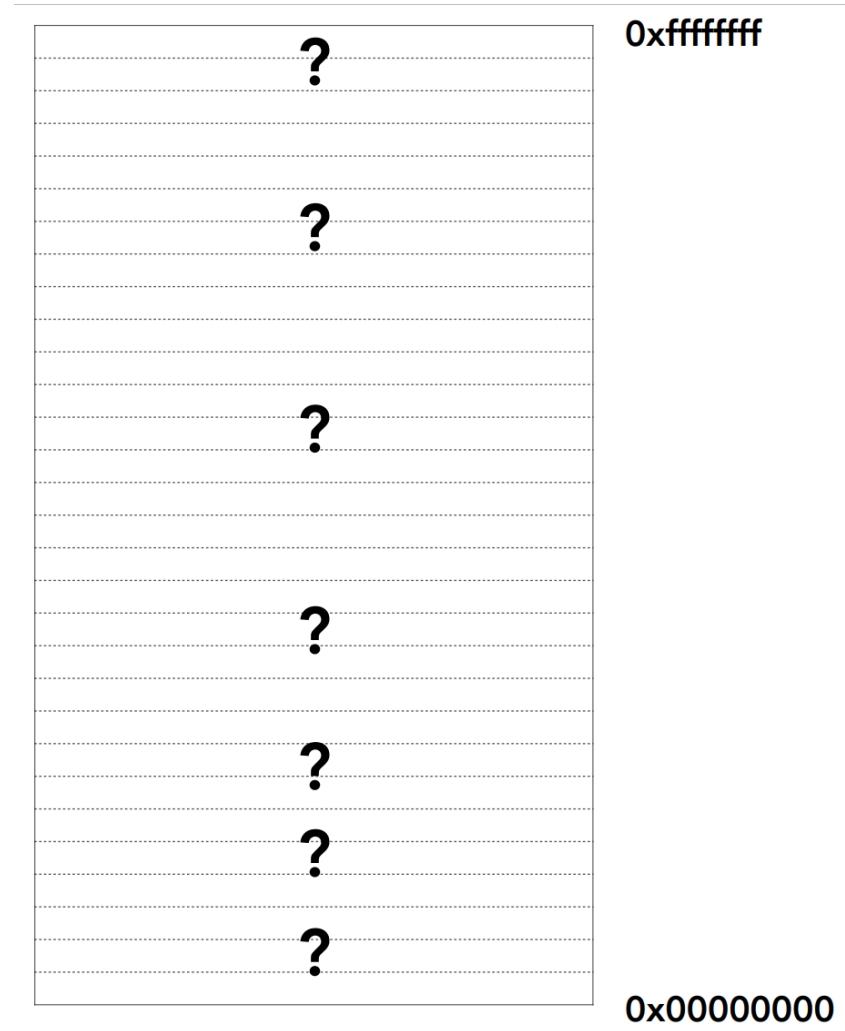
Applications Layout in Memory

What does a program look like in memory?



Process Manager

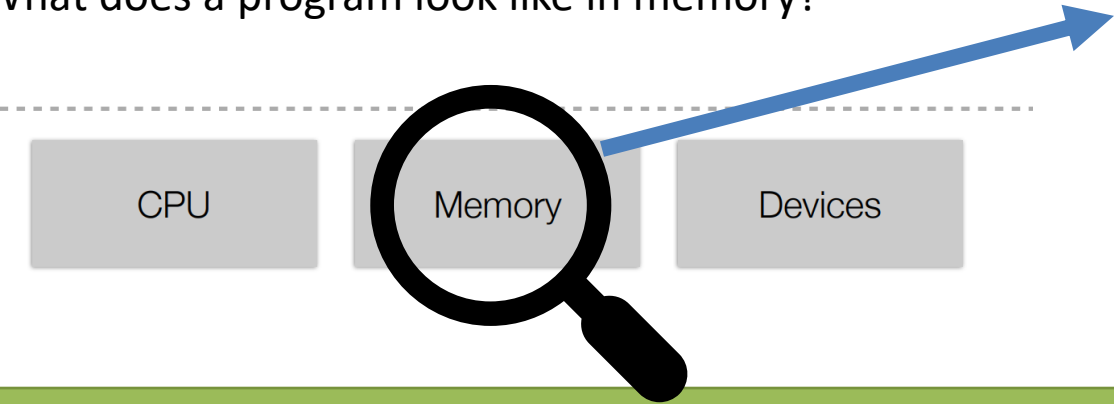
- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

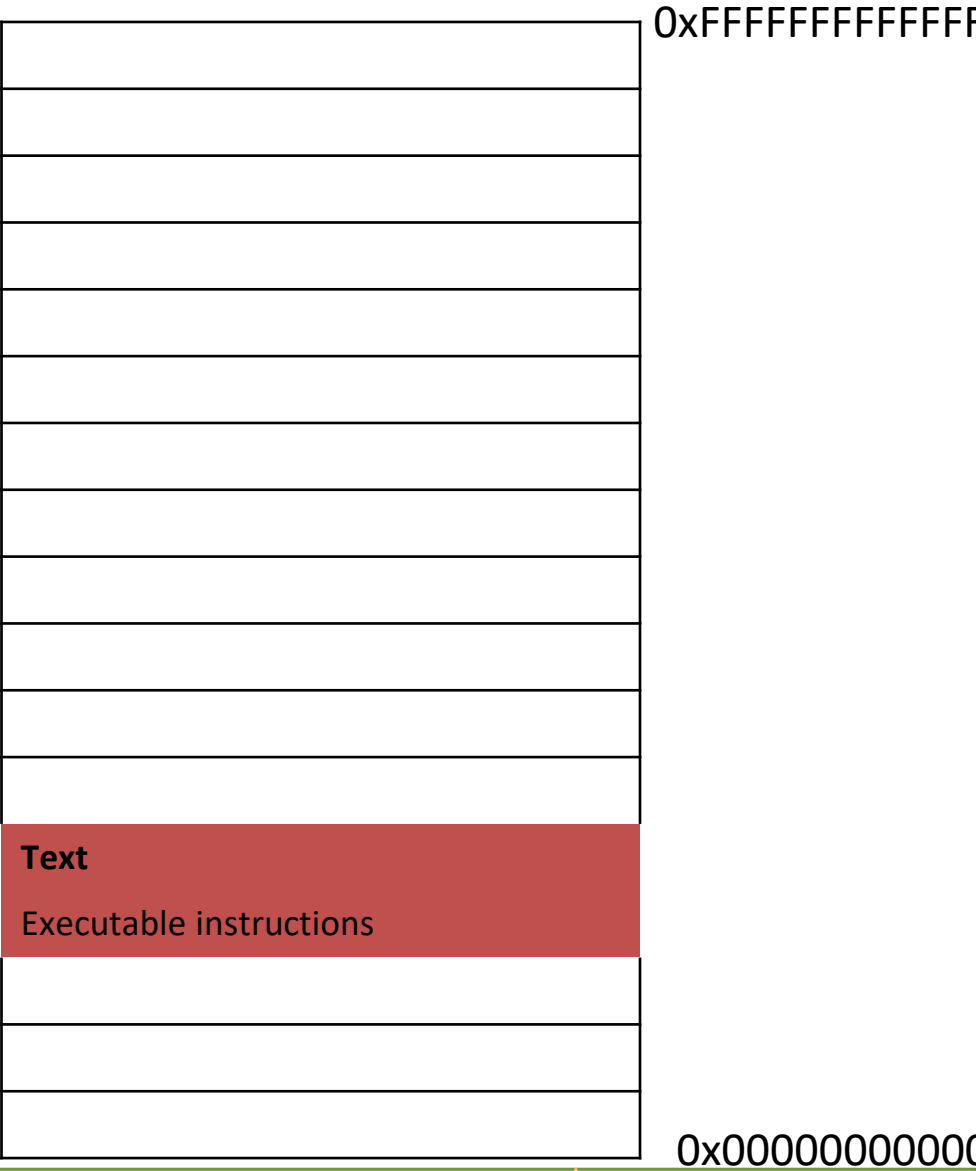
Text Segment- binary executable instructions for the process

What does a program look like in memory?



Process Manager

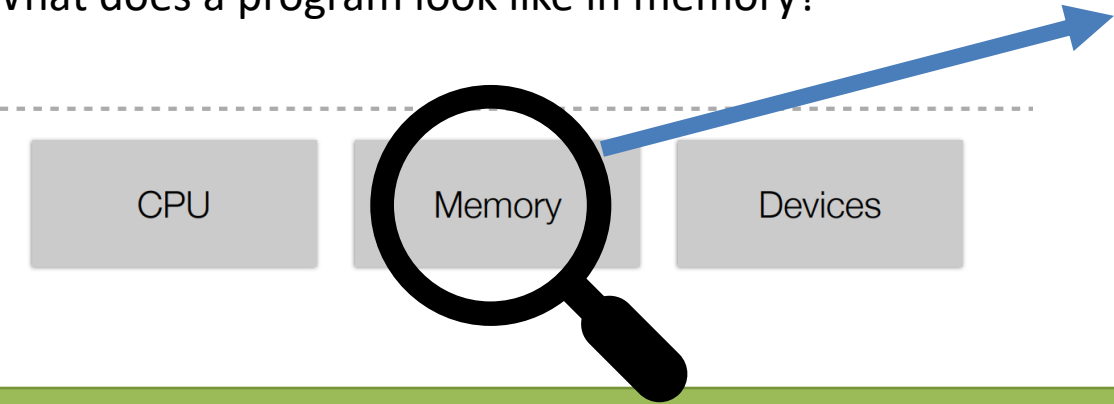
- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

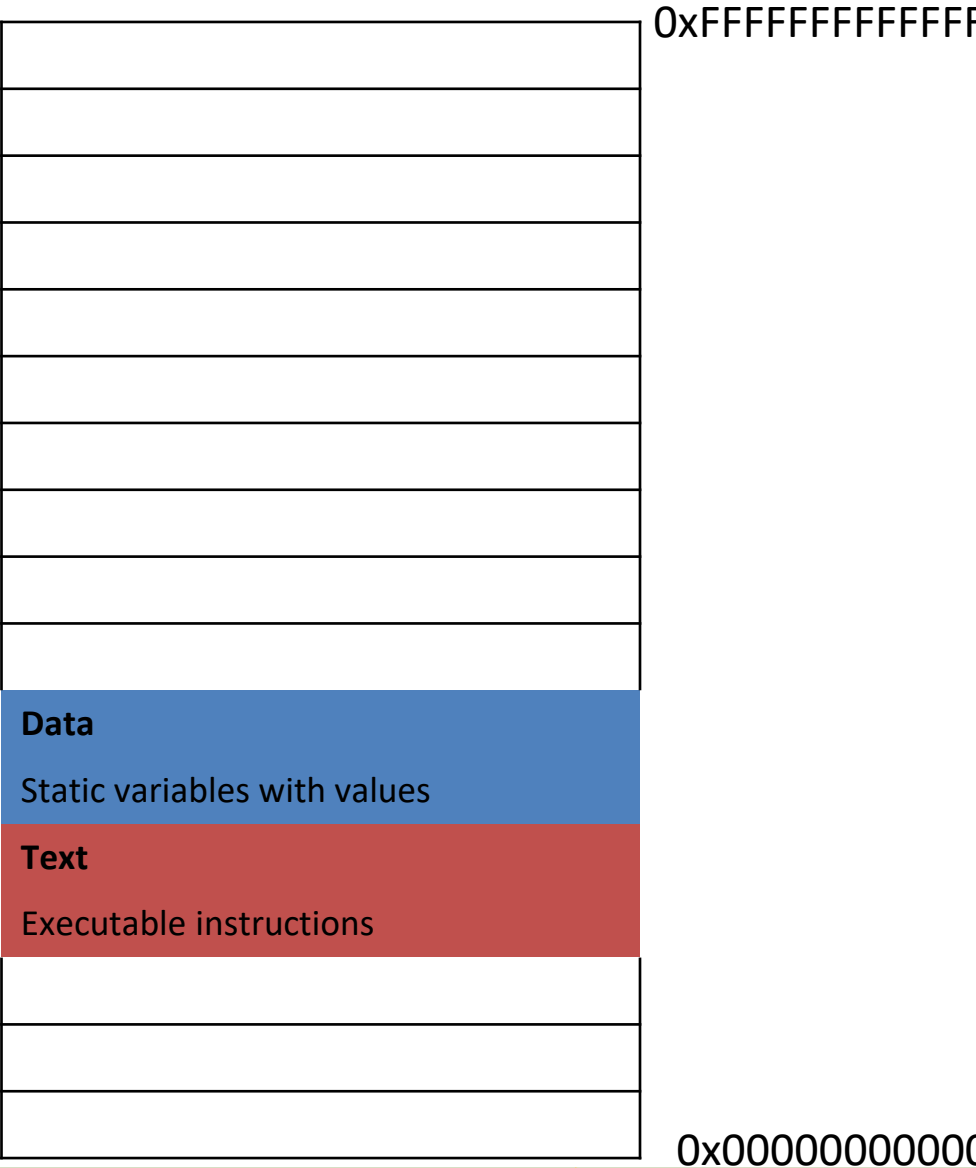
Data Segment- Static variables initialized by the programmer

What does a program look like in memory?



Process Manager

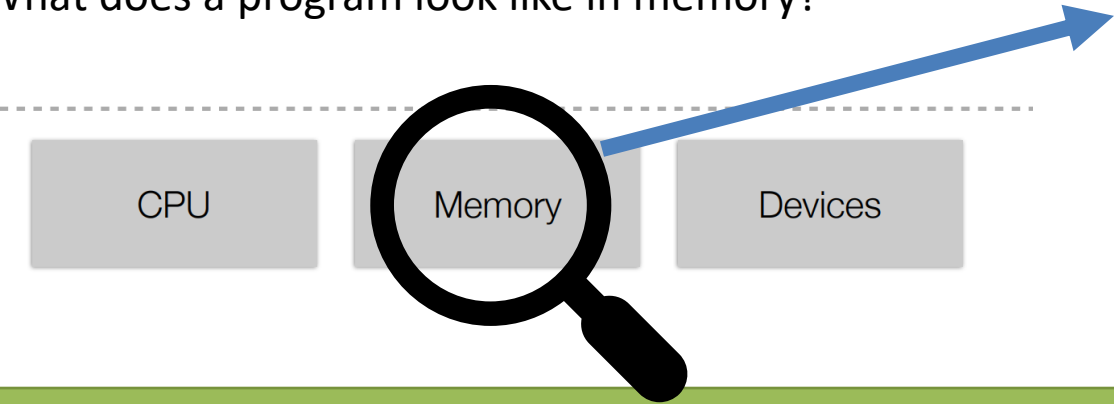
- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

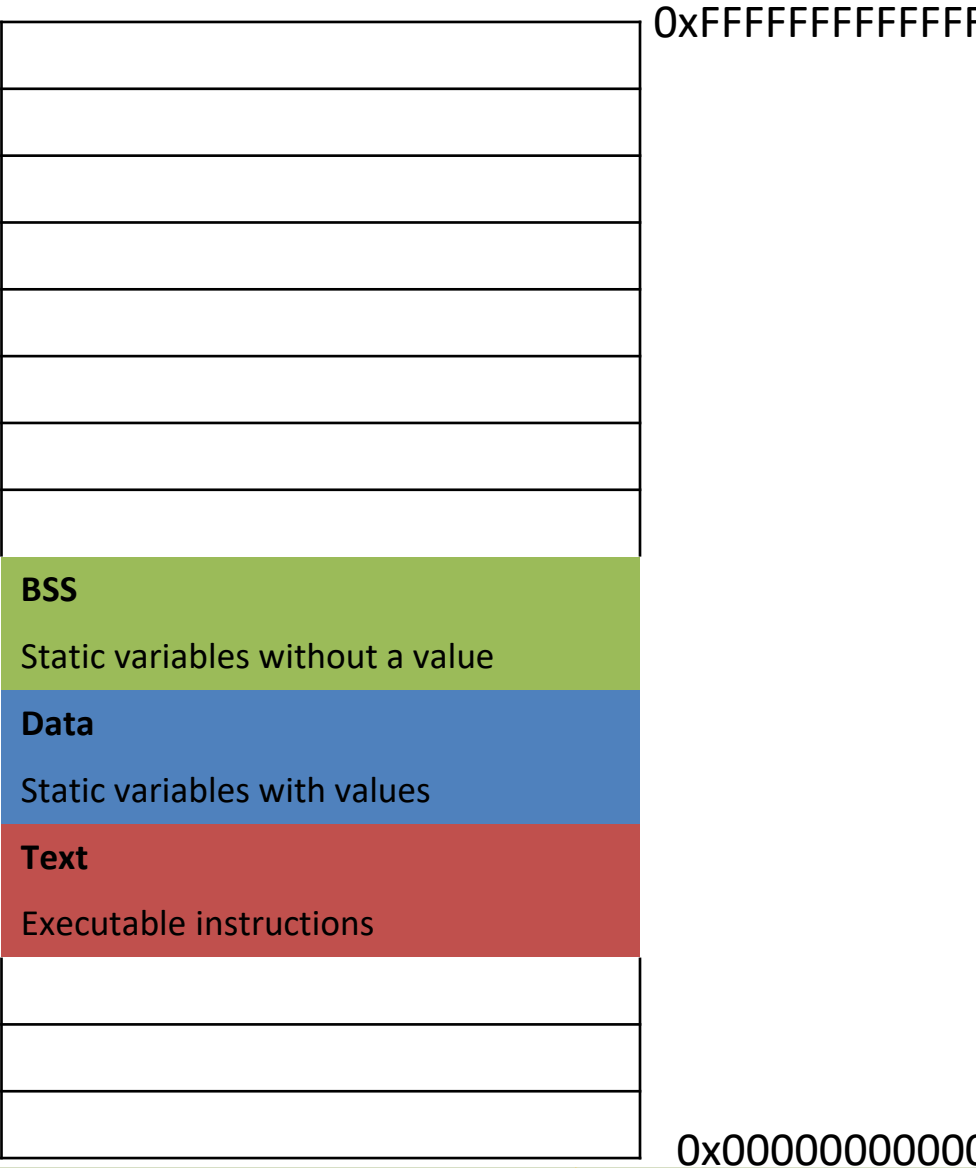
BSS Segment- contains statically allocated variables that are declared, but have not been assigned a value yet

What does a program look like in memory?



Process Manager

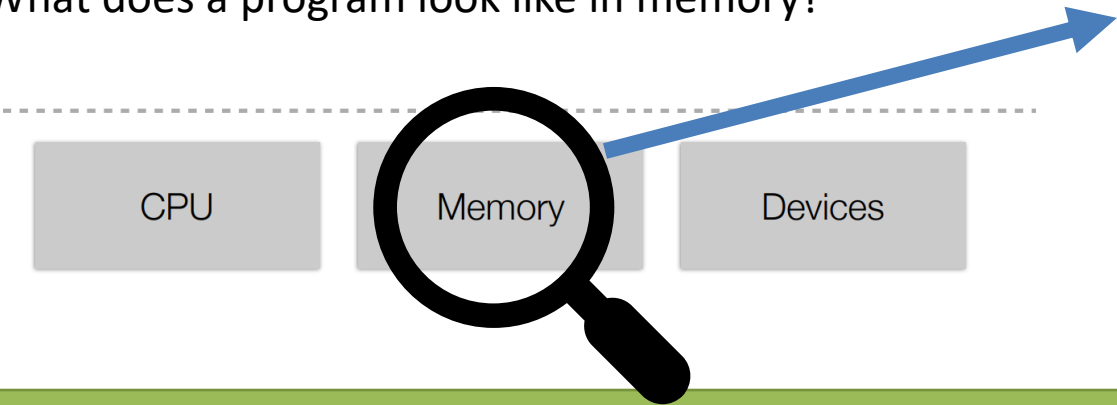
- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

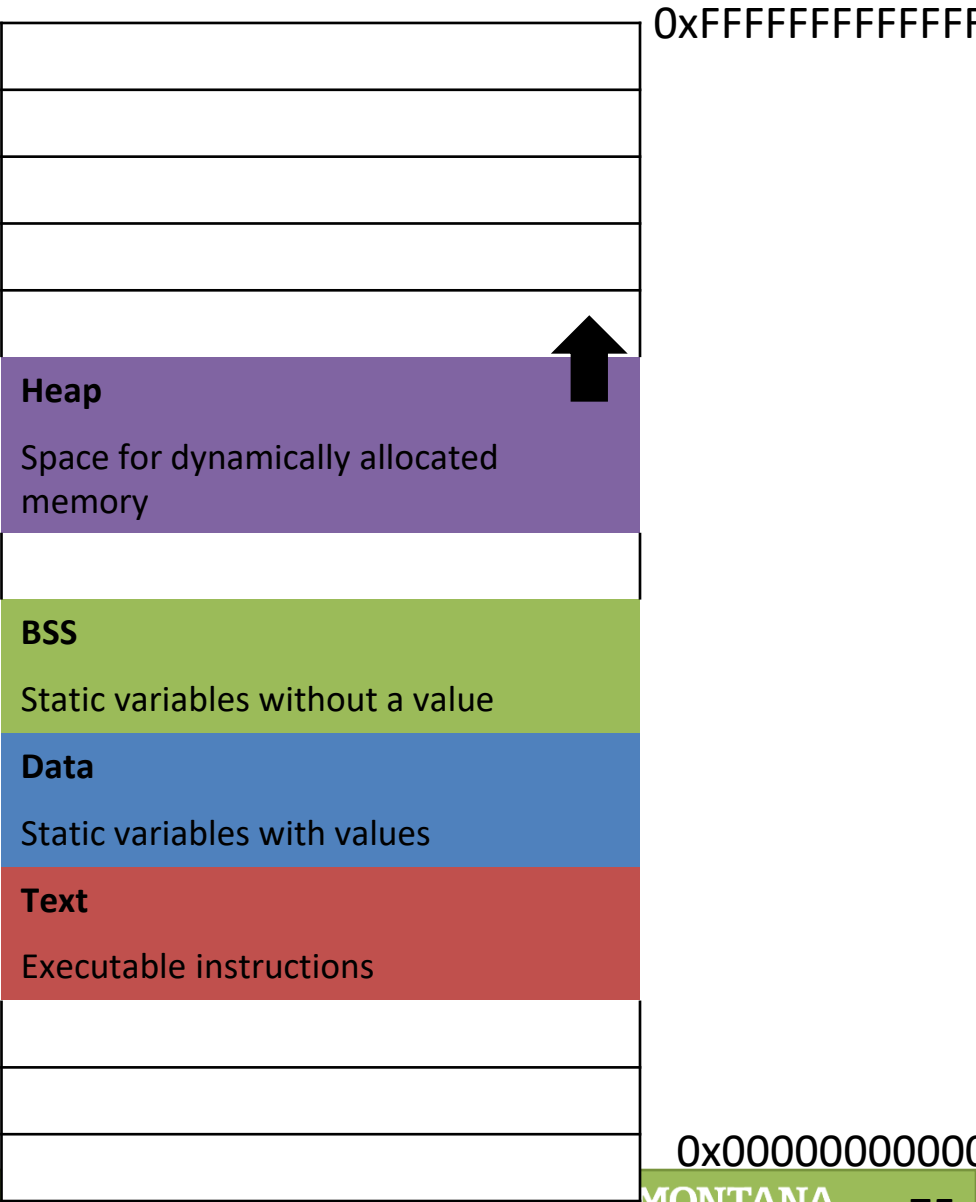
Heap- memory set aside for dynamic allocation (e.g. malloc). Grows “up” as more memory is allocated

What does a program look like in memory?



Process Manager

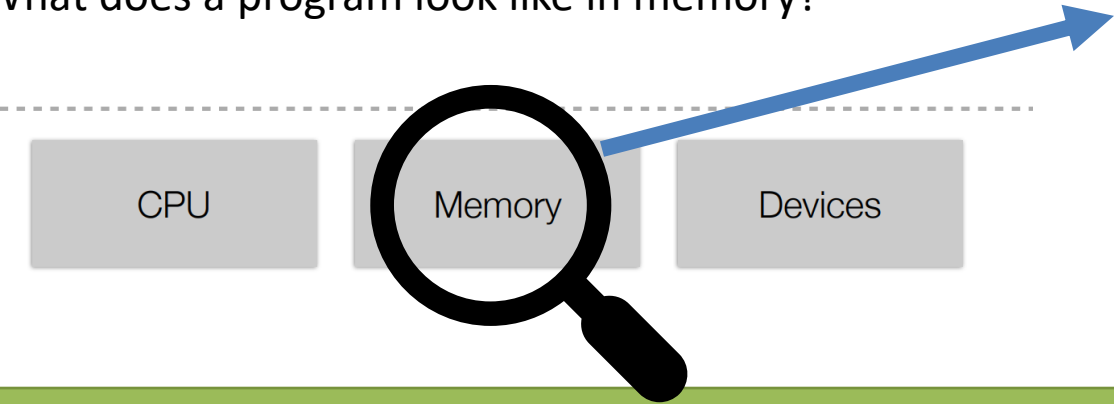
- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

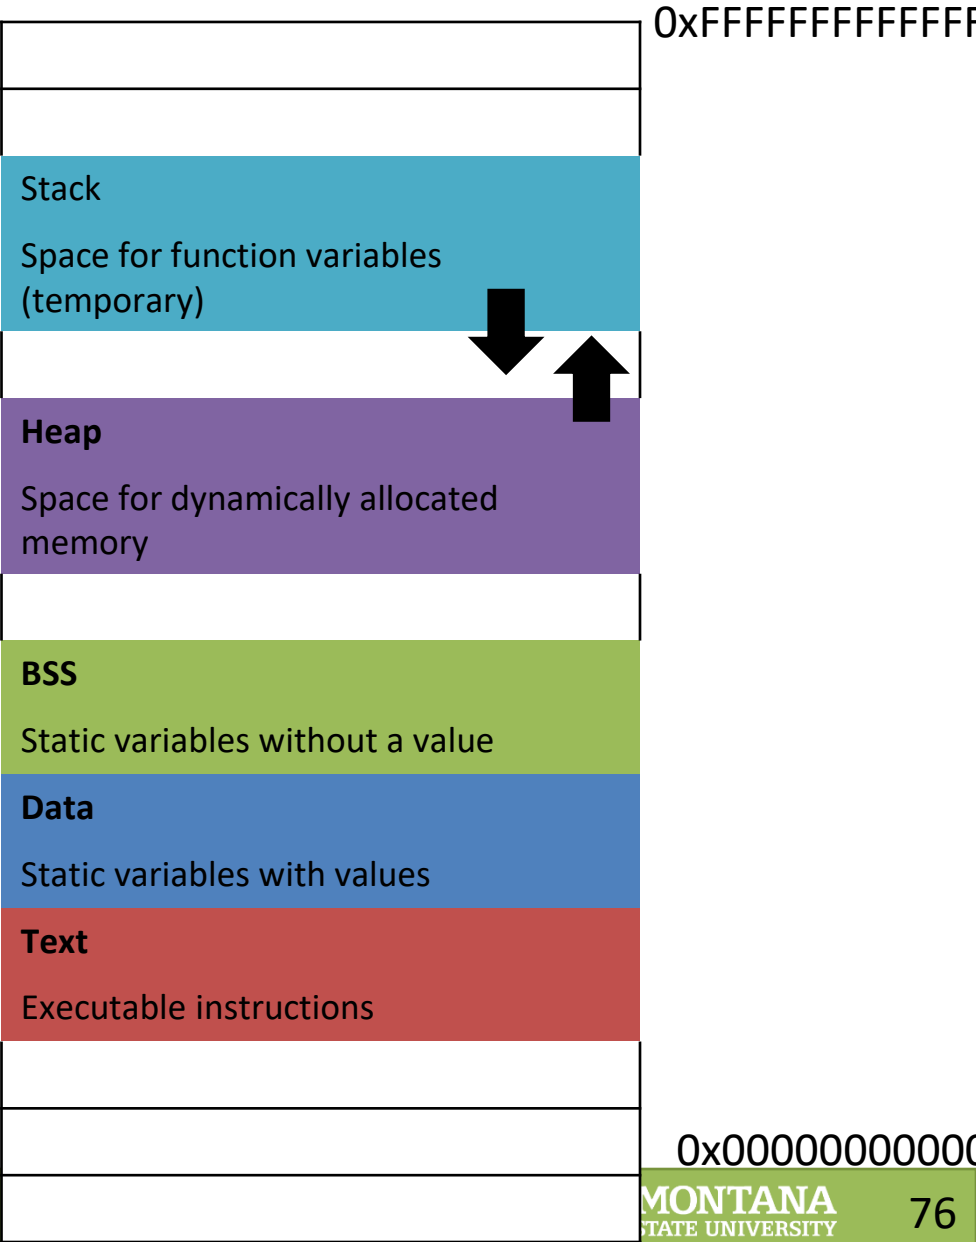
Stack – memory for storing function variables.
Grows “down” as additional functions are called

What does a program look like in memory?



Process Manager

- Manages how processes are structured and how to handle many processes running at once

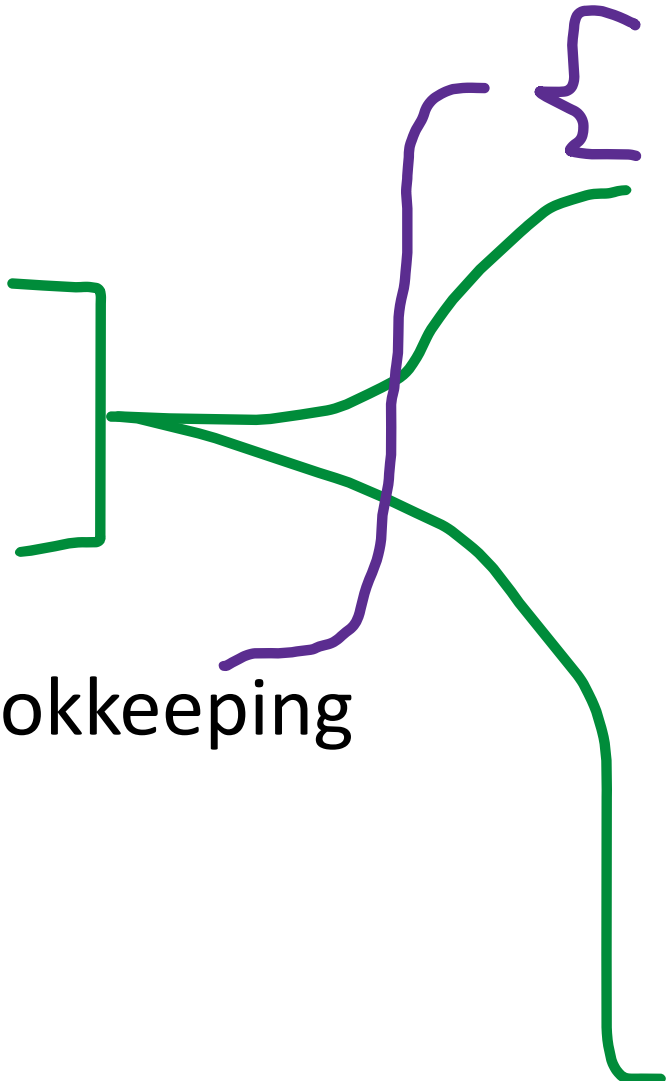


Applications Layout in Memory

1. Executable Code

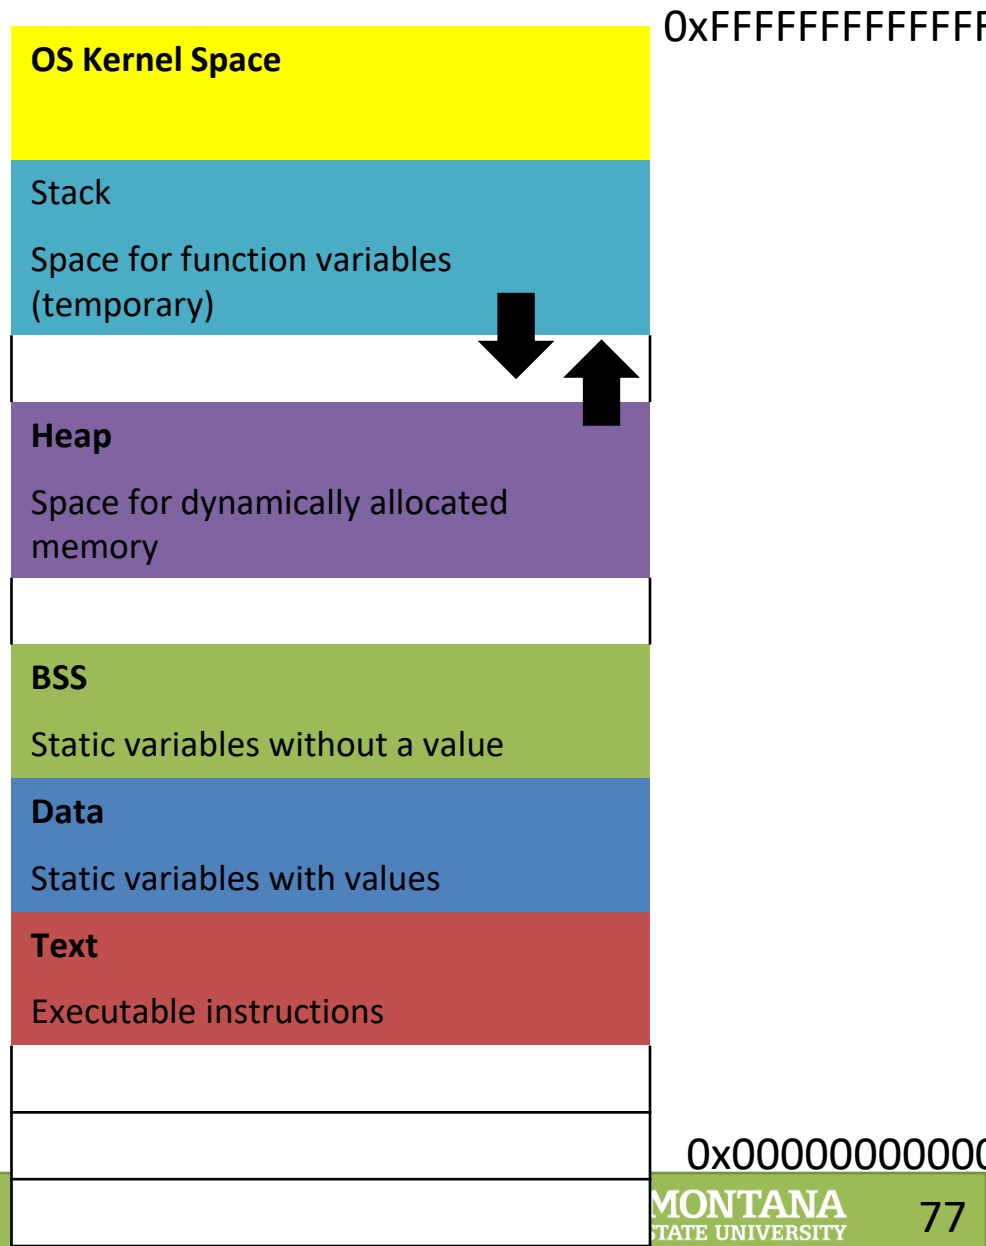
2. Associated Data

3. Execution Context/Bookkeeping information



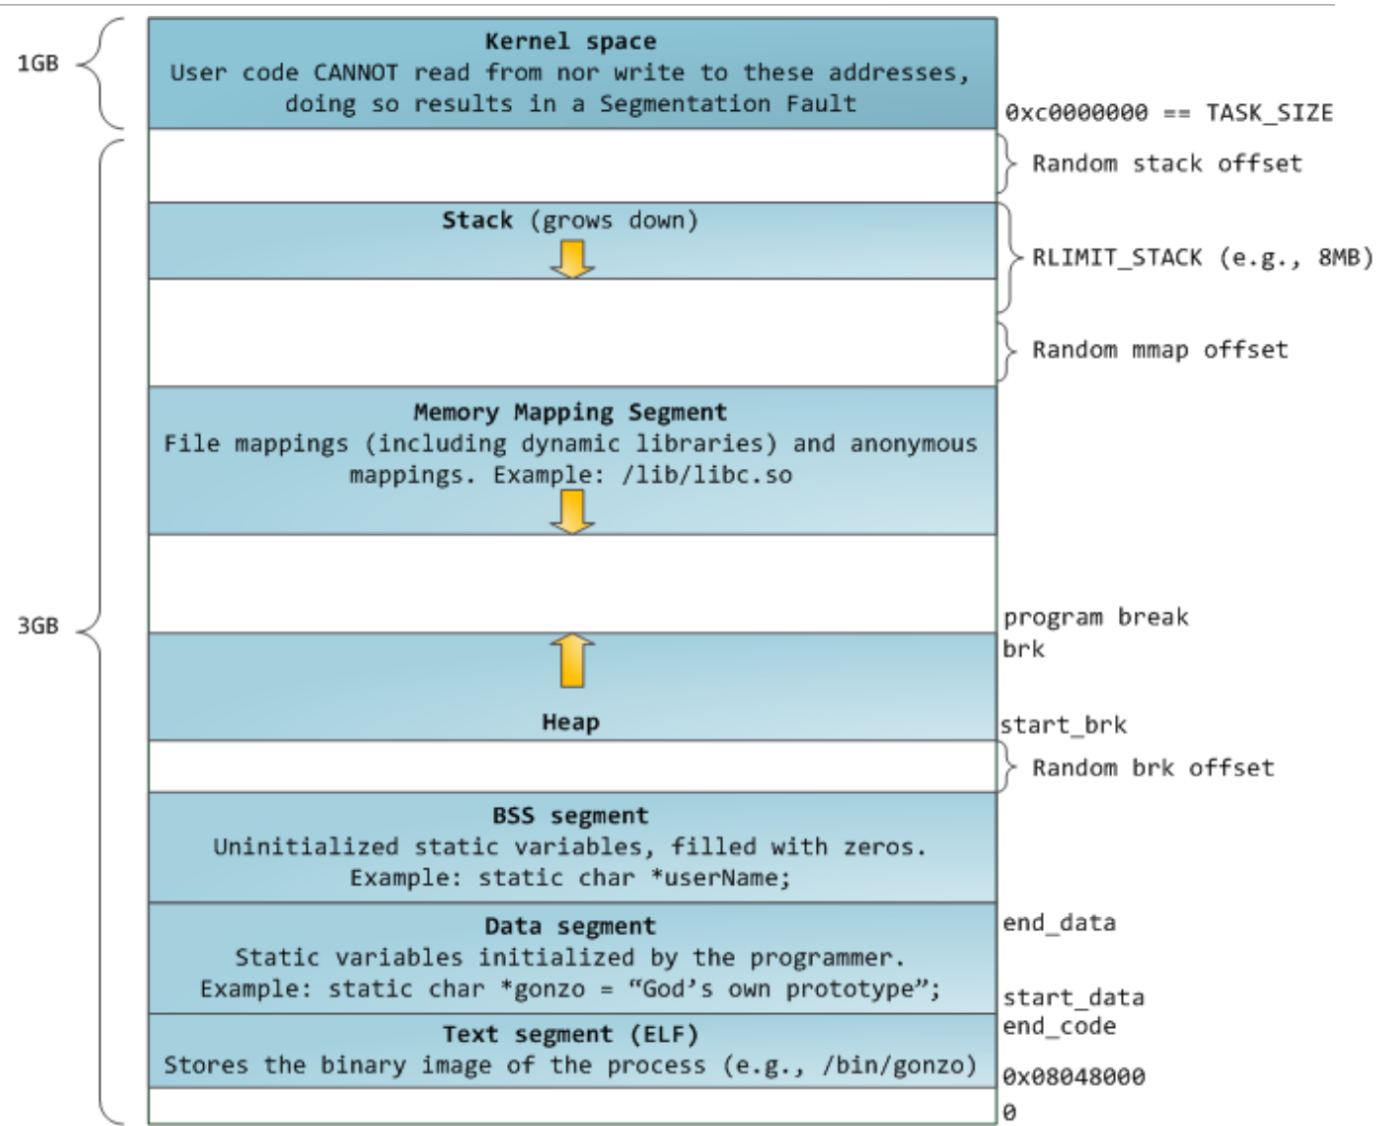
Process Manager

- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

Demo?



Applications Layout in Memory

Output of `pmap` (process mapping tool)

[illegible][illegible]

Applications Layout in Memory

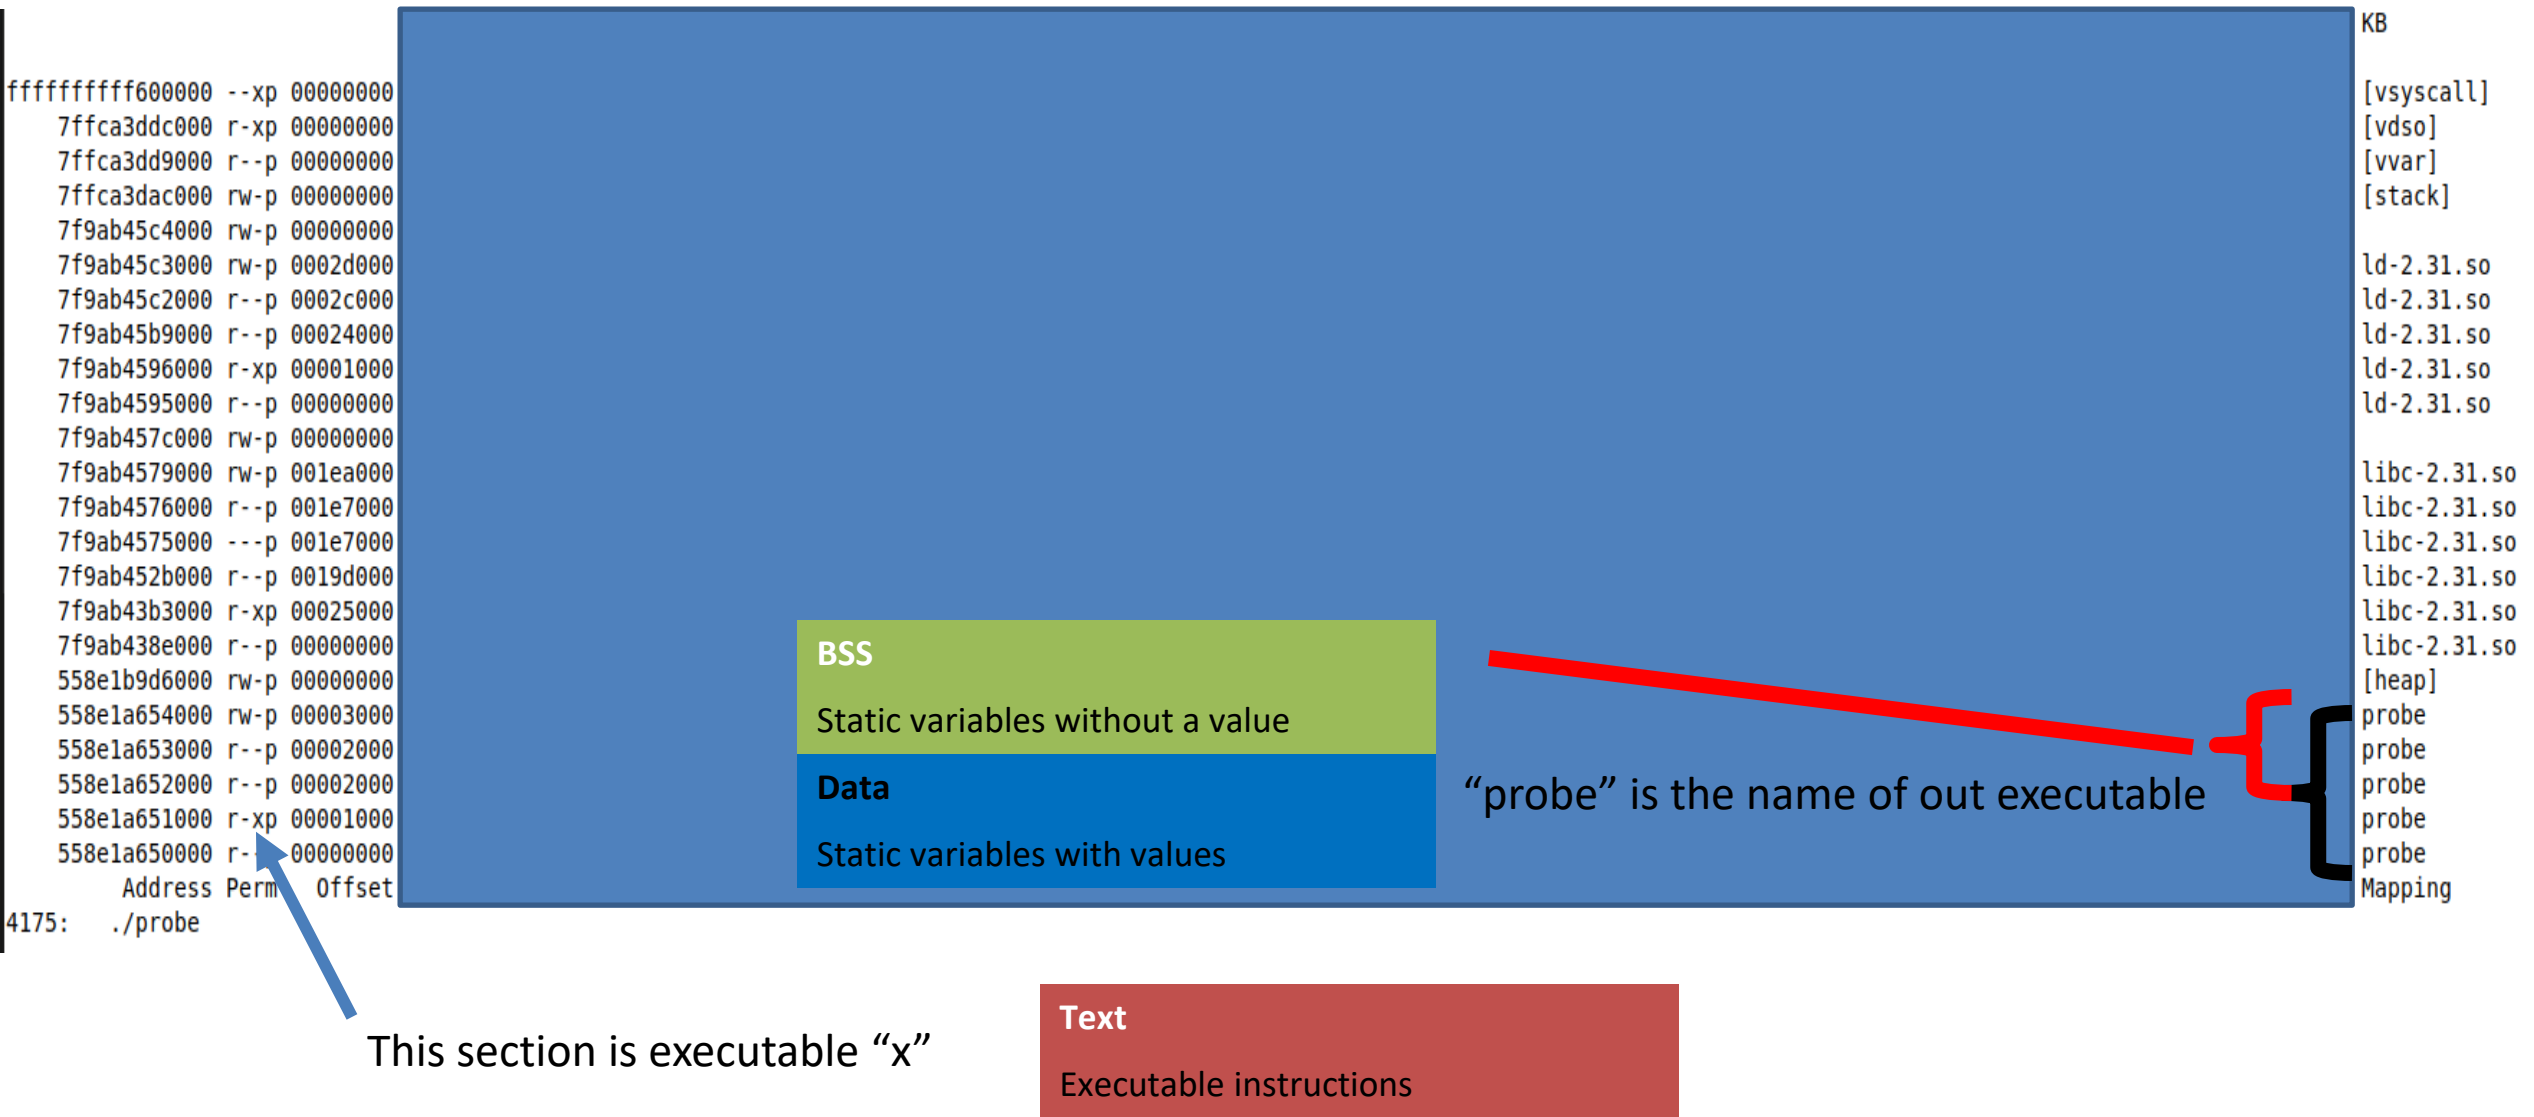
Output of `pmap` (process mapping tool)

			KB
fffffffff600000	--xp 00000000		[vsyscall]
7ffca3ddc000	r-xp 00000000		[vdso]
7ffca3dd9000	r--p 00000000		[vvar]
7ffca3dac000	rw-p 00000000		[stack]
7f9ab45c4000	rw-p 00000000		
7f9ab45c3000	rw-p 0002d000		ld-2.31.so
7f9ab45c2000	r--p 0002c000		ld-2.31.so
7f9ab45b9000	r--p 00024000		ld-2.31.so
7f9ab4596000	r-xp 00001000		ld-2.31.so
7f9ab4595000	r--p 00000000		ld-2.31.so
7f9ab457c000	rw-p 00000000		
7f9ab4579000	rw-p 001ea000		libc-2.31.so
7f9ab4576000	r--p 001e7000		libc-2.31.so
7f9ab4575000	---p 001e7000		libc-2.31.so
7f9ab452b000	r--p 0019d000		libc-2.31.so
7f9ab43b3000	r-xp 00025000		libc-2.31.so
7f9ab438e000	r--p 00000000		libc-2.31.so
558e1b9d6000	rw-p 00000000		[heap]
558e1a654000	rw-p 00003000		probe
558e1a653000	r--p 00002000		probe
558e1a652000	r--p 00002000		probe
558e1a651000	r-xp 00001000		probe
558e1a650000	r--p 00000000		probe
	Address Perm Offset		Mapping
4175:	./probe		

“probe” is the name of our executable

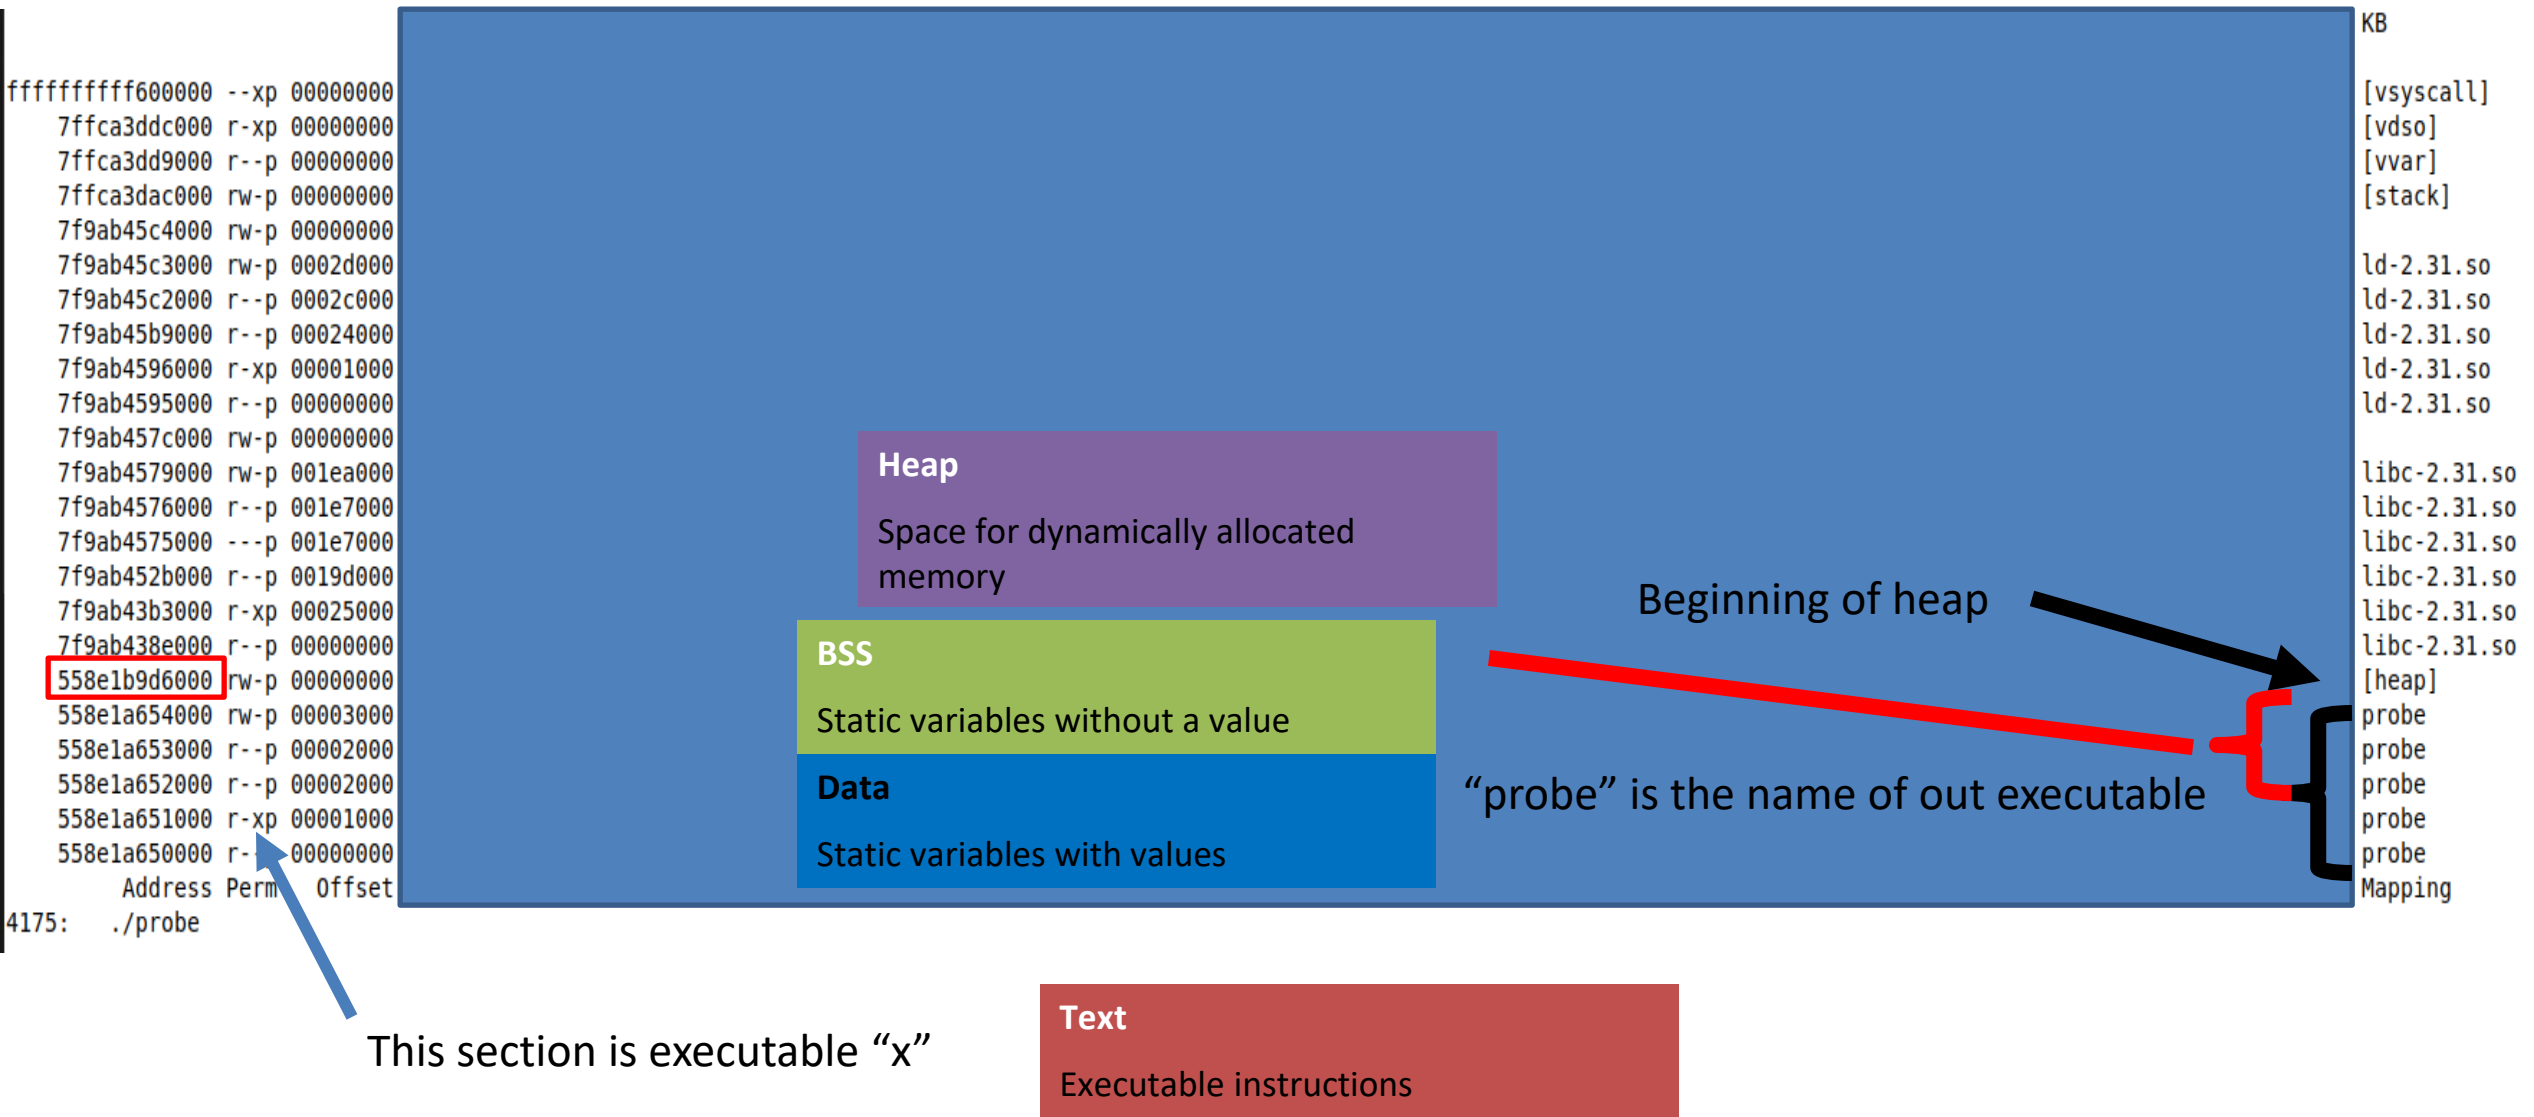
Applications Layout in Memory

Output of `pmap` (process mapping tool)



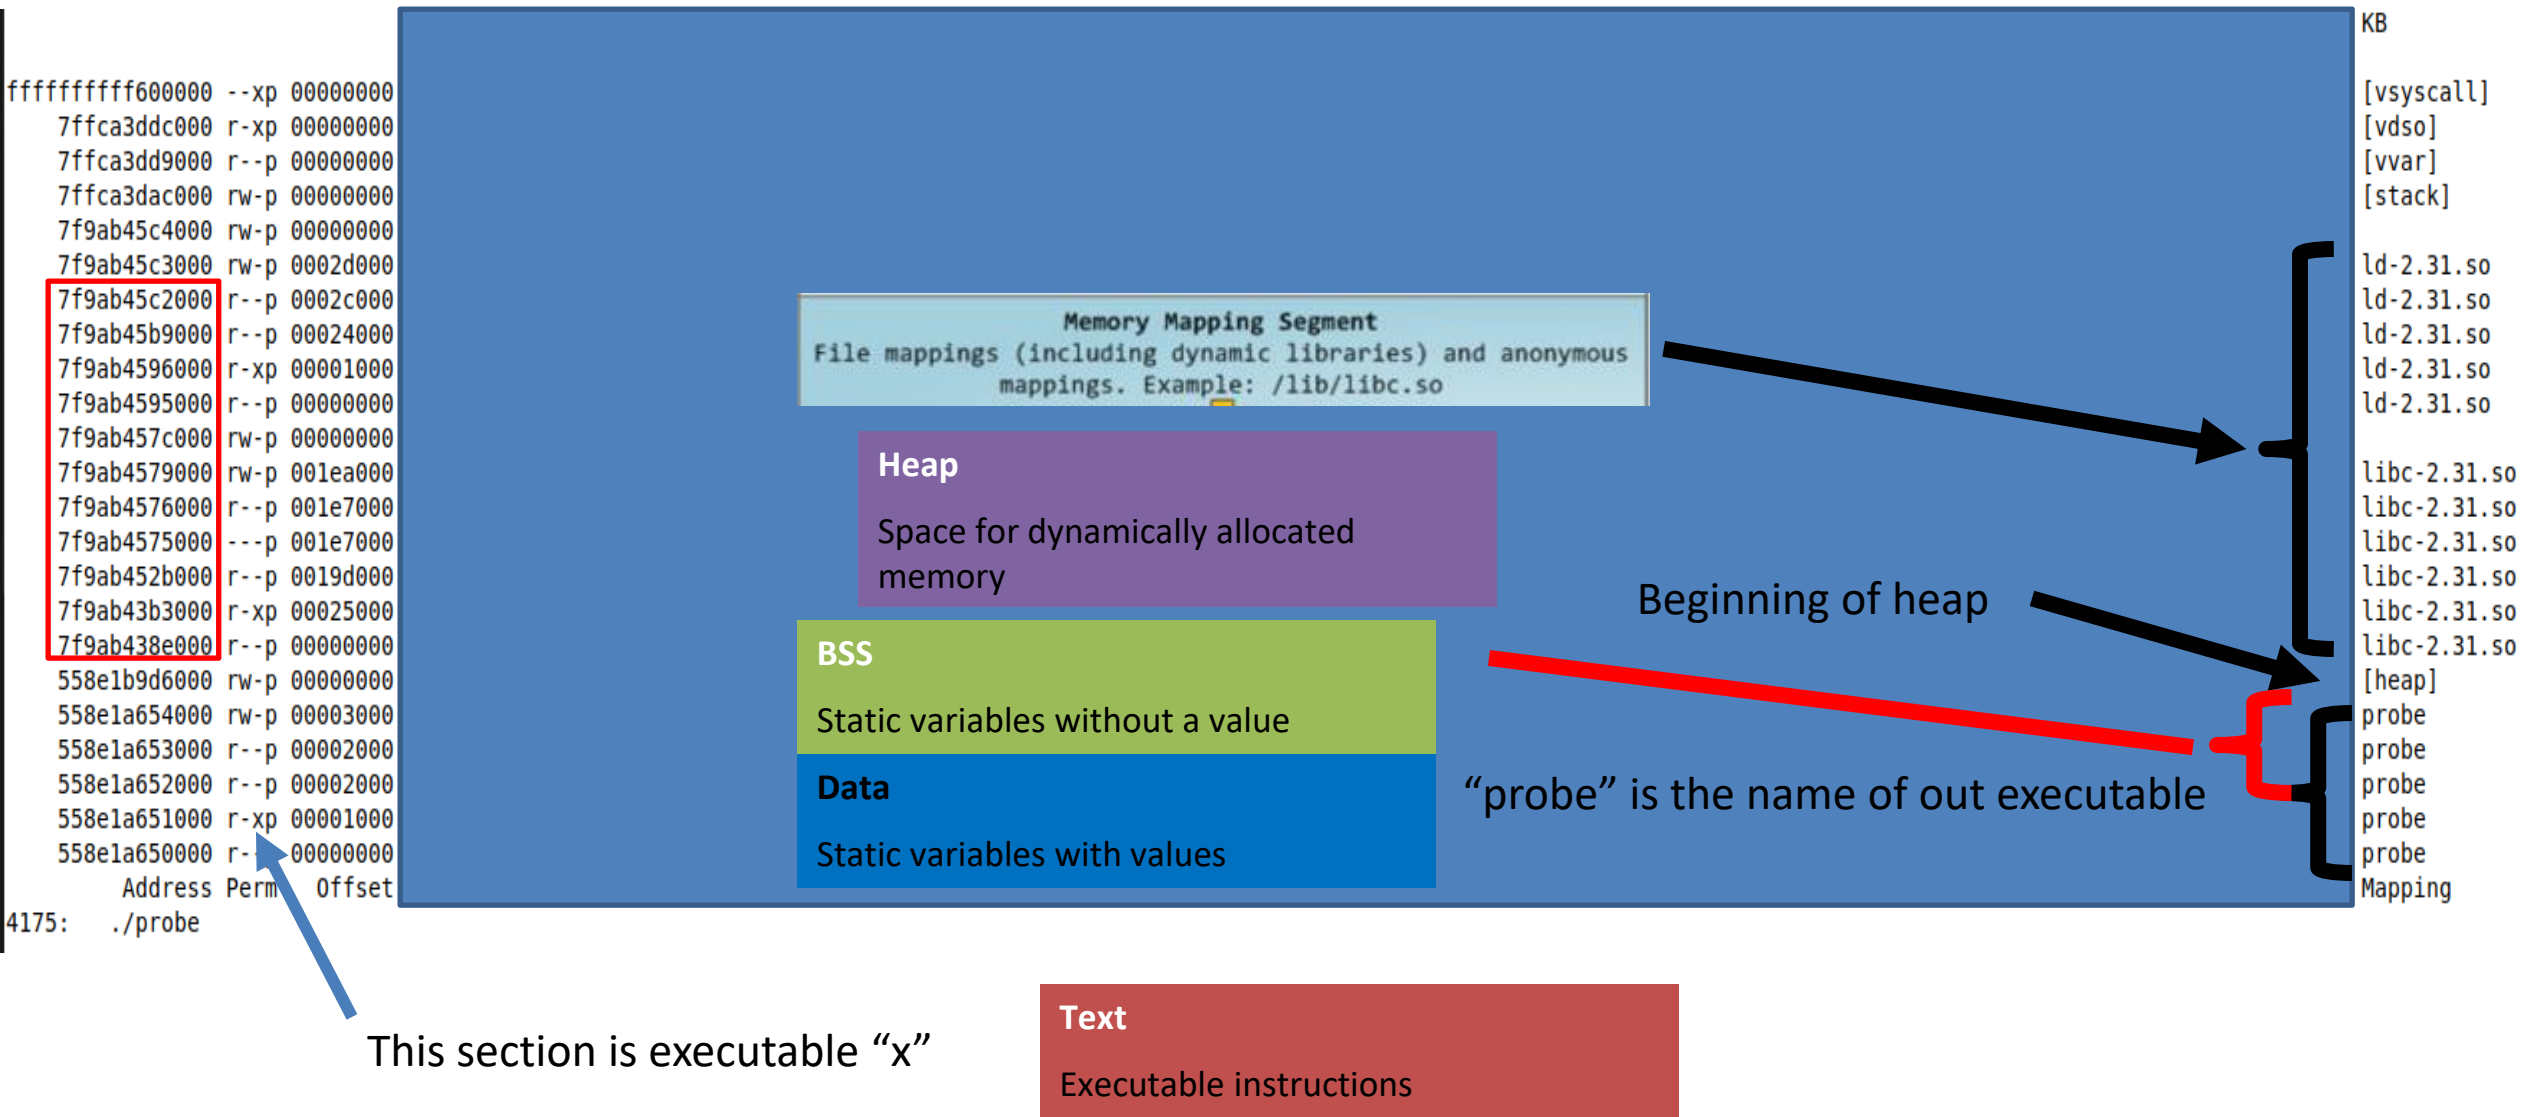
Applications Layout in Memory

Output of `pmap` (process mapping tool)



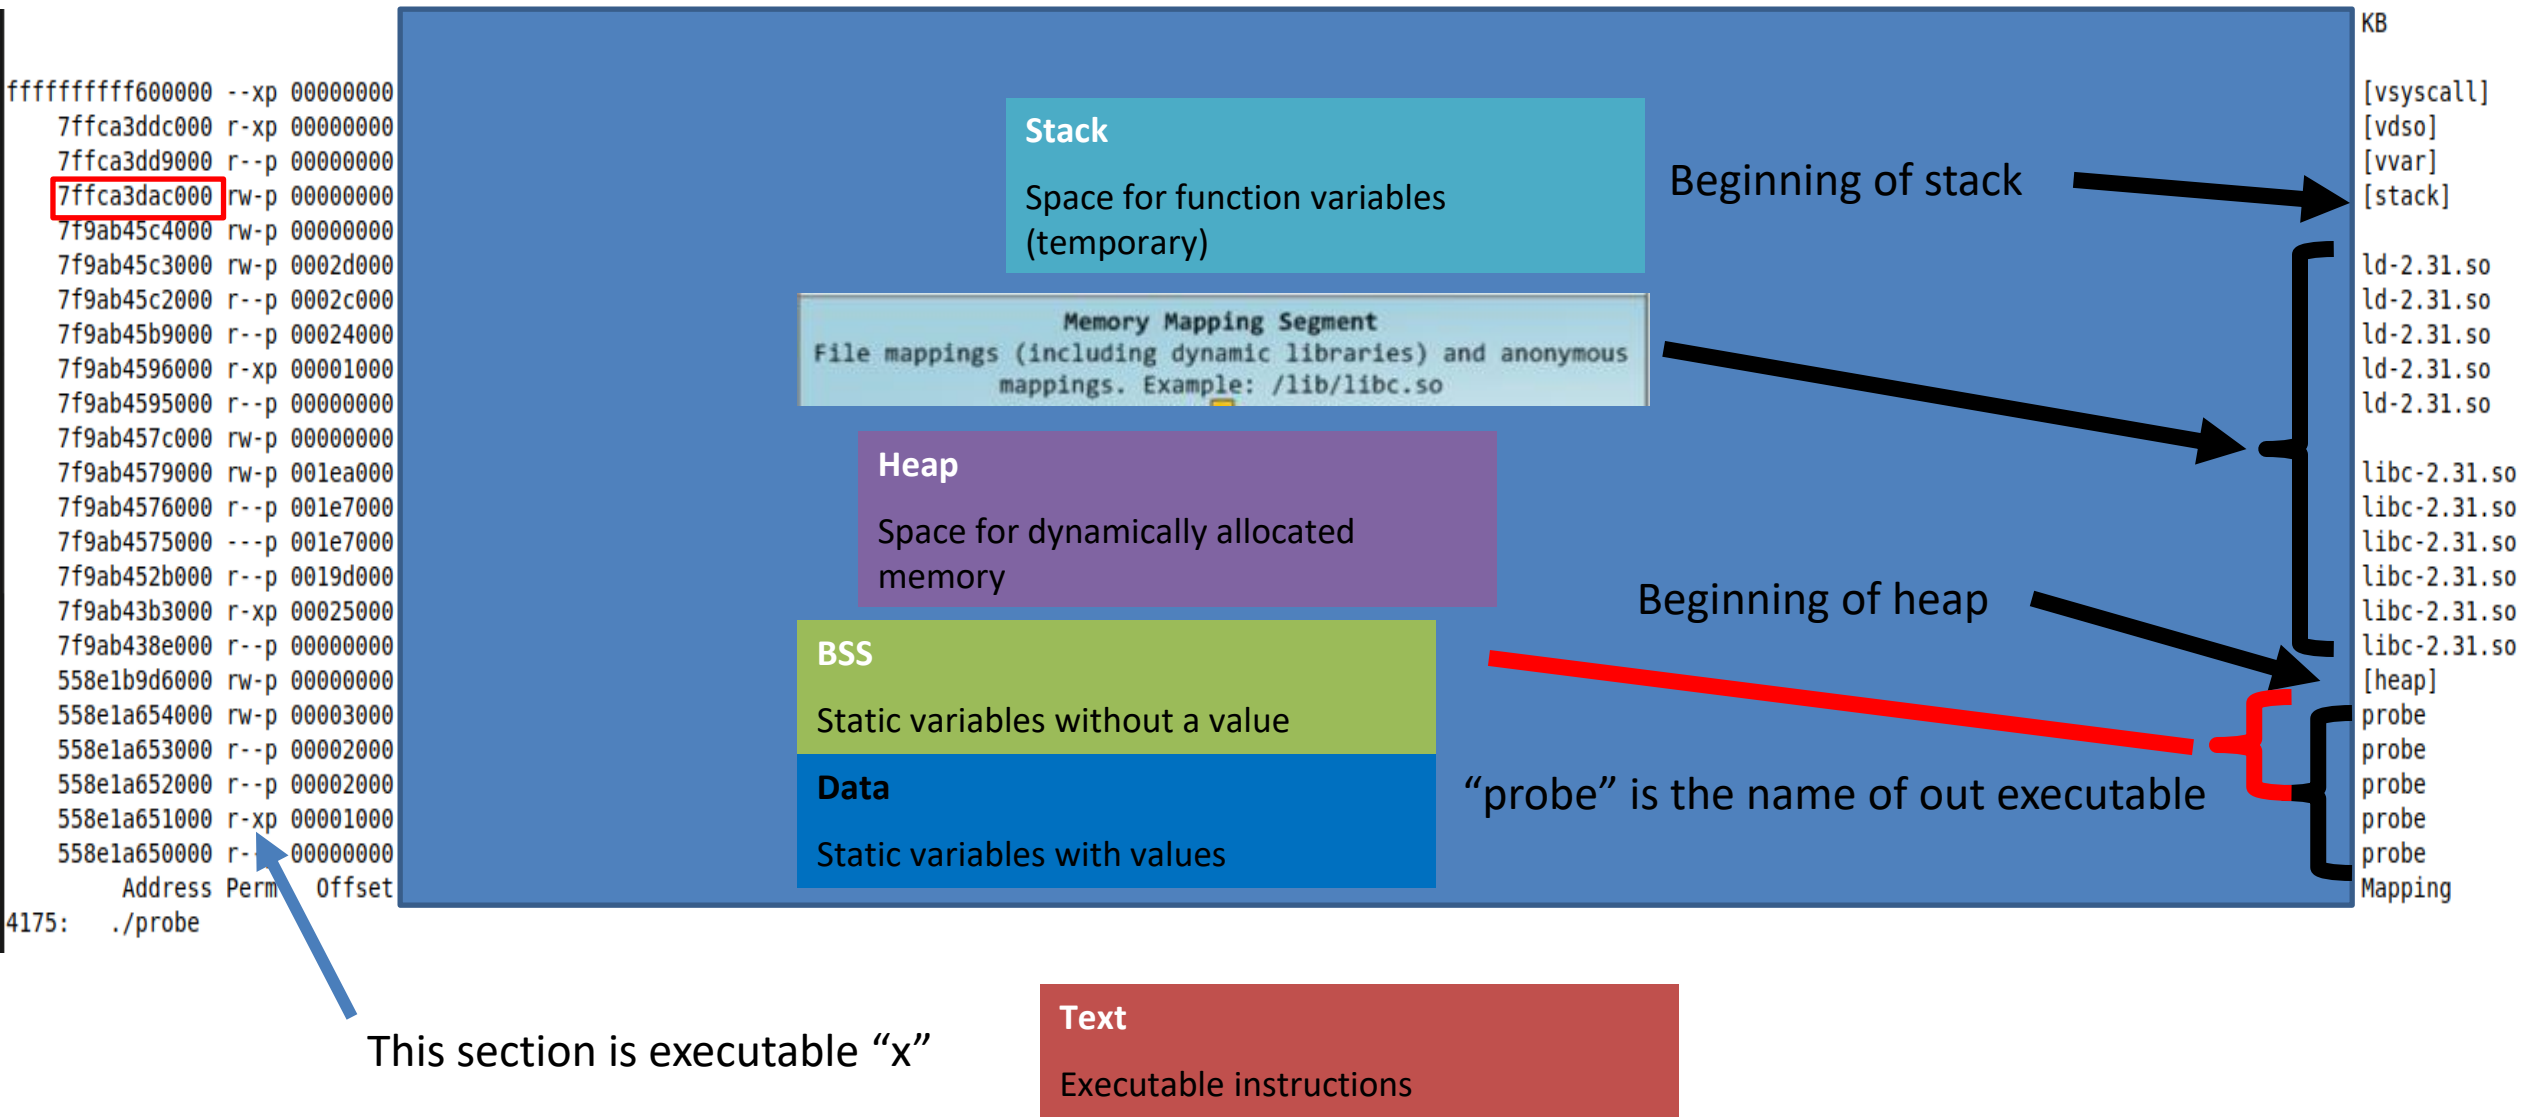
Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

When you allocate variables on the stack

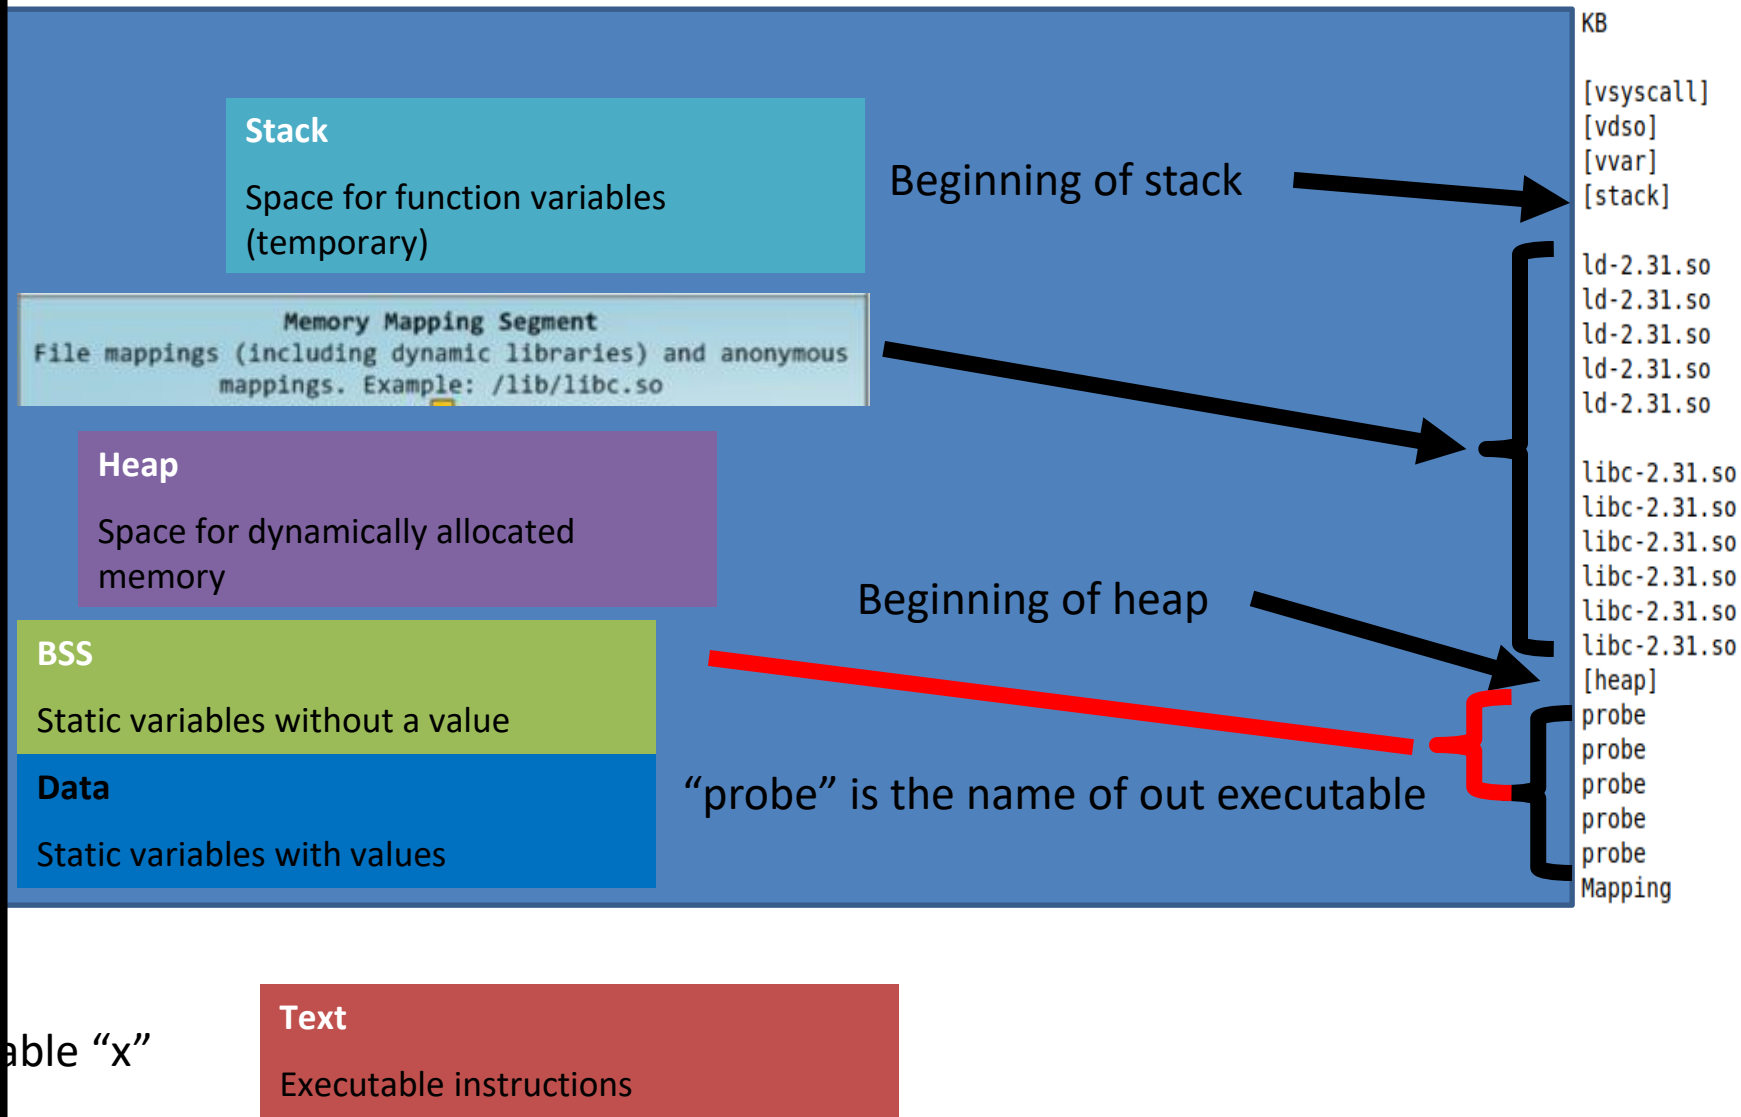


When you allocate variables on the heap



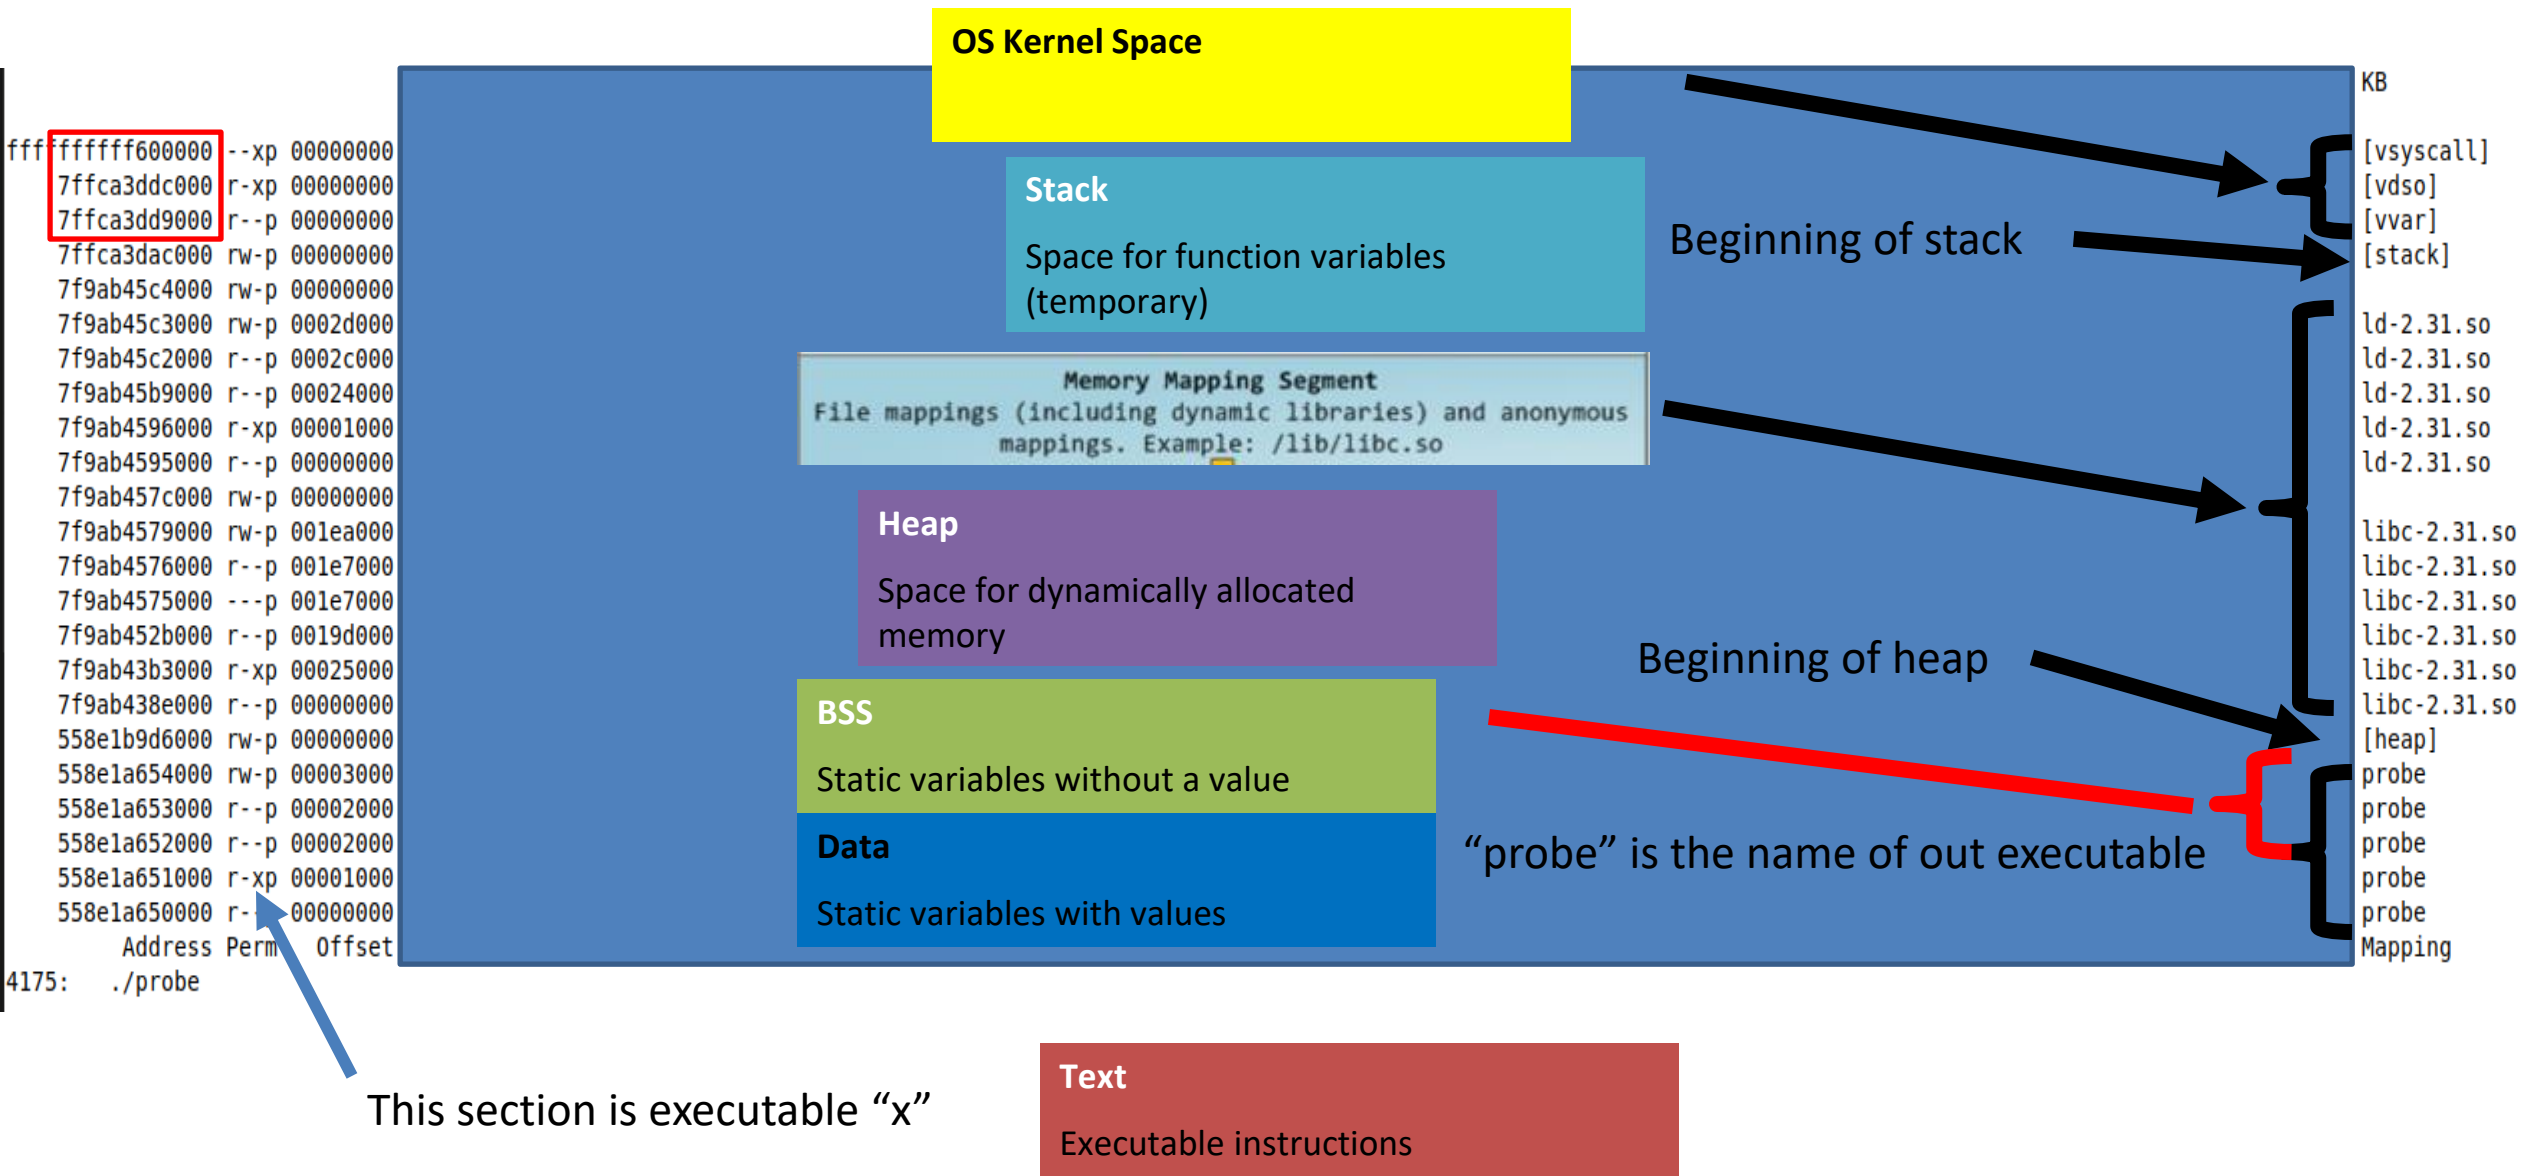
A new core memory!

Output of `pmap` (process mapping tool)



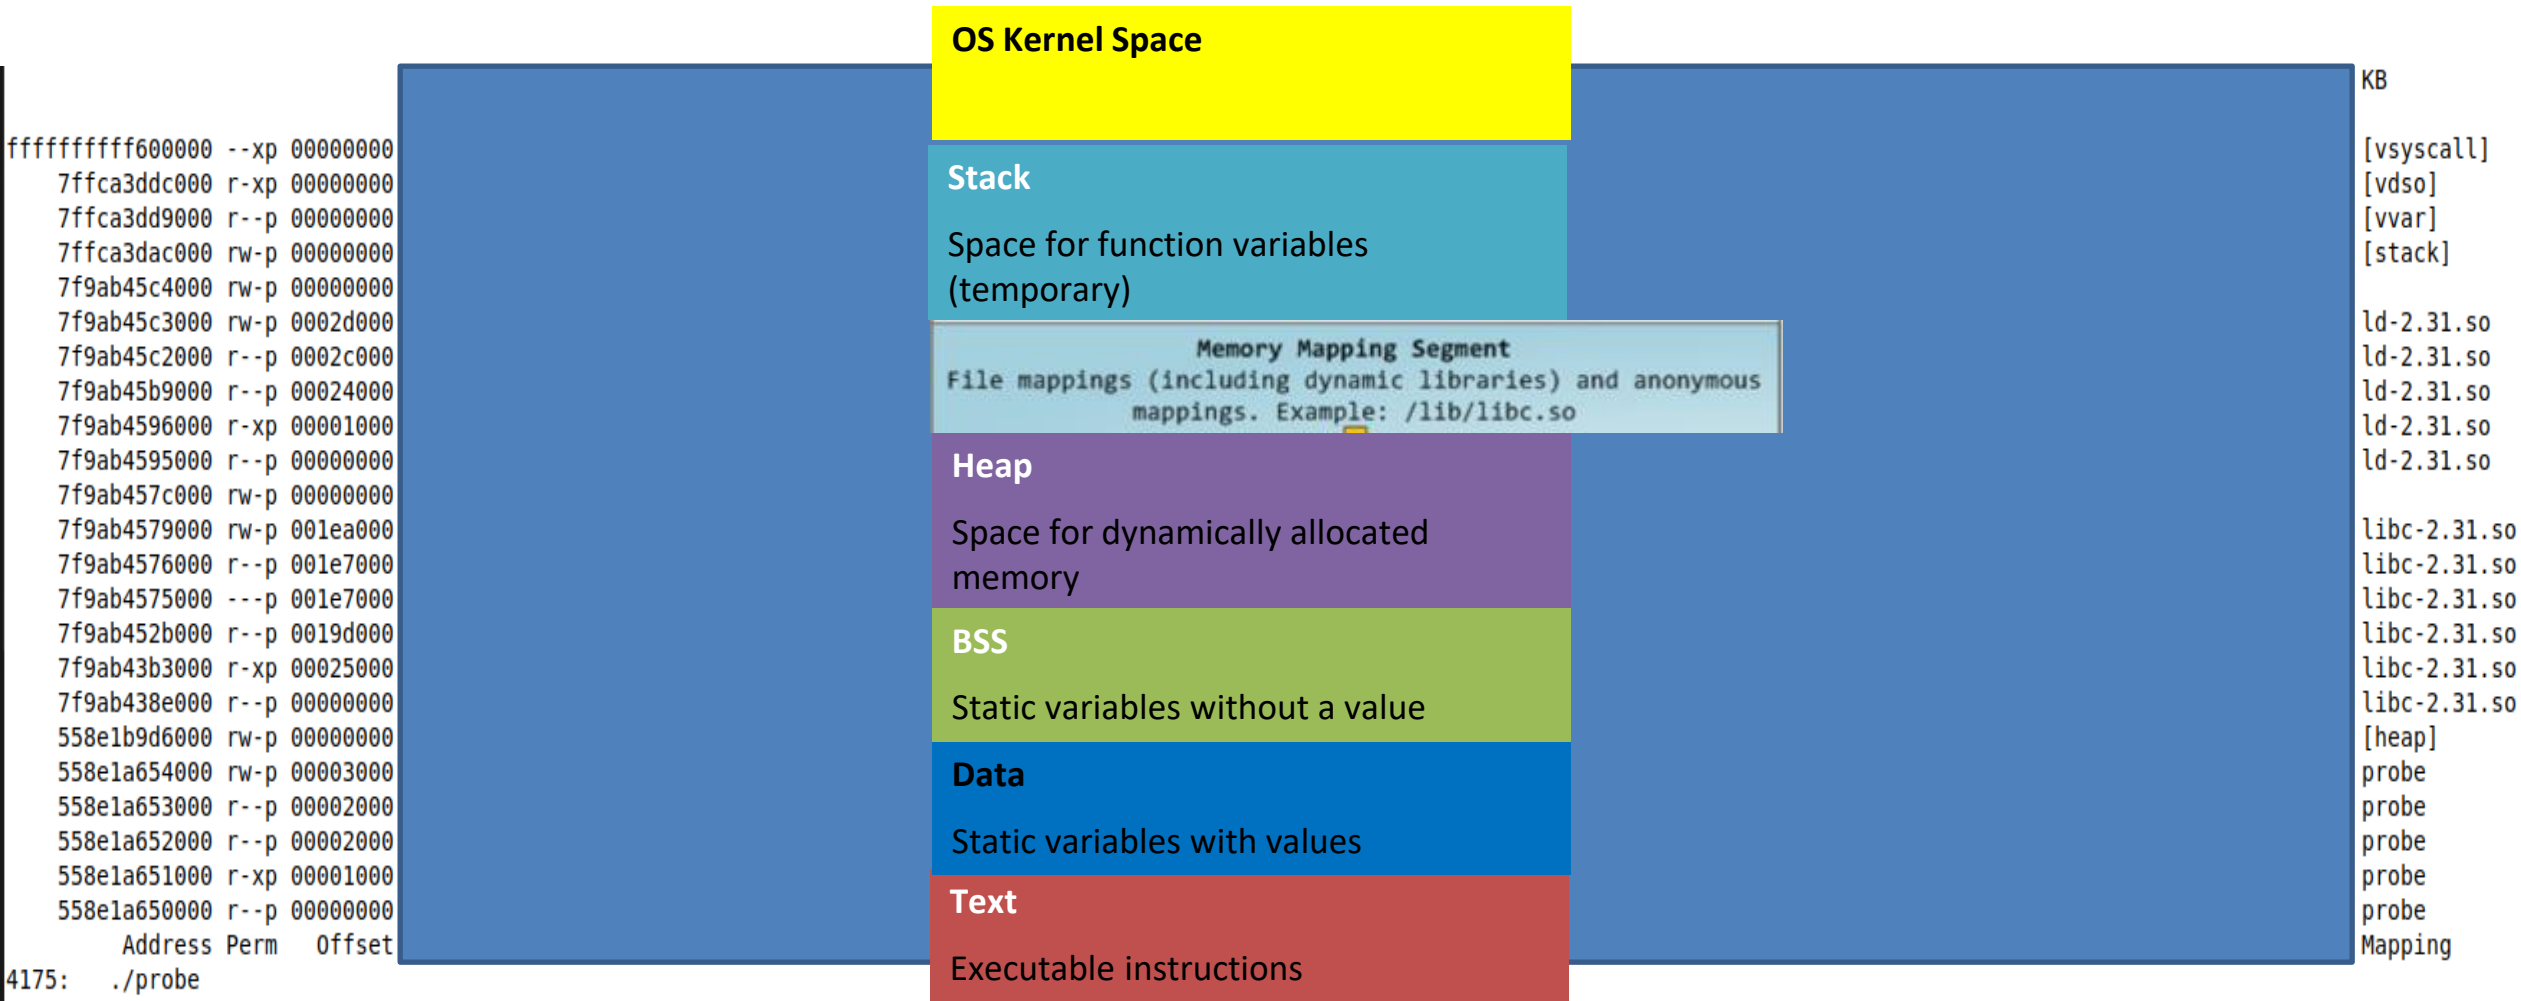
Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

Output of pmap (process mapping tool)

```
fffffffff600000 --xp 00000000
7ffca3ddc000 r-xp 00000000
7ffca3dd9000 r--p 00000000
7ffca3dac000 rw-p 00000000
7f9ab45c4000 rw-p 00000000
7f9ab45c3000 rw-p 0002d000
7f9ab45c2000 r--p 0002c000
7f9ab45b9000 r--p 00024000
7f9ab4596000 r-xp 00001000
7f9ab4595000 r--p 00000000
7f9ab457c000 rw-p 00000000
7f9ab4579000 rw-p 001ea000
7f9ab4576000 r--p 001e7000
7f9ab4575000 ---p 001e7000
7f9ab452b000 r--p 0019d000
7f9ab43b3000 r-xp 00025000
7f9ab438e000 r--p 00000000
558e1b9d6000 rw-p 00000000
558e1a654000 rw-p 00003000
558e1a653000 r--p 00002000
558e1a652000 r--p 00002000
558e1a651000 r-xp 00001000
558e1a650000 r--p 00000000
Address Perm Offset
```

```
-> the address of main      = 0x558e1a651249
-> the address of printf    = 0x7f9ab43f2e10
-> the address of getenv    = 0x7f9ab43d7020
-> a stack address         = 0x7ffca3dcb3b0
-> a global address        = 0x558e1a6540c4
-> the argv address        = 0x7ffca3dcb4f8
-> argv[0]                 = 0x7ffca3dcc45f
    value is [./probe]
-> the environ address     = 0x7ffca3dcb508
-> the envp address        = 0x7ffca3dcb508
-> getenv("PWD")           = 0x7ffca3dcc5ff
    value is [/home/seed/os-review]
-> a heap address          = 0x558e1b9d66b0
```

```
KB
[vsyscall]
[vdso]
[vvar]
[stack]
ld-2.31.so
ld-2.31.so
ld-2.31.so
ld-2.31.so
ld-2.31.so
libc-2.31.so
libc-2.31.so
libc-2.31.so
libc-2.31.so
libc-2.31.so
libc-2.31.so
[heap]
probe
probe
probe
probe
probe
Mapping
```

```
4175: ./probe
```


Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "main" located in memory?

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7ffca3dc3f2e10

-> the address of getenv = 0x7ffca3dc3d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

-> argv[0] value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

ffffffffff600000 --xp 00000000

7ffca3ddc000 r-xp 00000000

7ffca3dd9000 r--p 00000000

7ffca3dac000 rw-p 00000000

7f9ab45c4000 rw-p 00000000

7f9ab45c3000 rw-p 0002d000

7f9ab45c2000 r--p 0002c000

7f9ab45b9000 r--p 00024000

7f9ab4596000 r-xp 00001000

7f9ab4595000 r--p 00000000

7f9ab457c000 rw-p 00000000

7f9ab4579000 rw-p 001ea000

7f9ab4576000 r--p 001e7000

7f9ab4575000 ---p 001e7000

7f9ab452b000 r--p 0019d000

7f9ab43b3000 r-xp 00025000

7f9ab438e000 r--p 00000000

558e1b9d6000 rw-p 00000000

558e1a654000 rw-p 00003000

558e1a653000 r--p 00002000

558e1a652000 r--p 00002000

558e1a651000 r-xp 00001000

558e1a650000 r--p 00000000

Address Perm Offset

4175: ./probe

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe

probe

probe

probe

Mapping

Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "main" located in memory?

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7f9ab43f2e10

-> the address of getenv = 0x7f9ab43d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

-> argv[0] value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

558e1a651000 r-xp 00001000

Address Perm Offset

4175: ./probe

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe

probe

probe

Mapping

`main` is code in our program, so it goes inside the text segment

Applications Layout in Memory

Output of pmap (process mapping tool)

Where is "printf" located in memory?

```

-> the address of main      = 0x558e1a651249
-> the address of printf    = 0x7ff9ab43f2e10
-> the address of getenv    = 0x7ff9ab43d7020
-> a stack address         = 0x7ffca3dcb3b0
-> a global address        = 0x558e1a6540c4
-> the argv address        = 0x7ffca3dcb4f8
-> argv[0]                 = 0x7ffca3dcc45f
    value is [./probe]
-> the environ address     = 0x7ffca3dcb508
-> the envp address        = 0x7ffca3dcb508
-> getenv("PWD")           = 0x7ffca3dcc5ff
    value is [/home/seed/os-review]
-> a heap address          = 0x558e1b9d66b0

```

Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "printf" located in memory?

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7f9ab43f2e10

-> the address of getenv = 0x7f9ab43d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

-> argv[0] value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

7f9ab43b3000 r-xp 00025000

4175: ./probe

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe

probe

probe

probe

probe

Mapping

`printf` is executable code from a shared library (libc) so we are in the memory mapping segment!

Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "argv" located in memory?

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7f9ab43f2e10

-> the address of getenv = 0x7f9ab43d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

argv[0]

value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") = 0x7ffca3dcc5ff

value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

4175: ./probe

Address Perm Offset

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe

probe

probe

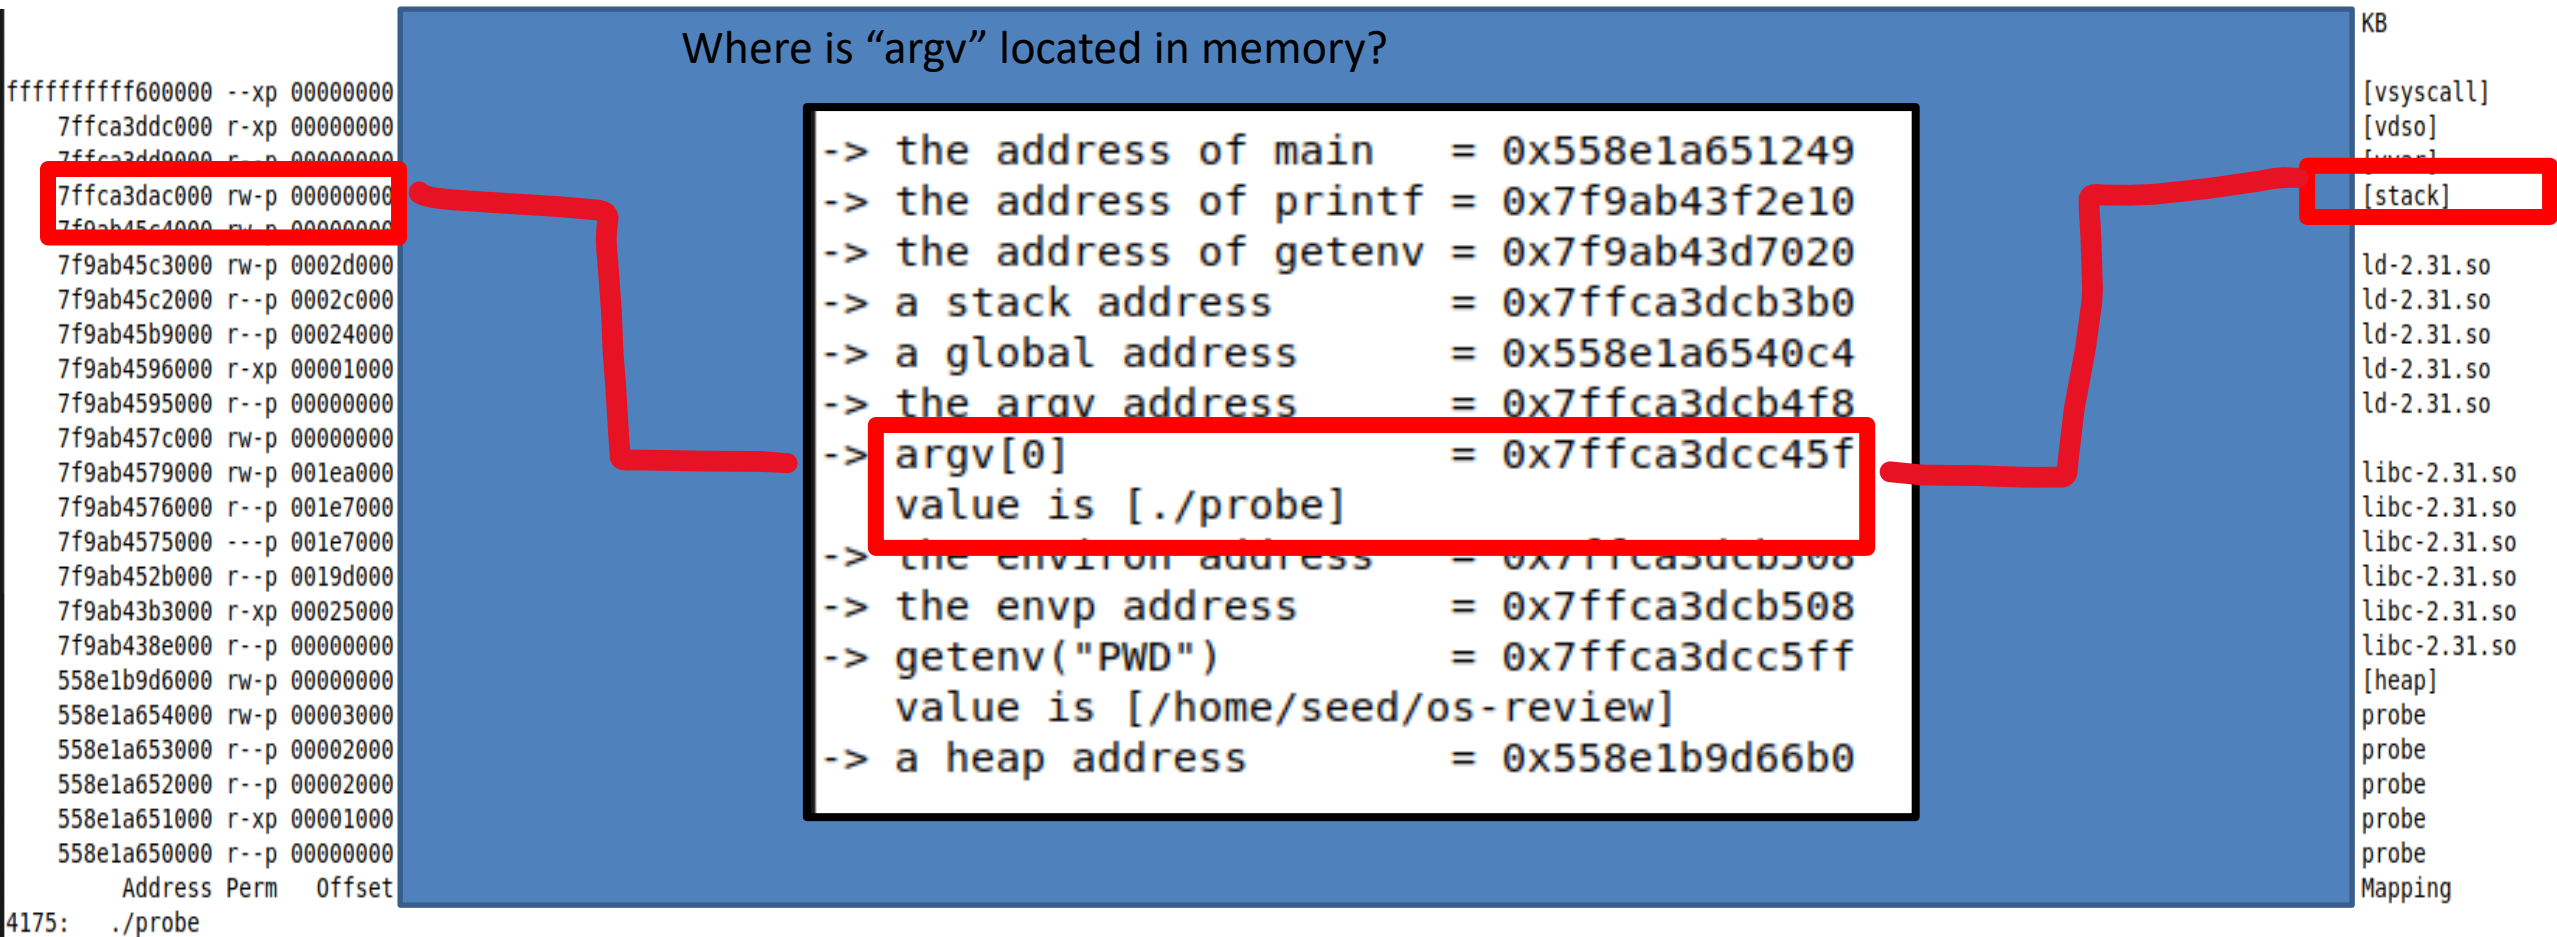
probe

Mapping

`argv` is an array that holds the command line parameters passed into this program

Applications Layout in Memory

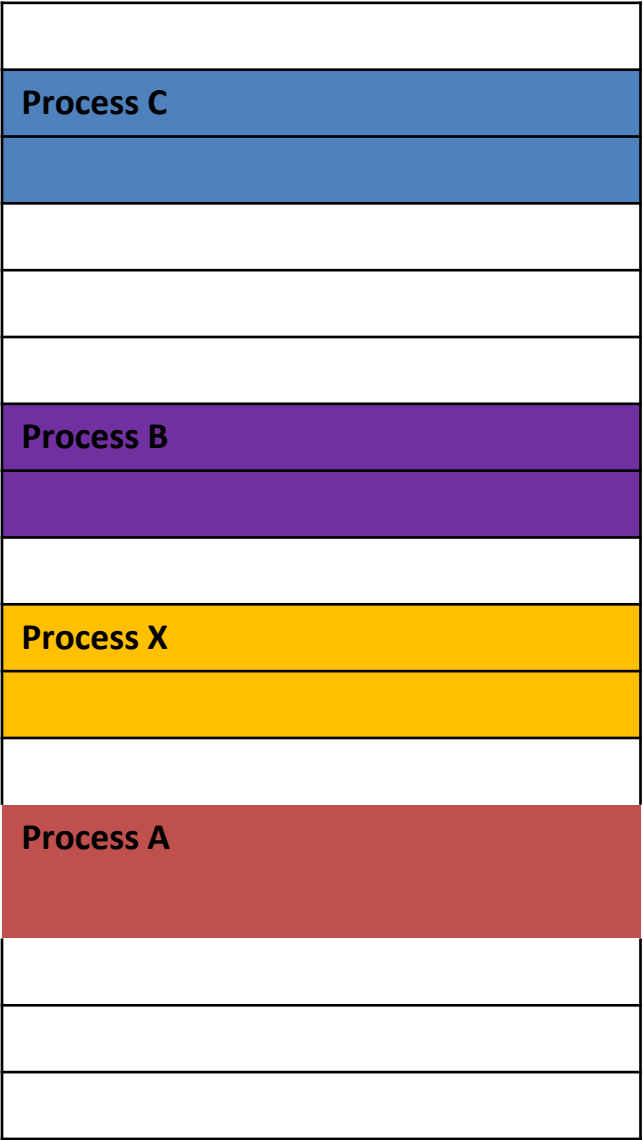
Output of `pmap` (process mapping tool)



`argv` is the argument to the main function, so we are in the stack!

Applications Layout in Memory

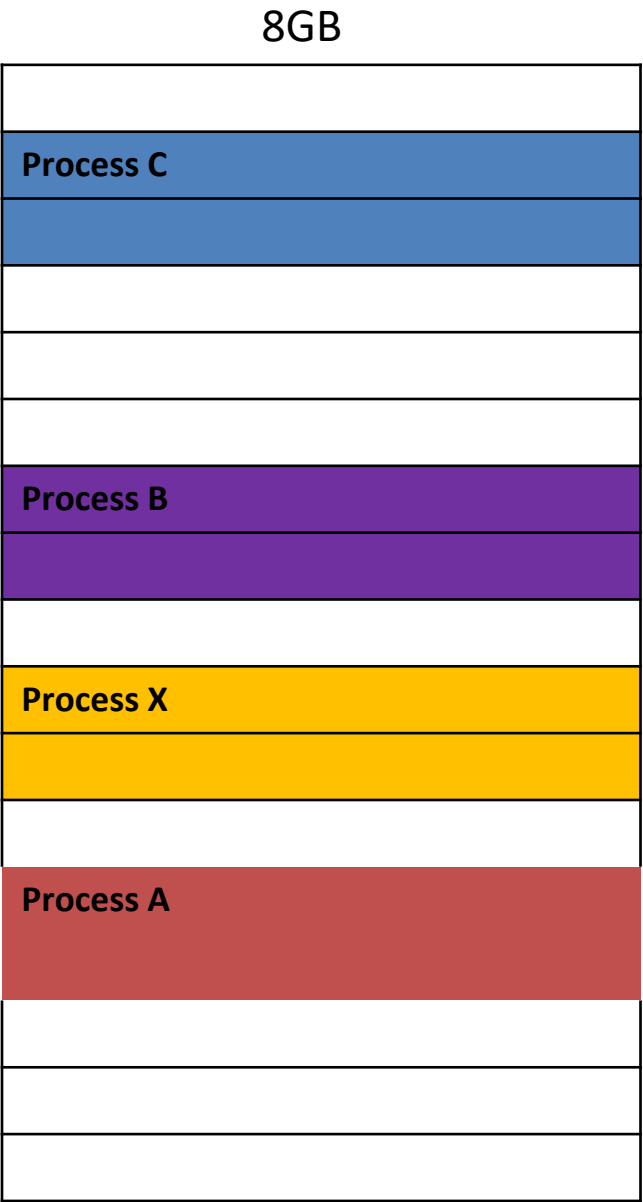
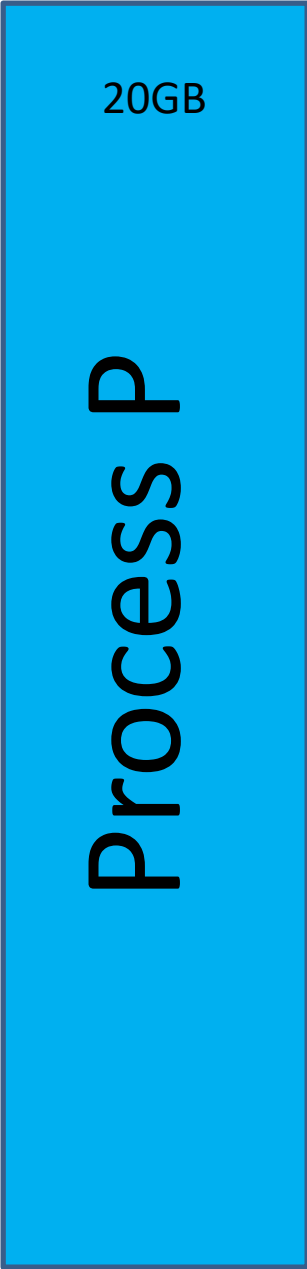
We have many programs that are actively running on our computer



Applications Layout in Memory

We have many programs that are actively running on our computer

What if we have a program that is bigger than out entire main memory?

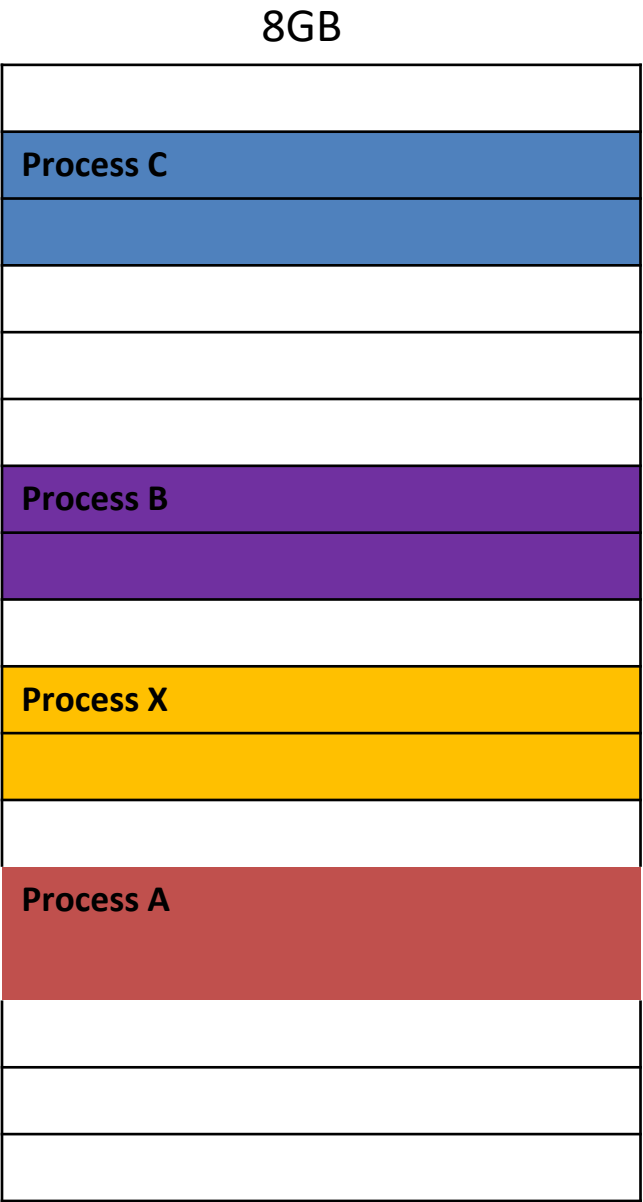
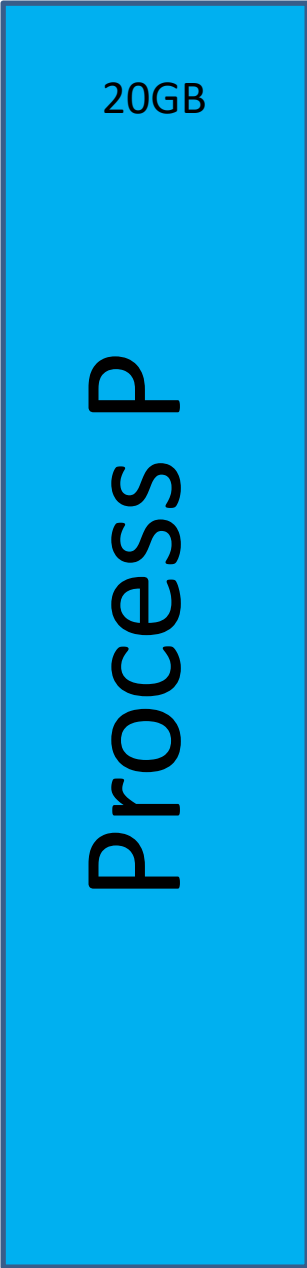


Applications Layout in Memory

We have many programs that are actively running on our computer

What if we have a program that is bigger than out entire main memory?

Does our computer crash?



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage



Secondary Storage

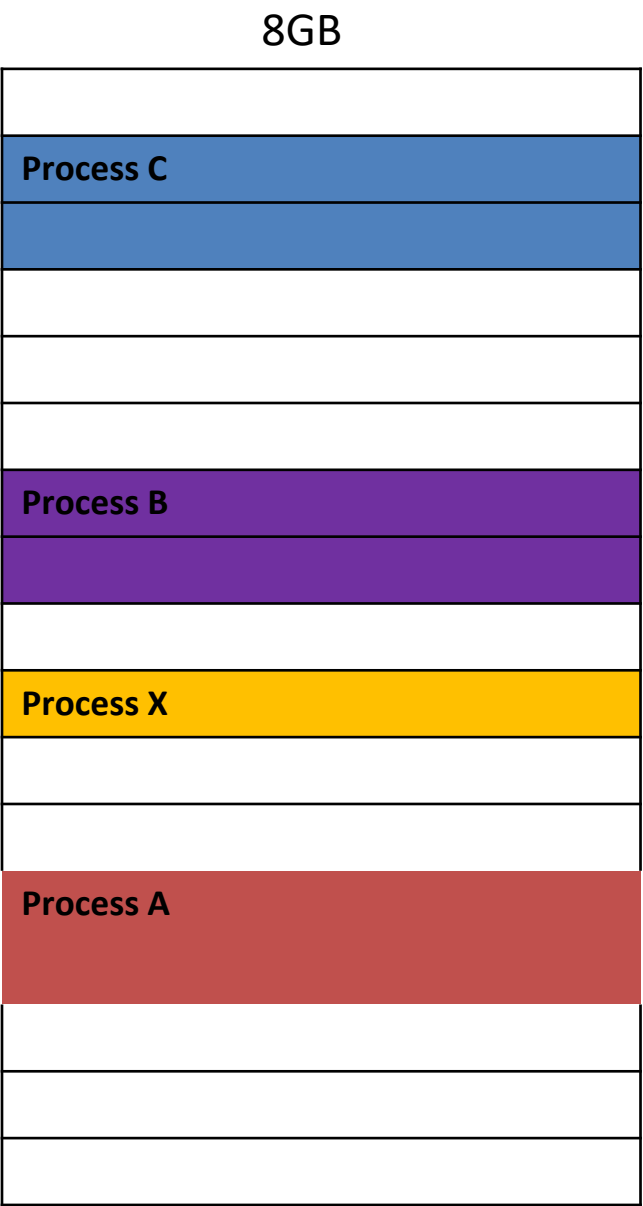
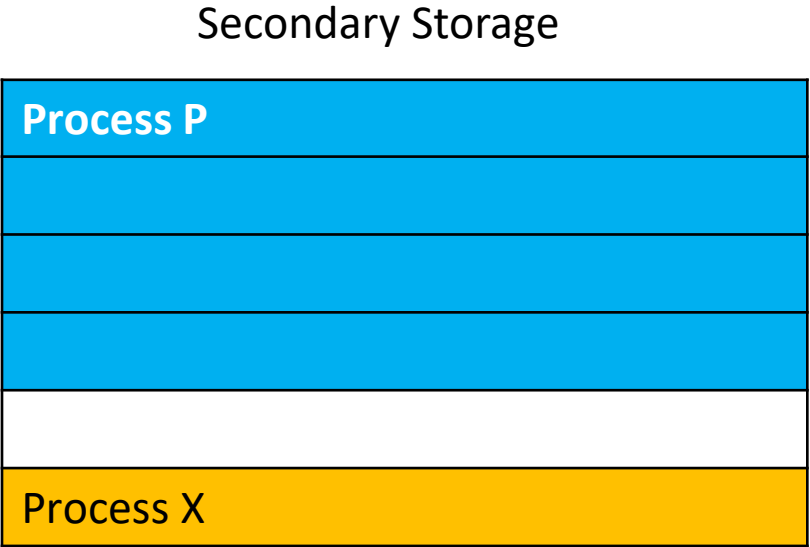
8GB

Process C
Process B
Process X
Process A

Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

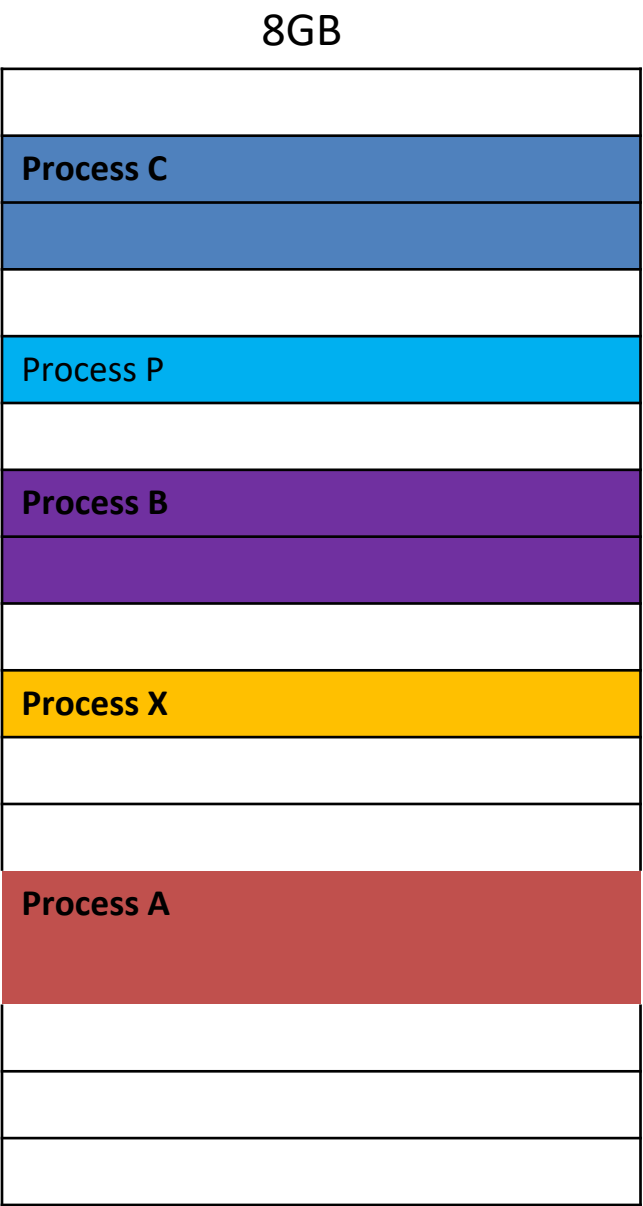
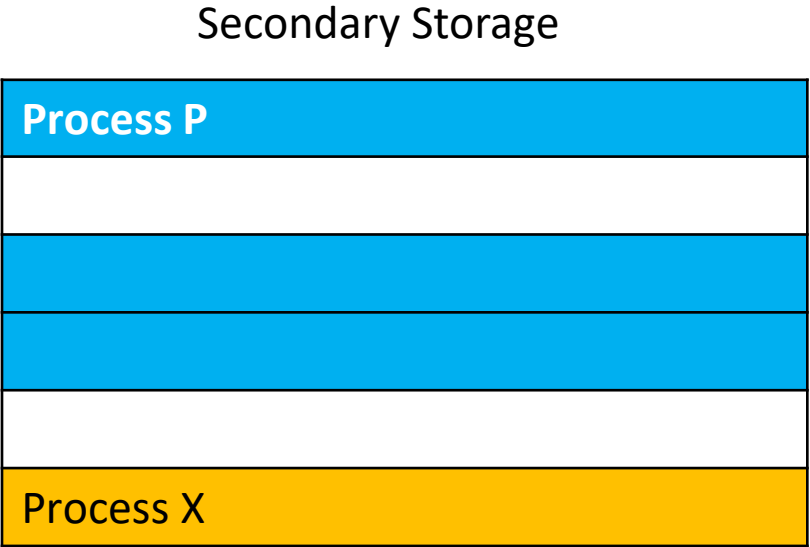
We split the process into smaller **pages**. Load pages into memory only when needed



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

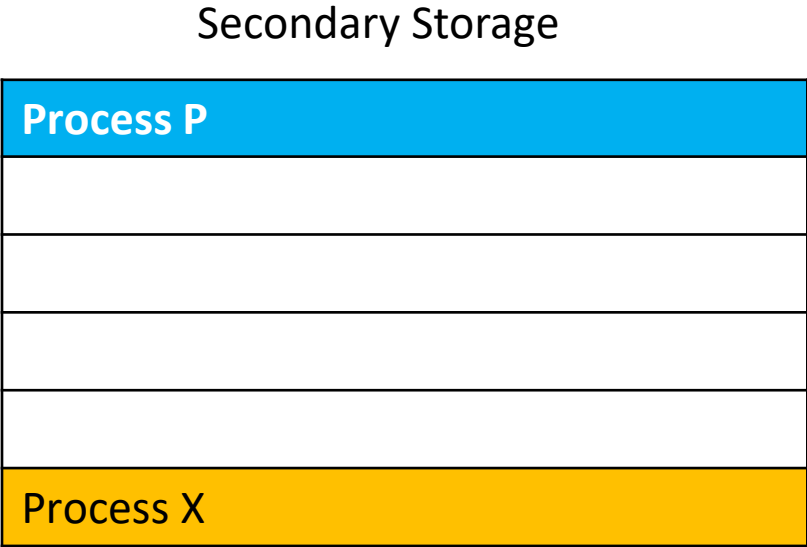
We split the process into smaller **pages**. Load pages into memory only when needed



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

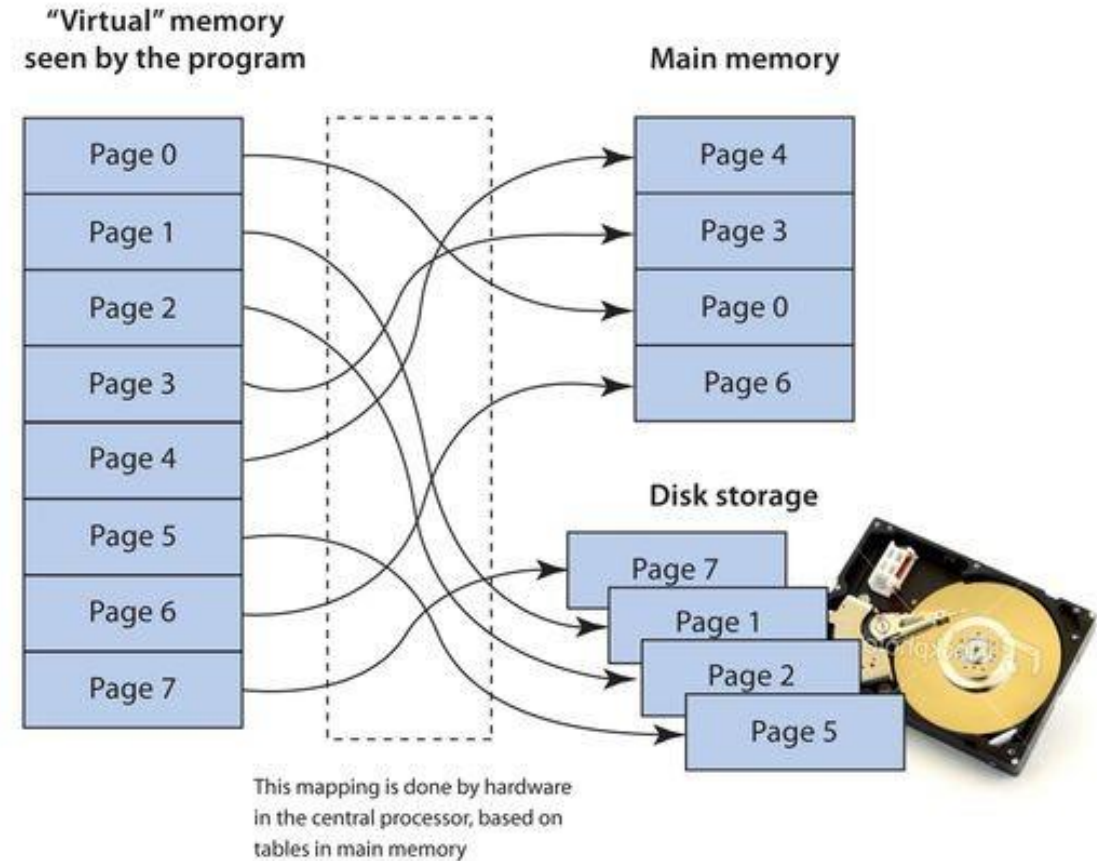
We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed

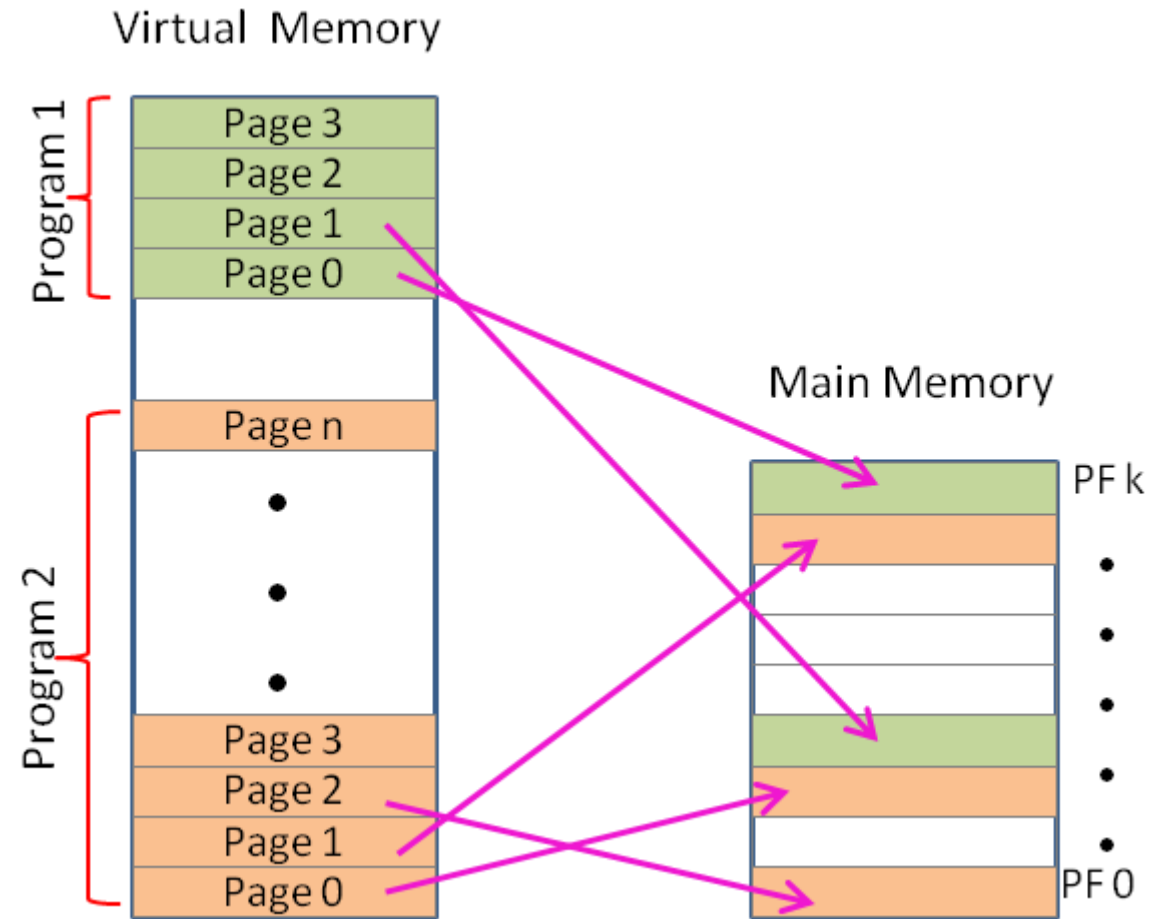


Constantly swapping stuff in and out of main memory

Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed

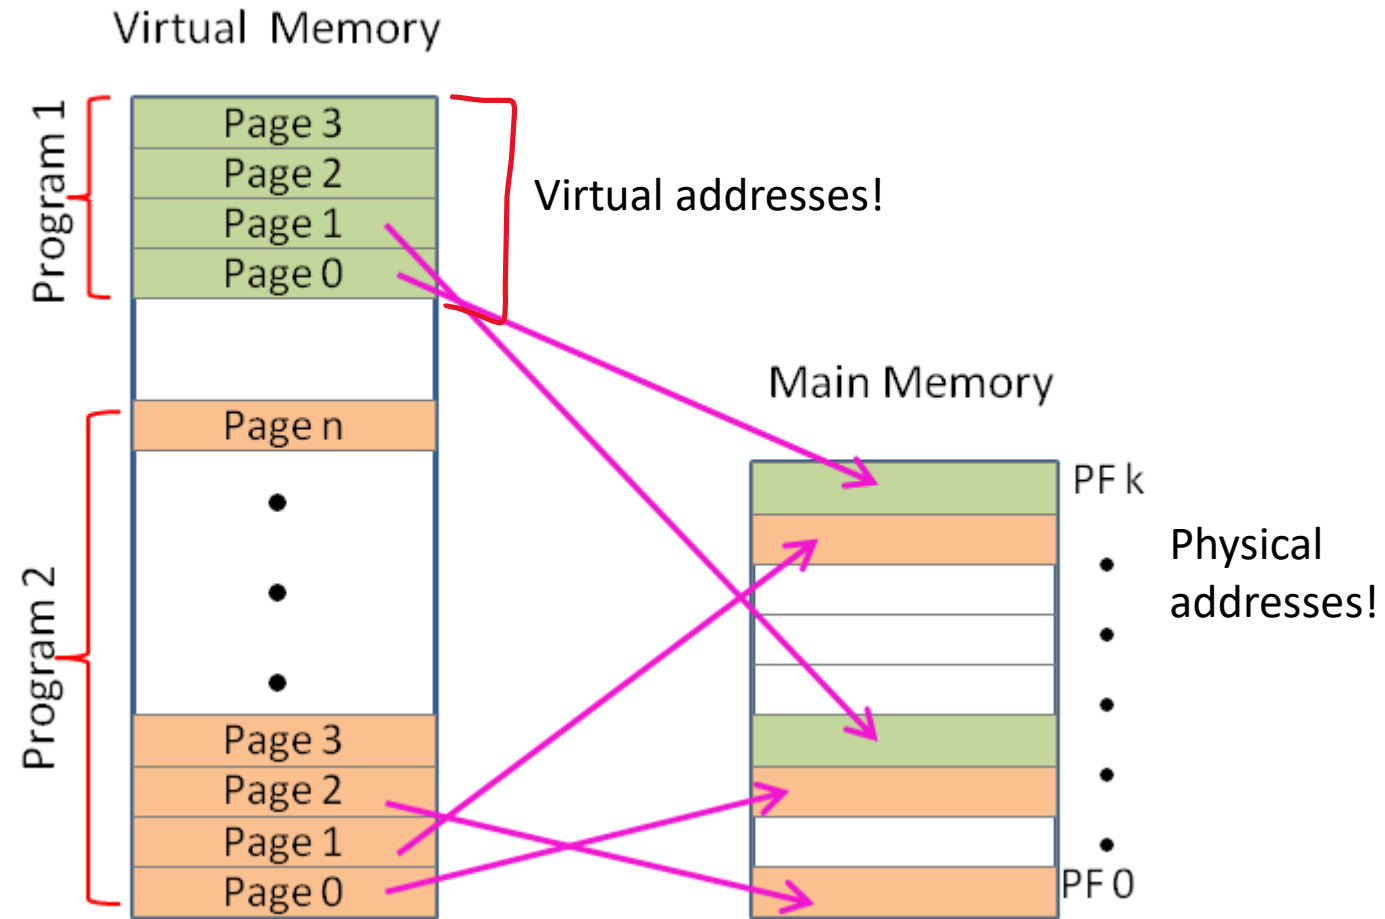


A process in memory is not contiguous

Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed



A process in memory is not contiguous

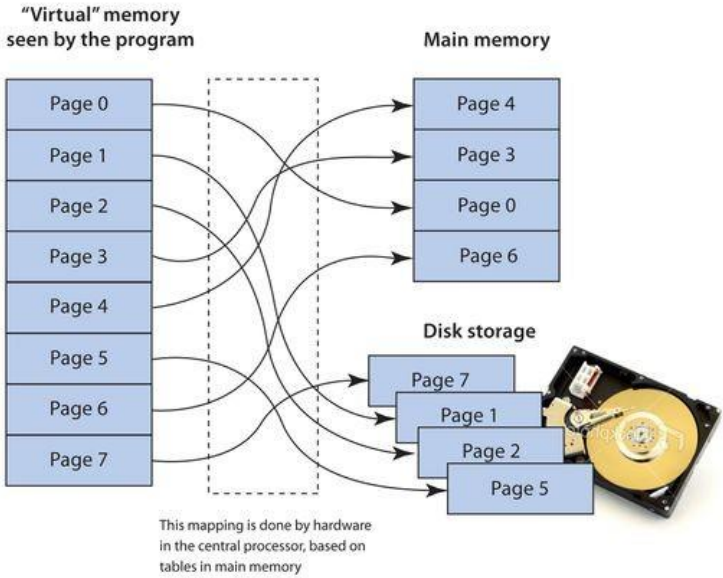
In probe.c, we are seeing virtual addresses!

Internal fragmentation vs external fragmentation

OS Review

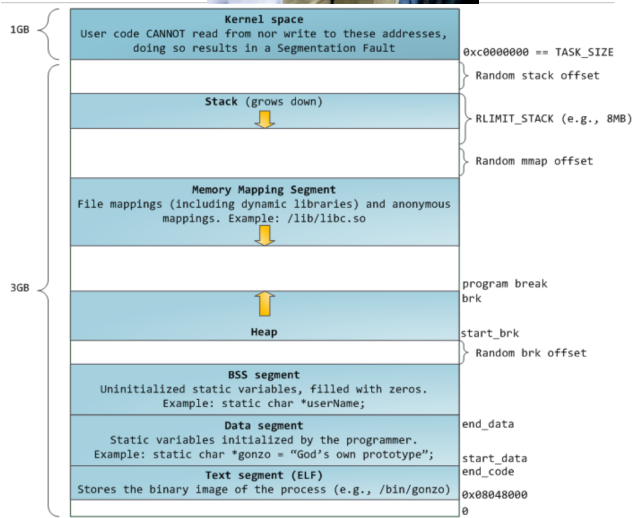
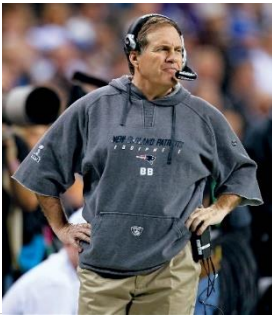
Memory Manager

- Manages how physical memory is utilized



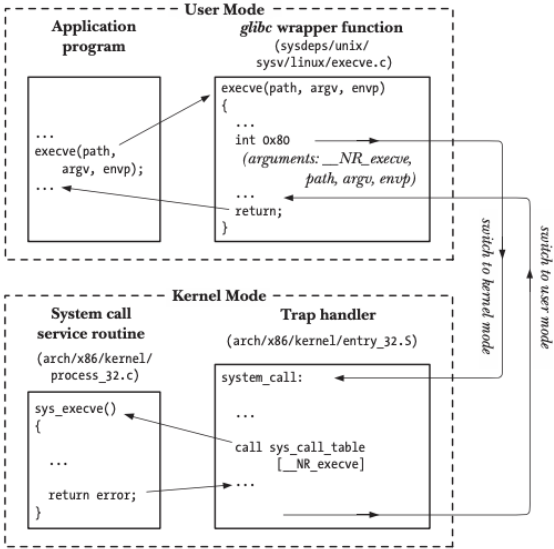
Process Manager

- Manages how processes are structured and how to handle many processes running at once



Interface Manager

- Manages communication between apps and hardware





Traffic Manager

- Manages which programs should be executed by the CPU

Process A (Ready)

Process B (Urgent)

Process C (Ready)

Process D (Blocked)

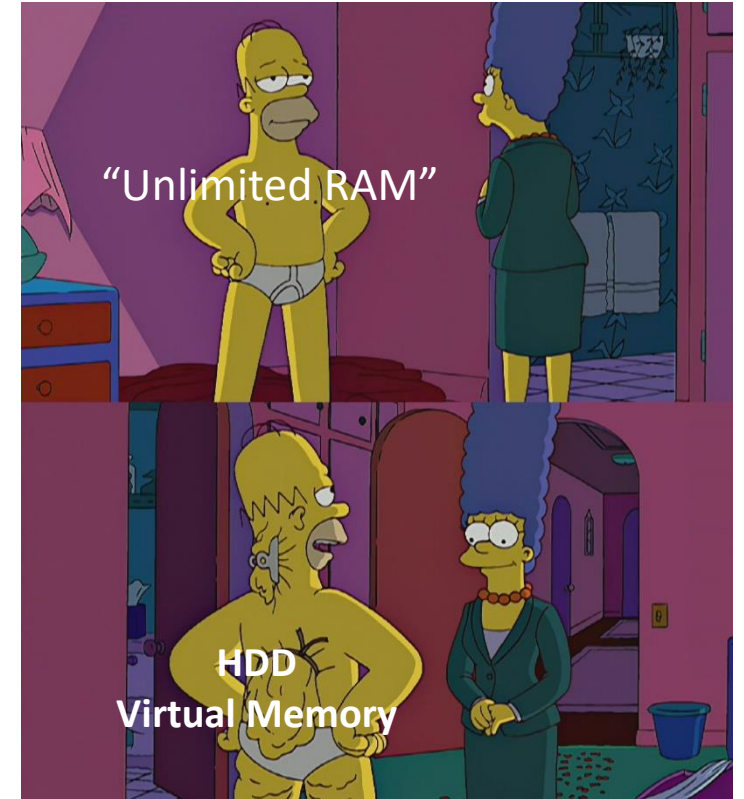


CPU



Illusion Manager

- Gives applications the illusion that they have infinite storage and resources



The jobs of an Operating System

1. Process Manager
“The Coach”

2. Interface Manager
“The Bouncer”

3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”

