

CSCI 466: Networks

FTP, Peer 2 Peer Networks (P2P), Content Distribution Networks (CDN)


Reese Pearsall
Fall 2024

Announcements

Wireshark Lab 1 due **tonight** at 11:59 PM

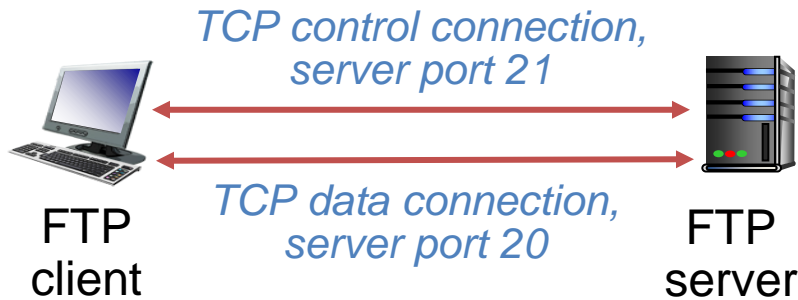
- Everyone needs to submit something to D2L, even if you work with partners

i hope you have a good weekend



FTP

File Transfer Protocol (FTP)- protocol used for transferring files from server to client



- FTP communicates over two connections
 - Port 21 for control information
 - Port 20 for data
- Differences from HTTP
 - Control communication “out-of-band”
 - Server maintains per client state: authentication, current directory

- **FTP procedure:**


1. FTP client contacts FTP server at port 21, using TCP
2. Client authorized over control connection
3. Client browses remote directory, sends commands over control connection
4. When server receives file transfer command, server opens 2nd TCP data connection (for file) to client
5. After transferring one file, server closes data connection

Why use a separate control connection?

WinSCP

P2P Networks


- No always-on server
- Arbitrary end systems directly communicate
- Peers are intermittently connected and change IP addresses



Time to distribute a file of size F to N clients

$$D_{P2P} = \max \left\{ \frac{F}{u_s}, \frac{F}{d_{\min}}, \frac{NF}{u_s + \sum_{i=1}^N u_i} \right\}$$

Client-Server Architecture




- Server can upload data at rate u_s
- Clients download data at rates d_1, d_2, \dots, d_N

$$D_{CS} = \max \left\{ \frac{NF}{u_s}, \frac{F}{d_{\min}} \right\}$$


Time to distribute a file of size F to N clients

P2P Networks

Client-Server Architecture



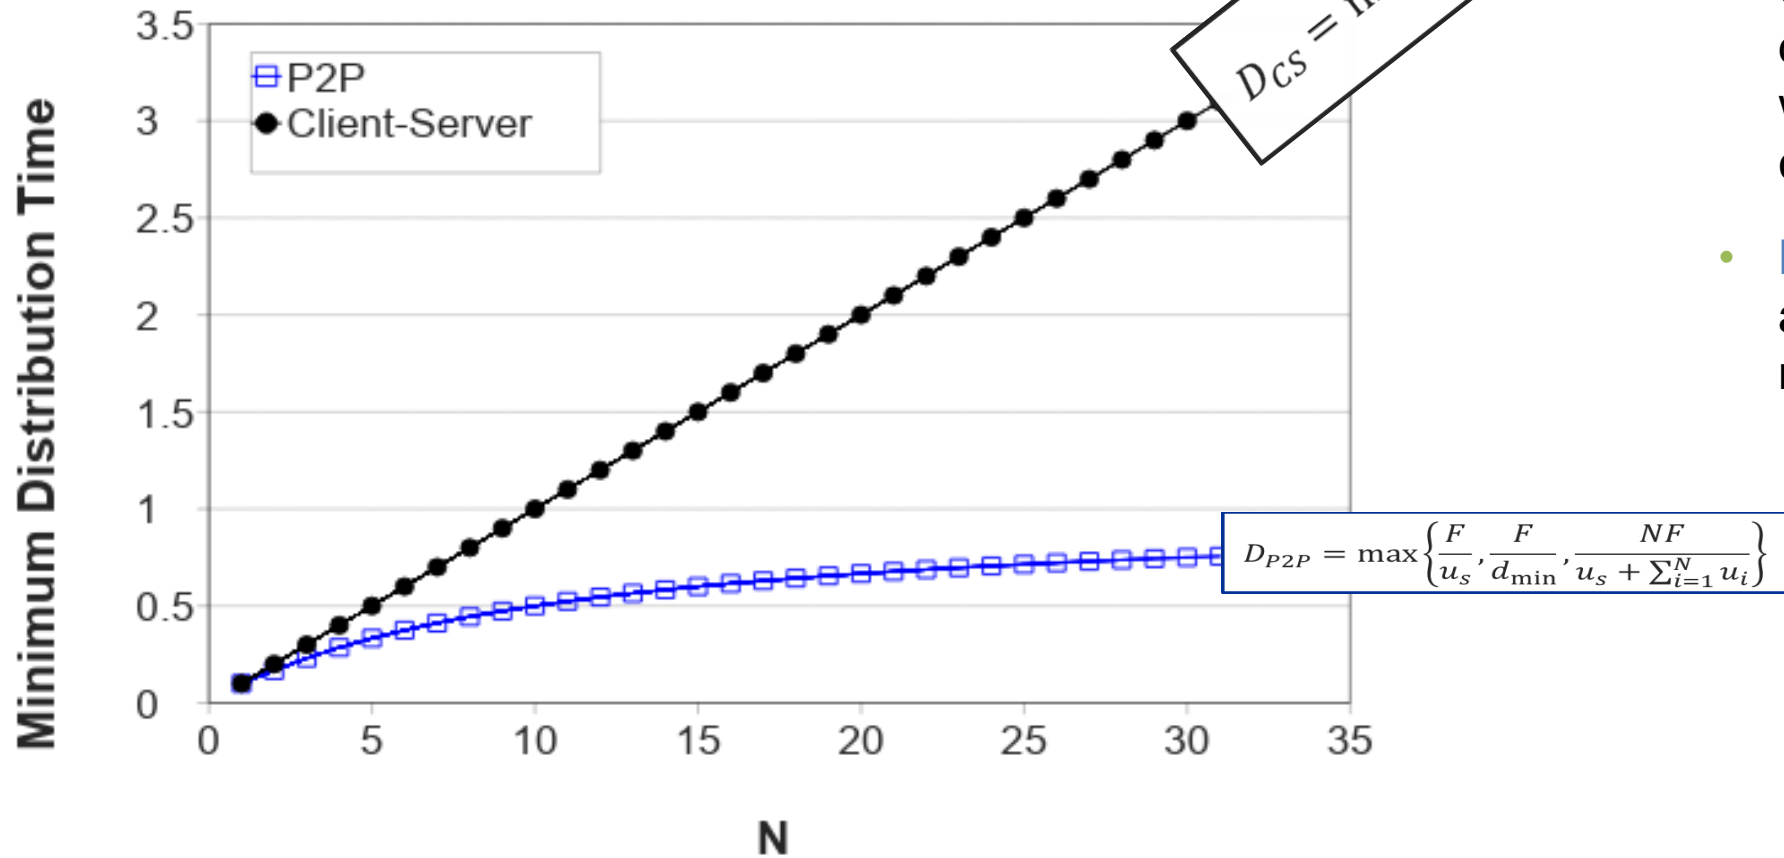
Existing clients have to share resources with new users



New peers are both a client and a server. There will not be a negative impact on current peers

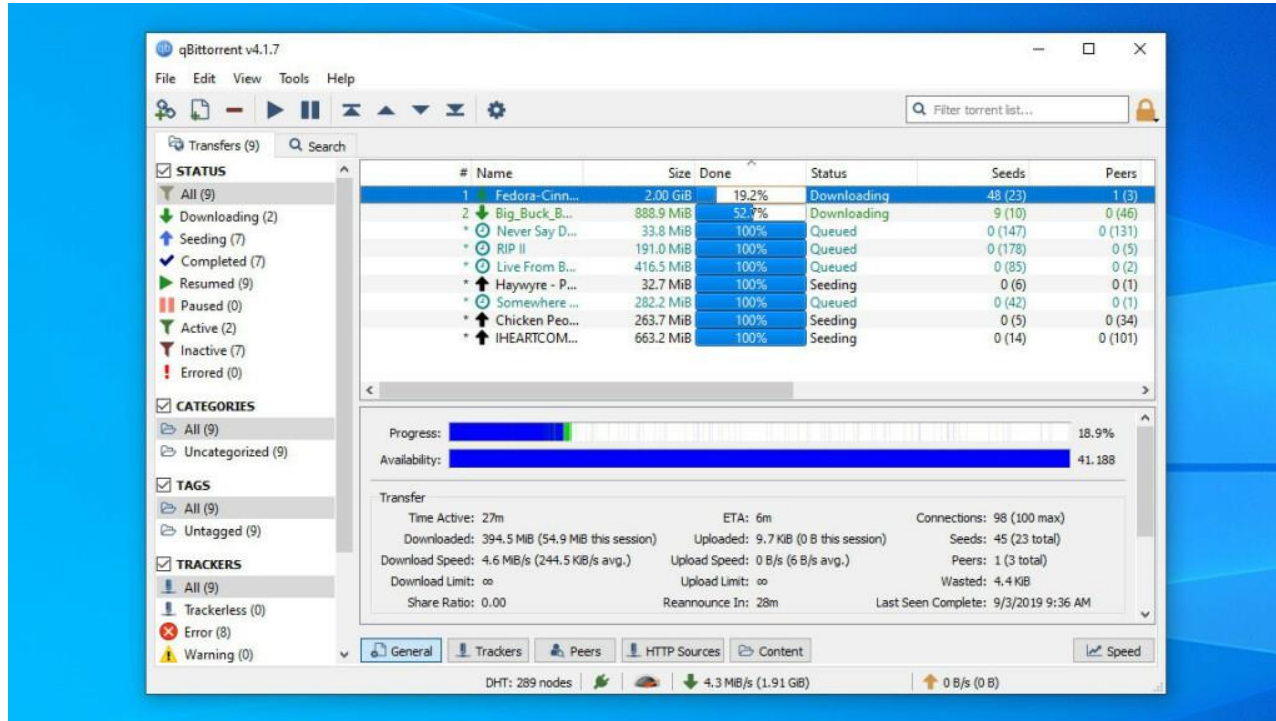
P2P architectures are *self-scaling*

P2P Networks



- **Client server** distribution time grows with the number of clients
- **P2P** distribution time approaches 1 hour as number of clients grows

P2P file distribution: BitTorrent

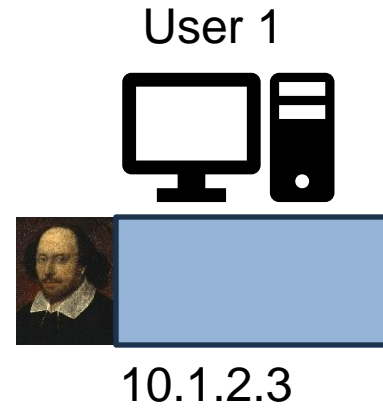


Service for sharing files over the internet in a decentralized fashion

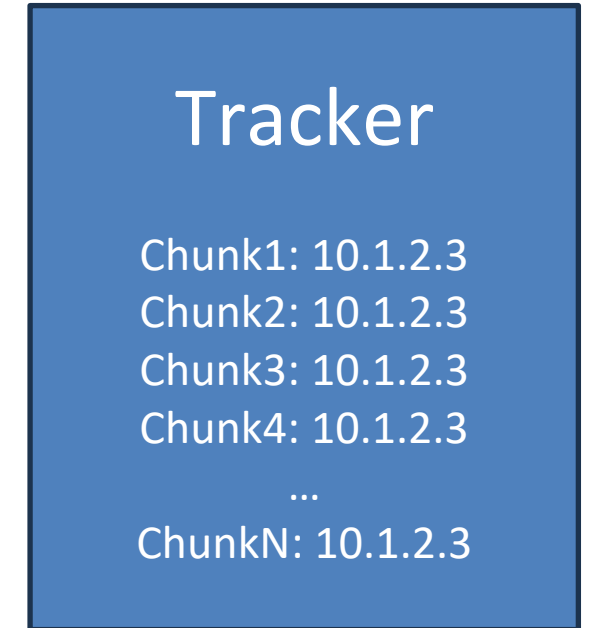
(BitTorrent, uTorrent, qBittorrent are all popular torrent clients)

User 1


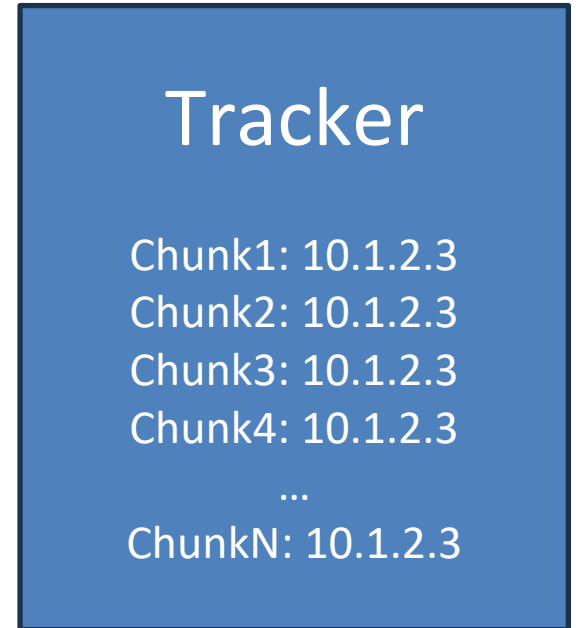
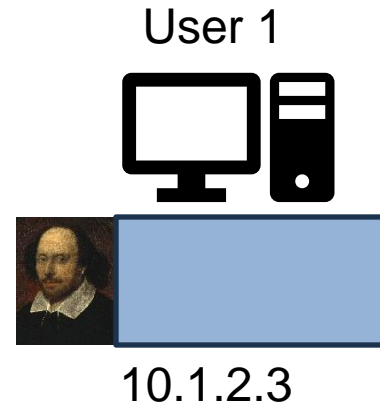




A **Tracker** is made, typically located at the person who created the torrent




A **Tracker** is made, typically located at the person who created the torrent



User 2 wants the file, so they download a **TorrentFile** which allows them to contact the Tracker and join the P2P

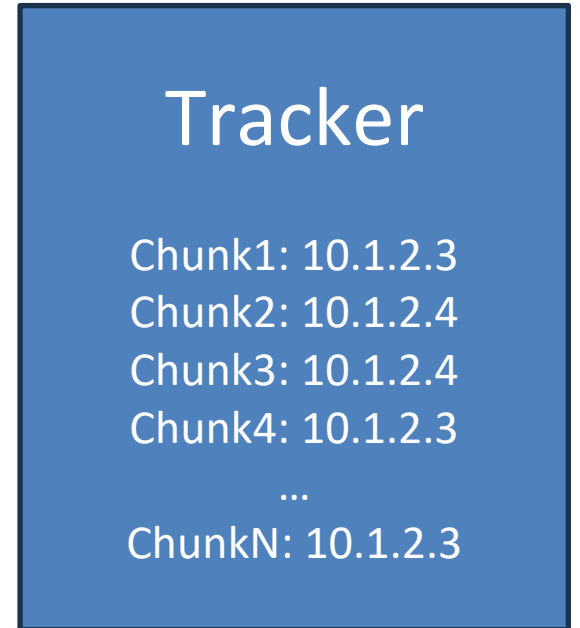

A **Tracker** is made, typically located at the person who created the torrent



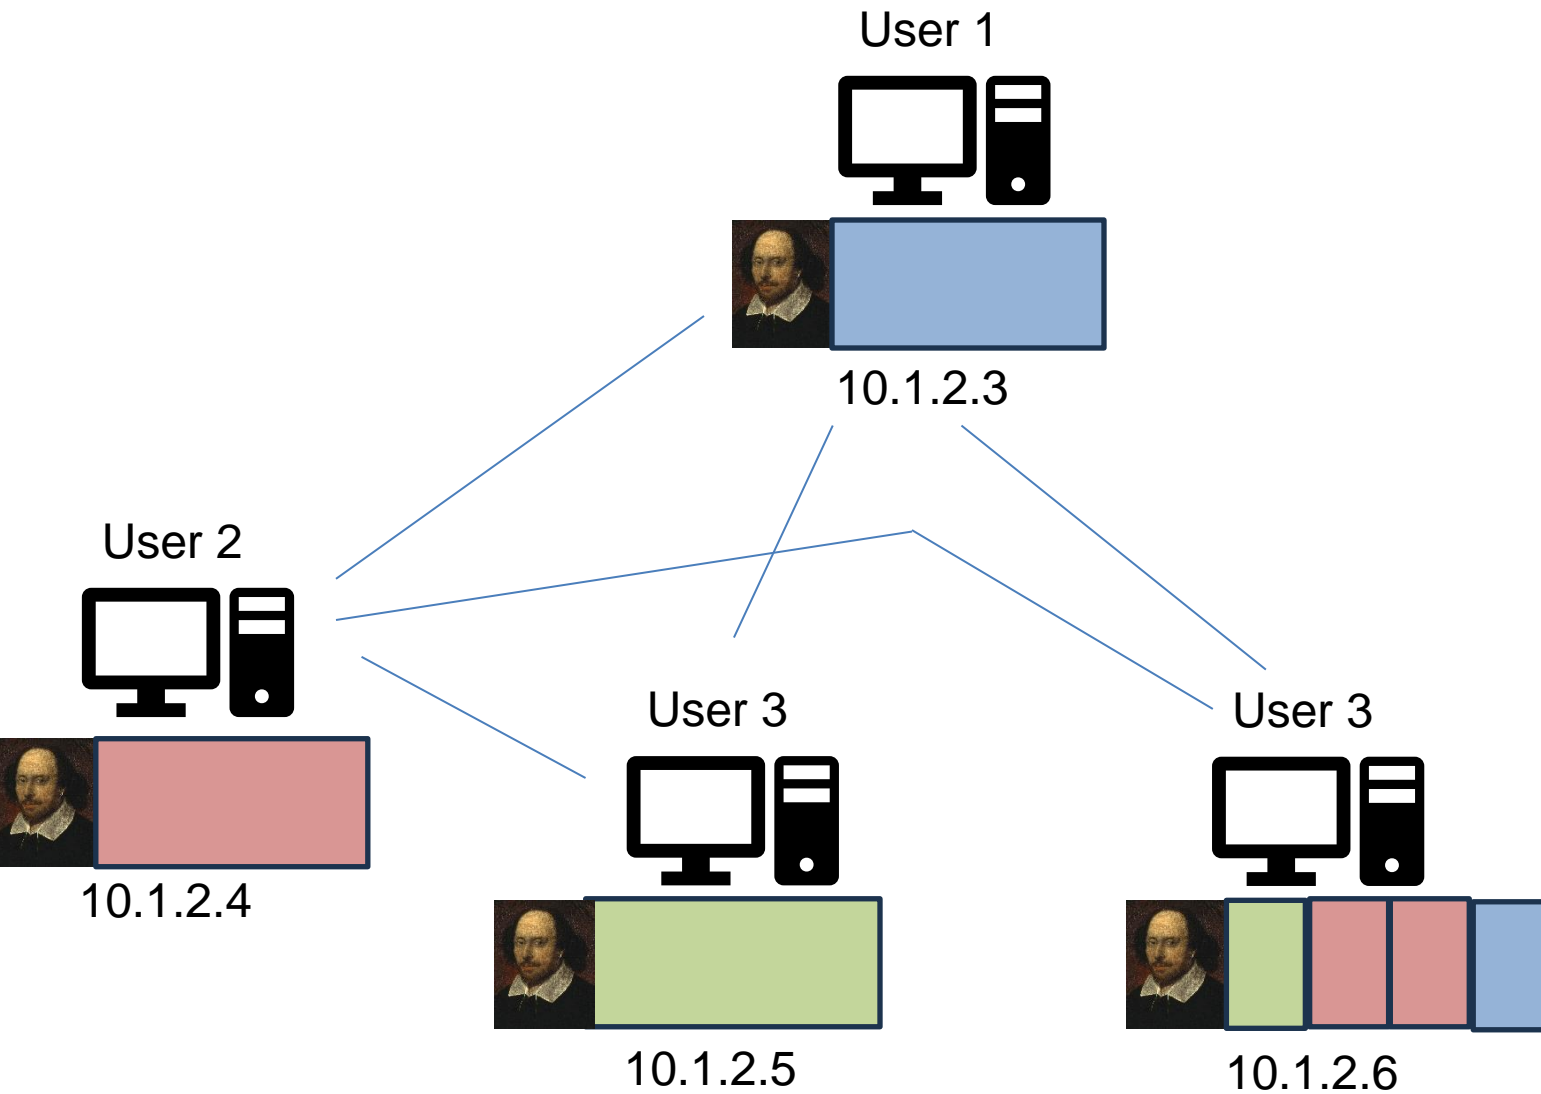
Tracker

Chunk1: 10.1.2.3
Chunk2: 10.1.2.4
Chunk3: 10.1.2.4
Chunk4: 10.1.2.3
...
ChunkN: 10.1.2.3

A **Tracker** is made, typically located at the person who created the torrent



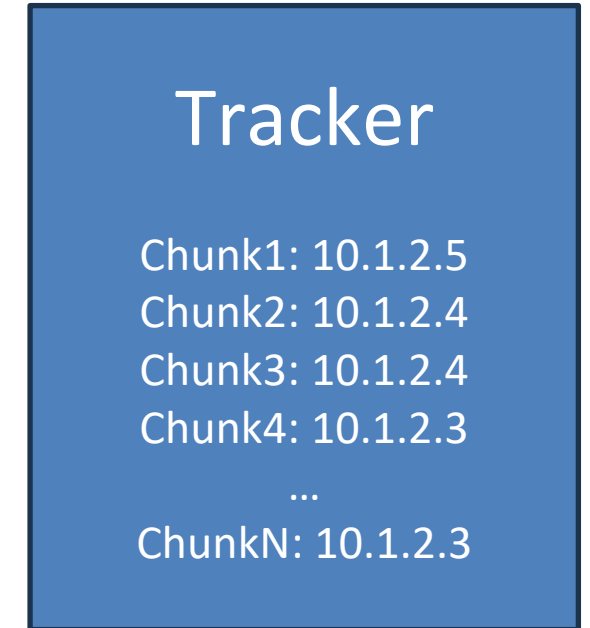
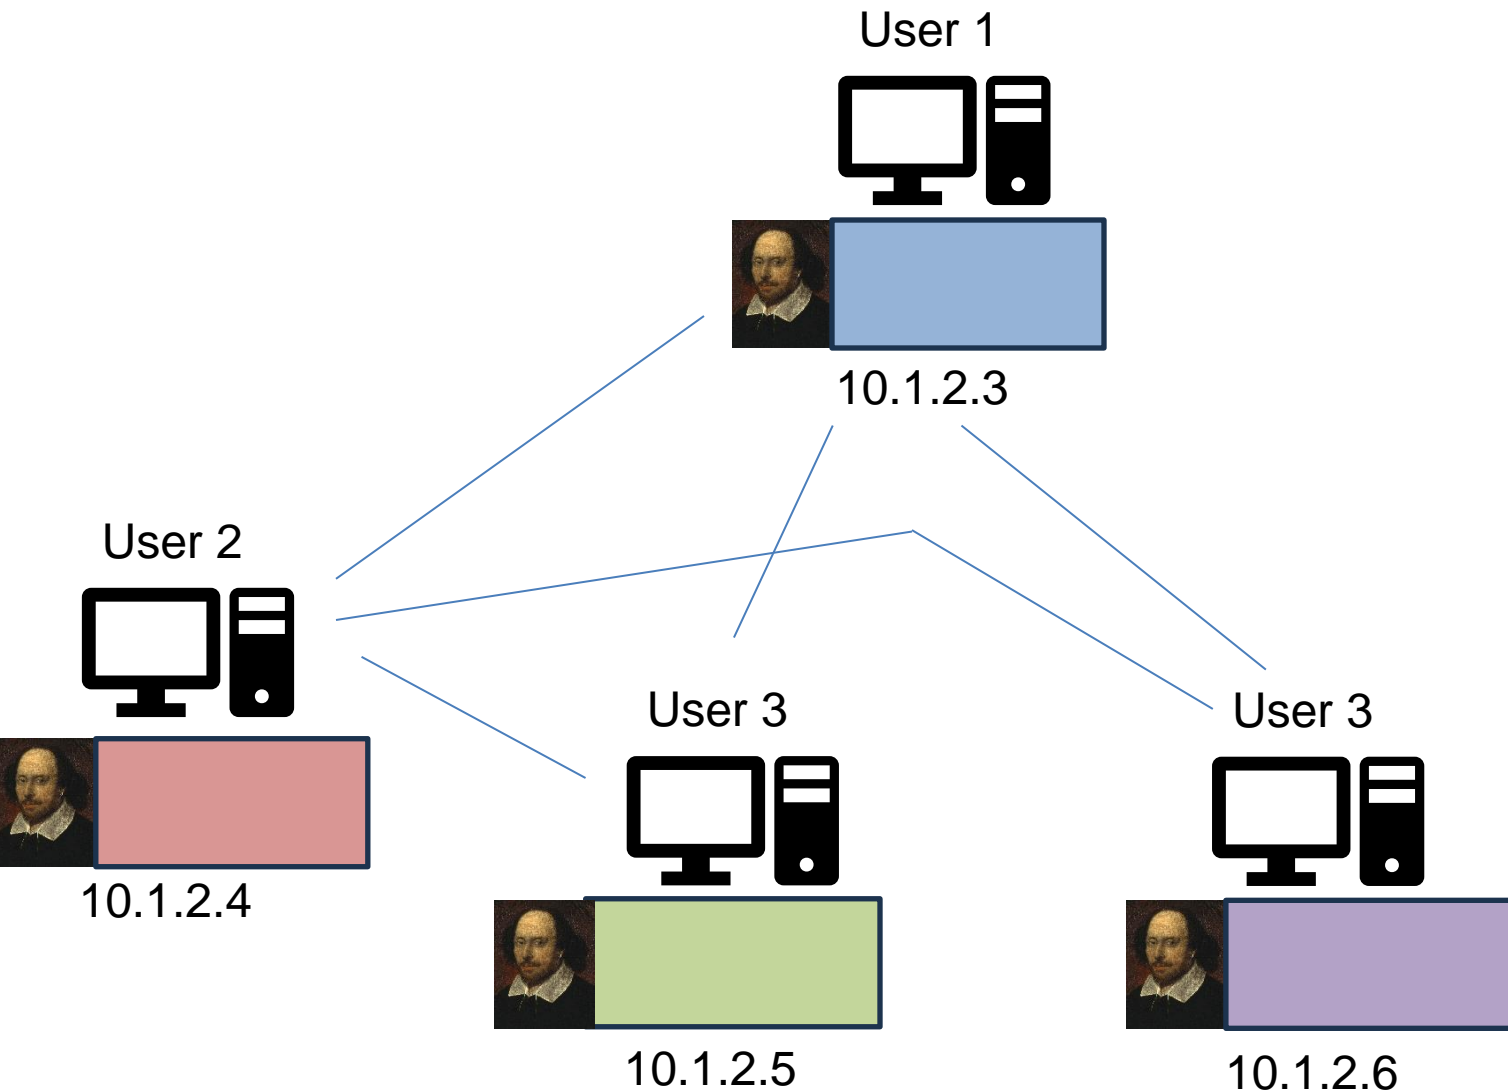
A **Tracker** is made, typically located at the person who created the torrent



Tracker

Chunk1: 10.1.2.5
Chunk2: 10.1.2.4
Chunk3: 10.1.2.4
Chunk4: 10.1.2.3
...
ChunkN: 10.1.2.3

A **Tracker** is made, typically located at the person who created the torrent




As more users join the network, it grows the capability of the networks and chunks are more scattered

P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file




P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent

torrent: group of peers exchanging chunks of a file




peer joining torrent:

- has no chunks, but will accumulate them over time from other peers
- registers with tracker to get list of peers, connects to subset of peers (“neighbors”)

P2P file distribution: BitTorrent

- Files are divided into chunks
- Peers in torrent send/receive file chunks

tracker: tracks peers participating in torrent


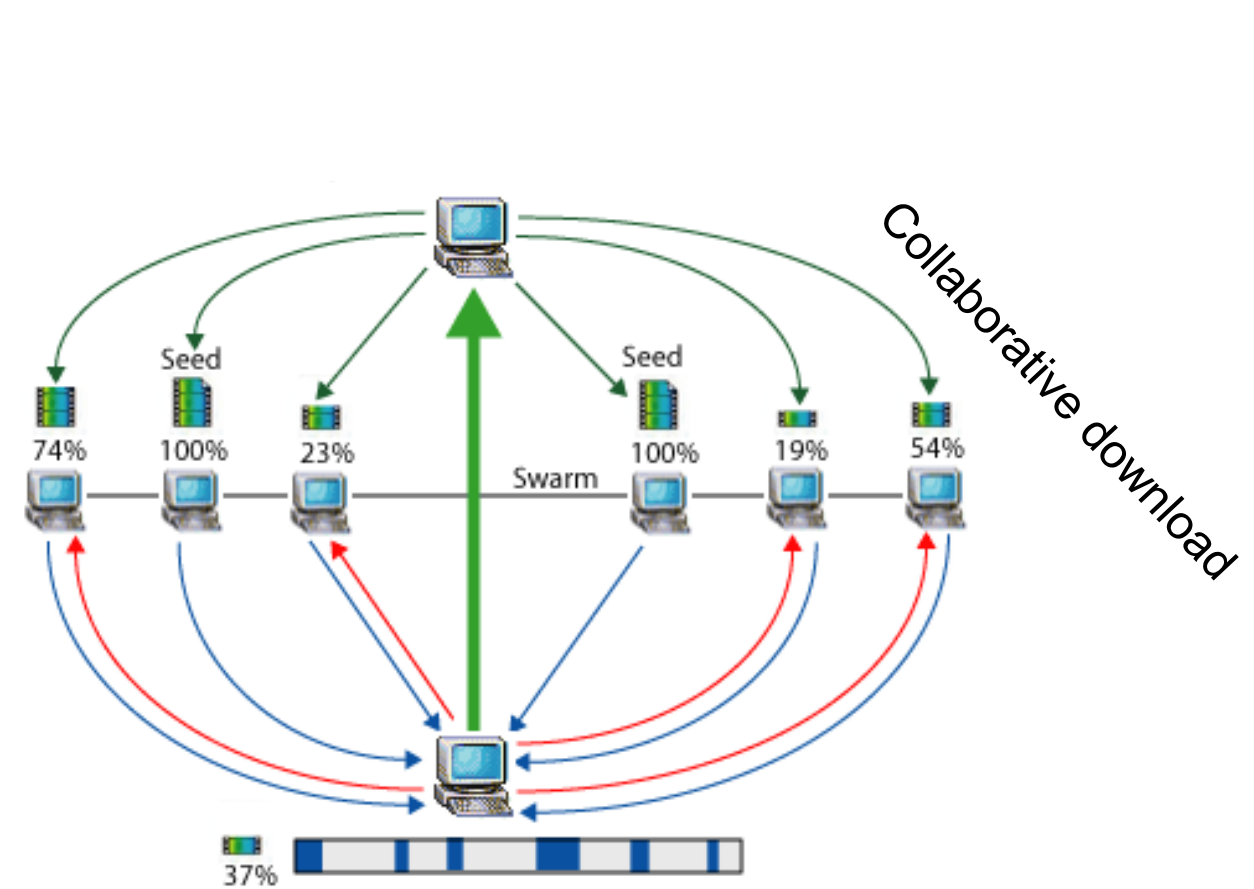


Alice arrives ...
... obtains list
of peers from tracker
... and begins exchanging
file chunks with peers in torrent


torrent: group of peers exchanging chunks of a file

- while downloading, peer uploads chunks to other peers
- peer may change peers with whom it exchanges chunks
- **churn:** peers may come and go
- once peer has entire file, it may (selfishly) leave or (altruistically) remain in torrent

P2P file distribution: BitTorrent




P2P file distribution: BitTorrent




Which chunks to download?

- Rarest first
- Try to equalize the number of copies of each chunk in the system




P2P file distribution: BitTorrent



From which peers to download?


- Allow download to N peers with highest upload rates
- Allows peers with similar upload rates to find each other




P2P file distribution: BitTorrent

requesting chunks:

- at any given time, different peers have different subsets of file chunks
- periodically, Alice asks each peer for list of chunks that they have
- Alice requests missing chunks from peers, rarest first



P2P file distribution: BitTorrent



higher upload rate: find better trading partners, get file faster !

sending chunks: tit-for-tat

- Alice sends chunks to those four peers currently sending her chunks *at highest rate*
 - other peers are *choked* by Alice (do not receive chunks from her)
 - re-evaluate top 4 every 10 secs
- every 30 secs: randomly select another peer, starts sending chunks
 - “optimistically unchoke” this peer
 - newly chosen peer may join top 4

BitTorrent is referred to as an **unstructured P2P**

Security Concerns?

- IP Address
- Malware




CDNs

- video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
 - ~1B YouTube users, ~75M Netflix users
- challenge: scale - how to reach ~1B users?
 - single mega-video server won't work (why?)
- challenge: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- *solution*: distributed, application-level infrastructure



CDNs


- Caching
 - Save previously delivered data
 - Subsequent requests served from cache on the browser, or in the access network
- Applications
 - Reduce response time for client request
 - Reduce ISP traffic costs
- Content distribution networks
 - Distributed caches
 - Web objects addressed to CDN server
 - CDN server fetches from content provider on first access



- Conditional GET
 - Cache: specify date of cached copy in HTTP request
If-modified-since: <date>
 - Server: response contains no object if cached copy is up-to-date:
HTTP/1.0 304 Not Modified

CDNs

- Challenge:
 - How to stream content to millions of users?
- Option 1:
 - Single mega-datacenter
 - Pros: Simple
 - Cons:
 - Single point of failure
 - Point of network congestion
 - Long path to distant clients
 - Multiple copies of video sent over outgoing link



CDNs

- Option 2:
 - Store/serve multiple copies of videos at multiple geographically distributed sites


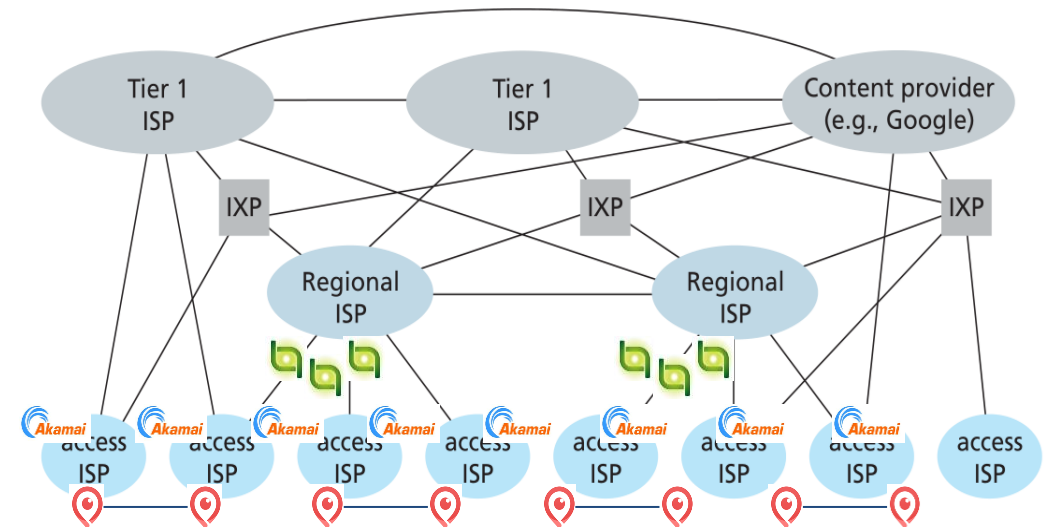
A **content delivery network (CDN)** refers to a geographically distributed group of cache servers which work together to provide fast delivery of Internet content. **(not a web host)**



34 DNS lookups

204 HTTP requests

520 KB of data downloaded



56% of domains resolve to a CDN


CDNs



[https://**cdn**.discordapp.com/attachments/1017630358621143110/1017849417216294972/unknown.png](https://cdn.discordapp.com/attachments/1017630358621143110/1017849417216294972/unknown.png)

CDNs

- Challenge: how does CDN DNS select “good” CDN node to stream to client
 - Pick CDN node **geographically** closest to client’s local DNS
 - Pick CDN node with **shortest delay** (or min # hops) to client (CDN nodes periodically ping access ISPs, reporting results to CDN DNS)



CDNs

In HTTP streaming, video is stored at an HTTP server and retrieved with a GET request

Frames are sent to a client buffered and played back after a certain threshold

All clients receiving the same encoding of the video despite widely different bandwidth

CDNs

In HTTP steaming, video is stored at an HTTP server and retrieved with a GET request

Frames are sent to a client buffered and played back after a certain threshold

All clients receiving the same encoding of the video despite widely different bandwidth

Dynamic Adaptive Streaming over HTTP (DASH)- video is encoded into several different versions


- Different encoding rate, different quality, etc

Clients dynamically request chunks of video segments every few seconds via GET requests

- Is the current bandwidth good? → Retrieve good quality
- Is the current bandwidth bad? → Retrieve ok quality

CDNs


When streaming, you are consistently issuing GET requests



#capybara #capivara

Capybara eating half a watermelon Full Video

2,145,989 61K DISLIKE SHARE DOWNLOAD CLIP SAVE ...

 **Crazy Cody's Creatures**
110K subscribers

[SUBSCRIBE](#)

Help Support my Channel by becoming a Patron <https://www.patreon.com/CrazyCody> Ever wonder what its like to live with a capybara???

Ask your questions below, I read them all!

Elements Console Sources **Network** Performance Memory Application Security Lighthouse


Filter ☐ Invert ☐ Hide data URLs **All** Fetch/XHR JS CSS Img Media Font Doc WS Wasm Manifest Other ☐ Has blocked cookies

☐ Blocked Requests ☐ 3rd-party requests

Name	Status	Type	Initiator	Size	Time	Waterfall
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.5 MB	48 ms	
watchtime?ns=yt&el=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	30 ms	
qoe?fmt=399&afmt=251&cpn=D1QCo4_7KrXNdhs4...	204	xhr	VM685:201	39 B	28 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	256 kB	20 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	2.1 MB	50 ms	
log_event?alt=json&key=AlzaSyAO_FJ2SlqU8Q4STEh...	200	xhr	VM685:201	77 B	46 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.4 MB	56 ms	
watchtime?ns=yt&el=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	31 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	2.1 MB	55 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	391 kB	23 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.6 MB	32 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.3 MB	53 ms	
watchtime?ns=yt&el=detailpage&cpn=D1QCo4_7KrX...	204	xhr	VM685:201	35 B	32 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	2.1 MB	60 ms	
videoplayback?expire=1663023840&ei=gGYfY7iCCoO...	200	fetch	VM685:223	1.6 MB	65 ms	
data:image/png;base...	200	png	Other	(memory cache)	0 ms	


18 requests 14.8 MB transferred 14.8 MB resources

CDNs



Studio master versions are uploaded to a private Amazon cloud


CDNs



Studio master versions are uploaded to a private Amazon cloud

Videos are processed into many different formats, allowing for DASH

CDNs



Studio master versions are uploaded to a private Amazon cloud

Videos are processed into many different formats, allowing for DASH

Versions are uploaded to Netflix's CDNs