

# CSCI 466: Networks

Network Layer – Routing (Control Plane)

Reese Pearsall  
Fall 2024

# Announcements

## No class on Friday

- **Workday for PA2 (I'll be in my office still if you need help)**

PA2 due on Sunday at 11:59 PM

- Make sure everything is inside of a /PA2 folder in your repo
- Everyone needs to submit video demo link to D2L


ACM is hosting a guest speaker tonight at 5PM in Barnard Hall 347

- Speaker has experience working at HP, Microsoft, Amazon
- Great opportunity to learn about industry, careers, etc

[https://www.bozemandailychronicle.com/news/montana\\_state\\_university/exponent/kind-of-scary-montana-state-officials-confirm-learning-platform-tracks-student-locations/article\\_740f212e-859d-11ef-a536-771135b8a5d8.html](https://www.bozemandailychronicle.com/news/montana_state_university/exponent/kind-of-scary-montana-state-officials-confirm-learning-platform-tracks-student-locations/article_740f212e-859d-11ef-a536-771135b8a5d8.html)

# 'Kind of scary': Montana State officials confirm learning platform tracks student locations

Dom Lucero The MSU Exponent 22 hrs ago



A student accesses DZL while at a coffee shop.  
Maddi Hohner/The MSU Exponent

Buy Now

Quizzes > Quiz 3 > Attempt Logs

## Attempt Logs


Overview

Detailed

2471 items in the list.

Attempt	Event	Modified by	IP Address	Date ▲
Reese Pearsall (Attempt: 1)	Quiz Entry	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:13 PM
Reese Pearsall (Attempt: 1)	Response to Question 1 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:13 PM
Reese Pearsall (Attempt: 1)	Response to Question 2 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 3 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 4 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 5 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 6 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 7 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 8 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 9 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM
Reese Pearsall (Attempt: 1)	Response to Question 10 Saved	Reese Pearsall	153.90.118.85	Oct 3, 2024 3:14 PM

**Forwarding** refers to moving packets from a **router's input** to appropriate **router output**, and is implemented in the data plane.




Address range	Interface (output link)
128.11.52.0 – 128.11.52.255	1
153.90.2.0 – 153.90.2.255	2
153.90.2.87 – 153.90.2.89	3



Ideally, this output links are the most optimal path to get to the destination

**Routing** refers to determining the route taken by packets from **source** to **destination**, and is implemented in the control plane.



What is the best way to get from **u** to **z**?

### Routing Metrics:

- Shortest Path
- Highest Throughput Path
- *Minimum Number of Hops*
- Lowest Congested Path

# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms


# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms

**Centralized/Global**- we know the edge costs of the network

**Link State** algorithms


(Dijkstra's Algorithm)




We can compute the shortest path from one node, to all other nodes in polynomial time.

Once we know the shortest path from A to B, we can update routing tables to reflect that shortest path


## Shortest Path



## Shortest Path









Path with the smallest sum of edge weights.

What is the shortest path between **vertex 0** and **vertex 6**?




We are going to find the shortest path between vertex 0 and every other vertex, flooding out from 0.

Shortest Path



Distance from 0


0	?
1	?
2	?
3	?
4	?
5	?
6	?
7	?



Distance from 0

0	0
1	$\infty$
2	$\infty$
3	$\infty$
4	$\infty$
5	$\infty$
6	$\infty$
7	$\infty$

How can we keep track of routes?




Distance from 0

0	0
1	$\infty$
2	$\infty$
3	$\infty$
4	$\infty$
5	$\infty$
6	$\infty$
7	$\infty$

Previous vertex

0	-
1	
2	
3	
4	
5	
6	
7	

How can we keep track of routes?




Distance from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	$\infty$
5	$\infty$
6	1.51
7	0.60

Previous vertex

0	-
1	
2	0
3	7
4	
5	
6	3
7	2

How can we keep track of routes?


Distance  
from 0

0	0
1	$\infty$
2	$\infty$
3	$\infty$
4	$\infty$
5	$\infty$
6	$\infty$
7	$\infty$

Previous  
vertex

0	-
1	
2	0
3	7
4	
5	
6	3
7	2

If this is the shortest path from 0 to 6, what can we say about the shortest path from 0 to 3?

Distance  
from 0

0	0
1	$\infty$
2	$\infty$
3	$\infty$
4	$\infty$
5	$\infty$
6	$\infty$
7	$\infty$


Previous  
vertex

0	-
1	
2	
3	
4	
5	
6	
7	

Priority  
queuevertex  
(distance)

What can we reach from connected vertices and at what distance (from 0)?



Distance  
from 0

0	0
1	$\infty$
2	$\infty$
3	$\infty$
4	$\infty$
5	$\infty$
6	$\infty$
7	$\infty$

Previous  
vertex


0	-
1	
2	
3	
4	
5	
6	
7	

Priority  
queue

2 (0.26)
4 (0.38)

vertex  
(distance)

What can we reach from connected vertices and at what distance (from 0)?

Distance  
from 0

0	0
1	$\infty$
2	$\infty$
3	$\infty$
4	$\infty$
5	$\infty$
6	$\infty$
7	$\infty$

Previous  
vertex


0	-
1	
2	0
3	
4	0
5	
6	
7	

Priority  
queue

2 (0.26)
4 (0.38)

vertex  
(distance)

What can we reach from connected vertices and at what distance (from 0)?

Distance  
from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	$\infty$
6	$\infty$
7	$\infty$

Previous  
vertex

0	-
1	
2	0
3	
4	0
5	
6	
7	


Priority  
queue

2 (0.26)
4 (0.38)

vertex  
(distance)

What can we reach from connected vertices and at what distance (from 0)?

queue  
top = 2 (0.26)



Distance  
from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	$\infty$
6	$\infty$
7	$\infty$

Previous  
vertex

0	-
1	
2	0
3	
4	0
5	
6	
7	


Priority  
queue

4 (0.38)

vertex  
(distance)

What can we reach from connected vertices and at what distance (from 0)?

queue  
top = 2 (0.26)



Distance  
from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	$\infty$
6	$\infty$
7	$\infty$

Previous  
vertex

0	-
1	
2	0
3	
4	0
5	
6	
7	


Priority  
queue

4 (0.38)

vertex  
(distance)

What can we reach from connected vertices and at what distance (from 0)?

queue  
top = 2 (0.26)



Distance  
from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	$\infty$
6	$\infty$
7	0.60

Previous  
vertex

0	-
1	
2	0
3	
4	0
5	
6	
7	2

Priority  
queue


4 (0.38)
7 (0.60)

vertex  
(distance)

What can we reach from connected vertices and at what distance (from 0)?

Shortest Path

queue  
top = 4 (0.38)



Distance from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	$\infty$
6	$\infty$
7	0.60

Previous vertex


0	-
1	
2	0
3	
4	0
5	
6	
7	2

Priority queue

7 (0.60)

vertex  
(distance)

queue  
top = 4 (0.38)



Distance from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	$\infty$
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	
4	0
5	
6	
7	2

Priority queue

7 (0.60)


vertex  
(distance)

Add neighbors to queue/previous.



Shortest Path

queue  
top = 4 (0.38)



Add neighbors to queue/previous.

Distance from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous vertex


0	-
1	
2	0
3	
4	0
5	4
6	
7	2

Priority queue

7 (0.60)
5 (0.73)

vertex  
(distance)

queue  
top = 4 (0.38)



Add neighbors to queue/previous.

**We have another route to 7!**

Distance  
from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous  
vertex


0	-
1	
2	0
3	
4	0
5	4
6	
7	2

Priority  
queue

7 (0.60)
5 (0.73)

vertex  
(distance)

queue  
top = 4 (0.38)



Distance from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	
4	0
5	4
6	
7	2

Priority queue


7 (0.60)
5 (0.73)

vertex  
(distance)

Add neighbors to queue/previous.

**We have another route to 7! Check to see if it is shorter!**

queue  
top = 4 (0.38)



Distance from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	
4	0
5	4
6	
7	2

Priority queue


7 (0.60)
5 (0.73)

vertex  
(distance)

Add neighbors to queue/previous.

We have another route to 7! Check to see if it is shorter! It's not ( $0.38 + 0.37 = 0.75 > 0.60$ ).

queue  
top = 4 (0.38)

Distance  
from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous  
vertex

0	-
1	
2	0
3	
4	0
5	4
6	
7	2

Priority  
queue

7 (0.60)
5 (0.73)


vertex  
(distance)

**Rule:** When processing vertex  $v$ , only add/modify queue for neighbor  $u$  if and only if:  
 $\text{distance}[v] + \text{weight}(v, u) < \text{distance}[u]$



Shortest Path

queue  
top =



Repeat.

Distance from 0

0	0
1	$\infty$
2	0.26
3	$\infty$
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous vertex


0	-
1	
2	0
3	
4	0
5	4
6	
7	2

Priority queue

7 (0.60)
5 (0.73)

vertex  
(distance)

queue  
top = 7 (0.60)



Repeat.

Distance from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous vertex


0	-
1	
2	0
3	7
4	0
5	4
6	
7	2

Priority queue

5 (0.73)
3 (0.99)

vertex  
(distance)

queue  
top = 7 (0.60)



Repeat.

Distance from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	7
4	0
5	4
6	
7	2

Priority queue


5 (0.73)
3 (0.99)

vertex  
(distance)

We have another route to 5, and at cost 0.7 <



queue  
top = 7 (0.60)



Distance from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	0.38
5	0.73
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	7
4	0
5	4
6	
7	2


Priority queue

5 (0.73)
3 (0.99)

vertex  
(distance)

Repeat. **We have another route to 5, and at cost  $0.7 < 0.73$ .**  
i.e.,  $\text{distance}[v] + \text{weight}(v, u) < \text{distance}[u]$

queue  
top = 7 (0.60)



Distance from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	0.38
5	<del>0.73</del>
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	7
4	0
5	4
6	
7	2


Priority queue

**0.70**  
5 (0.73)  
3 (0.99)

**0.70**vertex  
(distance)

Repeat. **We have another route to 5, and at cost  $0.7 < 0.73$ .  
So updated queue/previous/distance.**

queue  
top = 7 (0.60)

Distance  
from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous  
vertex

0	-
1	
2	0
3	7
4	0
5	7
6	
7	2

Priority  
queue


5 (0.70)
3 (0.99)

vertex  
(distance)

Repeat. We have another route to 5, and at cost  $0.7 < 0.73$ .  
So updated queue/previous/distance.

Shortest Path

queue  
top = 7 (0.60)



Repeat.

Distance from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	7
4	0
5	7
6	
7	2


Priority queue

5 (0.70)
3 (0.99)

vertex  
(distance)

Shortest Path

queue  
top = 5 (0.70)



Repeat.

Distance from 0

0	0
1	$\infty$
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous vertex

0	-
1	
2	0
3	7
4	0
5	7
6	
7	2


Priority queue

3 (0.99)

vertex  
(distance)

Shortest Path

queue  
top = 5 (0.70)




Repeat.

	Distance from 0	Previous vertex	Priority queue
0	0	0	-
1	$\infty$	1	
2	0.26	2	0
3	0.99	3	7
4	0.38	4	0
5	0.70	5	7
6	$\infty$	6	
7	0.60	7	2

vertex  
(distance)

Shortest Path

queue  
top = 5 (0.70)



Repeat.

Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous vertex


0	-
1	5
2	0
3	7
4	0
5	7
6	
7	2

Priority queue

1 (0.92)
3 (0.99)

vertex  
(distance)

queue  
top = 5 (0.70)



Repeat.

What about neighbor 4?

Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	
7	2


Priority queue

1 (0.92)
3 (0.99)

vertex  
(distance)



queue  
top = 5 (0.70)



Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	
7	2

Priority queue

1 (0.92)
3 (0.99)


vertex  
(distance)

Repeat.

What about neighbor 4?  $\text{distance}[5] + \text{weight}(5, 4) = 0.70 + 0.35 = 1.05 \nless 0.38 = \text{distance}[4]$

# Shortest Path

queue  
top = 5 (0.70)



Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	
7	2

Priority queue

1 (0.92)
3 (0.99)


vertex  
(distance)

Repeat. What about neighbor 7?

$$\text{distance}[5] + \text{weight}(5, 7) = 0.70 + 0.28 = 0.98 \not\leq 0.60 = \text{distance}[7]$$

Shortest Path

queue  
top = 1 (0.92)



Repeat.

Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous vertex


0	-
1	5
2	0
3	7
4	0
5	7
6	
7	2

Priority queue

3 (0.99)

vertex  
(distance)

queue  
top = 1 (0.92)



Repeat.

What about neighbor 3?  $0.92 + 0.29 = 1.21 > 0.99$

Distance  
from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	$\infty$
7	0.60

Previous  
vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	
7	2


Priority  
queue

3 (0.99)
----------

vertex  
(distance)

Shortest Path

queue  
top = 3 (0.99)




Repeat.

	Distance from 0	Previous vertex	Priority queue
0	0	-	
1	0.92	5	
2	0.26	0	
3	0.99	7	
4	0.38	0	
5	0.70	7	
6	$\infty$		
7	0.60	2	

vertex  
(distance)

Shortest Path

queue  
top = 3 (0.99)



Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60

Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2


Priority queue

6 (1.51)
----------

vertex  
(distance)

Shortest Path

queue  
top = 6 (1.51)



Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60

Previous vertex


0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2

Priority queue

vertex  
(distance)

Shortest Path

queue  
top = 6 (1.51)



Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60

Previous vertex


0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2

Priority queue

vertex  
(distance)



queue  
top = 6 (1.51)



Repeat?

Neighbor 4?  $1.51 + 0.93 > 0.38$

Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60


Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2

Priority queue

vertex  
(distance)

queue  
top = 6 (1.51)



Repeat?

Neighbor 0?  $1.51 + 0.58 > 0$

Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60


Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2

Priority queue

vertex  
(distance)

queue  
top = 6 (1.51)



Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60

Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2


Priority queue

vertex  
(distance)

Repeat?

Neighbor 2?  $1.51 + 0.4 > 0.26$

queue  
top = 6 (1.51)



When are we done?

Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60


Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2

Priority queue

vertex  
(distance)

queue  
top = 6 (1.51)



Distance from 0

0	0
1	0.92
2	0.26
3	0.99
4	0.38
5	0.70
6	1.51
7	0.60

Previous vertex

0	-
1	5
2	0
3	7
4	0
5	7
6	3
7	2

Priority queue

vertex  
(distance)

When are we done?

When the queue is empty!


# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms

**Centralized/Global**- we know the edge costs of the network

**Link State** algorithms

(Dijkstra's Algorithm)




We can compute the shortest path from one node, to all other nodes in polynomial time.

Once we know the shortest path from A to B, we can update routing tables to reflect that shortest path

# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms



**Decentralized-** we do not know the edge costs of the entire network.

Only know edge costs to neighbors

**Distance Vector** algorithms

# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms

Based on *Bellman-Ford* (BF) equation (dynamic programming):

Bellman-Ford equation

Let  $D_x(y)$ : cost of least-cost path from  $x$  to  $y$ .

Then:

$$D_x(y) = \min_v \{ c_{x,v} + D_v(y) \}$$

$v$ 's estimated least-cost-path cost to  $y$

$\min$  taken over all neighbors  $v$  of  $x$


direct cost of link from  $x$  to  $v$



<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y			
	Z			

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y	2	0	1
	Z			

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y			
	Z	7	1	0




<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y			
	Z			

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y	2	0	1
	Z			

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y			
	Z	7	1	0

<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	7	1	0



X learns that Y can reach Z with only a cost of 1

Meaning:

The cost to get to Y from X (2) plus the cost to get from Y to Z (1) is lower than my direct path to Y (7), so I have now learned a new shortest path!

<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y			
	Z			


<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	7	1	0

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y	2	0	1
	Z			

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y	2	0	1
	Z	7	1	0

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y			
	Z	7	1	0


<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y	2	0	1
	Z	3	1	0



<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y			
	Z			

<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	7	1	0

<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	3	1	0



<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y	2	0	1
	Z			

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y	2	0	1
	Z	7	1	0

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	3	1	0

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y			
	Z	7	1	0


<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y	2	0	1
	Z	3	1	0

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	3	1	0

<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y			
	Z			

<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	7	1	0

<div>X</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	3	1	0



<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y	2	0	1
	Z			

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y	2	0	1
	Z	7	1	0

<div>Y</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	3	1	0

We can update our routing table and then make sure we forward packets to the correct destination in the path

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X			
	Y			
	Z	7	1	0

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	7
	Y	2	0	1
	Z	3	1	0

<div>Z</div>		Cost to		
		X	Y	Z
Cost From	X	0	2	3
	Y	2	0	1
	Z	3	1	0


# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms

**Centralized/Global**- we know the edge costs of the network

**Link State** algorithms


(Dijkstra's Algorithm)



**Decentralized**- we do not know the edge costs of the entire network.

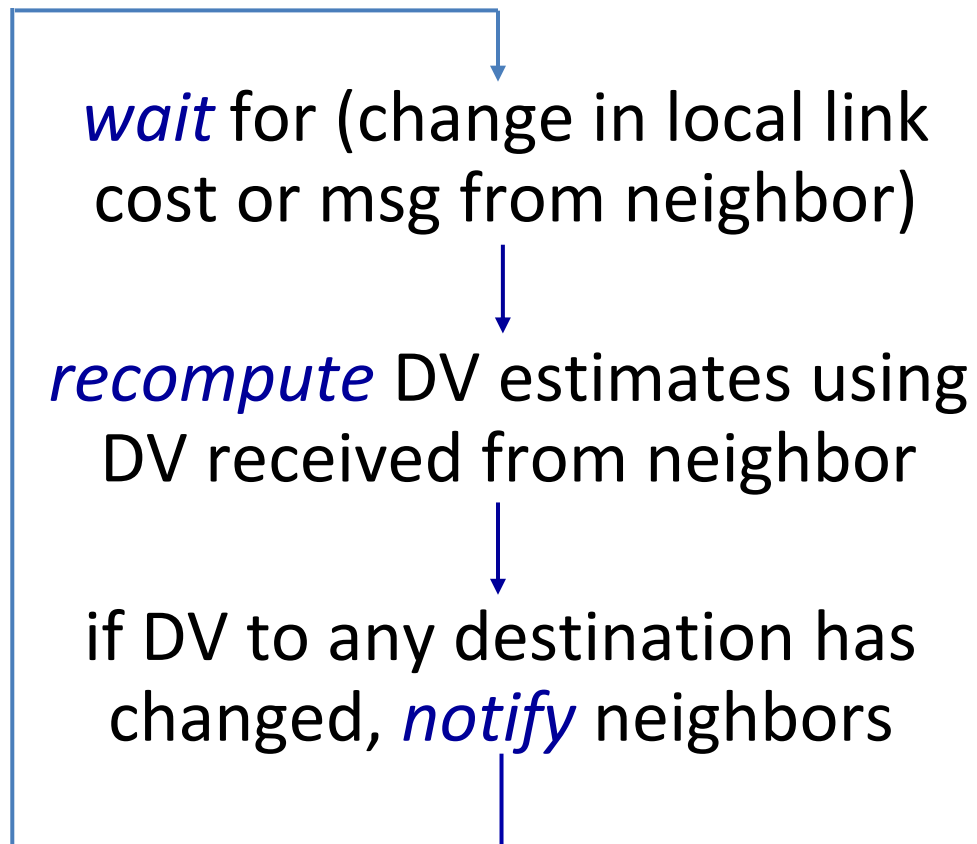
Only know edge costs to neighbors

**Distance Vector** algorithms



## Distance vector algorithm:

### each node:



**iterative, asynchronous:** each local iteration caused by:

- local link cost change
- DV update message from neighbor

**distributed, self-stopping:** each node notifies neighbors *only* when its DV changes

- neighbors then notify their neighbors – *only if necessary*
- no notification received, no actions taken!

# Distance vector: example




$t=0$

- All nodes have distance estimates to nearest neighbors (only)
- All nodes send their local distance vector to their neighbors

DV in a:

$D_a(a)=0$   
 $D_a(b)=8$   
 $D_a(c)=\infty$   
 $D_a(d)=1$   
 $D_a(e)=\infty$   
 $D_a(f)=\infty$   
 $D_a(g)=\infty$   
 $D_a(h)=\infty$   
 $D_a(i)=\infty$



A few asymmetries:

- missing link
- larger cost




# Distance vector example: iteration



$t=1$

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




# Distance vector example: iteration



$t=1$

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




# Distance vector example: iteration



$t=1$

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




# Distance vector example: iteration



$t=2$

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




# Distance vector example: iteration



$t=2$

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




# Distance vector example: iteration



$t=2$

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms

**Centralized/Global**- we know the edge costs of the network

**Link State** algorithms

(Dijkstra's Algorithm)




**These are not network protocols, these are simply general routing/shortest path algorithms**

**Decentralized**- we do not know the edge costs of the entire network.

Only know edge costs to neighbors

**Distance Vector** algorithms




# Routing tables are filled via **routing algorithms**

There are two types of routing algorithms

**Centralized/Global-** we know the edge costs of the network

Link State

OSPF




**Decentralized-** we do not know the edge costs of the network.

Only know neighbor

BGP

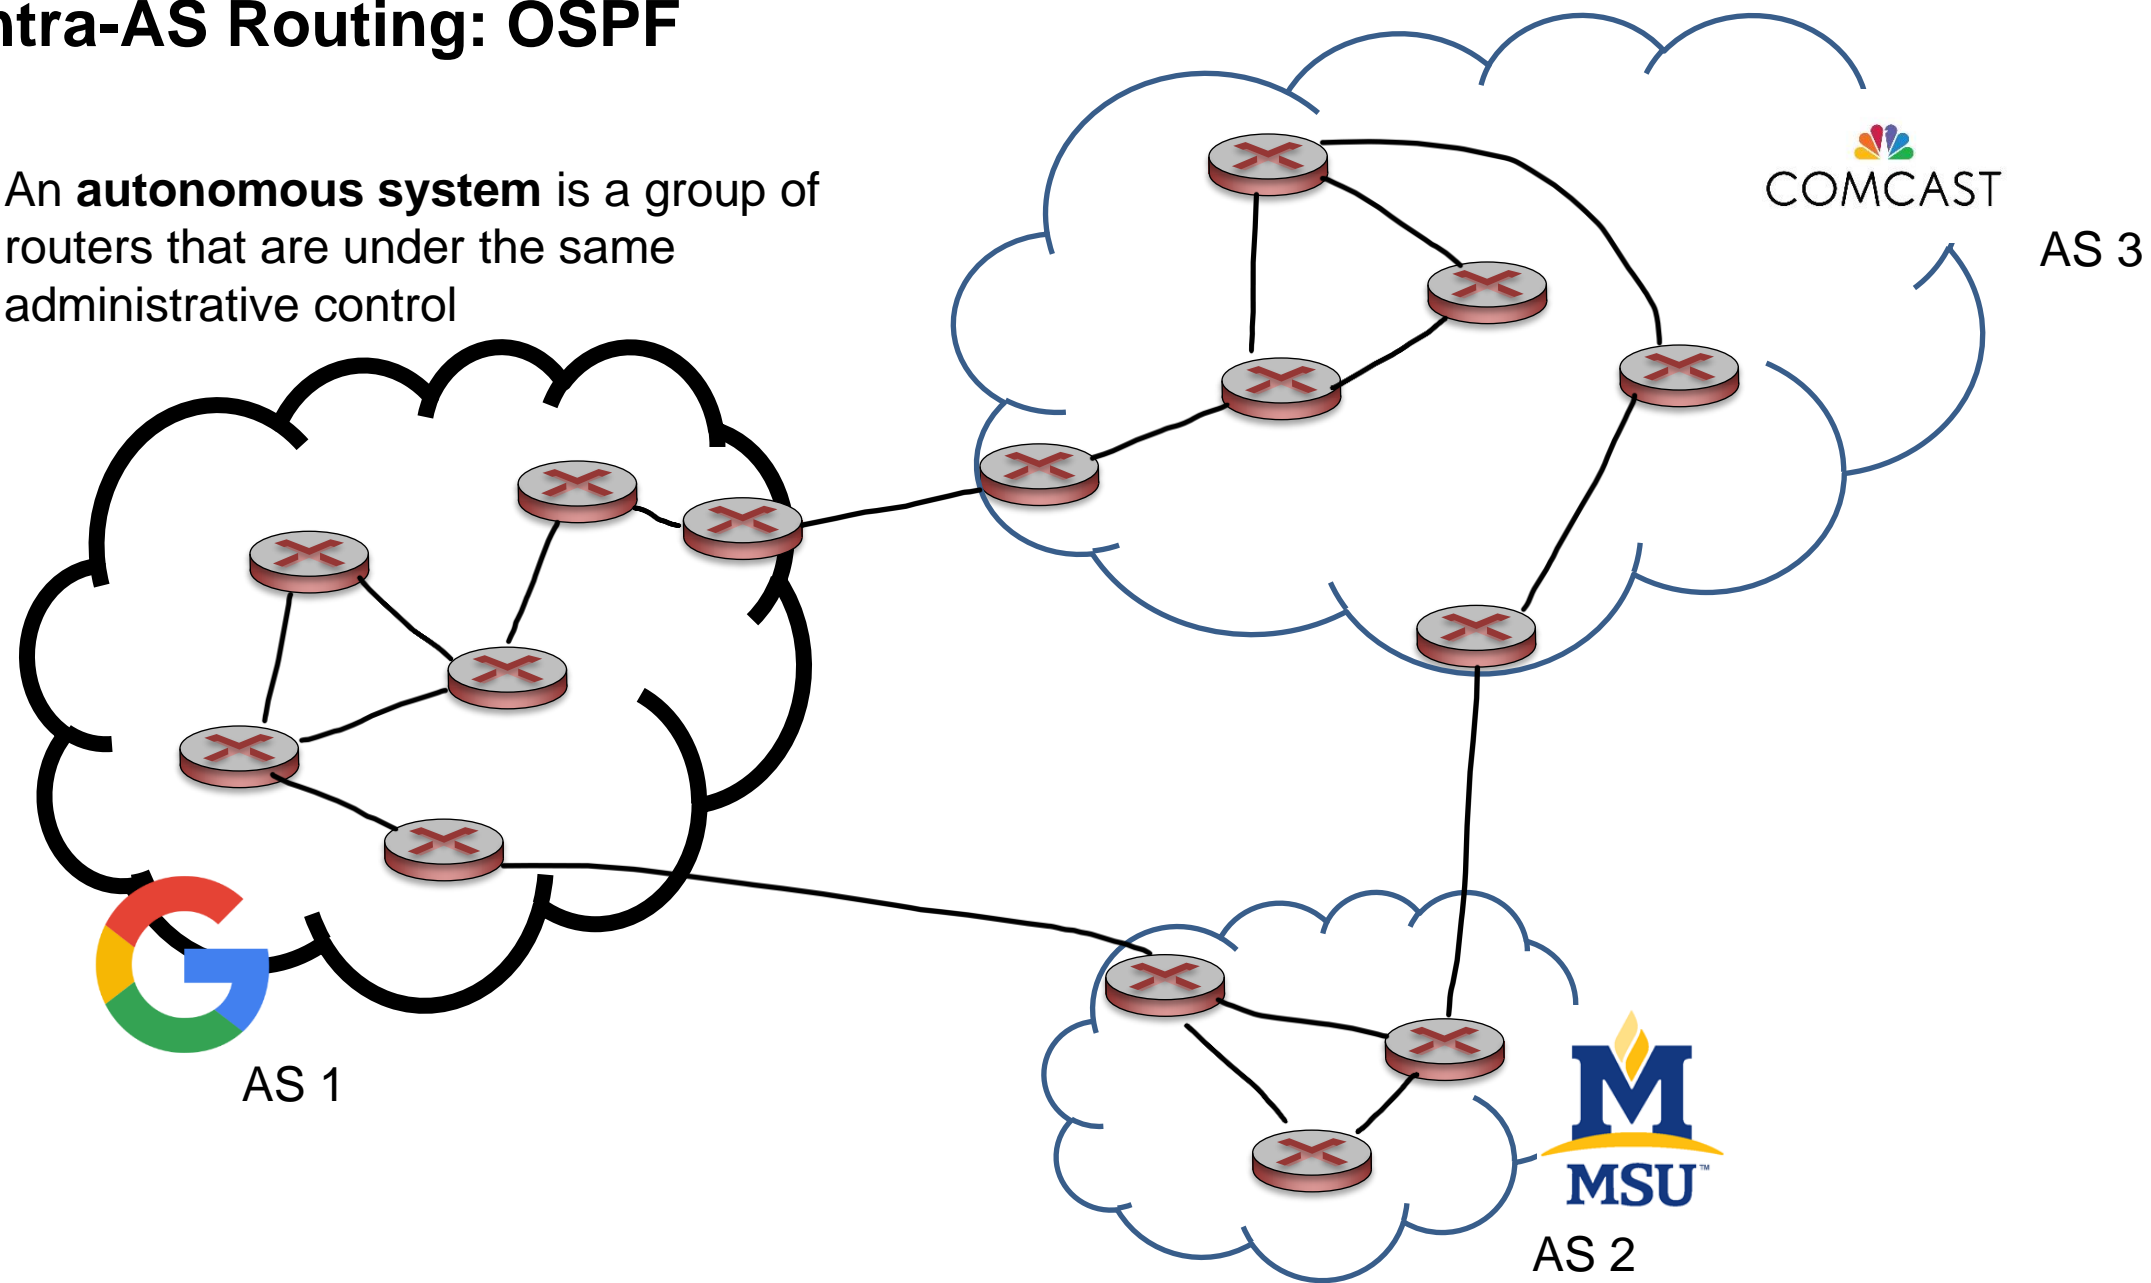
vector algorithms






# Intra-AS Routing: OSPF

An **autonomous system** is a group of routers that are under the same administrative control



# Intra-AS Routing: OSPF

An **autonomous system** is a group of routers that are under the same administrative control




## Open Shortest Path First


**OSPF** is a link-state protocol that uses flooding of link-state information and Dijkstra's least-cost algorithm

# Intra-AS Routing: OSPF

An **autonomous system** is a group of routers that are under the same administrative control



1. Each router constructs a topological map of the AS




## Open Shortest Path First


**OSPF** is a link-state protocol that uses flooding of link-state information and Dijkstra's least-cost algorithm

# Intra-AS Routing: OSPF

An **autonomous system** is a group of routers that are under the same administrative control



1. Each router constructs a topological map of the AS
2. Run Dijkstra's to determine shortest path to each subnet



## Open Shortest Path First


**OSPF** is a link-state protocol that uses flooding of link-state information and Dijkstra's least-cost algorithm

(Edge costs will be set by a network administrator)

If I wanted to find the path with the shortest amount of hops, what should edge cost be?

# Intra-AS Routing: OSPF


An **autonomous system** is a group of routers that are under the same administrative control



## Open Shortest Path First

**OSPF** is a link-state protocol that uses flooding of link-state information and Dijkstra's least-cost algorithm

- 1. Each router constructs a topological map of the AS
- 2. Run Dijkstra's to determine shortest path to each subnet



(Edge costs will be set by a network administrator)


(could set all edges to be a cost of 1)

- 3. Fill in routing table

A	1
B	2
C	3
...	...

# Routing Among the ISPs: BGP

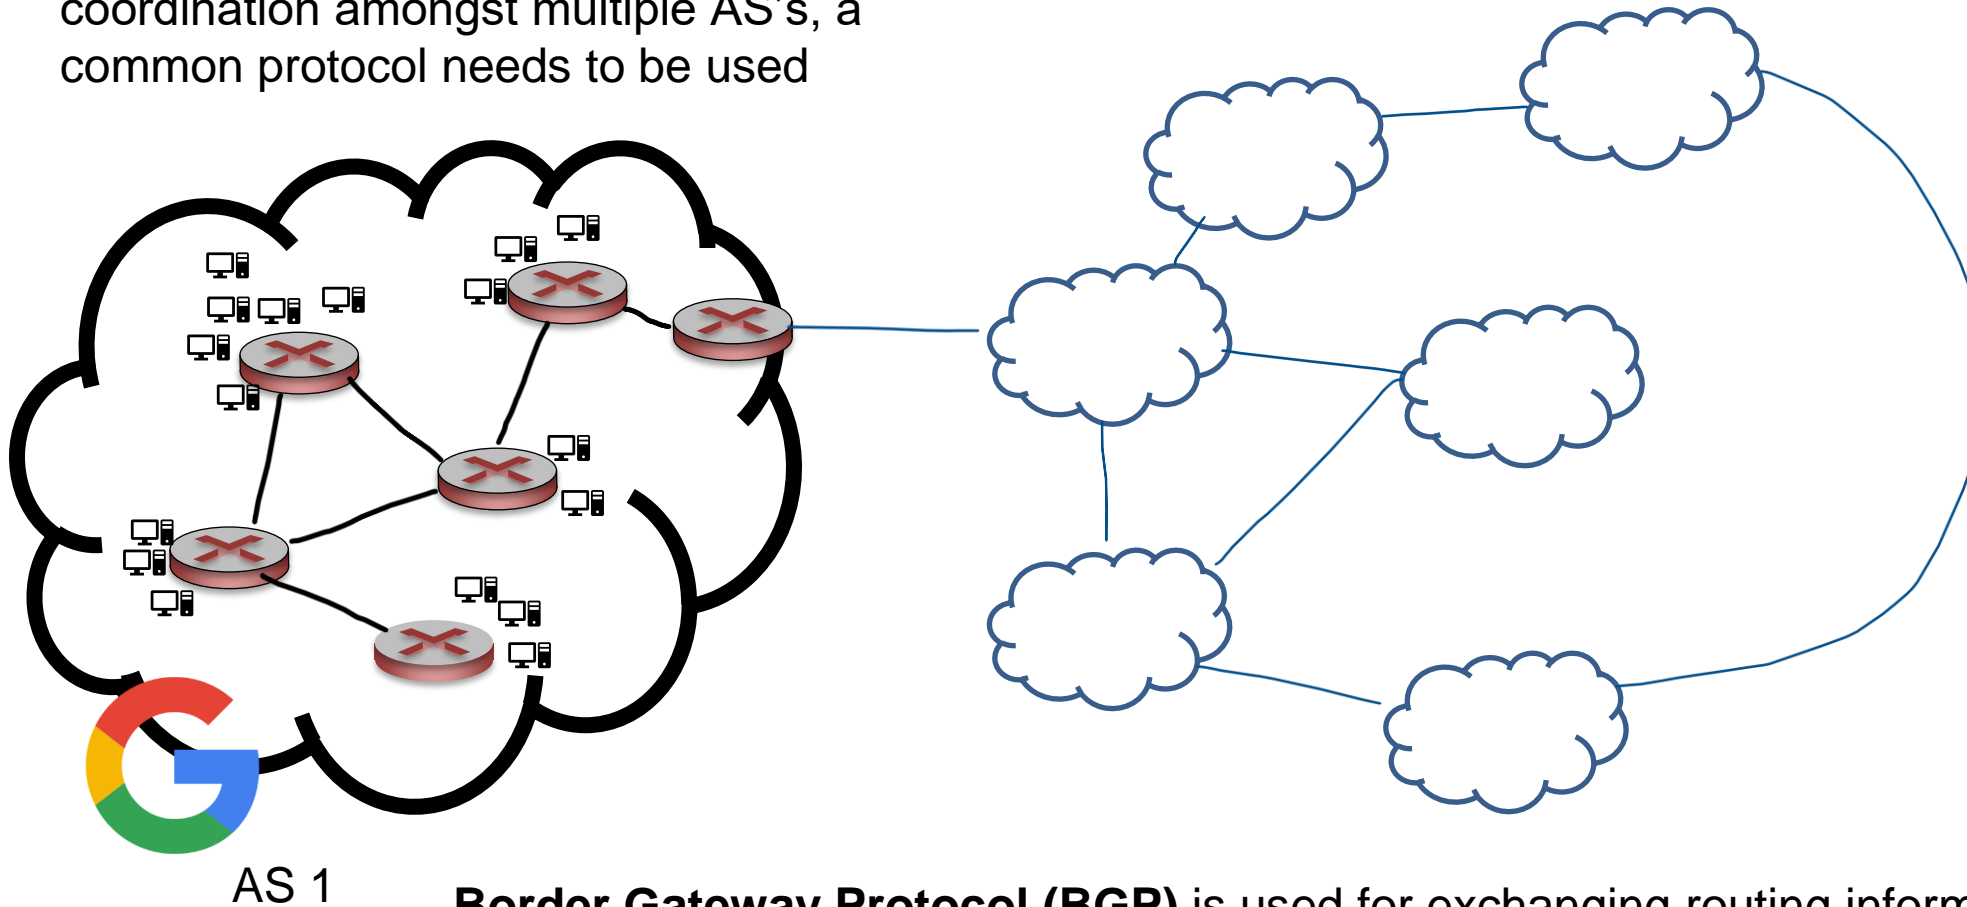
Inter-AS routing protocol involves coordination amongst multiple AS's, a common protocol needs to be used



**Border Gateway Protocol (BGP)** is used for exchanging routing information between AS

# Routing Among the ISPs: BGP

Inter-AS routing protocol involves coordination amongst multiple AS's, a common protocol needs to be used



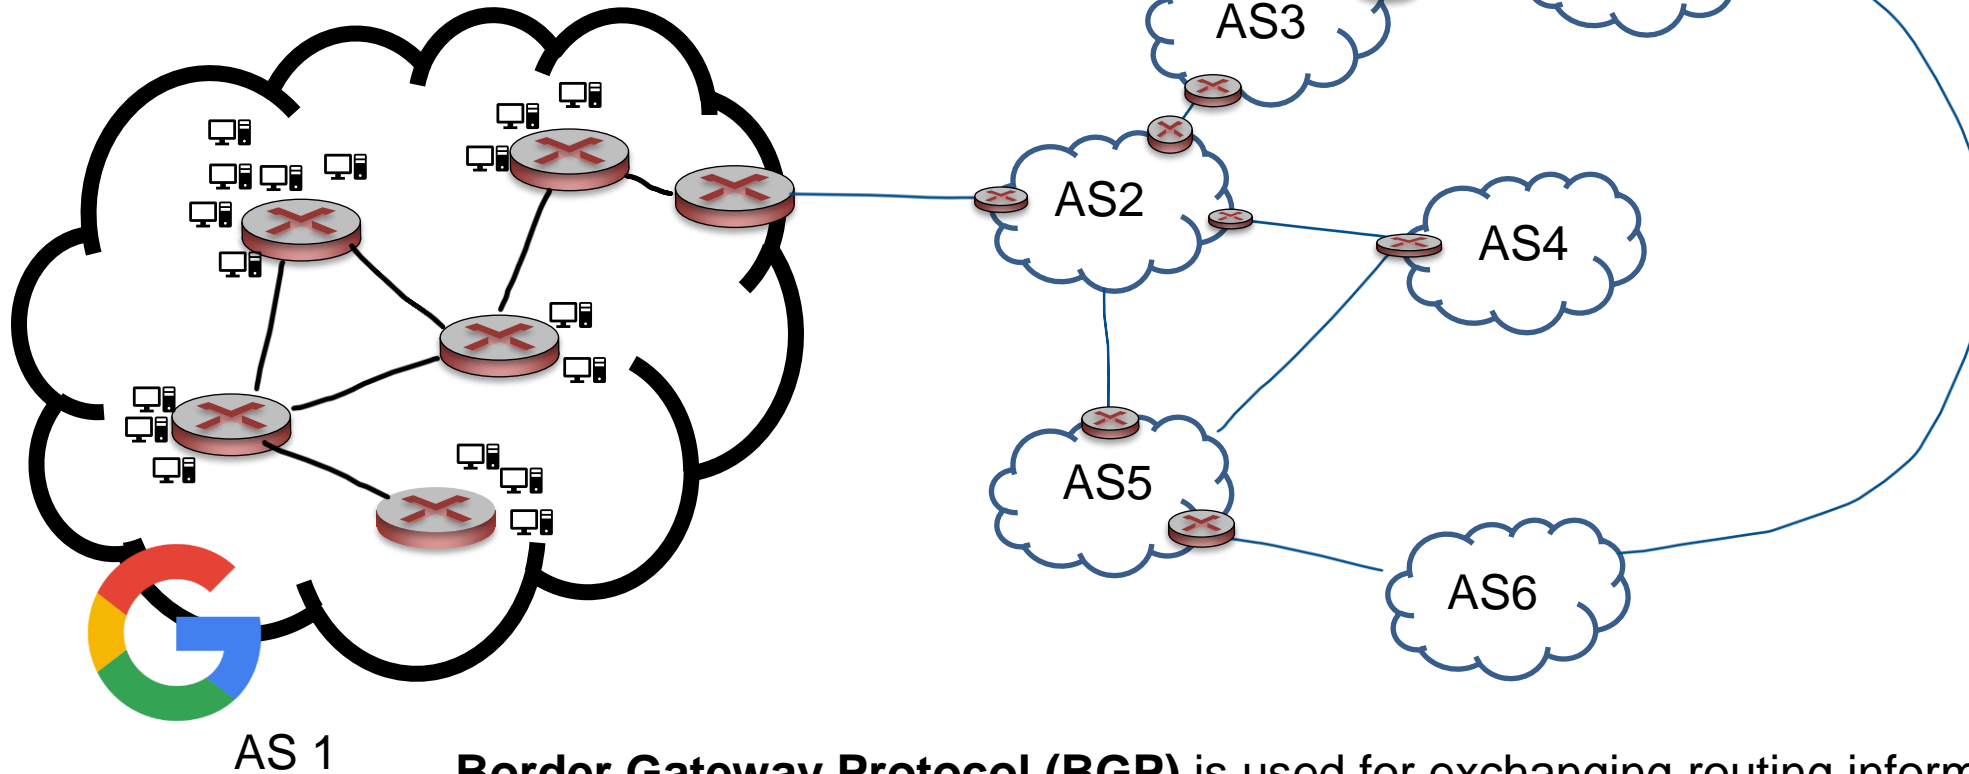
**Border Gateway Protocol (BGP)** is used for exchanging routing information between AS

BGP allows a router to tell other AS's that it exists and needs to be connected

# Routing Among the ISPs: BGP

AS consists of **gateway routers** and **internal routers**

Inter-AS routing protocol involves coordination amongst multiple AS's, a common protocol needs to be used



**Border Gateway Protocol (BGP)** is used for exchanging routing information between AS

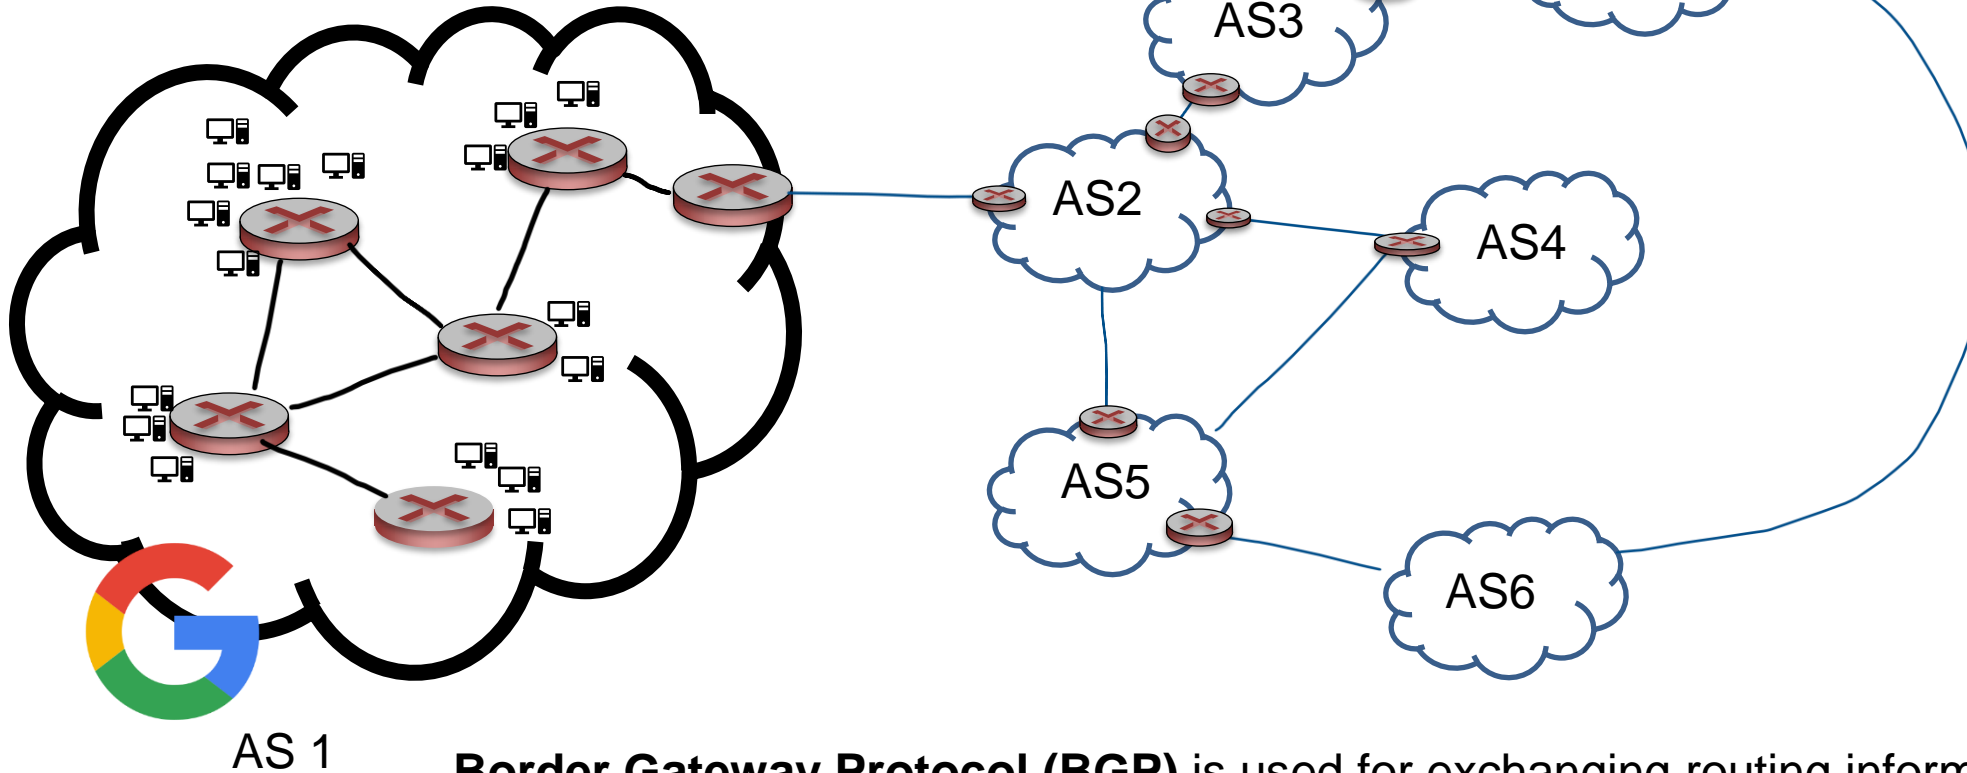
BGP allows a router to tell other AS's that it exists and needs to be connected



# Routing Among the ISPs: BGP

AS consists of **gateway routers** and **internal routers**

Inter-AS routing protocol involves coordination amongst multiple AS's, a common protocol needs to be used



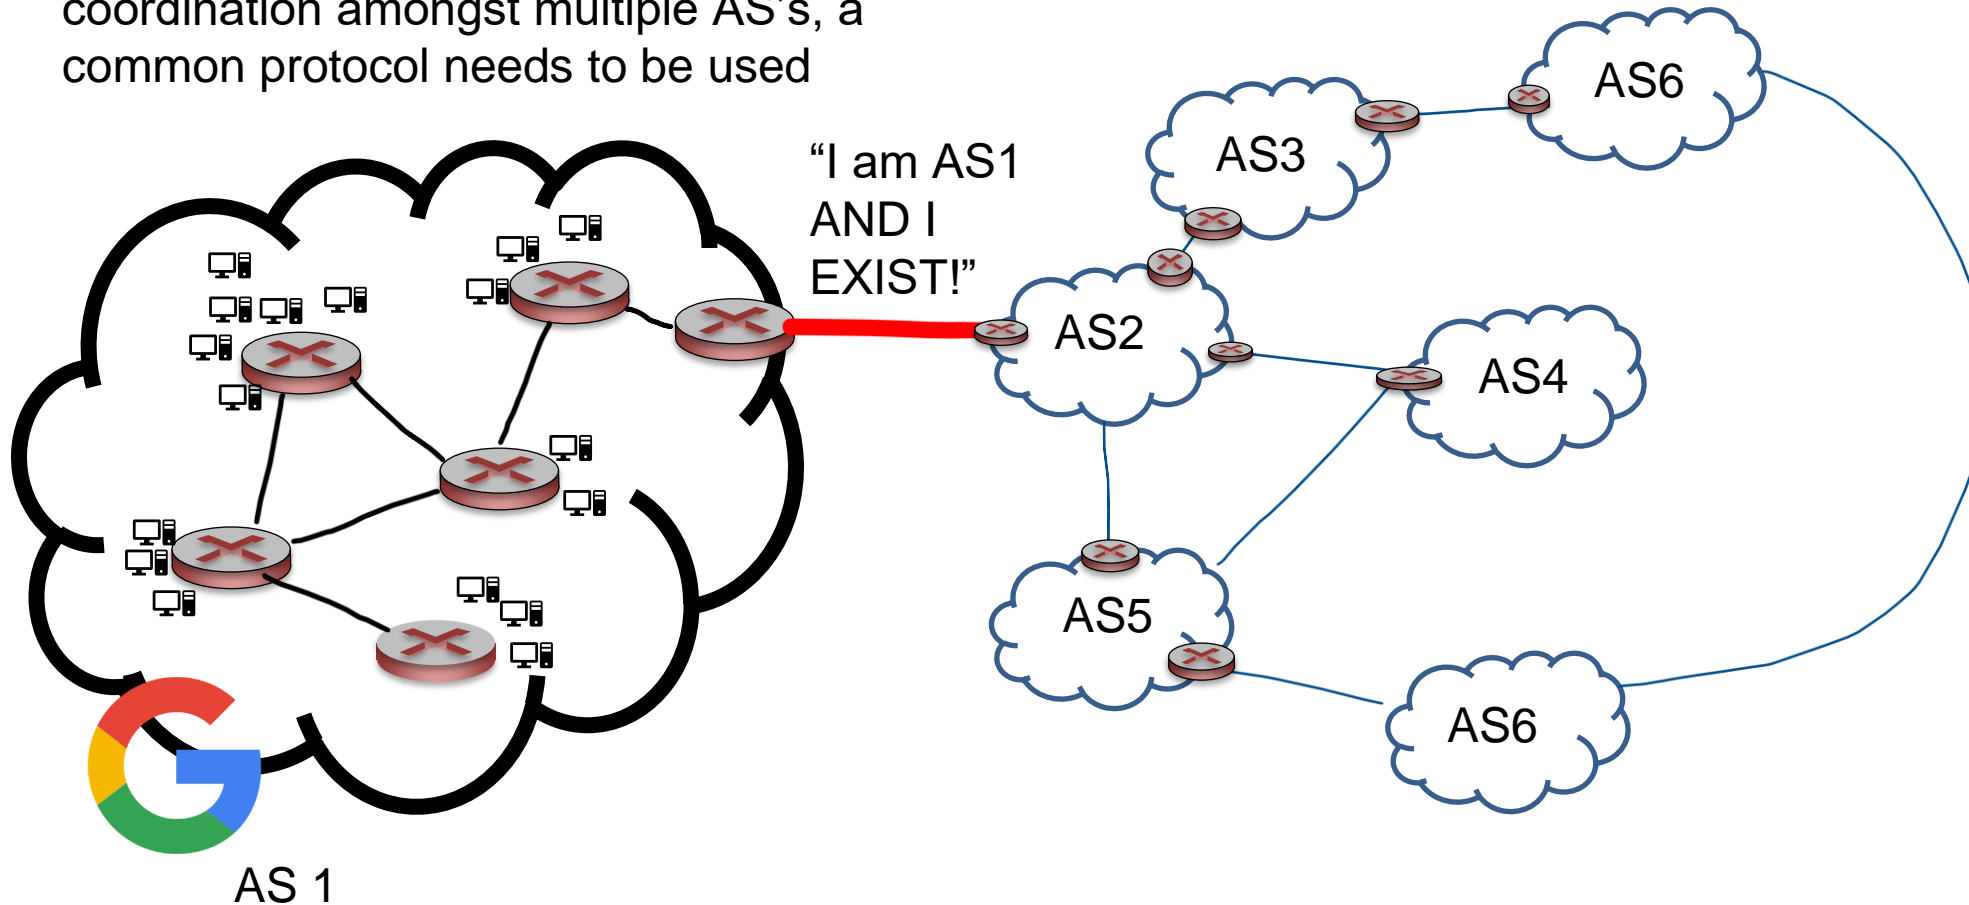
**Border Gateway Protocol (BGP)** is used for exchanging routing information between AS

BGP allows a router to tell other AS's that it exists and needs to be connected

# Routing Among the ISPs: BGP

AS consists of **gateway routers** and **internal routers**

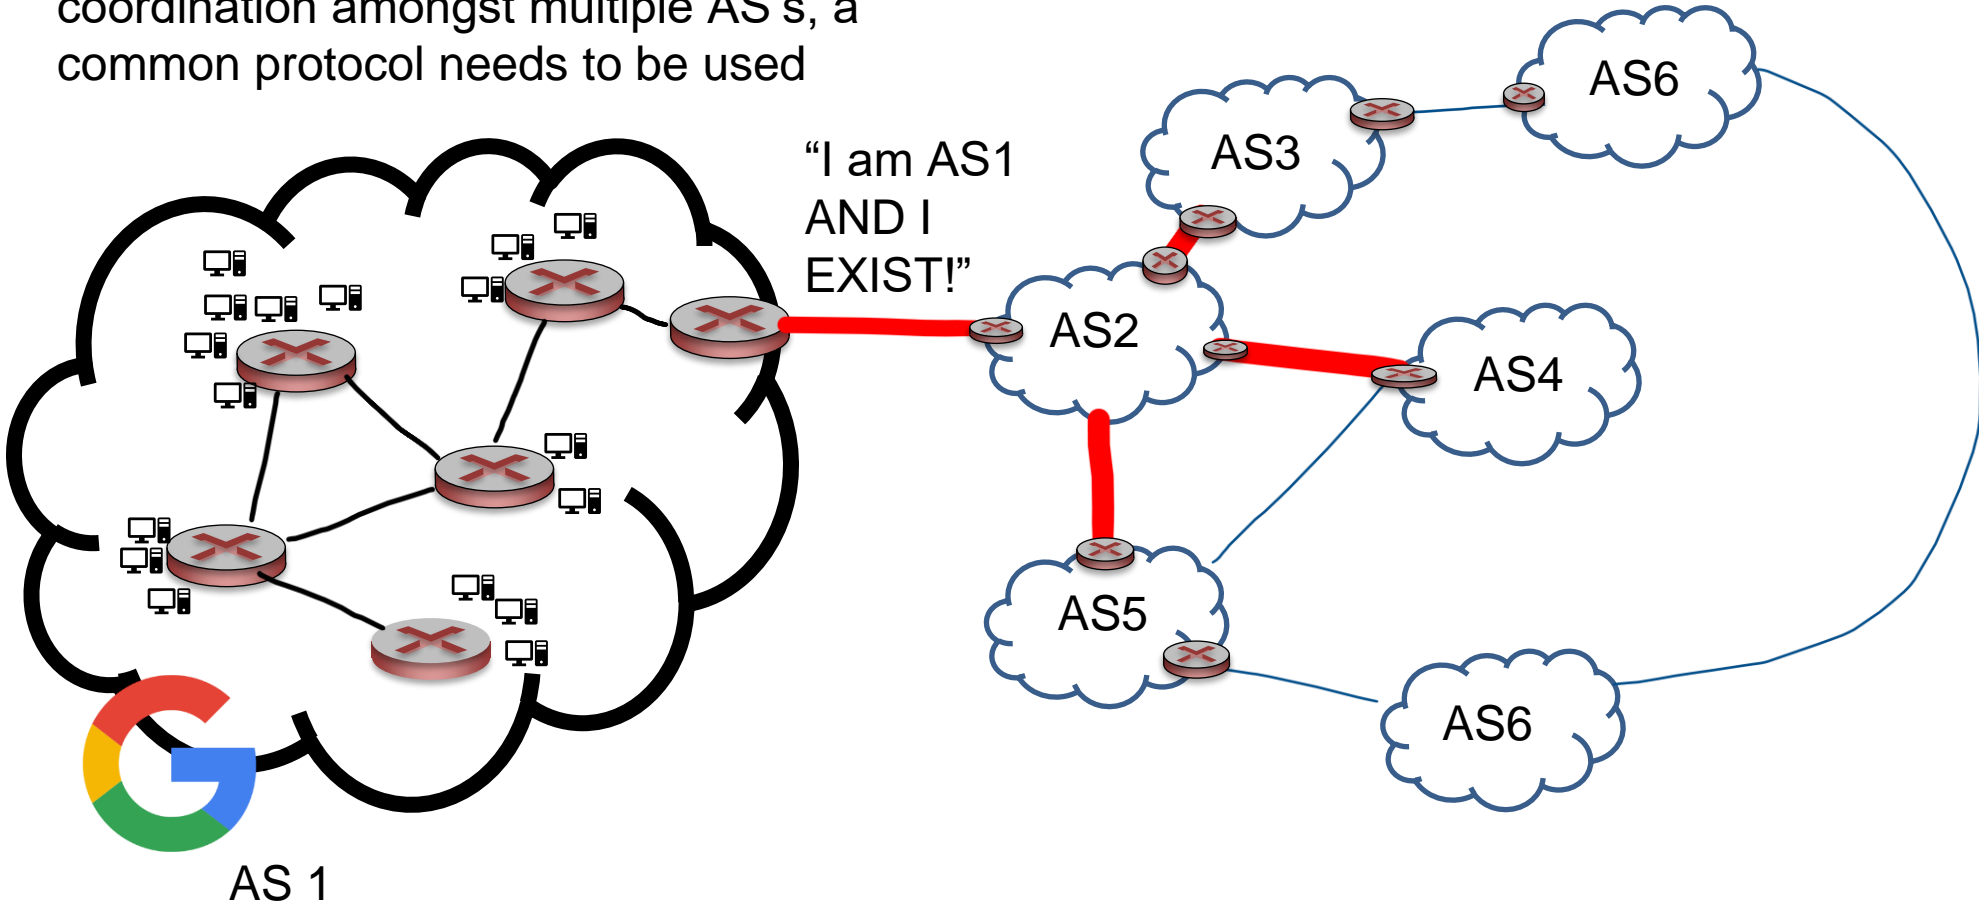
Inter-AS routing protocol involves coordination amongst multiple AS's, a common protocol needs to be used



# Routing Among the ISPs: BGP

AS consists of **gateway routers** and **internal routers**

Inter-AS routing protocol involves coordination amongst multiple AS's, a common protocol needs to be used



"A1 EXISTS AND FOUND THROUGH AS2"

# Internet inter-AS routing: BGP

- **BGP (Border Gateway Protocol):** *the* de facto inter-domain routing protocol
  - “glue that holds the Internet together”
- allows subnet to advertise its existence, and the destinations it can reach, to rest of Internet: *“I am here, here is who I can reach, and how”*
- BGP provides each AS a means to:
  - obtain destination network reachability info from neighboring ASes (**eBGP**)
  - determine routes to other networks based on reachability information and *policy*
  - propagate reachability information to all AS-internal routers (**iBGP**)
  - **advertise** (to neighboring networks) destination reachability info

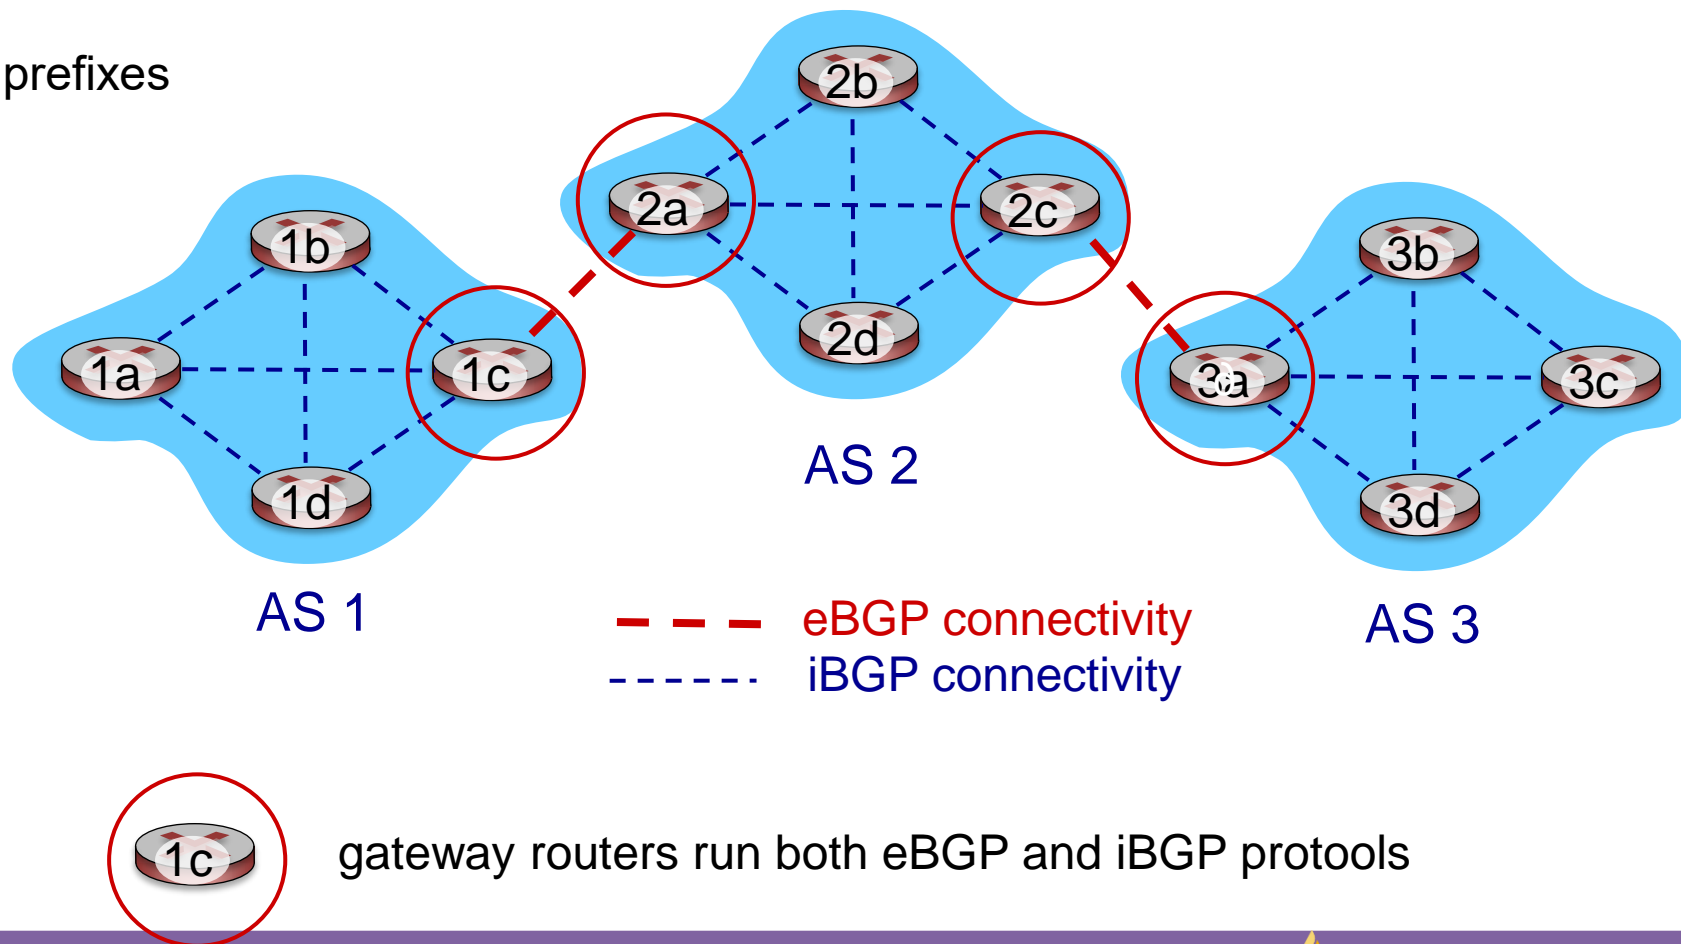
# Routing Among the ISPs: BGP

BGP is the routing protocol used for routing amongst different ISPs + AS

Two important functions

→ Obtain prefix reachability information from neighboring ASs (CIDR)

→ Determine the “best” routes to the prefixes



# Routing Among the ISPs: BGP

BGP is the routing protocol used for routing amongst different ISPs + AS

- Two important functions
- Obtain prefix reachability information from neighboring ASs
  - Determine the “best” routes to the prefixes


Prefix **X** connect

## External BGP (eBGP)

- 3a → 2c “Hey I have X”
- 2a → 1c “Hey AS 3 has X and I have AS3”

## Internal BGP (iBGP)

- 2c → 2b
  - 2c → 2d
  - 2c → 2a
-  gateway routers run both eBGP and iBGP protocols




# BGP basics

- **BGP session:** two BGP routers (“peers”) exchange BGP messages over semi-permanent TCP connection:
  - advertising *paths* to different destination network prefixes (BGP is a “path vector” protocol)

when AS3 gateway router 3a advertises path **AS3,X** to AS2 gateway router 2c:

AS3 *promises* to AS2 it will forward datagrams towards X



# BGP protocol messages


- BGP messages exchanged between peers over TCP connection
- BGP messages [RFC 4371]:
  - **OPEN**: opens TCP connection to remote BGP peer and authenticates sending BGP peer
  - **UPDATE**: advertises new path (or withdraws old)
  - **KEEPALIVE**: keeps connection alive in absence of UPDATES; also ACKs OPEN request
  - **NOTIFICATION**: reports errors in previous msg; also used to close connection




# Path attributes and BGP routes

- BGP advertised route: prefix + attributes
  - prefix: destination being advertised
  - two important attributes:
    - **AS-PATH**: list of ASes through which prefix advertisement has passed
    - **NEXT-HOP**: indicates specific internal-AS router to next-hop AS
- **policy-based routing**:
  - gateway receiving route advertisement uses *import policy* to accept/decline path (e.g., never route through AS Y).
  - AS policy also determines whether to *advertise* path to other neighboring ASes

# BGP path advertisement



- AS2 router 2c receives path advertisement **AS3,X** (via eBGP) from AS3 router 3a
- based on AS2 policy, AS2 router 2c accepts path AS3,X, propagates (via iBGP) to all AS2 routers
- based on AS2 policy, AS2 router 2a advertises (via eBGP) path **AS2, AS3, X** to AS1 router 1c





## Collectors

RouteViews is collecting BGP Updates at the following locations

### Exchanges

Host	MFG	Proto	Location
amsix.ams.routeviews.org	FRR	IPv4/6	AMS-IX Amsterdam, Netherlands
cix.atl.routeviews.org	FRR	IPv4/6	CIX-ATL Atlanta, Georgia
dectx.jhb.routeviews.org	FRR	IPv4/6	DE-CIX KUL, Johor Bahru, Malaysia
iraq-ixp.bgw.routeviews.org	FRR	IPv4/6	IRAQ-IXP Baghdad, Iraq
pacwave.lax.routeviews.org	FRR	IPv4/6	Pacific Wave, Los Angeles, California
pit.scl.routeviews.org	FRR	IPv4/6	PIT Chile Santiago, Santiago, Chile
pitmx.gro.routeviews.org	FRR	IPv4/6	PIT Chile MX, Querétaro, Mexico
route-views.routeviews.org	Cisco	IPv4	U of Oregon, Eugene Oregon

<https://www.routeviews.org/routeviews/index.php/collectors/>

```
route-views>show ip bgp sum
BGP router identifier 128.223.51.103, local AS number 6447
BGP table version is 355532718, main routing table version 355532718
Path RPKI states: 9671162 valid, 8728431 not found, 5720 invalid
990349 network entries using 245606552 bytes of memory
18405313 path entries using 2208637560 bytes of memory
2919344/172592 BGP path/bestpath attribute entries using 723997312 bytes of memory
2802466 BGP AS-PATH entries using 142729518 bytes of memory
227141 BGP community entries using 38426562 bytes of memory
2503 BGP extended community entries using 162876 bytes of memory
0 BGP route-map cache entries using 0 bytes of memory
0 BGP filter-list cache entries using 0 bytes of memory
BGP using 3359560380 total bytes of memory
BGP activity 15050631/13844711 prefixes, 883710380/861987137 paths, scan interval 60 secs
```

Lots of network and path entries on this BGP router

The list of neighbors it interacts with

Neighbor / PfxRef	V	AS	MsgRcvd	MsgSent	TblVer	InQ	OutQ	Up/Down	State
4.68.4.46	4	3356	15253358	590045	355538393	0	0	9w0d	9
46903									
12.0.1.63	4	7018	29842482	89389	355538393	0	0	8w0d	9
47328									
37.139.139.17	4	57866	271682	822	355538393	0	0	06:15:50	95
0829									
64.71.137.241	4	6939	8581620	44533	355538393	0	0	4w0d	96
9900									
77.39.192.30	4	20912	7062039	162429	355538393	0	0	2w3d	94
9277									
89.149.178.10	4	3257	5761037	3436	355538393	0	0	1w3d	94
7649									
91.218.184.60	4	49788	17612113	330474	355538393	0	0	9w0d	9
51340									
94.142.247.3	4	8283	51450625	330485	355538393	0	0	9w0d	9
53214									
114.31.199.16	4	4826	19029982	355588	355538393	0	0	9w0d	9
73382									
132.198.255.253	4	1351	3839141	33365	355538393	0	0	3w0d	97
3404									
140.192.8.16	4	20130	19891315	141651	355538393	0	0	6w3d	9
73729									
144.228.241.130	4	1239	199819	40395	355538393	0	0	6w0d	3
0597									
154.11.12.212	4	852	25322844	454901	355538393	0	0	9w0d	9
49731									
162.250.137.254	4	4901	72911065	400400	355538393	0	0	21w6d	9
52703									
153.251.163.3	4	53367	2456062	245361	355538393	0	0	9w0d	16

# ICMP (Internet Control Message Protocol)

used by hosts & routers to  
communicate network-level  
information

error reporting: unreachable  
host, network, port, protocol  
echo request/reply (used by  
ping)

network-layer “above” IP:

ICMP msgs carried in IP  
datagrams

**ICMP message:** type, code plus  
first 8 bytes of IP datagram  
causing error

<u>Type</u>	<u>Code</u>	<u>description</u>
0	0	echo reply (ping)
3	0	dest. network unreachable
3	1	dest host unreachable
3	2	dest protocol unreachable
3	3	dest port unreachable
3	6	dest network unknown
3	7	dest host unknown
4	0	source quench (congestion control - not used)
8	0	echo request (ping)
9	0	route advertisement
10	0	router discovery
11	0	TTL expired
12	0	bad IP header