

CSCI 476: Computer Security

Operating Systems, Processes, and `forking()`

Reese Pearsall
Fall 2024

Announcements

NO CLASS ON THURSDAY

(I'll still be around if you need help with anything)


Lab 0 due on Sunday 9/8 @ 11:59 PM

(All assignments will be due on Sundays)

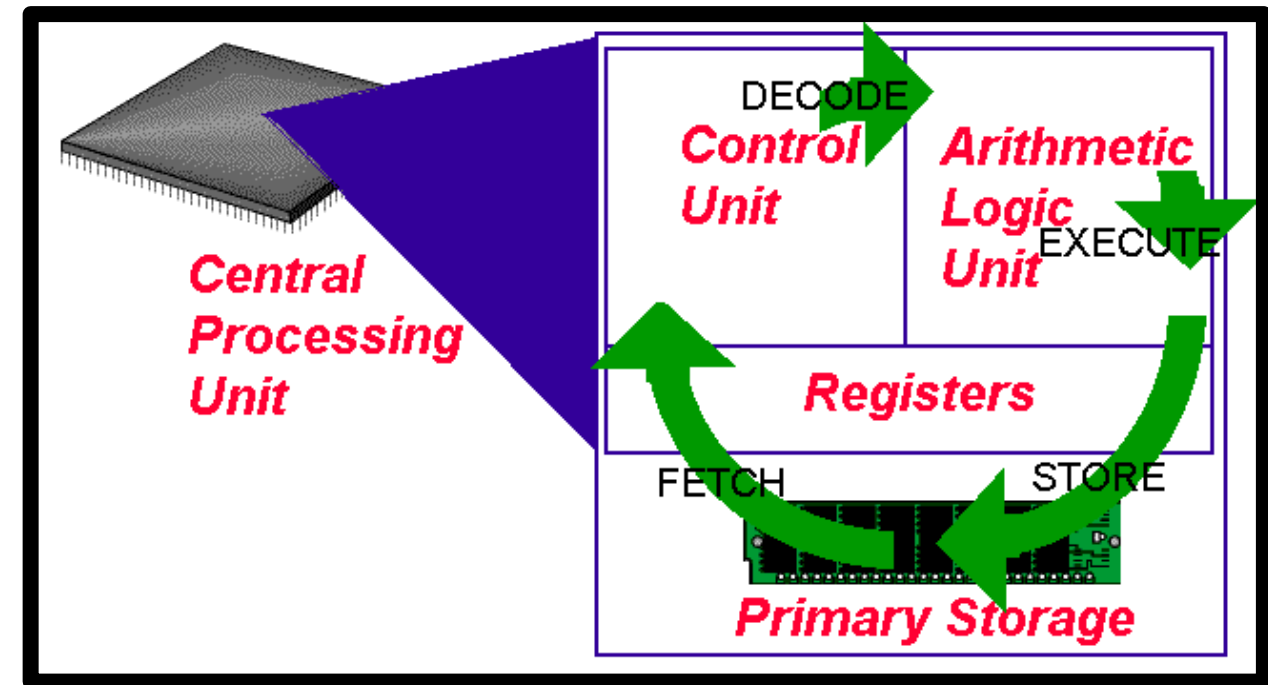
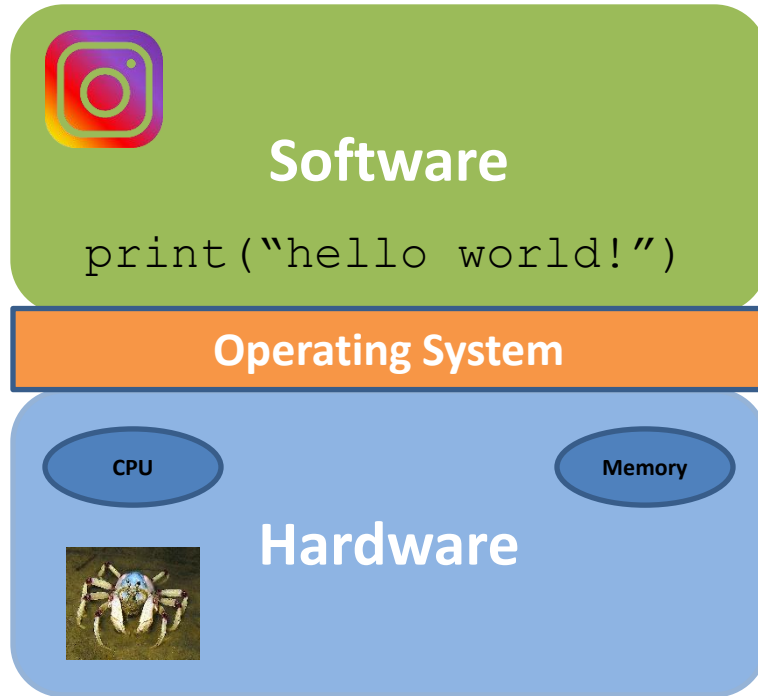
If you have an M1/M2 chip, or if you are still struggling with your VM,
check in with me this week

To understand the technical aspects of security, we must have a good understanding of how ~~computers~~ work

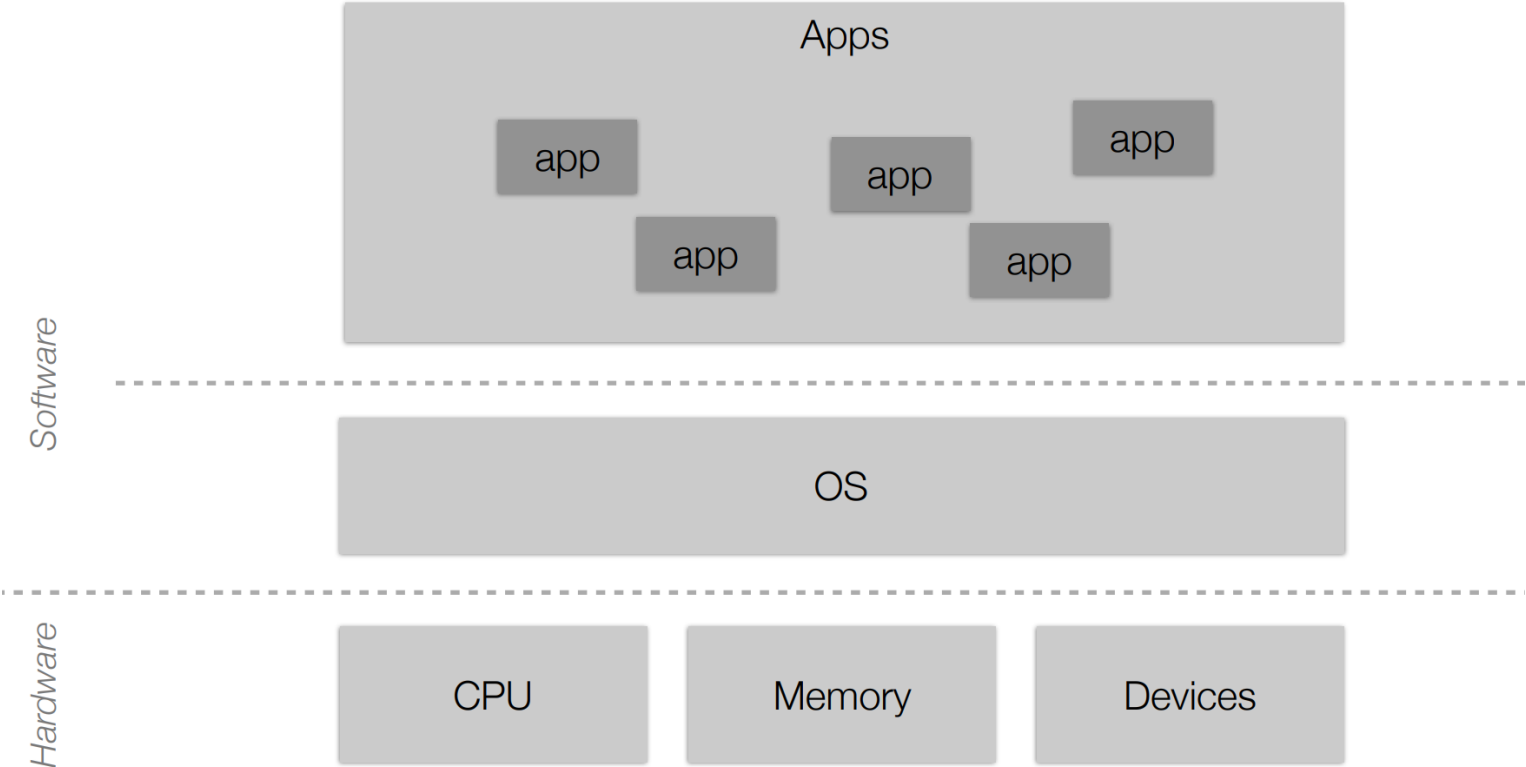
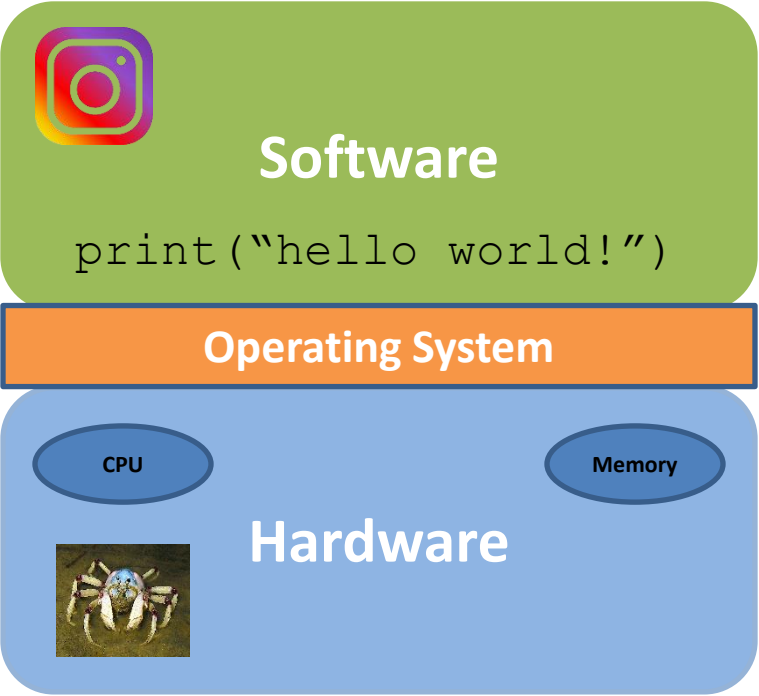
operating systems



The Operating System



The Operating System



The jobs of an Operating System

1. Process Manager
“The Coach”

2. Interface Manager
“The Bouncer”

3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”



The jobs of an Operating System

1. Process Manager
“The Coach”

2. Interface Manager
“The Bouncer”

3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”

*This will be the focus
of the first half of
lecture*



Source code to binary

```
#include <stdio.h>


int main() {
    printf("Hello W0rld! \n");

    int x = 0;
    int y = 3;

    int z = x + y;

    printf("%d %d %d \n", x, y, z);
    return 0;
}
```

Preprocessor



- Removal of comments
- Expand Macros

Compiler

```
0000000000000000 <main>:
0: f3 0f 1e fa          endbr64
4: 55                   push    %rbp
5: 48 89 e5             mov     %rsp,%rbp
8: 48 83 ec 10          sub     $0x10,%rsp
c: 48 8d 3d 00 00 00 00 lea     0x0(%rip),%rdi    # 13 <main+0x13>
13: e8 00 00 00 00      callq  18 <main+0x18>
18: c7 45 f4 00 00 00 00 movl    $0x0,-0x0(%rbp)
1f: c7 45 f8 03 00 00 00 movl    $0x3,-0x8(%rbp)
26: 8b 55 f4             mov     -0xc(%rbp),%edx
29: 8b 45 f8             mov     -0x8(%rbp),%eax
2c: 01 d0               add     %edx,%eax
2e: 89 45 fc             mov     %eax,-0x4(%rbp)
31: 8b 4d fc             mov     -0x4(%rbp),%ecx
34: 8b 55 f8             mov     -0x8(%rbp),%edx
37: 8b 45 f4             mov     -0xc(%rbp),%eax
3a: 89 c6               mov     %eax,%esi
3c: 48 8d 3d 00 00 00 00 lea     0x0(%rip),%rdi    # 43 <main+0x43>
43: b8 00 00 00 00      mov     $0x0,%eax
48: e8 00 00 00 00      callq  4d <main+0x4d>
4d: b8 00 00 00 00      mov     $0x0,%eax
52: c9                   leaveq  %eax
53: c3                   retq
```

Assembler

- Converted to assembly code
- .s file

./hello_world



Linker

```
1 00000000 00000100 0000000000000000
2 01011110 00001100 11000010 0000000000000010
3 11101111 00010110 00000000000000101
4 11101111 10011110 0000000000001011
5 11111000 10101101 11011111 0000000000010010
6 01100010 11011111 0000000000010101
7 11101111 00000010 11111011 0000000000010111
8 11110100 10101101 11011111 0000000000011110
9 00000011 10100010 11011111 0000000000100001
10 11101111 00000010 11111011 0000000000100100
11 01111110 11110100 10101101
12 11111000 10101110 11000101 0000000000101011
13 00000110 10100010 11111011 0000000000110001
14 11101111 00000010 11111011 0000000000110100
15 01010000 11010100 0000000000111011
16 00000100 0000000000111101
```


What happens when we run `./hello_world` ?

It gets turned into a **process**

A **process** is an instance of a running program on a computer

Processes										
Performance App history Startup Users Details Services										
Name	Status	37% CPU	54% Memory	1% Disk	1% Network	17% GPU	GPU engine	Power usage	Power usage t...	
> Firefox (42)		6.5%	1,304.5 MB	0.5 MB/s	3.1 Mbps	9.0%	GPU 0 - Video Decode	High	Very low	
> Google Chrome (14)		0.8%	484.9 MB	0 MB/s	0 Mbps	0%		Low	Very low	
> Discord (32 bit) (6)		4.3%	328.8 MB	0 MB/s	8.7 Mbps	6.6%	GPU 0 - Video Encode	Moderate	Very low	
> Search		5.0%	185.9 MB	0.2 MB/s	0.8 Mbps	0%	GPU 0 - 3D	Moderate	Very low	
> Antimalware Service Executable		3.8%	178.2 MB	0.1 MB/s	0 Mbps	0%		Moderate	Very low	
Google Chrome		0%	175.4 MB	0 MB/s	0 Mbps	0%		Very low	Very low	
Slack		0%	95.5 MB	0 MB/s	0 Mbps	0%		Very low	Very low	
Steam Client WebHelper		0%	89.1 MB	0 MB/s	0 Mbps	0%		Very low	Very low	
Google Chrome		0%	82.6 MB	0 MB/s	0 Mbps	0%		Very low	Very low	
> Microsoft PowerPoint (32 bit) (2)		0.1%	69.3 MB	0 MB/s	0 Mbps	0%		Very low	Very low	
SteelSeries GG Core		0.2%	67.7 MB	0 MB/s	0 Mbps	0%		Very low	Very low	
Steam Client WebHelper		0%	66.1 MB	0 MB/s	0 Mbps	0%		Very low	Very low	

A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

1. Executable Code

2. Associated Data

3. Execution Context/Bookkeeping information

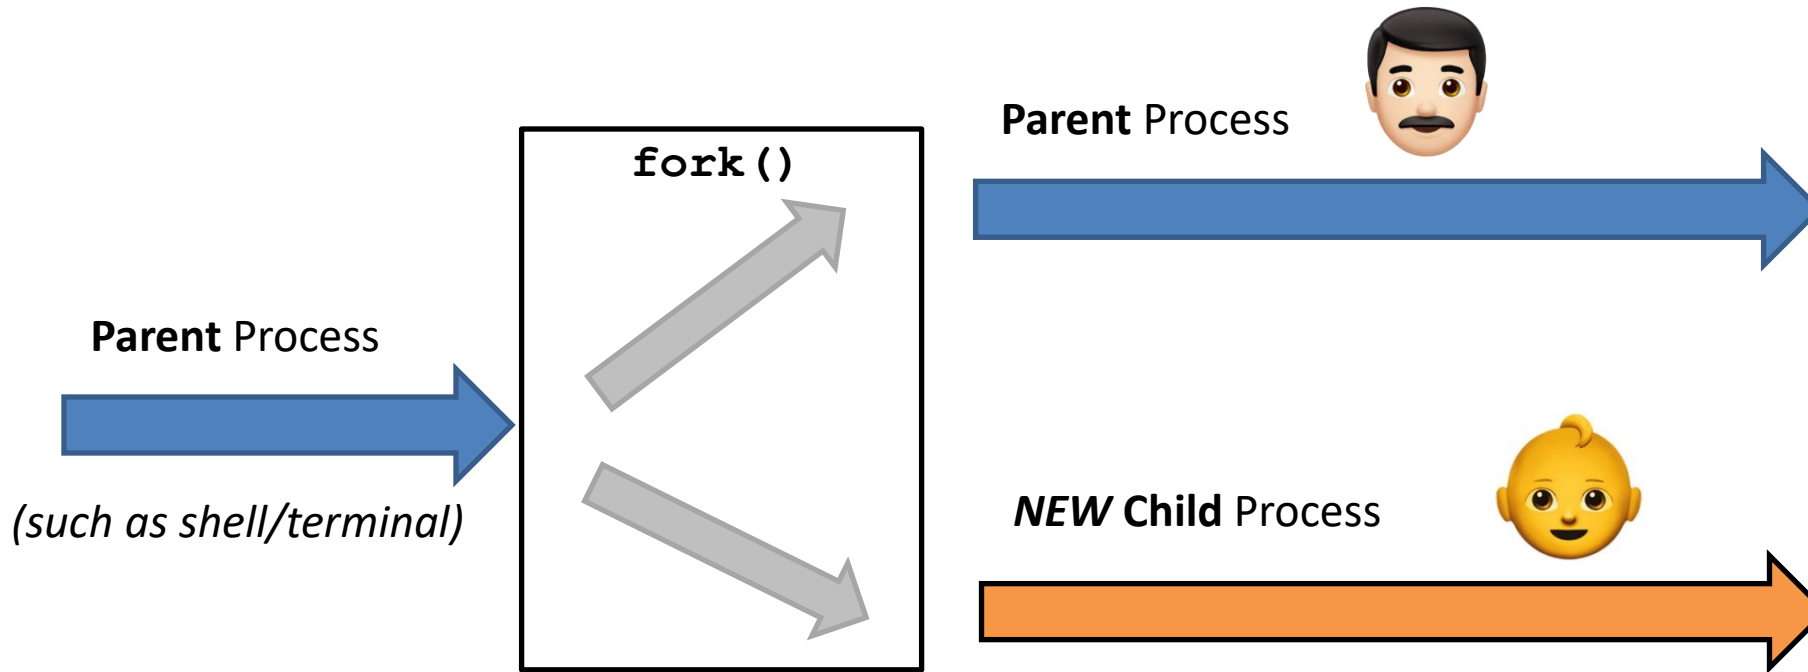
(info that the OS needs to handle the process)

Main Memory

Process A Information
Process A Data
Process A Executable Code
Process B Information
Process B Data
Process B Executable Code

Ok, but how do we *actually* create a process?

- In the Unix family (and others), we use **fork ()** to create a new process




fork () duplicates a process so that instead of one process, you get two!

fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;  
  
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }  
  
    sleep(1);  
    printf("I'm the parent.");  
  
    return 0;  
}
```

We check the return value
of **fork()**!





fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;
```

We check the return value
of **fork()** !

```
    pid = fork() ;  
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }  
  
    sleep(1);  
    printf("I'm the parent.");  
  
    return 0;  
}
```



1. Remember, **fork()** creates two process that are both actively running

fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;
```


```
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }
```

```
    sleep(1);  
    printf("I'm the parent.);
```

```
    return 0;
```

```
}
```

We check the return value
of **fork()** !



2. **fork()** always returns 0 for the child process, the parent process jumps to the code after the if statement

fork() duplicates a process so that instead of one process, you get two!

How can we tell the parent and child apart?

```
int main(void) {  
    int pid;
```


We check the return value
of **fork()** !

```
    pid = fork() ;  
    if (0 == pid) {  
        // I'm the child  
        printf("Hi, I'm the child. \n");  
    }
```

```
    sleep(1);  
    printf("I'm the parent.);
```

```
    return 0;
```

```
}
```



parent

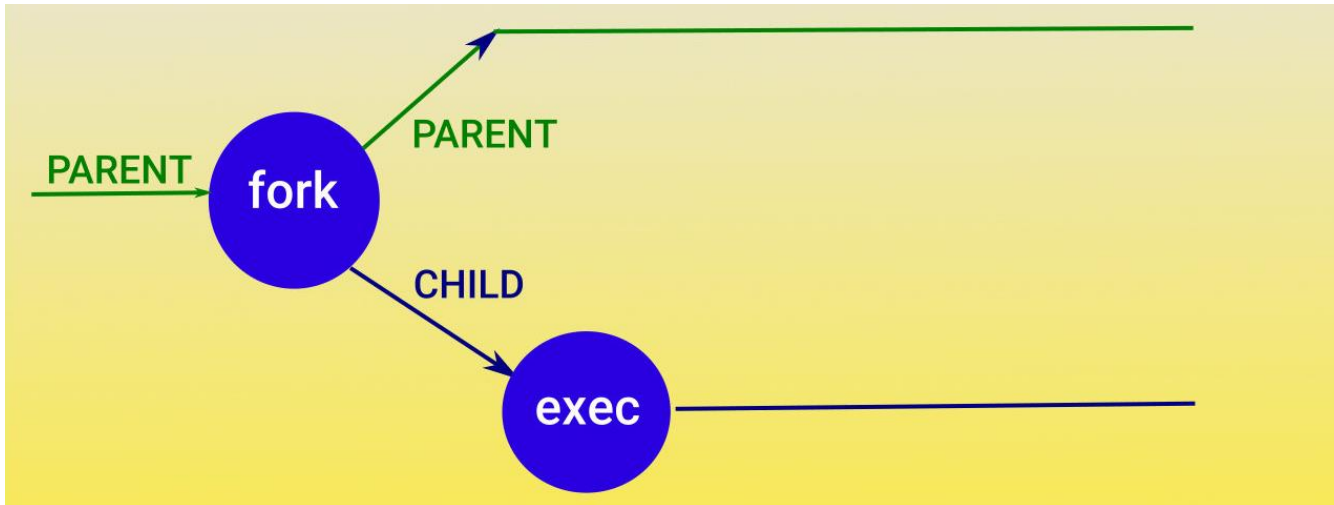
3. **fork()** always returns 0 for the child process, so the child process will execute the code in the `if` statement

Demo?

fork1.c

Issue: We want our child process to run an entirely new program
(hello_world c program)

We use the **exec()** family of functions to execute a different program

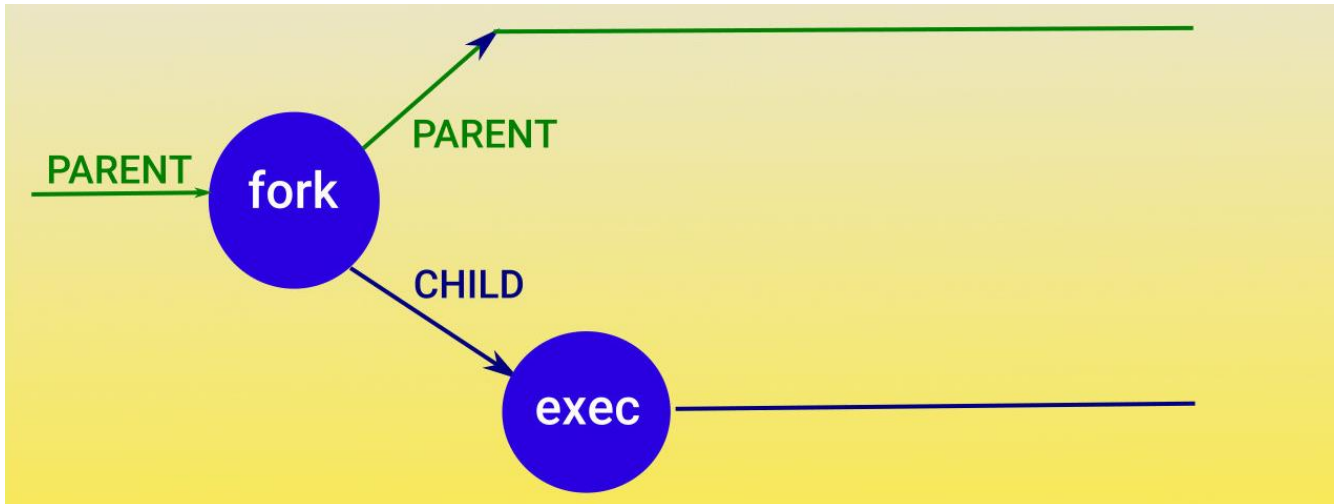


There are many different forms of the **exec()** function call

```
char *name[2];  
name[0] = "../hello";  
name[1] = NULL;  
execve(name[0], name, NULL);
```

Issue: We want our child process to run an entirely new program
(hello_world c program)

We use the **exec()** family of functions to execute a different program



There are many different forms of the **exec()** function call

```
char *name[2];  
name[0] = "../hello";  
name[1] = NULL;  
execve(name[0], name, NULL);
```

This will invoke a program called `hello`

Fork() and Exec()

```
int main(void) {  
    int pid;  
  
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
  
        char *name[2];  
        name[0] = "./hello";  
        name[1] = NULL;  
        execve(name[0], name, NULL);  
  
        _exit(0);  
    }  
    sleep(1);  
    printf("I'm the parent. My child has pid %d\n", pid);  
  
    return 0;  
}
```

Fork() and Exec()

```
int main(void) {  
    int pid;
```

```
    pid = fork();  
    if (0 == pid) {
```

```
        // I'm the child
```

```
        char *name[2];  
        name[0] = "./hello";  
        name[1] = NULL;  
        execve(name[0], name, NULL);
```

```
        _exit(0);
```

```
    }
```

```
    sleep(1);  
    printf("I'm the parent. My child has pid %d\n", pid);
```

```
    return 0;
```

```
}
```

Child code

Parent code

Fork() and Exec()

```
int main(void) {  
    int pid;  
  
    pid = fork();  
    if (0 == pid) {  
        // I'm the child  
  
        char *name[2];  
        name[0] = "./hello";  
        name[1] = NULL;  
        execve(name[0], name, NULL);  
  
        _exit(0);  
    }  
    sleep(1);  
    printf("I'm the parent. My child has pid %d\n", pid);  
  
    return 0;  
}
```

output

```
[01/25/23]seed@VM:~$ ./forkexec  
Hello from the C program!  
I'm the parent. My child has pid 33578
```

Demo?

forkandexec.c

Tl;dr


The programs we run get turned into a **process**

fork() is used to create a new process

- The parent process is typically the shell/terminal, and waits for the child process to finish
- The child process runs **exec()** to run our program

Contents	
9.4	Process Primitives
9.4.1	Having Children
9.4.2	Watching Your Children Die
9.4.3	Running New Programs
9.4.4	A Bit of History: vfork()
9.4.5	Killing Yourself
9.4.6	Killing Others
9.4.7	Dumping Core
9.5	Simple Children

you can kill children with the `kill()` function or `kill` command



```
#include <sys/types.h>
#include <unistd.h>
```

```
int main()
{
    while(1) {
        fork();
    }
    return 0;
}
```

Any ideas what might happen?

```
#include <sys/types.h>
#include <unistd.h>
```

```
int main()
{
    while(1) {
        fork();
    }
    return 0;
}
```



“Oh, these forks() aren’t homemade. They were made in factory. A **fork()** **bomb** factory. This is a **fork() bomb**”



A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

1. Executable Code

2. Associated Data

3. Execution Context/Bookkeeping information

(info that the OS needs to handle the process)

Main Memory

Process A Information
Process A Data
Process A Executable Code
Process B Information
Process B Data
Process B Executable Code

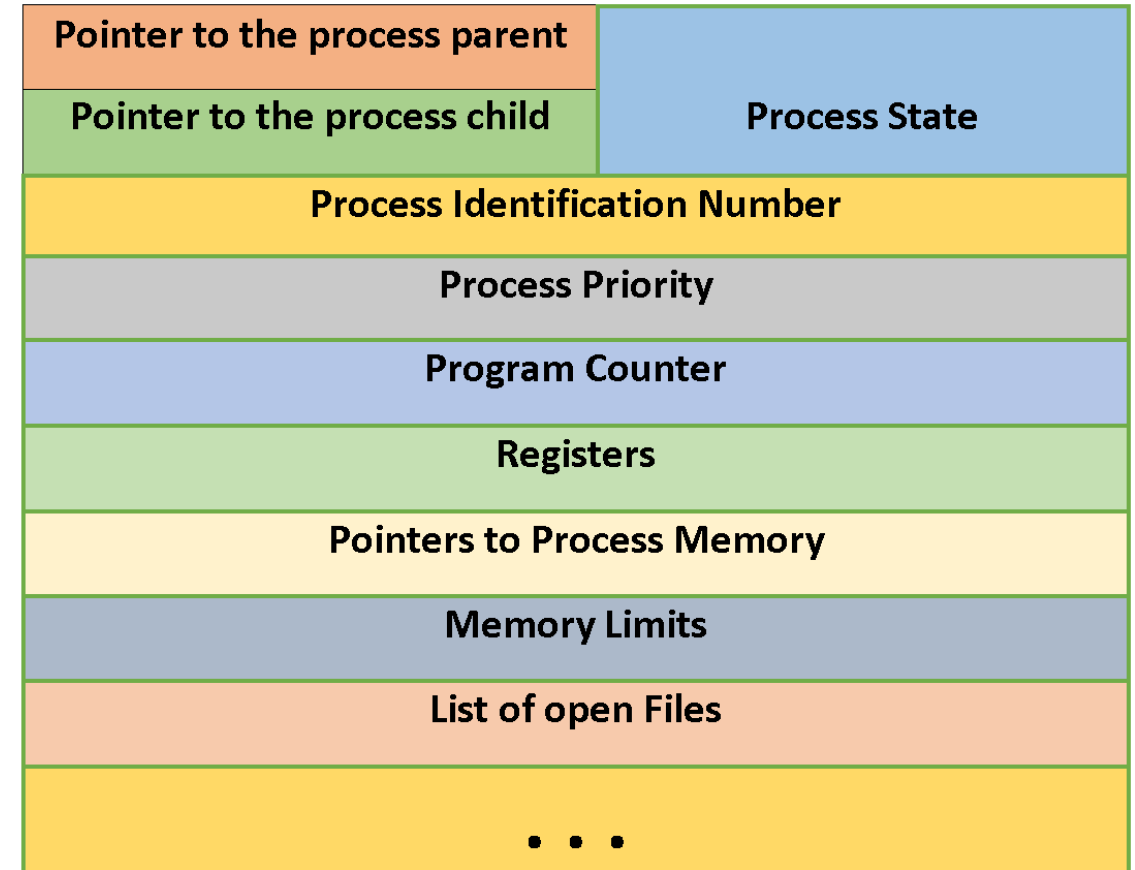
3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

Example PCB:



Created by NotesJam

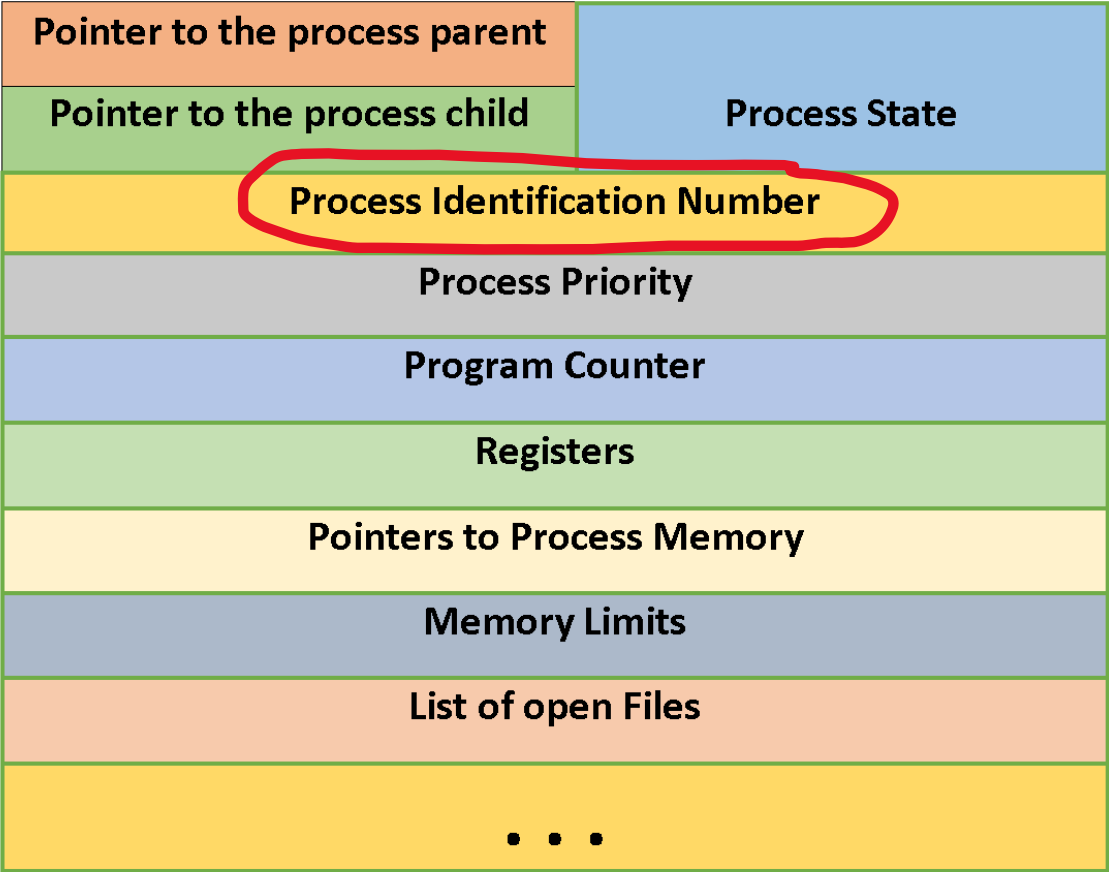
3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**
 - Simply just a data structure that holds information
 - The name of this varies by OS

Every process has a unique process ID (PID)

Process Name	User	% CPU	ID	Memory	Disk read toti D
at-spi2-registr	seed	0	1870	196.0 KiB	120.0 KiB
at-spi-bus-launcher	seed	0	1779	292.0 KiB	28.0 KiB
bash	seed	0	16245	1.6 MiB	3.1 MiB
bash	seed	0	20664	1.8 MiB	72.7 MiB
dbus-daemon	seed	0	1560	1.5 MiB	420.0 KiB

Example PCB:



We can use the PID to search for process, kill process, fork new process, etc

Created by Notes Jam

3. Execution Context/Bookkeeping information

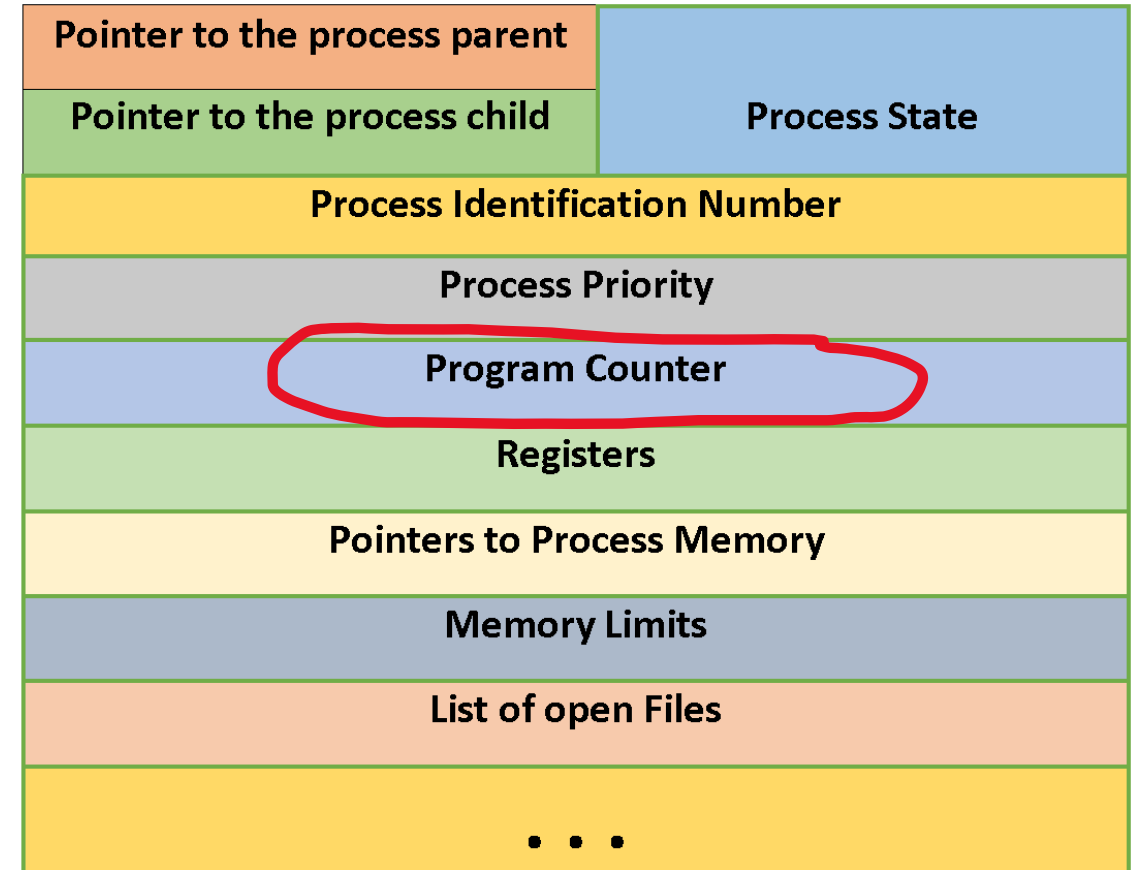
- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

Each process has a program counter (PC), which tells the CPU the next instruction to run in the process

Example PCB:



Created by NotesJam

3. Execution Context/Bookkeeping information

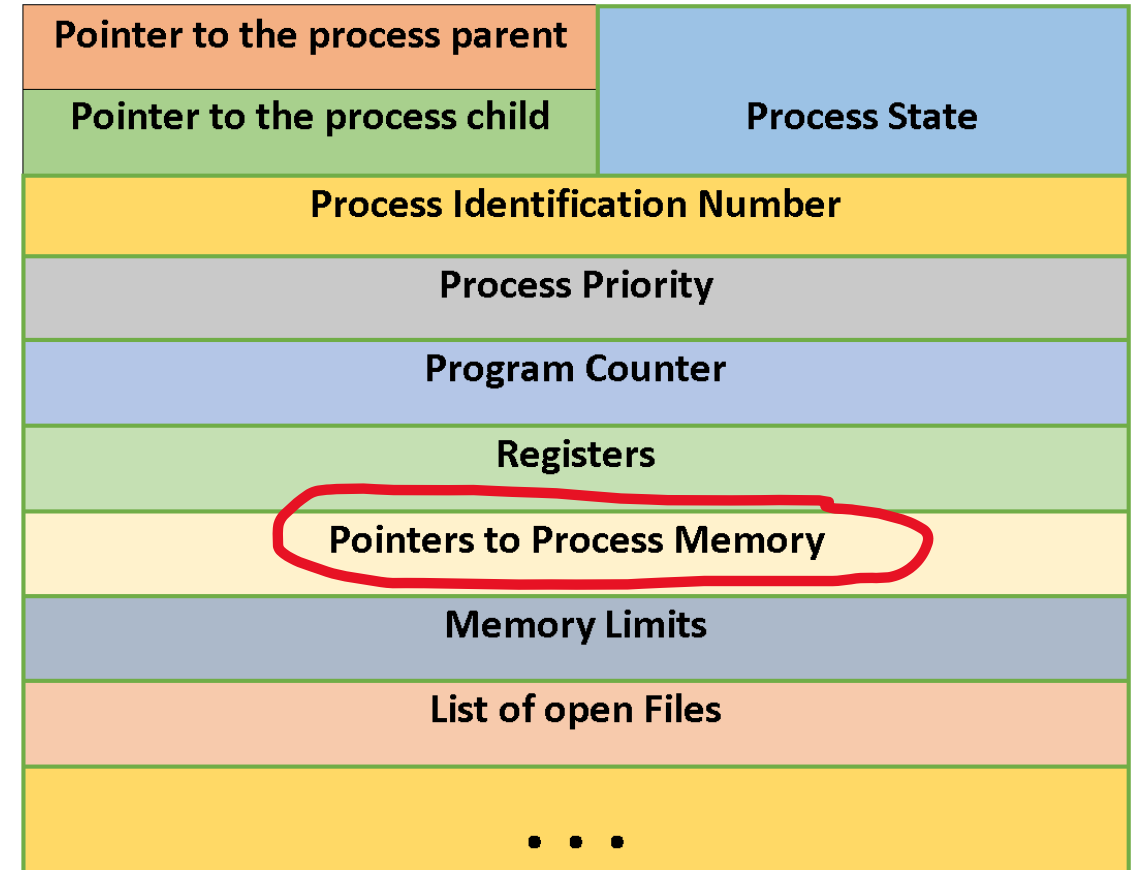
- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

PCB also maintains locations for the process Data and Code

Example PCB:



Created by NotesJam

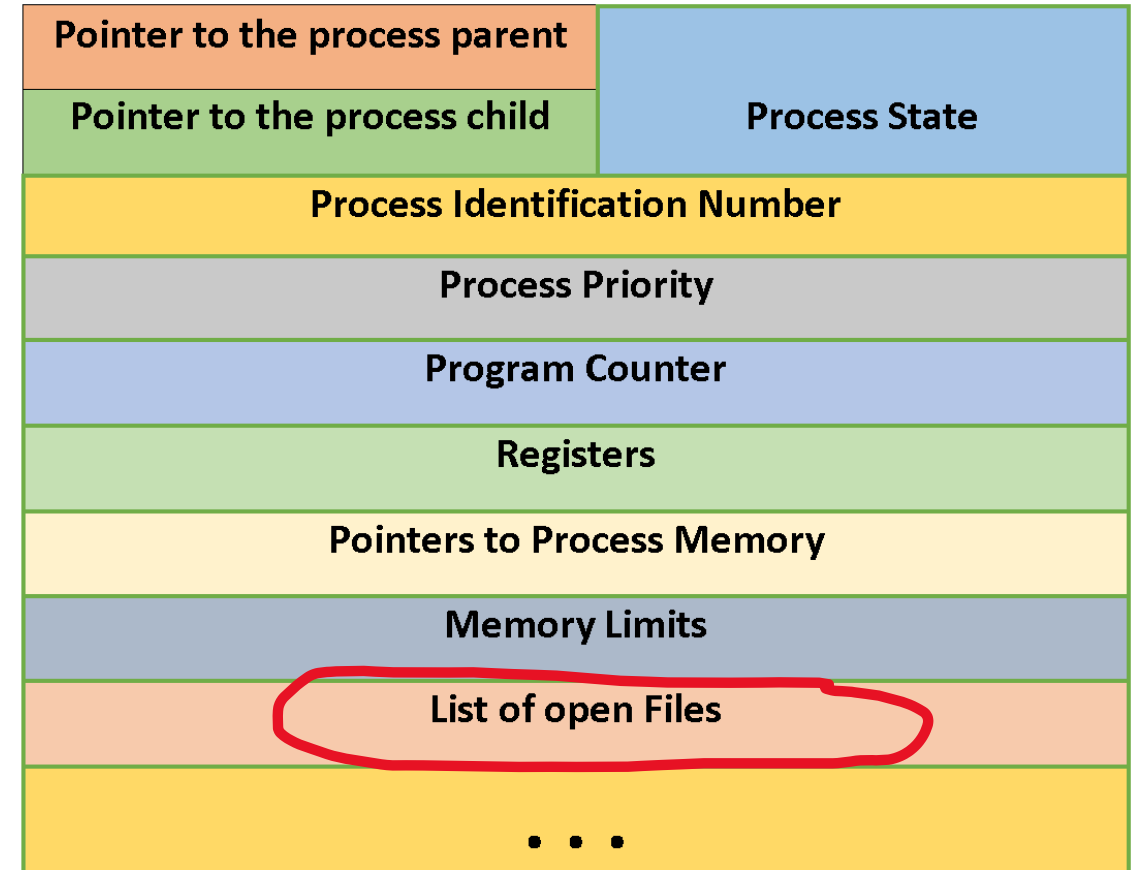
3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

Example PCB:



Created by NotesJam

3. Execution Context/Bookkeeping information

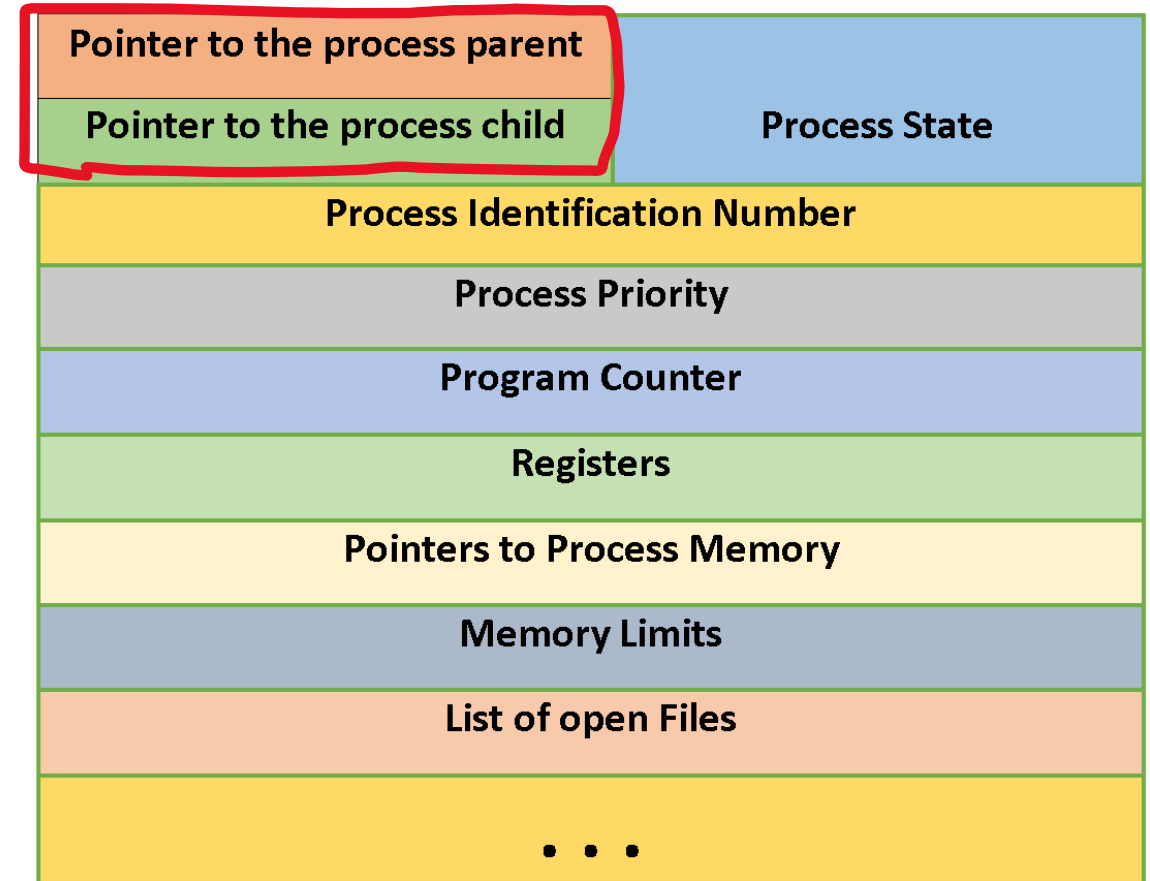
- Each process has a **Process Control Block (PCB)**

→ Simply just a data structure that holds information

→ The name of this varies by OS

PCB keeps track of who their parent is, and any child process (good parenting)

Example PCB:



Created by NotesJam

3. Execution Context/Bookkeeping information

- Each process has a **Process Control Block (PCB)**

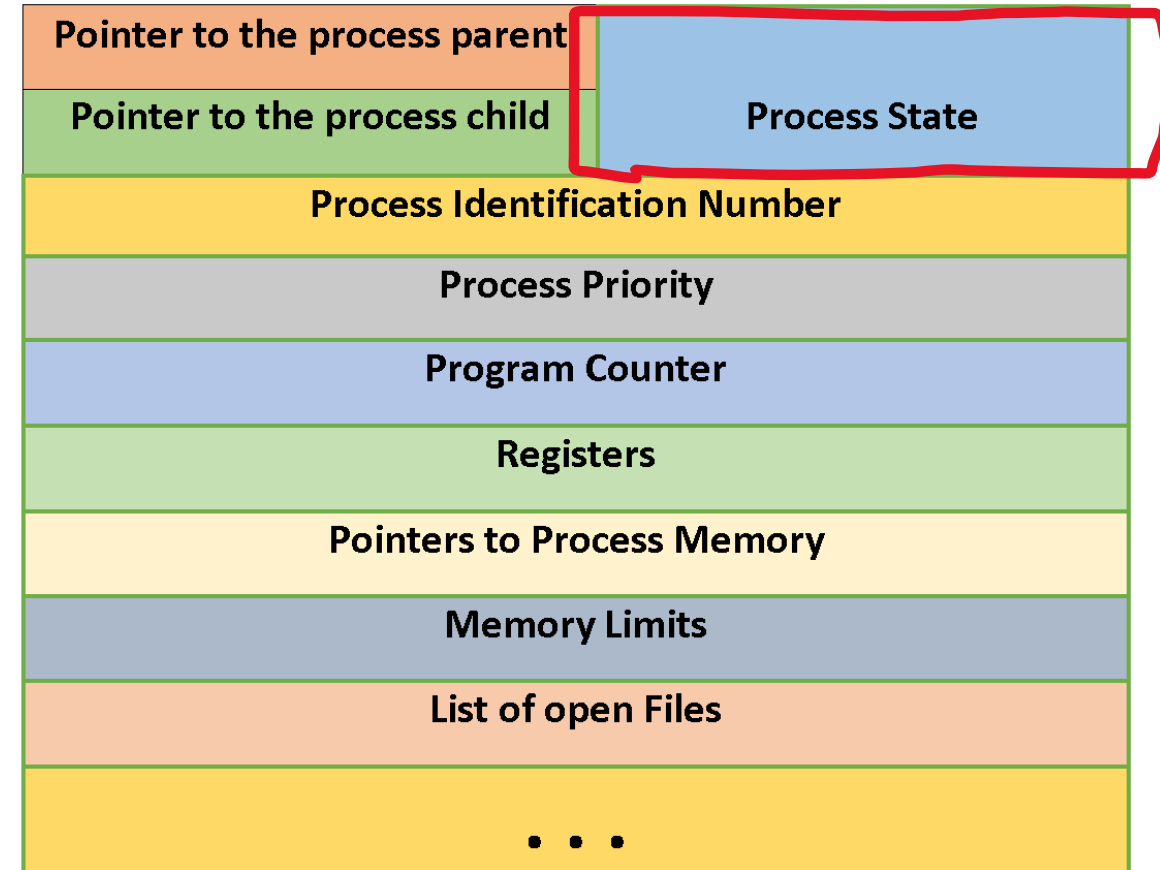
→ Simply just a data structure that holds information

→ The name of this varies by OS

A process goes through many **states**

- Active (running)**
- Blocked**
- Waiting**
- Suspended**

Example PCB:



Created by NotesJam

A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

1. Executable Code

2. Associated Data

3. Execution Context/Bookkeeping information

(info that the OS needs to handle the process)

We will talk about what goes here shortly

Pointer to the process parent	Process State
Pointer to the process child	
Process Identification Number	
Process Priority	
Program Counter	
Registers	
Pointers to Process Memory	
Memory Limits	
List of open Files	
...	

Created by NotesJam

The jobs of an Operating System

1. Process Manager

“The Coach”

The OS manages many active processes all at once, and they must create processes, manage current process, and control which processes do what



A **process** is an instance of a running program on a computer

All processes have the following data while they are running:

1. Executable Code

2. Associated Data


3. Execution Context/Bookkeeping information

(info that the OS needs to handle the process)

`./hello_world`

Fork() and exec()

Program is now
running as a
process



Pointer to the process parent	Process State
Pointer to the process child	
Process Identification Number	
Process Priority	
Program Counter	
Registers	
Pointers to Process Memory	
Memory Limits	
List of open Files	
...	

Created by NotesJam

Demo time!

```
int main(void) {
    int pid;

    pid = fork();
    if (0 == pid) {
        // I'm the child
        printf("Hi, I'm the child. \n");
    }

    sleep(1);
    // we could wait() here
    printf("I'm the parent.");

    return 0;
}
```

```
int main(void) {
    int pid;

    pid = fork();
    if (0 == pid) {
        // I'm the child

        char *name[2];
        name[0] = "./hello";
        name[1] = NULL;
        execve(name[0], name, NULL);

        _exit(0);
    }
    sleep(1);
    printf("I'm the parent. My child

    return 0;
}
```

The jobs of an Operating System

1. Process Manager
“The Coach”

2. Interface Manager
“The Bouncer”

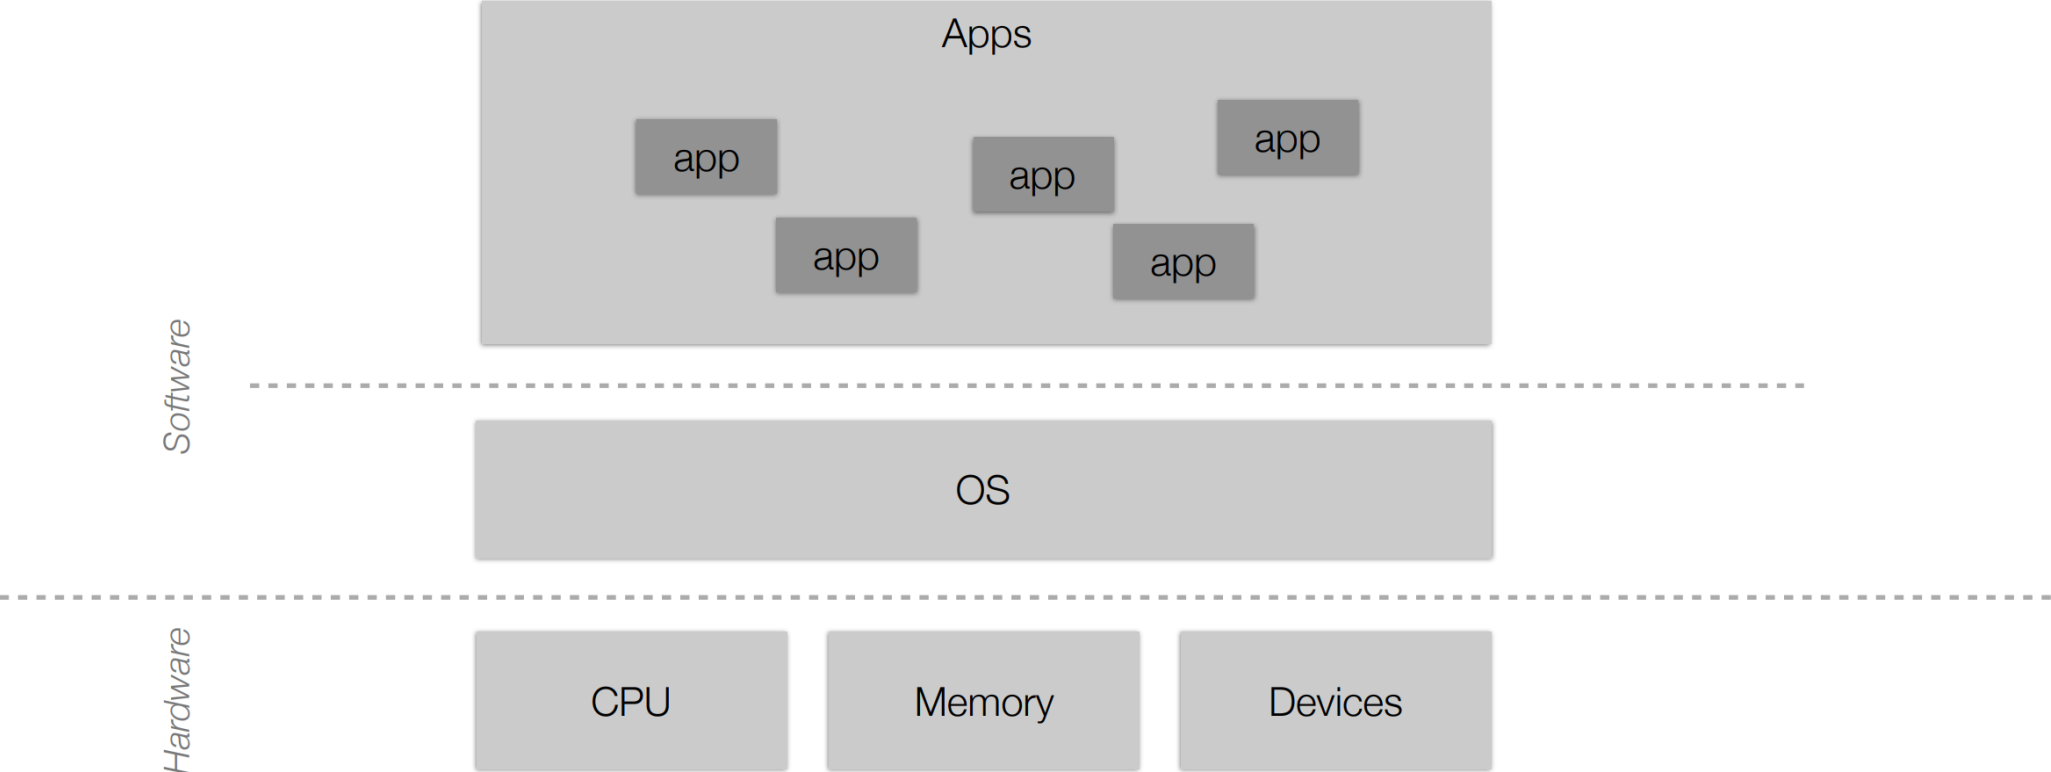
3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”



Operating Systems Review

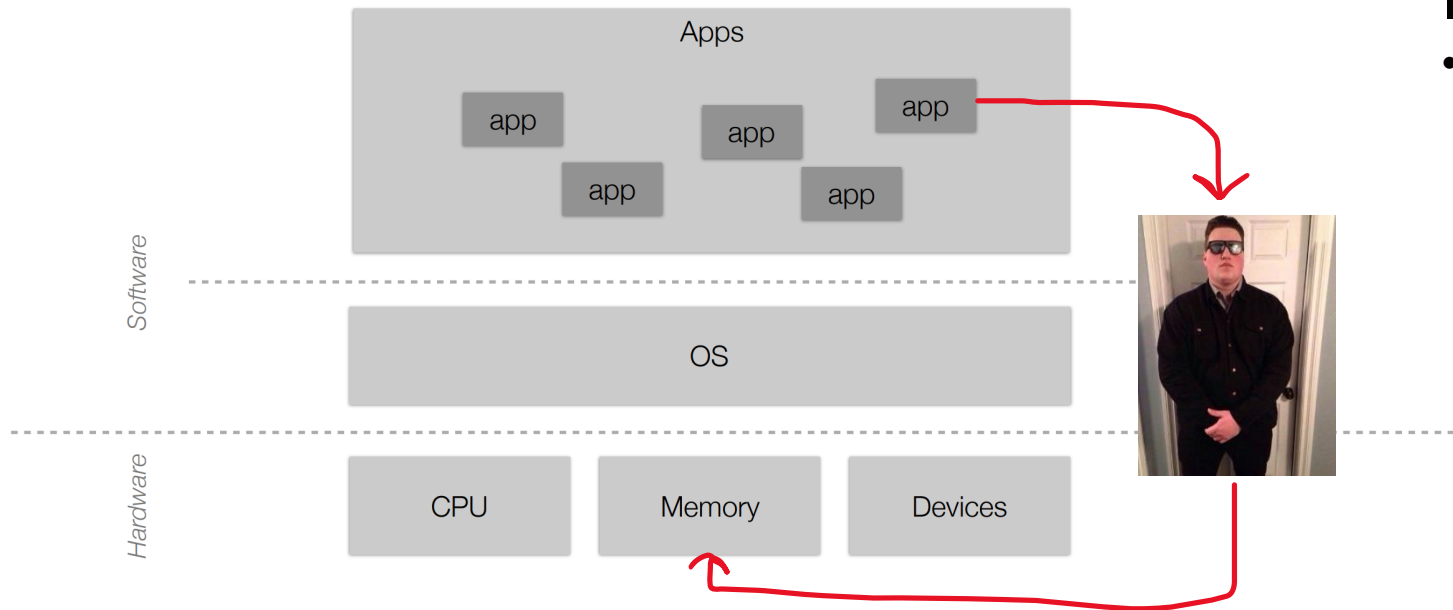


Operating Systems Review

Responsibilities of the OS?

Interface Manager

- Manages communication between apps and hardware



Operating Systems Review

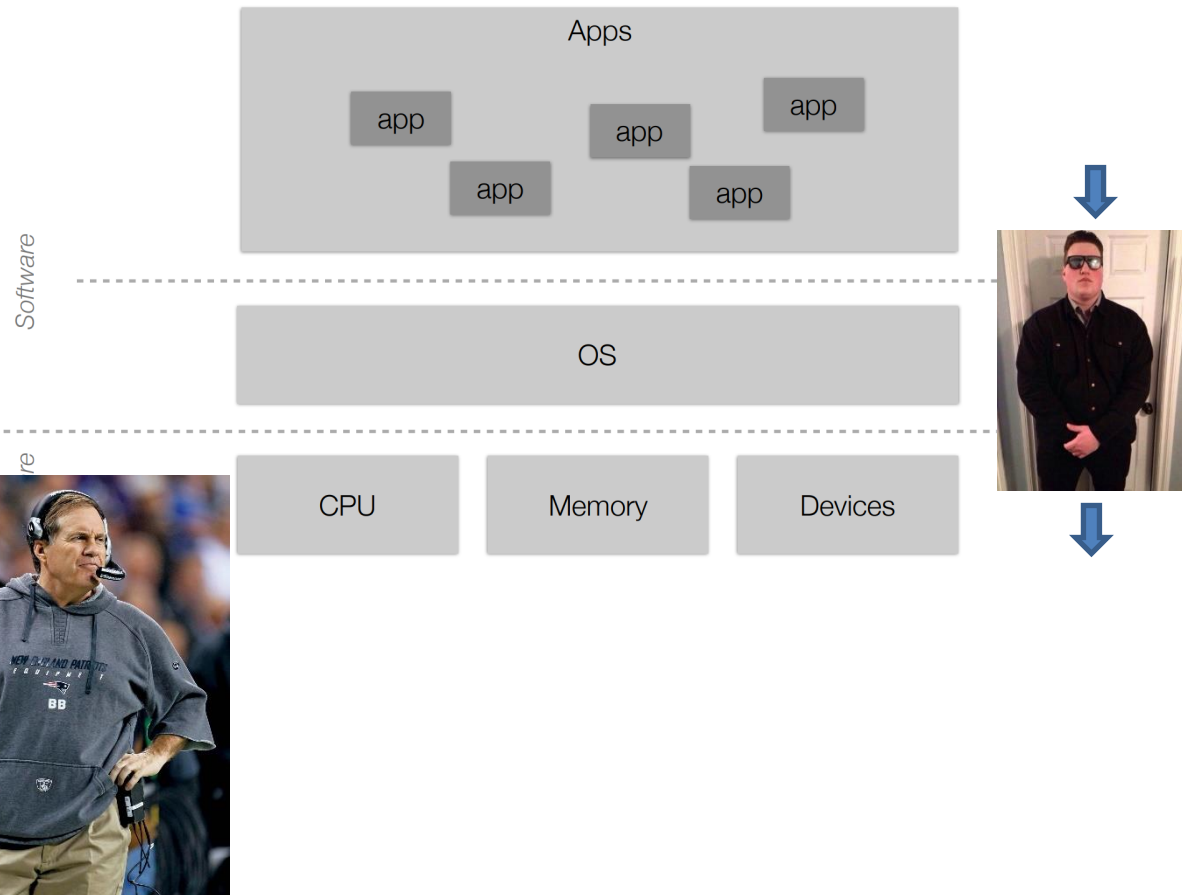
Responsibilities of the OS?

Interface Manager

- Manages communication between apps and hardware

Process Manager

- Manages how processes are structured and how to handle many processes running at once



Operating Systems Review

Responsibilities of the OS?

Interface Manager

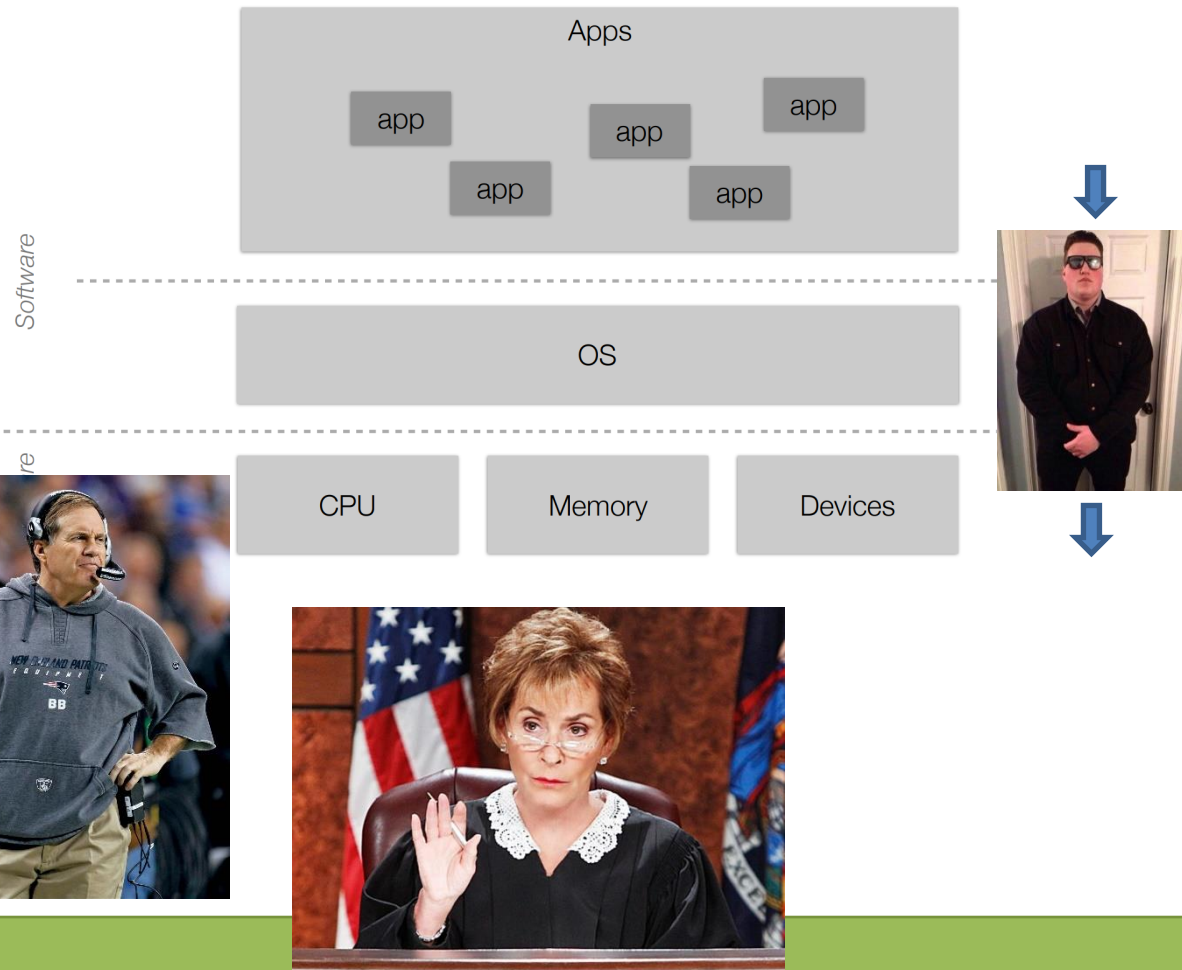
- Manages communication between apps and hardware

Process Manager

- Manages how processes are structured and how to handle many processes running at once

Traffic Manager

- Manages which programs should be executed by the CPU



Operating Systems Review

Responsibilities of the OS?

Interface Manager

- Manages communication between apps and hardware

Process Manager


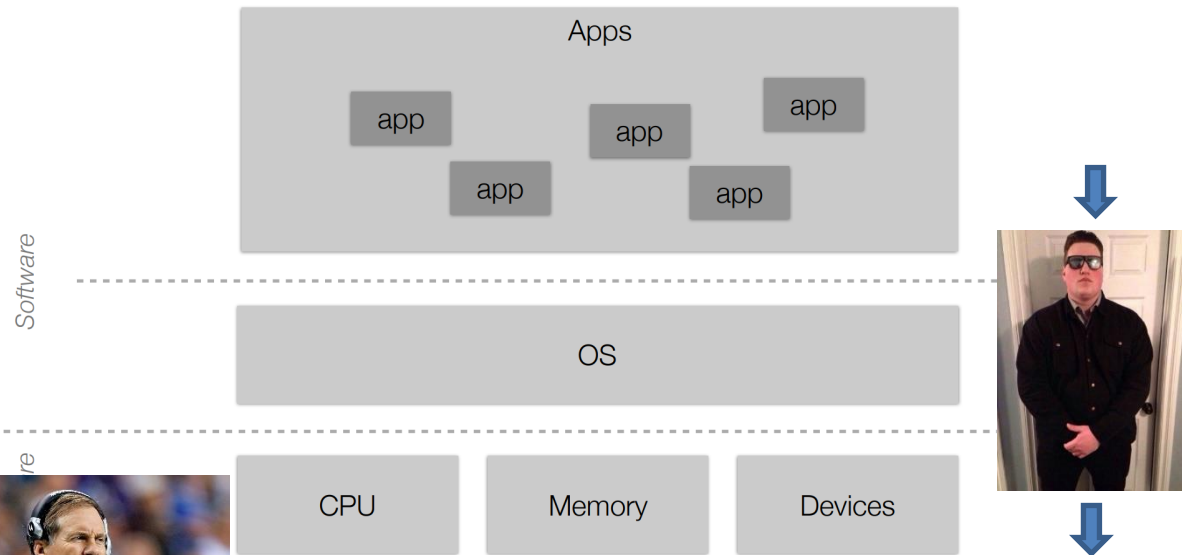
- Manages how processes are structured and how to handle many processes running at once

Traffic Manager

- Manages which programs should be executed by the CPU

Memory Manager

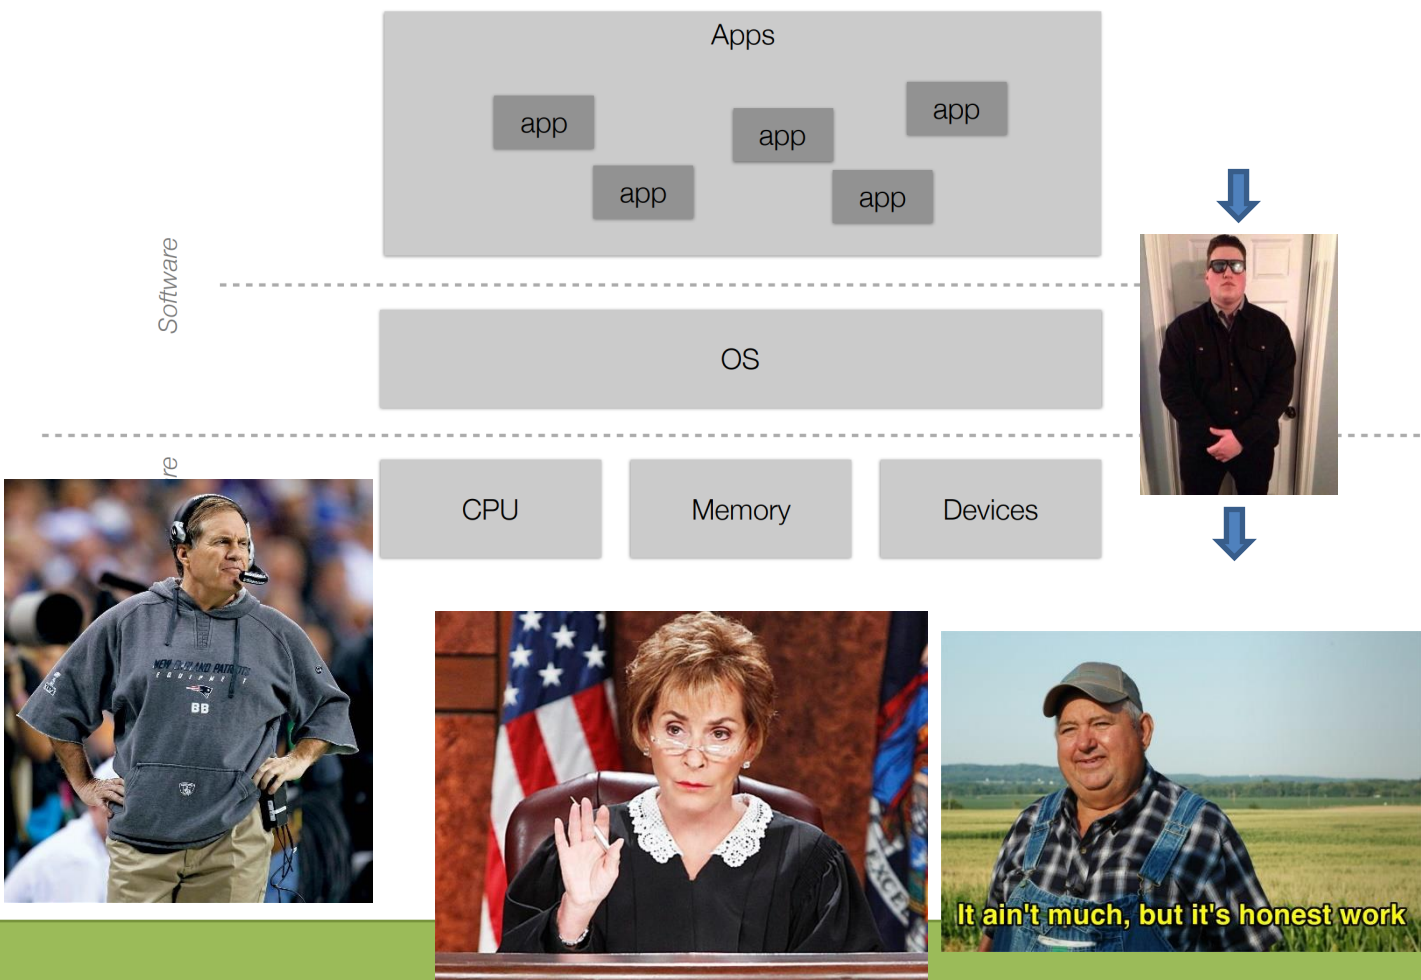
- Manages how physical memory is utilized



It ain't much, but it's honest work

Operating Systems Review

Responsibilities of the OS?



Interfacing

- Managing hardware resources

Providing

- Managing memory

Traffic

- Managing the flow of data

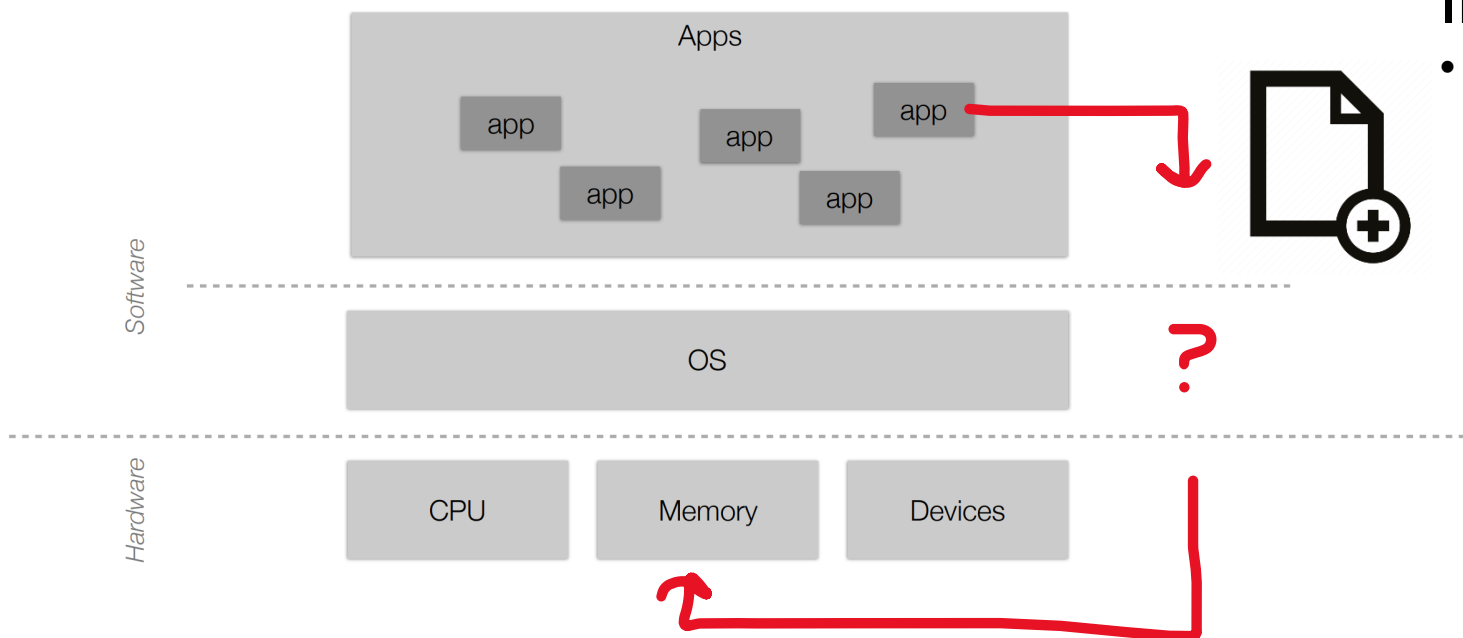
Memory

- Managing memory allocation



Operating Systems Review

Responsibilities of the OS?



Interface Manager

- Manages communication between apps and hardware

How does an application get access to a computer's resources?



Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
int main(void)
{
    printf("Hello, World!\n");

    return 0;
}
```

Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources


```
int main(void)
{
    printf("Hello, World!\n");

    return 0;
}
```



```
int main(void)
{
    write(1, "Hello, World!\n", 14);

    return 0;
}
```



```
int main(void)
{
    syscall(SYS_write, 1, "Hello, World!\n", 14);

    return 0;
}
```


Number	Name	Description
1	exit	terminate process execution
2	fork	fork a child process
3	read	read data from a file or socket
4	write	write data to a file or socket
5	open	open a file or socket
6	close	close a file or socket
37	kill	send a kill signal
90	old_mmap	map memory
91	munmap	unmap memory
301	socket	create a socket
303	connect	connect a socket

Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

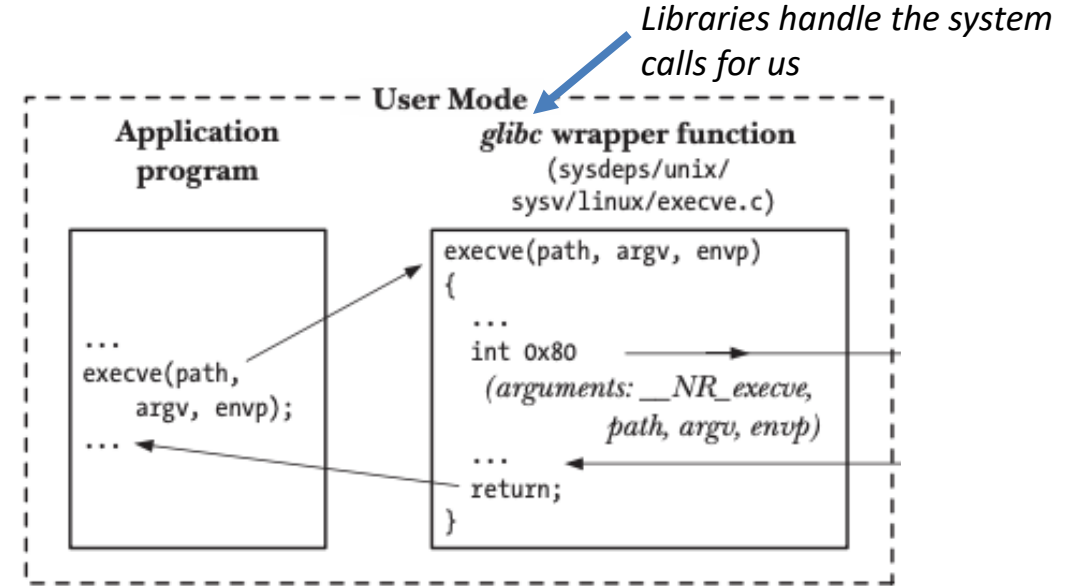


Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

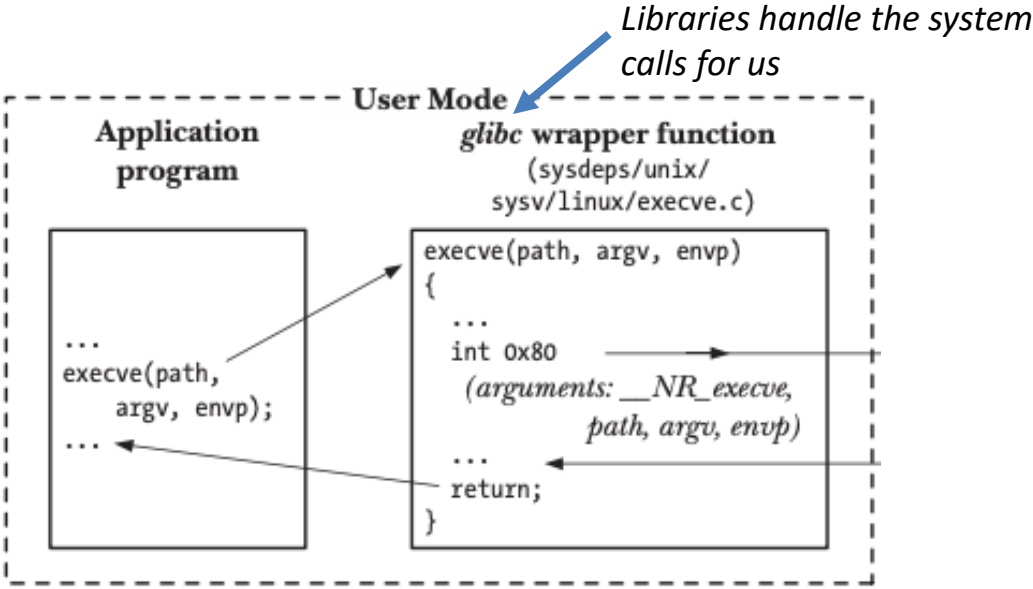


Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```



The operating system have hundreds of different syscalls, and different syscalls have different parameters, we need a way to distinguish them


- EAX
- EBX
- ECX
- EDX

Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```



EAX	System Call Number	
EBX	Address of “/bin/bc”	
ECX	0 or 1	Environment variables
EDX	INT 0x80	send trap to kernel and invoke the syscall

The operating system have hundreds of different syscalls, and different syscalls have different parameters, we need a way to distinguish them

The OS will look at the values at certain registers!

Syscalls


Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

Demo:
bc.c

- EAX** System Call Number
- EBX** Address of “/bin/bc”
- ECX** 0 or 1 Environment variables
- EDX** INT 0x80 send trap to kernel and invoke the syscall




normies seeing
calculator open on
its own

programmers
seeing calculator
open on its own




Syscalls



All applications run in user mode.

The code has no ability to directly access hardware
Code running in user mode must use API/syscalls to access hardware and memory

Syscalls




All applications run in user mode.

The code has no ability to directly access hardware
Code running in user mode must use API/syscalls to access hardware and memory

Code running in kernel-mode has complete, unrestricted access to computer resources

Reserved for the lowest-level trusted functions of the operating system

Syscalls



The collective functionality and services of the OS that manages the computer and its resources is called the **kernel**

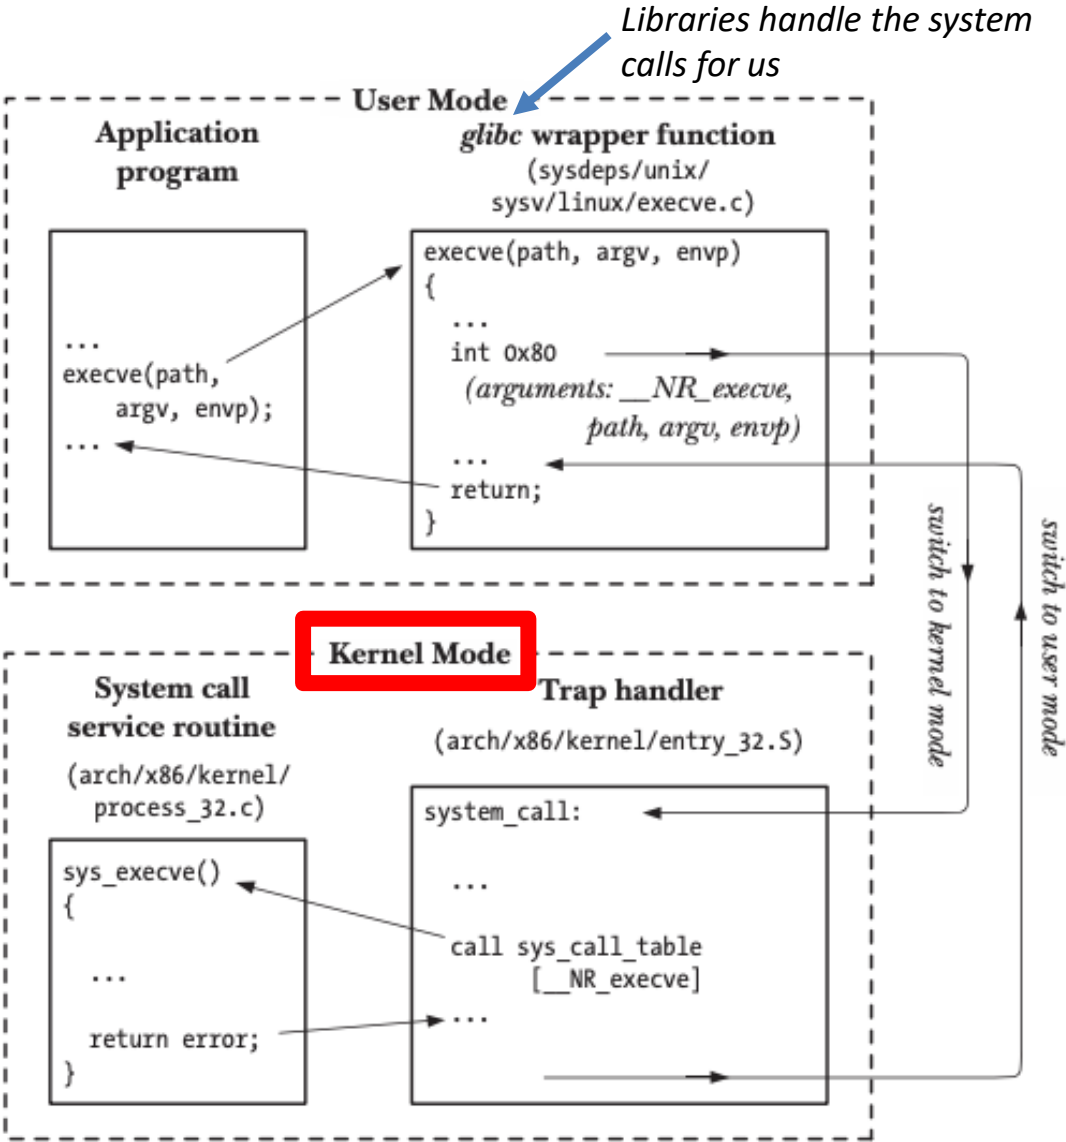
Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

- EAX** System Call Number
- EBX** Address of “/bin/bc”
- ECX** 0 or 1 Environment variables
- EDX** INT 0x80 send trap to kernel and invoke the syscall



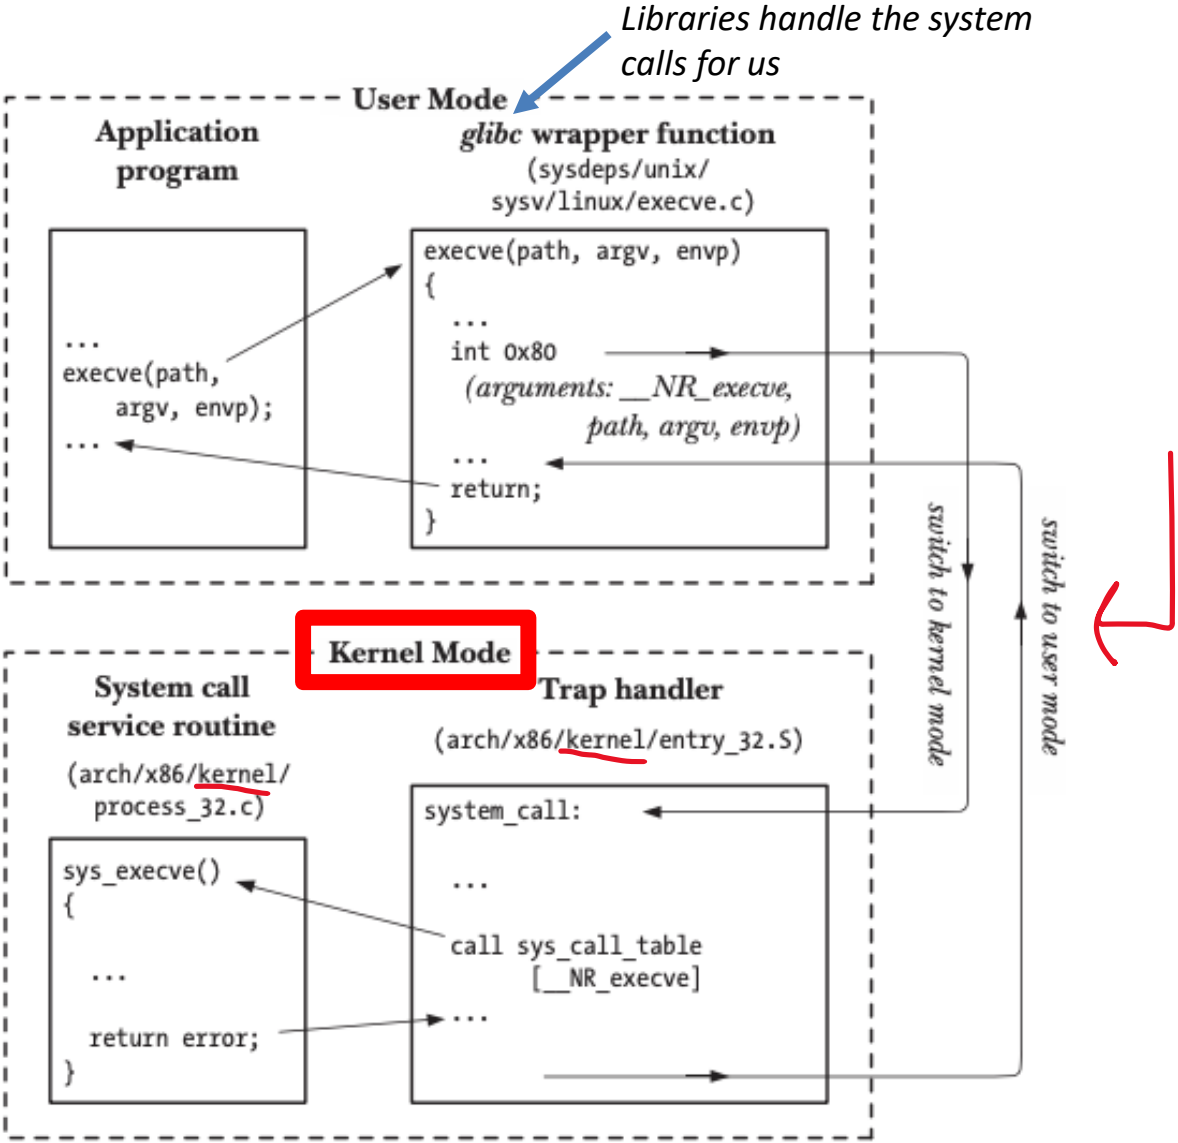
System calls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    char *name[2];
    name[0] = "/bin/bc";
    name[1] = NULL;
    execve(name[0], name, NULL);
    return 0;
} syscall
```

- EAX** System Call Number
- EBX** Address of “/bin/bc”
- ECX** 0 or 1 Environment variables
- EDX** INT 0x80 send trap to kernel and invoke the syscall



Syscalls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources

NR	syscall name	references	%eax	arg0 (%ebx)	arg1 (%ecx)	arg2 (%edx)	arg3 (%esi)	arg4 (%edi)	arg5 (%ebp)
0	restart_syscall	man/ cs/	0x00	-	-	-	-	-	-
1	exit	man/ cs/	0x01	int error_code	-	-	-	-	-
2	fork	man/ cs/	0x02	-	-	-	-	-	-
3	read	man/ cs/	0x03	unsigned int fd	char *buf	size_t count	-	-	-
4	write	man/ cs/	0x04	unsigned int fd	const char *buf	size_t count	-	-	-
5	open	man/ cs/	0x05	const char *filename	int flags	umode_t mode	-	-	-
6	close	man/ cs/	0x06	unsigned int fd	-	-	-	-	-
7	waitpid	man/ cs/	0x07	pid_t pid	int *stat_addr	int options	-	-	-
8	creat	man/ cs/	0x08	const char *pathname	umode_t mode	-	-	-	-
9	link	man/ cs/	0x09	const char *oldname	const char *newname	-	-	-	-
10	unlink	man/ cs/	0x0a	const char *pathname	-	-	-	-	-
11	execve	man/ cs/	0x0b	const char *filename	const char *const *argv	const char *const *envp	-	-	-
12	chdir	man/ cs/	0x0c	const char *filename	-	-	-	-	-

EDX

INT 0x80

send trap to kernel and
invoke the syscall

System calls

Applications evoke operating system defined functions, or **system calls (syscalls)**, to access computing resources


NR	syscall name	references	%eax	arg0 (%ebx)	arg1 (%ecx)	arg2 (%edx)	arg3 (%esi)	arg4 (%edi)	arg5 (%ebp)
0	restart_syscall	man/ cs/	0x00	-	-	-	-	-	-
1	exit	man/ cs/	0x01	int error_code	-	-	-	-	-
2	fork	man/ cs/	0x02	-	-	-	-	-	-
3	read	man/ cs/	0x03	unsigned int fd	char *buf	size_t count	-	-	-
4	https://chromium.googlesource.com/chromiumos/docs/+/_/master/constants/syscalls.md#x86-32_bit								
5									
6									
7									
8									
9									
10	unlink	man/ cs/	0x0a	const char *pathname	-	-	-	-	-
11	execve	man/ cs/	0x0b	const char *filename	const char *const *argv	const char *const *envp	-	-	-
12	chdir	man/ cs/	0x0c	const char *filename	-	-	-	-	-

EDX

INT 0x80

send trap to kernel and invoke the syscall

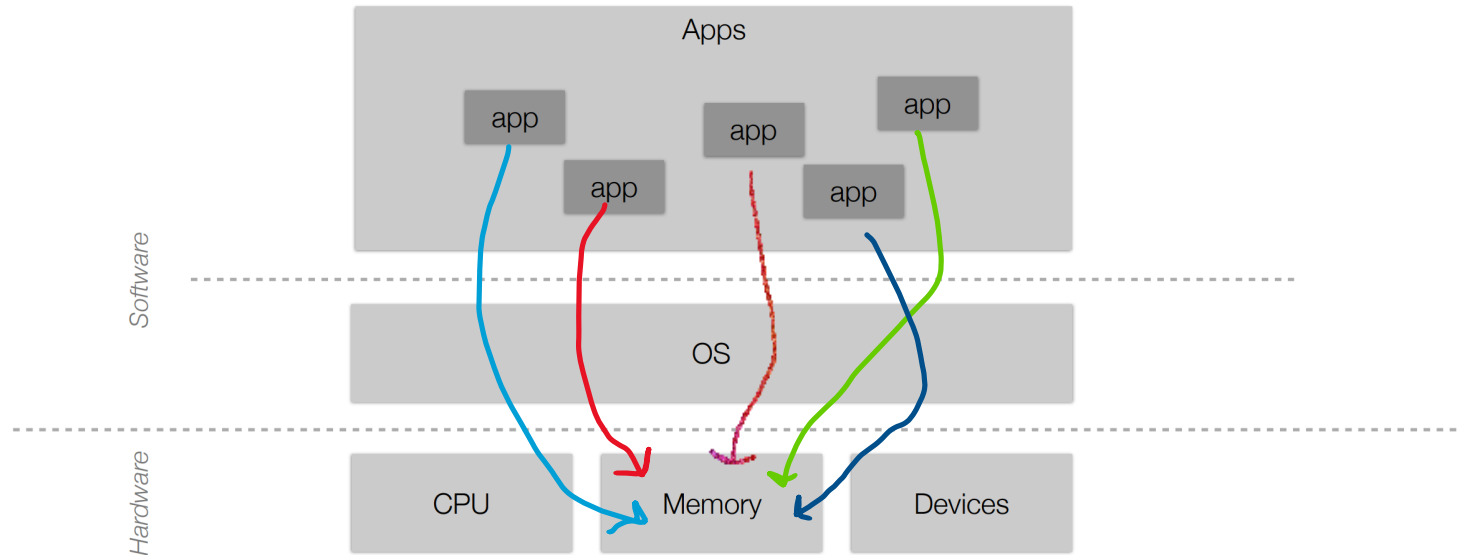
Syscalls



Applications Layout in Memory

Process Manager

- Manages how processes are structured and how to handle many processes running at once

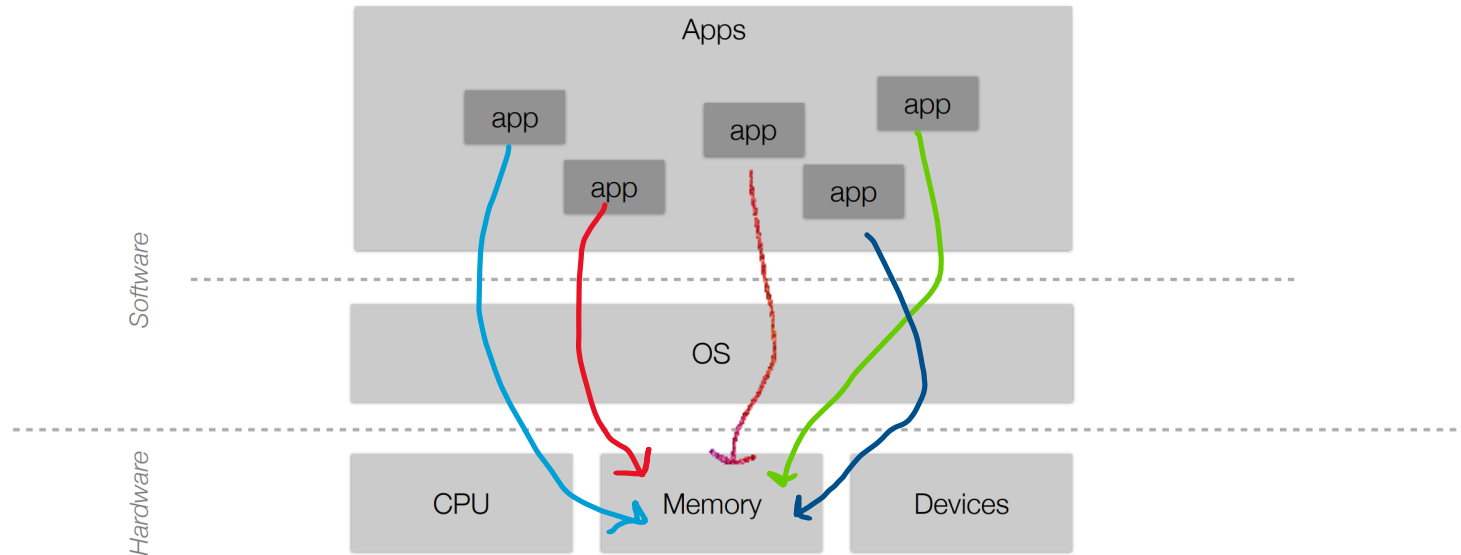


How does a **program** get loaded into memory?

Applications Layout in Memory

Process Manager

- Manages how processes are structured and how to handle many processes running at once



How does a **program** get loaded into memory?

An active program running on a computer is called a **process**

Applications Layout in Memory

What does this look like?


1. Executable Code

2. Associated Data

3. Execution Context/Bookkeeping information

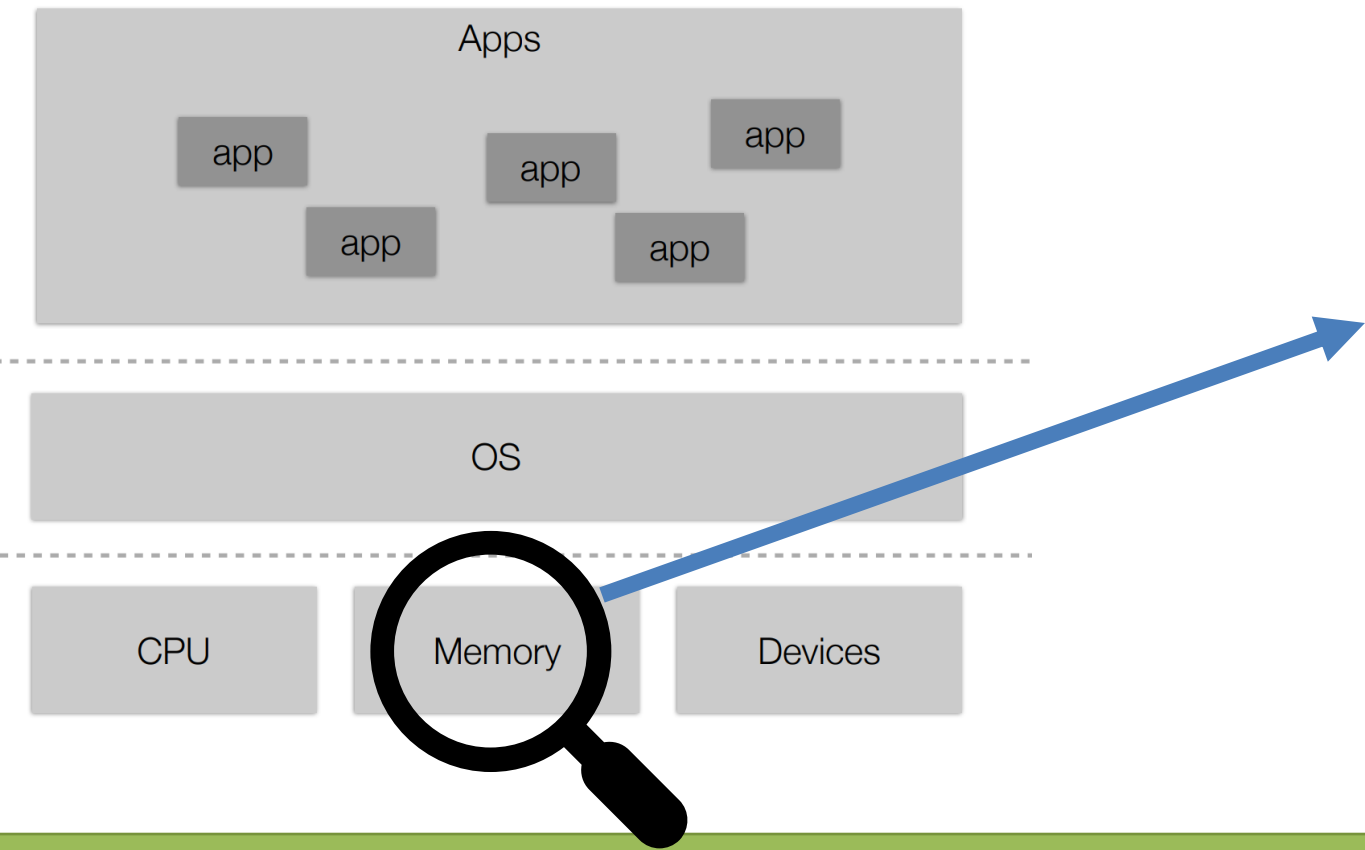
Process Manager

- Manages how processes are structured and how to handle many processes running at once




Applications Layout in Memory

What does a program look like in memory?



Process Manager

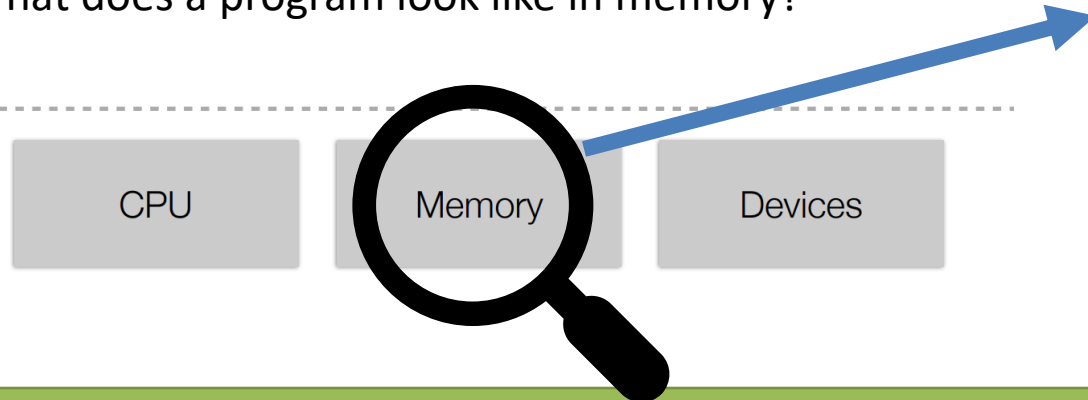
- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

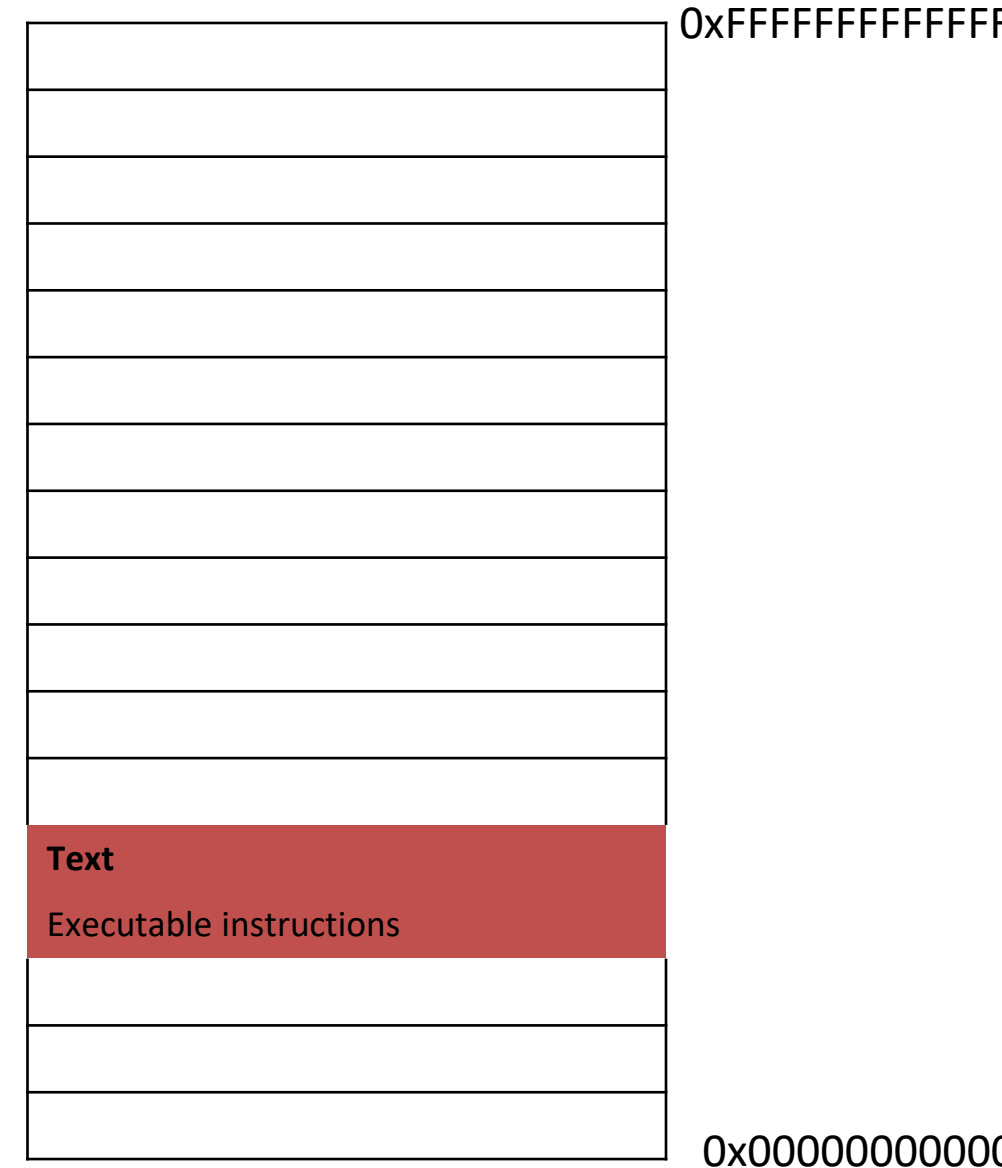
Text Segment- binary executable instructions for the process

What does a program look like in memory?



Process Manager


- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory


Data Segment- Static variables initialized by the programmer

What does a program look like in memory?



Process Manager


- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

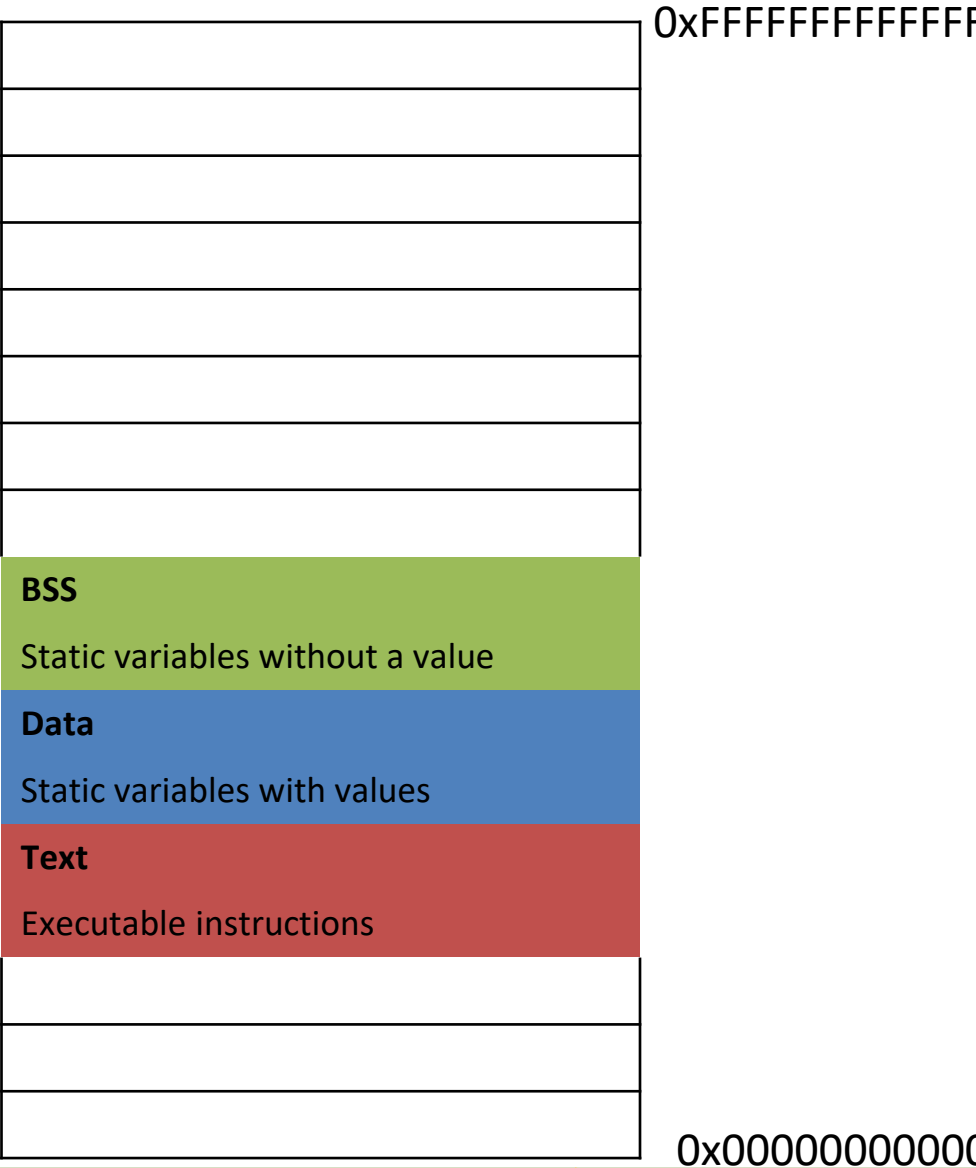
BSS Segment- contains statically allocated variables that are declared, but have not been assigned a value yet

What does a program look like in memory?



Process Manager

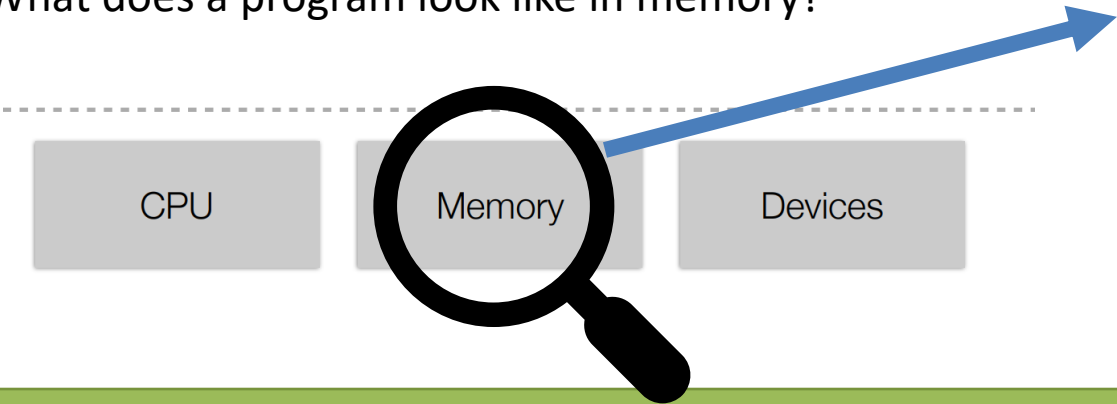
- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

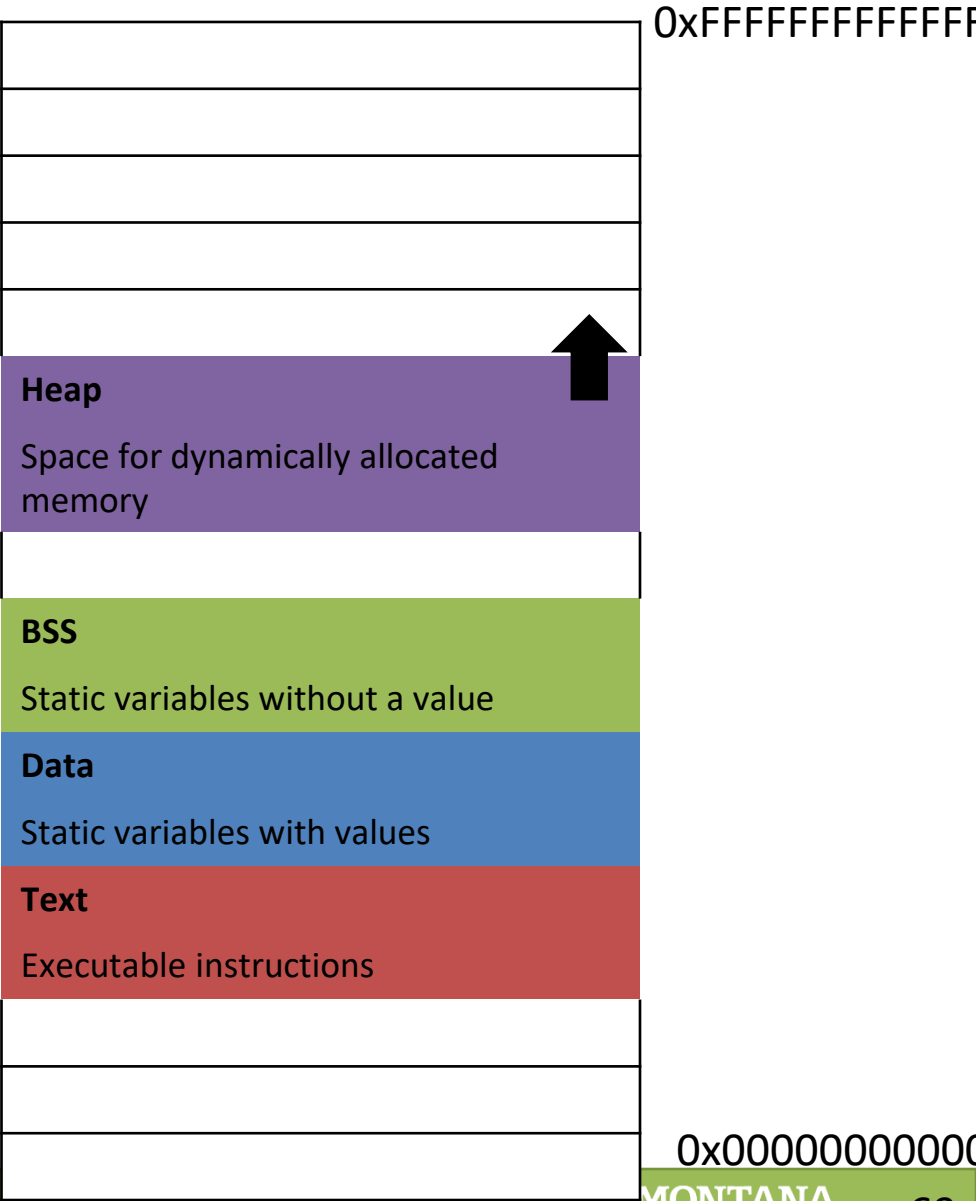
Heap- memory set aside for dynamic allocation (e.g. malloc). Grows “up” as more memory is allocated

What does a program look like in memory?



Process Manager


- Manages how processes are structured and how to handle many processes running at once



Applications Layout in Memory

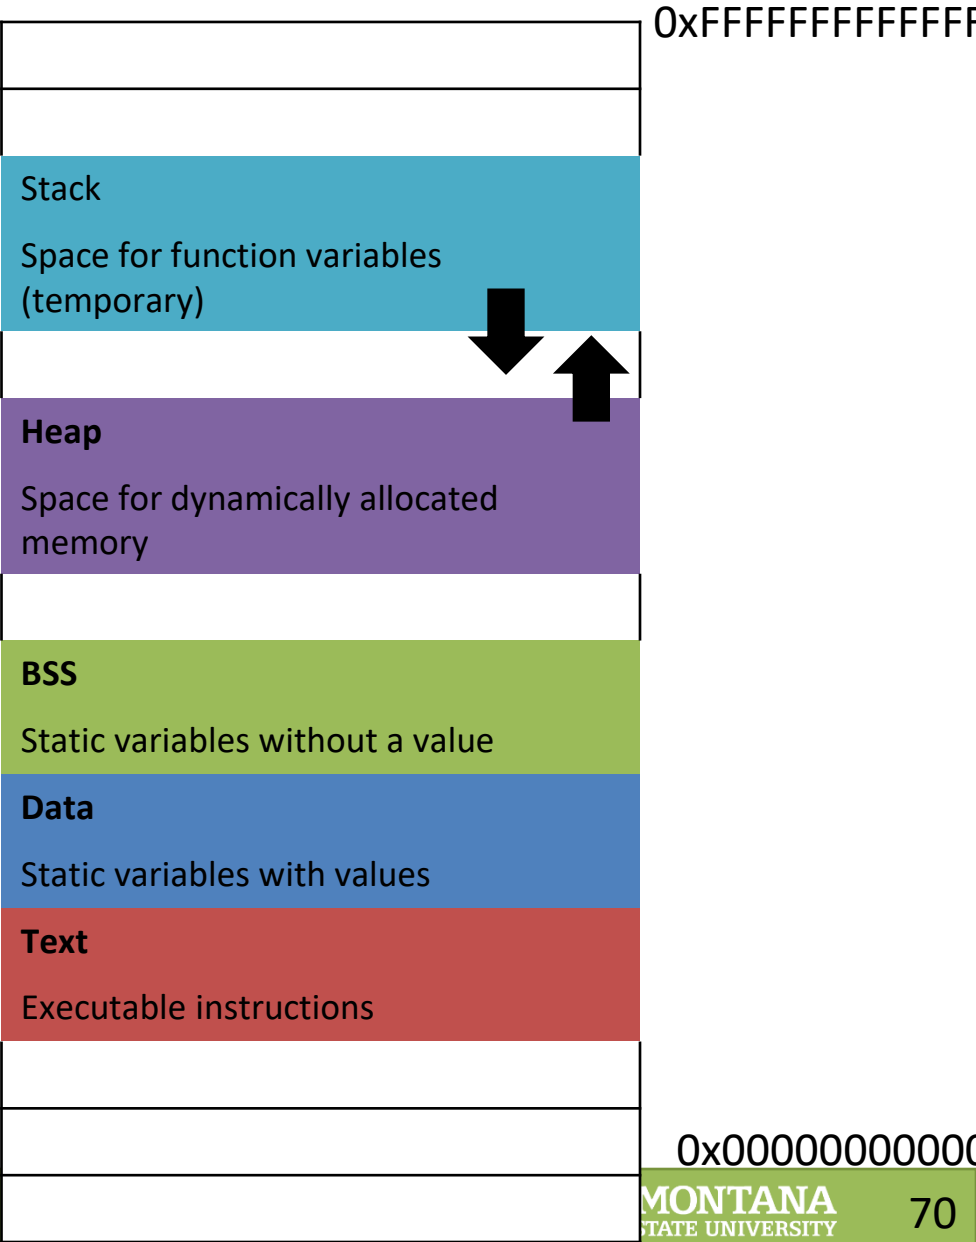
Stack – memory for storing function variables.
Grows “down” as additional functions are called

What does a program look like in memory?



Process Manager

- Manages how processes are structured and how to handle many processes running at once




Applications Layout in Memory

1. Executable Code

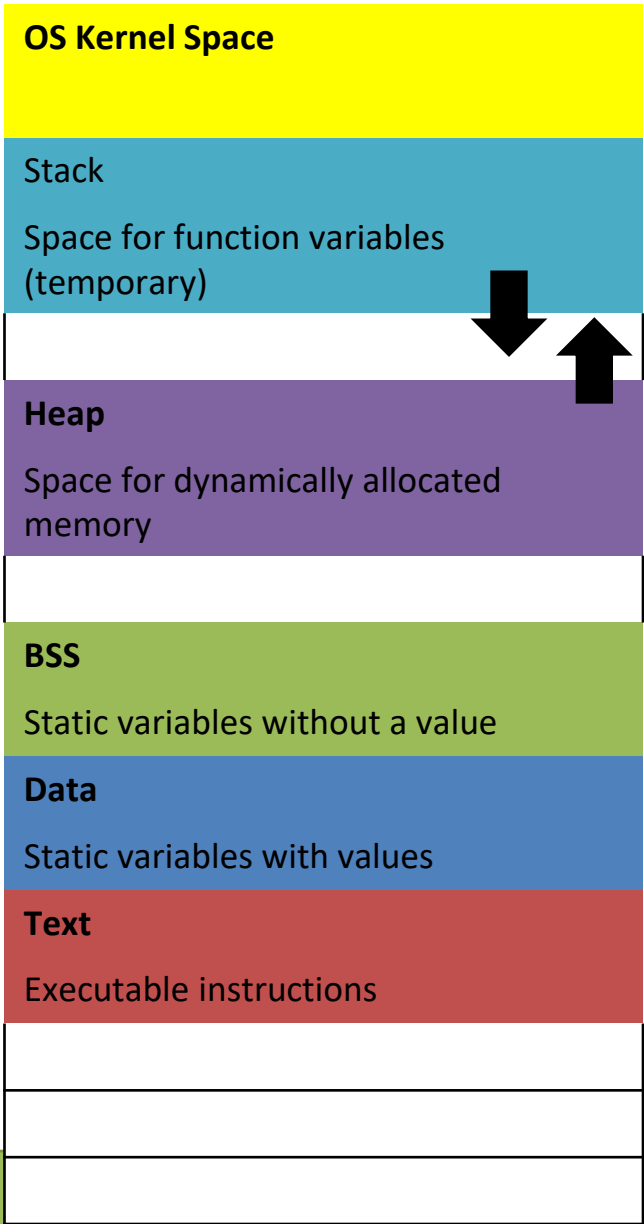
2. Associated Data

3. Execution Context/Bookkeeping information



Process Manager

- Manages how processes are structured and how to handle many processes running at once



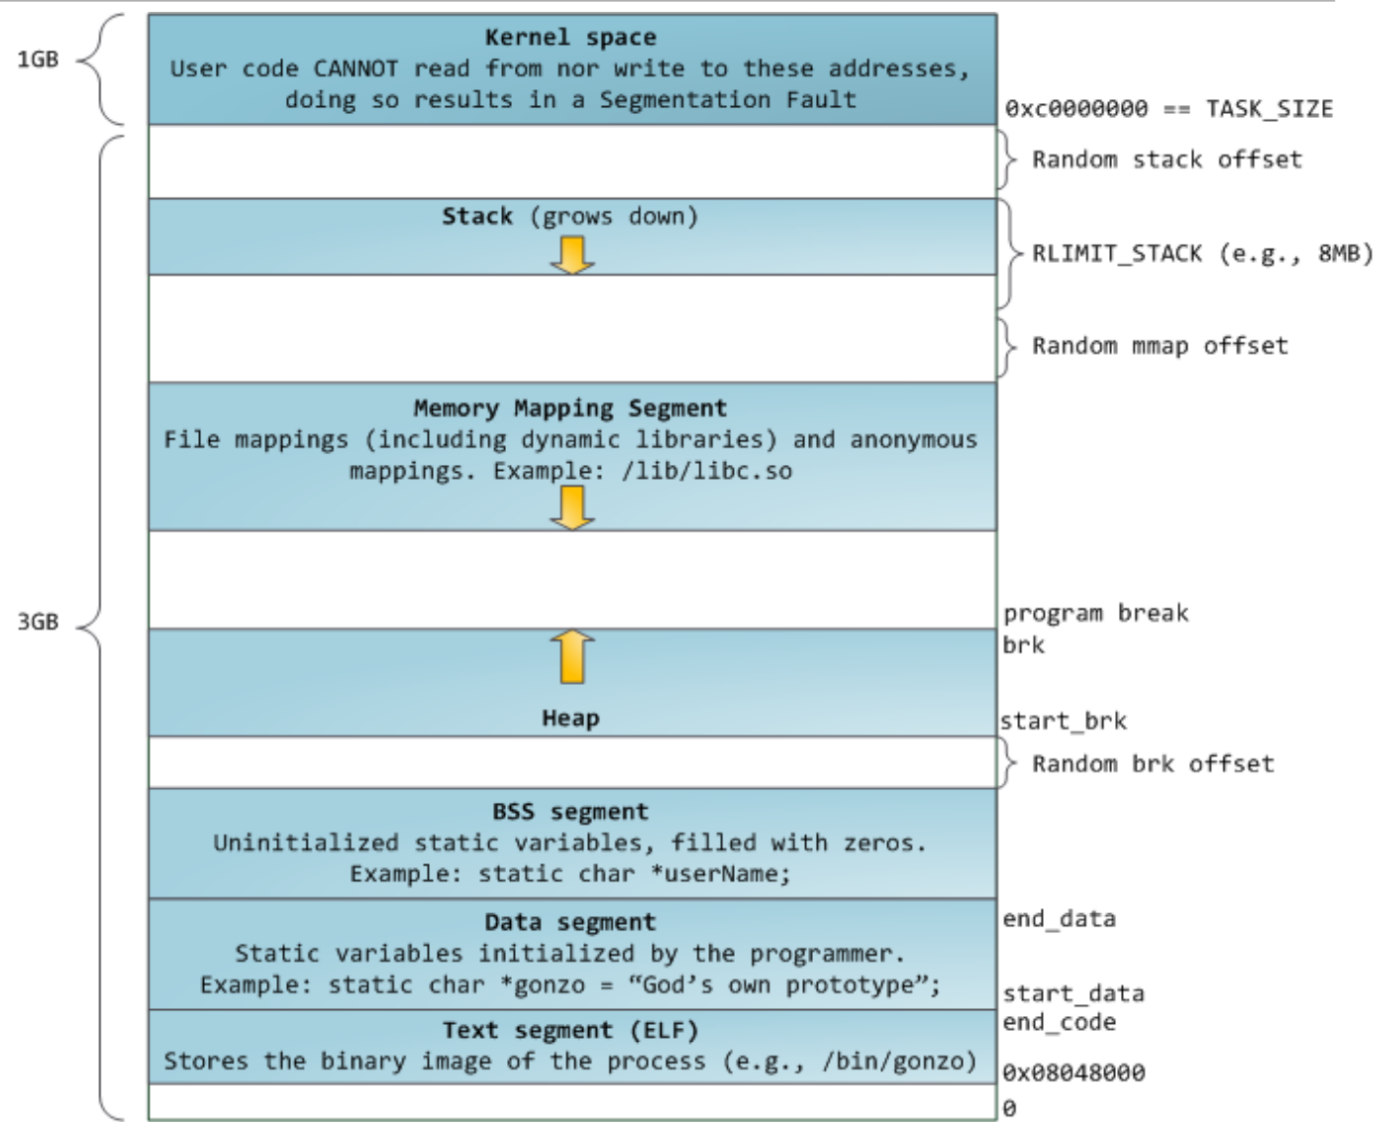
0xFFFFFFFFFFFFFFFF

0x0000000000000000

Applications Layout in Memory

Demo?

Makefile Demo



Applications Layout in Memory

Output of `pmap` (process mapping tool)

[illegible][illegible]

Applications Layout in Memory

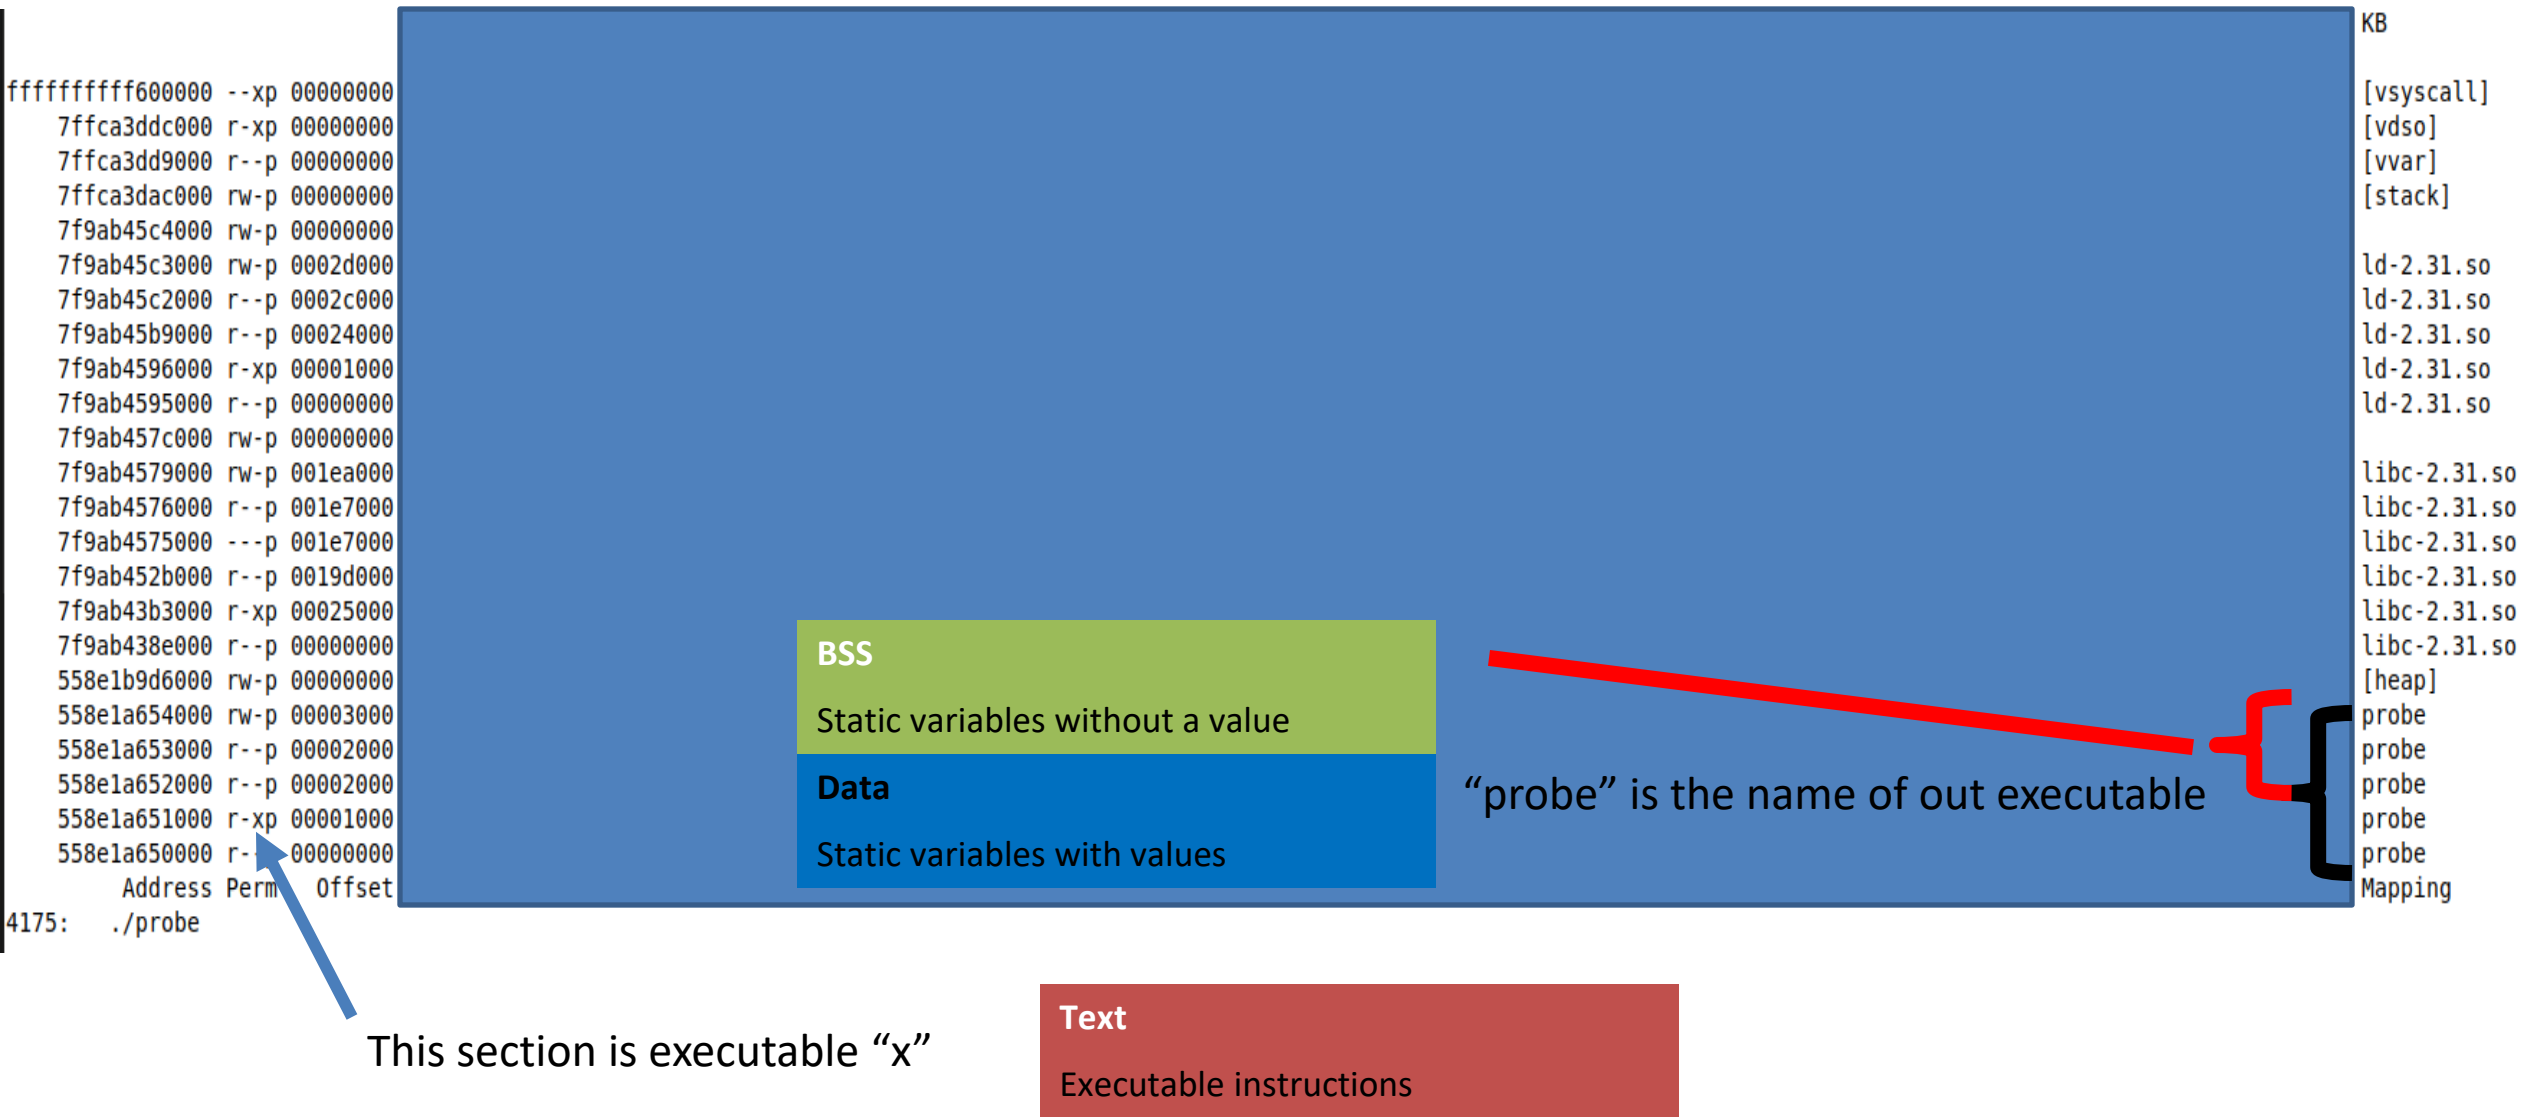
Output of `pmap` (process mapping tool)

			KB
fffffffff600000	--xp 00000000		[vsyscall]
7ffca3ddc000	r-xp 00000000		[vdso]
7ffca3dd9000	r--p 00000000		[vvar]
7ffca3dac000	rw-p 00000000		[stack]
7f9ab45c4000	rw-p 00000000		
7f9ab45c3000	rw-p 0002d000		ld-2.31.so
7f9ab45c2000	r--p 0002c000		ld-2.31.so
7f9ab45b9000	r--p 00024000		ld-2.31.so
7f9ab4596000	r-xp 00001000		ld-2.31.so
7f9ab4595000	r--p 00000000		ld-2.31.so
7f9ab457c000	rw-p 00000000		
7f9ab4579000	rw-p 001ea000		libc-2.31.so
7f9ab4576000	r--p 001e7000		libc-2.31.so
7f9ab4575000	---p 001e7000		libc-2.31.so
7f9ab452b000	r--p 0019d000		libc-2.31.so
7f9ab43b3000	r-xp 00025000		libc-2.31.so
7f9ab438e000	r--p 00000000		libc-2.31.so
558e1b9d6000	rw-p 00000000		[heap]
558e1a654000	rw-p 00003000		probe
558e1a653000	r--p 00002000		probe
558e1a652000	r--p 00002000		probe
558e1a651000	r-xp 00001000		probe
558e1a650000	r--p 00000000		probe
	Address Perm Offset		Mapping
4175: ./probe			

“probe” is the name of our executable

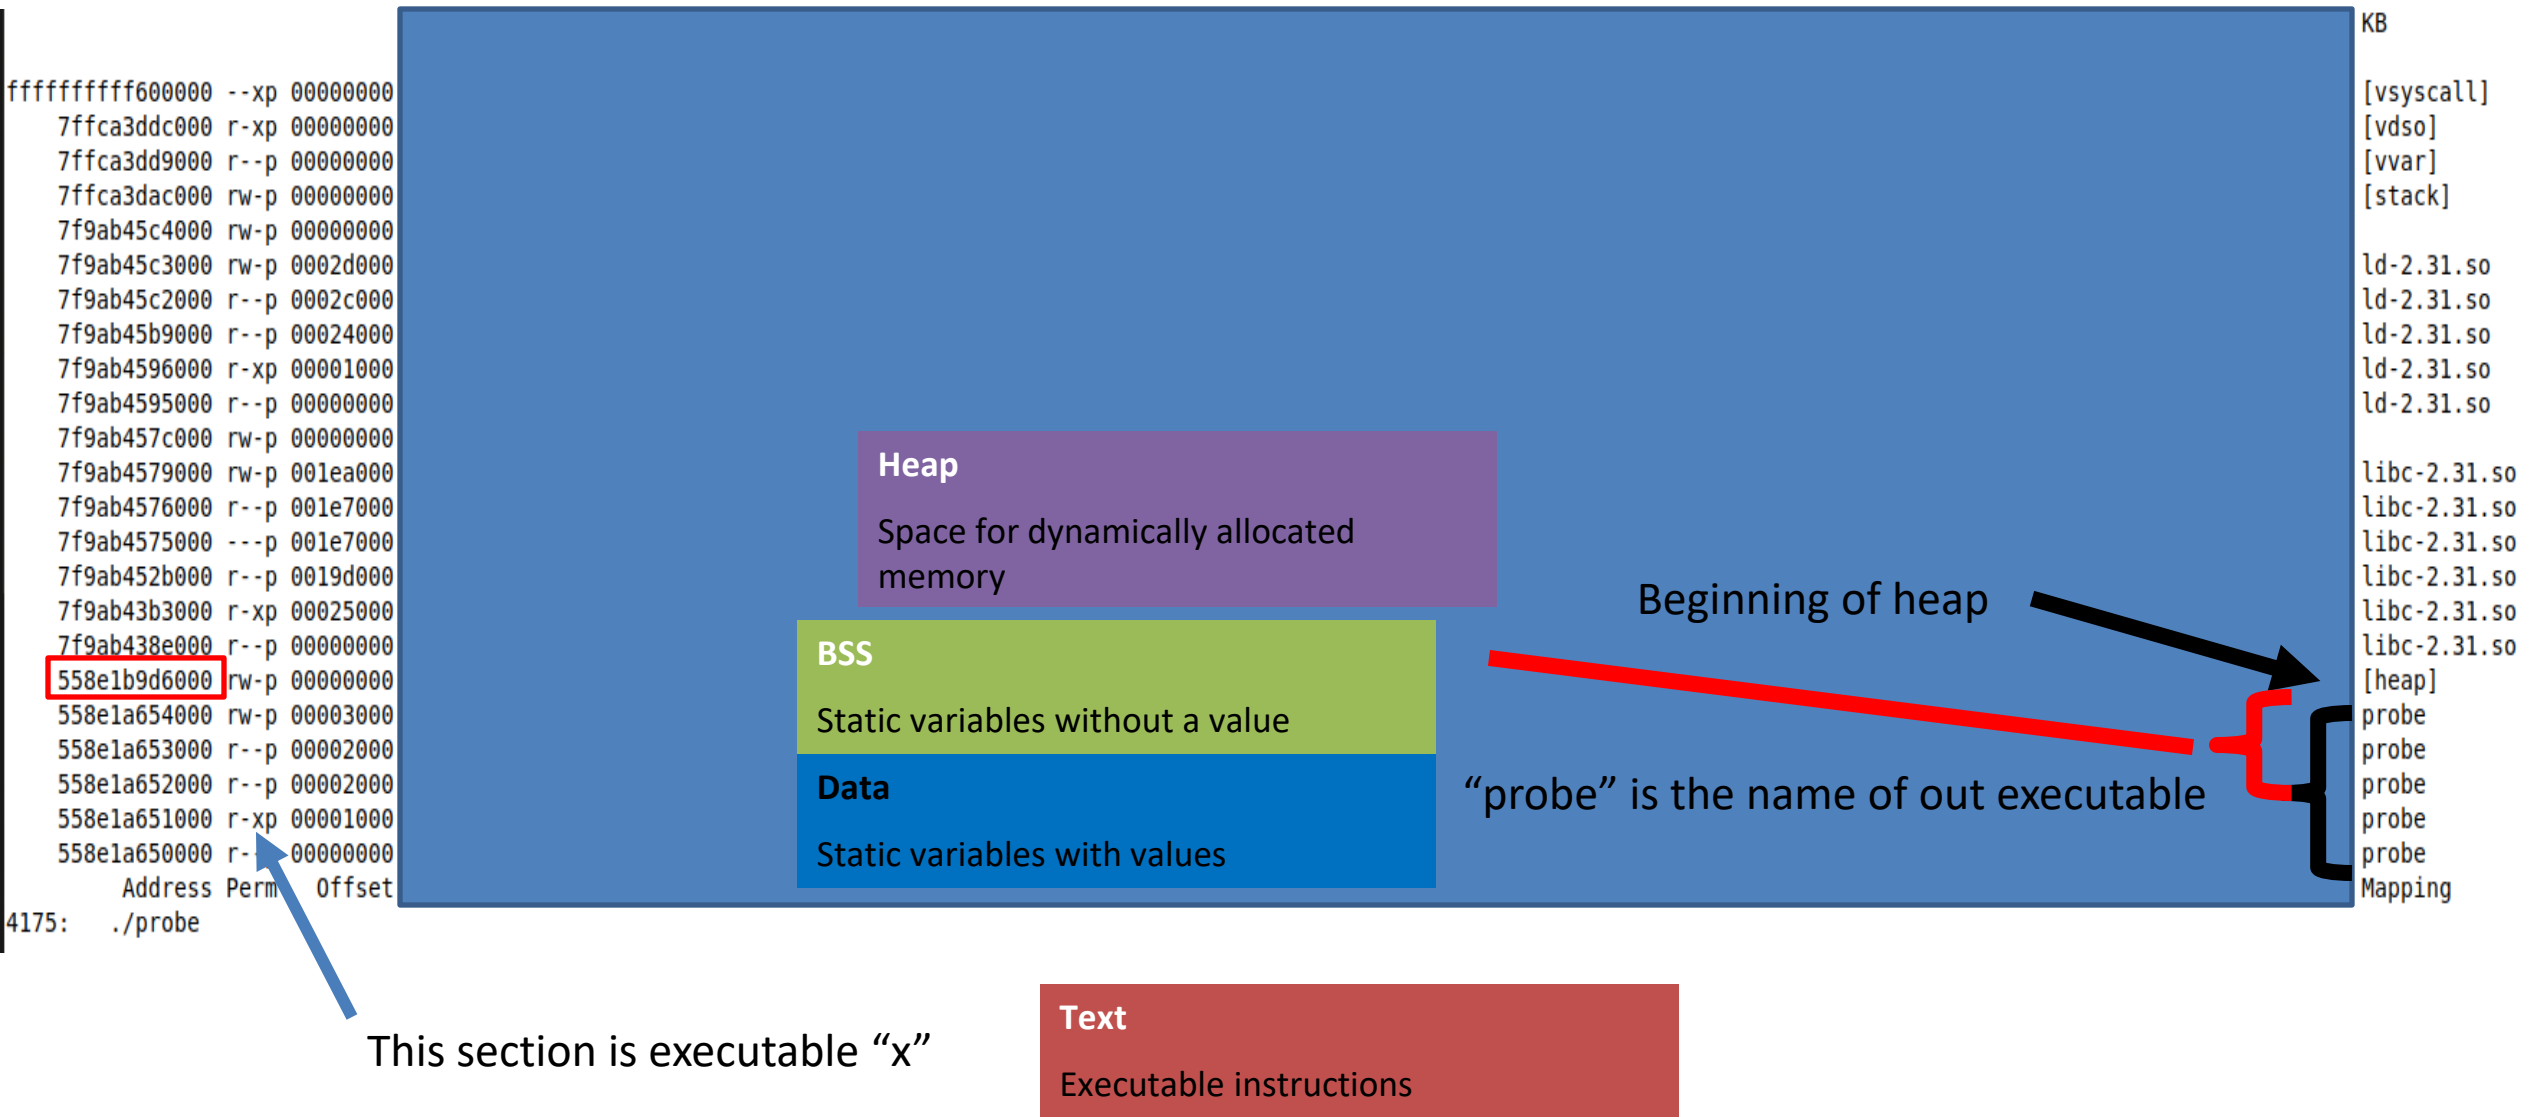
Applications Layout in Memory

Output of `pmap` (process mapping tool)



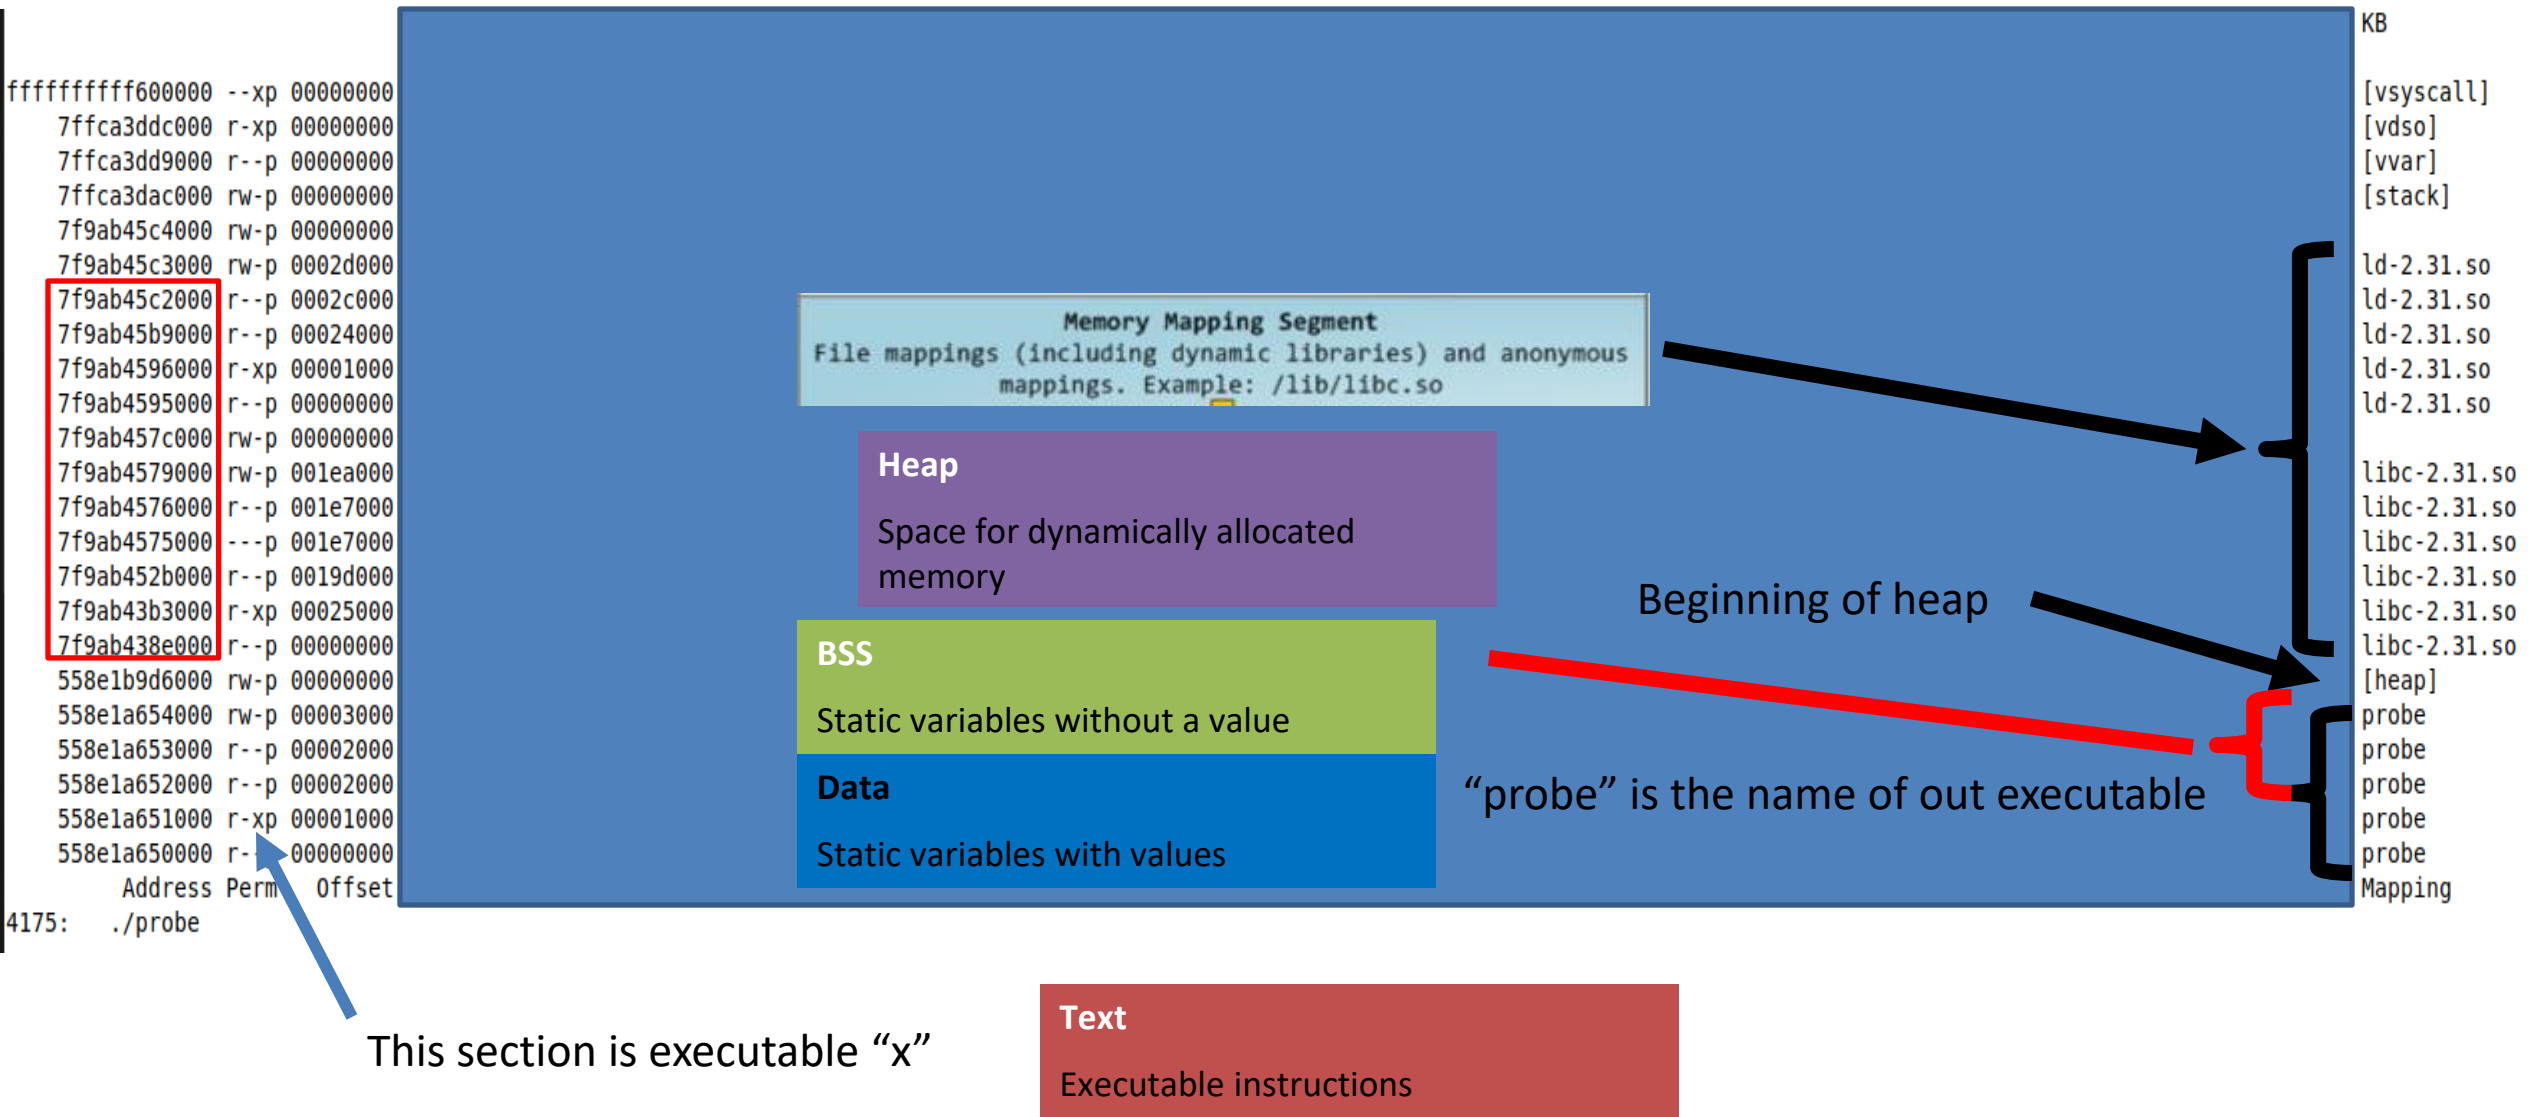
Applications Layout in Memory

Output of `pmap` (process mapping tool)



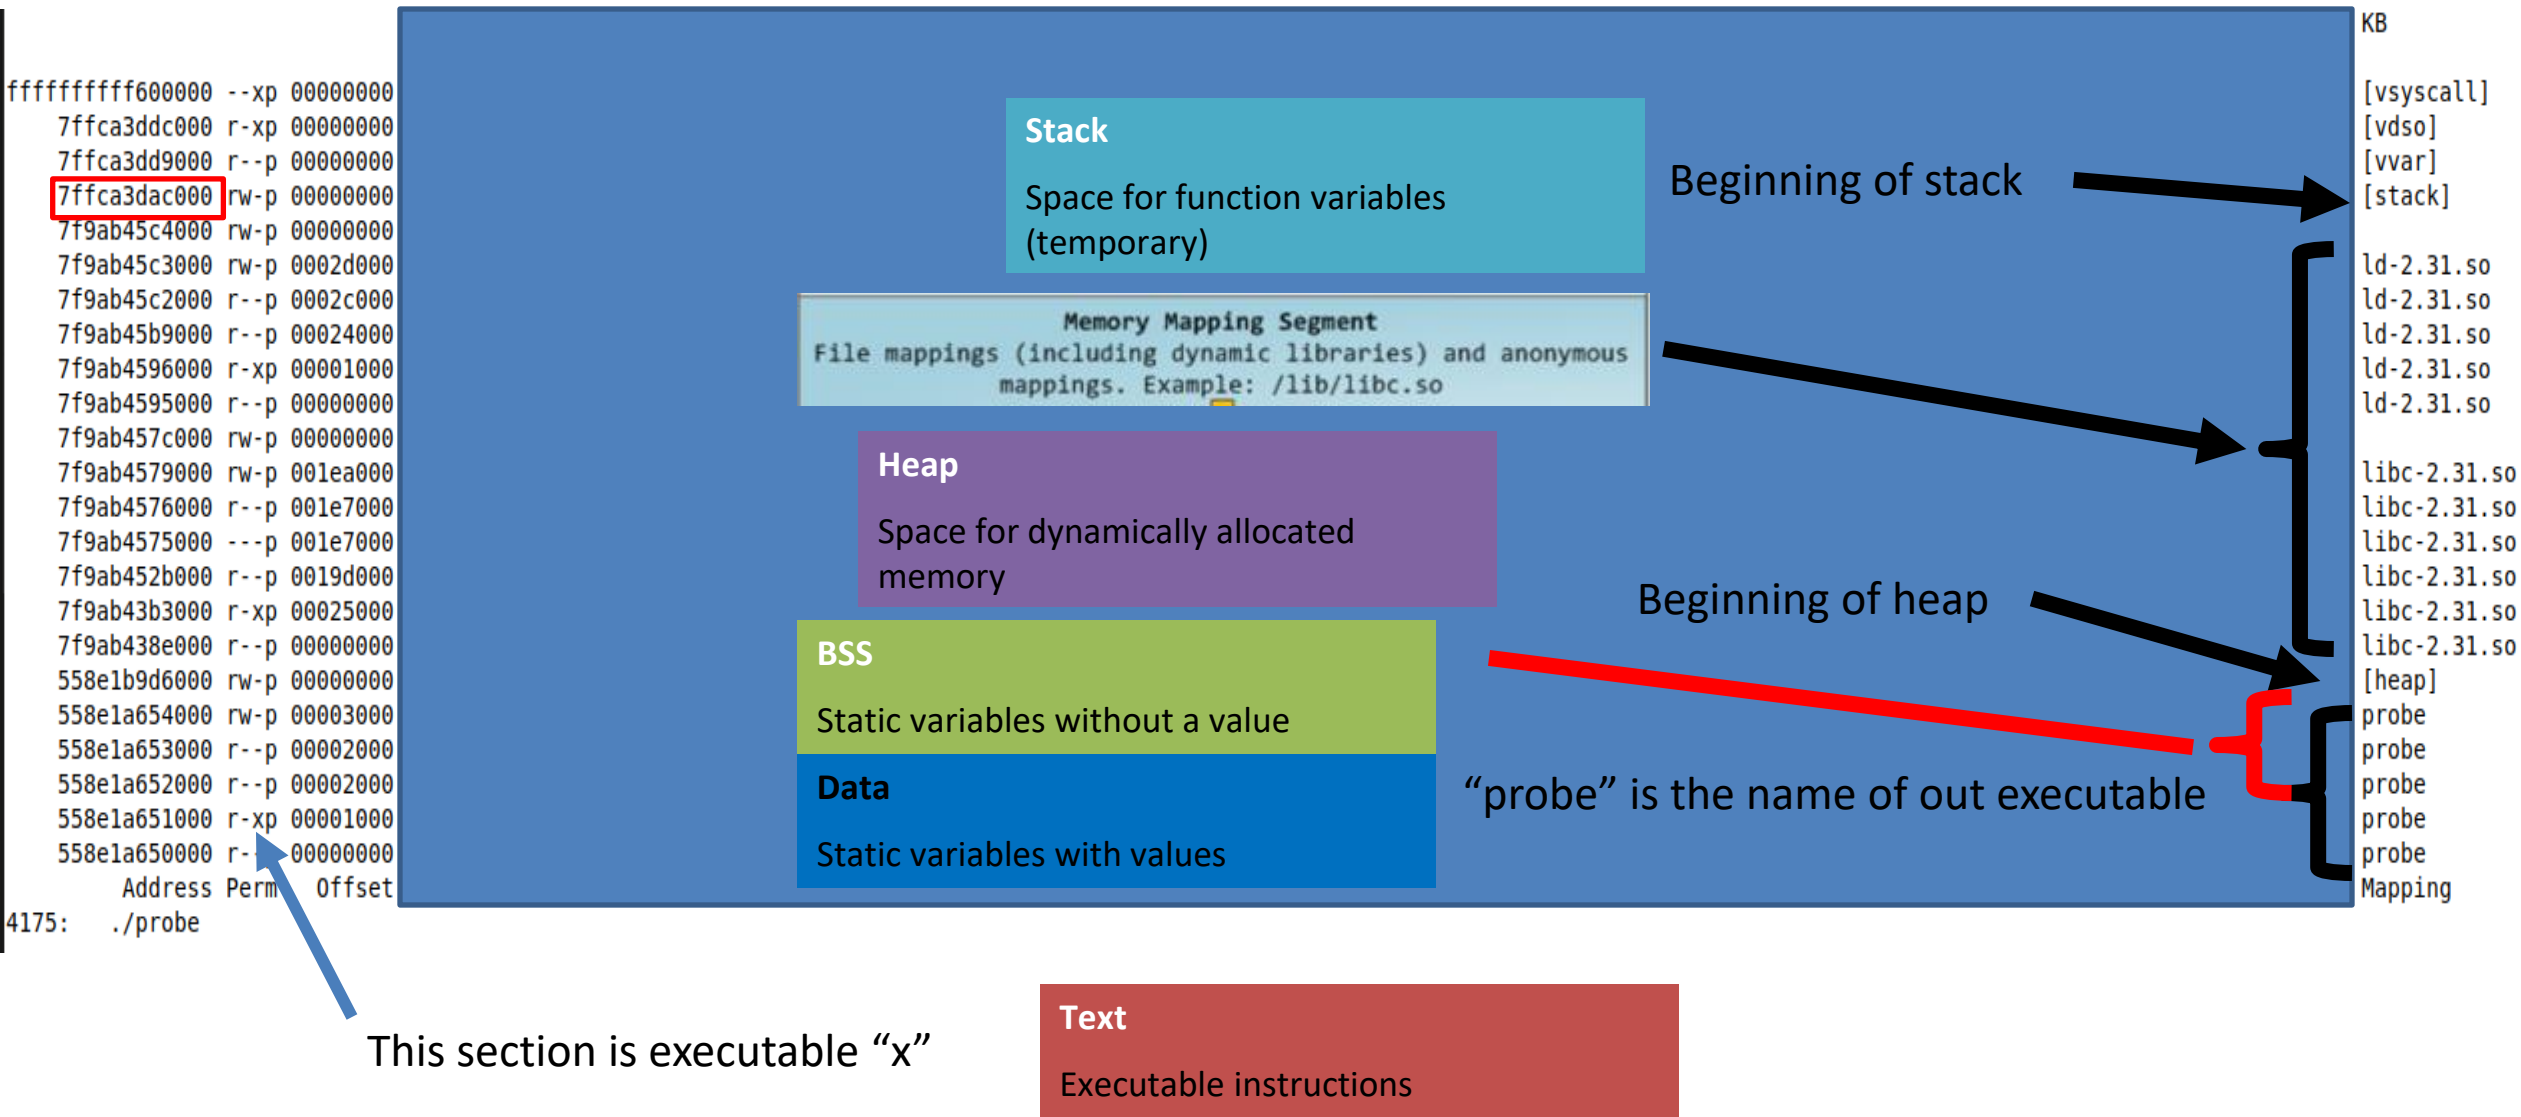
Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

Output of `pmap` (process mapping tool)

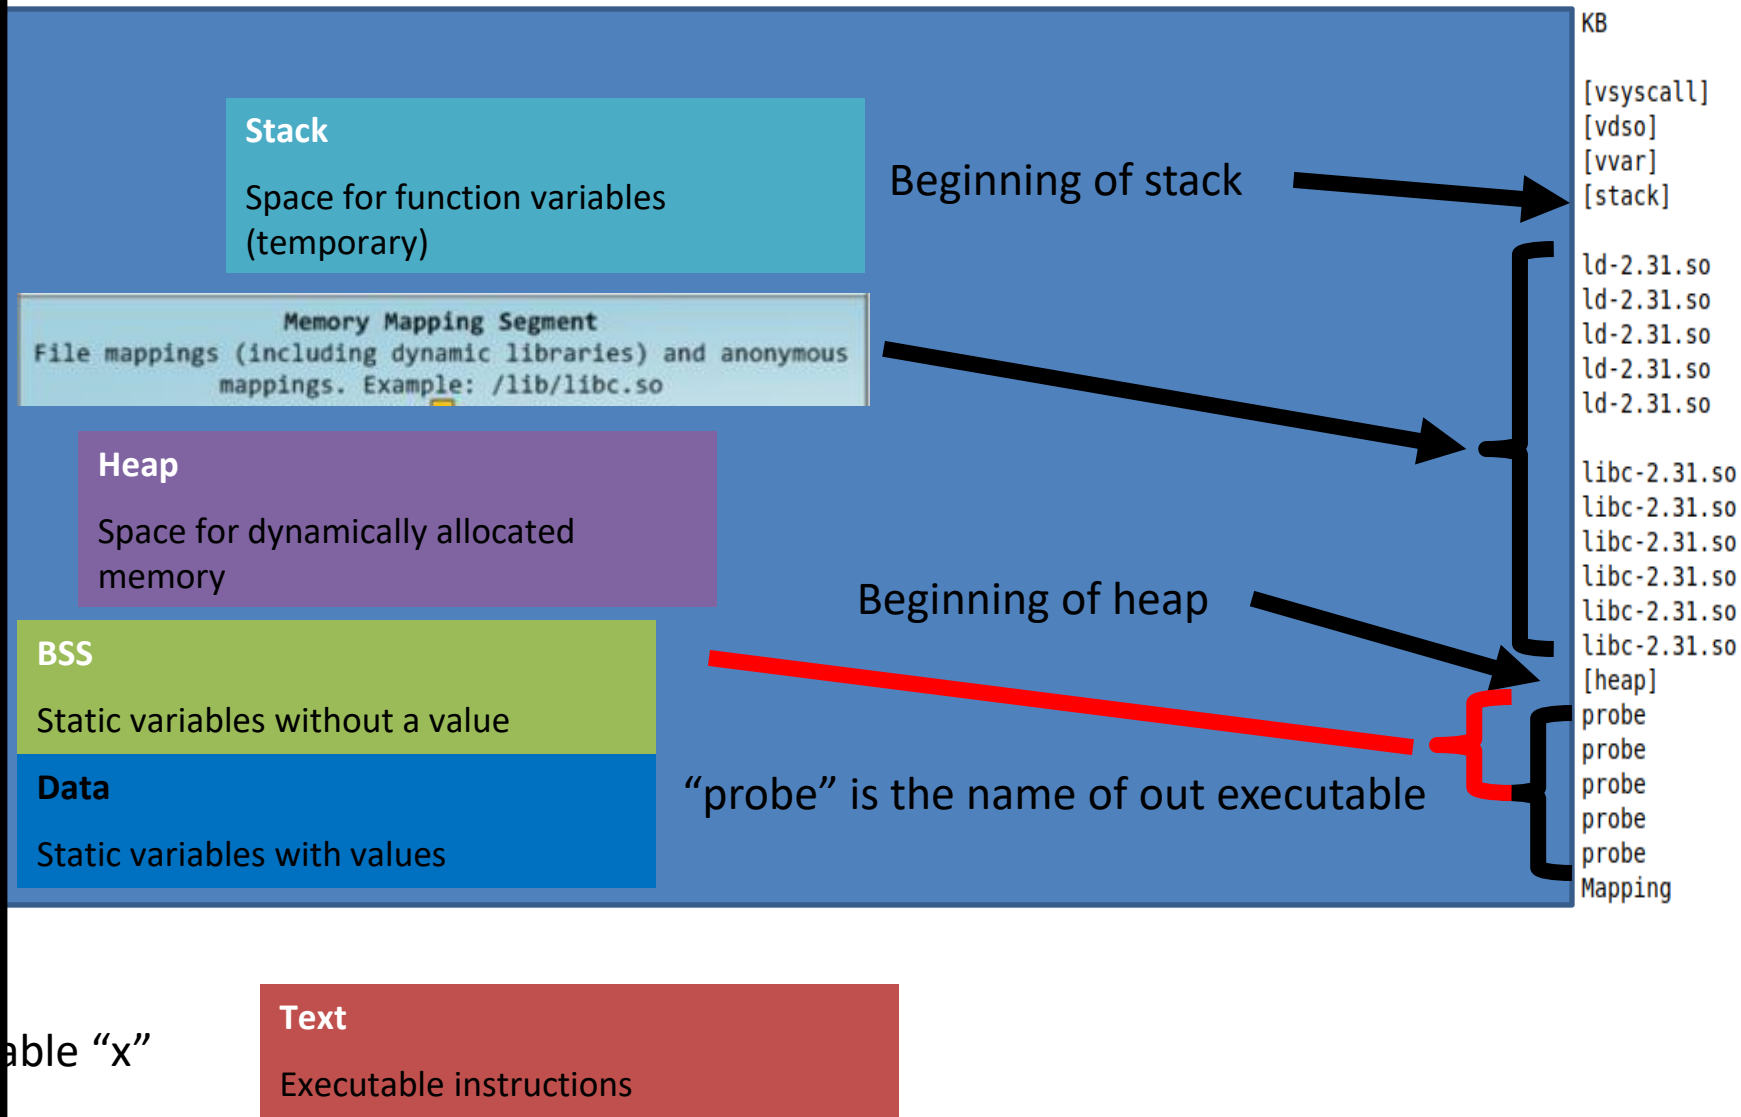
When you allocate variables on the stack



When you allocate variables on the heap

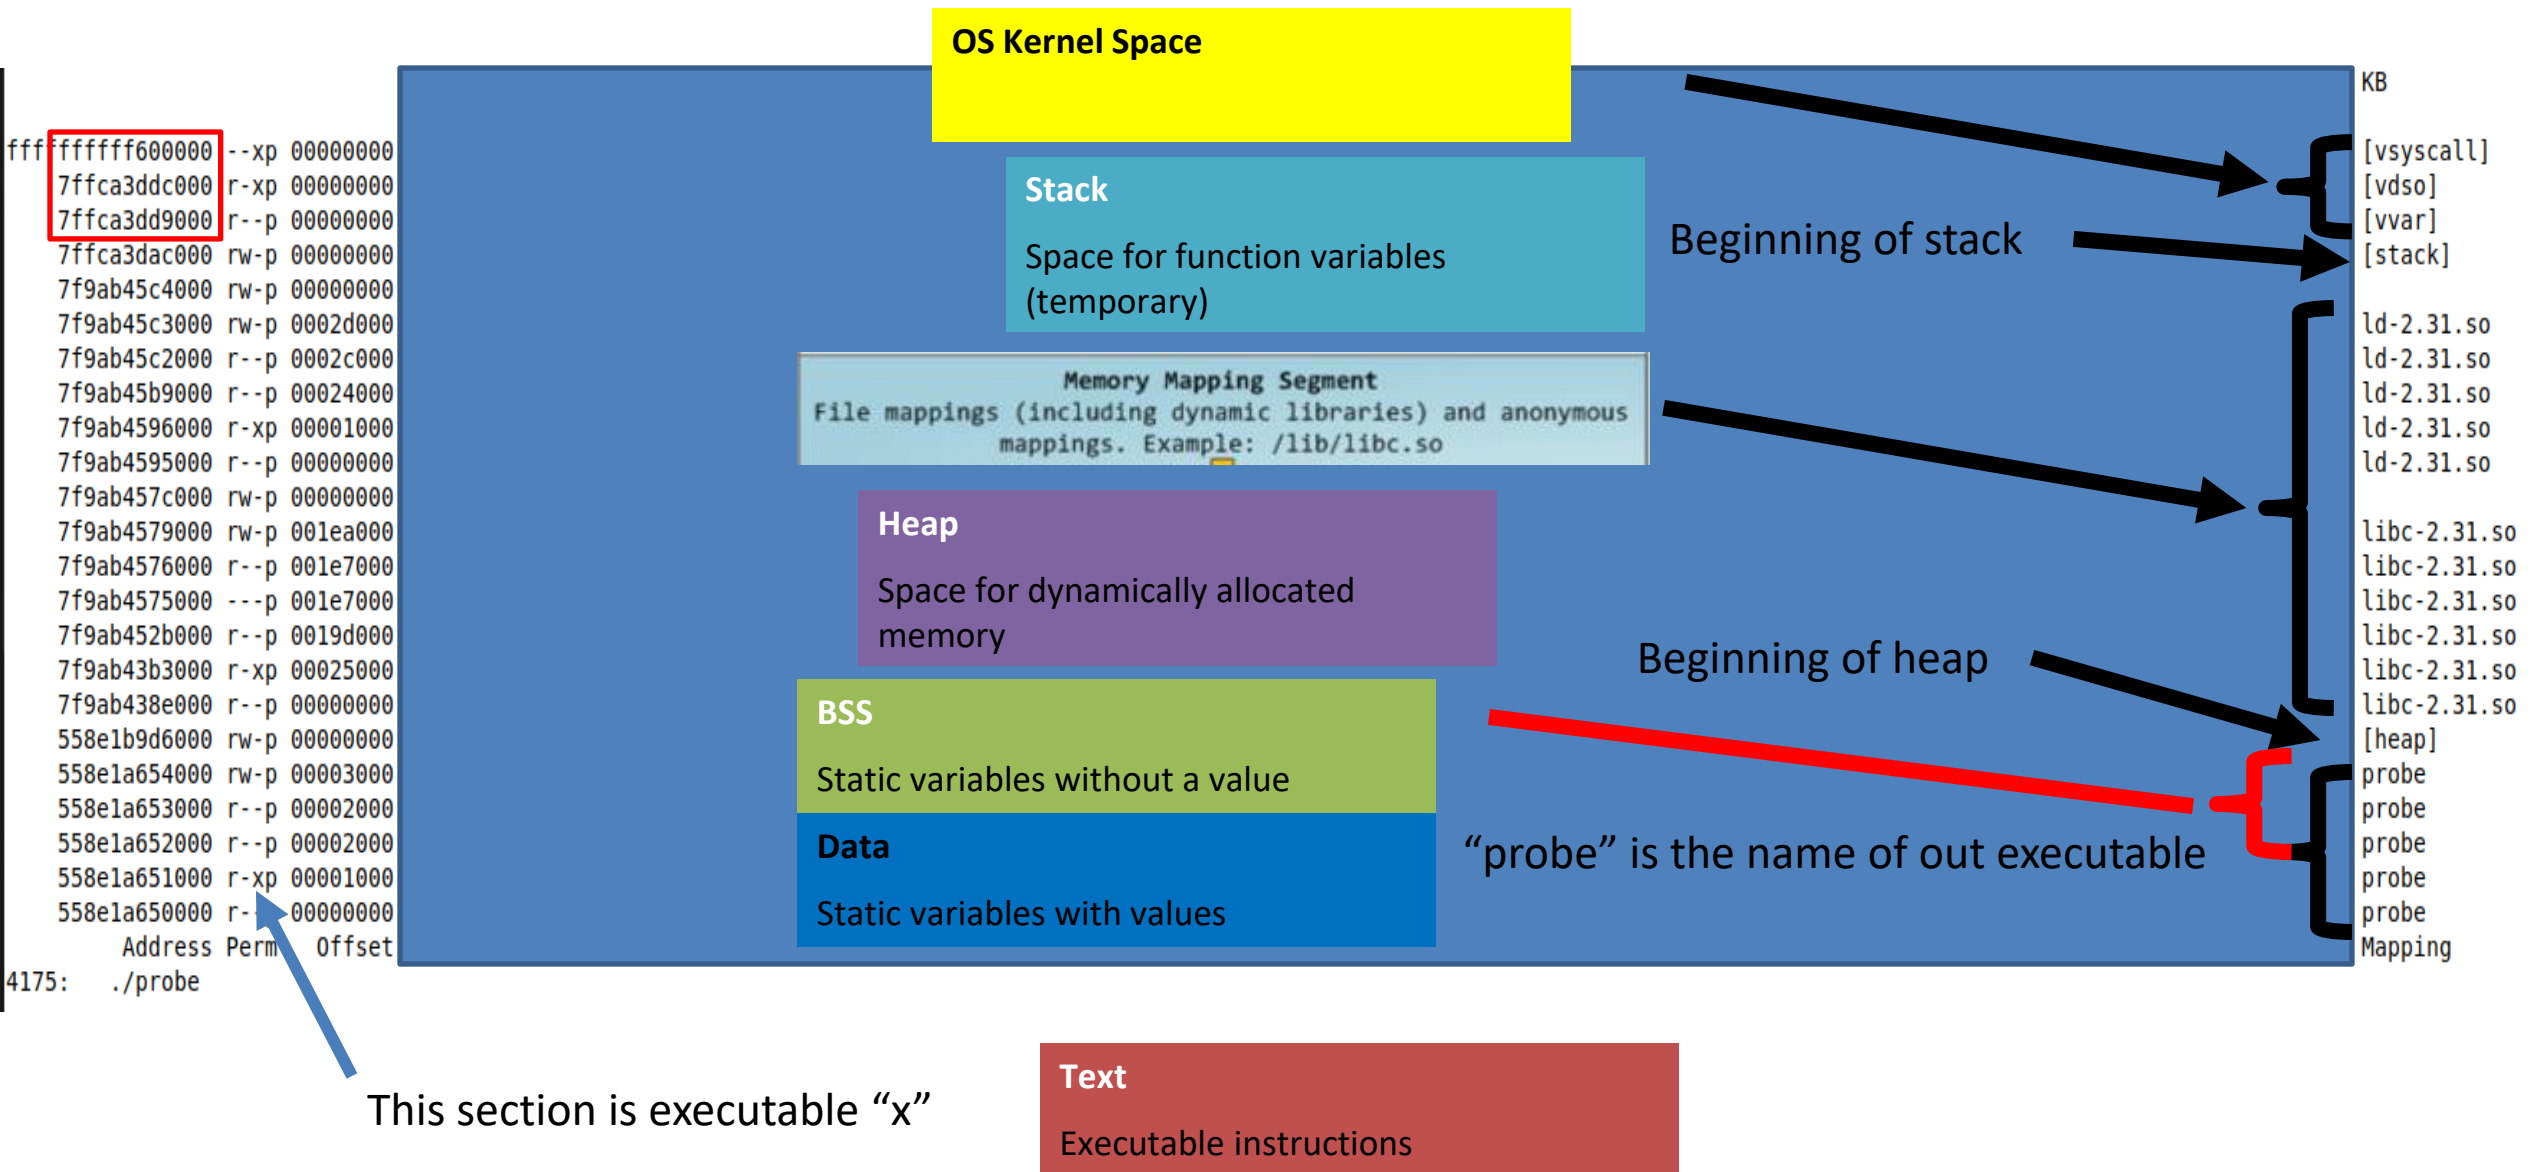


A new core memory!



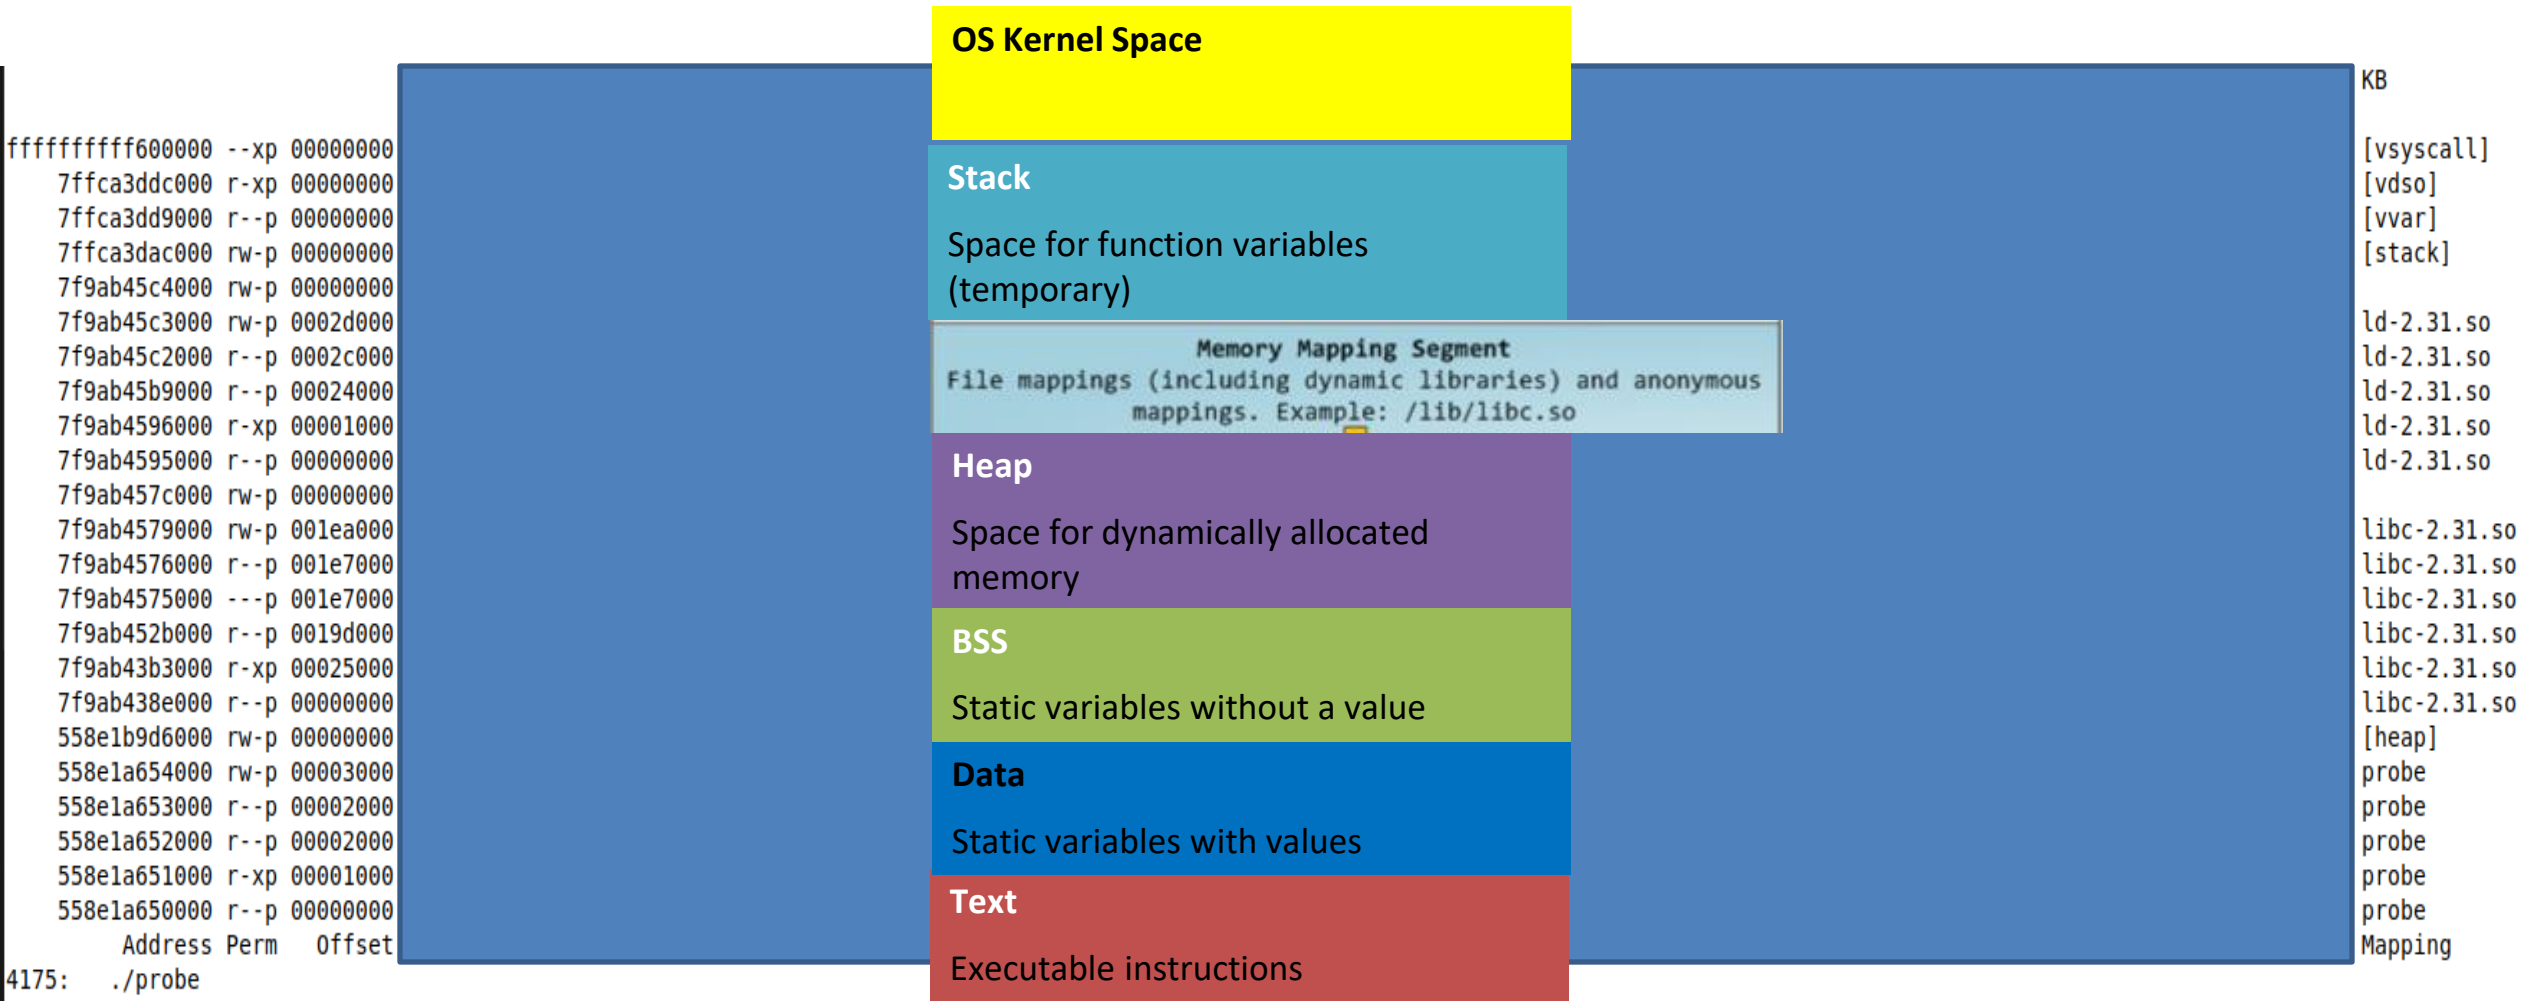
Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

Output of `pmap` (process mapping tool)



Applications Layout in Memory

Output of pmap (process mapping tool)

ffffffffff600000 --xp 00000000

7ffca3ddc000 r-xp 00000000

7ffca3dd9000 r--p 00000000

7ffca3dac000 rw-p 00000000

7f9ab45c4000 rw-p 00000000

7f9ab45c3000 rw-p 0002d000

7f9ab45c2000 r--p 0002c000

7f9ab45b9000 r--p 00024000

7f9ab4596000 r-xp 00001000

7f9ab4595000 r--p 00000000

7f9ab457c000 rw-p 00000000

7f9ab4579000 rw-p 001ea000

7f9ab4576000 r--p 001e7000

7f9ab4575000 ---p 001e7000

7f9ab452b000 r--p 0019d000

7f9ab43b3000 r-xp 00025000

7f9ab438e000 r--p 00000000

558e1b9d6000 rw-p 00000000

558e1a654000 rw-p 00003000

558e1a653000 r--p 00002000

558e1a652000 r--p 00002000

558e1a651000 r-xp 00001000

558e1a650000 r--p 00000000

Address Perm Offset

4175: ./probe

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7f9ab43f2e10

-> the address of getenv = 0x7f9ab43d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

-> argv[0] = 0x7ffca3dcc45f

value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") = 0x7ffca3dcc5ff

value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe


probe

probe

probe

probe

Mapping

 MONTANA
STATE UNIVERSITY

82

Applications Layout in Memory

Output of pmap (process mapping tool)

Where is "main" located in memory?

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7ffca3dc3f2e10

-> the address of getenv = 0x7ffca3dc3d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

-> argv[0] value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

ffffffffff600000 --xp 00000000

7ffca3ddc000 r-xp 00000000

7ffca3dd9000 r--p 00000000

7ffca3dac000 rw-p 00000000

7f9ab45c4000 rw-p 00000000

7f9ab45c3000 rw-p 0002d000

7f9ab45c2000 r--p 0002c000

7f9ab45b9000 r--p 00024000

7f9ab4596000 r-xp 00001000

7f9ab4595000 r--p 00000000

7f9ab457c000 rw-p 00000000

7f9ab4579000 rw-p 001ea000

7f9ab4576000 r--p 001e7000

7f9ab4575000 ---p 001e7000

7f9ab452b000 r--p 0019d000

7f9ab43b3000 r-xp 00025000

7f9ab438e000 r--p 00000000

558e1b9d6000 rw-p 00000000

558e1a654000 rw-p 00003000

558e1a653000 r--p 00002000

558e1a652000 r--p 00002000

558e1a651000 r-xp 00001000

558e1a650000 r--p 00000000

Address Perm Offset

4175: ./probe

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe

probe

probe

probe

Mapping

Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "main" located in memory?

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7f9ab43f2e10

-> the address of getenv = 0x7f9ab43d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

-> argv[0] value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

558e1a651000 r-xp 00001000

Address Perm Offset

4175: ./probe

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe

probe

probe

Mapping

`main` is code in our program, so it goes inside the text segment

Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "printf" located in memory?

```

-> the address of main      = 0x558e1a651249
-> the address of printf    = 0x7ffca3dc43f2e10
-> the address of getenv    = 0x7ffca3dc43d7020
-> a stack address         = 0x7ffca3dcb3b0
-> a global address        = 0x558e1a6540c4
-> the argv address        = 0x7ffca3dcb4f8
-> argv[0]                 = 0x7ffca3dcc45f
    value is [./probe]
-> the environ address     = 0x7ffca3dcb508
-> the envp address        = 0x7ffca3dcb508
-> getenv("PWD")           = 0x7ffca3dcc5ff
    value is [/home/seed/os-review]
-> a heap address          = 0x558e1b9d66b0

```

Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "printf" located in memory?

-> the address of main = 0x558e1a651249

-> the address of printf = 0x7f9ab43f2e10

-> the address of getenv = 0x7f9ab43d7020

-> a stack address = 0x7ffca3dcb3b0

-> a global address = 0x558e1a6540c4

-> the argv address = 0x7ffca3dcb4f8

-> argv[0] value is [./probe]

-> the environ address = 0x7ffca3dcb508

-> the envp address = 0x7ffca3dcb508

-> getenv("PWD") value is [/home/seed/os-review]

-> a heap address = 0x558e1b9d66b0

7f9ab43b3000 r-xp 00025000

KB

[vsyscall]

[vdso]

[vvar]

[stack]

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

ld-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

libc-2.31.so

[heap]

probe

probe

probe

probe

probe

Mapping

ffffffffff600000 --xp 00000000

7ffca3ddc000 r-xp 00000000

7ffca3dd9000 r--p 00000000

7ffca3dac000 rw-p 00000000

7f9ab45c4000 rw-p 00000000

7f9ab45c3000 rw-p 0002d000

7f9ab45c2000 r--p 0002c000

7f9ab45b9000 r--p 00024000

7f9ab4596000 r-xp 00001000

7f9ab4595000 r--p 00000000

7f9ab457c000 rw-p 00000000

7f9ab4579000 rw-p 001ea000

7f9ab4576000 r--p 001e7000

7f9ab4575000 ---p 001e7000

7f9ab452b000 r--p 0010d000

7f9ab438e000 r--p 00000000

558e1b9d6000 rw-p 00000000

558e1a654000 rw-p 00003000

558e1a653000 r--p 00002000

558e1a652000 r--p 00002000

558e1a651000 r-xp 00001000

558e1a650000 r--p 00000000

Address Perm Offset

4175: ./probe

`printf` is executable code from a shared library (libc) so we are in the memory mapping segment!

Applications Layout in Memory

Output of `pmap` (process mapping tool)

Where is "argv" located in memory?

```

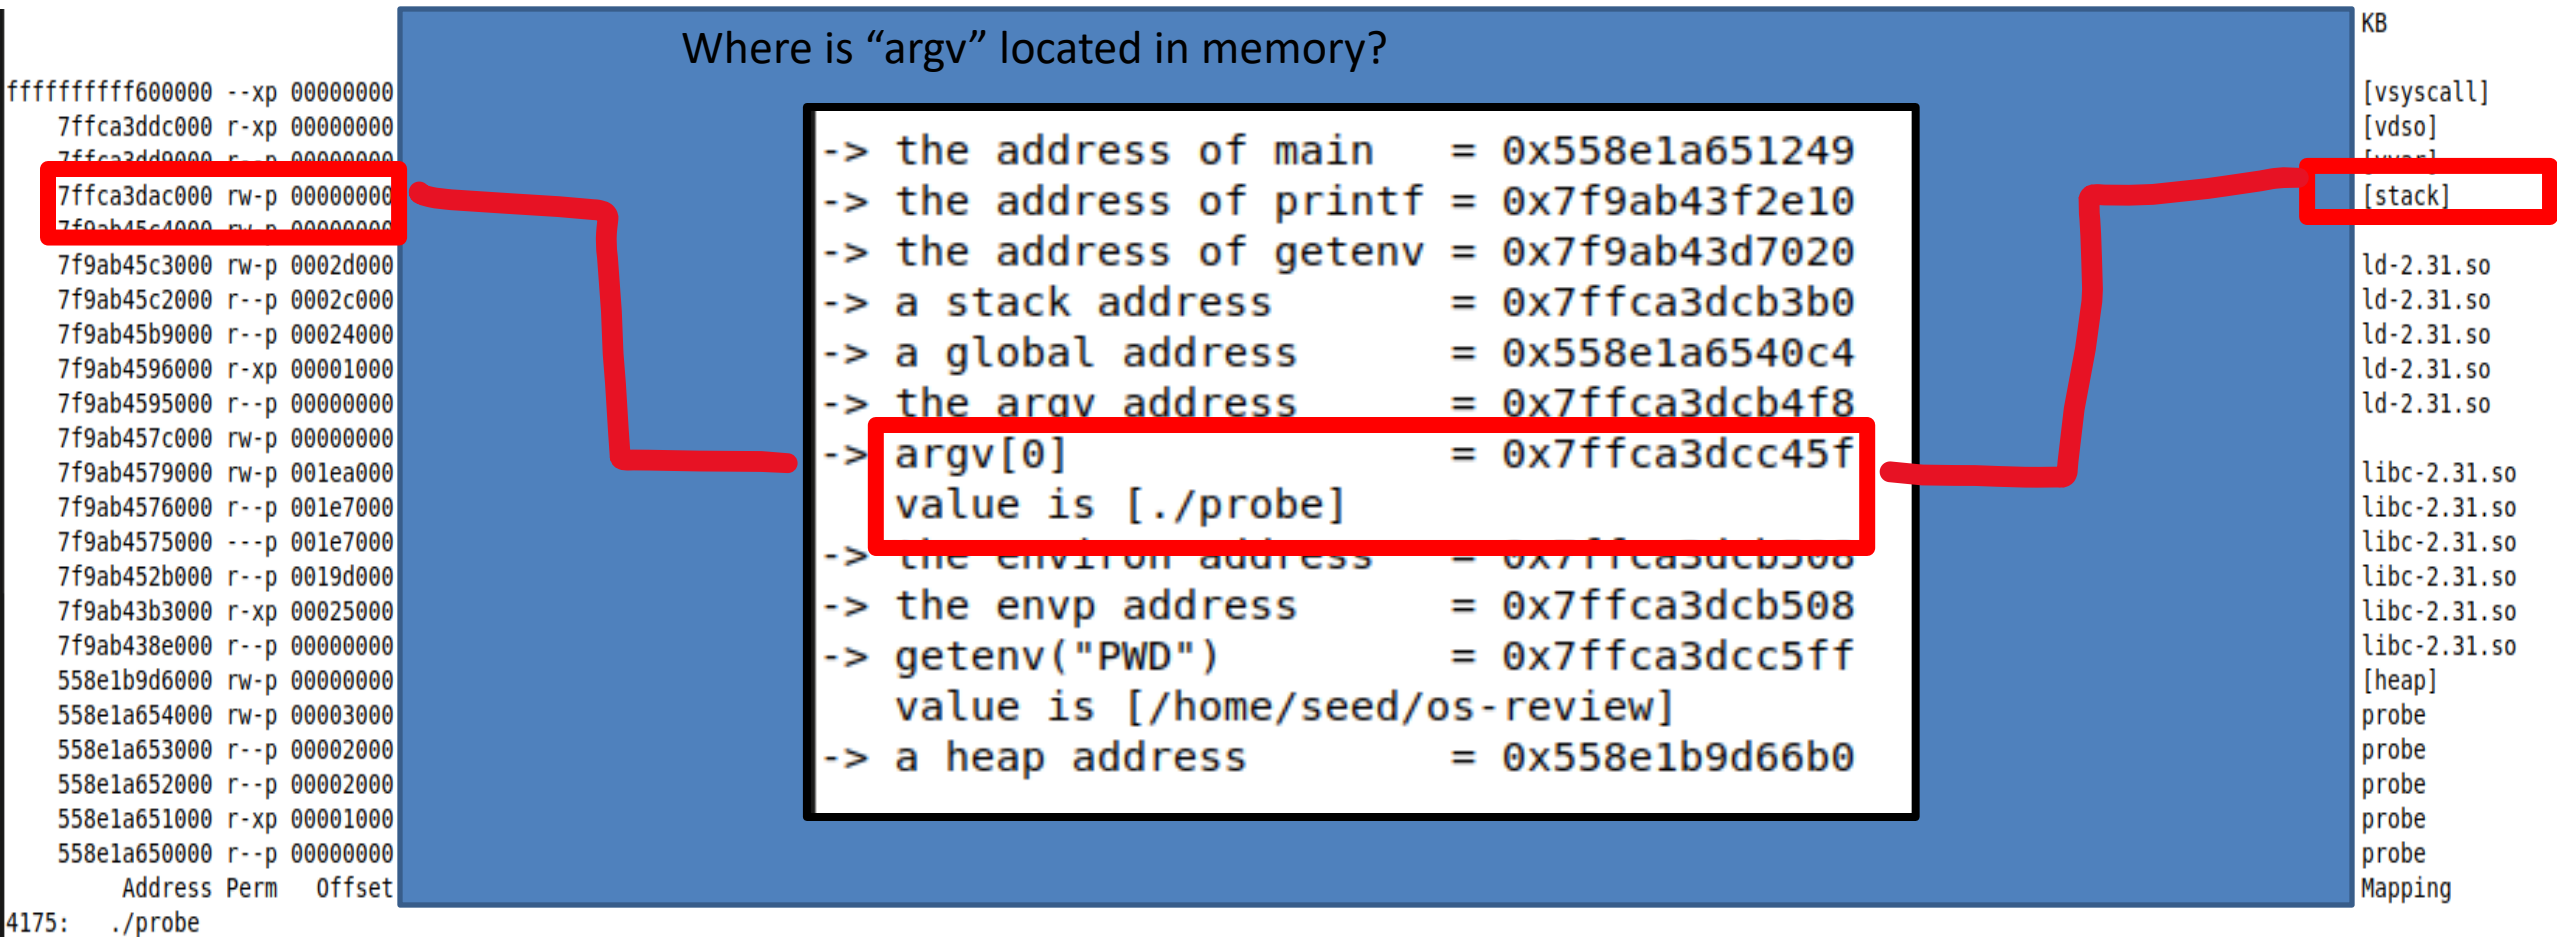
-> the address of main      = 0x558e1a651249
-> the address of printf    = 0x7f9ab43f2e10
-> the address of getenv    = 0x7f9ab43d7020
-> a stack address         = 0x7ffc3dcb3b0
-> a global address        = 0x558e1a6540c4
-> the argv address        = 0x7ffc3dcb4f8
-> argv[0]                 = 0x7ffc3dcc45f
    value is [./probe]
-> the environ address     = 0x7ffc3dcb508
-> the envp address        = 0x7ffc3dcb508
-> getenv("PWD")           = 0x7ffc3dcc5ff
    value is [/home/seed/os-review]
-> a heap address          = 0x558e1b9d66b0

```

`argv` is an array that holds the command line parameters passed into this program

Applications Layout in Memory


Output of `pmap` (process mapping tool)



`argv` is the argument to the main function, so we are in the stack!

Applications Layout in Memory

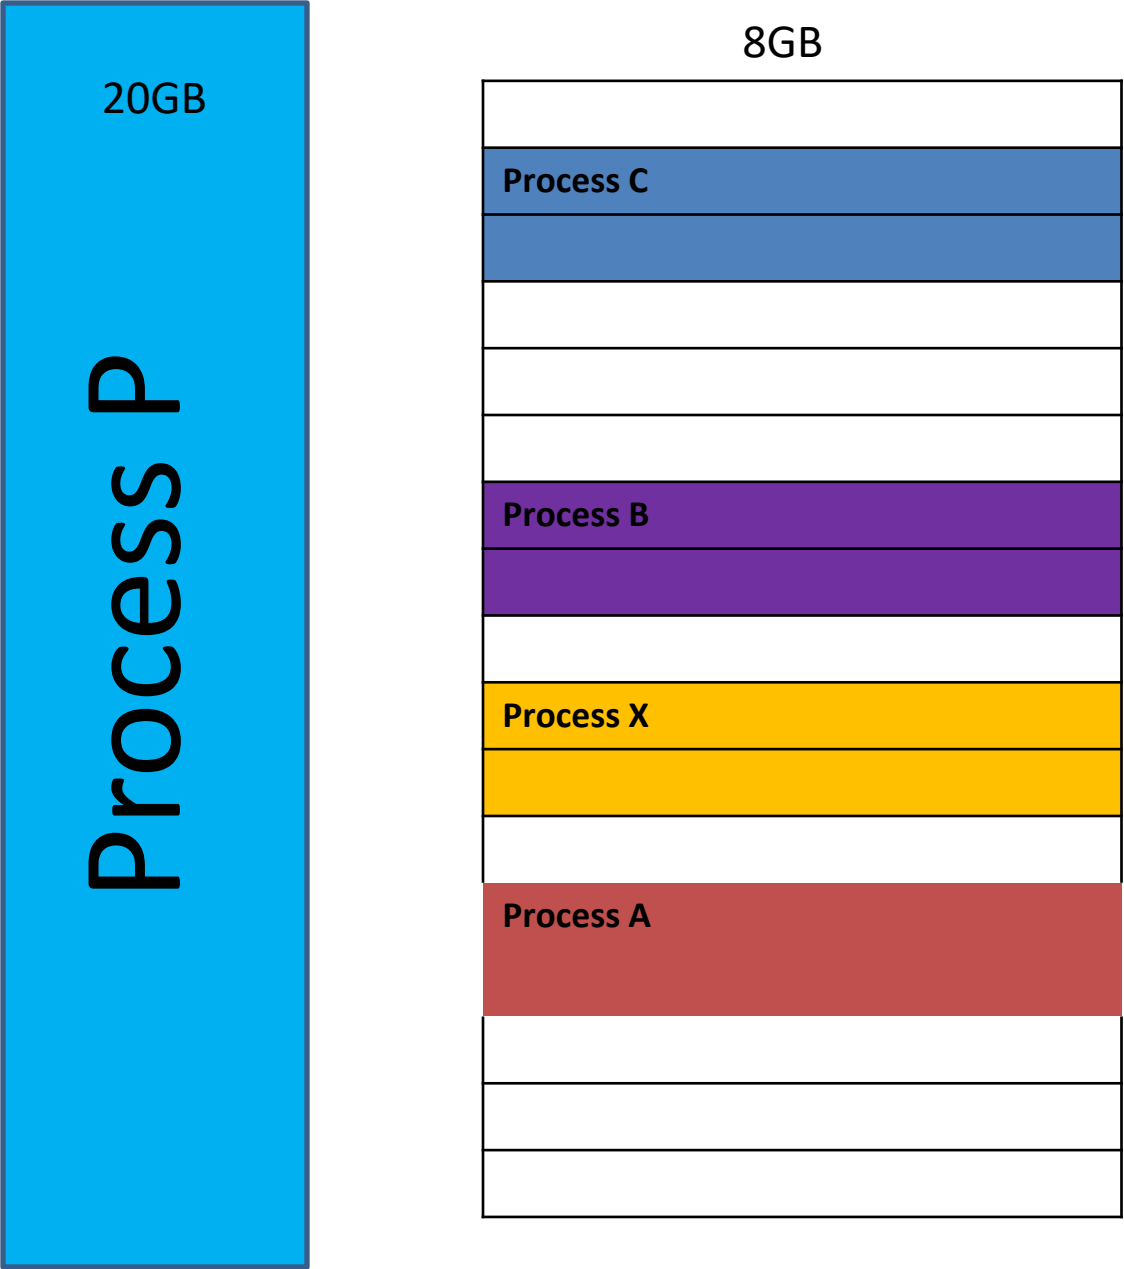
We have many programs that are actively running on our computer



Applications Layout in Memory

We have many programs that are actively running on our computer

What if we have a program that is bigger than out entire main memory?

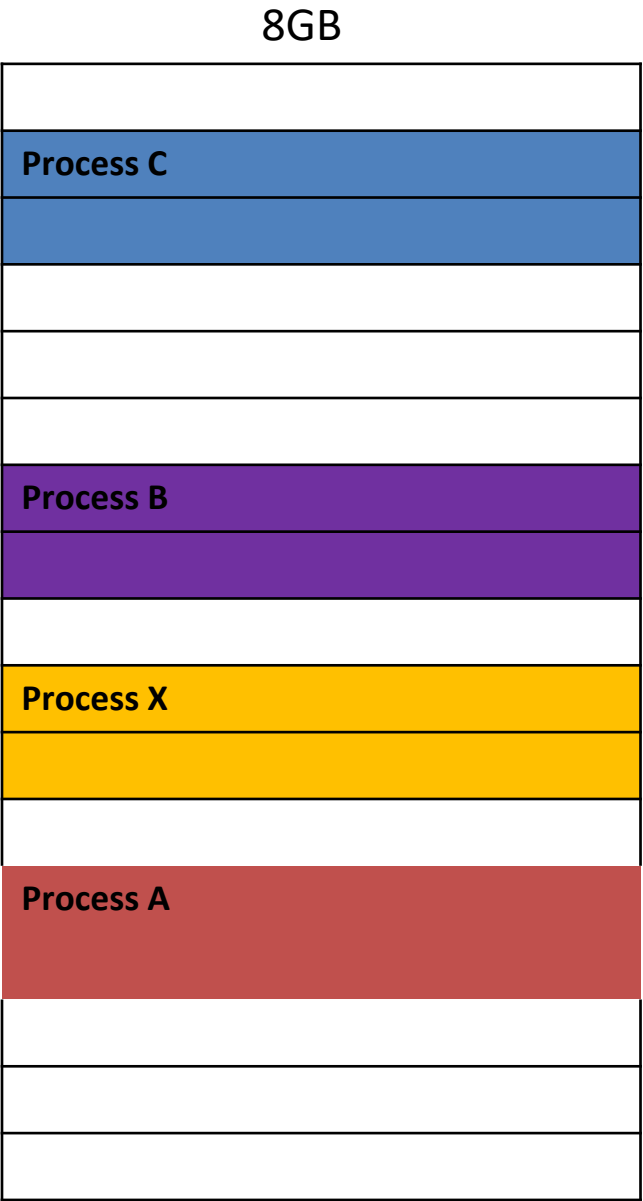



Applications Layout in Memory

We have many programs that are actively running on our computer

What if we have a program that is bigger than our entire main memory?

Does our computer crash?



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage



Secondary Storage



8GB

Process C
Process B
Process X
Process A

Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

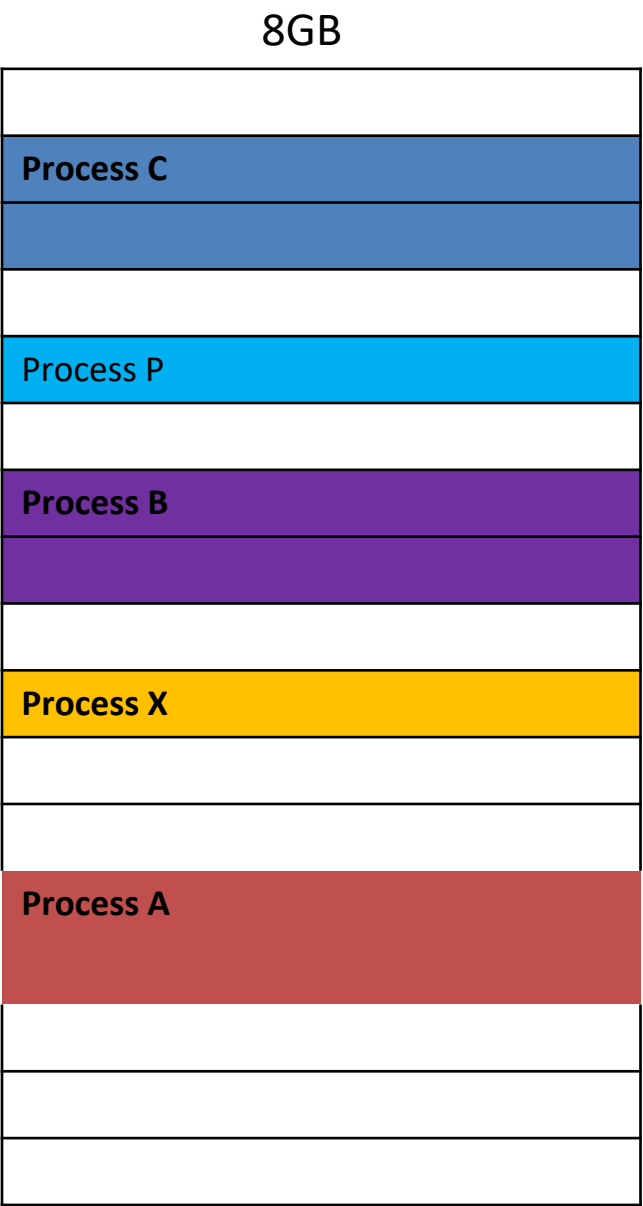

We split the process into smaller **pages**. Load pages into memory only when needed



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage


We split the process into smaller **pages**. Load pages into memory only when needed



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage


We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed



Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed

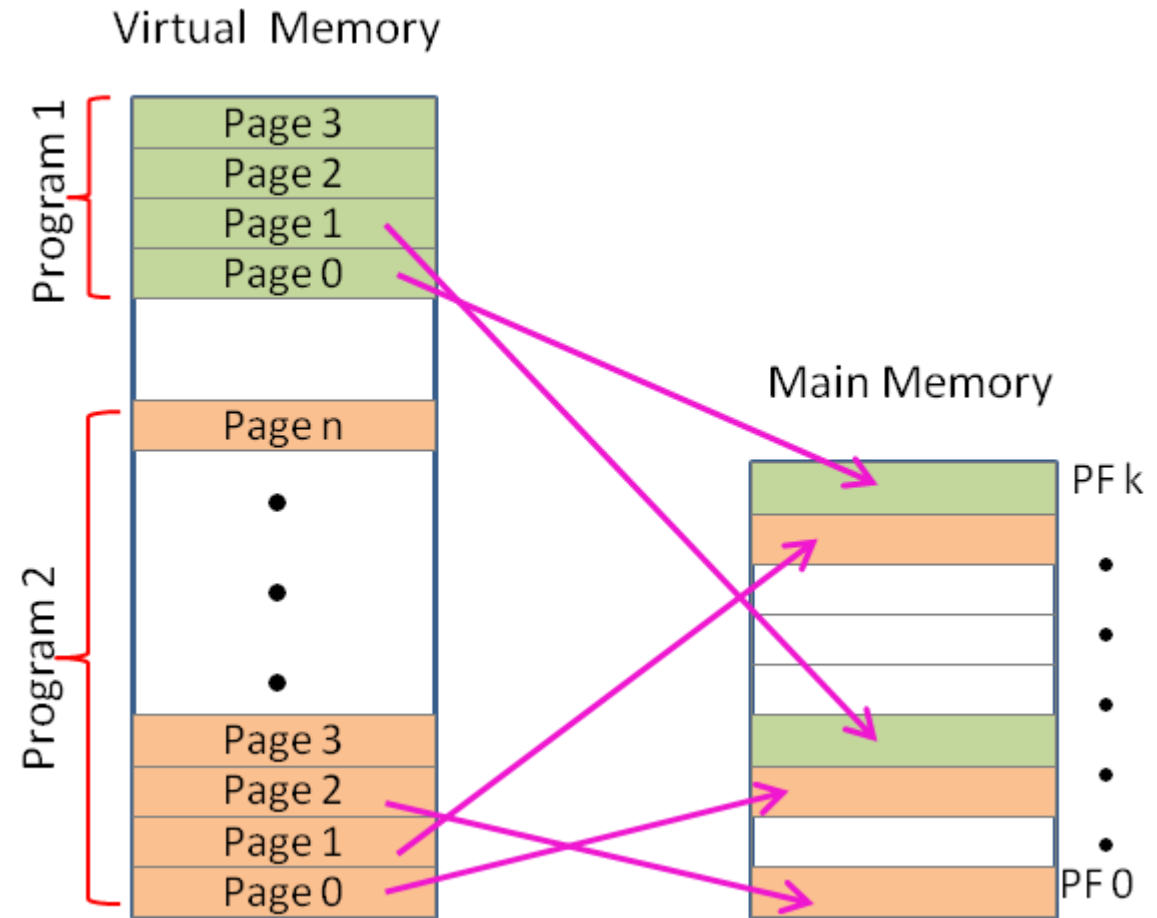


Constantly swapping stuff in and out of main memory

Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed

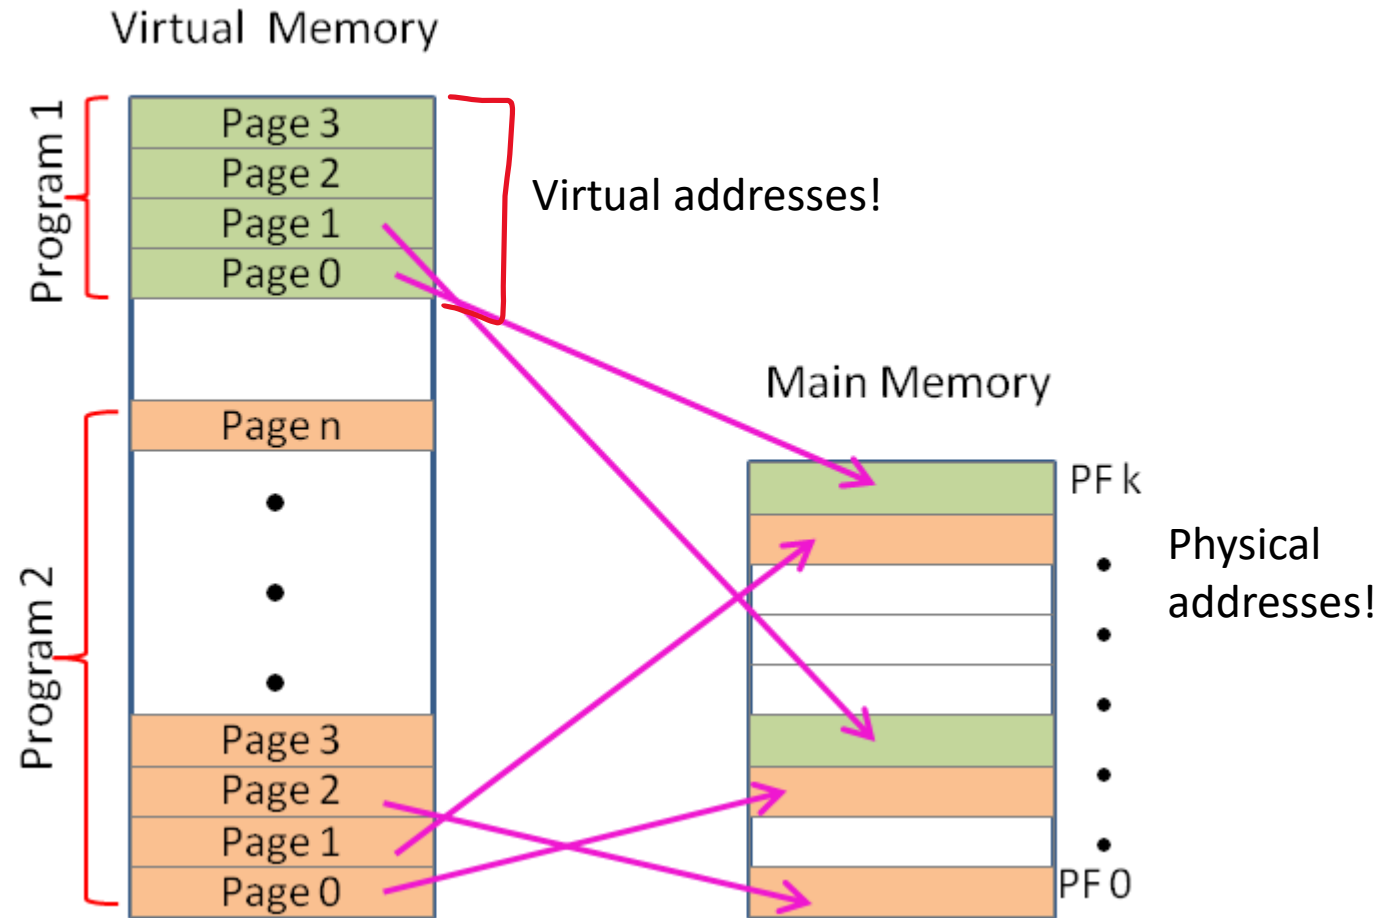


A process in memory is not contiguous

Memory management

Virtual Memory uses secondary storage to give programs the illusion that they have infinite storage

We split the process into smaller, fixed-size, **pages**. Load pages into memory only when needed



A process in memory is not contiguous

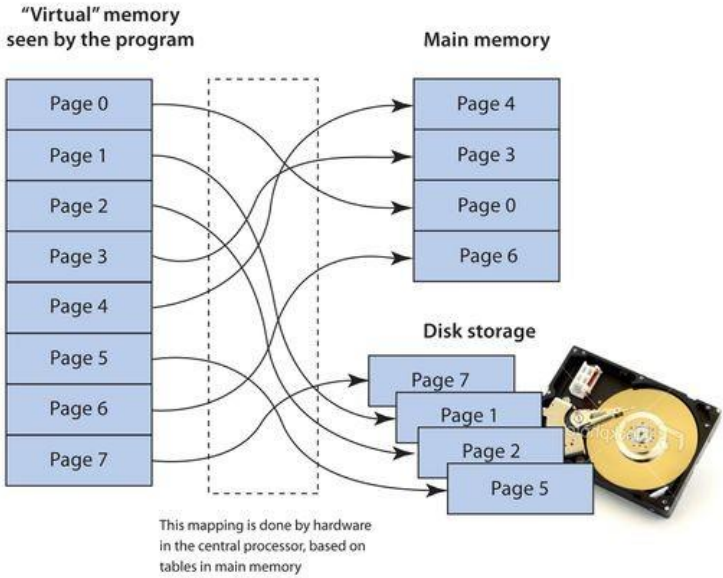
In probe.c, we are seeing virtual addresses!

Internal fragmentation vs external fragmentation

OS Review

Memory Manager

- Manages how physical memory is utilized



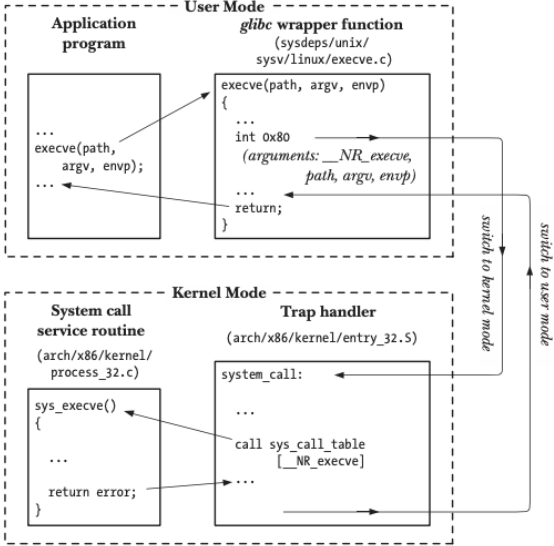
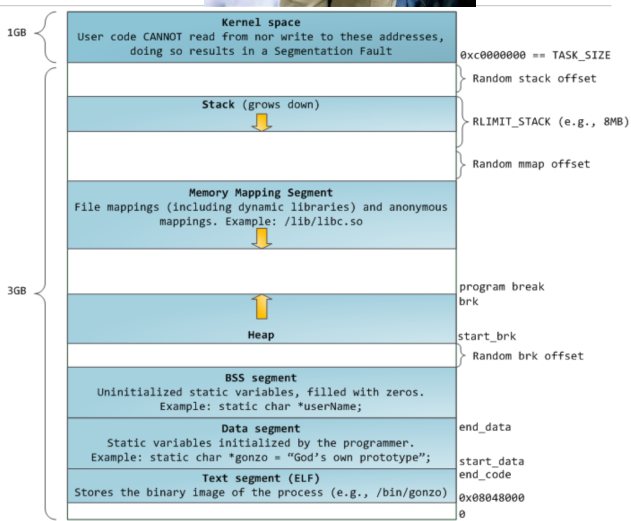
Process Manager

- Manages how processes are structured and how to handle many processes running at once



Interface Manager

- Manages communication between apps and hardware





Traffic Manager

- Manages which programs should be executed by the CPU

Process A (Ready)

Process B (Urgent)

Process C (Ready)

Process D (Blocked)




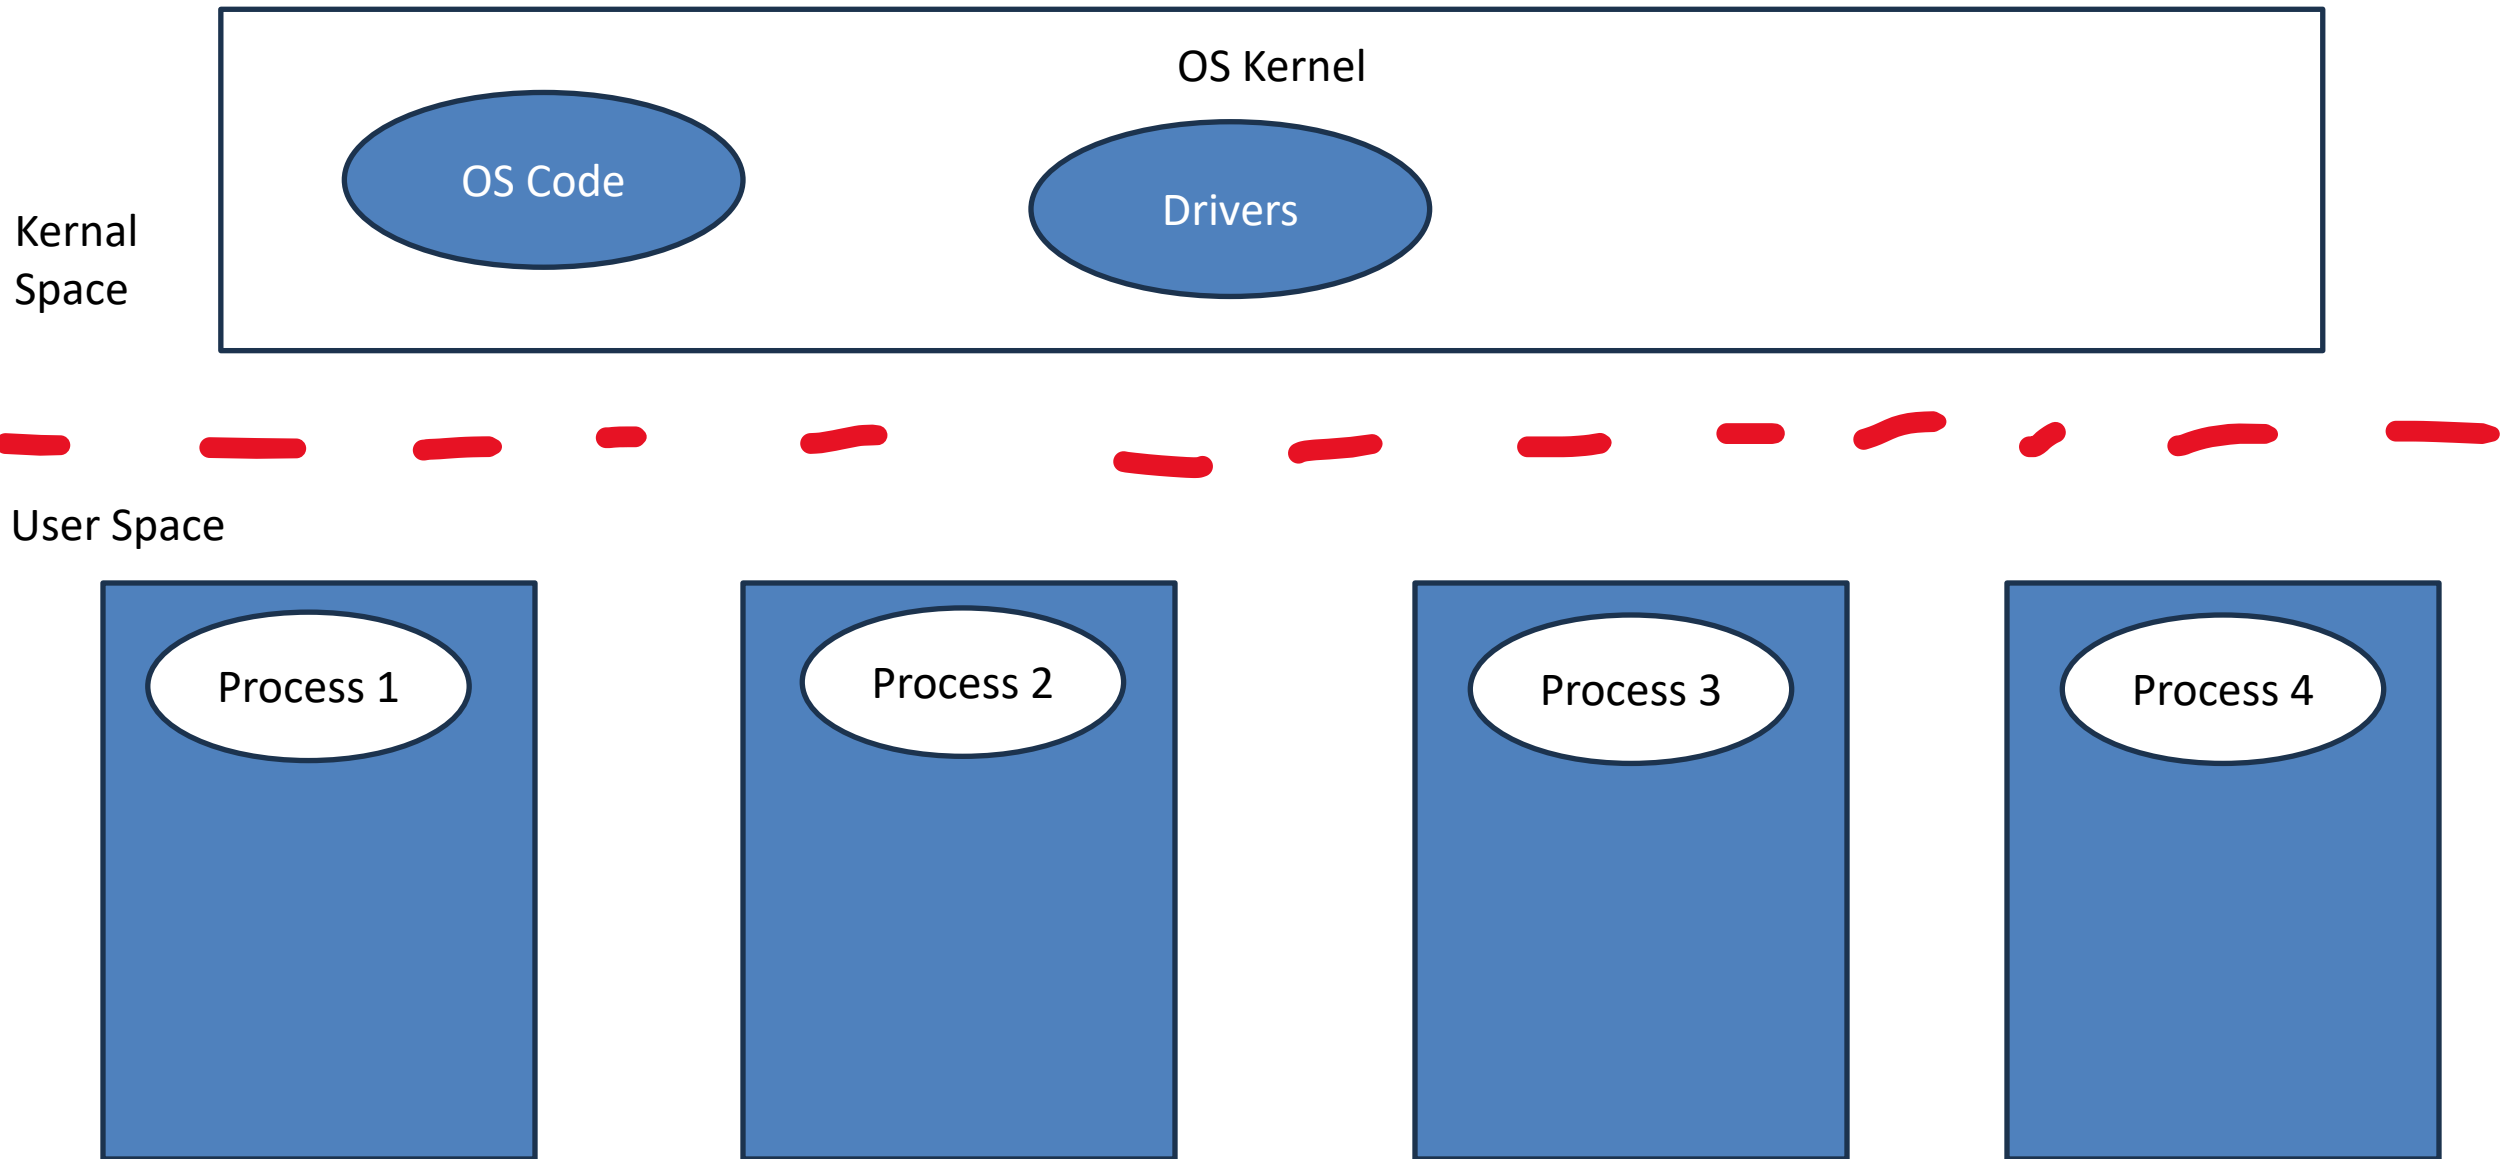
CPU



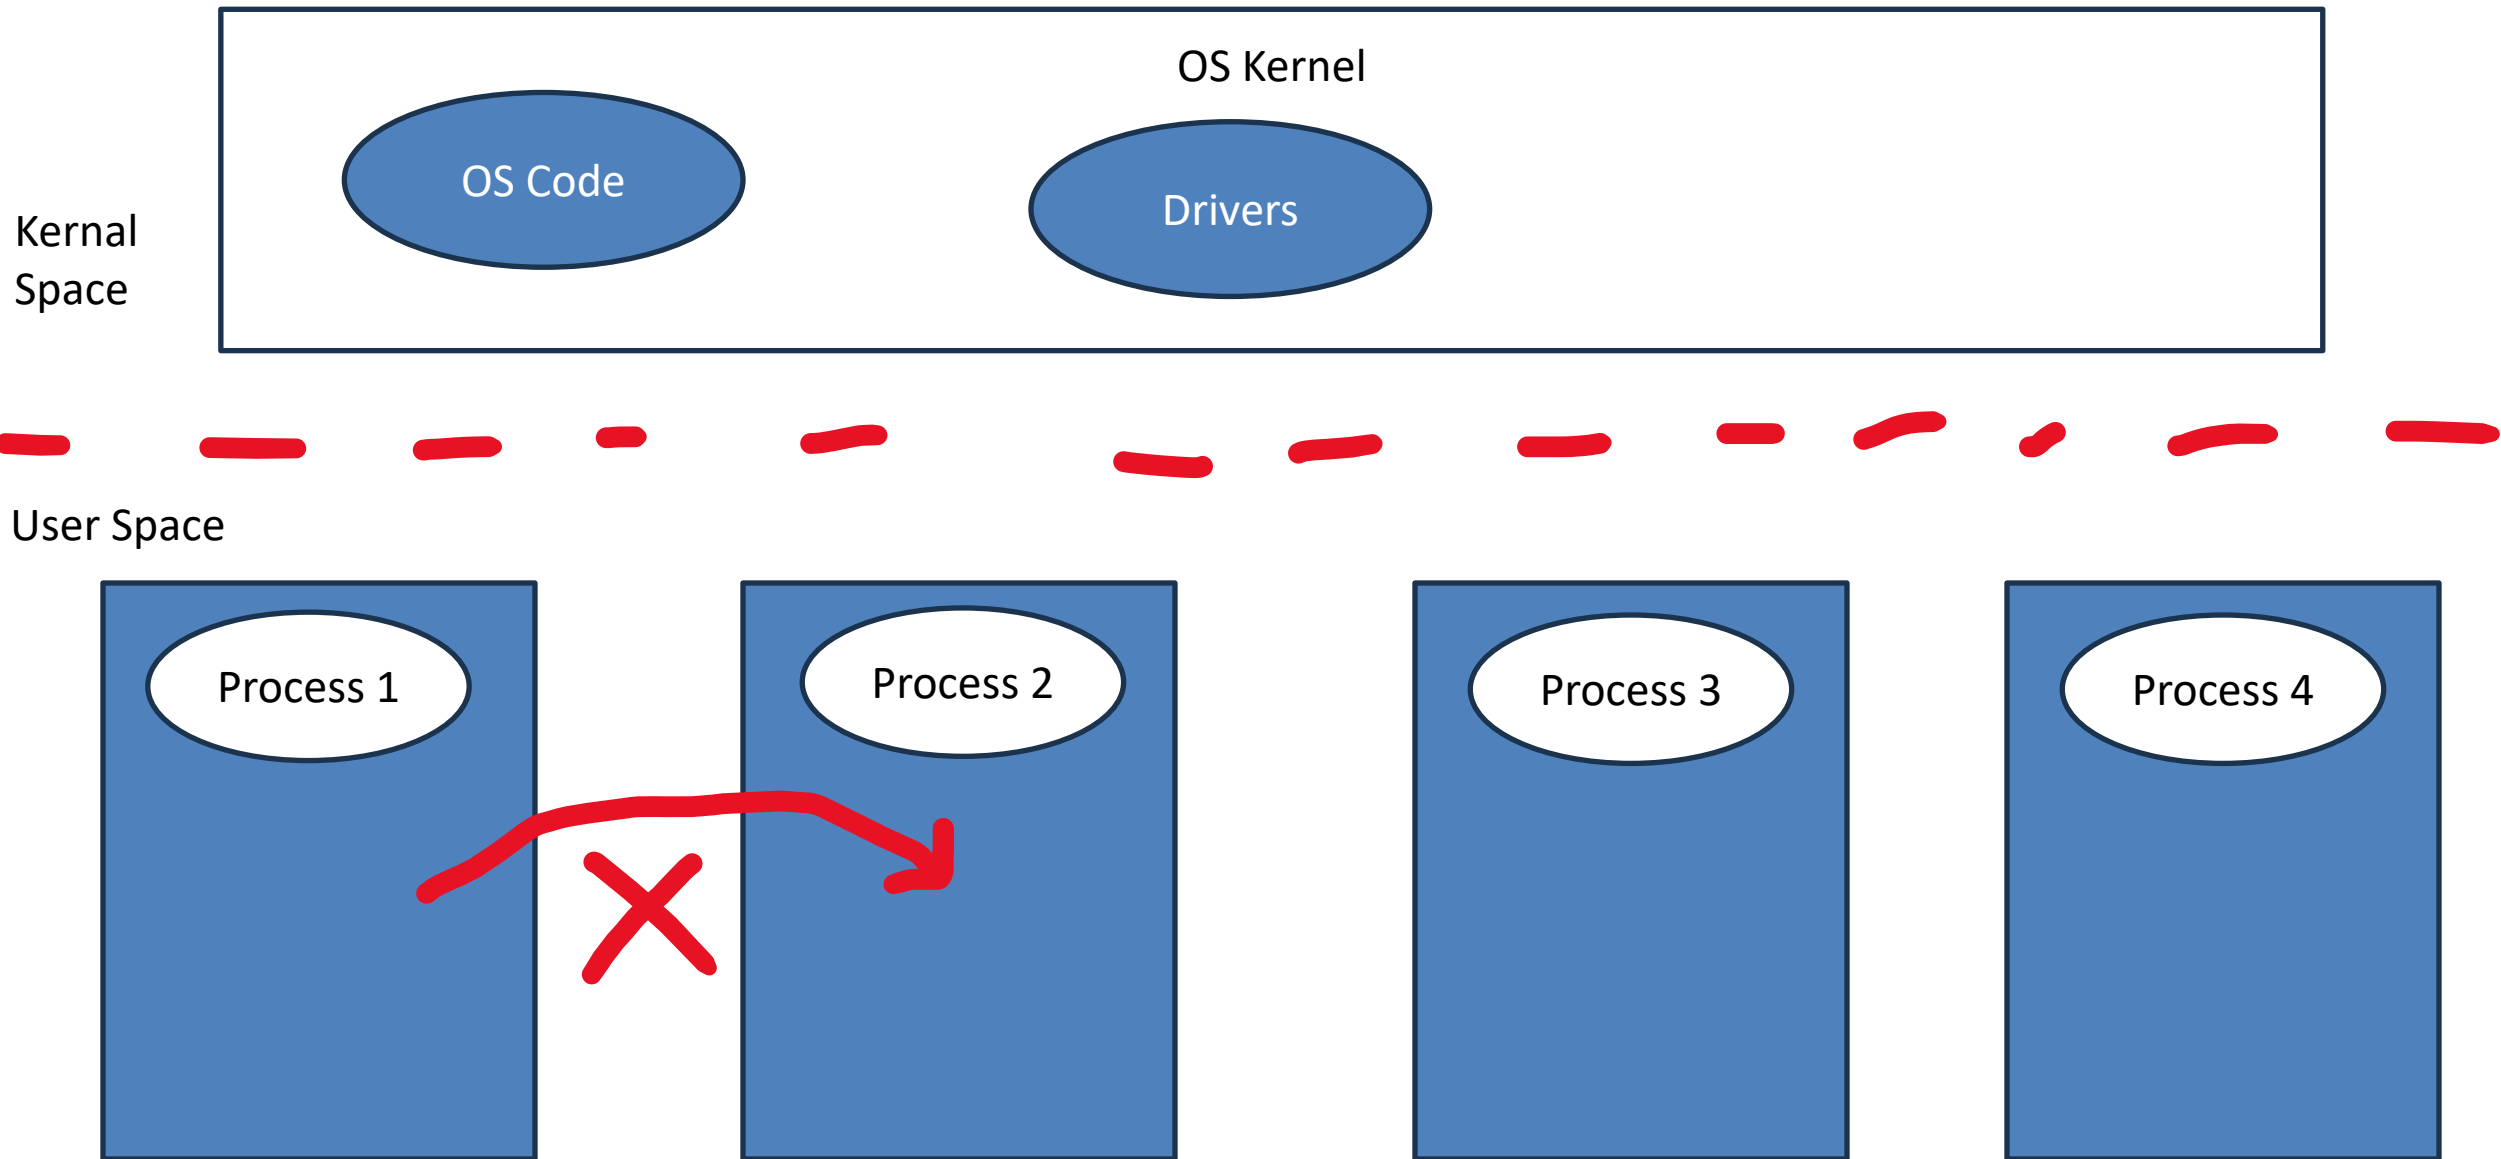
Illusion Manager

- Gives applications the illusion that they have infinite storage and resources

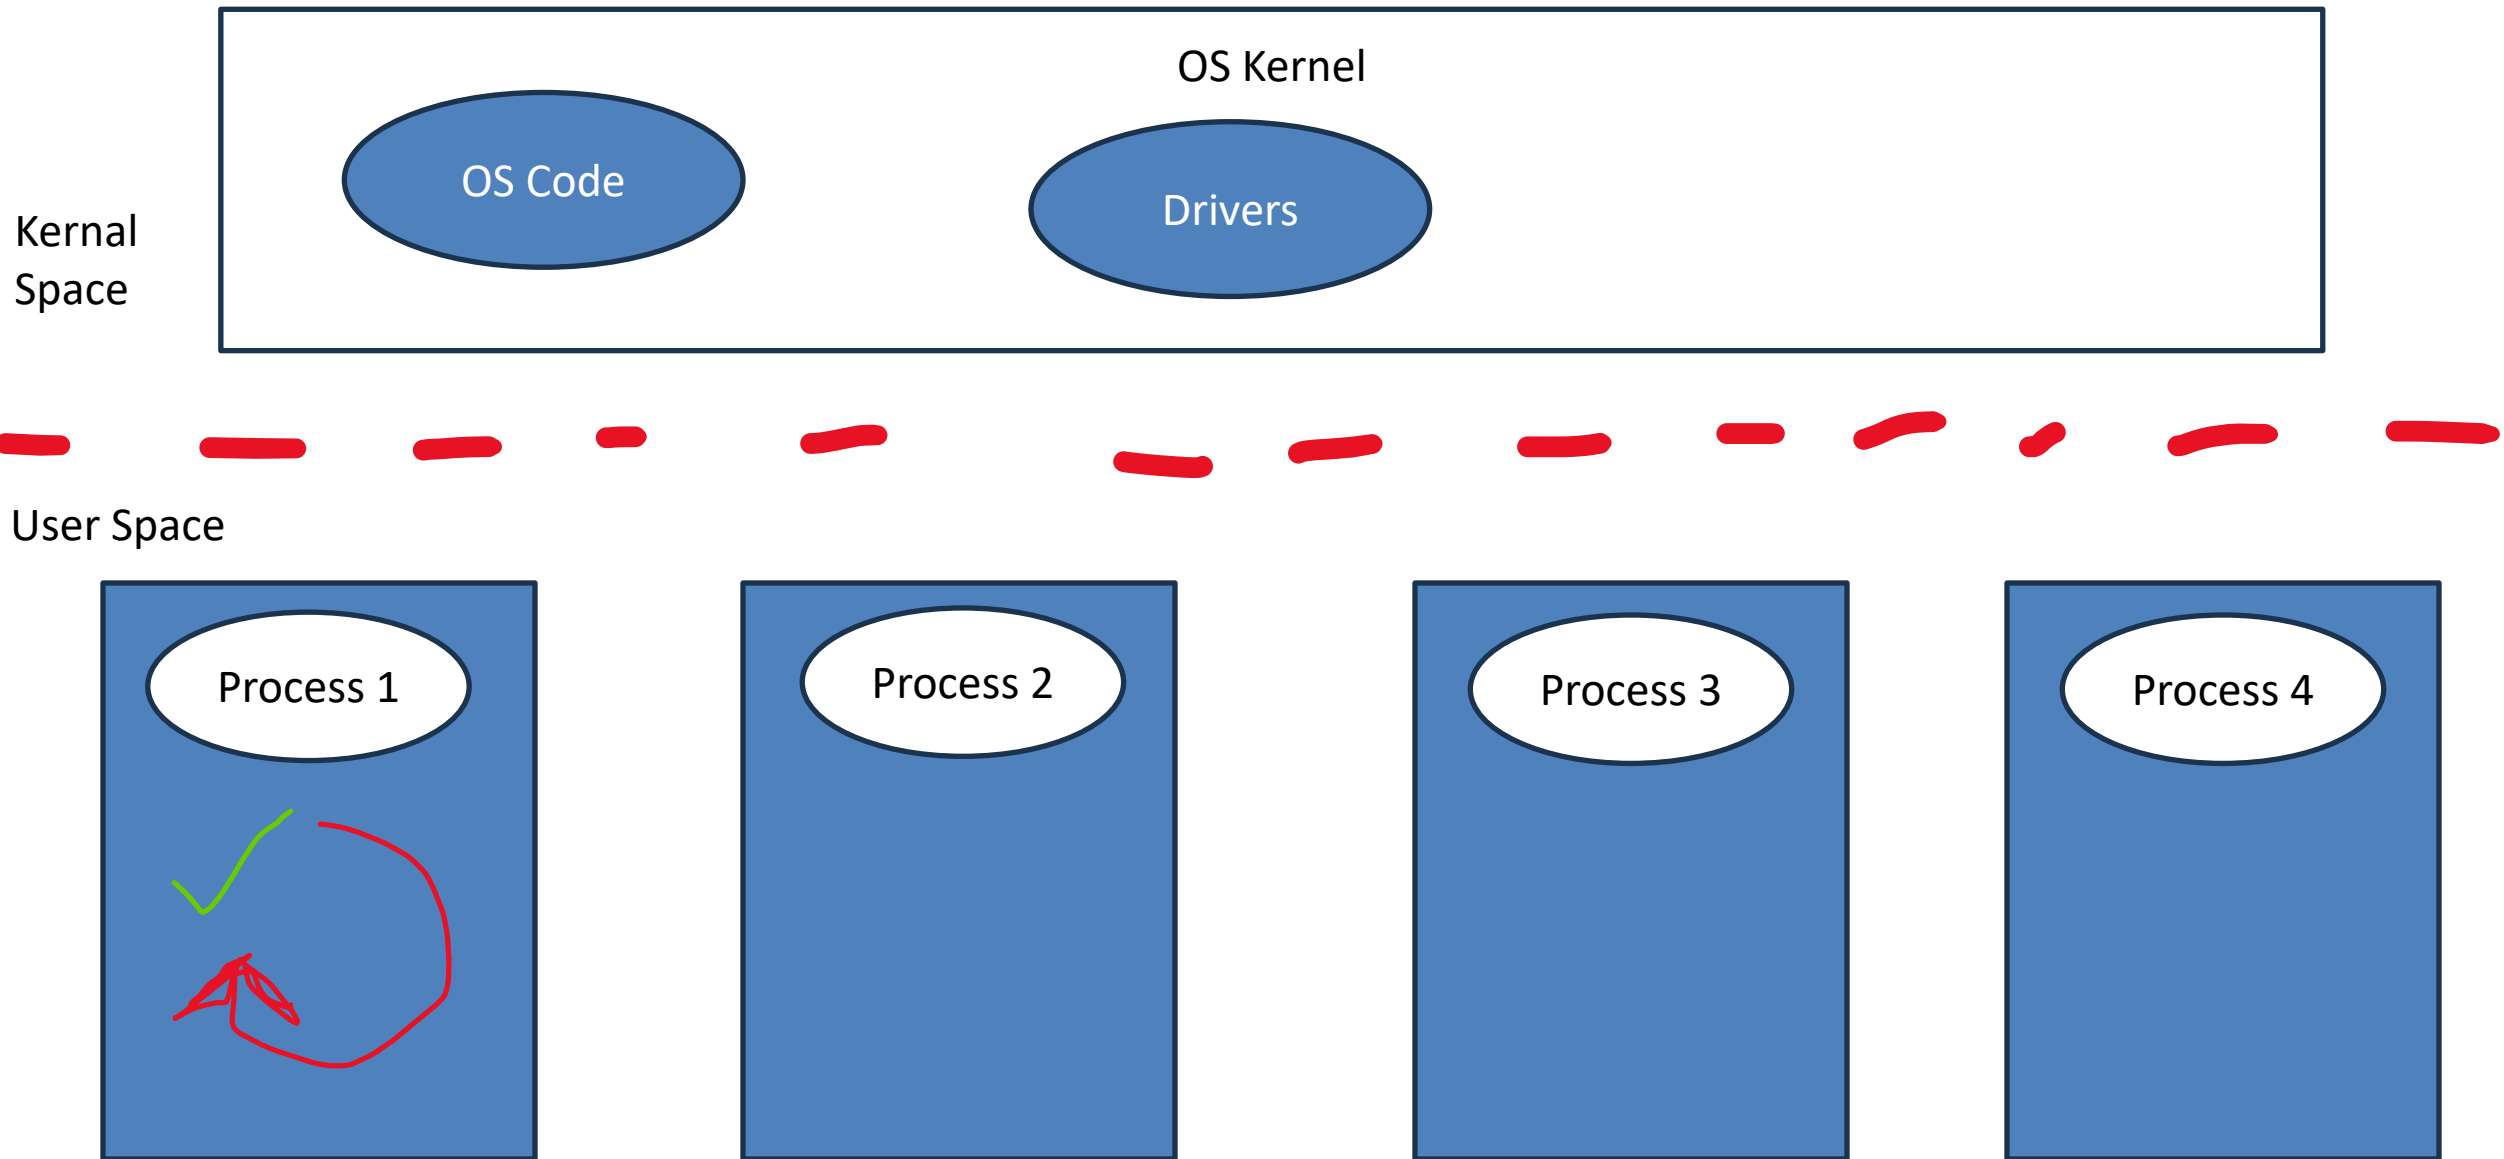




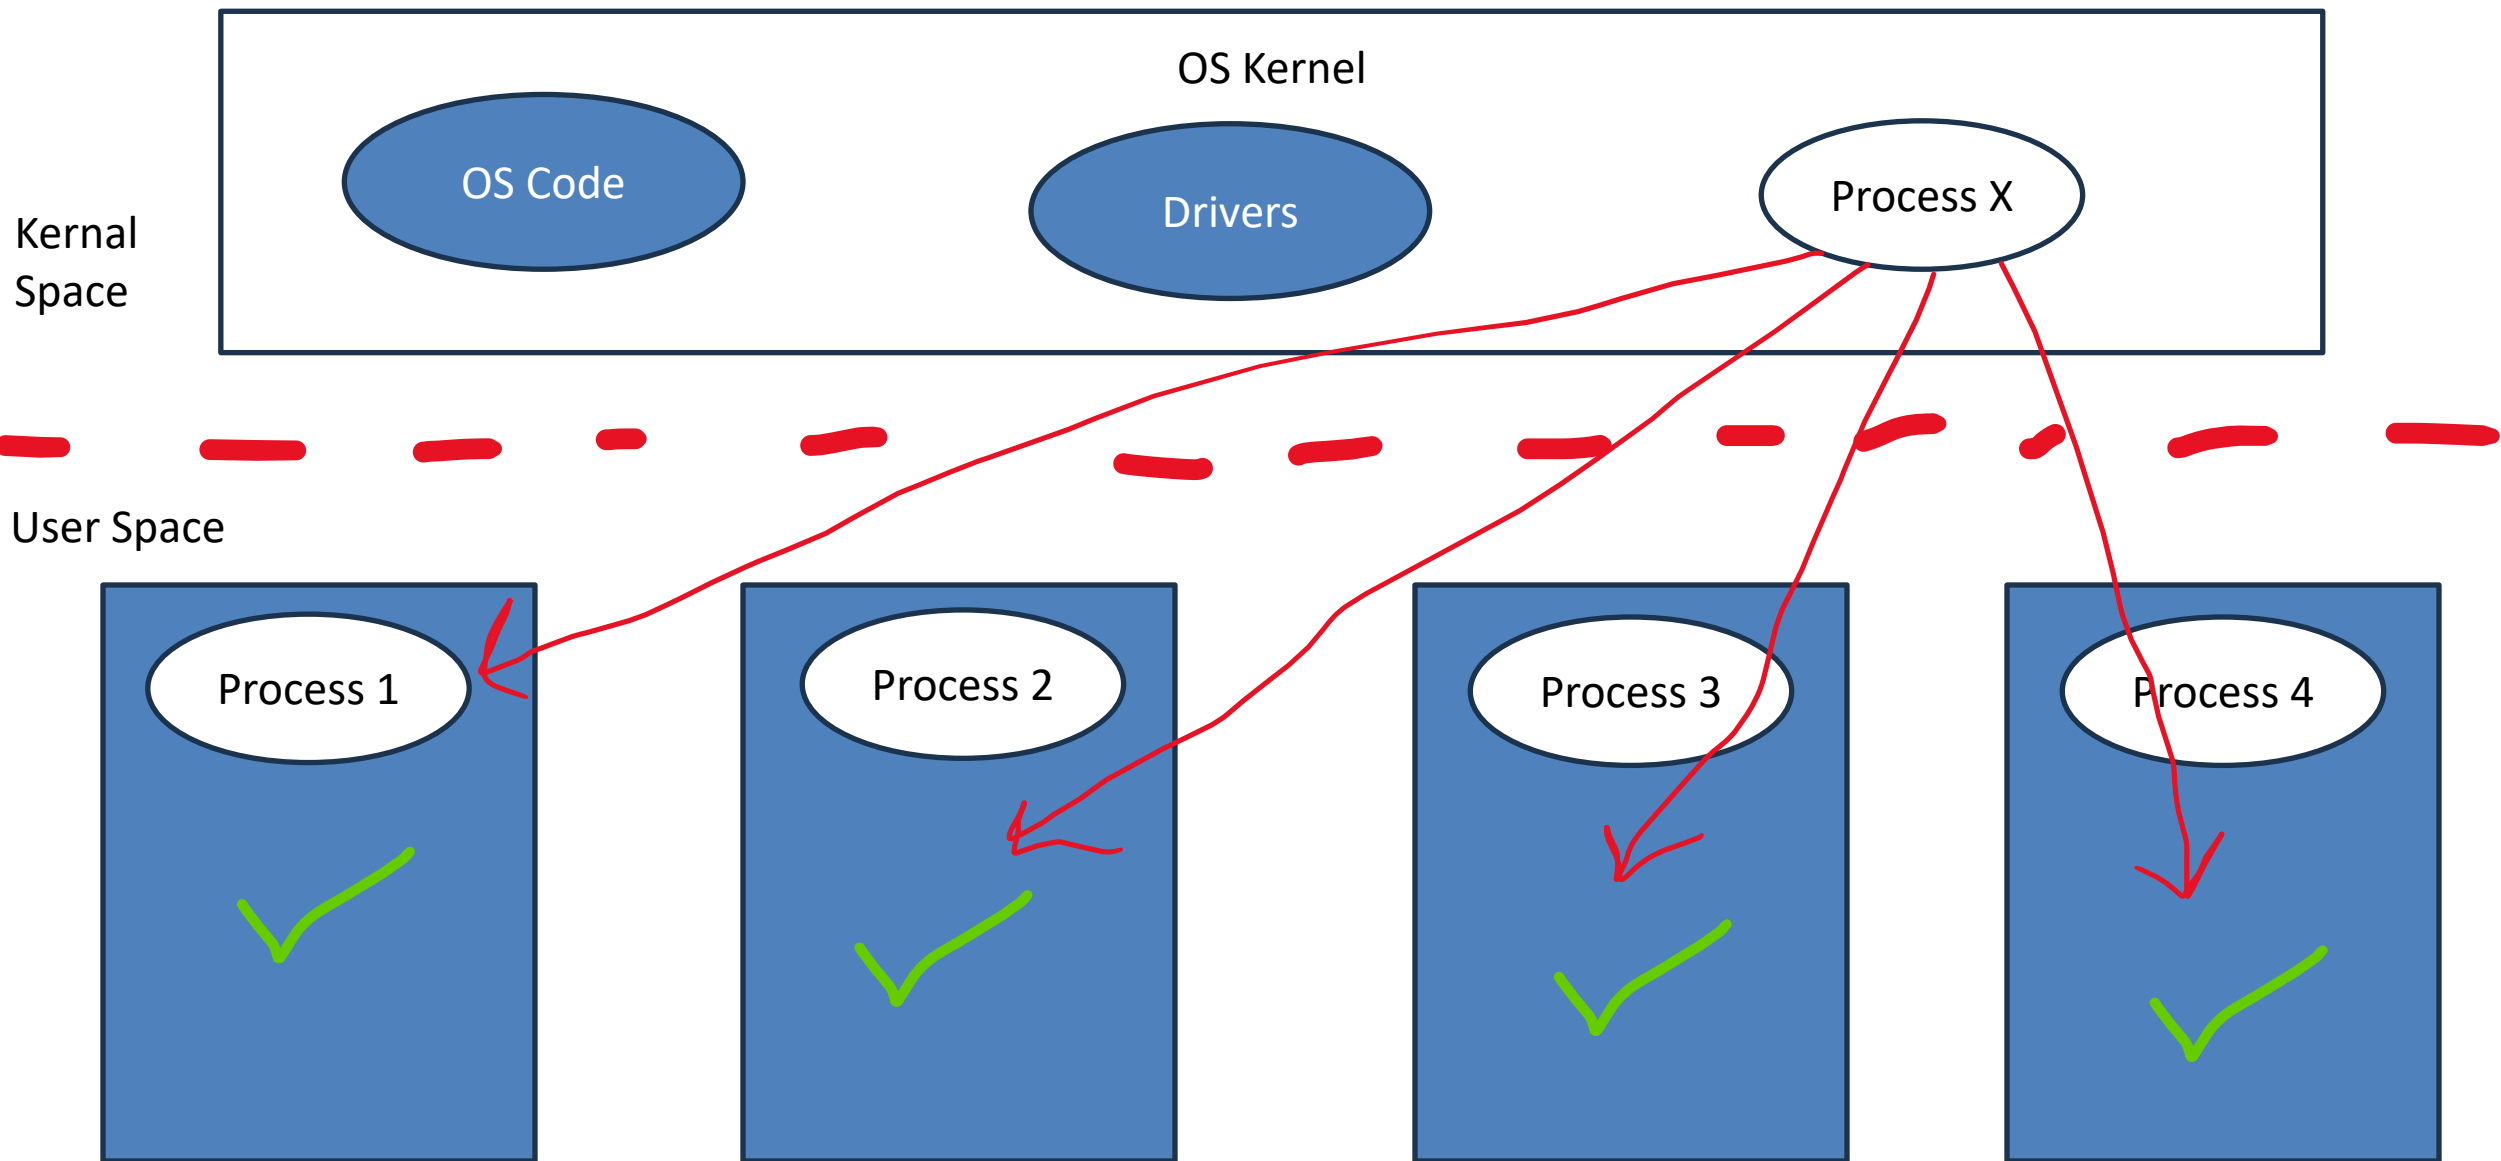
Processes in user space are **isolated**. This means that can normal processes not access other processes on the system*



Processes in user space are **isolated**. This means that can normal processes not access other processes on the system*



Processes in user space are **isolated**. This means that can normal processes not access other processes on the system*



A process running in **kernel mode** has access to **every** process (a big deal!!!)

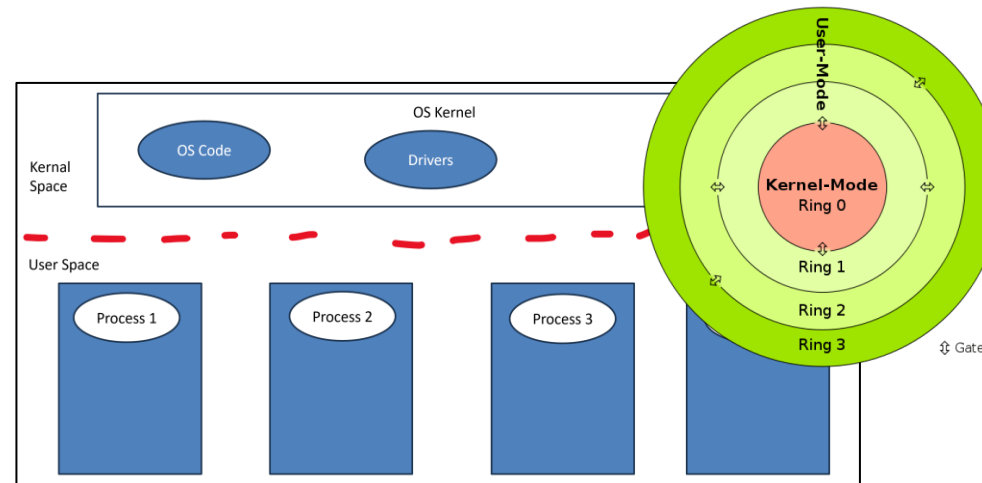
Case Study: Video Game Anticheat



The purpose of an anticheat program is to detect, prevent, and mitigate cheating in online games.

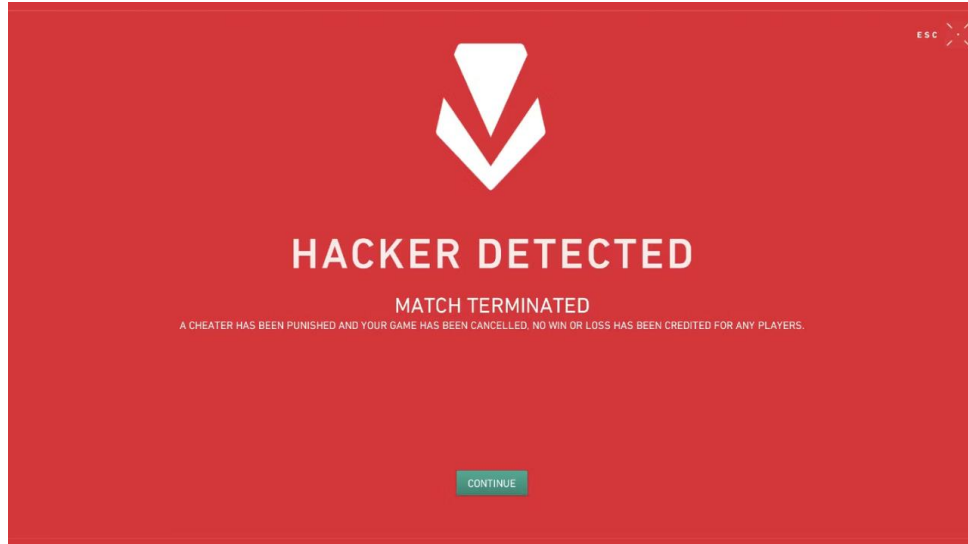
Active Cheating – A process currently running is giving a player unfair advantage while the game is running (aimbot)

Passive Cheating – A program that allows the player to change the game data, save data while the game is not running (hex editors)



Where should anticheat processes be running?

Case Study: Riot Vanguard



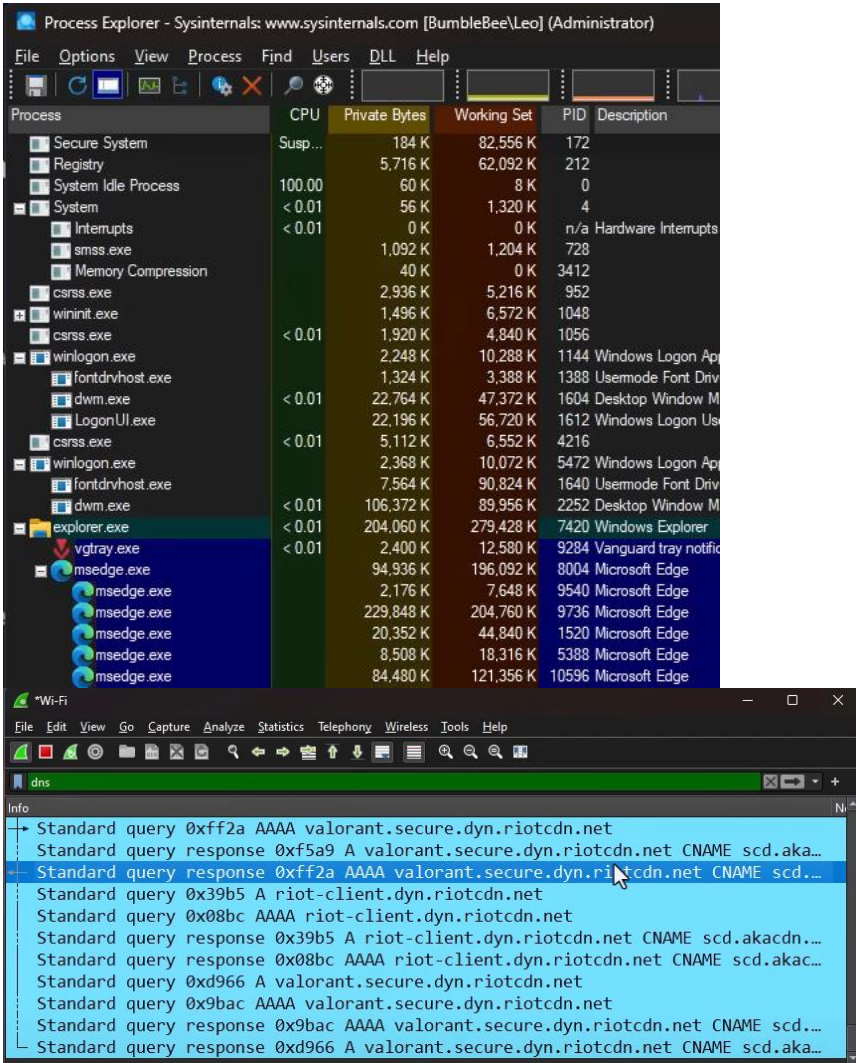
Riot Vanguard is the anticheat software for games made by Riot Games
→ League of Legends, Valorant

Riot Vanguard is a **kernel-level** anticheat

This means it has **very elevated** privileges, and can see almost everything on your system. There are different levels of kernel access, and Vanguard still requires **high** levels of access

Allegations have been made that Riot Vanguard is a suspicious program and that it is **spyware**

Case Study: Riot Vanguard



How to determine if something is malware?

- Static Analysis
- Dynamic Analysis

We can look at the network traffic being generated by Riot Vanguard, and we can look at the processes that it creates

Where it gets weird

Riot Vanguard is always running, even when you are not playing the game

If you stop the process, you have to restart your entire PC to play the game

Case Study: Riot Vanguard

Terms of Service

Last Modified: September 15, 2023

Greetings players,

SUS TOS

These terms of service (the "Terms") set out the terms and conditions by which Riot Games offers you access to use and enjoy our games, apps, websites and other services (the "Riot Services"). Riot Games is a global gaming company headquartered in Los Angeles with offices and operations around the world. When we say "Riot Games," we're referring to the Riot Games entity responsible for providing the Riot Services in your region (see Section 18, below) and these Terms are an agreement between you and that entity.

I agree to the Terms of Service, including the arbitration agreement and class action waiver in Section 17 to resolve any disputes. I have also read and acknowledge the Privacy Notice.

Riot Games is owned by [Tencent, a Chinese conglomerate](#). Tencent acquired a majority stake in Riot Games in 2011 and became the full owner in 2015. [↩](#)


Is it malware?

Kernel Level privilege is a very powerful (and scary) power to grant a process

The fact that it is constantly running, even when not playing the game, raises red flags

Riot Games is also owned by a Chinese company, which also raises many concerns

Many security experts condone the development of kernel-level anticheat



The jobs of an Operating System

1. Process Manager
“The Coach”

2. Interface Manager
“The Bouncer”

3. Memory Manager
“The Farmer”

4. Traffic Manager
“The Judge”

5. Illusion Manager
“The Illusionist”

