

ESOF 422:

Advanced Software Engineering: Cyber Practices

Secure by Design (Part 2)

Benefits of DDD, Immutability, Input Validation

Reese Pearsall
Spring 2025

Exam

If you *want* a third exam (finals week)

It would be optional

- If you don't take it, the average of your first two exams will be used

<https://etc.ch/FTqK>



Domain Expert

Doesn't know how to code

Expert in the field

Not concerned about business logic



Stakeholder

Doesn't know how to code

Might not know details of the field

Knows business logic



Programmer

Knows how to code

Might not know details of the field

Knows how to implement business logic

Domain-Driven Design: focus of modeling software to match a domain according to input from domain experts. Divide system into bounded contexts (domain primitives), each having their own model with strict constraints

```
public class Quantity {  
    private final int value;  
    public Quantity(final int value) throws Exception {  
        if(!inclusiveBetween(1,99)){  
            throw new Exception("Invalid Quantity");  
        }  
        this.value = value;  
    }  
}
```

Domain primitive enforce
domain rule validation at creation
time

Tightens our design by explicitly
stating requirements and
assumptions

Deeper modeling

Another important question:

Is the object **mutable** or **immutable** ?

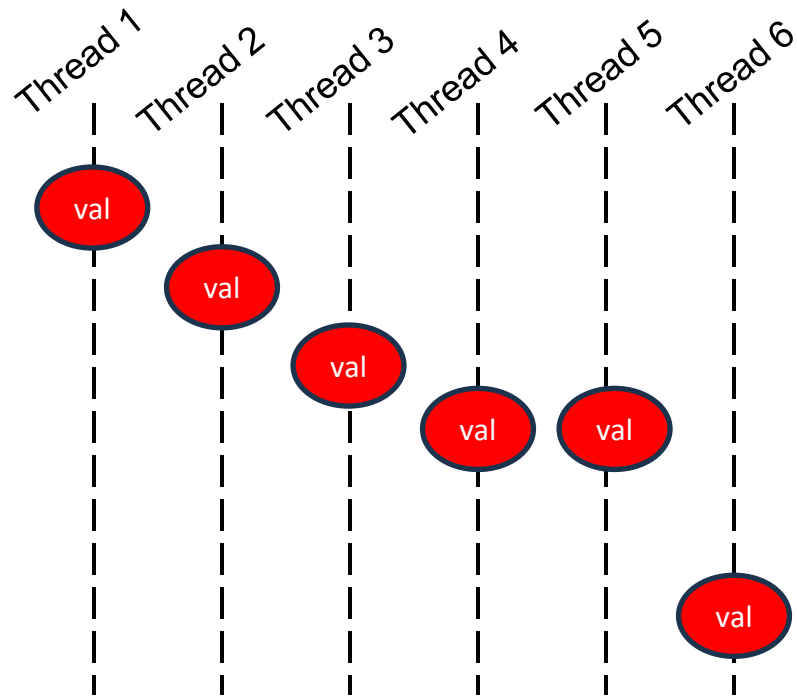
Mutable: allows an object to change (setters are used)

Immutable: object is not allowed to change

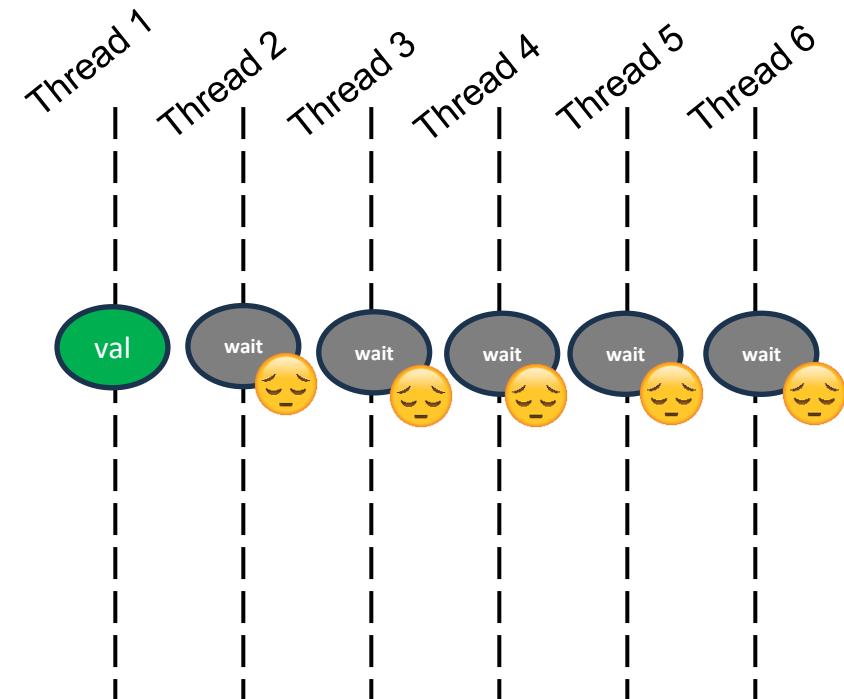
Immutability

Mutable: allows an object to change
Immutable: object is not allowed to change

Immutable objects are safe to share between threads



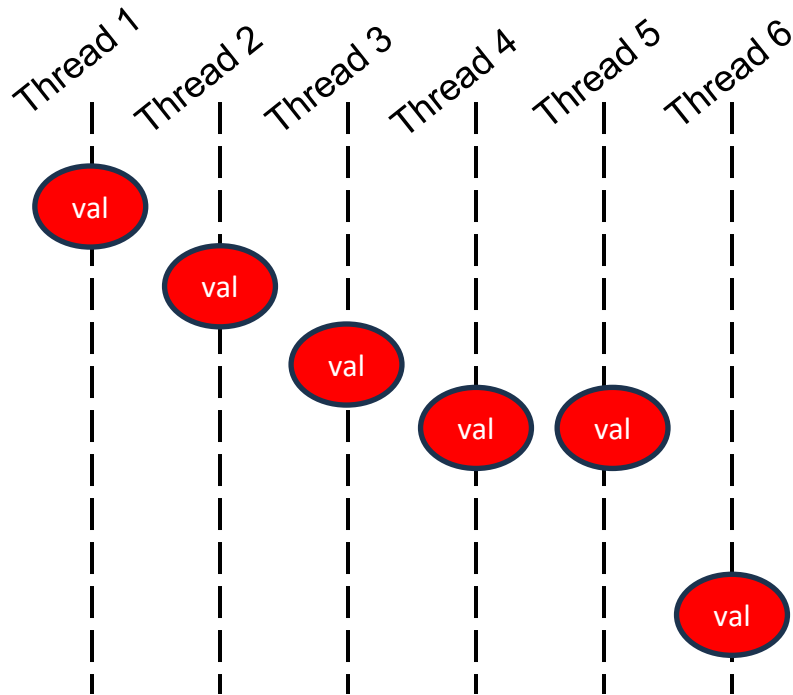
If an object is **mutable**, then *thread contention* is a problem (threads will have to wait until the previous thread is done with it)



Immutability

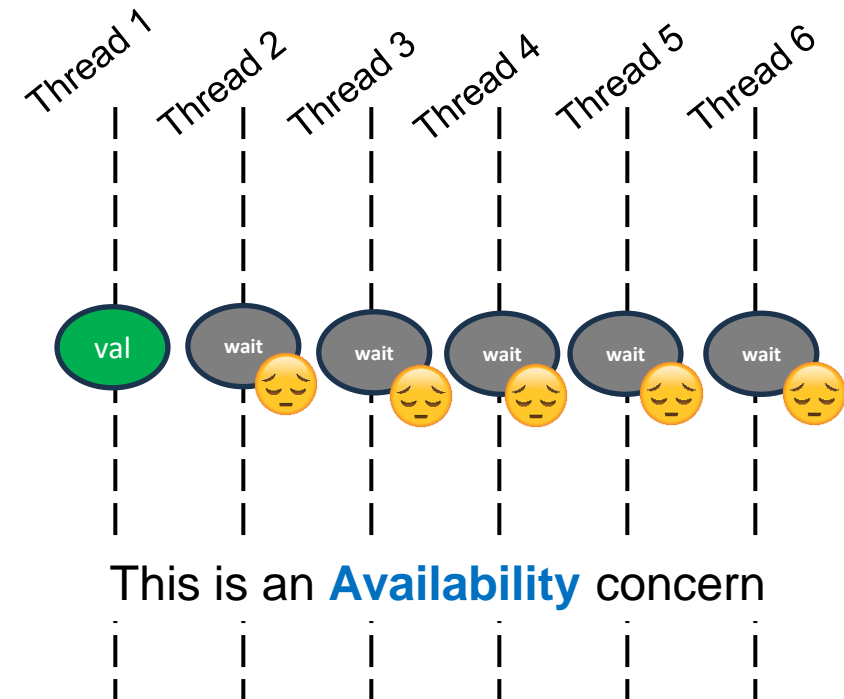


Immutable objects are safe to share between threads



Mutable: allows an object to change
Immutable: object is not allowed to change

If an object is **mutable**, then *thread contention* is a problem (threads will have to wait until the previous thread is done with it)



This is an **Availability** concern

```
public class Customer {  
    private static final int MIN_INVOICE_SCORE = 500;  
    private Id id;  
    private Name name;  
    private Order order;  
    private CreditScore creditScore;  
}
```

Customers with at least 500 credit score can pay by invoice (good thing). Customers with a credit score less than 500 must pay by credit card (bad thing?)

Credit score is based on payment history, and is a dynamic value

```
public class Customer {  
    private static final int MIN_INVOICE_SCORE = 500;  
    ...  
    private CreditScore creditScore;  
  
    public synchronized void setCreditScore(CreditScore creditScore) {  
        this.creditScore = creditScore;  
    }  
  
    public synchronized boolean isAcceptedForInvoicePayment() {  
        return creditScore.compute() > MIN_INVOICE_SCORE;  
    }  
}
```

This takes some time to compute...

Synchronized = only one thread is allowed to use method at a time

Many users = threads have to **wait**


```

public class Customer {
    private static final int MIN_INVOICE_SCORE = 500;
    ...
    private CreditScore creditScore;

    public synchronized void setCreditScore(CreditScore creditScore) {
        this.creditScore = creditScore;
    }

    public synchronized boolean isAcceptedForInvoicePayment() {
        return creditScore.compute() > MIN_INVOICE_SCORE;
    }
}

```

Problem	Category	Probable cause
Long waits and poor performance	Availability	System fails to access customer data in a reliable way and times out
Orders timing out at checkout	Availability	The system fails to retrieve necessary data to process the order in a timely fashion

```
public class Customer {  
    private static final int MIN_INVOICE_SCORE = 500;  
    ...  
    private CreditScore creditScore;  
  
    public synchronized CreditScore getCreditScore() {  
        return creditScore;  
    }  
  
    public synchronized void setCreditScore(CreditScore creditScore) {  
        this.creditScore = creditScore;  
    }  
  
    public synchronized boolean isAcceptedForInvoicePayment() {  
        return creditScore.compute() > MIN_INVOICE_SCORE;  
    }  
}
```

Let's look at another issue with mutable design

```
public class Customer {  
    private static final int MIN_INVOICE_SCORE = 500;  
    ...  
    private CreditScore creditScore;  
  
    public synchronized CreditScore getCreditScore() {  
        return creditScore;  
    }  
  
    public synchronized void setCreditScore(CreditScore creditScore) {  
        this.creditScore = creditScore;  
    }  
  
    public synchronized boolean isAcceptedForInvoicePayment() {  
        return creditScore.compute() > MIN_INVOICE_SCORE;  
    }  
}
```

We expect credit score to only be modified with the setter or the compute() method (synchronized)

```
public class Customer {  
    private static final int MIN_INVOICE_SCORE = 500;  
    ...  
    private CreditScore creditScore;  
  
    public synchronized CreditScore getCreditScore() {  
        return creditScore;  
    }  
  
    public synchronized void setCreditScore(CreditScore creditScore) {  
        this.creditScore = creditScore;  
    }  
  
    public synchronized boolean isAcceptedForInvoicePayment() {  
        return creditScore.compute() > MIN_INVOICE_SCORE;  
    }  
}
```

However, the `getCreditScore()` method returns a pointer to a mutable object!

```
public class Customer {  
    private static final int MIN_INVOICE_SCORE = 500;  
    ...  
    private CreditScore creditScore;  
  
    public synchronized CreditScore getCreditScore() {  
        return creditScore;  
    }  
  
    public synchronized void setCreditScore(CreditScore creditScore) {  
        this.creditScore = creditScore;  
    }  
  
    public synchronized boolean isAcceptedForInvoicePayment() {  
        return creditScore.compute() > MIN_INVOICE_SCORE;  
    }  
}
```

However, the `getCreditScore()` method returns a pointer to a mutable object!

This value can be modified outside of the class without requiring a lock! (scary)

```

public class Customer {
    private static final int MIN_INVOICE_SCORE = 500;
    ...
    private CreditScore creditScore;

    public synchronized CreditScore getCreditScore() {
        return creditScore;
    }
}

```

Lessons Learned:

The mutability of an object can impact the availability and integrity of your system

Problem	Category	Probable cause
Long waits and poor performance	Availability	System fails to access customer data in a reliable way and times out
Orders timing out at checkout	Availability	The system fails to retrieve necessary data to process the order in a timely fashion
Inconsistent payment options	Integrity	Credit score can be changed in an illegal way

```
public class Customer {  
    private static final int MIN_INVOICE_SCORE = 500;  
    ...  
    private final CreditScore creditScore;  
  
    public boolean isAcceptedForInvoicePayment() {  
        return creditScore.check();  
    }  
}
```

```
public class CreditScore {  
    private static final int MIN_INVOICE_SCORE = 500;  
    ...  
    private final int score;  
  
    public CreditScore(final int computedCreditScore){  
        //validation checks  
        this.score = computedCreditScore;  
    }  
    public boolean check() {  
        return score > MIN_INVOICE_SCORE;  
    }  
}
```

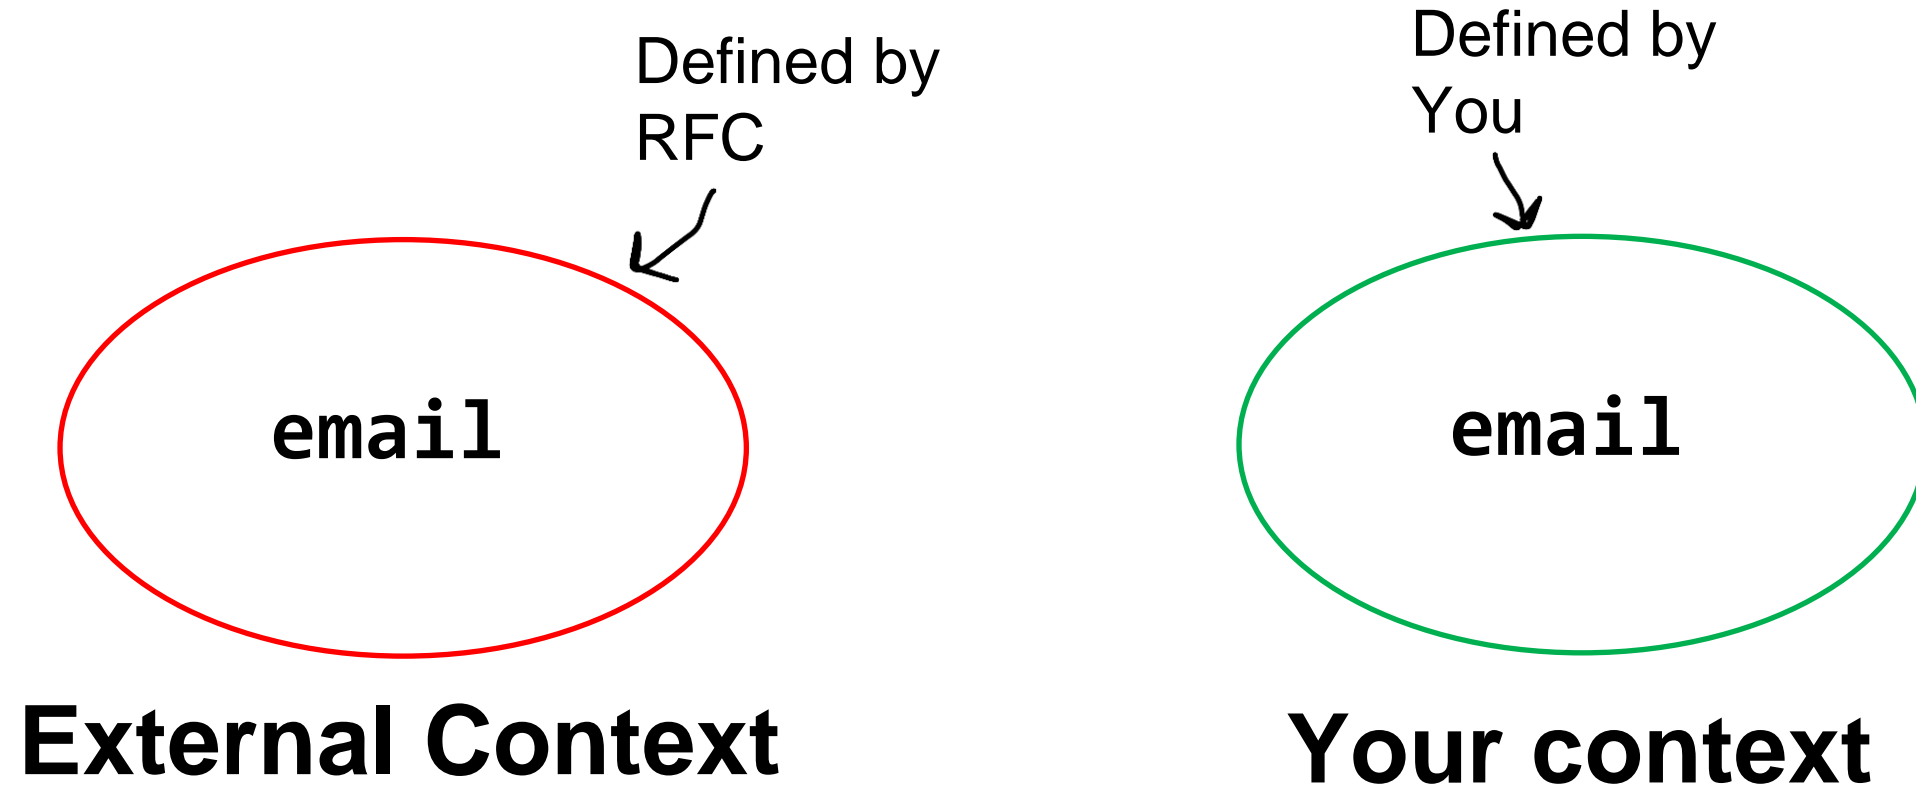
The **final** keyword in Java makes an object immutable

Choosing a design that favors **immutability**, the need for locks and protections against illegal change disappear

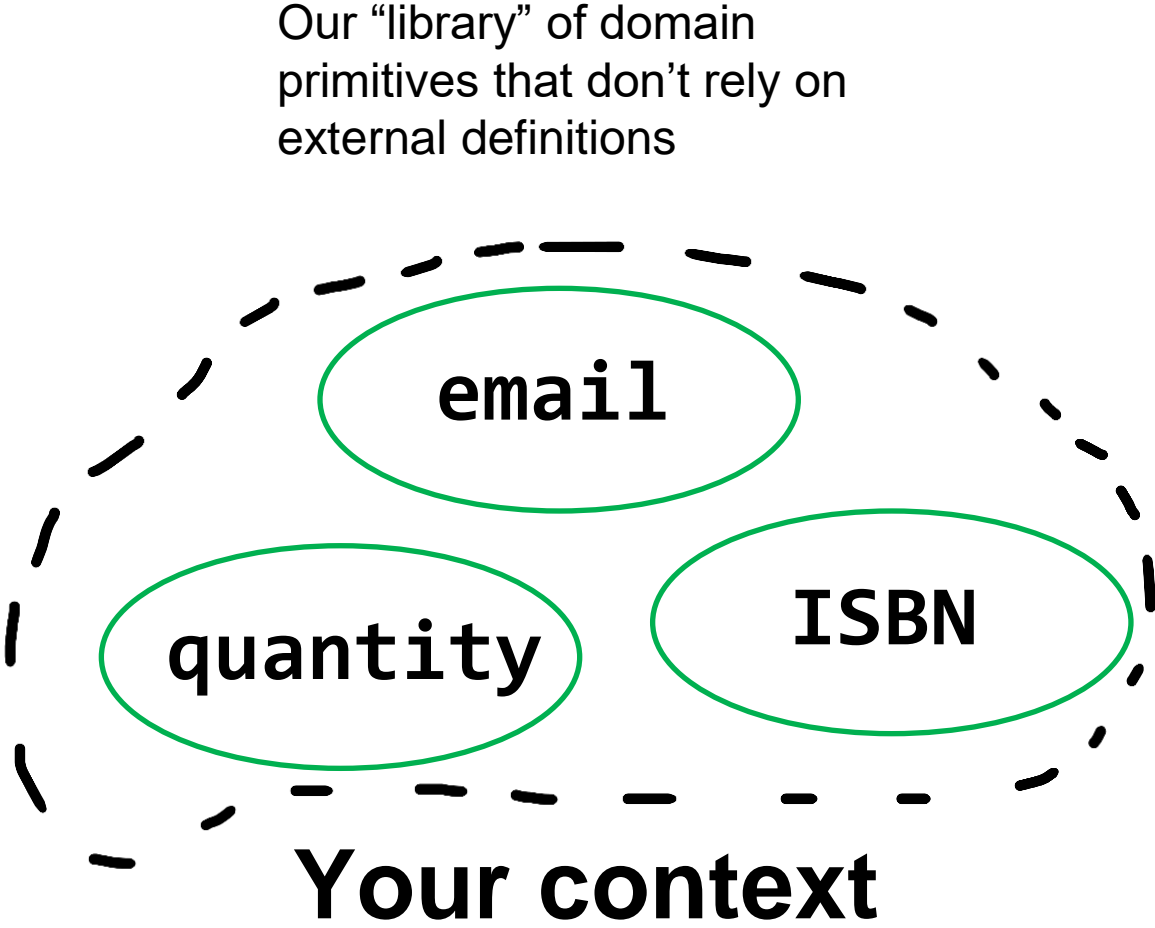

How to change Customer Data if its Immutable?

→ See “Entity Snapshot pattern”

External vs Internal Primitive



External vs Internal Primitive




Passing Primitives

```
public void sendAuditLogsToServerAt(java.net.InetAddress serverAddress) {  
  
}
```

By accepting a domain primitive (InternalAddress), it ensures the logs will be sent to a valid IP address

- Remember, the existence of a domain primitive means that it has to be valid!

Refactor



```
public void sendAuditLogsToServerAt(InternalAddress serverAddress) {  
    notNull(serverAddress)  
}
```

If you are building a public API, you should not reveal domain information

```

class Order {
    private BookRepository bookCatalog;
    private ArrayList<Object> items;
    private boolean paid = false;
    Inventory inventory;

    public void addItem(String isbn, int qty) {
        if (this.paid == false) {
            notNull(isbn);
            isTrue(isbn.length() == 10);
            isTrue(isbn.matches("[0-9X]+"));
            isTrue(isbn.matches("[0-9]{9}[0-9X]"));

            Book book = bookCatalog.findByISBN(isbn);
            if (inventory.availableBooks(isbn) >= qty) {
                items.add(new OrderLine(book, qty));
            }
        }
    }
    ...
}

```

```

class Order {
    private BookRepository bookCatalog;
    private ArrayList<Object> items;
    private boolean paid = false;
    Inventory inventory;

    public void addItem(String isbn, int qty) {
        if (this.paid == false) {
            notNull(isbn);
            assertTrue(isbn.length() == 10);
            assertTrue(isbn.matches("[0-9X]+"));
            assertTrue(isbn.matches("[0-9]{9}[0-9X]"));

            Book book = bookCatalog.findByISBN(isbn);
            if (inventory.availableBooks(isbn) >= qty) {
                items.add(new OrderLine(book, qty));
            }
        }
    }
    ...
}

```

This method does not treat ISBN and Quantity as domain primitives

It does validation checking on a few different things

It is missing checks for negative values!

```

class Order {
    private BookRepository bookCatalog;
    private ArrayList<Object> items;
    private boolean paid = false;
    Inventory inventory;

    public void addItem(String isbn, int qty) {
        if (this.paid == false) {
            notNull(isbn);
            assertTrue(isbn.length() == 10);
            assertTrue(isbn.matches("[0-9X]+"));
            assertTrue(isbn.matches("[0-9]{9}[0-9X]"));
            isNotNegative(qty)
            isLessThan(99)
            Book book = bookCatalog.findByISBN(isbn);
            if (inventory.availableBooks(isbn) >= qty) {
                items.add(new OrderLine(book, qty));
            }
        }
    }
    ...
}

```

If a method does validation checks on several different arguments, this method can become **cluttered**, which increases the chances of missing something

Tip: Leave domain validation to domain primitives

```

class Order {
    private BookRepository bookCatalog;
    private ArrayList<Object> items;
    private boolean paid = false;
    Inventory inventory;

    public void addItem(ISBN isbn, Quantity qty) {
        notNull(isbn);
        notNull(qty);
        if (this.paid == false) {
            Book book = bookCatalog.findByISBN(isbn);
            if (inventory.availableBooks(isbn) >= qty) {
                items.add(new OrderLine(book, qty));
            }
        }
    }
    ..
}

```

isbn and **qty** are
already valid before
entering method

Validation only needs
to happen in the
domain primitive class,
resulting in less
cluttering, and less
chance of missing
something

```

Class ISBN {
    private final String value;
    public ISBN(String isbn){
        //validation checks here
        value = isbn
    }
}

```

```

Class Quantity {
    private final int value;
    public ISBN(int q){
        //validation checks here
        value = q
    }
}

```

Input Checking

The **attack surface** of your system will usually always include areas of **untrusted user input**

The most severe vulnerabilities are usually due to lack of input checking or input validation

SQL Injections

Username or email address

`';DROP table Users where 0=0;--`

Password

[Forgot password?](#)

Sign in

XSS Attacks

Username or email address

`<script> //steal cookies </script>`

Password

[Forgot password?](#)

Sign in

Path Traversal

Username or email address

`“;cat ../../../../../../etc/shadow`

Password

[Forgot password?](#)

Sign in

Availability Attacks

Username or email address

`AAAAAAAAAAAAAAAAAAAA`

Password

[Forgot password?](#)

Sign in

Input Checking

Is this a “valid” input?

```
';DROP table Users where 0=0;--
```

99.9% of the time, this isn't valid

Bad

Unexpected inputs → Unexpected Behaviors → Potential Vulnerabilities

Good

Unexpected inputs → Input Validation → Only expected behaviors → More secure

Steps of input validation

1. **Origin** – Is the data from a legitimate sender?

- Check source IP address (Whitelists, Cloud configuration)
- Require use of API key or Access Token

Request URL

```
https://api.fantasydata.net/nfl/v2/JSON/DailyFantasyPlayers/2015-DEC-28
```

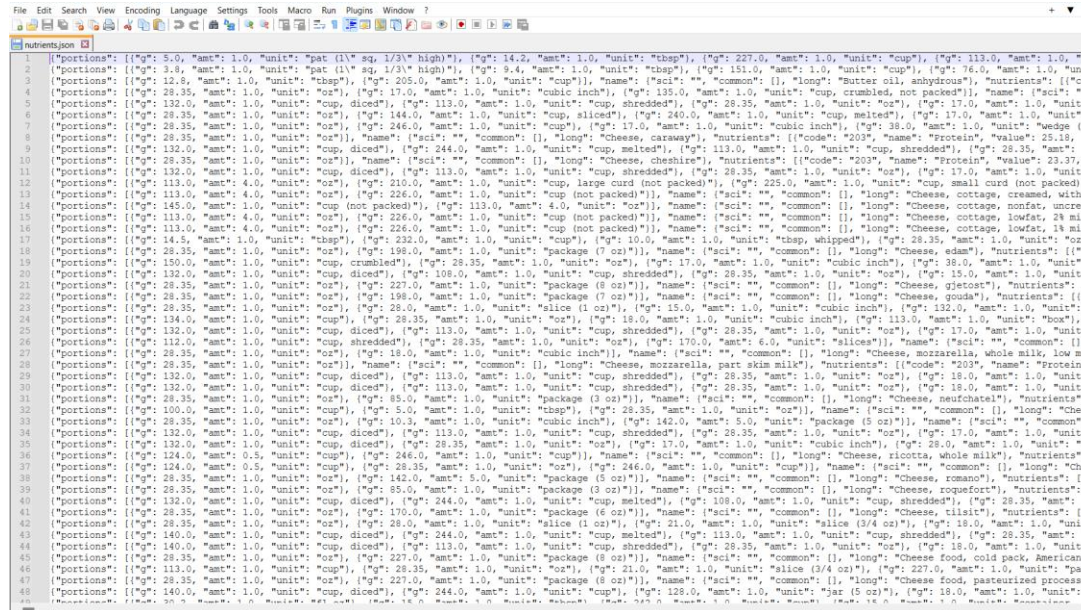
HTTP request

```
GET https://api.fantasydata.net/nfl/v2/JSON/DailyFantasyPlayers/2015-DEC-28 HTTP/1.1
Host: api.fantasydata.net
Ocp-Apim-Subscription-Key: .....
```




Steps of input validation

1. **Origin** – Is the data from a legitimate sender?
2. **Size** – Is the input reasonably large?
 - HTTP Content-Length header
 - Is 10 MB of JSON reasonable? What about 10 GB of JSON?



Everybody Gangsta until
the json File is 900GB



Steps of input validation

1. **Origin** – Is the data from a legitimate sender?
2. **Size** – Is the input reasonably large?
3. **Lexical Content** – Does it contain the right characters and encoding ?

; , < > -- Have special meanings in programming languages, which often makes them part of an attack. Do those characters make sense in an email address input box?

```
isTrue(isbn.matches("[0-9x]*"));
```

Filter out or encode special characters

<script>  <script>

Steps of input validation

1. **Origin** – Is the data from a legitimate sender?
2. **Size** – Is the input reasonably large?
3. **Lexical Content** – Does it contain the right characters and encoding ?
4. **Syntax**– Is the format right?

XML – Do all opening tags have an ending tag?

JSON – Are key value pairs correctly defined and follow correct JSON syntax ?

HTTP – Does it contain the HTTP method? Does it have the necessary headers?

Steps of input validation

1. **Origin** – Is the data from a legitimate sender?
2. **Size** – Is the input reasonably large?
3. **Lexical Content** – Does it contain the right characters and encoding ?
4. **Syntax**– Is the format right?
5. **Semantics**– Does it make sense?

Are the actual input values *valid* for the domain ?

If input is selecting an item to add to cart, does the item actually exist?


Taint Analysis

Tainted code analysis looks at the flow of potentially tainted input and flags potentially malicious before it is used

Username: input1

Password: input2

Email: input3



Data is used in an SQL query


```
SELECT * WHERE Username==input1 and Password==input2 and Email == input3
```

Taint Analysis

Tainted code analysis looks at the flow of potentially tainted input and flags potentially malicious before it is used

Username: input1 Password: input2 Email: input3

These are the **sources**
Areas where malicious code
may be introduced



Data is used in an SQL query


```
SELECT * WHERE Username==input1 and Password==input2 and Email == input3
```

Taint Analysis

Tainted code analysis looks at the flow of potentially tainted input and flags potentially malicious before it is used

Username: input1 Password: input2 Email: input3

These are the **sources**
Areas where malicious code may be introduced



Data is used in an SQL query

```
SELECT * WHERE Username==input1 and Password==input2 and Email == input3
```

These is a **sink**
Areas where input is used in application

Taint Analysis

Tainted code analysis looks at the flow of potentially tainted input and flags potentially malicious before it is used

Username: input1

Password: input2

Email: input3

These are the **sources**
Areas where malicious code may be introduced

Username input
validation

Password input
validation

Ideally, all user inputs should be **sanitized**

```
SELECT * WHERE Username==input1 and Password==input2 and Email == input3
```

These is a **sink**
Areas where input is used in application

Taint Analysis

Tainted code analysis looks at the flow of potentially tainted input and flags potentially malicious before it is used

Username: input1

Password: input2

Email: input3

Username input
validation

Password input
validation

These are the **sources**
Areas where malicious code
may be introduced

Ideally, all user inputs
should be **sanitized**

It's possible an input could
never be sanitized, which
creates security concern

```
SELECT * WHERE Username==input1 and Password==input2 and Email == input3
```

These is a **sink**
Areas where input is
used in application

Taint Analysis

Tainted code analysis looks at the flow of potentially tainted input and flags potentially malicious before it is used

Taint analysis keeps track of the flow of user input data, and makes sure all sources are sanitized

